



Core OS



About Me  
CTO/CO-FOUNDER  
systems engineer


@brandonphilips  
[github.com/philips](https://github.com/philips)

*What is CoreOS?*

# *What is CoreOS?*







# *Why build CoreOS?*

**The Datacenter as a Computer**  
*An Introduction to the Design of  
Warehouse-Scale Machines*

Luiz André Barroso and Urs Hölzle  
Google Inc.

*minimal*

reduce API contracts



kernel  
systemd  
etcd  
ssh  
docker

python  
java  
nginx  
mysql  
openssl

o distro distro distro distro distro distro

app

kernel  
systemd  
etc  
ssh  
docker

o distro distro distro distro distro distro

python  
java  
nginx  
mysql  
openssl

app

kernel  
systemd  
etc  
ssh  
docker

o distro distro distro distro distro distro

python  
openssl-A

app1

java  
openssl-B

app2

java  
openssl-B

app3

# CoreOS

o distro distro distro distro distro distro

python  
openssl-A

app1

java  
openssl-B

app2

java  
openssl-B

app3

# CoreOS

o distro distro distro distro distro distro

container

container

container

manual  
updates



# automatic updates



# automatic updates







*auto updates*

atomic with rollback



Update




**Data**


**A**

**B**







Downloading



Downloaded



Installed



Instance Hold



Complete



CoreOS

Applications

Packages

My Account ▾

+ Create a New Group

# CoreOS

## Production

f698ffbe-7488-43cf-97a1-66cf179609f



CHANNEL

 **Stable** ▸

INSTANCES

**69 instances**

AUTOMATIC UPDATES


**Enabled** ▸




UPDATE BEHAVIOR

**10 updates allowed per 15 minutes** ▸

## Version Breakdown

[View All Graphs](#)



 349.0.0  353.0.0  343.0.0  353.0.0  
 338.0.0

## Channels

[View Packages](#)



Alpha

361.0.0



Beta

357.0.0



Stable

353.0.0



[Add New Channel](#)

EC2 Autoscale

*containers*

run and isolate apps

*pid ns*

isolated pid 1

*user ns*

isolated uid 0

*network ns*

isolated netdev



*mount ns*

isolated /

*cgroups*

manage resources

*cgroups*


count resources

cAdvisor



*cgroups*

limit resources




docker engine

 Commit



 Github




 CI System



 Registry



Docker Pull



*configuration*

machines and clusters



*machine conf*

get into the cluster

```
#cloud-config
```

```
users:
```

- name: core  
    coreos-ssh-import-github: philips

```
coreos:
```

```
  etcd:
```

```
    discovery: https://discovery.etcd.io/ba09c
```

```
  units:
```

- name: etcd.service  
    command: start
- name: fleet.service  
    command: start

#cloud-config

users:

- name: core  
coreos-ssh-import-github: philips

coreos:

etcd:

discovery: <https://discovery.etcd.io/ba09c>

units:

- name: etcd.service  
command: start
- name: fleet.service  
command: start

```
#cloud-config
```

```
users:
```

- name: core  
    coreos-ssh-import-github: philips

```
coreos:
```

```
  etcd:
```

```
    discovery: https://discovery.etcd.io/ba09c
```

```
units:
```

- name: etcd.service  
    command: start
- name: fleet.service  
    command: start

```
#cloud-config
```

```
users:
```

- name: core  
coreos-ssh-import-github: philips

```
coreos:
```

```
etcd:
```

```
discovery: https://discovery.etcd.io/ba09c
```

```
units:
```

- name: etcd.service  
command: start
- name: fleet.service  
command: start

*cluster conf*

what is running

*services*

skydns, discoverd, confd

*cluster conf*

what should run



*coordination*

locksmith

*scheduling*

fleet, kubernetes

# etcd

/etc

distributed

open source software

highly available and reliable

sequentially consistent

watchable

exposed via HTTP

runtime reconfigurable

**-X GET**

Get Wait


**-X PUT**

Put Create CAS


**-X DELETE**

Delete CAD

# Available




# Available







# Available



# Unavailable



# Available




Leader




Follower

# Available




Leader




Follower

# Temporarily Unavailable




Leader



Follower

# Available




# Sequential Consistency


Index Time



GET A @ index 2 -> 2




GET A @ index 2 -> blocking







GET A @ index 2 -> 2



etcd guarantees that a get at index X will always return the same result.

Avoid thinking in terms of real time because with network latency the result is always out-of-date.

# Watchable Changes

HTTP Long-poll



> GET asdf?waitIndex=4&wait=true HTTP/1.1

> Accept: \*/\*

>

< HTTP/1.1 200 OK

< Content-Type: application/json

< X-Etcd-Index: 3

< X-Raft-Index: 97

< X-Raft-Term: 0

<

BLOCK



> GET asdf?waitIndex=4&wait=true HTTP/1.1

> Accept: \*/\*

>

< HTTP/1.1 200 OK

< Content-Type: application/json

< X-Etcd-Index: 3

< X-Raft-Index: 97

< X-Raft-Term: 0

<

```
{"action": "set", "node": {"key": "/asdf", "value": "foobar", "modifiedIndex": 4, "createdIndex": 4}}
```



> GET asdf?waitIndex=4&wait=true HTTP/1.1

> Accept: \*/\*

>

< HTTP/1.1 200 OK

< Content-Type: application/json

< X-Etcd-Index: 4

< X-Raft-Index: 516

< X-Raft-Term: 0

<

```
{"action": "set", "node": {"key": "/asdf", "value": "foobar", "modifiedIndex": 4, "createdIndex": 4}}
```

*scheduling*

getting work to servers

You



Scheduler API



Scheduler



Machine(s)



```
$ cat foo.service  
[Service]  
ExecStart=/usr/bin/sleep 500
```

```
$ fleetctl start foo.service  
Job foo.service launched on  
e1cd2bcd.../172.17.8.101
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

```
while true {  
    todo = diff(desState, curState)  
    schedule(todo)  
}
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

```
while true {  
  todo = diff(desState, curState)  
  schedule(todo)  
}
```

Talk tomorrow!

fleet: systemd for clusters

2:30pm Mississippi

# *Why use CoreOS?*

**The Datacenter as a Computer**  
*An Introduction to the Design of  
Warehouse-Scale Machines*

Luiz André Barroso and Urs Hölzle  
Google Inc.

*is it ready?*

stable released



*the future*

lots of work todo



Core OS

DEMO?