
kubernetes-python-client Documentation

Release

Kubernetes

Nov 01, 2017

Contents

| | | |
|----------|------------------------------|------------|
| 1 | Readme | 3 |
| 2 | Installation | 7 |
| 3 | Usage | 9 |
| 4 | kubernetes | 11 |
| 4.1 | kubernetes package | 11 |
| 5 | Contributing | 639 |
| 6 | Indices and tables | 641 |
| | Python Module Index | 643 |

Contents:

CHAPTER 1

Readme

Kubernetes Python Client

[![Build Status](https://travis-ci.org/kubernetes-incubator/client-python.svg?branch=master)](https://travis-ci.org/kubernetes-incubator/client-python) [![PyPI version](https://badge.fury.io/py/kubernetes.svg)](https://badge.fury.io/py/kubernetes) [![codecov](https://codecov.io/gh/kubernetes-incubator/client-python/branch/master/graph/badge.svg)](https://codecov.io/gh/kubernetes-incubator/client-python "Non-generated packages only") [![pypi supported versions](https://img.shields.io/pypi/pyversions/kubernetes.svg)](https://pypi.python.org/pypi/kubernetes) [![Client Capabilities](https://img.shields.io/badge/Kubernetes%20client-Silver-blue.svg?style=flat&colorB=C0C0C0&colorA=306CE8)](http://bit.ly/kubernetes-client-capabilities-badge) [![Client Support Level](https://img.shields.io/badge/kubernetes%20client-beta-green.svg?style=flat&colorA=306CE8)](http://bit.ly/kubernetes-client-support-badge)

Python client for the [kubernetes](http://kubernetes.io/) API.

Installation

From source:

```
` git clone --recursive https://github.com/kubernetes-incubator/client-python.  
git cd client-python python setup.py install `
```

From [PyPi](https://pypi.python.org/pypi/kubernetes/) directly:

```
` pip install kubernetes `
```

Example

list all pods:

```
“python from kubernetes import client, config
```

```
# Configs can be set in Configuration class directly or using helper utility config.load_kube_config()
```

```
v1 = client.CoreV1Api() print("Listing pods with their IPs:") ret = v1.list_pod_for_all_namespaces(watch=False) for  
i in ret.items:
```

```
    print("%s\t%s\t%s" % (i.status.pod_ip, i.metadata.namespace, i.metadata.name))
```

““

watch on namespace object:

“python from kubernetes import client, config, watch

Configs can be set in Configuration class directly or using helper utility config.load_kube_config()

v1 = client.CoreV1Api() count = 10 w = watch.Watch() for event in w.stream(v1.list_namespace, _request_timeout=60):

print("Event: %s %s" % (event['type'], event['object'].metadata.name)) count -= 1 if not count:

w.stop()

print("Ended.") ““

More examples can be found in examples folder. To run examples, run this command:

`shell python -m examples.example1`

(replace example1 with the example base filename)

Documentation

All APIs and Models' documentation can be found at the [Generated client's README file](kubernetes/README.md)

Compatibility

client-python follows [semver](<http://semver.org/>), so until the major version of client-python gets increased, your code will continue to work with explicitly supported versions of Kubernetes clusters.

Compatibility matrix

| Kubernetes 1.4 | Kubernetes 1.5 | Kubernetes 1.6 | Kubernetes 1.7 | Kubernetes 1.8 |

|-----|-----|-----|-----|-----| | client-python 1.0 | + | ✓ | - | - |
| - | | client-python 2.0 | + | + | ✓ | - | - | | client-python 3.0 | + | + | + | ✓ | ✓ | | client-python HEAD | + | + | + | + | ✓ |

Key:

- ✓ Exactly the same features / API objects in both client-python and the Kubernetes version.
- + client-python has features or api objects that may not be present in the Kubernetes cluster, but everything they have in common will work.
- The Kubernetes cluster has features the client-python library can't use (additional API objects, etc).

See the [CHANGELOG](./CHANGELOG.md) for a detailed description of changes between client-python versions.

Client version | Canonical source for OpenAPI spec | Maintenance status |

|-----|-----|-----| | 1.0 Alpha/Beta | Kubernetes main repo, 1.5 branch | | 1.0.x | Kubernetes main repo, 1.5 branch | ✓ | | 2.0 Alpha/Beta | Kubernetes main repo, 1.6 branch | | 2.0.x | Kubernetes main repo, 1.6 branch | ✓ | | 3.0 Alpha/Beta | Kubernetes main repo, 1.7 branch | | 3.0 | Kubernetes main repo, 1.7 branch | ✓ | | 4.0 Alpha/Beta | Kubernetes main repo, 1.8 branch | |

Key:

- ✓ Changes in main Kubernetes repo are manually ([should be automated](<https://github.com/kubernetes-incubator/client-python/issues/177>)) published to client-python when they are available.
- No longer maintained; please upgrade.

Note: There would be no maintenance for alpha/beta releases except the latest one.

Community, Support, Discussion

You can reach the maintainers of this project at [SIG API Machinery](<https://github.com/kubernetes/community/tree/master/sig-api-machinery>). If you have any problem with the package or any suggestions, please file an [issue](<https://github.com/kubernetes-incubator/client-python/issues>).

Code of Conduct

Participation in the Kubernetes community is governed by the [CNCF Code of Conduct](<https://github.com/cncf/foundation/blob/master/code-of-conduct.md>).

Kubernetes Incubator

This is a [Kubernetes Incubator project](<https://github.com/kubernetes/community/blob/master/incubator.md>).

- [SIG: sig-api-machinery](<https://github.com/kubernetes/community/tree/master/sig-api-machinery>)

Troubleshooting

SSLError on macOS

If you get an `SSLError`, you likely need to update your version of python. The version that ships with macOS may not be supported.

Install the latest version of python with [brew](<https://brew.sh/>):

```
` brew install python `
```

Once installed, you can query the version of OpenSSL like so:

```
` python -c "import ssl; print ssl.OPENSSL_VERSION" `
```

You'll need a version with OpenSSL version 1.0.0 or later.

Hostname doesn't match

If you get an `ssl.CertificateError` complaining about hostname match, your installed packages does not meet version [requirements](requirements.txt). Specifically check `ipaddress` and `urllib3` package versions to make sure they met requirements in requirements.txt file.

Why Exec/Attach calls doesn't work Starting from 4.0 release, we do not support directly calling `exec` or `attach` calls. you should use `stream` module to call them. so instead of `resp = api.connect_get_namespaced_pod_exec(name, ...` you should call `resp = stream(api.connect_get_namespaced_pod_exec, name,` See more at [exec example](examples/exec.py).

CHAPTER 2

Installation

At the command line:

```
$ pip install kubernetes
```

Or, if you have virtualenvwrapper installed:

```
$ mkvirtualenv kubernetes  
$ pip install kubernetes
```


CHAPTER 3

Usage

To use kubernetes-python-client in a project:

```
import kubernetes
```


4.1 kubernetes package

4.1.1 Subpackages

kubernetes.client package

Subpackages

kubernetes.client.apis package

Submodules

kubernetes.client.apis.apis_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.apis_api.ApisApi` (*api_client=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_versions (***kwargs*)

get available API versions This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_versions(async=True)` >>> `result = thread.get()`

:param async bool :return: V1APIGroupList

If the method is called asynchronously, returns the request thread.

get_api_versions_with_http_info (**kwargs)

get available API versions This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_versions_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIGroupList

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.apps_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.apps_api.**AppsApi** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (**kwargs)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_group(async=True) >>> result = thread.get()

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (**kwargs)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_group_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.apps_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.apps_v1beta1_api.**AppsV1beta1Api** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_controller_revision (*namespace, body, **kwargs*)

create a ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_controller_revision(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ControllerRevision body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

create_namespaced_controller_revision_with_http_info (*namespace, body, **kwargs*)

create a ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_controller_revision_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ControllerRevision body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment (*namespace, body, **kwargs*)

create a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `AppsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_rollback (*name, namespace, body, **kwargs*)

create rollback of a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_rollback(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DeploymentRollback (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `AppsV1beta1DeploymentRollback body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1DeploymentRollback`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_rollback_with_http_info (*name, namespace, body, **kwargs*)

create rollback of a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_rollback_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DeploymentRollback (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `AppsV1beta1DeploymentRollback body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1DeploymentRollback`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_with_http_info (*namespace, body, **kwargs*)

create a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `AppsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

create_namespaced_stateful_set (*namespace, body, **kwargs*)

create a StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_stateful_set(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1StatefulSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

create_namespaced_stateful_set_with_http_info (*namespace, body, **kwargs*)

create a StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_stateful_set_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1StatefulSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_controller_revision (*namespace, **kwargs*)

delete collection of ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_controller_revision(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field

is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_controller_revision_with_http_info(*namespace*,
***kwargs*)

delete collection of ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_controller_revision_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_deployment (*namespace*, ***kwargs*)

delete collection of Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_deployment(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_deployment_with_http_info (*namespace*, ***kwargs*)

delete collection of Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_deployment_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is

true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_stateful_set (*namespace, **kwargs*)

delete collection of StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_stateful_set(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of

a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_stateful_set_with_http_info (*namespace,*
***kwargs*)

delete collection of StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_stateful_set_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param bool async: :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_controller_revision (*name, namespace, body, **kwargs*)

delete a ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.delete_namespaced_controller_revision(name, namespace, body, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the ControllerRevision (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_controller_revision_with_http_info (*name, namespace, body,*
***kwargs*)

delete a ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_controller_revision_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ControllerRevision (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_deployment (*name, namespace, body, **kwargs*)

delete a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_deployment(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_deployment_with_http_info (*name, namespace, body, **kwargs*)

delete a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_deployment_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_stateful_set (*name, namespace, body, **kwargs*)

delete a StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_stateful_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_stateful_set_with_http_info (*name, namespace, body, **kwargs*)

delete a StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_stateful_set_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`:

Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_controller_revision_for_all_namespaces (**kwargs)

list or watch objects of kind ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_controller_revision_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes

to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ControllerRevisionList

If the method is called asynchronously, returns the request thread.

list_controller_revision_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_controller_revision_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ControllerRevisionList

If the method is called asynchronously, returns the request thread.

list_deployment_for_all_namespaces (**kwargs)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_deployment_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the

list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: AppsV1beta1DeploymentList

If the method is called asynchronously, returns the request thread.

list_deployment_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_deployment_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str

`resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `AppsV1beta1DeploymentList`

If the method is called asynchronously, returns the request thread.

`list_namespaced_controller_revision` (*namespace*, ***kwargs*)

list or watch objects of kind `ControllerRevision` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_controller_revision(namespace, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `namespace`: object name and auth scope, such as for teams and projects (required) :param bool `pretty`: If 'true', then the output is pretty printed. :param str `_continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1beta1ControllerRevisionList`

If the method is called asynchronously, returns the request thread.

`list_namespaced_controller_revision_with_http_info` (*namespace*, ***kwargs*)

list or watch objects of kind `ControllerRevision` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_controller_revision_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `namespace`: object name and auth scope, such as for teams and projects

(required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ControllerRevisionList

If the method is called asynchronously, returns the request thread.

list_namespaced_deployment (namespace, **kwargs)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_deployment(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and

clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `AppsV1beta1DeploymentList`

If the method is called asynchronously, returns the request thread.

`list_namespaced_deployment_with_http_info` (*namespace*, ***kwargs*)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_deployment_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `namespace`: object name and auth scope, such as for teams and projects (required) :param bool `pretty`: If 'true', then the output is pretty printed. :param str `_continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when `watch` is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify

resourceVersion. :return: AppsV1beta1DeploymentList

If the method is called asynchronously, returns the request thread.

list_namespaced_stateful_set (*namespace*, ***kwargs*)

list or watch objects of kind StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_stateful_set(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StatefulSetList

If the method is called asynchronously, returns the request thread.

list_namespaced_stateful_set_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_stateful_set_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is

true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StatefulSetList

If the method is called asynchronously, returns the request thread.

list_stateful_set_for_all_namespaces (**kwargs)

list or watch objects of kind StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_stateful_set_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list

result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StatefulSetList

If the method is called asynchronously, returns the request thread.

list_stateful_set_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_stateful_set_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StatefulSetList

If the method is called asynchronously, returns the request thread.

patch_namespaced_controller_revision (name, namespace, body, **kwargs)

partially update the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_controller_revision(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ControllerRevision (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required)

:param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ControllerRevision

If the method is called asynchronously, returns the request thread.

patch_namespaced_controller_revision_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_controller_revision_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ControllerRevision (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ControllerRevision

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment (*name, namespace, body, **kwargs*)

partially update the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_scale_with_http_info (*name, namespace, body, **kwargs*)

partially update scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_status (*name, namespace, body, **kwargs*)

partially update status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set (*name, namespace, body, **kwargs*)

partially update the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set_scale_with_http_info (*name, namespace, body, **kwargs*)

partially update scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set_status (*name, namespace, body, **kwargs*)

partially update status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

patch_namespaced_stateful_set_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_stateful_set_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

read_namespaced_controller_revision (*name, namespace, **kwargs*)

read the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_controller_revision(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ControllerRevision (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

read_namespaced_controller_revision_with_http_info (*name, namespace, **kwargs*)

read the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_controller_revision_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ControllerRevision (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment (*name, namespace, **kwargs*)

read the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_scale (*name, namespace, **kwargs*)

read scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_scale(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_scale_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_status (*name, namespace, **kwargs*)

read status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_with_http_info (*name, namespace, **kwargs*)

read the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `AppsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set (*name, namespace, **kwargs*)

read the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set_scale (*name, namespace, **kwargs*)

read scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set_scale(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set_scale_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `AppsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set_status (*name, namespace, **kwargs*)

read status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set_status_with_http_info (*name*, *namespace*, ***kwargs*)

read status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

read_namespaced_stateful_set_with_http_info (*name*, *namespace*, ***kwargs*)

read the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_stateful_set_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StatefulSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1StatefulSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_controller_revision (*name*, *namespace*, *body*, ***kwargs*)

replace the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_controller_revision(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ControllerRevision (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ControllerRevision body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

replace_namespaced_controller_revision_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

replace the specified ControllerRevision This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_controller_revision_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ControllerRevision (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ControllerRevision body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ControllerRevision`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment (*name*, *namespace*, *body*, ***kwargs*)

replace the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `AppsV1beta1Deployment body`:

(required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_scale (*name, namespace, body, **kwargs*)

replace scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.replace_namespaced_deployment_scale(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_scale_with_http_info (*name, namespace, body, **kwargs*)

replace scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.replace_namespaced_deployment_scale_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_status (*name, namespace, body, **kwargs*)

replace status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.replace_namespaced_deployment_status(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Deployment body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.replace_namespaced_deployment_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Deployment body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.replace_namespaced_deployment_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Deployment body:

(required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

replace_namespaced_stateful_set (*name, namespace, body, **kwargs*)

replace the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_stateful_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the StatefulSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1StatefulSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1StatefulSet

If the method is called asynchronously, returns the request thread.

replace_namespaced_stateful_set_scale (*name, namespace, body, **kwargs*)

replace scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_stateful_set_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_namespaced_stateful_set_scale_with_http_info (*name, namespace, body, **kwargs*)

replace scale of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_stateful_set_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param AppsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: AppsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_namespaced_stateful_set_status (*name, namespace, body, **kwargs*)

replace status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_stateful_set_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the StatefulSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1StatefulSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1StatefulSet

If the method is called asynchronously, returns the request thread.

replace_namespaced_stateful_set_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified StatefulSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_stateful_set_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the StatefulSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1StatefulSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1StatefulSet

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_stateful_set_with_http_info (name, namespace, body,
                                                    **kwargs)
    replace the specified StatefulSet This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.replace_namespaced_stateful_set_with_http_info(name, namespace, body, async=True) >>> result =
    thread.get()

    :param async bool :param str name: name of the StatefulSet (required) :param str namespace: object name
    and auth scope, such as for teams and projects (required) :param V1beta1StatefulSet body: (required)
    :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1StatefulSet
```

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.authentication_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.apis.authentication_api.AuthenticationApi (api_client=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

```
get_api_group (**kwargs)
    get information of a group This method makes a synchronous HTTP request by default. To make an
    asynchronous HTTP request, please pass async=True >>> thread = api.get_api_group(async=True) >>>
    result = thread.get()

    :param async bool :return: V1APIGroup
```

If the method is called asynchronously, returns the request thread.

```
get_api_group_with_http_info (**kwargs)
    get information of a group This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_group_with_http_info(async=True) >>> result = thread.get()

    :param async bool :return: V1APIGroup
```

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.authentication_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api (api_client=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

```
create_token_review (body, **kwargs)
```

create a TokenReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_token_review(body, async=True)`
 >>> `result = thread.get()`

:param `async` bool :param `V1beta1TokenReview body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1TokenReview`

If the method is called asynchronously, returns the request thread.

```
create_token_review_with_http_info (body, **kwargs)
```

create a TokenReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_token_review_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1beta1TokenReview body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1TokenReview`

If the method is called asynchronously, returns the request thread.

```
get_api_resources (**kwargs)
```

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIResourceList`

If the method is called asynchronously, returns the request thread.

```
get_api_resources_with_http_info (**kwargs)
```

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIResourceList`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.authorization_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.apis.authorization_api.AuthorizationApi (api_client=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

```
get_api_group (**kwargs)
```

get information of a group This method makes a synchronous HTTP request by default. To make an

asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.authorization_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api` (*api_client=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_local_subject_access_review (*namespace, body, **kwargs*)

create a `LocalSubjectAccessReview` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_local_subject_access_review(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1LocalSubjectAccessReview body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1LocalSubjectAccessReview`

If the method is called asynchronously, returns the request thread.

create_namespaced_local_subject_access_review_with_http_info (*namespace, body, **kwargs*)

create a `LocalSubjectAccessReview` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_local_subject_access_review_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1LocalSubjectAccessReview body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1LocalSubjectAccessReview`

If the method is called asynchronously, returns the request thread.

create_self_subject_access_review (*body, **kwargs*)

create a `SelfSubjectAccessReview` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_self_subject_access_review(body, async=True)` >>> `result = thread.get()`

:param async bool :param V1beta1SelfSubjectAccessReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SelfSubjectAccessReview

If the method is called asynchronously, returns the request thread.

create_self_subject_access_review_with_http_info (*body*, ***kwargs*)

create a SelfSubjectAccessReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_self_subject_access_review_with_http_info(body, async=True) >>> result = thread.get()

:param async bool :param V1beta1SelfSubjectAccessReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SelfSubjectAccessReview

If the method is called asynchronously, returns the request thread.

create_self_subject_rules_review (*body*, ***kwargs*)

create a SelfSubjectRulesReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_self_subject_rules_review(body, async=True) >>> result = thread.get()

:param async bool :param V1beta1SelfSubjectRulesReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SelfSubjectRulesReview

If the method is called asynchronously, returns the request thread.

create_self_subject_rules_review_with_http_info (*body*, ***kwargs*)

create a SelfSubjectRulesReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_self_subject_rules_review_with_http_info(body, async=True) >>> result = thread.get()

:param async bool :param V1beta1SelfSubjectRulesReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SelfSubjectRulesReview

If the method is called asynchronously, returns the request thread.

create_subject_access_review (*body*, ***kwargs*)

create a SubjectAccessReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_subject_access_review(body, async=True) >>> result = thread.get()

:param async bool :param V1beta1SubjectAccessReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SubjectAccessReview

If the method is called asynchronously, returns the request thread.

create_subject_access_review_with_http_info (*body*, ***kwargs*)

create a SubjectAccessReview This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_subject_access_review_with_http_info(body, async=True) >>> result = thread.get()

:param async bool :param V1beta1SubjectAccessReview body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1SubjectAccessReview

If the method is called asynchronously, returns the request thread.

get_api_resources (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

```
get_api_resources_with_http_info (**kwargs)
    get available resources This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_resources_with_http_info(async=True) >>> result = thread.get()
```

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.autoscaling_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.autoscaling_api.**AutoscalingApi** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

```
get_api_group (**kwargs)
    get information of a group This method makes a synchronous HTTP request by default. To make an
    asynchronous HTTP request, please pass async=True >>> thread = api.get_api_group(async=True) >>>
    result = thread.get()
```

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

```
get_api_group_with_http_info (**kwargs)
    get information of a group This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_group_with_http_info(async=True) >>> result = thread.get()
```

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.autoscaling_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.autoscaling_v1_api.**AutoscalingV1Api** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_horizontal_pod_autoscaler (*namespace*, *body*, ***kwargs*)

create a HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_horizontal_pod_autoscaler(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1HorizontalPodAutoscaler body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

create_namespaced_horizontal_pod_autoscaler_with_http_info (*namespace*, *body*, ***kwargs*)

create a HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_horizontal_pod_autoscaler_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1HorizontalPodAutoscaler body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_horizontal_pod_autoscaler (*namespace*, ***kwargs*)

delete collection of HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_horizontal_pod_autoscaler(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version

of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_horizontal_pod_autoscaler_with_http_info(*namespace*,
***kwargs*)

delete collection of HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_horizontal_pod_autoscaler_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_horizontal_pod_autoscaler(*name*, *namespace*, *body*, ***kwargs*)

delete a HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_namespaced_horizontal_pod_autoscaler(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_horizontal_pod_autoscaler_with_http_info (*name, namespace, body, **kwargs*)

delete a HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_horizontal_pod_autoscaler_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_horizontal_pod_autoscaler_for_all_namespaces (***kwargs*)

list or watch objects of kind HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_horizontal_pod_autoscaler_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1HorizontalPodAutoscalerList

If the method is called asynchronously, returns the request thread.

list_horizontal_pod_autoscaler_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_horizontal_pod_autoscaler_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether

more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1HorizontalPodAutoscalerList

If the method is called asynchronously, returns the request thread.

list_namespaced_horizontal_pod_autoscaler (*namespace, **kwargs*)

list or watch objects of kind HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_horizontal_pod_autoscaler(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify

resourceVersion. :return: V1HorizontalPodAutoscalerList

If the method is called asynchronously, returns the request thread.

list_namespaced_horizontal_pod_autoscaler_with_http_info (*namespace*,
***kwargs*)

list or watch objects of kind HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_horizontal_pod_autoscaler_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1HorizontalPodAutoscalerList

If the method is called asynchronously, returns the request thread.

patch_namespaced_horizontal_pod_autoscaler (*name*, *namespace*, *body*, ***kwargs*)

partially update the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_horizontal_pod_autoscaler(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the HorizontalPodAutoscaler (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

patch_namespaced_horizontal_pod_autoscaler_status (*name, namespace, body, **kwargs*)

partially update status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_horizontal_pod_autoscaler_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the HorizontalPodAutoscaler (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

patch_namespaced_horizontal_pod_autoscaler_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_horizontal_pod_autoscaler_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the HorizontalPodAutoscaler (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

patch_namespaced_horizontal_pod_autoscaler_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_horizontal_pod_autoscaler_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the HorizontalPodAutoscaler (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

read_namespaced_horizontal_pod_autoscaler (*name, namespace, **kwargs*)

read the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_horizontal_pod_autoscaler(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the HorizontalPodAutoscaler (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1HorizontalPodAutoscaler`

If the method is called asynchronously, returns the request thread.

read_namespaced_horizontal_pod_autoscaler_status (*name, namespace, **kwargs*)

read status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread`

```
= api.read_namespaced_horizontal_pod_autoscaler_status(name, namespace, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
read_namespaced_horizontal_pod_autoscaler_status_with_http_info (name,
                                                                    names-
                                                                    pace,
                                                                    **kwargs)
```

read status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_horizontal_pod_autoscaler_status_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
read_namespaced_horizontal_pod_autoscaler_with_http_info (name, namespace,
                                                            **kwargs)
```

read the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_horizontal_pod_autoscaler_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_horizontal_pod_autoscaler (name, namespace, body, **kwargs)
```

replace the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_horizontal_pod_autoscaler(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1HorizontalPodAutoscaler body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_horizontal_pod_autoscaler_status (name, namespace, body,
                                                       **kwargs)
```

replace status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_horizontal_pod_autoscaler_status(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param

V1HorizontalPodAutoscaler body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_horizontal_pod_autoscaler_status_with_http_info(name,
                                                                    names-
                                                                    pace,
                                                                    body,
                                                                    **kwargs)
```

replace status of the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_horizontal_pod_autoscaler_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1HorizontalPodAutoscaler body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_horizontal_pod_autoscaler_with_http_info(name, namespace,
                                                            body,
                                                            **kwargs)
```

replace the specified HorizontalPodAutoscaler This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_horizontal_pod_autoscaler_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the HorizontalPodAutoscaler (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1HorizontalPodAutoscaler body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: V1HorizontalPodAutoscaler

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.batch_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.batch_api.**BatchApi** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (**kwargs)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_group(async=True) >>> result = thread.get()

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

```
get_api_group_with_http_info (**kwargs)
    get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_group_with_http_info(async=True) >>> result = thread.get()
```

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.batch_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.batch_v1_api.**BatchV1Api** (api_client=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_job (namespace, body, **kwargs)

create a Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_job(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Job body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Job

If the method is called asynchronously, returns the request thread.

create_namespaced_job_with_http_info (namespace, body, **kwargs)

create a Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_job_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Job body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Job

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_job (namespace, **kwargs)

delete collection of Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_job(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is

true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_job_with_http_info(namespace, **kwargs)

delete collection of Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_job_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of

a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_job (*name, namespace, body, **kwargs*)

delete a Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_namespaced_job(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Job (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_job_with_http_info (*name, namespace, body, **kwargs*)

delete a Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_namespaced_job_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Job (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIResourceList`

If the method is called asynchronously, returns the request thread.

list_job_for_all_namespaces (**kwargs)

list or watch objects of kind Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_job_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: `V1JobList`

If the method is called asynchronously, returns the request thread.

list_job_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_job_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject

a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1JobList

If the method is called asynchronously, returns the request thread.

list_namespaced_job (*namespace*, ***kwargs*)

list or watch objects of kind Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_job(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field

is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1JobList

If the method is called asynchronously, returns the request thread.

list_namespaced_job_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_job_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1JobList

If the method is called asynchronously, returns the request thread.

patch_namespaced_job (*name, namespace, body, **kwargs*)

partially update the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_job(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

patch_namespaced_job_status (*name, namespace, body, **kwargs*)

partially update status of the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_job_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

patch_namespaced_job_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_job_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

patch_namespaced_job_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_job_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

read_namespaced_job (*name, namespace, **kwargs*)

read the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_job(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

read_namespaced_job_status (*name, namespace, **kwargs*)

read status of the specified Job This method makes a synchronous HTTP request by de-

fault. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_job_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

`read_namespaced_job_status_with_http_info` (*name, namespace, **kwargs*)

read status of the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_job_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

`read_namespaced_job_with_http_info` (*name, namespace, **kwargs*)

read the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_job_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_job` (*name, namespace, body, **kwargs*)

replace the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_job(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Job body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_job_status` (*name, namespace, body, **kwargs*)

replace status of the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_job_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Job (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Job body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Job`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_job_status_with_http_info` (*name, namespace, body, **kwargs*)

replace status of the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_job_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Job (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Job body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Job

If the method is called asynchronously, returns the request thread.

replace_namespaced_job_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Job This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_job_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Job (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Job body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Job

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.batch_v2alpha1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.apis.batch_v2alpha1_api.**BatchV2alpha1Api** (*api_client=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_cron_job (*namespace, body, **kwargs*)

create a CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_cron_job(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V2alpha1CronJob body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

create_namespaced_cron_job_with_http_info (*namespace, body, **kwargs*)

create a CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_cron_job_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V2alpha1CronJob body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_cron_job (*namespace, **kwargs*)

delete collection of CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_cron_job(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_cron_job_with_http_info(namespace, **kwargs)
delete collection of CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_cron_job_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that

can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_cron_job (*name, namespace, body, **kwargs*)

delete a CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_cron_job(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_cron_job_with_http_info (*name, namespace, body, **kwargs*)

delete a CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_cron_job_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set,

but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_cron_job_for_all_namespaces (**kwargs)

list or watch objects of kind CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_cron_job_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify

resourceVersion. :return: V2alpha1CronJobList

If the method is called asynchronously, returns the request thread.

list_cron_job_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_cron_job_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V2alpha1CronJobList

If the method is called asynchronously, returns the request thread.

list_namespaced_cron_job (namespace, **kwargs)

list or watch objects of kind CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_cron_job(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss

any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V2alpha1CronJobList

If the method is called asynchronously, returns the request thread.

list_namespaced_cron_job_with_http_info(namespace, **kwargs)

list or watch objects of kind CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_cron_job_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the

version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V2alpha1CronJobList

If the method is called asynchronously, returns the request thread.

patch_namespaced_cron_job (*name, namespace, body, **kwargs*)

partially update the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_cron_job(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

patch_namespaced_cron_job_status (*name, namespace, body, **kwargs*)

partially update status of the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_cron_job_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

patch_namespaced_cron_job_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_cron_job_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

patch_namespaced_cron_job_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_cron_job_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the CronJob (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V2alpha1CronJob

If the method is called asynchronously, returns the request thread.

read_namespaced_cron_job (*name, namespace, **kwargs*)

read the specified CronJob This method makes a synchronous HTTP request by default. To make an

asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_cron_job(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`read_namespaced_cron_job_status` (*name, namespace, **kwargs*)

read status of the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_cron_job_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`read_namespaced_cron_job_status_with_http_info` (*name, namespace, **kwargs*)

read status of the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_cron_job_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`read_namespaced_cron_job_with_http_info` (*name, namespace, **kwargs*)

read the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_cron_job_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_cron_job` (*name, namespace, body, **kwargs*)

replace the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_cron_job(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V2alpha1CronJob body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_cron_job_status` (*name, namespace, body, **kwargs*)

replace status of the specified CronJob This method makes a synchronous HTTP request by

default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_cron_job_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V2alpha1CronJob body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_cron_job_status_with_http_info` (*name, namespace, body, **kwargs*)

replace status of the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_cron_job_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V2alpha1CronJob body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_cron_job_with_http_info` (*name, namespace, body, **kwargs*)

replace the specified CronJob This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_cron_job_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the CronJob (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V2alpha1CronJob body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V2alpha1CronJob`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.certificates_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.certificates_api.CertificatesApi` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

`get_api_group` (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.


```
get_api_group_with_http_info (**kwargs)
    get information of a group This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_group_with_http_info(async=True) >>> result = thread.get()
```

:param async bool :return: V1APIGroup

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.certificates_v1alpha1_api module

kubernetes.client.apis.core_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.apis.core_api.CoreApi (api_client=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

```
get_api_versions (**kwargs)
    get available API versions This method makes a synchronous HTTP request by default. To make an
    asynchronous HTTP request, please pass async=True >>> thread = api.get_api_versions(async=True) >>>
    result = thread.get()
```

:param async bool :return: V1APIVersions

If the method is called asynchronously, returns the request thread.

```
get_api_versions_with_http_info (**kwargs)
    get available API versions This method makes a synchronous HTTP request by de-
    fault. To make an asynchronous HTTP request, please pass async=True >>> thread =
    api.get_api_versions_with_http_info(async=True) >>> result = thread.get()
```

:param async bool :return: V1APIVersions

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.core_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.apis.core_v1_api.CoreV1Api (api_client=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

connect_delete_namespaced_pod_proxy (*name, namespace, **kwargs*)

connect DELETE requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_pod_proxy_with_http_info (*name, namespace, **kwargs*)

connect DELETE requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_pod_proxy_with_path (*name, namespace, path, **kwargs*)

connect DELETE requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect DELETE requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_service_proxy (*name, namespace, **kwargs*)

connect DELETE requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_service_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_service_proxy_with_http_info (*name*, *namespace*,
***kwargs*)

connect DELETE requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_service_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_service_proxy_with_path (*name*, *namespace*, *path*,
***kwargs*)

connect DELETE requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_service_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_namespaced_service_proxy_with_path_with_http_info (*name*,
namespace,
path,
***kwargs*)

connect DELETE requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_node_proxy (*name*, ***kwargs*)

connect DELETE requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_node_proxy(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_node_proxy_with_http_info (*name*, ***kwargs*)

connect DELETE requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_node_proxy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect DELETE requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_node_proxy_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_delete_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect DELETE requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_delete_node_proxy_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_attach (*name*, *namespace*, ***kwargs*)

connect GET requests to attach of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_attach(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container in which to execute the command. Defaults to only container if there is only one container in the pod. :param `bool stderr`: Stderr if true indicates that stderr is to be redirected for the attach call. Defaults to true. :param `bool stdin`: Stdin if true, redirects the standard input stream of the pod for this call. Defaults to false. :param `bool stdout`: Stdout if true indicates that stdout is to be redirected for the attach call. Defaults to true. :param `bool tty`: TTY if true indicates that a tty will be allocated for the attach call. This is passed through the container runtime so the tty is allocated on the worker node by the container runtime. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_attach_with_http_info (*name*, *namespace*, ***kwargs*)

connect GET requests to attach of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_attach_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container in which to execute the command. Defaults to only container if there is only one container in the pod. :param `bool stderr`: Stderr if true indicates that stderr is to be redirected for the attach call. Defaults to true. :param `bool stdin`: Stdin if true, redirects the standard input stream of the pod for this call. Defaults to false. :param

bool stdout: Stdout if true indicates that stdout is to be redirected for the attach call. Defaults to true.
:param bool tty: TTY if true indicates that a tty will be allocated for the attach call. This is passed through the container runtime so the tty is allocated on the worker node by the container runtime. Defaults to false.
:return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_exec (*name, namespace, **kwargs*)

connect GET requests to exec of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_exec(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str command: Command is the remote command to execute. argv array. Not executed within a shell. :param str container: Container in which to execute the command. Defaults to only container if there is only one container in the pod. :param bool stderr: Redirect the standard error stream of the pod for this call. Defaults to true. :param bool stdin: Redirect the standard input stream of the pod for this call. Defaults to false. :param bool stdout: Redirect the standard output stream of the pod for this call. Defaults to true. :param bool tty: TTY if true indicates that a tty will be allocated for the exec call. Defaults to false. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_exec_with_http_info (*name, namespace, **kwargs*)

connect GET requests to exec of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_exec_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str command: Command is the remote command to execute. argv array. Not executed within a shell. :param str container: Container in which to execute the command. Defaults to only container if there is only one container in the pod. :param bool stderr: Redirect the standard error stream of the pod for this call. Defaults to true. :param bool stdin: Redirect the standard input stream of the pod for this call. Defaults to false. :param bool stdout: Redirect the standard output stream of the pod for this call. Defaults to true. :param bool tty: TTY if true indicates that a tty will be allocated for the exec call. Defaults to false. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_portforward (*name, namespace, **kwargs*)

connect GET requests to portforward of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_portforward(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param int ports: List of ports to forward Required when using WebSockets :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_portforward_with_http_info (*name, namespace, **kwargs*)

connect GET requests to portforward of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_portforward_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param int ports: List of ports to forward Required

when using WebSockets :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_proxy (*name, namespace, **kwargs*)

connect GET requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_proxy_with_http_info (*name, namespace, **kwargs*)

connect GET requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_proxy_with_path (*name, namespace, path, **kwargs*)

connect GET requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect GET requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_service_proxy (*name, namespace, **kwargs*)

connect GET requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_service_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_service_proxy_with_http_info (*name*, *namespace*,
***kwargs*)

connect GET requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_service_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_service_proxy_with_path (*name*, *namespace*, *path*,
***kwargs*)

connect GET requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_service_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_namespaced_service_proxy_with_path_with_http_info (*name*,
namespace, *path*,
***kwargs*)

connect GET requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_node_proxy (*name*, ***kwargs*)

connect GET requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_node_proxy(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_node_proxy_with_http_info (*name*, ***kwargs*)

connect GET requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_node_proxy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect GET requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_node_proxy_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_get_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect GET requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_get_node_proxy_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_pod_proxy (*name*, *namespace*, ***kwargs*)

connect HEAD requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_head_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_pod_proxy_with_http_info (*name*, *namespace*, ***kwargs*)

connect HEAD requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_head_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_pod_proxy_with_path (*name*, *namespace*, *path*, ***kwargs*)

connect HEAD requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_head_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param

str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect HEAD requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_service_proxy (*name, namespace, **kwargs*)

connect HEAD requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_namespaced_service_proxy(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_service_proxy_with_http_info (*name, namespace, **kwargs*)

connect HEAD requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_namespaced_service_proxy_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_service_proxy_with_path (*name, namespace, path, **kwargs*)

connect HEAD requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_namespaced_service_proxy_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_namespaced_service_proxy_with_path_with_http_info (*name*,
names-
pace, *path*,
***kwargs*)

connect HEAD requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_node_proxy (*name*, ***kwargs*)

connect HEAD requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_node_proxy(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_node_proxy_with_http_info (*name*, ***kwargs*)

connect HEAD requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_node_proxy_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect HEAD requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_node_proxy_with_path(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_head_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect HEAD requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_head_node_proxy_with_path_with_http_info(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_options_namespaced_pod_proxy (*name*, *namespace*, ***kwargs*)

connect OPTIONS requests to proxy of Pod This method makes a synchronous HTTP request

by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

`connect_options_namespaced_pod_proxy_with_http_info` (*name*, *namespace*,
***kwargs*)

connect OPTIONS requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

`connect_options_namespaced_pod_proxy_with_path` (*name*, *namespace*, *path*,
***kwargs*)

connect OPTIONS requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

`connect_options_namespaced_pod_proxy_with_path_with_http_info` (*name*,
names-
pace, path,
***kwargs*)

connect OPTIONS requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

`connect_options_namespaced_service_proxy` (*name*, *namespace*, ***kwargs*)

connect OPTIONS requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_service_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_namespaced_service_proxy_with_http_info (*name, namespace, **kwargs*)

connect OPTIONS requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_service_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_namespaced_service_proxy_with_path (*name, namespace, path, **kwargs*)

connect OPTIONS requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_service_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_namespaced_service_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect OPTIONS requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_node_proxy (*name, **kwargs*)

connect OPTIONS requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_node_proxy(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_node_proxy_with_http_info (*name*, ***kwargs*)

connect OPTIONS requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_node_proxy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect OPTIONS requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_node_proxy_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_options_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect OPTIONS requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_options_node_proxy_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_pod_proxy (*name*, *namespace*, ***kwargs*)

connect PATCH requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_pod_proxy_with_http_info (*name*, *namespace*, ***kwargs*)

connect PATCH requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_pod_proxy_with_path (*name*, *namespace*, *path*, ***kwargs*)

connect PATCH requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect PATCH requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_patch_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_service_proxy (*name, namespace, **kwargs*)

connect PATCH requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_patch_namespaced_service_proxy(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_service_proxy_with_http_info (*name, namespace, **kwargs*)

connect PATCH requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_patch_namespaced_service_proxy_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_patch_namespaced_service_proxy_with_path (*name, namespace, path, **kwargs*)

connect PATCH requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_patch_namespaced_service_proxy_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy.

```
//localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path
is _search?q=user:kimchy. :return: str
```

If the method is called asynchronously, returns the request thread.

```
connect_patch_namespaced_service_proxy_with_path_with_http_info (name,
                                                                    names-
                                                                    pace,
                                                                    path,
                                                                    **kwargs)
```

connect PATCH requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

```
connect_patch_node_proxy (name, **kwargs)
```

connect PATCH requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_node_proxy(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

```
connect_patch_node_proxy_with_http_info (name, **kwargs)
```

connect PATCH requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_node_proxy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

```
connect_patch_node_proxy_with_path (name, path, **kwargs)
```

connect PATCH requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_node_proxy_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

```
connect_patch_node_proxy_with_path_with_http_info (name, path, **kwargs)
```

connect PATCH requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_patch_node_proxy_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_attach (*name, namespace, **kwargs*)

connect POST requests to attach of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_attach(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container in which to execute the command. Defaults to only container if there is only one container in the pod. :param `bool stderr`: Stderr if true indicates that stderr is to be redirected for the attach call. Defaults to true. :param `bool stdin`: Stdin if true, redirects the standard input stream of the pod for this call. Defaults to false. :param `bool stdout`: Stdout if true indicates that stdout is to be redirected for the attach call. Defaults to true. :param `bool tty`: TTY if true indicates that a tty will be allocated for the attach call. This is passed through the container runtime so the tty is allocated on the worker node by the container runtime. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_attach_with_http_info (*name, namespace, **kwargs*)

connect POST requests to attach of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_attach_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container in which to execute the command. Defaults to only container if there is only one container in the pod. :param `bool stderr`: Stderr if true indicates that stderr is to be redirected for the attach call. Defaults to true. :param `bool stdin`: Stdin if true, redirects the standard input stream of the pod for this call. Defaults to false. :param `bool stdout`: Stdout if true indicates that stdout is to be redirected for the attach call. Defaults to true. :param `bool tty`: TTY if true indicates that a tty will be allocated for the attach call. This is passed through the container runtime so the tty is allocated on the worker node by the container runtime. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_exec (*name, namespace, **kwargs*)

connect POST requests to exec of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_exec(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str command`: Command is the remote command to execute. `argv` array. Not executed within a shell. :param `str container`: Container in which to execute the command. Defaults to only container if there is only one container in the pod. :param `bool stderr`: Redirect the standard error stream of the pod for this call. Defaults to true. :param `bool stdin`: Redirect the standard input stream of the pod for this call. Defaults to false. :param `bool stdout`: Redirect the standard output stream of the pod for this call. Defaults to true. :param `bool tty`: TTY if true indicates that a tty will be allocated for the exec call. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_exec_with_http_info (*name, namespace, **kwargs*)

connect POST requests to exec of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_exec_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str command: Command is the remote command to execute. argv array. Not executed within a shell. :param str container: Container in which to execute the command. Defaults to only container if there is only one container in the pod. :param bool stderr: Redirect the standard error stream of the pod for this call. Defaults to true. :param bool stdin: Redirect the standard input stream of the pod for this call. Defaults to false. :param bool stdout: Redirect the standard output stream of the pod for this call. Defaults to true. :param bool tty: TTY if true indicates that a tty will be allocated for the exec call. Defaults to false. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_portforward (*name, namespace, **kwargs*)

connect POST requests to portforward of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_portforward(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param int ports: List of ports to forward Required when using WebSockets :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_portforward_with_http_info (*name, namespace, **kwargs*)

connect POST requests to portforward of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_portforward_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param int ports: List of ports to forward Required when using WebSockets :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_proxy (*name, namespace, **kwargs*)

connect POST requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_proxy_with_http_info (*name, namespace, **kwargs*)

connect POST requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_post_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_proxy_with_path (*name, namespace, path, **kwargs*)

connect POST requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.connect_post_namespaced_pod_proxy_with_path(name, namespace, path, async=True) >>> result = thread.get()
```

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect POST requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to pod. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_service_proxy (*name, namespace, **kwargs*)

connect POST requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_namespaced_service_proxy(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_service_proxy_with_http_info (*name, namespace, **kwargs*)

connect POST requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_namespaced_service_proxy_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_service_proxy_with_path (*name, namespace, path, **kwargs*)

connect POST requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_namespaced_service_proxy_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the part of URLs that include service endpoints, suffixes, and pa-

parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_namespaced_service_proxy_with_path_with_http_info (*name*,
namespace,
path,
***kwargs*)

connect POST requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :param str path2: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_node_proxy (*name*, ***kwargs*)

connect POST requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_node_proxy(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_node_proxy_with_http_info (*name*, ***kwargs*)

connect POST requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_node_proxy_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect POST requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_node_proxy_with_path(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_post_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect POST requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.connect_post_node_proxy_with_path_with_http_info(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :param str path2: Path is the URL path to use for the current proxy request to node. :return: str

If the method is called asynchronously, returns the request thread.

connect_put_namespaced_pod_proxy (*name, namespace, **kwargs*)

connect PUT requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_pod_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_namespaced_pod_proxy_with_http_info (*name, namespace, **kwargs*)

connect PUT requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_pod_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_namespaced_pod_proxy_with_path (*name, namespace, path, **kwargs*)

connect PUT requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_pod_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_namespaced_pod_proxy_with_path_with_http_info (*name, namespace, path, **kwargs*)

connect PUT requests to proxy of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_pod_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to pod. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_namespaced_service_proxy (*name, namespace, **kwargs*)

connect PUT requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_service_proxy(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

```
connect_put_namespaced_service_proxy_with_http_info (name, namespace,
                                                         **kwargs)
```

connect PUT requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_service_proxy_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

```
connect_put_namespaced_service_proxy_with_path (name, namespace, path,
                                                  **kwargs)
```

connect PUT requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_service_proxy_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

```
connect_put_namespaced_service_proxy_with_path_with_http_info (name,
                                                                    names-
                                                                    pace, path,
                                                                    **kwargs)
```

connect PUT requests to proxy of Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_namespaced_service_proxy_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the part of URLs that include service endpoints, suffixes, and parameters to use for the current proxy request to service. For example, the whole request URL is http://localhost/api/v1/namespaces/kube-system/services/elasticsearch-logging/_search?q=user:kimchy. Path is `_search?q=user:kimchy`. :return: `str`

If the method is called asynchronously, returns the request thread.

```
connect_put_node_proxy (name, **kwargs)
```

connect PUT requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_node_proxy(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_node_proxy_with_http_info (*name*, ***kwargs*)

connect PUT requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_node_proxy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_node_proxy_with_path (*name*, *path*, ***kwargs*)

connect PUT requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_node_proxy_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

connect_put_node_proxy_with_path_with_http_info (*name*, *path*, ***kwargs*)

connect PUT requests to proxy of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.connect_put_node_proxy_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :param `str path2`: Path is the URL path to use for the current proxy request to node. :return: `str`

If the method is called asynchronously, returns the request thread.

create_namespace (*body*, ***kwargs*)

create a Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespace(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

create_namespace_with_http_info (*body*, ***kwargs*)

create a Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespace_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

create_namespaced_binding (*namespace*, *body*, ***kwargs*)

create a Binding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_binding(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Binding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Binding`

If the method is called asynchronously, returns the request thread.

create_namespaced_binding_with_http_info (*namespace, body, **kwargs*)

create a Binding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_binding_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Binding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Binding`

If the method is called asynchronously, returns the request thread.

create_namespaced_config_map (*namespace, body, **kwargs*)

create a ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_config_map(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ConfigMap body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

create_namespaced_config_map_with_http_info (*namespace, body, **kwargs*)

create a ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_config_map_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ConfigMap body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

create_namespaced_endpoints (*namespace, body, **kwargs*)

create Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_endpoints(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Endpoints body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

create_namespaced_endpoints_with_http_info (*namespace, body, **kwargs*)

create Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_endpoints_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Endpoints body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

create_namespaced_event (*namespace, body, **kwargs*)

create an Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_event(namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Event body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

create_namespaced_event_with_http_info (*namespace, body, **kwargs*)

create an Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_event_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Event body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

create_namespaced_limit_range (*namespace, body, **kwargs*)

create a LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_limit_range(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1LimitRange body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

create_namespaced_limit_range_with_http_info (*namespace, body, **kwargs*)

create a LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_limit_range_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1LimitRange body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

create_namespaced_persistent_volume_claim (*namespace, body, **kwargs*)

create a PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_persistent_volume_claim(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1PersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

create_namespaced_persistent_volume_claim_with_http_info (*namespace, body, **kwargs*)

create a PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_persistent_volume_claim_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1PersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

create_namespaced_pod (*namespace, body, **kwargs*)

create a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Pod body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_binding (*name, namespace, body, **kwargs*)

create binding of a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_binding(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Binding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Binding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Binding`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_binding_with_http_info (*name, namespace, body, **kwargs*)

create binding of a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_binding_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Binding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Binding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Binding`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_eviction (*name, namespace, body, **kwargs*)

create eviction of a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_eviction(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Eviction (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Eviction body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Eviction`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_eviction_with_http_info (*name, namespace, body, **kwargs*)

create eviction of a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_eviction_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Eviction (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Eviction body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Eviction`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_template (*namespace, body, **kwargs*)

create a PodTemplate This method makes a synchronous HTTP request by de-

fault. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_template(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1PodTemplate body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`create_namespaced_pod_template_with_http_info` (*namespace, body, **kwargs*)

create a `PodTemplate` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_template_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1PodTemplate body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`create_namespaced_pod_with_http_info` (*namespace, body, **kwargs*)

create a `Pod` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Pod body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`create_namespaced_replication_controller` (*namespace, body, **kwargs*)

create a `ReplicationController` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_replication_controller(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

`create_namespaced_replication_controller_with_http_info` (*namespace, body, **kwargs*)

create a `ReplicationController` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_replication_controller_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

`create_namespaced_resource_quota` (*namespace, body, **kwargs*)

create a `ResourceQuota` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_resource_quota(namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1ResourceQuota body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

create_namespaced_resource_quota_with_http_info (*namespace, body, **kwargs*)
 create a ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_resource_quota_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1ResourceQuota body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

create_namespaced_secret (*namespace, body, **kwargs*)
 create a Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_secret(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Secret body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Secret

If the method is called asynchronously, returns the request thread.

create_namespaced_secret_with_http_info (*namespace, body, **kwargs*)
 create a Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_secret_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Secret body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Secret

If the method is called asynchronously, returns the request thread.

create_namespaced_service (*namespace, body, **kwargs*)
 create a Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_service(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Service body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Service

If the method is called asynchronously, returns the request thread.

create_namespaced_service_account (*namespace, body, **kwargs*)
 create a ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_service_account(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1ServiceAccount body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ServiceAccount

If the method is called asynchronously, returns the request thread.

create_namespaced_service_account_with_http_info (*namespace, body, **kwargs*)

create a ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_service_account_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ServiceAccount body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

create_namespaced_service_with_http_info (*namespace, body, **kwargs*)

create a Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_service_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Service body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

create_node (*body, **kwargs*)

create a Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_node(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

create_node_with_http_info (*body, **kwargs*)

create a Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_node_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

create_persistent_volume (*body, **kwargs*)

create a PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_persistent_volume(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

create_persistent_volume_with_http_info (*body, **kwargs*)

create a PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_persistent_volume_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_config_map (*namespace*, ***kwargs*)

delete collection of ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_config_map(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_config_map_with_http_info (*namespace*, ***kwargs*)

delete collection of ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_config_map_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss

any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_endpoints (*namespace*, ***kwargs*)

delete collection of Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_endpoints(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the

version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_endpoints_with_http_info (*namespace*, ***kwargs*)

delete collection of Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_endpoints_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_event (*namespace*, ***kwargs*)

delete collection of Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_event(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_event_with_http_info (*namespace*, ***kwargs*)
delete collection of Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_event_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that

can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_limit_range (namespace, **kwargs)

delete collection of LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_limit_range(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_limit_range_with_http_info(namespace, ***kwargs*)

delete collection of LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_limit_range_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_persistent_volume_claim(namespace, ***kwargs*)

delete collection of PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_persistent_volume_claim(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined,

clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_persistent_volume_claim_with_http_info(*namespace*,
***kwargs*)

delete collection of PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_persistent_volume_claim_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available.

Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod(namespace, **kwargs)

delete collection of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_pod(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod_template (*namespace*, ***kwargs*)

delete collection of PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_pod_template(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod_template_with_http_info (*namespace*, ***kwargs*)

delete collection of PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_pod_template_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is

true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod_with_http_info(namespace, **kwargs)

delete collection of Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_pod_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of

a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_replication_controller(namespace, **kwargs)

delete collection of ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_namespaced_replication_controller(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_replication_controller_with_http_info(namespace, **kwargs)

delete collection of ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =

```
api.delete_collection_namespaced_replication_controller_with_http_info(namespace, async=True) >>>
result = thread.get()
```

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_resource_quota (*namespace*, ***kwargs*)

delete collection of ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_resource_quota(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return

for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_resource_quota_with_http_info (*namespace*,
***kwargs*)

delete collection of ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_resource_quota_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then

the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_secret (namespace, **kwargs)

delete collection of Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_secret(namespace, async=True) >>> result = thread.get()

:param bool async: :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_secret_with_http_info (namespace, **kwargs)

delete collection of Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_secret_with_http_info(namespace, async=True) >>> result = thread.get()

:param bool async: :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined,

clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_service_account (*namespace*, ***kwargs*)

delete collection of ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_service_account(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is

specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_service_account_with_http_info (*namespace*,
***kwargs*)

delete collection of ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_service_account_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_node (**kwargs)

delete collection of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_node(async=True)`
>>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_node_with_http_info (**kwargs)

delete collection of Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_node_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are in-

cluded in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_persistent_volume (***kwargs*)

delete collection of PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_persistent_volume(async=True)` >>> `result = thread.get()`

:param bool async :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then

the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_persistent_volume_with_http_info (**kwargs)

delete collection of PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_persistent_volume_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespace (name, body, **kwargs)

delete a Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespace(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Namespace (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool

orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespace_with_http_info (*name, body, **kwargs*)

delete a Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespace_with_http_info(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Namespace (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_config_map (*name, namespace, body, **kwargs*)

delete a ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_config_map(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ConfigMap (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_config_map_with_http_info (*name, namespace, body, **kwargs*)

delete a ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_config_map_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ConfigMap (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates

delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_endpoints (*name, namespace, body, **kwargs*)

delete Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_endpoints(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Endpoints (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_endpoints_with_http_info (*name, namespace, body, **kwargs*)

delete Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_endpoints_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Endpoints (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_event (*name, namespace, body, **kwargs*)

delete an Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_event(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Event (required) :param str namespace: object name and

auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_event_with_http_info (*name, namespace, body, **kwargs*)

delete an Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_event_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Event (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_limit_range (*name, namespace, body, **kwargs*)

delete a LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_limit_range(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_limit_range_with_http_info (*name, namespace, body, **kwargs*)

delete a LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.delete_namespaced_limit_range_with_http_info(name, namespace, body, async=True) >>> result = thread.get()
```

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_persistent_volume_claim (*name, namespace, body, **kwargs*)

delete a PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_persistent_volume_claim(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_persistent_volume_claim_with_http_info (*name, namespace, body, **kwargs*)

delete a PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_persistent_volume_claim_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod (*name, namespace, body, **kwargs*)

delete a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_pod(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod_template (*name, namespace, body, **kwargs*)

delete a PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_pod_template(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod_template_with_http_info (*name, namespace, body, **kwargs*)

delete a PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_pod_template_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or Or-

phanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod_with_http_info (*name, namespace, body, **kwargs*)

delete a Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_pod_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_replication_controller (*name, namespace, body, **kwargs*)

delete a ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_replication_controller(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_replication_controller_with_http_info (*name, namespace, body, **kwargs*)

delete a ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_replication_controller_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool`

orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_resource_quota (*name, namespace, body, **kwargs*)

delete a ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_resource_quota(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_resource_quota_with_http_info (*name, namespace, body, **kwargs*)

delete a ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_resource_quota_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_secret (*name, namespace, body, **kwargs*)

delete a Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_secret(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Secret (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in

seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_secret_with_http_info (*name, namespace, body, **kwargs*)

delete a Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_secret_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Secret (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_service (*name, namespace, **kwargs*)

delete a Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_service(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If ‘true’, then the output is pretty printed. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_service_account (*name, namespace, body, **kwargs*)

delete a ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_service_account(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ServiceAccount (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If ‘true’, then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or Or-

phanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_service_account_with_http_info (*name*, *namespace*, *body*,
***kwargs*)

delete a ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_service_account_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_service_with_http_info (*name*, *namespace*, ***kwargs*)

delete a Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_node (*name*, *body*, ***kwargs*)

delete a Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_node(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_node_with_http_info (*name*, *body*, ***kwargs*)

delete a Node This method makes a synchronous HTTP request by default. To make an asyn-

chronous HTTP request, please pass `async=True` >>> `thread = api.delete_node_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_persistent_volume` (*name*, *body*, ***kwargs*)

delete a PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_persistent_volume(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_persistent_volume_with_http_info` (*name*, *body*, ***kwargs*)

delete a PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_persistent_volume_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`get_api_resources` (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources(async=True)` >>>

```
result = thread.get()
```

```
:param async bool :return: V1APIResourceList
```

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources_with_http_info(async=True)` >>> `result = thread.get()`

```
:param async bool :return: V1APIResourceList
```

If the method is called asynchronously, returns the request thread.

list_component_status (**kwargs)

list objects of kind ComponentStatus This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_component_status(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ComponentStatusList

If the method is called asynchronously, returns the request thread.

list_component_status_with_http_info (**kwargs)

list objects of kind ComponentStatus This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_component_status_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ComponentStatusList

If the method is called asynchronously, returns the request thread.

list_config_map_for_all_namespaces (**kwargs)

list or watch objects of kind ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_config_map_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of

the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ConfigMapList

If the method is called asynchronously, returns the request thread.

list_config_map_for_all_namespaces_with_http_info (kwargs)**

list or watch objects of kind ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_config_map_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ConfigMapList

If the method is called asynchronously, returns the request thread.

list_endpoints_for_all_namespaces (**kwargs)

list or watch objects of kind Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_endpoints_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EndpointsList

If the method is called asynchronously, returns the request thread.

list_endpoints_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_endpoints_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the

list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EndpointsList

If the method is called asynchronously, returns the request thread.

list_event_for_all_namespaces (**kwargs)

list or watch objects of kind Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_event_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EventList

If the method is called asynchronously, returns the request thread.

list_event_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_event_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EventList

If the method is called asynchronously, returns the request thread.

list_limit_range_for_all_namespaces (**kwargs)

list or watch objects of kind LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_limit_range_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value

returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1LimitRangeList

If the method is called asynchronously, returns the request thread.

list_limit_range_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_limit_range_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list

result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1LimitRangeList

If the method is called asynchronously, returns the request thread.

list_namespace (**kwargs)

list or watch objects of kind Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespace(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1NamespaceList

If the method is called asynchronously, returns the request thread.

list_namespace_with_http_info (**kwargs)

list or watch objects of kind Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespace_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server de-

finied, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1NamespaceList

If the method is called asynchronously, returns the request thread.

list_namespaced_config_map(namespace, **kwargs)

list or watch objects of kind ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_config_map(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is

specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ConfigMapList

If the method is called asynchronously, returns the request thread.

list_namespaced_config_map_with_http_info(namespace, **kwargs)

list or watch objects of kind ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_config_map_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ConfigMapList

If the method is called asynchronously, returns the request thread.

list_namespaced_endpoints (*namespace*, ***kwargs*)

list or watch objects of kind Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_endpoints(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given `rv`. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1EndpointsList`

If the method is called asynchronously, returns the request thread.

list_namespaced_endpoints_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_endpoints_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are

included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EndpointsList

If the method is called asynchronously, returns the request thread.

list_namespaced_event (*namespace, **kwargs*)

list or watch objects of kind Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_event(namespace, async=True) >>> result = thread.get()

:param bool async :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version

of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EventList

If the method is called asynchronously, returns the request thread.

list_namespaced_event_with_http_info (namespace, **kwargs)

list or watch objects of kind Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_event_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1EventList

If the method is called asynchronously, returns the request thread.

list_namespaced_limit_range (namespace, **kwargs)

list or watch objects of kind LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_limit_range(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined,

clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1LimitRangeList

If the method is called asynchronously, returns the request thread.

list_namespaced_limit_range_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_limit_range_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is

specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1LimitRangeList

If the method is called asynchronously, returns the request thread.

list_namespaced_persistent_volume_claim(namespace, **kwargs)

list or watch objects of kind PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_persistent_volume_claim(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeClaimList

If the method is called asynchronously, returns the request thread.

list_namespaced_persistent_volume_claim_with_http_info (*namespace*,
***kwargs*)

list or watch objects of kind PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_persistent_volume_claim_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeClaimList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod (*namespace*, ***kwargs*)

list or watch objects of kind Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_pod(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss

any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod_template (*namespace*, ***kwargs*)

list or watch objects of kind PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_pod_template(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the

version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodTemplateList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod_template_with_http_info (namespace, **kwargs)

list or watch objects of kind PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_pod_template_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodTemplateList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod_with_http_info (namespace, **kwargs)

list or watch objects of kind Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_pod_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects

(required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodList

If the method is called asynchronously, returns the request thread.

list_namespaced_replication_controller (namespace, **kwargs)

list or watch objects of kind ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_replication_controller(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and

clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1ReplicationControllerList`

If the method is called asynchronously, returns the request thread.

`list_namespaced_replication_controller_with_http_info` (*namespace*, ***kwargs*)

list or watch objects of kind `ReplicationController` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_replication_controller_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `namespace`: object name and auth scope, such as for teams and projects (required) :param str `pretty`: If 'true', then the output is pretty printed. :param str `_continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes

to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ReplicationControllerList

If the method is called asynchronously, returns the request thread.

list_namespaced_resource_quota (namespace, **kwargs)

list or watch objects of kind ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_resource_quota(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ResourceQuotaList

If the method is called asynchronously, returns the request thread.

list_namespaced_resource_quota_with_http_info (namespace, **kwargs)

list or watch objects of kind ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_resource_quota_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating

the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ResourceQuotaList

If the method is called asynchronously, returns the request thread.

list_namespaced_secret (namespace, **kwargs)

list or watch objects of kind Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_secret(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes

referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1SecretList

If the method is called asynchronously, returns the request thread.

list_namespaced_secret_with_http_info (namespace, **kwargs)

list or watch objects of kind Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_secret_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1SecretList

If the method is called asynchronously, returns the request thread.

list_namespaced_service (namespace, **kwargs)

list or watch objects of kind Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =


```
api.list_namespaced_service(namespace, async=True) >>> result = thread.get()
```

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceList

If the method is called asynchronously, returns the request thread.

list_namespaced_service_account (*namespace*, ***kwargs*)

list or watch objects of kind ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_service_account(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that

can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceAccountList

If the method is called asynchronously, returns the request thread.

list_namespaced_service_account_with_http_info (namespace, **kwargs)

list or watch objects of kind ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_service_account_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceAccountList

If the method is called asynchronously, returns the request thread.

list_namespaced_service_with_http_info (namespace, **kwargs)

list or watch objects of kind Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_service_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceList

If the method is called asynchronously, returns the request thread.

list_node (**kwargs)

list or watch objects of kind Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_node(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating

the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1NodeList

If the method is called asynchronously, returns the request thread.

list_node_with_http_info (**kwargs)

list or watch objects of kind Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_node_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of

a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1NodeList

If the method is called asynchronously, returns the request thread.

list_persistent_volume (**kwargs)

list or watch objects of kind PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_persistent_volume(async=True)` >>> `result = thread.get()`

:param bool async :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeList

If the method is called asynchronously, returns the request thread.

list_persistent_volume_claim_for_all_namespaces (**kwargs)

list or watch objects of kind PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_persistent_volume_claim_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param bool async :param str _continue: The continue option should be set when retrieving more results

from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeClaimList

If the method is called asynchronously, returns the request thread.

list_persistent_volume_claim_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_persistent_volume_claim_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of

the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeClaimList

If the method is called asynchronously, returns the request thread.

list_persistent_volume_with_http_info (**kwargs)

list or watch objects of kind PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_persistent_volume_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PersistentVolumeList

If the method is called asynchronously, returns the request thread.

list_pod_for_all_namespaces (**kwargs)

list or watch objects of kind Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_pod_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodList

If the method is called asynchronously, returns the request thread.

list_pod_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_pod_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the

list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodList

If the method is called asynchronously, returns the request thread.

list_pod_template_for_all_namespaces (**kwargs)

list or watch objects of kind PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_pod_template_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodTemplateList

If the method is called asynchronously, returns the request thread.

list_pod_template_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_pod_template_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1PodTemplateList

If the method is called asynchronously, returns the request thread.

list_replication_controller_for_all_namespaces (**kwargs)

list or watch objects of kind ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_replication_controller_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value

returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ReplicationControllerList

If the method is called asynchronously, returns the request thread.

list_replication_controller_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_replication_controller_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list

result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ReplicationControllerList

If the method is called asynchronously, returns the request thread.

list_resource_quota_for_all_namespaces (**kwargs)

list or watch objects of kind ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_resource_quota_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ResourceQuotaList

If the method is called asynchronously, returns the request thread.

list_resource_quota_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_resource_quota_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous

query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is `true`. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is `true`. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `pretty`: If 'true', then the output is pretty printed. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1ResourceQuotaList`

If the method is called asynchronously, returns the request thread.

list_secret_for_all_namespaces (***kwargs*)

list or watch objects of kind `Secret` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_secret_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param bool `async` :param str `_continue`: The `continue` option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the `continue` value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is `true`. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is `true`. The server guarantees that the objects

returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1SecretList

If the method is called asynchronously, returns the request thread.

list_secret_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_secret_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1SecretList

If the method is called asynchronously, returns the request thread.

list_service_account_for_all_namespaces (**kwargs)

list or watch objects of kind ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_service_account_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceAccountList

If the method is called asynchronously, returns the request thread.

`list_service_account_for_all_namespaces_with_http_info(kwargs)`**

list or watch objects of kind ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_service_account_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting

a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceAccountList

If the method is called asynchronously, returns the request thread.

list_service_for_all_namespaces (**kwargs)

list or watch objects of kind Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_service_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes

to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceList

If the method is called asynchronously, returns the request thread.

list_service_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_service_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1ServiceList

If the method is called asynchronously, returns the request thread.

patch_namespace (name, body, **kwargs)

partially update the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespace(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Namespace (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Namespace

If the method is called asynchronously, returns the request thread.

patch_namespace_status (name, body, **kwargs)

partially update status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespace_status(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Namespace (required) :param object body: (required)
:param str pretty: If 'true', then the output is pretty printed. :return: V1Namespace

If the method is called asynchronously, returns the request thread.

patch_namespace_status_with_http_info (*name, body, **kwargs*)

partially update status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespace_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Namespace (required) :param object body: (required)
:param str pretty: If 'true', then the output is pretty printed. :return: V1Namespace

If the method is called asynchronously, returns the request thread.

patch_namespace_with_http_info (*name, body, **kwargs*)

partially update the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespace_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Namespace (required) :param object body: (required)
:param str pretty: If 'true', then the output is pretty printed. :return: V1Namespace

If the method is called asynchronously, returns the request thread.

patch_namespaced_config_map (*name, namespace, body, **kwargs*)

partially update the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_config_map(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ConfigMap (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ConfigMap

If the method is called asynchronously, returns the request thread.

patch_namespaced_config_map_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_config_map_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ConfigMap (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ConfigMap

If the method is called asynchronously, returns the request thread.

patch_namespaced_endpoints (*name, namespace, body, **kwargs*)

partially update the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_endpoints(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Endpoints (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Endpoints

If the method is called asynchronously, returns the request thread.

patch_namespaced_endpoints_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.patch_namespaced_endpoints_with_http_info(name, namespace, body, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the Endpoints (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Endpoints

If the method is called asynchronously, returns the request thread.

patch_namespaced_event (*name, namespace, body, **kwargs*)

partially update the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_event(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Event (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

patch_namespaced_event_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_event_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Event (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

patch_namespaced_limit_range (*name, namespace, body, **kwargs*)

partially update the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_limit_range(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

patch_namespaced_limit_range_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_limit_range_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

patch_namespaced_persistent_volume_claim (*name, namespace, body, **kwargs*)

partially update the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_persistent_volume_claim(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

patch_namespaced_persistent_volume_claim_status (*name, namespace, body, **kwargs*)

partially update status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_persistent_volume_claim_status(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

patch_namespaced_persistent_volume_claim_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_persistent_volume_claim_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

patch_namespaced_persistent_volume_claim_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_persistent_volume_claim_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod (*name, namespace, body, **kwargs*)

partially update the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_pod(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Pod

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod_status (*name, namespace, body, **kwargs*)

partially update status of the specified Pod This method makes a synchronous HTTP request

by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_pod_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_pod_status_with_http_info` (*name, namespace, body, **kwargs*)

partially update status of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_pod_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_pod_template` (*name, namespace, body, **kwargs*)

partially update the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_pod_template(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_pod_template_with_http_info` (*name, namespace, body, **kwargs*)

partially update the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_pod_template_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_pod_with_http_info` (*name, namespace, body, **kwargs*)

partially update the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_pod_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_replication_controller` (*name, namespace, body, **kwargs*)

partially update the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ReplicationController (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_scale_with_http_info (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_status (*name, namespace, body, **kwargs*)

partially update status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ReplicationController (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ReplicationController (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_replication_controller_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicationController (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

patch_namespaced_resource_quota (*name, namespace, body, **kwargs*)

partially update the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_resource_quota(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

patch_namespaced_resource_quota_status (*name, namespace, body, **kwargs*)

partially update status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_resource_quota_status(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

patch_namespaced_resource_quota_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_resource_quota_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

patch_namespaced_resource_quota_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_resource_quota_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

patch_namespaced_secret (*name, namespace, body, **kwargs*)

partially update the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_secret(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Secret (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Secret`

If the method is called asynchronously, returns the request thread.

patch_namespaced_secret_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_secret_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Secret (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Secret`

If the method is called asynchronously, returns the request thread.

patch_namespaced_service (*name, namespace, body, **kwargs*)

partially update the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

patch_namespaced_service_account (*name, namespace, body, **kwargs*)

partially update the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service_account(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

patch_namespaced_service_account_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service_account_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

patch_namespaced_service_status (*name, namespace, body, **kwargs*)

partially update status of the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Service

If the method is called asynchronously, returns the request thread.

patch_namespaced_service_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Service

If the method is called asynchronously, returns the request thread.

patch_namespaced_service_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_service_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Service

If the method is called asynchronously, returns the request thread.

patch_node (*name, body, **kwargs*)

partially update the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_node(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Node

If the method is called asynchronously, returns the request thread.

patch_node_status (*name, body, **kwargs*)

partially update status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_node_status(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Node

If the method is called asynchronously, returns the request thread.

patch_node_status_with_http_info (*name, body, **kwargs*)

partially update status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_node_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Node

If the method is called asynchronously, returns the request thread.

patch_node_with_http_info (*name, body, **kwargs*)

partially update the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_node_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

patch_persistent_volume (*name, body, **kwargs*)

partially update the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_persistent_volume(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

patch_persistent_volume_status (*name, body, **kwargs*)

partially update status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_persistent_volume_status(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

patch_persistent_volume_status_with_http_info (*name, body, **kwargs*)

partially update status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_persistent_volume_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

patch_persistent_volume_with_http_info (*name, body, **kwargs*)

partially update the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_persistent_volume_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_pod (*name, namespace, **kwargs*)

proxy DELETE requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_delete_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_pod_with_http_info (*name, namespace, **kwargs*)

proxy DELETE requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.proxy_delete_namespaced_pod_with_http_info(name, namespace, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_pod_with_path (*name, namespace, path, **kwargs*)

proxy DELETE requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_delete_namespaced_pod_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_pod_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy DELETE requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_delete_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_service (*name, namespace, **kwargs*)

proxy DELETE requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_delete_namespaced_service(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_service_with_http_info (*name, namespace, **kwargs*)

proxy DELETE requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_delete_namespaced_service_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_service_with_path (*name, namespace, path, **kwargs*)

proxy DELETE requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_delete_namespaced_service_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required)

:return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy DELETE requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_delete_namespaced_service_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_node (*name, **kwargs*)

proxy DELETE requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_delete_node(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_node_with_http_info (*name, **kwargs*)

proxy DELETE requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_delete_node_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_node_with_path (*name, path, **kwargs*)

proxy DELETE requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_delete_node_with_path(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_delete_node_with_path_with_http_info (*name, path, **kwargs*)

proxy DELETE requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_delete_node_with_path_with_http_info(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_get_namespaced_pod (*name, namespace, **kwargs*)

proxy GET requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_get_namespaced_pod(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_pod_with_http_info` (*name, namespace, **kwargs*)

proxy GET requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_pod_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_pod_with_path` (*name, namespace, path, **kwargs*)

proxy GET requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_pod_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_pod_with_path_with_http_info` (*name, namespace, path, **kwargs*)

proxy GET requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_service` (*name, namespace, **kwargs*)

proxy GET requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_service(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_service_with_http_info` (*name, namespace, **kwargs*)

proxy GET requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_get_namespaced_service_with_path` (*name, namespace, path, **kwargs*)

proxy GET requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_service_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required)
:return: str

If the method is called asynchronously, returns the request thread.

proxy_get_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy GET requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_namespaced_service_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required)
:return: str

If the method is called asynchronously, returns the request thread.

proxy_get_node (*name, **kwargs*)

proxy GET requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_node(name, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_get_node_with_http_info (*name, **kwargs*)

proxy GET requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_node_with_http_info(name, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_get_node_with_path (*name, path, **kwargs*)

proxy GET requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_node_with_path(name, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_get_node_with_path_with_http_info (*name, path, **kwargs*)

proxy GET requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_get_node_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_pod (*name, namespace, **kwargs*)

proxy HEAD requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_pod_with_http_info (*name, namespace, **kwargs*)

proxy HEAD requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_head_namespaced_pod_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_pod_with_path (*name, namespace, path, **kwargs*)

proxy HEAD requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_head_namespaced_pod_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_pod_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy HEAD requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_head_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_service (*name, namespace, **kwargs*)

proxy HEAD requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_head_namespaced_service(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_service_with_http_info (*name, namespace, **kwargs*)

proxy HEAD requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.proxy_head_namespaced_service_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_service_with_path (*name, namespace, path, **kwargs*)

proxy HEAD requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_namespaced_service_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required)
:return: `str`

If the method is called asynchronously, returns the request thread.

proxy_head_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy HEAD requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_namespaced_service_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required)
:return: `str`

If the method is called asynchronously, returns the request thread.

proxy_head_node (*name, **kwargs*)

proxy HEAD requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_node(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_head_node_with_http_info (*name, **kwargs*)

proxy HEAD requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_node_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_head_node_with_path (*name, path, **kwargs*)

proxy HEAD requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_node_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_head_node_with_path_with_http_info (*name, path, **kwargs*)

proxy HEAD requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_head_node_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_pod (*name, namespace, **kwargs*)

proxy OPTIONS requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_pod_with_http_info (*name, namespace, **kwargs*)

proxy OPTIONS requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_pod_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_pod_with_path (*name, namespace, path, **kwargs*)

proxy OPTIONS requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_pod_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_pod_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy OPTIONS requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_service (*name, namespace, **kwargs*)

proxy OPTIONS requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_service(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_service_with_http_info (*name, namespace, **kwargs*)

proxy OPTIONS requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_options_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_service_with_path (*name, namespace, path, **kwargs*)

proxy OPTIONS requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_namespaced_service_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy OPTIONS requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_namespaced_service_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_node (*name, **kwargs*)

proxy OPTIONS requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_node(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_node_with_http_info (*name, **kwargs*)

proxy OPTIONS requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_node_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_node_with_path (*name, path, **kwargs*)

proxy OPTIONS requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_node_with_path(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_options_node_with_path_with_http_info (*name, path, **kwargs*)

proxy OPTIONS requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_options_node_with_path_with_http_info(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_pod (*name, namespace, **kwargs*)

proxy PATCH requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_patch_namespaced_pod(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_pod_with_http_info (*name, namespace, **kwargs*)

proxy PATCH requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_patch_namespaced_pod_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_pod_with_path (*name, namespace, path, **kwargs*)

proxy PATCH requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_patch_namespaced_pod_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_pod_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy PATCH requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_patch_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_service (*name, namespace, **kwargs*)

proxy PATCH requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_patch_namespaced_service(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_service_with_http_info (*name, namespace, **kwargs*)

proxy PATCH requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_service_with_path (*name, namespace, path, **kwargs*)

proxy PATCH requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_namespaced_service_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_patch_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy PATCH requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_namespaced_service_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_patch_node (*name, **kwargs*)

proxy PATCH requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_node(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_patch_node_with_http_info (*name, **kwargs*)

proxy PATCH requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_node_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_patch_node_with_path (*name, path, **kwargs*)

proxy PATCH requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_node_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_patch_node_with_path_with_http_info` (*name, path, **kwargs*)

proxy PATCH requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_patch_node_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_post_namespaced_pod` (*name, namespace, **kwargs*)

proxy POST requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_post_namespaced_pod_with_http_info` (*name, namespace, **kwargs*)

proxy POST requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_namespaced_pod_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_post_namespaced_pod_with_path` (*name, namespace, path, **kwargs*)

proxy POST requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_namespaced_pod_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_post_namespaced_pod_with_path_with_http_info` (*name, namespace, path, **kwargs*)

proxy POST requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_post_namespaced_service` (*name, namespace, **kwargs*)

proxy POST requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_namespaced_service(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_namespaced_service_with_http_info (*name, namespace, **kwargs*)

proxy POST requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_namespaced_service_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_namespaced_service_with_path (*name, namespace, path, **kwargs*)

proxy POST requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_namespaced_service_with_path(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_namespaced_service_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy POST requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_namespaced_service_with_path_with_http_info(name, namespace, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_node (*name, **kwargs*)

proxy POST requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_node(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_node_with_http_info (*name, **kwargs*)

proxy POST requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_node_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_node_with_path (*name, path, **kwargs*)

proxy POST requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.proxy_post_node_with_path(name, path, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_post_node_with_path_with_http_info (*name, path, **kwargs*)

proxy POST requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_post_node_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Node (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_put_namespaced_pod (*name, namespace, **kwargs*)

proxy PUT requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_put_namespaced_pod_with_http_info (*name, namespace, **kwargs*)

proxy PUT requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_pod_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_put_namespaced_pod_with_path (*name, namespace, path, **kwargs*)

proxy PUT requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_pod_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_put_namespaced_pod_with_path_with_http_info (*name, namespace, path, **kwargs*)

proxy PUT requests to Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_pod_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Pod (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str path: path to the resource (required) :return: str

If the method is called asynchronously, returns the request thread.

proxy_put_namespaced_service (*name, namespace, **kwargs*)

proxy PUT requests to Service This method makes a synchronous HTTP request by de-

fault. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_service(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_put_namespaced_service_with_http_info` (*name, namespace, **kwargs*)

proxy PUT requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_put_namespaced_service_with_path` (*name, namespace, path, **kwargs*)

proxy PUT requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_service_with_path(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_put_namespaced_service_with_path_with_http_info` (*name, namespace, path, **kwargs*)

proxy PUT requests to Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_namespaced_service_with_path_with_http_info(name, namespace, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_put_node` (*name, **kwargs*)

proxy PUT requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_node(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

`proxy_put_node_with_http_info` (*name, **kwargs*)

proxy PUT requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_node_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_put_node_with_path (*name, path, **kwargs*)

proxy PUT requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_node_with_path(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

proxy_put_node_with_path_with_http_info (*name, path, **kwargs*)

proxy PUT requests to Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.proxy_put_node_with_path_with_http_info(name, path, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str path`: path to the resource (required) :return: `str`

If the method is called asynchronously, returns the request thread.

read_component_status (*name, **kwargs*)

read the specified ComponentStatus This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_component_status(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ComponentStatus (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ComponentStatus`

If the method is called asynchronously, returns the request thread.

read_component_status_with_http_info (*name, **kwargs*)

read the specified ComponentStatus This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_component_status_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ComponentStatus (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ComponentStatus`

If the method is called asynchronously, returns the request thread.

read_namespace (*name, **kwargs*)

read the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespace(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

read_namespace_status (*name, **kwargs*)

read status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespace_status(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

read_namespace_status_with_http_info (*name, **kwargs*)

read status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespace_status_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

read_namespace_with_http_info (*name, **kwargs*)

read the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespace_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

read_namespaced_config_map (*name, namespace, **kwargs*)

read the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_config_map(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ConfigMap (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

read_namespaced_config_map_with_http_info (*name, namespace, **kwargs*)

read the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_config_map_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ConfigMap (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

read_namespaced_endpoints (*name, namespace, **kwargs*)

read the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_endpoints(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Endpoints (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

read_namespaced_endpoints_with_http_info (*name, namespace, **kwargs*)

read the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_endpoints_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Endpoints (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

read_namespaced_event (*name, namespace, **kwargs*)

read the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_event(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Event (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Event`

If the method is called asynchronously, returns the request thread.

read_namespaced_event_with_http_info (*name, namespace, **kwargs*)

read the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_event_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Event (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Event`

If the method is called asynchronously, returns the request thread.

read_namespaced_limit_range (*name, namespace, **kwargs*)

read the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_limit_range(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the LimitRange (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1LimitRange`

If the method is called asynchronously, returns the request thread.

read_namespaced_limit_range_with_http_info (*name, namespace, **kwargs*)

read the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_limit_range_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the LimitRange (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output

is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

read_namespaced_persistent_volume_claim (*name, namespace, **kwargs*)

read the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_persistent_volume_claim(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

read_namespaced_persistent_volume_claim_status (*name, namespace, **kwargs*)

read status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_persistent_volume_claim_status(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

read_namespaced_persistent_volume_claim_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_persistent_volume_claim_status_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

read_namespaced_persistent_volume_claim_with_http_info (*name, namespace, **kwargs*)

read the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_persistent_volume_claim_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

read_namespaced_pod (*name, namespace, **kwargs*)

read the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

read_namespaced_pod_log (*name, namespace, **kwargs*)

read log of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_log(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container for which to stream logs. Defaults to only container if there is one container in the pod. :param `bool follow`: Follow the log stream of the pod. Defaults to false. :param `int limit_bytes`: If set, the number of bytes to read from the server before terminating the log output. This may not display a complete final line of logging, and may return slightly more or slightly less than the specified limit. :param `str pretty`: If 'true', then the output is pretty printed. :param `bool previous`: Return previous terminated container logs. Defaults to false. :param `int since_seconds`: A relative time in seconds before the current time from which to show logs. If this value precedes the time a pod was started, only logs since the pod start will be returned. If this value is in the future, no logs will be returned. Only one of `sinceSeconds` or `sinceTime` may be specified. :param `int tail_lines`: If set, the number of lines from the end of the logs to show. If not specified, logs are shown from the creation of the container or `sinceSeconds` or `sinceTime` :param `bool timestamps`: If true, add an RFC3339 or RFC3339Nano timestamp at the beginning of every line of log output. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

read_namespaced_pod_log_with_http_info (*name, namespace, **kwargs*)

read log of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_log_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str container`: The container for which to stream logs. Defaults to only container if there is one container in the pod. :param `bool follow`: Follow the log stream of the pod. Defaults to false. :param `int limit_bytes`: If set, the number of bytes to read from the server before terminating the log output. This may not display a complete final line of logging, and may return slightly more or slightly less than the specified limit. :param `str pretty`: If 'true', then the output is pretty printed. :param `bool previous`: Return previous terminated container logs. Defaults to false. :param `int since_seconds`: A relative time in seconds before the current time from which to show logs. If this value precedes the time a pod was started, only logs since the pod start will be returned. If this value is in the future, no logs will be returned. Only one of `sinceSeconds` or `sinceTime` may be specified. :param `int tail_lines`: If set, the number of lines from the end of the logs to show. If not specified, logs are shown from the creation of the container or `sinceSeconds` or `sinceTime` :param `bool timestamps`: If true, add an RFC3339 or RFC3339Nano timestamp at the beginning of every line of log output. Defaults to false. :return: `str`

If the method is called asynchronously, returns the request thread.

read_namespaced_pod_status (*name, namespace, **kwargs*)

read status of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`read_namespaced_pod_status_with_http_info` (*name, namespace, **kwargs*)

read status of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

`read_namespaced_pod_template` (*name, namespace, **kwargs*)

read the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_template(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`read_namespaced_pod_template_with_http_info` (*name, namespace, **kwargs*)

read the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_template_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

`read_namespaced_pod_with_http_info` (*name, namespace, **kwargs*)

read the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller (*name, namespace, **kwargs*)

read the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_scale (*name, namespace, **kwargs*)

read scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller_scale(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller_scale_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Scale`

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_status (*name, namespace, **kwargs*)

read status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the

output is pretty printed. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_with_http_info (*name*, *namespace*, ***kwargs*)

read the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_replication_controller_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: V1ReplicationController

If the method is called asynchronously, returns the request thread.

read_namespaced_resource_quota (*name*, *namespace*, ***kwargs*)

read the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_resource_quota(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

read_namespaced_resource_quota_status (*name*, *namespace*, ***kwargs*)

read status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_resource_quota_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

read_namespaced_resource_quota_status_with_http_info (*name*, *namespace*, ***kwargs*)

read status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_resource_quota_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

read_namespaced_resource_quota_with_http_info (*name*, *namespace*, ***kwargs*)

read the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_resource_quota_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ResourceQuota (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1ResourceQuota

If the method is called asynchronously, returns the request thread.

read_namespaced_secret (*name, namespace, **kwargs*)

read the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_secret(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Secret (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1Secret

If the method is called asynchronously, returns the request thread.

read_namespaced_secret_with_http_info (*name, namespace, **kwargs*)

read the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_secret_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Secret (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1Secret

If the method is called asynchronously, returns the request thread.

read_namespaced_service (*name, namespace, **kwargs*)

read the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_service(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Service (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1Service

If the method is called asynchronously, returns the request thread.

read_namespaced_service_account (*name, namespace, **kwargs*)

read the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_service_account(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ServiceAccount (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1ServiceAccount

If the method is called asynchronously, returns the request thread.

read_namespaced_service_account_with_http_info (*name, namespace, **kwargs*)

read the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_service_account_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

read_namespaced_service_status (*name, namespace, **kwargs*)

read status of the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_service_status(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

read_namespaced_service_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_service_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

read_namespaced_service_with_http_info (*name, namespace, **kwargs*)

read the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_service_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

read_node (*name, **kwargs*)

read the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_node(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

read_node_status (*name*, ***kwargs*)

read status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_node_status(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

read_node_status_with_http_info (*name*, ***kwargs*)

read status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_node_status_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

read_node_with_http_info (*name*, ***kwargs*)

read the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_node_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

read_persistent_volume (*name*, ***kwargs*)

read the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_persistent_volume(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

read_persistent_volume_status (*name*, ***kwargs*)

read status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_persistent_volume_status(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

read_persistent_volume_status_with_http_info (*name*, ***kwargs*)

read status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_persistent_volume_status_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

read_persistent_volume_with_http_info (*name*, ***kwargs*)

read the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_persistent_volume_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

replace_namespace (*name*, *body*, ***kwargs*)

replace the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespace_finalize (*name*, *body*, ***kwargs*)

replace finalize of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace_finalize(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespace_finalize_with_http_info (*name*, *body*, ***kwargs*)

replace finalize of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace_finalize_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespace_status (*name*, *body*, ***kwargs*)

replace status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace_status(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespace_status_with_http_info (*name*, *body*, ***kwargs*)

replace status of the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespace_with_http_info (*name, body, **kwargs*)

replace the specified Namespace This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespace_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Namespace (required) :param `V1Namespace body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Namespace`

If the method is called asynchronously, returns the request thread.

replace_namespaced_config_map (*name, namespace, body, **kwargs*)

replace the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_config_map(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ConfigMap (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ConfigMap body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

replace_namespaced_config_map_with_http_info (*name, namespace, body, **kwargs*)

replace the specified ConfigMap This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_config_map_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ConfigMap (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ConfigMap body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ConfigMap`

If the method is called asynchronously, returns the request thread.

replace_namespaced_endpoints (*name, namespace, body, **kwargs*)

replace the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_endpoints(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Endpoints (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Endpoints body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

replace_namespaced_endpoints_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Endpoints This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_endpoints_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Endpoints (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Endpoints body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Endpoints`

If the method is called asynchronously, returns the request thread.

replace_namespaced_event (*name, namespace, body, **kwargs*)

replace the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_event(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Event (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Event body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

replace_namespaced_event_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Event This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_event_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Event (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1Event body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1Event

If the method is called asynchronously, returns the request thread.

replace_namespaced_limit_range (*name, namespace, body, **kwargs*)

replace the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_limit_range(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1LimitRange body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

replace_namespaced_limit_range_with_http_info (*name, namespace, body, **kwargs*)

replace the specified LimitRange This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_limit_range_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the LimitRange (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1LimitRange body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1LimitRange

If the method is called asynchronously, returns the request thread.

replace_namespaced_persistent_volume_claim (*name, namespace, body, **kwargs*)

replace the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_persistent_volume_claim(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1PersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1PersistentVolumeClaim

If the method is called asynchronously, returns the request thread.

replace_namespaced_persistent_volume_claim_status (*name, namespace, body, **kwargs*)

replace status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_persistent_volume_claim_status(name, namespace, body, async=True) >>> result = thread.get()

```
:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str
namespace: object name and auth scope, such as for teams and projects (required) :param
VIPersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: VIPersistentVolumeClaim
```

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_persistent_volume_claim_status_with_http_info (name,
                                                                    names-
                                                                    pace,
                                                                    body,
                                                                    **kwargs)
```

replace status of the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_persistent_volume_claim_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

```
:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str
namespace: object name and auth scope, such as for teams and projects (required) :param
VIPersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: VIPersistentVolumeClaim
```

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_persistent_volume_claim_with_http_info (name, namespace,
                                                            body, **kwargs)
```

replace the specified PersistentVolumeClaim This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_persistent_volume_claim_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

```
:param async bool :param str name: name of the PersistentVolumeClaim (required) :param str
namespace: object name and auth scope, such as for teams and projects (required) :param
VIPersistentVolumeClaim body: (required) :param str pretty: If 'true', then the output is pretty printed.
:return: VIPersistentVolumeClaim
```

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_pod (name, namespace, body, **kwargs)
```

replace the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod(name, namespace, body, async=True)` >>> `result = thread.get()`

```
:param async bool :param str name: name of the Pod (required) :param str namespace: object name and
auth scope, such as for teams and projects (required) :param VIPod body: (required) :param str pretty: If
'true', then the output is pretty printed. :return: VIPod
```

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_pod_status (name, namespace, body, **kwargs)
```

replace status of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_status(name, namespace, body, async=True)` >>> `result = thread.get()`

```
:param async bool :param str name: name of the Pod (required) :param str namespace: object name and
auth scope, such as for teams and projects (required) :param VIPod body: (required) :param str pretty: If
'true', then the output is pretty printed. :return: VIPod
```

If the method is called asynchronously, returns the request thread.

replace_namespaced_pod_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Pod body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

replace_namespaced_pod_template (*name, namespace, body, **kwargs*)

replace the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_template(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1PodTemplate body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

replace_namespaced_pod_template_with_http_info (*name, namespace, body, **kwargs*)

replace the specified PodTemplate This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_template_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodTemplate (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1PodTemplate body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PodTemplate`

If the method is called asynchronously, returns the request thread.

replace_namespaced_pod_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Pod This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Pod (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Pod body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Pod`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replication_controller (*name, namespace, body, **kwargs*)

replace the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.


```
replace_namespaced_replication_controller_scale (name, namespace, body,  
                                                **kwargs)
```

replace scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Scale`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_replication_controller_scale_with_http_info (name,  
                                                                names-  
                                                                pace,  
                                                                body,  
                                                                **kwargs)
```

replace scale of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Scale`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_replication_controller_status (name, namespace, body,  
                                                **kwargs)
```

replace status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_replication_controller_status_with_http_info (name,  
                                                                names-  
                                                                pace,  
                                                                body,  
                                                                **kwargs)
```

replace status of the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ReplicationController`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replication_controller_with_http_info (*name, namespace, body, **kwargs*)

replace the specified ReplicationController This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replication_controller_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicationController (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `VIReplicationController body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `VIReplicationController`

If the method is called asynchronously, returns the request thread.

replace_namespaced_resource_quota (*name, namespace, body, **kwargs*)

replace the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_resource_quota(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `VIResourceQuota body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `VIResourceQuota`

If the method is called asynchronously, returns the request thread.

replace_namespaced_resource_quota_status (*name, namespace, body, **kwargs*)

replace status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_resource_quota_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `VIResourceQuota body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `VIResourceQuota`

If the method is called asynchronously, returns the request thread.

replace_namespaced_resource_quota_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_resource_quota_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `VIResourceQuota body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `VIResourceQuota`

If the method is called asynchronously, returns the request thread.

replace_namespaced_resource_quota_with_http_info (*name, namespace, body, **kwargs*)

replace the specified ResourceQuota This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_resource_quota_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ResourceQuota (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `VIResourceQuota body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `VIResourceQuota`

If the method is called asynchronously, returns the request thread.

replace_namespaced_secret (*name, namespace, body, **kwargs*)

replace the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_secret(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Secret (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Secret body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Secret`

If the method is called asynchronously, returns the request thread.

replace_namespaced_secret_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Secret This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_secret_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Secret (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Secret body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Secret`

If the method is called asynchronously, returns the request thread.

replace_namespaced_service (*name, namespace, body, **kwargs*)

replace the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Service body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

replace_namespaced_service_account (*name, namespace, body, **kwargs*)

replace the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service_account(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ServiceAccount body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

replace_namespaced_service_account_with_http_info (*name, namespace, body, **kwargs*)

replace the specified ServiceAccount This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service_account_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ServiceAccount (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1ServiceAccount body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1ServiceAccount`

If the method is called asynchronously, returns the request thread.

replace_namespaced_service_status (*name, namespace, body, **kwargs*)

replace status of the specified Service This method makes a synchronous HTTP request by

default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Service body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_service_status_with_http_info` (*name, namespace, body, **kwargs*)

replace status of the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Service body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_service_with_http_info` (*name, namespace, body, **kwargs*)

replace the specified Service This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_service_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Service (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1Service body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Service`

If the method is called asynchronously, returns the request thread.

`replace_node` (*name, body, **kwargs*)

replace the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_node(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

`replace_node_status` (*name, body, **kwargs*)

replace status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_node_status(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

`replace_node_status_with_http_info` (*name, body, **kwargs*)

replace status of the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_node_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

replace_node_with_http_info (*name, body, **kwargs*)

replace the specified Node This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_node_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Node (required) :param `V1Node body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1Node`

If the method is called asynchronously, returns the request thread.

replace_persistent_volume (*name, body, **kwargs*)

replace the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_persistent_volume(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

replace_persistent_volume_status (*name, body, **kwargs*)

replace status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_persistent_volume_status(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

replace_persistent_volume_status_with_http_info (*name, body, **kwargs*)

replace status of the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_persistent_volume_status_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

replace_persistent_volume_with_http_info (*name, body, **kwargs*)

replace the specified PersistentVolume This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_persistent_volume_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PersistentVolume (required) :param `V1PersistentVolume body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1PersistentVolume`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.extensions_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.extensions_api.ExtensionsApi` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.extensions_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api` (*api_client=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_daemon_set (*namespace, body, **kwargs*)

create a DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_daemon_set(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

create_namespaced_daemon_set_with_http_info (*namespace, body, **kwargs*)

create a DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_daemon_set_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment (*namespace, body, **kwargs*)

create a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_rollback (*name, namespace, body, **kwargs*)

create rollback of a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_rollback(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DeploymentRollback (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1DeploymentRollback body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1DeploymentRollback`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_rollback_with_http_info (*name, namespace, body, **kwargs*)

create rollback of a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_rollback_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DeploymentRollback (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1DeploymentRollback body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1DeploymentRollback`

If the method is called asynchronously, returns the request thread.

create_namespaced_deployment_with_http_info (*namespace, body, **kwargs*)

create a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_deployment_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

create_namespaced_ingress (*namespace, body, **kwargs*)

create an Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_ingress(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

create_namespaced_ingress_with_http_info (*namespace, body, **kwargs*)

create an Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_ingress_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

create_namespaced_network_policy (*namespace, body, **kwargs*)

create a NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_network_policy(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1NetworkPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

create_namespaced_network_policy_with_http_info (*namespace, body, **kwargs*)

create a NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_network_policy_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1NetworkPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

create_namespaced_replica_set (*namespace, body, **kwargs*)

create a ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_replica_set(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ReplicaSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ReplicaSet`

If the method is called asynchronously, returns the request thread.

create_namespaced_replica_set_with_http_info (*namespace, body, **kwargs*)

create a ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_replica_set_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ReplicaSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ReplicaSet`

If the method is called asynchronously, returns the request thread.

create_pod_security_policy (*body, **kwargs*)

create a PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_pod_security_policy(body, async=True)` >>> `result = thread.get()`

:param async bool :param V1beta1PodSecurityPolicy body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodSecurityPolicy

If the method is called asynchronously, returns the request thread.

create_pod_security_policy_with_http_info (*body*, ***kwargs*)

create a PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_pod_security_policy_with_http_info(body, async=True)` >>> `result = thread.get()`

:param async bool :param V1beta1PodSecurityPolicy body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodSecurityPolicy

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_daemon_set (*namespace*, ***kwargs*)

delete collection of DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_daemon_set(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_daemon_set_with_http_info (*namespace*, ***kwargs*)

delete collection of DaemonSet This method makes a synchronous HTTP request by de-

fault. To make an asynchronous HTTP request, please pass `async=True` `>>> thread = api.delete_collection_namespaced_daemon_set_with_http_info(namespace, async=True) >>> result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given `rv`. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_collection_namespaced_deployment` (`namespace`, `**kwargs`)

delete collection of Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` `>>> thread = api.delete_collection_namespaced_deployment(namespace, async=True) >>> result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by

their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_deployment_with_http_info (*namespace*,
***kwargs*)

delete collection of Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_deployment_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version

of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_ingress (namespace, **kwargs)

delete collection of Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_ingress(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_ingress_with_http_info (namespace, **kwargs)

delete collection of Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_ingress_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue

option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_network_policy (*namespace*, ***kwargs*)

delete collection of NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_network_policy(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available.

Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_network_policy_with_http_info (*namespace*,
***kwargs*)

delete collection of NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.delete_collection_namespaced_network_policy_with_http_info(namespace, async=True) >>> result = thread.get()

:param bool async: :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify

resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_replica_set (*namespace*, ***kwargs*)

delete collection of ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_replica_set(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_replica_set_with_http_info (*namespace*, ***kwargs*)

delete collection of ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_replica_set_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or

a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_pod_security_policy (**kwargs)

delete collection of PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_pod_security_policy(async=True)` >>> `result = thread.get()`

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes

referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_pod_security_policy_with_http_info (**kwargs)

delete collection of PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_pod_security_policy_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_daemon_set (name, namespace, body, **kwargs)

delete a DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_daemon_set(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_daemon_set_with_http_info (*name, namespace, body, **kwargs*)

delete a DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_daemon_set_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_deployment (*name, namespace, body, **kwargs*)

delete a Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_deployment(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_deployment_with_http_info (*name, namespace, body, **kwargs*)

delete a Deployment This method makes a synchronous HTTP request by default.

To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_deployment_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool `orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the `metadata.finalizers` and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_namespaced_ingress` (*name, namespace, body, **kwargs*)

delete an Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_ingress(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool `orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the `metadata.finalizers` and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_namespaced_ingress_with_http_info` (*name, namespace, body, **kwargs*)

delete an Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_ingress_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool `orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the `metadata.finalizers` and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_network_policy (*name, namespace, body, **kwargs*)

delete a NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_network_policy(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_network_policy_with_http_info (*name, namespace, body, **kwargs*)

delete a NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_network_policy_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_namespaced_replica_set (*name, namespace, body, **kwargs*)

delete a ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_replica_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicaSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`:

Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

`delete_namespaced_replica_set_with_http_info` (*name, namespace, body, **kwargs*)

delete a ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_replica_set_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicaSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

`delete_pod_security_policy` (*name, body, **kwargs*)

delete a PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_pod_security_policy(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

`delete_pod_security_policy_with_http_info` (*name, body, **kwargs*)

delete a PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_pod_security_policy_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will

be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_daemon_set_for_all_namespaces (**kwargs)

list or watch objects of kind DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_daemon_set_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1DaemonSetList

If the method is called asynchronously, returns the request thread.

list_daemon_set_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_daemon_set_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1DaemonSetList

If the method is called asynchronously, returns the request thread.

list_deployment_for_all_namespaces (**kwargs)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_deployment_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value

returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: ExtensionsV1beta1DeploymentList

If the method is called asynchronously, returns the request thread.

list_deployment_for_all_namespaces_with_http_info (kwargs)**

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_deployment_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list

result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: ExtensionsV1beta1DeploymentList

If the method is called asynchronously, returns the request thread.

list_ingress_for_all_namespaces (**kwargs)

list or watch objects of kind Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_ingress_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1IngressList

If the method is called asynchronously, returns the request thread.

list_ingress_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_ingress_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous

query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is `true`. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If `true`, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is `true`. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `pretty`: If `'true'`, then the output is pretty printed. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1beta1IngressList`

If the method is called asynchronously, returns the request thread.

list_namespaced_daemon_set (*namespace*, ***kwargs*)

list or watch objects of kind `DaemonSet` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_daemon_set(namespace, async=True)` >>> `result = thread.get()`

:param bool `async` :param str `namespace`: object name and auth scope, such as for teams and projects (required) :param str `pretty`: If `'true'`, then the output is pretty printed. :param str `_continue`: The `continue` option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the `continue` value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is `true`. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If `true`, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is

specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1DaemonSetList

If the method is called asynchronously, returns the request thread.

list_namespaced_daemon_set_with_http_info(namespace, **kwargs)

list or watch objects of kind DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_daemon_set_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1DaemonSetList

If the method is called asynchronously, returns the request thread.

list_namespaced_deployment (*namespace*, ***kwargs*)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_deployment(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `ExtensionsV1beta1DeploymentList`

If the method is called asynchronously, returns the request thread.

list_namespaced_deployment_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_deployment_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are

included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: ExtensionsV1beta1DeploymentList

If the method is called asynchronously, returns the request thread.

list_namespaced_ingress (namespace, **kwargs)

list or watch objects of kind Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_ingress(namespace, async=True)` >>> `result = thread.get()`

:param bool async: :param str namespace: object name and auth scope, such as for teams and projects (required) :param bool pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version

of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1IngressList

If the method is called asynchronously, returns the request thread.

list_namespaced_ingress_with_http_info (namespace, **kwargs)

list or watch objects of kind Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_ingress_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1IngressList

If the method is called asynchronously, returns the request thread.

list_namespaced_network_policy (namespace, **kwargs)

list or watch objects of kind NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_network_policy(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined,

clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1NetworkPolicyList

If the method is called asynchronously, returns the request thread.

list_namespaced_network_policy_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_network_policy_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is

specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1NetworkPolicyList

If the method is called asynchronously, returns the request thread.

list_namespaced_replica_set (namespace, **kwargs)

list or watch objects of kind ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_replica_set(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ReplicaSetList

If the method is called asynchronously, returns the request thread.

list_namespaced_replica_set_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_namespaced_replica_set_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given `rv`. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1beta1ReplicaSetList`

If the method is called asynchronously, returns the request thread.

list_network_policy_for_all_namespaces (***kwargs*)

list or watch objects of kind NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_network_policy_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum

number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1NetworkPolicyList

If the method is called asynchronously, returns the request thread.

list_network_policy_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_network_policy_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what

we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1NetworkPolicyList

If the method is called asynchronously, returns the request thread.

list_pod_security_policy (**kwargs)

list or watch objects of kind PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_pod_security_policy(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1PodSecurityPolicyList

If the method is called asynchronously, returns the request thread.

list_pod_security_policy_with_http_info (**kwargs)

list or watch objects of kind PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_pod_security_policy_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating

the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1PodSecurityPolicyList

If the method is called asynchronously, returns the request thread.

list_replica_set_for_all_namespaces (**kwargs)

list or watch objects of kind ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_replica_set_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If

objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ReplicaSetList

If the method is called asynchronously, returns the request thread.

list_replica_set_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_replica_set_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1ReplicaSetList

If the method is called asynchronously, returns the request thread.

patch_namespaced_daemon_set (name, namespace, body, **kwargs)

partially update the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_daemon_set(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name

and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_daemon_set_status (*name, namespace, body, **kwargs*)

partially update status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_daemon_set_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_daemon_set_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_daemon_set_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_daemon_set_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_daemon_set_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment (*name, namespace, body, **kwargs*)

partially update the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_scale_with_http_info (*name*, *namespace*, *body*,
***kwargs*)

partially update scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_status (*name*, *namespace*, *body*, ***kwargs*)

partially update status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_status_with_http_info (*name*, *namespace*, *body*,
***kwargs*)

partially update status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

patch_namespaced_deployment_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

partially update the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_deployment_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

patch_namespaced_ingress (*name*, *namespace*, *body*, ***kwargs*)

partially update the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_ingress(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

patch_namespaced_ingress_status (*name*, *namespace*, *body*, ***kwargs*)

partially update status of the specified Ingress This method makes a synchronous HTTP request

by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_ingress_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_ingress_status_with_http_info` (*name, namespace, body,*
***kwargs*)

partially update status of the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_ingress_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_ingress_with_http_info` (*name, namespace, body, **kwargs*)

partially update the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_ingress_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_network_policy` (*name, namespace, body, **kwargs*)

partially update the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_network_policy(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_network_policy_with_http_info` (*name, namespace, body,*
***kwargs*)

partially update the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_network_policy_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

`patch_namespaced_replica_set` (*name, namespace, body, **kwargs*)

partially update the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replica_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_replica_set_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_replica_set_scale(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_replica_set_scale_with_http_info (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_replica_set_scale_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

patch_namespaced_replica_set_status (*name, namespace, body, **kwargs*)

partially update status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_replica_set_status(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_replica_set_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_replica_set_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_replica_set_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_replica_set_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty:

If ‘true’, then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_dummy_scale (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_dummy_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If ‘true’, then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

patch_namespaced_replication_controller_dummy_scale_with_http_info (*name, namespace, body, **kwargs*)

partially update scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_namespaced_replication_controller_dummy_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `object body`: (required) :param `str pretty`: If ‘true’, then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

patch_pod_security_policy (*name, body, **kwargs*)

partially update the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_pod_security_policy(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `object body`: (required) :param `str pretty`: If ‘true’, then the output is pretty printed. :return: `V1beta1PodSecurityPolicy`

If the method is called asynchronously, returns the request thread.

patch_pod_security_policy_with_http_info (*name, body, **kwargs*)

partially update the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_pod_security_policy_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `object body`: (required) :param `str pretty`: If ‘true’, then the output is pretty printed. :return: `V1beta1PodSecurityPolicy`

If the method is called asynchronously, returns the request thread.

read_namespaced_daemon_set (*name, namespace, **kwargs*)

read the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_daemon_set(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DaemonSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If ‘true’, then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific

fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

read_namespaced_daemon_set_status (*name, namespace, **kwargs*)

read status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_daemon_set_status(name, namespace, async=True)` >>> `result = thread.get()`

:param bool async :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

read_namespaced_daemon_set_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_daemon_set_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param bool async :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

read_namespaced_daemon_set_with_http_info (*name, namespace, **kwargs*)

read the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_daemon_set_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param bool async :param str name: name of the DaemonSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1DaemonSet

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment (*name, namespace, **kwargs*)

read the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment(name, namespace, async=True)` >>> `result = thread.get()`

:param bool async :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: ExtensionsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_scale (*name, namespace, **kwargs*)

read scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_deployment_scale(name, namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_deployment_scale_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_status (*name, namespace, **kwargs*)

read status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_deployment_status(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_deployment_status_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

read_namespaced_deployment_with_http_info (*name, namespace, **kwargs*)

read the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_deployment_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Deployment (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: ExtensionsV1beta1Deployment

If the method is called asynchronously, returns the request thread.

read_namespaced_ingress (*name, namespace, **kwargs*)

read the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_ingress(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Ingress (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1Ingress

If the method is called asynchronously, returns the request thread.

read_namespaced_ingress_status (*name, namespace, **kwargs*)

read status of the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_ingress_status(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Ingress (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1Ingress

If the method is called asynchronously, returns the request thread.

read_namespaced_ingress_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_ingress_status_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Ingress (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1Ingress

If the method is called asynchronously, returns the request thread.

read_namespaced_ingress_with_http_info (*name, namespace, **kwargs*)

read the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_ingress_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Ingress (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1Ingress

If the method is called asynchronously, returns the request thread.

read_namespaced_network_policy (*name, namespace, **kwargs*)

read the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_network_policy(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the NetworkPolicy (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1NetworkPolicy

If the method is called asynchronously, returns the request thread.

read_namespaced_network_policy_with_http_info (*name, namespace, **kwargs*)

read the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =

```
api.read_namespaced_network_policy_with_http_info(name, namespace, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the NetworkPolicy (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1NetworkPolicy

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set (*name, namespace, **kwargs*)

read the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replica_set(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set_scale (*name, namespace, **kwargs*)

read scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replica_set_scale(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replica_set_scale_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set_status (*name, namespace, **kwargs*)

read status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replica_set_status(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set_status_with_http_info (*name, namespace, **kwargs*)

read status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =

```
api.read_namespaced_replica_set_status_with_http_info(name, namespace, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

read_namespaced_replica_set_with_http_info (*name, namespace, **kwargs*)

read the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replica_set_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_dummy_scale (*name, namespace, **kwargs*)

read scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replication_controller_dummy_scale(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_namespaced_replication_controller_dummy_scale_with_http_info (*name, namespace, **kwargs*)

read scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_replication_controller_dummy_scale_with_http_info(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

read_pod_security_policy (*name, **kwargs*)

read the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_pod_security_policy(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodSecurityPolicy (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1PodSecurityPolicy

If the method is called asynchronously, returns the request thread.

read_pod_security_policy_with_http_info (*name*, ***kwargs*)

read the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_pod_security_policy_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1PodSecurityPolicy`

If the method is called asynchronously, returns the request thread.

replace_namespaced_daemon_set (*name*, *namespace*, *body*, ***kwargs*)

replace the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_daemon_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DaemonSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_daemon_set_status (*name*, *namespace*, *body*, ***kwargs*)

replace status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_daemon_set_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DaemonSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_daemon_set_status_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

replace status of the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_daemon_set_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DaemonSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_daemon_set_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

replace the specified DaemonSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_daemon_set_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the DaemonSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1DaemonSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1DaemonSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment (*name, namespace, body, **kwargs*)

replace the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_scale (*name, namespace, body, **kwargs*)

replace scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_scale_with_http_info (*name, namespace, body, **kwargs*)

replace scale of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_status (*name, namespace, body, **kwargs*)

replace status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

replace_namespaced_deployment_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Deployment This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_deployment_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Deployment (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Deployment body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Deployment`

If the method is called asynchronously, returns the request thread.

replace_namespaced_ingress (*name, namespace, body, **kwargs*)

replace the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_ingress(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

replace_namespaced_ingress_status (*name, namespace, body, **kwargs*)

replace status of the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_ingress_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

replace_namespaced_ingress_status_with_http_info (*name, namespace, body, **kwargs*)

replace status of the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_ingress_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

replace_namespaced_ingress_with_http_info (*name, namespace, body, **kwargs*)

replace the specified Ingress This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_ingress_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Ingress (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1Ingress body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1Ingress`

If the method is called asynchronously, returns the request thread.

replace_namespaced_network_policy (*name, namespace, body, **kwargs*)

replace the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_network_policy(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1NetworkPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

replace_namespaced_network_policy_with_http_info (*name, namespace, body, **kwargs*)

replace the specified NetworkPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_network_policy_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the NetworkPolicy (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1NetworkPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1NetworkPolicy`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set (*name, namespace, body, **kwargs*)

replace the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replica_set(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ReplicaSet (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1ReplicaSet body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1ReplicaSet`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set_scale (*name, namespace, body, **kwargs*)

replace scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replica_set_scale(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set_scale_with_http_info (*name, namespace, body, **kwargs*)

replace scale of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_replica_set_scale_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Scale (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `ExtensionsV1beta1Scale body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `ExtensionsV1beta1Scale`

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set_status (*name, namespace, body, **kwargs*)

replace status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread`

```
= api.replace_namespaced_replica_set_status(name, namespace, body, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1ReplicaSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set_status_with_http_info (*name, namespace, body,*
***kwargs*)

replace status of the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_replica_set_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1ReplicaSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

replace_namespaced_replica_set_with_http_info (*name, namespace, body, **kwargs*)

replace the specified ReplicaSet This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_replica_set_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ReplicaSet (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1beta1ReplicaSet body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1ReplicaSet

If the method is called asynchronously, returns the request thread.

replace_namespaced_replication_controller_dummy_scale (*name, namespace, body,*
***kwargs*)

replace scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_replication_controller_dummy_scale(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param ExtensionsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_namespaced_replication_controller_dummy_scale_with_http_info (*name,*
names-
pace,
body,
***kwargs*)

replace scale of the specified ReplicationControllerDummy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.replace_namespaced_replication_controller_dummy_scale_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Scale (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param ExtensionsV1beta1Scale body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: ExtensionsV1beta1Scale

If the method is called asynchronously, returns the request thread.

replace_pod_security_policy (*name, body, **kwargs*)
 replace the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_pod_security_policy(name, body, async=True)` >>> `result = thread.get()`
 :param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `V1beta1PodSecurityPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodSecurityPolicy`

If the method is called asynchronously, returns the request thread.

replace_pod_security_policy_with_http_info (*name, body, **kwargs*)
 replace the specified PodSecurityPolicy This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_pod_security_policy_with_http_info(name, body, async=True)` >>> `result = thread.get()`
 :param `async` bool :param `str name`: name of the PodSecurityPolicy (required) :param `V1beta1PodSecurityPolicy body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodSecurityPolicy`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.logs_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.logs_api.LogsApi` (*api_client=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

log_file_handler (*logpath, **kwargs*)
 This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.log_file_handler(logpath, async=True)` >>> `result = thread.get()`
 :param `async` bool :param `str logpath`: path to the log (required) :return: `None`

If the method is called asynchronously, returns the request thread.

log_file_handler_with_http_info (*logpath, **kwargs*)
 This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.log_file_handler_with_http_info(logpath, async=True)` >>> `result = thread.get()`
 :param `async` bool :param `str logpath`: path to the log (required) :return: `None`

If the method is called asynchronously, returns the request thread.

log_file_list_handler (***kwargs*)
 This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.log_file_list_handler(async=True)` >>> `result = thread.get()`
 :param `async` bool :return: `None`

If the method is called asynchronously, returns the request thread.

log_file_list_handler_with_http_info (**kwargs)

This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.log_file_list_handler_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: None

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.policy_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.policy_api.PolicyApi` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (**kwargs)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (**kwargs)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.policy_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

create_namespaced_pod_disruption_budget (*namespace*, *body*, ***kwargs*)

create a PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_disruption_budget(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

create_namespaced_pod_disruption_budget_with_http_info (*namespace*, *body*, ***kwargs*)

create a PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_pod_disruption_budget_with_http_info(namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod_disruption_budget (*namespace*, ***kwargs*)

delete collection of PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_pod_disruption_budget(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what

we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_pod_disruption_budget_with_http_info (*namespace*, ***kwargs*)

delete collection of PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_pod_disruption_budget_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod_disruption_budget (*name*, *namespace*, *body*, ***kwargs*)

delete a PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_pod_disruption_budget(name, namespace, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body:

(required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_pod_disruption_budget_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

delete a PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_pod_disruption_budget_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (***kwargs*)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod_disruption_budget (*namespace*, ***kwargs*)

list or watch objects of kind PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_pod_disruption_budget(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue

option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1PodDisruptionBudgetList

If the method is called asynchronously, returns the request thread.

list_namespaced_pod_disruption_budget_with_http_info (*namespace*, ***kwargs*)

list or watch objects of kind PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_namespaced_pod_disruption_budget_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and

clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: V1beta1PodDisruptionBudgetList

If the method is called asynchronously, returns the request thread.

list_pod_disruption_budget_for_all_namespaces (**kwargs)

list or watch objects of kind PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_pod_disruption_budget_for_all_namespaces(async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `_continue`: The `continue` option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the `continue` value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a `limit` may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `pretty`: If 'true', then the output is pretty printed. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: V1beta1PodDisruptionBudgetList

If the method is called asynchronously, returns the request thread.

list_pod_disruption_budget_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.list_pod_disruption_budget_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1PodDisruptionBudgetList

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod_disruption_budget (name, namespace, body, **kwargs)

partially update the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread = api.patch_namespaced_pod_disruption_budget(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod_disruption_budget_status (name, namespace, body, **kwargs)

partially update status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> thread =

```
api.patch_namespaced_pod_disruption_budget_status(name, namespace, body, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod_disruption_budget_status_with_http_info (*name, namespace, body, **kwargs*)

partially update status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_pod_disruption_budget_status_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

patch_namespaced_pod_disruption_budget_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_pod_disruption_budget_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

read_namespaced_pod_disruption_budget (*name, namespace, **kwargs*)

read the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_pod_disruption_budget(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param bool exact: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param bool export: Should this value be exported. Export strips fields that a user can not specify. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

read_namespaced_pod_disruption_budget_status (*name, namespace, **kwargs*)

read status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_pod_disruption_budget_status(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the PodDisruptionBudget (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1PodDisruptionBudget

If the method is called asynchronously, returns the request thread.

```
read_namespaced_pod_disruption_budget_status_with_http_info(name,
                                                                namespace,
                                                                **kwargs)
```

read status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_disruption_budget_status_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

```
read_namespaced_pod_disruption_budget_with_http_info(name, namespace,
                                                       **kwargs)
```

read the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_pod_disruption_budget_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_pod_disruption_budget(name, namespace, body, **kwargs)
```

replace the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_disruption_budget(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_pod_disruption_budget_status(name, namespace, body,
                                                  **kwargs)
```

replace status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_disruption_budget_status(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

```
replace_namespaced_pod_disruption_budget_status_with_http_info(name,
                                                                namespace,
                                                                body,
                                                                **kwargs)
```

replace status of the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_disruption_budget_status_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed.
:return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

`replace_namespaced_pod_disruption_budget_with_http_info` (*name, namespace, body, **kwargs*)

replace the specified PodDisruptionBudget This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_pod_disruption_budget_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the PodDisruptionBudget (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1beta1PodDisruptionBudget body`: (required) :param `str pretty`: If 'true', then the output is pretty printed.
:return: `V1beta1PodDisruptionBudget`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.rbac_authorization_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.rbac_authorization_api.RbacAuthorizationApi` (*api_client=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.rbac_authorization_v1alpha1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api` (`api_`
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
 Ref: <https://github.com/swagger-api/swagger-codegen>

create_cluster_role (*body*, ****kwargs**)

create a ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_cluster_role(body, async=True)`
 >>> `result = thread.get()`

:param `async` bool :param `V1alpha1ClusterRole` *body*: (required) :param `str` *pretty*: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRole`

If the method is called asynchronously, returns the request thread.

create_cluster_role_binding (*body*, ****kwargs**)

create a ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_cluster_role_binding(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1alpha1ClusterRoleBinding` *body*: (required) :param `str` *pretty*: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRoleBinding`

If the method is called asynchronously, returns the request thread.

create_cluster_role_binding_with_http_info (*body*, ****kwargs**)

create a ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_cluster_role_binding_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1alpha1ClusterRoleBinding` *body*: (required) :param `str` *pretty*: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRoleBinding`

If the method is called asynchronously, returns the request thread.

create_cluster_role_with_http_info (*body*, ****kwargs**)

create a ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_cluster_role_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1alpha1ClusterRole` *body*: (required) :param `str` *pretty*: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRole`

If the method is called asynchronously, returns the request thread.

create_namespaced_role (*namespace*, *body*, ****kwargs**)

create a Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_namespaced_role(namespace, body, async=True)`
 >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1alpha1Role body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1Role

If the method is called asynchronously, returns the request thread.

create_namespaced_role_binding (*namespace, body, **kwargs*)

create a RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_role_binding(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1alpha1RoleBinding body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1RoleBinding

If the method is called asynchronously, returns the request thread.

create_namespaced_role_binding_with_http_info (*namespace, body, **kwargs*)

create a RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_role_binding_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1alpha1RoleBinding body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1RoleBinding

If the method is called asynchronously, returns the request thread.

create_namespaced_role_with_http_info (*namespace, body, **kwargs*)

create a Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.create_namespaced_role_with_http_info(namespace, body, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1alpha1Role body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1Role

If the method is called asynchronously, returns the request thread.

delete_cluster_role (*name, body, **kwargs*)

delete a ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_cluster_role(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRole (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_cluster_role_binding (*name, body, **kwargs*)

delete a ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an

asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_cluster_role_binding(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRoleBinding (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_cluster_role_binding_with_http_info` (*name, body, **kwargs*)

delete a ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_cluster_role_binding_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRoleBinding (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_cluster_role_with_http_info` (*name, body, **kwargs*)

delete a ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_cluster_role_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRole (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_collection_cluster_role` (***kwargs*)

delete collection of ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread =`

```
api.delete_collection_cluster_role(async=True) >>> result = thread.get()
```

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_cluster_role_binding (**kwargs)

delete collection of ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_collection_cluster_role_binding(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and

clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_collection_cluster_role_binding_with_http_info` (kwargs)**

delete collection of ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_cluster_role_binding_with_http_info(async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `pretty`: If 'true', then the output is pretty printed. :param str `_continue`: The `continue` option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the `continue` value from a previous query result with identical query parameters (except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a `limit` may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the `limit` argument and will return all of the available results. If `limit` is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a `limit` - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using `limit` to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_cluster_role_with_http_info (**kwargs)

delete collection of ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_cluster_role_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_role (namespace, **kwargs)

delete collection of Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_role(namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their

fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_role_binding (*namespace*, ***kwargs*)

delete collection of RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_role_binding(namespace, async=True)` >>> `result = thread.get()`

:param bool async :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str

resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_role_binding_with_http_info (*namespace*,
***kwargs*)

delete collection of RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_role_binding_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_collection_namespaced_role_with_http_info (*namespace*, ***kwargs*)

delete collection of Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_namespaced_role_with_http_info(namespace, async=True)` >>> `result = thread.get()`

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_role (*name, namespace, body, **kwargs*)

delete a Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.delete_namespaced_role(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Role (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param V1DeleteOptions body: (required) :param str pretty: If 'true', then the output is pretty printed. :param int grace_period_seconds: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

delete_namespaced_role_binding (*name, namespace, body, **kwargs*)

`delete a RoleBinding` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_role_binding(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_namespaced_role_binding_with_http_info` (*name, namespace, body, **kwargs*)

`delete a RoleBinding` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_role_binding_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

`delete_namespaced_role_with_http_info` (*name, namespace, body, **kwargs*)

`delete a Role` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_namespaced_role_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Role (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1DeleteOptions body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `int grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param `bool orphan_dependents`: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or PropagationPolicy may be set, but not both. :param `str propagation_policy`: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

get_api_resources (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIResourceList`

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_resources_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIResourceList`

If the method is called asynchronously, returns the request thread.

list_cluster_role (**kwargs)

list or watch objects of kind `ClusterRole` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_cluster_role(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1alpha1ClusterRoleList`

If the method is called asynchronously, returns the request thread.

list_cluster_role_binding (**kwargs)

list or watch objects of kind ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_cluster_role_binding(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: `V1alpha1ClusterRoleBindingList`

If the method is called asynchronously, returns the request thread.

list_cluster_role_binding_with_http_info (**kwargs)

list or watch objects of kind ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_cluster_role_binding_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str pretty`: If 'true', then the output is pretty printed. :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a

list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1ClusterRoleBindingList

If the method is called asynchronously, returns the request thread.

list_cluster_role_with_http_info (kwargs)**

list or watch objects of kind ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_cluster_role_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given

rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1ClusterRoleList

If the method is called asynchronously, returns the request thread.

list_namespaced_role (namespace, **kwargs)

list or watch objects of kind Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_role(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleList

If the method is called asynchronously, returns the request thread.

list_namespaced_role_binding (namespace, **kwargs)

list or watch objects of kind RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_role_binding(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or

a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleBindingList

If the method is called asynchronously, returns the request thread.

list_namespaced_role_binding_with_http_info(namespace, **kwargs)

list or watch objects of kind RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_role_binding_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted

after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleBindingList

If the method is called asynchronously, returns the request thread.

list_namespaced_role_with_http_info(namespace, **kwargs)

list or watch objects of kind Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_namespaced_role_with_http_info(namespace, async=True) >>> result = thread.get()

:param async bool :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleList

If the method is called asynchronously, returns the request thread.

list_role_binding_for_all_namespaces(**kwargs)

list or watch objects of kind RoleBinding This method makes a synchronous HTTP request

by default. To make an asynchronous HTTP request, please pass `async=True` `>>> thread = api.list_role_binding_for_all_namespaces(async=True) >>> result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str pretty`: If 'true', then the output is pretty printed. :param `str resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param `int timeout_seconds`: Timeout for the list/watch call. :param `bool watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: `V1alpha1RoleBindingList`

If the method is called asynchronously, returns the request thread.

`list_role_binding_for_all_namespaces_with_http_info` (kwargs)**

list or watch objects of kind RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` `>>> thread = api.list_role_binding_for_all_namespaces_with_http_info(async=True) >>> result = thread.get()`

:param `async` bool :param `str _continue`: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param `str field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool include_uninitialized`: If true, partially initialized resources are included in the response. :param `str label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested

objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleBindingList

If the method is called asynchronously, returns the request thread.

list_role_for_all_namespaces (**kwargs)

list or watch objects of kind Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_role_for_all_namespaces(async=True) >>> result = thread.get()

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify

resourceVersion. :return: V1alpha1RoleList

If the method is called asynchronously, returns the request thread.

list_role_for_all_namespaces_with_http_info (**kwargs)

list or watch objects of kind Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.list_role_for_all_namespaces_with_http_info(async=True)` >>> `result = thread.get()`

:param async bool :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str pretty: If 'true', then the output is pretty printed. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1alpha1RoleList

If the method is called asynchronously, returns the request thread.

patch_cluster_role (name, body, **kwargs)

partially update the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_cluster_role(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ClusterRole (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRole

If the method is called asynchronously, returns the request thread.

patch_cluster_role_binding (name, body, **kwargs)

partially update the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_cluster_role_binding(name, body, async=True)` >>> `result = thread.get()`

:param async bool :param str name: name of the ClusterRoleBinding (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRoleBinding

If the method is called asynchronously, returns the request thread.

patch_cluster_role_binding_with_http_info (*name, body, **kwargs*)

partially update the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_cluster_role_binding_with_http_info(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRoleBinding (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRoleBinding

If the method is called asynchronously, returns the request thread.

patch_cluster_role_with_http_info (*name, body, **kwargs*)

partially update the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_cluster_role_with_http_info(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRole (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRole

If the method is called asynchronously, returns the request thread.

patch_namespaced_role (*name, namespace, body, **kwargs*)

partially update the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_role(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Role (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1Role

If the method is called asynchronously, returns the request thread.

patch_namespaced_role_binding (*name, namespace, body, **kwargs*)

partially update the specified RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_role_binding(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the RoleBinding (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1RoleBinding

If the method is called asynchronously, returns the request thread.

patch_namespaced_role_binding_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_namespaced_role_binding_with_http_info(name, namespace, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the RoleBinding (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1RoleBinding

If the method is called asynchronously, returns the request thread.

patch_namespaced_role_with_http_info (*name, namespace, body, **kwargs*)

partially update the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread =

```
api.patch_namespaced_role_with_http_info(name, namespace, body, async=True) >>> result =
thread.get()
```

:param async bool :param str name: name of the Role (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1Role

If the method is called asynchronously, returns the request thread.

read_cluster_role (*name*, ***kwargs*)

read the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_cluster_role(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRole (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRole

If the method is called asynchronously, returns the request thread.

read_cluster_role_binding (*name*, ***kwargs*)

read the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_cluster_role_binding(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRoleBinding (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRoleBinding

If the method is called asynchronously, returns the request thread.

read_cluster_role_binding_with_http_info (*name*, ***kwargs*)

read the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_cluster_role_binding_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRoleBinding (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRoleBinding

If the method is called asynchronously, returns the request thread.

read_cluster_role_with_http_info (*name*, ***kwargs*)

read the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_cluster_role_with_http_info(name, async=True) >>> result = thread.get()

:param async bool :param str name: name of the ClusterRole (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1ClusterRole

If the method is called asynchronously, returns the request thread.

read_namespaced_role (*name*, *namespace*, ***kwargs*)

read the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.read_namespaced_role(name, namespace, async=True) >>> result = thread.get()

:param async bool :param str name: name of the Role (required) :param str namespace: object name and auth scope, such as for teams and projects (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1alpha1Role

If the method is called asynchronously, returns the request thread.

read_namespaced_role_binding (*name*, *namespace*, ***kwargs*)

read the specified RoleBinding This method makes a synchronous HTTP request by de-

fault. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_role_binding(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1RoleBinding`

If the method is called asynchronously, returns the request thread.

`read_namespaced_role_binding_with_http_info` (*name, namespace, **kwargs*)

read the specified RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_role_binding_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1RoleBinding`

If the method is called asynchronously, returns the request thread.

`read_namespaced_role_with_http_info` (*name, namespace, **kwargs*)

read the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_namespaced_role_with_http_info(name, namespace, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Role (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1Role`

If the method is called asynchronously, returns the request thread.

`replace_cluster_role` (*name, body, **kwargs*)

replace the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_cluster_role(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRole (required) :param `V1alpha1ClusterRole body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRole`

If the method is called asynchronously, returns the request thread.

`replace_cluster_role_binding` (*name, body, **kwargs*)

replace the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_cluster_role_binding(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRoleBinding (required) :param `V1alpha1ClusterRoleBinding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRoleBinding`

If the method is called asynchronously, returns the request thread.

`replace_cluster_role_binding_with_http_info` (*name, body, **kwargs*)

replace the specified ClusterRoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_cluster_role_binding_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRoleBinding (required) :param `V1alpha1ClusterRoleBinding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRoleBinding`

If the method is called asynchronously, returns the request thread.

replace_cluster_role_with_http_info (*name*, *body*, ***kwargs*)

replace the specified ClusterRole This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_cluster_role_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the ClusterRole (required) :param `V1alpha1ClusterRole body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1ClusterRole`

If the method is called asynchronously, returns the request thread.

replace_namespaced_role (*name*, *namespace*, *body*, ***kwargs*)

replace the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_role(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Role (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1alpha1Role body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1Role`

If the method is called asynchronously, returns the request thread.

replace_namespaced_role_binding (*name*, *namespace*, *body*, ***kwargs*)

replace the specified RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_role_binding(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1alpha1RoleBinding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1RoleBinding`

If the method is called asynchronously, returns the request thread.

replace_namespaced_role_binding_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

replace the specified RoleBinding This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_role_binding_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the RoleBinding (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1alpha1RoleBinding body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1RoleBinding`

If the method is called asynchronously, returns the request thread.

replace_namespaced_role_with_http_info (*name*, *namespace*, *body*, ***kwargs*)

replace the specified Role This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_namespaced_role_with_http_info(name, namespace, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the Role (required) :param `str namespace`: object name and auth scope, such as for teams and projects (required) :param `V1alpha1Role body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1alpha1Role`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.storage_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.storage_api.StorageApi` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

get_api_group (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

get_api_group_with_http_info (***kwargs*)

get information of a group This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_api_group_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `V1APIGroup`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.storage_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.
Ref: <https://github.com/swagger-api/swagger-codegen>

create_storage_class (*body, **kwargs*)

create a StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_storage_class(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1beta1StorageClass` `body`: (required) :param `str` `pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

create_storage_class_with_http_info (*body*, ***kwargs*)

create a StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.create_storage_class_with_http_info(body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `V1beta1StorageClass` *body*: (required) :param `str` *pretty*: If 'true', then the output is pretty printed. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

delete_collection_storage_class (***kwargs*)

delete collection of StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_storage_class(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str` *pretty*: If 'true', then the output is pretty printed. :param `str` *_continue*: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the continue field. This field is not supported when `watch` is true. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param `str` *field_selector*: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param `bool` *include_uninitialized*: If true, partially initialized resources are included in the response. :param `str` *label_selector*: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param `int` *limit*: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if `watch` is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param `str` *resource_version*: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given *rv*. :param `int` *timeout_seconds*: Timeout for the list/watch call. :param `bool` *watch*: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_collection_storage_class_with_http_info (***kwargs*)

delete collection of StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_collection_storage_class_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :param `str` *pretty*: If 'true', then the output is pretty printed. :param `str` *_continue*: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters

(except for the value of `continue`) and the server may reject a `continue` value it does not recognize. If the specified `continue` value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 `ResourceExpired` error indicating the client must restart their list without the `continue` field. This field is not supported when `watch` is `true`. Clients may start a watch from the last `resourceVersion` value returned by the server and not miss any modifications. :param str `field_selector`: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool `include_uninitialized`: If true, partially initialized resources are included in the response. :param str `label_selector`: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int `limit`: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the `continue` field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the `continue` field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the `continue` field is empty, clients may assume that no more results are available. This field is not supported if `watch` is `true`. The server guarantees that the objects returned when using `continue` will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str `resource_version`: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int `timeout_seconds`: Timeout for the list/watch call. :param bool `watch`: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify `resourceVersion`. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_storage_class (*name, body, **kwargs*)

delete a `StorageClass` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_storage_class(name, body, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `name`: name of the `StorageClass` (required) :param `V1DeleteOptions` `body`: (required) :param bool `pretty`: If 'true', then the output is pretty printed. :param int `grace_period_seconds`: The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool `orphan_dependents`: Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the "orphan" finalizer will be added to/removed from the object's finalizers list. Either this field or `PropagationPolicy` may be set, but not both. :param str `propagation_policy`: Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: `V1Status`

If the method is called asynchronously, returns the request thread.

delete_storage_class_with_http_info (*name, body, **kwargs*)

delete a `StorageClass` This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.delete_storage_class_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param bool `async`: :param str `name`: name of the `StorageClass` (required) :param `V1DeleteOptions` `body`: (required) :param bool `pretty`: If 'true', then the output is pretty printed. :param int `grace_period_seconds`:

The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately. :param bool orphan_dependents: Deprecated: please use the PropagationPolicy, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or PropagationPolicy may be set, but not both. :param str propagation_policy: Whether and how garbage collection will be performed. Either this field or OrphanDependents may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy. :return: V1Status

If the method is called asynchronously, returns the request thread.

get_api_resources (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

get_api_resources_with_http_info (**kwargs)

get available resources This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.get_api_resources_with_http_info(async=True) >>> result = thread.get()

:param async bool :return: V1APIResourceList

If the method is called asynchronously, returns the request thread.

list_storage_class (**kwargs)

list or watch objects of kind StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_storage_class(async=True) >>> result = thread.get()

:param async bool :param str pretty: If ‘true’, then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the

version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StorageClassList

If the method is called asynchronously, returns the request thread.

list_storage_class_with_http_info (**kwargs)

list or watch objects of kind StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.list_storage_class_with_http_info(async=True) >>> result = thread.get()

:param async bool :param str pretty: If 'true', then the output is pretty printed. :param str _continue: The continue option should be set when retrieving more results from the server. Since this value is server defined, clients may only use the continue value from a previous query result with identical query parameters (except for the value of continue) and the server may reject a continue value it does not recognize. If the specified continue value is no longer valid whether due to expiration (generally five to fifteen minutes) or a configuration change on the server the server will respond with a 410 ResourceExpired error indicating the client must restart their list without the continue field. This field is not supported when watch is true. Clients may start a watch from the last resourceVersion value returned by the server and not miss any modifications. :param str field_selector: A selector to restrict the list of returned objects by their fields. Defaults to everything. :param bool include_uninitialized: If true, partially initialized resources are included in the response. :param str label_selector: A selector to restrict the list of returned objects by their labels. Defaults to everything. :param int limit: limit is a maximum number of responses to return for a list call. If more items exist, the server will set the *continue* field on the list metadata to a value that can be used with the same initial query to retrieve the next set of results. Setting a limit may return fewer than the requested amount of items (up to zero items) in the event all requested objects are filtered out and clients should only use the presence of the continue field to determine whether more results are available. Servers may choose not to support the limit argument and will return all of the available results. If limit is specified and the continue field is empty, clients may assume that no more results are available. This field is not supported if watch is true. The server guarantees that the objects returned when using continue will be identical to issuing a single list call without a limit - that is, no objects created, modified, or deleted after the first request is issued will be included in any subsequent continued requests. This is sometimes referred to as a consistent snapshot, and ensures that a client that is using limit to receive smaller chunks of a very large result can ensure they see all possible objects. If objects are updated during a chunked list the version of the object that was present at the time the first list result was calculated is returned. :param str resource_version: When specified with a watch call, shows changes that occur after that particular version of a resource. Defaults to changes from the beginning of history. When specified for list: - if unset, then the result is returned from remote storage based on quorum-read flag; - if it's 0, then we simply return what we currently have in cache, no guarantee; - if set to non zero, then the result is at least as fresh as given rv. :param int timeout_seconds: Timeout for the list/watch call. :param bool watch: Watch for changes to the described resources and return them as a stream of add, update, and remove notifications. Specify resourceVersion. :return: V1beta1StorageClassList

If the method is called asynchronously, returns the request thread.

patch_storage_class (name, body, **kwargs)

partially update the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass async=True >>> thread = api.patch_storage_class(name, body, async=True) >>> result = thread.get()

:param async bool :param str name: name of the StorageClass (required) :param object body: (required) :param str pretty: If 'true', then the output is pretty printed. :return: V1beta1StorageClass

If the method is called asynchronously, returns the request thread.

patch_storage_class_with_http_info (*name, body, **kwargs*)

partially update the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.patch_storage_class_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StorageClass (required) :param `object body`: (required)
:param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

read_storage_class (*name, **kwargs*)

read the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_storage_class(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StorageClass (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

read_storage_class_with_http_info (*name, **kwargs*)

read the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.read_storage_class_with_http_info(name, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StorageClass (required) :param `str pretty`: If 'true', then the output is pretty printed. :param `bool exact`: Should the export be exact. Exact export maintains cluster-specific fields like 'Namespace'. :param `bool export`: Should this value be exported. Export strips fields that a user can not specify. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

replace_storage_class (*name, body, **kwargs*)

replace the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_storage_class(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StorageClass (required) :param `V1beta1StorageClass body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

replace_storage_class_with_http_info (*name, body, **kwargs*)

replace the specified StorageClass This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.replace_storage_class_with_http_info(name, body, async=True)` >>> `result = thread.get()`

:param `async` bool :param `str name`: name of the StorageClass (required) :param `V1beta1StorageClass body`: (required) :param `str pretty`: If 'true', then the output is pretty printed. :return: `V1beta1StorageClass`

If the method is called asynchronously, returns the request thread.

kubernetes.client.apis.version_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.apis.version_api.VersionApi` (*api_client=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

Ref: <https://github.com/swagger-api/swagger-codegen>

get_code (***kwargs*)

get the code version This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_code(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `VersionInfo`

If the method is called asynchronously, returns the request thread.

get_code_with_http_info (***kwargs*)

get the code version This method makes a synchronous HTTP request by default. To make an asynchronous HTTP request, please pass `async=True` >>> `thread = api.get_code_with_http_info(async=True)` >>> `result = thread.get()`

:param `async` bool :return: `VersionInfo`

If the method is called asynchronously, returns the request thread.

Module contents

`kubernetes.client.models` package

Submodules

`kubernetes.client.models.intstr_int_or_string` module

`kubernetes.client.models.resource_quantity` module

`kubernetes.client.models.runtime_raw_extension` module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.runtime_raw_extension.RuntimeRawExtension` (*raw=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'raw': 'Raw'}

raw

Gets the raw of this RuntimeRawExtension. Raw is the underlying serialization of this object.

Returns The raw of this RuntimeRawExtension.

Return type `str`

```

swagger_types = {'raw': 'str'}

to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model

```

`kubernetes.client.models.unversioned_api_group` module

`kubernetes.client.models.unversioned_api_group_list` module

`kubernetes.client.models.unversioned_api_resource` module

`kubernetes.client.models.unversioned_api_resource_list` module

`kubernetes.client.models.unversioned_api_versions` module

`kubernetes.client.models.unversioned_group_version_for_discovery` module

`kubernetes.client.models.unversioned_label_selector` module

`kubernetes.client.models.unversioned_label_selector_requirement` module

`kubernetes.client.models.unversioned_list_meta` module

`kubernetes.client.models.unversioned_server_address_by_client_cidr` module

`kubernetes.client.models.unversioned_status` module

`kubernetes.client.models.unversioned_status_cause` module

`kubernetes.client.models.unversioned_status_details` module

`kubernetes.client.models.unversioned_time` module

`kubernetes.client.models.v1_attached_volume` module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.client.models.v1_attached_volume.V1AttachedVolume(device_path=None,
                                                                    name=None)

```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```

attribute_map = {'device_path': 'devicePath', 'name': 'name'}

```

device_path

Gets the device_path of this V1AttachedVolume. DevicePath represents the device path where the volume should be available

Returns The device_path of this V1AttachedVolume.

Return type str

name

Gets the name of this V1AttachedVolume. Name of the attached volume

Returns The name of this V1AttachedVolume.

Return type str

swagger_types = {'device_path': 'str', 'name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_aws_elastic_block_store_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AWSElasticBlockStoreVolumeSource

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'fs_type': 'fsType', 'volume_id': 'volumeID', 'partition': 'partition'}

fs_type

Gets the fs_type of this V1AWSElasticBlockStoreVolumeSource. Filesystem type of the volume that you want to mount. Tip: Ensure that the filesystem type is supported by the host operating system. Examples: "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified. More info: <https://kubernetes.io/docs/concepts/storage/volumes#awselasticblockstore>

Returns The fs_type of this V1AWSElasticBlockStoreVolumeSource.

Return type str

partition

Gets the partition of this V1AWSElasticBlockStoreVolumeSource. The partition in the volume that you want to mount. If omitted, the default is to mount by volume name. Examples: For volume /dev/sda1, you specify the partition as "1". Similarly, the volume partition for /dev/sda is "0" (or you can leave the property empty).

Returns The partition of this V1AWSElasticBlockStoreVolumeSource.

Return type int

read_only

Gets the read_only of this V1AWSElasticBlockStoreVolumeSource. Specify “true” to force and set the ReadOnly property in VolumeMounts to “true”. If omitted, the default is “false”. More info: <https://kubernetes.io/docs/concepts/storage/volumes#awselasticblockstore>

Returns The read_only of this V1AWSElasticBlockStoreVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'fs_type': 'str', 'volume_id': 'str', 'partition': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

volume_id

Gets the volume_id of this V1AWSElasticBlockStoreVolumeSource. Unique ID of the persistent disk resource in AWS (Amazon EBS volume). More info: <https://kubernetes.io/docs/concepts/storage/volumes#awselasticblockstore>

Returns The volume_id of this V1AWSElasticBlockStoreVolumeSource.

Return type str

kubernetes.client.models.v1_azure_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource(cached_mode=None, disk_name=None, disk_uri=None, fs_type=None, kind=None, read_only=None):
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'disk_uri': 'diskURI', 'read_only': 'readOnly', 'kind': 'kind', 'cached_mode': 'cachingMode', 'fs_

caching_mode

Gets the caching_mode of this V1AzureDiskVolumeSource. Host Caching mode: None, Read Only, Read Write.

Returns The caching_mode of this V1AzureDiskVolumeSource.

Return type str

disk_name

Gets the disk_name of this V1AzureDiskVolumeSource. The Name of the data disk in the blob storage

Returns The disk_name of this V1AzureDiskVolumeSource.

Return type str

disk_uri

Gets the disk_uri of this V1AzureDiskVolumeSource. The URI the data disk in the blob storage

Returns The disk_uri of this V1AzureDiskVolumeSource.

Return type str

fs_type

Gets the fs_type of this V1AzureDiskVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Ex. "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified.

Returns The fs_type of this V1AzureDiskVolumeSource.

Return type str

kind

Gets the kind of this V1AzureDiskVolumeSource. Expected values Shared: multiple blob disks per storage account Dedicated: single blob disk per storage account Managed: azure managed data disk (only in managed availability set). defaults to shared

Returns The kind of this V1AzureDiskVolumeSource.

Return type str

read_only

Gets the read_only of this V1AzureDiskVolumeSource. Defaults to false (read/write). ReadOnly here will force the ReadOnly setting in VolumeMounts.

Returns The read_only of this V1AzureDiskVolumeSource.

Return type bool

swagger_types = {'disk_uri': 'str', 'read_only': 'bool', 'kind': 'str', 'caching_mode': 'str', 'fs_type': 'str', 'disk_name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_azure_file_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_azure_file_volume_source.V1AzureFileVolumeSource (read_only=No

se-

cret_name=No

share_name=No

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'secret_name': 'secretName', 'share_name': 'shareName'}

read_only

Gets the read_only of this V1AzureFileVolumeSource. Defaults to false (read/write). ReadOnly here will force the ReadOnly setting in VolumeMounts.

Returns The read_only of this V1AzureFileVolumeSource.

Return type bool

secret_name

Gets the secret_name of this V1AzureFileVolumeSource. the name of secret that contains Azure Storage Account Name and Key

Returns The secret_name of this V1AzureFileVolumeSource.

Return type str

share_name

Gets the share_name of this V1AzureFileVolumeSource. Share Name

Returns The share_name of this V1AzureFileVolumeSource.

Return type str

swagger_types = {'read_only': 'bool', 'secret_name': 'str', 'share_name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_binding.V1Binding(*api_version=None, kind=None, metadata=None, target=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Binding. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Binding.

Return type str

attribute_map = {'kind': 'kind', 'target': 'target', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1Binding. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Binding.

Return type str

metadata

Gets the metadata of this V1Binding. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Binding.

Return type *V1ObjectMeta*

swagger_types = {'kind': 'str', 'target': 'V1ObjectReference', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

target

Gets the target of this V1Binding. The target object that you want to bind to the standard object.

Returns The target of this V1Binding.

Return type *V1ObjectReference*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_capabilities module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_capabilities.V1Capabilities` (*add=None*,
drop=None)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

add

Gets the add of this V1Capabilities. Added capabilities

Returns The add of this V1Capabilities.

Return type `list[str]`

attribute_map = {'add': 'add', 'drop': 'drop'}

drop

Gets the drop of this V1Capabilities. Removed capabilities

Returns The drop of this V1Capabilities.

Return type `list[str]`

swagger_types = {'add': 'list[str]', 'drop': 'list[str]'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_ceph_fs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource (monitors=None,
                                                                    path=None,
                                                                    read_only=None,
                                                                    se-
                                                                    cret_file=None,
                                                                    se-
                                                                    cret_ref=None,
                                                                    user=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'secret_file': 'secretFile', 'secret_ref': 'secretRef', 'user': 'user', 'path': 'p

monitors

Gets the monitors of this V1CephFSVolumeSource. Required: Monitors is a collection of Ceph monitors
More info: <https://releases.k8s.io/HEAD/examples/volumes/cephfs/README.md#how-to-use-it>

Returns The monitors of this V1CephFSVolumeSource.

Return type list[str]

path

Gets the path of this V1CephFSVolumeSource. Optional: Used as the mounted root, rather than the full Ceph tree, default is /

Returns The path of this V1CephFSVolumeSource.

Return type str

read_only

Gets the read_only of this V1CephFSVolumeSource. Optional: Defaults to false (read/write). Read-Only here will force the ReadOnly setting in VolumeMounts. More info: <https://releases.k8s.io/HEAD/examples/volumes/cephfs/README.md#how-to-use-it>

Returns The read_only of this V1CephFSVolumeSource.

Return type bool

secret_file

Gets the secret_file of this V1CephFSVolumeSource. Optional: SecretFile is the path to key ring for User, default is /etc/ceph/user.secret More info: <https://releases.k8s.io/HEAD/examples/volumes/cephfs/README.md#how-to-use-it>

Returns The secret_file of this V1CephFSVolumeSource.

Return type str

secret_ref

Gets the secret_ref of this V1CephFSVolumeSource. Optional: SecretRef is reference to the authentication secret for User, default is empty. More info: <https://releases.k8s.io/HEAD/examples/volumes/cephfs/README.md#how-to-use-it>

Returns The secret_ref of this V1CephFSVolumeSource.

Return type *V1LocalObjectReference*

swagger_types = {'read_only': 'bool', 'secret_file': 'str', 'secret_ref': 'V1LocalObjectReference', 'user': 'str', 'path':

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

user

Gets the user of this V1CephFSVolumeSource. Optional: User is the rados user name, default is admin
More info: <https://releases.k8s.io/HEAD/examples/volumes/cephfs/README.md#how-to-use-it>

Returns The user of this V1CephFSVolumeSource.

Return type str

kubernetes.client.models.v1_cinder_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource (*fs_type=None, read_only=None, volume_id=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'fs_type': 'fsType', 'volume_id': 'volumeID'}

fs_type

Gets the fs_type of this V1CinderVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Examples: "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified. More info: <https://releases.k8s.io/HEAD/examples/mysql-cinder-pd/README.md>

Returns The fs_type of this V1CinderVolumeSource.

Return type str

read_only

Gets the read_only of this V1CinderVolumeSource. Optional: Defaults to false (read/write). Read-Only here will force the ReadOnly setting in VolumeMounts. More info: <https://releases.k8s.io/HEAD/examples/mysql-cinder-pd/README.md>

Returns The read_only of this V1CinderVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'fs_type': 'str', 'volume_id': 'str'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

volume_id

Gets the volume_id of this V1CinderVolumeSource. volume id used to identify the volume in cinder More info: <https://releases.k8s.io/HEAD/examples/mysql-cinder-pd/README.md>

Returns The volume_id of this V1CinderVolumeSource.

Return type str

kubernetes.client.models.v1_component_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_component_condition.V1ComponentCondition (error=None,
                                                                              mes-
                                                                              sage=None,
                                                                              sta-
                                                                              tus=None,
                                                                              type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'status': 'status', 'message': 'message', 'type': 'type', 'error': 'error'}

error

Gets the error of this V1ComponentCondition. Condition error code for a component. For example, a health check error code.

Returns The error of this V1ComponentCondition.

Return type str

message

Gets the message of this V1ComponentCondition. Message about the condition for a component. For example, information about a health check.

Returns The message of this V1ComponentCondition.

Return type str

status

Gets the status of this V1ComponentCondition. Status of the condition for a component. Valid values for "Healthy": "True", "False", or "Unknown".

Returns The status of this V1ComponentCondition.

Return type str

swagger_types = {'status': 'str', 'message': 'str', 'type': 'str', 'error': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1ComponentCondition. Type of condition for a component. Valid value: "Healthy"

Returns The type of this V1ComponentCondition.

Return type str

kubernetes.client.models.v1_component_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_component_status.V1ComponentStatus (api_version=None,
                                                                    condi-
                                                                    tions=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ComponentStatus. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ComponentStatus.

Return type str

attribute_map = {'kind': 'kind', 'conditions': 'conditions', 'api_version': 'apiVersion', 'metadata': 'metadata'}

conditions

Gets the conditions of this V1ComponentStatus. List of component conditions observed

Returns The conditions of this V1ComponentStatus.

Return type list[V1ComponentCondition]

kind

Gets the kind of this V1ComponentStatus. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ComponentStatus.

Return type str

metadata

Gets the metadata of this V1ComponentStatus. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ComponentStatus.

Return type V1ObjectMeta

swagger_types = {'kind': 'str', 'conditions': 'list[V1ComponentCondition]', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()
Returns the string representation of the model

kubernetes.client.models.v1_component_status_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_component_status_list.V1ComponentStatusList (*api_version=None, items=None, kind=None, meta-data=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ComponentStatusList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ComponentStatusList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1ComponentStatusList. List of ComponentStatus objects.

Returns The items of this V1ComponentStatusList.

Return type list[V1ComponentStatus]

kind

Gets the kind of this V1ComponentStatusList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ComponentStatusList.

Return type str

metadata

Gets the metadata of this V1ComponentStatusList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1ComponentStatusList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1ComponentStatus]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()
Returns the model properties as a dict


```
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1_config_map module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_config_map.V1ConfigMap (api_version=None,
                                                         data=None,    kind=None,
                                                         metadata=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ConfigMap. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ConfigMap.

Return type str

```
attribute_map = {'kind': 'kind', 'data': 'data', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

data

Gets the data of this V1ConfigMap. Data contains the configuration data. Each key must consist of alphanumeric characters, '-', '_' or '.'.

Returns The data of this V1ConfigMap.

Return type dict(str, str)

kind

Gets the kind of this V1ConfigMap. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ConfigMap.

Return type str

metadata

Gets the metadata of this V1ConfigMap. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ConfigMap.

Return type *V1ObjectMeta*

```
swagger_types = {'kind': 'str', 'data': 'dict(str, str)', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

```
to_dict()
    Returns the model properties as a dict
```

```
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1_config_map_key_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector (key=None,  
name=None,  
optional=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'optional': 'optional', 'name': 'name', 'key': 'key'}

key

Gets the key of this V1ConfigMapKeySelector. The key to select.

Returns The key of this V1ConfigMapKeySelector.

Return type str

name

Gets the name of this V1ConfigMapKeySelector. Name of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1ConfigMapKeySelector.

Return type str

optional

Gets the optional of this V1ConfigMapKeySelector. Specify whether the ConfigMap or it's key must be defined

Returns The optional of this V1ConfigMapKeySelector.

Return type bool

swagger_types = {'optional': 'bool', 'name': 'str', 'key': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_config_map_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_config_map_list.V1ConfigMapList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ConfigMapList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ConfigMapList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1ConfigMapList. Items is the list of ConfigMaps.

Returns The items of this V1ConfigMapList.

Return type list[V1ConfigMap]

kind

Gets the kind of this V1ConfigMapList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ConfigMapList.

Return type str

metadata

Gets the metadata of this V1ConfigMapList. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ConfigMapList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1ConfigMap]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_config_map_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource (default_mode=  
items=None,  
name=None,  
op-  
tional=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'items': 'items', 'default_mode': 'defaultMode', 'optional': 'optional', 'name': 'name'}

default_mode

Gets the default_mode of this V1ConfigMapVolumeSource. Optional: mode bits to use on created files by default. Must be a value between 0 and 0777. Defaults to 0644. Directories within the path are not affected by this setting. This might be in conflict with other options that affect the file mode, like fsGroup, and the result can be other mode bits set.

Returns The default_mode of this V1ConfigMapVolumeSource.

Return type int

items

Gets the items of this V1ConfigMapVolumeSource. If unspecified, each key-value pair in the Data field of the referenced ConfigMap will be projected into the volume as a file whose name is the key and content is the value. If specified, the listed keys will be projected into the specified paths, and unlisted keys will not be present. If a key is specified which is not present in the ConfigMap, the volume setup will error unless it is marked optional. Paths must be relative and may not contain the '..' path or start with '..'.

Returns The items of this V1ConfigMapVolumeSource.

Return type list[V1KeyToPath]

name

Gets the name of this V1ConfigMapVolumeSource. Name of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1ConfigMapVolumeSource.

Return type str

optional

Gets the optional of this V1ConfigMapVolumeSource. Specify whether the ConfigMap or it's keys must be defined

Returns The optional of this V1ConfigMapVolumeSource.

Return type bool

swagger_types = {'items': 'list[V1KeyToPath]', 'default_mode': 'int', 'optional': 'bool', 'name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_container module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_container.V1Container (args=None, command=None,
env=None, env_from=None,
image=None, image_pull_policy=None,
lifecycle=None, liveness_probe=None,
name=None, ports=None, readiness_probe=None,
resources=None, security_context=None,
stdin=None, stdin_once=None, termination_message_path=None,
termination_message_policy=None,
tty=None, volume_mounts=None, working_dir=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

args

Gets the args of this V1Container. Arguments to the entrypoint. The docker image's CMD is used if this is not provided. Variable references \$(VAR_NAME) are expanded using the container's environment. If a variable cannot be resolved, the reference in the input string will be unchanged. The \$(VAR_NAME) syntax can be escaped with a double \$\$, ie: \$\$\$(VAR_NAME). Escaped references will never be expanded, regardless of whether the variable exists or not. Cannot be updated. More info: <https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/#running-a-command-in-a-shell>

Returns The args of this V1Container.

Return type list[str]

attribute_map = {'image_pull_policy': 'imagePullPolicy', 'tty': 'tty', 'security_context': 'securityContext', 'stdin_once': 'stdinOnce'}

command

Gets the command of this V1Container. Entrypoint array. Not executed within a shell. The docker image's ENTRYPOINT is used if this is not provided. Variable references \$(VAR_NAME) are expanded using the container's environment. If a variable cannot be resolved, the reference in the input string will be unchanged. The \$(VAR_NAME) syntax can be escaped with a double \$\$, ie: \$\$\$(VAR_NAME). Escaped references will never be expanded, regardless of whether the variable exists or not. Cannot be updated. More info: <https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/#running-a-command-in-a-shell>

Returns The command of this V1Container.

Return type list[str]

env

Gets the env of this V1Container. List of environment variables to set in the container. Cannot be updated.

Returns The env of this V1Container.

Return type list[V1EnvVar]

env_from

Gets the env_from of this V1Container. List of sources to populate environment variables in the container. The keys defined within a source must be a C_IDENTIFIER. All invalid keys will be reported as an event

when the container is starting. When a key exists in multiple sources, the value associated with the last source will take precedence. Values defined by an Env with a duplicate key will take precedence. Cannot be updated.

Returns The env_from of this V1Container.

Return type list[V1EnvFromSource]

image

Gets the image of this V1Container. Docker image name. More info: <https://kubernetes.io/docs/concepts/containers/images> This field is optional to allow higher level config management to default or override container images in workload controllers like Deployments and StatefulSets.

Returns The image of this V1Container.

Return type str

image_pull_policy

Gets the image_pull_policy of this V1Container. Image pull policy. One of Always, Never, IfNotPresent. Defaults to Always if :latest tag is specified, or IfNotPresent otherwise. Cannot be updated. More info: <https://kubernetes.io/docs/concepts/containers/images#updating-images>

Returns The image_pull_policy of this V1Container.

Return type str

lifecycle

Gets the lifecycle of this V1Container. Actions that the management system should take in response to container lifecycle events. Cannot be updated.

Returns The lifecycle of this V1Container.

Return type *V1Lifecycle*

liveness_probe

Gets the liveness_probe of this V1Container. Periodic probe of container liveness. Container will be restarted if the probe fails. Cannot be updated. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#container-probes>

Returns The liveness_probe of this V1Container.

Return type *V1Probe*

name

Gets the name of this V1Container. Name of the container specified as a DNS_LABEL. Each container in a pod must have a unique name (DNS_LABEL). Cannot be updated.

Returns The name of this V1Container.

Return type str

ports

Gets the ports of this V1Container. List of ports to expose from the container. Exposing a port here gives the system additional information about the network connections a container uses, but is primarily informational. Not specifying a port here DOES NOT prevent that port from being exposed. Any port which is listening on the default “0.0.0.0” address inside a container will be accessible from the network. Cannot be updated.

Returns The ports of this V1Container.

Return type list[V1ContainerPort]

readiness_probe

Gets the `readiness_probe` of this `V1Container`. Periodic probe of container service readiness. Container will be removed from service endpoints if the probe fails. Cannot be updated. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#container-probes>

Returns The `readiness_probe` of this `V1Container`.

Return type *V1Probe*

resources

Gets the resources of this `V1Container`. Compute Resources required by this container. Cannot be updated. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#resources>

Returns The resources of this `V1Container`.

Return type *V1ResourceRequirements*

security_context

Gets the `security_context` of this `V1Container`. Security options the pod should run with. More info: <https://kubernetes.io/docs/concepts/policy/security-context/> More info: https://git.k8s.io/community/contributors/design-proposals/security_context.md

Returns The `security_context` of this `V1Container`.

Return type *V1SecurityContext*

stdin

Gets the `stdin` of this `V1Container`. Whether this container should allocate a buffer for `stdin` in the container runtime. If this is not set, reads from `stdin` in the container will always result in EOF. Default is false.

Returns The `stdin` of this `V1Container`.

Return type bool

stdin_once

Gets the `stdin_once` of this `V1Container`. Whether the container runtime should close the `stdin` channel after it has been opened by a single attach. When `stdin` is true the `stdin` stream will remain open across multiple attach sessions. If `stdinOnce` is set to true, `stdin` is opened on container start, is empty until the first client attaches to `stdin`, and then remains open and accepts data until the client disconnects, at which time `stdin` is closed and remains closed until the container is restarted. If this flag is false, a container processes that reads from `stdin` will never receive an EOF. Default is false

Returns The `stdin_once` of this `V1Container`.

Return type bool

swagger_types = {'image_pull_policy': 'str', 'tty': 'bool', 'security_context': 'V1SecurityContext', 'stdin_once': 'bool'}

termination_message_path

Gets the `termination_message_path` of this `V1Container`. Optional: Path at which the file to which the container's termination message will be written is mounted into the container's filesystem. Message written is intended to be brief final status, such as an assertion failure message. Will be truncated by the node if greater than 4096 bytes. The total message length across all containers will be limited to 12kb. Defaults to `/dev/termination-log`. Cannot be updated.

Returns The `termination_message_path` of this `V1Container`.

Return type str

termination_message_policy

Gets the `termination_message_policy` of this `V1Container`. Indicate how the termination message should be populated. File will use the contents of `terminationMessagePath` to populate the container status message on both success and failure. `FallbackToLogsOnError` will use the last chunk of container log output

if the termination message file is empty and the container exited with an error. The log output is limited to 2048 bytes or 80 lines, whichever is smaller. Defaults to File. Cannot be updated.

Returns The `termination_message_policy` of this `V1Container`.

Return type str

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

tty

Gets the `tty` of this `V1Container`. Whether this container should allocate a TTY for itself, also requires 'stdin' to be true. Default is false.

Returns The `tty` of this `V1Container`.

Return type bool

volume_mounts

Gets the `volume_mounts` of this `V1Container`. Pod volumes to mount into the container's filesystem. Cannot be updated.

Returns The `volume_mounts` of this `V1Container`.

Return type list[[V1VolumeMount](#)]

working_dir

Gets the `working_dir` of this `V1Container`. Container's working directory. If not specified, the container runtime's default will be used, which might be configured in the container image. Cannot be updated.

Returns The `working_dir` of this `V1Container`.

Return type str

kubernetes.client.models.v1_container_image module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_container_image.V1ContainerImage` (*names=None, size_bytes=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'size_bytes': 'sizeBytes', 'names': 'names'}

names

Gets the `names` of this `V1ContainerImage`. Names by which this image is known. e.g. ["gcr.io/google_containers/hyperkube:v1.0.7", "dockerhub.io/google_containers/hyperkube:v1.0.7"]

Returns The `names` of this `V1ContainerImage`.

Return type list[str]

size_bytes

Gets the `size_bytes` of this `V1ContainerImage`. The size of the image in bytes.

Returns The size_bytes of this V1ContainerImage.

Return type int

swagger_types = {'size_bytes': 'int', 'names': 'list[str]}'

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_container_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_container_port.V1ContainerPort (container_port=None,
                                                                    host_ip=None,
                                                                    host_port=None,
                                                                    name=None,
                                                                    protocol=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'host_port': 'hostPort', 'protocol': 'protocol', 'host_ip': 'hostIP', 'name': 'name', 'container_port':

container_port

Gets the container_port of this V1ContainerPort. Number of port to expose on the pod's IP address. This must be a valid port number, 0 < x < 65536.

Returns The container_port of this V1ContainerPort.

Return type int

host_ip

Gets the host_ip of this V1ContainerPort. What host IP to bind the external port to.

Returns The host_ip of this V1ContainerPort.

Return type str

host_port

Gets the host_port of this V1ContainerPort. Number of port to expose on the host. If specified, this must be a valid port number, 0 < x < 65536. If HostNetwork is specified, this must match ContainerPort. Most containers do not need this.

Returns The host_port of this V1ContainerPort.

Return type int

name

Gets the name of this V1ContainerPort. If specified, this must be an IANA_SVC_NAME and unique within the pod. Each named port in a pod must have a unique name. Name for the port that can be referred to by services.

Returns The name of this V1ContainerPort.

Return type str

protocol

Gets the protocol of this V1ContainerPort. Protocol for port. Must be UDP or TCP. Defaults to “TCP”.

Returns The protocol of this V1ContainerPort.

Return type str

swagger_types = {'host_port': 'int', 'protocol': 'str', 'host_ip': 'str', 'name': 'str', 'container_port': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_container_state module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_container_state.V1ContainerState (*running=None, terminated=None, waiting=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'terminated': 'terminated', 'running': 'running', 'waiting': 'waiting'}

running

Gets the running of this V1ContainerState. Details about a running container

Returns The running of this V1ContainerState.

Return type *V1ContainerStateRunning*

swagger_types = {'terminated': 'V1ContainerStateTerminated', 'running': 'V1ContainerStateRunning', 'waiting': 'V1ContainerStateWaiting'}

terminated

Gets the terminated of this V1ContainerState. Details about a terminated container

Returns The terminated of this V1ContainerState.

Return type *V1ContainerStateTerminated*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

waiting

Gets the waiting of this V1ContainerState. Details about a waiting container

Returns The waiting of this V1ContainerState.

Return type *V1ContainerStateWaiting*

kubernetes.client.models.v1_container_state_running module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_container_state_running.V1ContainerStateRunning` (*started_at=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'started_at': 'startedAt'}

started_at

Gets the started_at of this V1ContainerStateRunning. Time at which the container was last (re-)started

Returns The started_at of this V1ContainerStateRunning.

Return type datetime

swagger_types = {'started_at': 'datetime'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_container_state_terminated module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated` (*container_id=None, exit_code=None, finished_at=None, message=None, reason=None, signal=None, signal_code=None, started_at=None*)
 Bases: `object`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'finished_at': 'finishedAt', 'signal': 'signal', 'container_id': 'containerID', 'exit_code': 'exitCode', 'message': 'message', 'reason': 'reason', 'signal_code': 'signalCode', 'started_at': 'startedAt'}

container_id

Gets the container_id of this V1ContainerStateTerminated. Container's ID in the format 'docker://<container_id>'

Returns The container_id of this V1ContainerStateTerminated.

Return type str

exit_code

Gets the exit_code of this V1ContainerStateTerminated. Exit status from the last termination of the container

Returns The exit_code of this V1ContainerStateTerminated.

Return type int

finished_at

Gets the finished_at of this V1ContainerStateTerminated. Time at which the container last terminated

Returns The finished_at of this V1ContainerStateTerminated.

Return type datetime

message

Gets the message of this V1ContainerStateTerminated. Message regarding the last termination of the container

Returns The message of this V1ContainerStateTerminated.

Return type str

reason

Gets the reason of this V1ContainerStateTerminated. (brief) reason from the last termination of the container

Returns The reason of this V1ContainerStateTerminated.

Return type str

signal

Gets the signal of this V1ContainerStateTerminated. Signal from the last termination of the container

Returns The signal of this V1ContainerStateTerminated.

Return type int

started_at

Gets the started_at of this V1ContainerStateTerminated. Time at which previous execution of the container started

Returns The started_at of this V1ContainerStateTerminated.

Return type datetime

swagger_types = {'finished_at': 'datetime', 'signal': 'int', 'container_id': 'str', 'exit_code': 'int', 'reason': 'str', 'mess

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_container_state_waiting module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_container_state_waiting.V1ContainerStateWaiting(message=None, reason=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'message': 'message', 'reason': 'reason'}

message

Gets the message of this V1ContainerStateWaiting. Message regarding why the container is not yet running.

Returns The message of this V1ContainerStateWaiting.

Return type str

reason

Gets the reason of this V1ContainerStateWaiting. (brief) reason the container is not yet running.

Returns The reason of this V1ContainerStateWaiting.

Return type str

swagger_types = {'message': 'str', 'reason': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_container_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_container_status.V1ContainerStatus(container_id=None, image_id=None, image_name=None, image_pull_policy=None, last_state=None, name=None, ready=None, restart_count=None, state=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'last_state': 'lastState', 'image_id': 'imageID', 'state': 'state', 'name': 'name', 'ready': 'ready', 'im

container_id

Gets the container_id of this V1ContainerStatus. Container's ID in the format 'docker://<container_id>'.

Returns The container_id of this V1ContainerStatus.

Return type str

image

Gets the image of this V1ContainerStatus. The image the container is running. More info: <https://kubernetes.io/docs/concepts/containers/images>

Returns The image of this V1ContainerStatus.

Return type str

image_id

Gets the image_id of this V1ContainerStatus. ImageID of the container's image.

Returns The image_id of this V1ContainerStatus.

Return type str

last_state

Gets the last_state of this V1ContainerStatus. Details about the container's last termination condition.

Returns The last_state of this V1ContainerStatus.

Return type *V1ContainerState*

name

Gets the name of this V1ContainerStatus. This must be a DNS_LABEL. Each container in a pod must have a unique name. Cannot be updated.

Returns The name of this V1ContainerStatus.

Return type str

ready

Gets the ready of this V1ContainerStatus. Specifies whether the container has passed its readiness probe.

Returns The ready of this V1ContainerStatus.

Return type bool

restart_count

Gets the restart_count of this V1ContainerStatus. The number of times the container has been restarted, currently based on the number of dead containers that have not yet been removed. Note that this is calculated from dead containers. But those containers are subject to garbage collection. This value will get capped at 5 by GC.

Returns The restart_count of this V1ContainerStatus.

Return type int

state

Gets the state of this V1ContainerStatus. Details about the container's current condition.

Returns The state of this V1ContainerStatus.

Return type *V1ContainerState*

swagger_types = {'last_state': 'V1ContainerState', 'image_id': 'str', 'state': 'V1ContainerState', 'name': 'str', 'ready'

```

to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_cross_version_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1CrossVersionObjectReference. API version of the referent

Returns The api_version of this V1CrossVersionObjectReference.

Return type `str`

attribute_map = {'kind': 'kind', 'name': 'name', 'api_version': 'apiVersion'}

kind

Gets the kind of this V1CrossVersionObjectReference. Kind of the referent; More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1CrossVersionObjectReference.

Return type `str`

name

Gets the name of this V1CrossVersionObjectReference. Name of the referent; More info: <http://kubernetes.io/docs/user-guide/identifiers#names>

Returns The name of this V1CrossVersionObjectReference.

Return type `str`

swagger_types = {'kind': 'str', 'name': 'str', 'api_version': 'str'}

```

to_dict()

```

Returns the model properties as a dict

```

to_str()

```

Returns the string representation of the model

kubernetes.client.models.v1_daemon_endpoint module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_daemon_endpoint.V1DaemonEndpoint` (*port=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'port': 'Port'}

port

Gets the port of this V1DaemonEndpoint. Port number of the given endpoint.

Returns The port of this V1DaemonEndpoint.

Return type `int`

swagger_types = {'port': 'int'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_delete_options module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_delete_options.V1DeleteOptions` (*api_version=None, grace_period_seconds=None, kind=None, orphan_dependents=None, preconditions=None, propagation_policy=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1DeleteOptions. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1DeleteOptions.

Return type `str`

attribute_map = {'kind': 'kind', 'propagation_policy': 'propagationPolicy', 'orphan_dependents': 'orphanDependents'}

grace_period_seconds

Gets the grace_period_seconds of this V1DeleteOptions. The duration in seconds before the object should be deleted. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period for the specified type will be used. Defaults to a per object value if not specified. zero means delete immediately.

Returns The `grace_period_seconds` of this `V1DeleteOptions`.

Return type `int`

kind

Gets the kind of this `V1DeleteOptions`. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this `V1DeleteOptions`.

Return type `str`

orphan_dependents

Gets the `orphan_dependents` of this `V1DeleteOptions`. Deprecated: please use the `PropagationPolicy`, this field will be deprecated in 1.7. Should the dependent objects be orphaned. If true/false, the “orphan” finalizer will be added to/removed from the object’s finalizers list. Either this field or `PropagationPolicy` may be set, but not both.

Returns The `orphan_dependents` of this `V1DeleteOptions`.

Return type `bool`

preconditions

Gets the preconditions of this `V1DeleteOptions`. Must be fulfilled before a deletion is carried out. If not possible, a 409 Conflict status will be returned.

Returns The preconditions of this `V1DeleteOptions`.

Return type *`V1Preconditions`*

propagation_policy

Gets the `propagation_policy` of this `V1DeleteOptions`. Whether and how garbage collection will be performed. Either this field or `OrphanDependents` may be set, but not both. The default policy is decided by the existing finalizer set in the metadata.finalizers and the resource-specific default policy.

Returns The `propagation_policy` of this `V1DeleteOptions`.

Return type `str`

swagger_types = {'kind': 'str', 'propagation_policy': 'str', 'orphan_dependents': 'bool', 'preconditions': 'V1Precond

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_downward_api_volume_file module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile (field_ref=None,
mode=None,
path=None,
re-
source_field_r
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'field_ref': 'fieldRef', 'mode': 'mode', 'resource_field_ref': 'resourceFieldRef'}

field_ref

Gets the field_ref of this V1DownwardAPIVolumeFile. Required: Selects a field of the pod: only annotations, labels, name and namespace are supported.

Returns The field_ref of this V1DownwardAPIVolumeFile.

Return type *V1ObjectFieldSelector*

mode

Gets the mode of this V1DownwardAPIVolumeFile. Optional: mode bits to use on this file, must be a value between 0 and 0777. If not specified, the volume defaultMode will be used. This might be in conflict with other options that affect the file mode, like fsGroup, and the result can be other mode bits set.

Returns The mode of this V1DownwardAPIVolumeFile.

Return type int

path

Gets the path of this V1DownwardAPIVolumeFile. Required: Path is the relative path name of the file to be created. Must not be absolute or contain the '..' path. Must be utf-8 encoded. The first item of the relative path must not start with '..'

Returns The path of this V1DownwardAPIVolumeFile.

Return type str

resource_field_ref

Gets the resource_field_ref of this V1DownwardAPIVolumeFile. Selects a resource of the container: only resources limits and requests (limits.cpu, limits.memory, requests.cpu and requests.memory) are currently supported.

Returns The resource_field_ref of this V1DownwardAPIVolumeFile.

Return type *V1ResourceFieldSelector*

swagger_types = {'path': 'str', 'field_ref': 'V1ObjectFieldSelector', 'mode': 'int', 'resource_field_ref': 'V1ResourceF

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_downward_api_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource` (*default_items=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'items': 'items', 'default_mode': 'defaultMode'}

default_mode

Gets the default_mode of this V1DownwardAPIVolumeSource. Optional: mode bits to use on created files by default. Must be a value between 0 and 0777. Defaults to 0644. Directories within the path are not affected by this setting. This might be in conflict with other options that affect the file mode, like fsGroup, and the result can be other mode bits set.

Returns The default_mode of this V1DownwardAPIVolumeSource.

Return type `int`

items

Gets the items of this V1DownwardAPIVolumeSource. Items is a list of downward API volume file

Returns The items of this V1DownwardAPIVolumeSource.

Return type `list[V1DownwardAPIVolumeFile]`

swagger_types = {'items': 'list[V1DownwardAPIVolumeFile]', 'default_mode': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_empty_dir_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_empty_dir_volume_source.V1EmptyDirVolumeSource` (*medium=None, size_limit=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'size_limit': 'sizeLimit', 'medium': 'medium'}

medium

Gets the medium of this V1EmptyDirVolumeSource. What type of storage medium should back this directory. The default is "" which means to use the node's default medium. Must be an empty string (default) or Memory. More info: <https://kubernetes.io/docs/concepts/storage/volumes#emptydir>

Returns The medium of this V1EmptyDirVolumeSource.

Return type `str`

size_limit

Gets the size_limit of this V1EmptyDirVolumeSource. Total amount of local storage required for this EmptyDir volume. The size limit is also applicable for memory medium. The maximum usage on memory medium EmptyDir would be the minimum value between the SizeLimit specified here and the sum of

memory limits of all containers in a pod. The default is nil which means that the limit is undefined. More info: <http://kubernetes.io/docs/user-guide/volumes#emptydir>

Returns The size_limit of this V1EmptyDirVolumeSource.

Return type str

swagger_types = {'size_limit': 'str', 'medium': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_endpoint_address module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_endpoint_address.V1EndpointAddress (hostname=None,
                                                                    ip=None,
                                                                    node_name=None,
                                                                    tar-
                                                                    get_ref=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'target_ref': 'targetRef', 'ip': 'ip', 'hostname': 'hostname', 'node_name': 'nodeName'}

hostname

Gets the hostname of this V1EndpointAddress. The Hostname of this endpoint

Returns The hostname of this V1EndpointAddress.

Return type str

ip

Gets the ip of this V1EndpointAddress. The IP of this endpoint. May not be loopback (127.0.0.0/8), link-local (169.254.0.0/16), or link-local multicast ((224.0.0.0/24). IPv6 is also accepted but not fully supported on all platforms. Also, certain kubernetes components, like kube-proxy, are not IPv6 ready.

Returns The ip of this V1EndpointAddress.

Return type str

node_name

Gets the node_name of this V1EndpointAddress. Optional: Node hosting this endpoint. This can be used to determine endpoints local to a node.

Returns The node_name of this V1EndpointAddress.

Return type str

swagger_types = {'target_ref': 'V1ObjectReference', 'ip': 'str', 'hostname': 'str', 'node_name': 'str'}

target_ref

Gets the target_ref of this V1EndpointAddress. Reference to object providing the endpoint.

Returns The target_ref of this V1EndpointAddress.

Return type *V1ObjectReference*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_endpoint_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_endpoint_port.V1EndpointPort (name=None,
                                                                port=None, protocol=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'protocol': 'protocol', 'name': 'name', 'port': 'port'}

name

Gets the name of this V1EndpointPort. The name of this port (corresponds to ServicePort.Name). Must be a DNS_LABEL. Optional only if one port is defined.

Returns The name of this V1EndpointPort.

Return type str

port

Gets the port of this V1EndpointPort. The port number of the endpoint.

Returns The port of this V1EndpointPort.

Return type int

protocol

Gets the protocol of this V1EndpointPort. The IP protocol for this port. Must be UDP or TCP. Default is TCP.

Returns The protocol of this V1EndpointPort.

Return type str

swagger_types = {'protocol': 'str', 'name': 'str', 'port': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_endpoint_subset module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset (addresses=None,  
not_ready_addresses=None,  
ports=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

addresses

Gets the addresses of this V1EndpointSubset. IP addresses which offer the related ports that are marked as ready. These endpoints should be considered safe for load balancers and clients to utilize.

Returns The addresses of this V1EndpointSubset.

Return type list[V1EndpointAddress]

attribute_map = {'ports': 'ports', 'addresses': 'addresses', 'not_ready_addresses': 'notReadyAddresses'}

not_ready_addresses

Gets the not_ready_addresses of this V1EndpointSubset. IP addresses which offer the related ports but are not currently marked as ready because they have not yet finished starting, have recently failed a readiness check, or have recently failed a liveness check.

Returns The not_ready_addresses of this V1EndpointSubset.

Return type list[V1EndpointAddress]

ports

Gets the ports of this V1EndpointSubset. Port numbers available on the related IP addresses.

Returns The ports of this V1EndpointSubset.

Return type list[V1EndpointPort]

swagger_types = {'ports': 'list[V1EndpointPort]', 'addresses': 'list[V1EndpointAddress]', 'not_ready_addresses': 'lis

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_endpoints module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_endpoints.V1Endpoints (api_version=None,
                                                         kind=None, metadata=None,
                                                         subsets=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Endpoints. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Endpoints.

Return type str

```
attribute_map = {'kind': 'kind', 'subsets': 'subsets', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1Endpoints. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Endpoints.

Return type str

metadata

Gets the metadata of this V1Endpoints. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Endpoints.

Return type *V1ObjectMeta*

subsets

Gets the subsets of this V1Endpoints. The set of all endpoints is the union of all subsets. Addresses are placed into subsets according to the IPs they share. A single address with multiple ports, some of which are ready and some of which are not (because they come from different containers) will result in the address being displayed in different subsets for the different ports. No address will appear in both Addresses and NotReadyAddresses in the same subset. Sets of addresses and ports that comprise a service.

Returns The subsets of this V1Endpoints.

Return type list[*V1EndpointSubset*]

```
swagger_types = {'kind': 'str', 'subsets': 'list[V1EndpointSubset]', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_endpoints_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_endpoints_list.V1EndpointsList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1EndpointsList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1EndpointsList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1EndpointsList. List of endpoints.

Returns The items of this V1EndpointsList.

Return type list[V1Endpoints]

kind

Gets the kind of this V1EndpointsList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1EndpointsList.

Return type str

metadata

Gets the metadata of this V1EndpointsList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1EndpointsList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Endpoints]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_env_var module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>


```
class kubernetes.client.models.v1_env_var.V1EnvVar (name=None, value=None,
                                                    value_from=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'value_from': 'valueFrom', 'name': 'name', 'value': 'value'}
```

name

Gets the name of this V1EnvVar. Name of the environment variable. Must be a C_IDENTIFIER.

Returns The name of this V1EnvVar.

Return type str

```
swagger_types = {'value_from': 'V1EnvVarSource', 'name': 'str', 'value': 'str'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

value

Gets the value of this V1EnvVar. Variable references \$(VAR_NAME) are expanded using the previous defined environment variables in the container and any service environment variables. If a variable cannot be resolved, the reference in the input string will be unchanged. The \$(VAR_NAME) syntax can be escaped with a double \$\$, ie: \$\$\$(VAR_NAME). Escaped references will never be expanded, regardless of whether the variable exists or not. Defaults to "".

Returns The value of this V1EnvVar.

Return type str

value_from

Gets the value_from of this V1EnvVar. Source for the environment variable's value. Cannot be used if value is not empty.

Returns The value_from of this V1EnvVar.

Return type *V1EnvVarSource*

kubernetes.client.models.v1_env_var_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_env_var_source.V1EnvVarSource (config_map_key_ref=None,
                                                                field_ref=None,
                                                                re-
                                                                source_field_ref=None,
                                                                se-
                                                                cret_key_ref=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'secret_key_ref': 'secretKeyRef', 'config_map_key_ref': 'configMapKeyRef', 'field_ref': 'fieldRef'}
```

config_map_key_ref

Gets the config_map_key_ref of this V1EnvVarSource. Selects a key of a ConfigMap.

Returns The config_map_key_ref of this V1EnvVarSource.

Return type *V1ConfigMapKeySelector*

field_ref

Gets the field_ref of this V1EnvVarSource. Selects a field of the pod: supports metadata.name, metadata.namespace, metadata.labels, metadata.annotations, spec.nodeName, spec.serviceAccountName, status.hostIP, status.podIP.

Returns The field_ref of this V1EnvVarSource.

Return type *V1ObjectFieldSelector*

resource_field_ref

Gets the resource_field_ref of this V1EnvVarSource. Selects a resource of the container: only resources limits and requests (limits.cpu, limits.memory, limits.ephemeral-storage, requests.cpu, requests.memory and requests.ephemeral-storage) are currently supported.

Returns The resource_field_ref of this V1EnvVarSource.

Return type *V1ResourceFieldSelector*

secret_key_ref

Gets the secret_key_ref of this V1EnvVarSource. Selects a key of a secret in the pod's namespace

Returns The secret_key_ref of this V1EnvVarSource.

Return type *V1SecretKeySelector*

swagger_types = {'secret_key_ref': 'V1SecretKeySelector', 'config_map_key_ref': 'V1ConfigMapKeySelector', 'field_ref': 'V1ObjectFieldSelector', 'resource_field_ref': 'V1ResourceFieldSelector'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_event module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_event.V1Event (api_version=None, count=None,
first_timestamp=None, involved_object=None, kind=None,
last_timestamp=None, message=None,
metadata=None, reason=None,
source=None, type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Event. APIVersion defines the versioned schema of this representation of an

object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Event.

Return type str

attribute_map = {'last_timestamp': 'lastTimestamp', 'reason': 'reason', 'first_timestamp': 'firstTimestamp', 'message': 'message'}

count

Gets the count of this V1Event. The number of times this event has occurred.

Returns The count of this V1Event.

Return type int

first_timestamp

Gets the first_timestamp of this V1Event. The time at which the event was first recorded. (Time of server receipt is in TypeMeta.)

Returns The first_timestamp of this V1Event.

Return type datetime

involved_object

Gets the involved_object of this V1Event. The object that this event is about.

Returns The involved_object of this V1Event.

Return type *V1ObjectReference*

kind

Gets the kind of this V1Event. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Event.

Return type str

last_timestamp

Gets the last_timestamp of this V1Event. The time at which the most recent occurrence of this event was recorded.

Returns The last_timestamp of this V1Event.

Return type datetime

message

Gets the message of this V1Event. A human-readable description of the status of this operation.

Returns The message of this V1Event.

Return type str

metadata

Gets the metadata of this V1Event. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Event.

Return type *V1ObjectMeta*

reason

Gets the reason of this V1Event. This should be a short, machine understandable string that gives the reason for the transition into the object's current status.

Returns The reason of this V1Event.

Return type str

source

Gets the source of this V1Event. The component reporting this event. Should be a short machine understandable string.

Returns The source of this V1Event.

Return type *V1EventSource*

swagger_types = {'last_timestamp': 'datetime', 'reason': 'str', 'first_timestamp': 'datetime', 'message': 'str', 'count':

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1Event. Type of this event (Normal, Warning), new types could be added in the future

Returns The type of this V1Event.

Return type str

kubernetes.client.models.v1_event_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_event_list.V1EventList (api_version=None,
                                                         items=None,    kind=None,
                                                         metadata=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1EventList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1EventList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1EventList. List of events

Returns The items of this V1EventList.

Return type list[*V1Event*]

kind

Gets the kind of this V1EventList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1EventList.

Return type str

metadata

Gets the metadata of this V1EventList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1EventList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Event]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_event_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_event_source.V1EventSource (*component=None, host=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'host': 'host', 'component': 'component'}

component

Gets the component of this V1EventSource. Component from which the event is generated.

Returns The component of this V1EventSource.

Return type str

host

Gets the host of this V1EventSource. Node name on which the event is generated.

Returns The host of this V1EventSource.

Return type str

swagger_types = {'host': 'str', 'component': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_exec_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_exec_action.V1ExecAction` (*command=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'command': 'command'}

command

Gets the command of this V1ExecAction. Command is the command line to execute inside the container, the working directory for the command is root ('/') in the container's filesystem. The command is simply exec'd, it is not run inside a shell, so traditional shell instructions ('!', etc) won't work. To use a shell, you need to explicitly call out to that shell. Exit status of 0 is treated as live/healthy and non-zero is unhealthy.

Returns The command of this V1ExecAction.

Return type list[str]

swagger_types = {'command': 'list[str]'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_fc_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource` (*fs_type=None, lun=None, read_only=None, target_ww_ns=None, wwids=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'fs_type': 'fsType', 'target_ww_ns': 'targetWWNs', 'lun': 'lun', 'wwids':

fs_type

Gets the fs_type of this V1FCVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Ex. "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified.

Returns The fs_type of this V1FCVolumeSource.

Return type str

lun

Gets the lun of this V1FCVolumeSource. Optional: FC target lun number

Returns The lun of this V1FCVolumeSource.

Return type int

read_only

Gets the read_only of this V1FCVolumeSource. Optional: Defaults to false (read/write). ReadOnly here will force the ReadOnly setting in VolumeMounts.

Returns The read_only of this V1FCVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'fs_type': 'str', 'target_ww_ns': 'list[str]', 'lun': 'int', 'wwids': 'list[str]'}

target_ww_ns

Gets the target_ww_ns of this V1FCVolumeSource. Optional: FC target worldwide names (WWNs)

Returns The target_ww_ns of this V1FCVolumeSource.

Return type list[str]

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

wwids

Gets the wwids of this V1FCVolumeSource. Optional: FC volume world wide identifiers (wwids) Either wwids or combination of targetWWNs and lun must be set, but not both simultaneously.

Returns The wwids of this V1FCVolumeSource.

Return type list[str]

kubernetes.client.models.v1_flex_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource (driver=None,
                                                                    fs_type=None,
                                                                    op-
                                                                    tions=None,
                                                                    read_only=None,
                                                                    se-
                                                                    cret_ref=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'secret_ref': 'secretRef', 'fs_type': 'fsType', 'driver': 'driver', 'options': 'options'}

driver

Gets the driver of this V1FlexVolumeSource. Driver is the name of the driver to use for this volume.

Returns The driver of this V1FlexVolumeSource.

Return type str

fs_type

Gets the fs_type of this V1FlexVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Ex. “ext4”, “xfs”, “ntfs”. The default filesystem depends on FlexVolume script.

Returns The fs_type of this V1FlexVolumeSource.

Return type str

options

Gets the options of this V1FlexVolumeSource. Optional: Extra command options if any.

Returns The options of this V1FlexVolumeSource.

Return type dict(str, str)

read_only

Gets the read_only of this V1FlexVolumeSource. Optional: Defaults to false (read/write). ReadOnly here will force the ReadOnly setting in VolumeMounts.

Returns The read_only of this V1FlexVolumeSource.

Return type bool

secret_ref

Gets the secret_ref of this V1FlexVolumeSource. Optional: SecretRef is reference to the secret object containing sensitive information to pass to the plugin scripts. This may be empty if no secret object is specified. If the secret object contains more than one secret, all secrets are passed to the plugin scripts.

Returns The secret_ref of this V1FlexVolumeSource.

Return type *V1LocalObjectReference*

swagger_types = {'read_only': 'bool', 'secret_ref': 'V1LocalObjectReference', 'fs_type': 'str', 'driver': 'str', 'options': 'dict(str, str)'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_flocker_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource` (*dataset_name=None, dataset_uuid=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'dataset_uuid': 'datasetUUID', 'dataset_name': 'datasetName'}

dataset_name

Gets the dataset_name of this V1FlockerVolumeSource. Name of the dataset stored as metadata -> name on the dataset for Flocker should be considered as deprecated

Returns The dataset_name of this V1FlockerVolumeSource.

Return type str

dataset_uuid

Gets the dataset_uuid of this V1FlockerVolumeSource. UUID of the dataset. This is unique identifier of a Flocker dataset

Returns The dataset_uuid of this V1FlockerVolumeSource.

Return type str

swagger_types = {'dataset_uuid': 'str', 'dataset_name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_gce_persistent_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'pd_name': 'pdName', 'fs_type': 'fsType', 'partition': 'partition'}

fs_type

Gets the fs_type of this V1GCEPersistentDiskVolumeSource. Filesystem type of the volume that you want to mount. Tip: Ensure that the filesystem type is supported by the host operating system. Examples: "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified. More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The fs_type of this V1GCEPersistentDiskVolumeSource.

Return type str

partition

Gets the partition of this V1GCEPersistentDiskVolumeSource. The partition in the volume that you want to mount. If omitted, the default is to mount by volume name. Examples: For volume /dev/sda1, you specify the partition as "1". Similarly, the volume partition for /dev/sda is "0" (or you can leave the property empty). More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The partition of this V1GCEPersistentDiskVolumeSource.

Return type int

pd_name

Gets the pd_name of this V1GCEPersistentDiskVolumeSource. Unique name of the PD resource in GCE. Used to identify the disk in GCE. More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The pd_name of this V1GCEPersistentDiskVolumeSource.

Return type str

read_only

Gets the read_only of this V1GCEPersistentDiskVolumeSource. ReadOnly here will force the Read-Only setting in VolumeMounts. Defaults to false. More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The read_only of this V1GCEPersistentDiskVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'pd_name': 'str', 'fs_type': 'str', 'partition': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_git_repo_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource (directory=None,
                                                                                   repos-
                                                                                   i-
                                                                                   tory=None,
                                                                                   re-
                                                                                   vi-
                                                                                   sion=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'directory': 'directory', 'repository': 'repository', 'revision': 'revision'}

directory

Gets the directory of this V1GitRepoVolumeSource. Target directory name. Must not contain or start with '..'. If '.' is supplied, the volume directory will be the git repository. Otherwise, if specified, the volume will contain the git repository in the subdirectory with the given name.

Returns The directory of this V1GitRepoVolumeSource.

Return type str

repository

Gets the repository of this V1GitRepoVolumeSource. Repository URL

Returns The repository of this V1GitRepoVolumeSource.

Return type str

revision

Gets the revision of this V1GitRepoVolumeSource. Commit hash for the specified revision.

Returns The revision of this V1GitRepoVolumeSource.

Return type str

swagger_types = {'directory': 'str', 'repository': 'str', 'revision': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_glusterfs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterfsVolumeSource (*endpoints=None, path=None, read_only=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'path': 'path', 'endpoints': 'endpoints'}

endpoints

Gets the endpoints of this V1GlusterfsVolumeSource. EndpointsName is the endpoint name that details Glusterfs topology. More info: <https://releases.k8s.io/HEAD/examples/volumes/glusterfs/README.md#create-a-pod>

Returns The endpoints of this V1GlusterfsVolumeSource.

Return type str

path

Gets the path of this V1GlusterfsVolumeSource. Path is the Glusterfs volume path. More info: <https://releases.k8s.io/HEAD/examples/volumes/glusterfs/README.md#create-a-pod>

Returns The path of this V1GlusterfsVolumeSource.

Return type str

read_only

Gets the read_only of this V1GlusterfsVolumeSource. ReadOnly here will force the Glusterfs volume to be mounted with read-only permissions. Defaults to false. More info: <https://releases.k8s.io/HEAD/examples/volumes/glusterfs/README.md#create-a-pod>

Returns The read_only of this V1GlusterfsVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'path': 'str', 'endpoints': 'str'}

```

to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_handler module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.client.models.v1_handler.V1Handler(_exec=None,      http_get=None,
                                                    tcp_socket=None)

```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```

attribute_map = {'_exec': 'exec', 'http_get': 'httpGet', 'tcp_socket': 'tcpSocket'}

```

http_get

Gets the http_get of this V1Handler. HTTPGet specifies the http request to perform.

Returns The http_get of this V1Handler.

Return type *V1HTTPGetAction*

```

swagger_types = {'_exec': 'V1ExecAction', 'http_get': 'V1HTTPGetAction', 'tcp_socket': 'V1TCPSocketAction'}

```

tcp_socket

Gets the tcp_socket of this V1Handler. TCP socket specifies an action involving a TCP port. TCP hooks not yet supported

Returns The tcp_socket of this V1Handler.

Return type *V1TCPSocketAction*

```

to_dict()

```

Returns the model properties as a dict

```

to_str()

```

Returns the string representation of the model

kubernetes.client.models.v1_horizontal_pod_autoscaler module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler (api_version=api_version,
kind=Non,
meta-data=Non,
spec=Non,
sta-tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1HorizontalPodAutoscaler. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1HorizontalPodAutoscaler.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1HorizontalPodAutoscaler. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1HorizontalPodAutoscaler.

Return type str

metadata

Gets the metadata of this V1HorizontalPodAutoscaler. Standard object metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1HorizontalPodAutoscaler.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1HorizontalPodAutoscaler. behaviour of autoscaler. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>.

Returns The spec of this V1HorizontalPodAutoscaler.

Return type *V1HorizontalPodAutoscalerSpec*

status

Gets the status of this V1HorizontalPodAutoscaler. current information about the autoscaler.

Returns The status of this V1HorizontalPodAutoscaler.

Return type *V1HorizontalPodAutoscalerStatus*

swagger_types = {'status': 'V1HorizontalPodAutoscalerStatus', 'kind': 'str', 'spec': 'V1HorizontalPodAutoscalerSpec'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_horizontal_pod_autoscaler_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1HorizontalPodAutoscalerList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1HorizontalPodAutoscalerList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1HorizontalPodAutoscalerList. list of horizontal pod autoscaler objects.

Returns The items of this V1HorizontalPodAutoscalerList.

Return type list[V1HorizontalPodAutoscaler]

kind

Gets the kind of this V1HorizontalPodAutoscalerList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1HorizontalPodAutoscalerList.

Return type str

metadata

Gets the metadata of this V1HorizontalPodAutoscalerList. Standard list metadata.

Returns The metadata of this V1HorizontalPodAutoscalerList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1HorizontalPodAutoscaler]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_horizontal_pod_autoscaler_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'scale_target_ref': 'scaleTargetRef', 'min_replicas': 'minReplicas', 'target_cpu_utilization_percent

max_replicas

Gets the max_replicas of this V1HorizontalPodAutoscalerSpec. upper limit for the number of pods that can be set by the autoscaler; cannot be smaller than MinReplicas.

Returns The max_replicas of this V1HorizontalPodAutoscalerSpec.

Return type `int`

min_replicas

Gets the min_replicas of this V1HorizontalPodAutoscalerSpec. lower limit for the number of pods that can be set by the autoscaler, default 1.

Returns The min_replicas of this V1HorizontalPodAutoscalerSpec.

Return type `int`

scale_target_ref

Gets the scale_target_ref of this V1HorizontalPodAutoscalerSpec. reference to scaled resource; horizontal pod autoscaler will learn the current resource consumption and will set the desired number of pods by using its Scale subresource.

Returns The scale_target_ref of this V1HorizontalPodAutoscalerSpec.

Return type *V1CrossVersionObjectReference*

swagger_types = {'scale_target_ref': 'V1CrossVersionObjectReference', 'min_replicas': 'int', 'target_cpu_utilization

target_cpu_utilization_percentage

Gets the target_cpu_utilization_percentage of this V1HorizontalPodAutoscalerSpec. target average CPU utilization (represented as a percentage of requested CPU) over all the pods; if not specified the default autoscaling policy will be used.

Returns The target_cpu_utilization_percentage of this V1HorizontalPodAutoscalerSpec.

Return type `int`

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_horizontal_pod_autoscaler_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerS

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'observed_generation': 'observedGeneration', 'last_scale_time': 'lastScaleTime', 'current_cpu_utili

current_cpu_utilization_percentage

Gets the current_cpu_utilization_percentage of this V1HorizontalPodAutoscalerStatus. current average CPU utilization over all pods, represented as a percentage of requested CPU, e.g. 70 means that an average pod is using now 70% of its requested CPU.

Returns The current_cpu_utilization_percentage of this V1HorizontalPodAutoscalerStatus.

Return type int

current_replicas

Gets the current_replicas of this V1HorizontalPodAutoscalerStatus. current number of replicas of pods managed by this autoscaler.

Returns The current_replicas of this V1HorizontalPodAutoscalerStatus.

Return type int

desired_replicas

Gets the desired_replicas of this V1HorizontalPodAutoscalerStatus. desired number of replicas of pods managed by this autoscaler.

Returns The desired_replicas of this V1HorizontalPodAutoscalerStatus.

Return type int

last_scale_time

Gets the last_scale_time of this V1HorizontalPodAutoscalerStatus. last time the HorizontalPodAutoscaler scaled the number of pods; used by the autoscaler to control how often the number of pods is changed.

Returns The last_scale_time of this V1HorizontalPodAutoscalerStatus.

Return type datetime

observed_generation

Gets the observed_generation of this V1HorizontalPodAutoscalerStatus. most recent generation observed by this autoscaler.

Returns The observed_generation of this V1HorizontalPodAutoscalerStatus.

Return type int


```
swagger_types = {'observed_generation': 'int', 'last_scale_time': 'datetime', 'current_cpu_utilization_percentage': 'int'}
```

to_dict()
Returns the model properties as a dict

to_str()
Returns the string representation of the model

kubernetes.client.models.v1_host_path_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource` (*path=None, type=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'type': 'type'}

path

Gets the path of this V1HostPathVolumeSource. Path of the directory on the host. If the path is a symlink, it will follow the link to the real path. More info: <https://kubernetes.io/docs/concepts/storage/volumes#hostpath>

Returns The path of this V1HostPathVolumeSource.

Return type `str`

swagger_types = {'path': 'str', 'type': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1HostPathVolumeSource. Type for HostPath Volume Defaults to "" More info: <https://kubernetes.io/docs/concepts/storage/volumes#hostpath>

Returns The type of this V1HostPathVolumeSource.

Return type `str`

kubernetes.client.models.v1_http_get_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_http_get_action.V1HTTPGetAction (host=None,
                                                                    http_headers=None,
                                                                    path=None,
                                                                    port=None,
                                                                    scheme=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'host': 'host', 'scheme': 'scheme', 'port': 'port', 'http_headers': 'httpHeaders'}

host

Gets the host of this V1HTTPGetAction. Host name to connect to, defaults to the pod IP. You probably want to set “Host” in httpHeaders instead.

Returns The host of this V1HTTPGetAction.

Return type str

http_headers

Gets the http_headers of this V1HTTPGetAction. Custom headers to set in the request. HTTP allows repeated headers.

Returns The http_headers of this V1HTTPGetAction.

Return type list[V1HTTPHeader]

path

Gets the path of this V1HTTPGetAction. Path to access on the HTTP server.

Returns The path of this V1HTTPGetAction.

Return type str

port

Gets the port of this V1HTTPGetAction. Name or number of the port to access on the container. Number must be in the range 1 to 65535. Name must be an IANA_SVC_NAME.

Returns The port of this V1HTTPGetAction.

Return type object

scheme

Gets the scheme of this V1HTTPGetAction. Scheme to use for connecting to the host. Defaults to HTTP.

Returns The scheme of this V1HTTPGetAction.

Return type str

swagger_types = {'path': 'str', 'host': 'str', 'scheme': 'str', 'port': 'object', 'http_headers': 'list[V1HTTPHeader]'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_http_header module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_http_header.V1HTTPHeader (name=None,
                                                           value=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'name': 'name', 'value': 'value'}

name

Gets the name of this V1HTTPHeader. The header field name

Returns The name of this V1HTTPHeader.

Return type str

swagger_types = {'name': 'str', 'value': 'str'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

value

Gets the value of this V1HTTPHeader. The header field value

Returns The value of this V1HTTPHeader.

Return type str

kubernetes.client.models.v1_iscsi_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource (chap_auth_discovery=None,
                                                                              chap_auth_session=None,
                                                                              fs_type=None,
                                                                              initiator_name=None,
                                                                              iqn=None,
                                                                              iscsi_interface=None,
                                                                              lun=None,
                                                                              portals=None,
                                                                              read_only=None,
                                                                              secret_ref=None,
                                                                              target_portal=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'portals': 'portals', 'secret_ref': 'secretRef', 'fs_type': 'fsType', 'target_p

chap_auth_discovery

Gets the chap_auth_discovery of this V1IscsiVolumeSource. whether support iSCSI Discovery CHAP authentication

Returns The chap_auth_discovery of this V1IscsiVolumeSource.

Return type bool

chap_auth_session

Gets the chap_auth_session of this V1IscsiVolumeSource. whether support iSCSI Session CHAP authentication

Returns The chap_auth_session of this V1IscsiVolumeSource.

Return type bool

fs_type

Gets the fs_type of this V1IscsiVolumeSource. Filesystem type of the volume that you want to mount. Tip: Ensure that the filesystem type is supported by the host operating system. Examples: "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified. More info: <https://kubernetes.io/docs/concepts/storage/volumes#iscsi>

Returns The fs_type of this V1IscsiVolumeSource.

Return type str

initiator_name

Gets the initiator_name of this V1IscsiVolumeSource. Custom iSCSI initiator name. If initiatorName is specified with iscsiInterface simultaneously, new iSCSI interface <target portal>:<volume name> will be created for the connection.

Returns The initiator_name of this V1IscsiVolumeSource.

Return type str

iqn

Gets the iqn of this V1IscsiVolumeSource. Target iSCSI Qualified Name.

Returns The iqn of this V1IscsiVolumeSource.

Return type str

iscsi_interface

Gets the iscsi_interface of this V1IscsiVolumeSource. Optional: Defaults to 'default' (tcp). iSCSI interface name that uses an iSCSI transport.

Returns The iscsi_interface of this V1IscsiVolumeSource.

Return type str

lun

Gets the lun of this V1IscsiVolumeSource. iSCSI target lun number.

Returns The lun of this V1IscsiVolumeSource.

Return type int

portals

Gets the portals of this V1IscsiVolumeSource. iSCSI target portal List. The portal is either an IP or ip_addr:port if the port is other than default (typically TCP ports 860 and 3260).

Returns The portals of this V1IscsiVolumeSource.

Return type list[str]

read_only

Gets the read_only of this V1ISCSIVolumeSource. ReadOnly here will force the ReadOnly setting in VolumeMounts. Defaults to false.

Returns The read_only of this V1ISCSIVolumeSource.

Return type bool

secret_ref

Gets the secret_ref of this V1ISCSIVolumeSource. CHAP secret for iSCSI target and initiator authentication

Returns The secret_ref of this V1ISCSIVolumeSource.

Return type *V1LocalObjectReference*

swagger_types = {'read_only': 'bool', 'portals': 'list[str]', 'secret_ref': 'V1LocalObjectReference', 'fs_type': 'str', 'target_portal': 'str'}

target_portal

Gets the target_portal of this V1ISCSIVolumeSource. iSCSI target portal. The portal is either an IP or ip_addr:port if the port is other than default (typically TCP ports 860 and 3260).

Returns The target_portal of this V1ISCSIVolumeSource.

Return type str

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_job module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_job.V1Job` (*api_version=None, kind=None, metadata=None, spec=None, status=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Job. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Job.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1Job. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Job.

Return type str

metadata

Gets the metadata of this V1Job. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Job.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Job. Specification of the desired behavior of a job. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1Job.

Return type *V1JobSpec*

status

Gets the status of this V1Job. Current status of a job. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1Job.

Return type *V1JobStatus*

swagger_types = {'status': 'V1JobStatus', 'kind': 'str', 'spec': 'V1JobSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_job_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_job_condition.V1JobCondition(last_probe_time=None,
                                                                last_transition_time=None,
                                                                message=None,
                                                                reason=None,
                                                                status=None,
                                                                type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'status': 'status', 'last_transition_time': 'lastTransitionTime', 'reason': 'reason', 'message': 'message'}

last_probe_time

Gets the last_probe_time of this V1JobCondition. Last time the condition was checked.

Returns The last_probe_time of this V1JobCondition.

Return type datetime

last_transition_time

Gets the last_transition_time of this V1JobCondition. Last time the condition transit from one status to another.

Returns The last_transition_time of this V1JobCondition.

Return type datetime

message

Gets the message of this V1JobCondition. Human readable message indicating details about last transition.

Returns The message of this V1JobCondition.

Return type str

reason

Gets the reason of this V1JobCondition. (brief) reason for the condition's last transition.

Returns The reason of this V1JobCondition.

Return type str

status

Gets the status of this V1JobCondition. Status of the condition, one of True, False, Unknown.

Returns The status of this V1JobCondition.

Return type str

swagger_types = {'status': 'str', 'last_transition_time': 'datetime', 'reason': 'str', 'message': 'str', 'type': 'str', 'last_

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1JobCondition. Type of job condition, Complete or Failed.

Returns The type of this V1JobCondition.

Return type str

kubernetes.client.models.v1_job_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_job_list.V1JobList (*api_version=None, items=None, kind=None, metadata=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1JobList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1JobList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1JobList. items is the list of Jobs.

Returns The items of this V1JobList.

Return type list[V1Job]

kind

Gets the kind of this V1JobList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1JobList.

Return type str

metadata

Gets the metadata of this V1JobList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1JobList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Job]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_job_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_job_spec.V1JobSpec (active_deadline_seconds=None,
                                                       backoff_limit=None,           com-
                                                       pletions=None,             man-
                                                       ual_selector=None,        paral-
                                                       lelism=None,              selector=None,
                                                       template=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

active_deadline_seconds

Gets the active_deadline_seconds of this V1JobSpec. Specifies the duration in seconds relative to the startTime that the job may be active before the system tries to terminate it; value must be positive integer

Returns The active_deadline_seconds of this V1JobSpec.

Return type int

attribute_map = {'backoff_limit': 'backoffLimit', 'completions': 'completions', 'manual_selector': 'manualSelector',

backoff_limit

Gets the backoff_limit of this V1JobSpec. Specifies the number of retries before marking this job failed. Defaults to 6

Returns The backoff_limit of this V1JobSpec.

Return type int

completions

Gets the completions of this V1JobSpec. Specifies the desired number of successfully finished pods the job should be run with. Setting to nil means that the success of any pod signals the success of all pods, and allows parallelism to have any positive value. Setting to 1 means that parallelism is limited to 1 and the success of that pod signals the success of the job. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>

Returns The completions of this V1JobSpec.

Return type int

manual_selector

Gets the manual_selector of this V1JobSpec. manualSelector controls generation of pod labels and pod selectors. Leave *manualSelector* unset unless you are certain what you are doing. When false or unset, the system pick labels unique to this job and appends those labels to the pod template. When true, the user is responsible for picking unique labels and specifying the selector. Failure to pick a unique label may cause this and other jobs to not function correctly. However, You may see *manualSelector=true* in jobs that were created with the old *extensions/v1beta1* API. More info: <https://git.k8s.io/community/contributors/design-proposals/selector-generation.md>

Returns The manual_selector of this V1JobSpec.

Return type bool

parallelism

Gets the parallelism of this V1JobSpec. Specifies the maximum desired number of pods the job should run at any given time. The actual number of pods running in steady state will be less than this number when $((\text{spec.completions} - \text{.status.successful}) < \text{.spec.parallelism})$, i.e. when the work left to do is less than max parallelism. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>

Returns The parallelism of this V1JobSpec.

Return type int

selector

Gets the selector of this V1JobSpec. A label query over pods that should match the pod count. Normally, the system sets this field for you. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#label-selectors>

Returns The selector of this V1JobSpec.

Return type V1LabelSelector

swagger_types = {'backoff_limit': 'int', 'completions': 'int', 'manual_selector': 'bool', 'selector': 'V1LabelSelector',

template

Gets the template of this V1JobSpec. Describes the pod that will be created when executing a job. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>

Returns The template of this V1JobSpec.

Return type *V1PodTemplateSpec*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_job_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_job_status.V1JobStatus (active=None, completion_time=None, conditions=None, failed=None, start_time=None, succeeded=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

active

Gets the active of this V1JobStatus. The number of actively running pods.

Returns The active of this V1JobStatus.

Return type int

attribute_map = {'succeeded': 'succeeded', 'failed': 'failed', 'start_time': 'startTime', 'completion_time': 'completionTime'}

completion_time

Gets the completion_time of this V1JobStatus. Represents time when the job was completed. It is not guaranteed to be set in happens-before order across separate operations. It is represented in RFC3339 form and is in UTC.

Returns The completion_time of this V1JobStatus.

Return type datetime

conditions

Gets the conditions of this V1JobStatus. The latest available observations of an object's current state. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>

Returns The conditions of this V1JobStatus.

Return type list[V1JobCondition]

failed

Gets the failed of this V1JobStatus. The number of pods which reached phase Failed.

Returns The failed of this V1JobStatus.

Return type int

start_time

Gets the start_time of this V1JobStatus. Represents time when the job was acknowledged by the job controller. It is not guaranteed to be set in happens-before order across separate operations. It is represented in RFC3339 form and is in UTC.

Returns The start_time of this V1JobStatus.

Return type datetime

succeeded

Gets the succeeded of this V1JobStatus. The number of pods which reached phase Succeeded.

Returns The succeeded of this V1JobStatus.

Return type int

swagger_types = {'succeeded': 'int', 'failed': 'int', 'start_time': 'datetime', 'completion_time': 'datetime', 'active': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_key_to_path module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_key_to_path.V1KeyToPath(*key=None, mode=None, path=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'mode': 'mode', 'key': 'key'}

key

Gets the key of this V1KeyToPath. The key to project.

Returns The key of this V1KeyToPath.

Return type str

mode

Gets the mode of this V1KeyToPath. Optional: mode bits to use on this file, must be a value between 0 and 0777. If not specified, the volume defaultMode will be used. This might be in conflict with other options that affect the file mode, like fsGroup, and the result can be other mode bits set.

Returns The mode of this V1KeyToPath.

Return type int

path

Gets the path of this V1KeyToPath. The relative path of the file to map the key to. May not be an absolute path. May not contain the path element '..'. May not start with the string '..'.

Returns The path of this V1KeyToPath.

Return type str

swagger_types = {'path': 'str', 'mode': 'int', 'key': 'str'}

to_dict()

Returns the model properties as a dict

```
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1_lifecycle module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_lifecycle.V1Lifecycle (post_start=None,
                                                         pre_stop=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'post_start': 'postStart', 'pre_stop': 'preStop'}
```

post_start

Gets the post_start of this V1Lifecycle. PostStart is called immediately after a container is created. If the handler fails, the container is terminated and restarted according to its restart policy. Other management of the container blocks until the hook completes. More info: <https://kubernetes.io/docs/concepts/containers/container-lifecycle-hooks/#container-hooks>

Returns The post_start of this V1Lifecycle.

Return type *V1Handler*

pre_stop

Gets the pre_stop of this V1Lifecycle. PreStop is called immediately before a container is terminated. The container is terminated after the handler completes. The reason for termination is passed to the handler. Regardless of the outcome of the handler, the container is eventually terminated. Other management of the container blocks until the hook completes. More info: <https://kubernetes.io/docs/concepts/containers/container-lifecycle-hooks/#container-hooks>

Returns The pre_stop of this V1Lifecycle.

Return type *V1Handler*

```
swagger_types = {'post_start': 'V1Handler', 'pre_stop': 'V1Handler'}
```

```
to_dict()
    Returns the model properties as a dict
```

```
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1_limit_range module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_limit_range.V1LimitRange (api_version=None,
                                                            kind=None,      meta-
                                                            data=None, spec=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1LimitRange. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1LimitRange.

Return type str

```
attribute_map = {'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1LimitRange. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1LimitRange.

Return type str

metadata

Gets the metadata of this V1LimitRange. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1LimitRange.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1LimitRange. Spec defines the limits enforced. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1LimitRange.

Return type *V1LimitRangeSpec*

```
swagger_types = {'kind': 'str', 'spec': 'V1LimitRangeSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_limit_range_item module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem (default=None,
                                                                    de-
                                                                    fault_request=None,
                                                                    max=None,
                                                                    max_limit_request_ratio=None,
                                                                    min=None,
                                                                    type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'default_request': 'defaultRequest', 'min': 'min', 'default': 'default', 'max': 'max', 'max_limit_req

default

Gets the default of this V1LimitRangeItem. Default resource requirement limit value by resource name if resource limit is omitted.

Returns The default of this V1LimitRangeItem.

Return type dict(str, str)

default_request

Gets the default_request of this V1LimitRangeItem. DefaultRequest is the default resource requirement request value by resource name if resource request is omitted.

Returns The default_request of this V1LimitRangeItem.

Return type dict(str, str)

max

Gets the max of this V1LimitRangeItem. Max usage constraints on this kind by resource name.

Returns The max of this V1LimitRangeItem.

Return type dict(str, str)

max_limit_request_ratio

Gets the max_limit_request_ratio of this V1LimitRangeItem. MaxLimitRequestRatio if specified, the named resource must have a request and limit that are both non-zero where limit divided by request is less than or equal to the enumerated value; this represents the max burst for the named resource.

Returns The max_limit_request_ratio of this V1LimitRangeItem.

Return type dict(str, str)

min

Gets the min of this V1LimitRangeItem. Min usage constraints on this kind by resource name.

Returns The min of this V1LimitRangeItem.

Return type dict(str, str)

swagger_types = {'default_request': 'dict(str, str)', 'min': 'dict(str, str)', 'default': 'dict(str, str)', 'max': 'dict(str, str)'

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1LimitRangeItem. Type of resource that this limit applies to.

Returns The type of this V1LimitRangeItem.

Return type str

kubernetes.client.models.v1_limit_range_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_limit_range_list.V1LimitRangeList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1LimitRangeList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1LimitRangeList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1LimitRangeList. Items is a list of LimitRange objects. More info: https://git.k8s.io/community/contributors/design-proposals/admission_control_limit_range.md

Returns The items of this V1LimitRangeList.

Return type list[V1LimitRange]

kind

Gets the kind of this V1LimitRangeList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1LimitRangeList.

Return type str

metadata

Gets the metadata of this V1LimitRangeList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1LimitRangeList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1LimitRange]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_limit_range_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec` (*limits=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'limits': 'limits'}

limits

Gets the limits of this V1LimitRangeSpec. Limits is the list of LimitRangeItem objects that are enforced.

Returns The limits of this V1LimitRangeSpec.

Return type `list[V1LimitRangeItem]`

swagger_types = {'limits': 'list[V1LimitRangeItem]'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_load_balancer_ingress module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_load_balancer_ingress.V1LoadBalancerIngress` (*hostname=None, ip=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'ip': 'ip', 'hostname': 'hostname'}

hostname

Gets the hostname of this V1LoadBalancerIngress. Hostname is set for load-balancer ingress points that are DNS based (typically AWS load-balancers)

Returns The hostname of this V1LoadBalancerIngress.

Return type `str`

ip

Gets the ip of this V1LoadBalancerIngress. IP is set for load-balancer ingress points that are IP based (typically GCE or OpenStack load-balancers)

Returns The ip of this V1LoadBalancerIngress.

Return type `str`


```
swagger_types = {'ip': 'str', 'hostname': 'str'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.v1_load_balancer_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus` (*ingress=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'ingress': 'ingress'}
```

ingress

Gets the ingress of this V1LoadBalancerStatus. Ingress is a list containing ingress points for the load-balancer. Traffic intended for the service should be sent to these ingress points.

Returns The ingress of this V1LoadBalancerStatus.

Return type `list[V1LoadBalancerIngress]`

```
swagger_types = {'ingress': 'list[V1LoadBalancerIngress]'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.v1_local_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_local_object_reference.V1LocalObjectReference` (*name=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'name': 'name'}
```

name

Gets the name of this V1LocalObjectReference. Name of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1LocalObjectReference.

Return type str

swagger_types = {'name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_namespace module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_namespace.V1Namespace (api_version=None,
                                                         kind=None, metadata=None,
                                                         spec=None, status=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Namespace. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Namespace.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1Namespace. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Namespace.

Return type str

metadata

Gets the metadata of this V1Namespace. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Namespace.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Namespace. Spec defines the behavior of the Namespace. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1Namespace.

Return type *V1NamespaceSpec*

status

Gets the status of this V1Namespace. Status describes the current status of a Namespace. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1Namespace.

Return type *V1NamespaceStatus*

swagger_types = {'status': 'V1NamespaceStatus', 'kind': 'str', 'spec': 'V1NamespaceSpec', 'api_version': 'str', 'meta

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_namespace_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_namespace_list.V1NamespaceList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1NamespaceList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1NamespaceList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1NamespaceList. Items is the list of Namespace objects in the list. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/namespaces/>

Returns The items of this V1NamespaceList.

Return type list[V1Namespace]

kind

Gets the kind of this V1NamespaceList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1NamespaceList.

Return type str

metadata

Gets the metadata of this V1NamespaceList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1NamespaceList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Namespace]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_namespace_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_namespace_spec.V1NamespaceSpec` (*finalizers=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'finalizers': 'finalizers'}

finalizers

Gets the finalizers of this V1NamespaceSpec. Finalizers is an opaque list of values that must be empty to permanently remove object from storage. More info: <https://git.k8s.io/community/contributors/design-proposals/namespaces.md#finalizers>

Returns The finalizers of this V1NamespaceSpec.

Return type list[str]

swagger_types = {'finalizers': 'list[str]'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_namespace_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_namespace_status.V1NamespaceStatus (phase=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'phase': 'phase'}
```

phase

Gets the phase of this V1NamespaceStatus. Phase is the current lifecycle phase of the namespace. More info: <https://git.k8s.io/community/contributors/design-proposals/namespaces.md#phases>

Returns The phase of this V1NamespaceStatus.

Return type str

```
swagger_types = {'phase': 'str'}
```

```
to_dict ()
```

Returns the model properties as a dict

```
to_str ()
```

Returns the string representation of the model

kubernetes.client.models.v1_nfs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_nfs_volume_source.V1NFSSource (path=None,
                                                                    read_only=None,
                                                                    server=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'read_only': 'readOnly', 'path': 'path', 'server': 'server'}
```

path

Gets the path of this V1NFSSource. Path that is exported by the NFS server. More info: <https://kubernetes.io/docs/concepts/storage/volumes#nfs>

Returns The path of this V1NFSSource.

Return type str

read_only

Gets the read_only of this V1NFSSource. ReadOnly here will force the NFS export to be mounted with read-only permissions. Defaults to false. More info: <https://kubernetes.io/docs/concepts/storage/volumes#nfs>

Returns The read_only of this V1NFSSource.

Return type bool

server

Gets the server of this V1NFSSource. Server is the hostname or IP address of the NFS server. More info: <https://kubernetes.io/docs/concepts/storage/volumes#nfs>

Returns The server of this V1NFSSource.

Return type str

swagger_types = {'read_only': 'bool', 'path': 'str', 'server': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_node module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_node.V1Node(*api_version=None, kind=None, metadata=None, spec=None, status=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Node. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Node.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1Node. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Node.

Return type str

metadata

Gets the metadata of this V1Node. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Node.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Node. Spec defines the behavior of a node. <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1Node.

Return type *V1NodeSpec*

status

Gets the status of this V1Node. Most recently observed status of the node. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1Node.

Return type *V1NodeStatus*

swagger_types = {'status': 'V1NodeStatus', 'kind': 'str', 'spec': 'V1NodeSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_node_address module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_node_address.V1NodeAddress` (*address=None, type=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

address

Gets the address of this V1NodeAddress. The node address.

Returns The address of this V1NodeAddress.

Return type `str`

attribute_map = {'type': 'type', 'address': 'address'}

swagger_types = {'type': 'str', 'address': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1NodeAddress. Node address type, one of Hostname, ExternalIP or InternalIP.

Returns The type of this V1NodeAddress.

Return type `str`

kubernetes.client.models.v1_node_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_node_condition.V1NodeCondition (last_heartbeat_time=None,
                                                                    last_transition_time=None,
                                                                    message=None,
                                                                    reason=None,
                                                                    status=None,
                                                                    type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'last_heartbeat_time': 'lastHeartbeatTime', 'status': 'status', 'last_transition_time': 'lastTransitionTime'}

last_heartbeat_time

Gets the last_heartbeat_time of this V1NodeCondition. Last time we got an update on a given condition.

Returns The last_heartbeat_time of this V1NodeCondition.

Return type datetime

last_transition_time

Gets the last_transition_time of this V1NodeCondition. Last time the condition transit from one status to another.

Returns The last_transition_time of this V1NodeCondition.

Return type datetime

message

Gets the message of this V1NodeCondition. Human readable message indicating details about last transition.

Returns The message of this V1NodeCondition.

Return type str

reason

Gets the reason of this V1NodeCondition. (brief) reason for the condition's last transition.

Returns The reason of this V1NodeCondition.

Return type str

status

Gets the status of this V1NodeCondition. Status of the condition, one of True, False, Unknown.

Returns The status of this V1NodeCondition.

Return type str

swagger_types = {'last_heartbeat_time': 'datetime', 'status': 'str', 'last_transition_time': 'datetime', 'reason': 'str', 'type': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1NodeCondition. Type of node condition.

Returns The type of this V1NodeCondition.

Return type str

kubernetes.client.models.v1_node_daemon_endpoints module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_node_daemon_endpoints.V1NodeDaemonEndpoints` (*kubelet_endpoint=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'kubelet_endpoint': 'kubeletEndpoint'}

kubelet_endpoint

Gets the kubelet_endpoint of this V1NodeDaemonEndpoints. Endpoint on which Kubelet is listening.

Returns The kubelet_endpoint of this V1NodeDaemonEndpoints.

Return type *V1DaemonEndpoint*

swagger_types = {'kubelet_endpoint': 'V1DaemonEndpoint'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_node_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_node_list.V1NodeList` (*api_version=None, items=None, kind=None, metadata=None*)
 Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1NodeList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1NodeList.

Return type `str`

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1NodeList. List of nodes

Returns The items of this V1NodeList.

Return type `list[V1Node]`

kind

Gets the kind of this V1NodeList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1NodeList.

Return type `str`

metadata

Gets the metadata of this V1NodeList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1NodeList.

Return type `V1ListMeta`

swagger_types = {'items': 'list[V1Node]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_node_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_node_spec.V1NodeSpec(config_source=None, external_id=None, pod_cidr=None, provider_id=None, taints=None, unschedulable=None)
```

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'provider_id': 'providerID', 'taints': 'taints', 'config_source': 'configSource', 'unschedulable': 'unschedulable'}

config_source

Gets the config_source of this V1NodeSpec. If specified, the source to get node configuration from The DynamicKubeletConfig feature gate must be enabled for the Kubelet to use this field

Returns The config_source of this V1NodeSpec.

Return type `V1NodeConfigSource`

external_id

Gets the external_id of this V1NodeSpec. External ID of the node assigned by some machine database (e.g. a cloud provider). Deprecated.

Returns The external_id of this V1NodeSpec.

Return type `str`

`pod_cidr`

Gets the `pod_cidr` of this `V1NodeSpec`. PodCIDR represents the pod IP range assigned to the node.

Returns The `pod_cidr` of this `V1NodeSpec`.

Return type `str`

`provider_id`

Gets the `provider_id` of this `V1NodeSpec`. ID of the node assigned by the cloud provider in the format: `<ProviderName>://<ProviderSpecificNodeID>`

Returns The `provider_id` of this `V1NodeSpec`.

Return type `str`

swagger_types = {'provider_id': 'str', 'taints': 'list[V1Taint]', 'config_source': 'V1NodeConfigSource', 'unschedulable': 'bool'}

`taints`

Gets the taints of this `V1NodeSpec`. If specified, the node's taints.

Returns The taints of this `V1NodeSpec`.

Return type `list[V1Taint]`

`to_dict()`

Returns the model properties as a dict

`to_str()`

Returns the string representation of the model

`unschedulable`

Gets the `unschedulable` of this `V1NodeSpec`. `Unschedulable` controls node schedulability of new pods. By default, node is schedulable. More info: <https://kubernetes.io/docs/concepts/nodes/node/#manual-node-administration>

Returns The `unschedulable` of this `V1NodeSpec`.

Return type `bool`

kubernetes.client.models.v1_node_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_node_status.V1NodeStatus (addresses=None,
                                                             allocatable=None,
                                                             capacity=None,
                                                             conditions=None,
                                                             daemon_endpoints=None,
                                                             images=None,
                                                             node_info=None,
                                                             phase=None,
                                                             volumes_attached=None,
                                                             volumes_in_use=None)
```

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

addresses

Gets the addresses of this V1NodeStatus. List of addresses reachable to the node. Queried from cloud provider, if available. More info: <https://kubernetes.io/docs/concepts/nodes/node/#addresses>

Returns The addresses of this V1NodeStatus.

Return type list[V1NodeAddress]

allocatable

Gets the allocatable of this V1NodeStatus. Allocatable represents the resources of a node that are available for scheduling. Defaults to Capacity.

Returns The allocatable of this V1NodeStatus.

Return type dict(str, str)

attribute_map = {'volumes_attached': 'volumesAttached', 'phase': 'phase', 'node_info': 'nodeInfo', 'daemon_endpoi

capacity

Gets the capacity of this V1NodeStatus. Capacity represents the total resources of a node. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#capacity>

Returns The capacity of this V1NodeStatus.

Return type dict(str, str)

conditions

Gets the conditions of this V1NodeStatus. Conditions is an array of current observed node conditions. More info: <https://kubernetes.io/docs/concepts/nodes/node/#condition>

Returns The conditions of this V1NodeStatus.

Return type list[V1NodeCondition]

daemon_endpoints

Gets the daemon_endpoints of this V1NodeStatus. Endpoints of daemons running on the Node.

Returns The daemon_endpoints of this V1NodeStatus.

Return type V1NodeDaemonEndpoints

images

Gets the images of this V1NodeStatus. List of container images on this node

Returns The images of this V1NodeStatus.

Return type list[V1ContainerImage]

node_info

Gets the node_info of this V1NodeStatus. Set of ids/uuids to uniquely identify the node. More info: <https://kubernetes.io/docs/concepts/nodes/node/#info>

Returns The node_info of this V1NodeStatus.

Return type V1NodeSystemInfo

phase

Gets the phase of this V1NodeStatus. NodePhase is the recently observed lifecycle phase of the node. More info: <https://kubernetes.io/docs/concepts/nodes/node/#phase> The field is never populated, and now is deprecated.

Returns The phase of this V1NodeStatus.

Return type str

swagger_types = {'volumes_attached': 'list[V1AttachedVolume]', 'phase': 'str', 'node_info': 'V1NodeSystemInfo', 'd

to_dict()
Returns the model properties as a dict

to_str()
Returns the string representation of the model

volumes_attached
Gets the volumes_attached of this V1NodeStatus. List of volumes that are attached to the node.
Returns The volumes_attached of this V1NodeStatus.
Return type list[V1AttachedVolume]

volumes_in_use
Gets the volumes_in_use of this V1NodeStatus. List of attachable volumes in use (mounted) by the node.
Returns The volumes_in_use of this V1NodeStatus.
Return type list[str]

kubernetes.client.models.v1_node_system_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo (architecture=None,
                                                                    boot_id=None,
                                                                    con-
                                                                    tainer_runtime_version=None,
                                                                    ker-
                                                                    nel_version=None,
                                                                    kube_proxy_version=None,
                                                                    kubelet_version=None,
                                                                    ma-
                                                                    chine_id=None,
                                                                    operat-
                                                                    ing_system=None,
                                                                    os_image=None,
                                                                    sys-
                                                                    tem_uuid=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

architecture
Gets the architecture of this V1NodeSystemInfo. The Architecture reported by the node
Returns The architecture of this V1NodeSystemInfo.
Return type str

attribute_map = {'kernel_version': 'kernelVersion', 'operating_system': 'operatingSystem', 'os_image': 'osImage', 'a

boot_id
Gets the boot_id of this V1NodeSystemInfo. Boot ID reported by the node.
Returns The boot_id of this V1NodeSystemInfo.

Return type str

container_runtime_version

Gets the container_runtime_version of this V1NodeSystemInfo. ContainerRuntime Version reported by the node through runtime remote API (e.g. docker://1.5.0).

Returns The container_runtime_version of this V1NodeSystemInfo.

Return type str

kernel_version

Gets the kernel_version of this V1NodeSystemInfo. Kernel Version reported by the node from ‘uname -r’ (e.g. 3.16.0-0.bpo.4-amd64).

Returns The kernel_version of this V1NodeSystemInfo.

Return type str

kube_proxy_version

Gets the kube_proxy_version of this V1NodeSystemInfo. KubeProxy Version reported by the node.

Returns The kube_proxy_version of this V1NodeSystemInfo.

Return type str

kubelet_version

Gets the kubelet_version of this V1NodeSystemInfo. Kubelet Version reported by the node.

Returns The kubelet_version of this V1NodeSystemInfo.

Return type str

machine_id

Gets the machine_id of this V1NodeSystemInfo. MachineID reported by the node. For unique machine identification in the cluster this field is preferred. Learn more from man(5) machine-id: <http://man7.org/linux/man-pages/man5/machine-id.5.html>

Returns The machine_id of this V1NodeSystemInfo.

Return type str

operating_system

Gets the operating_system of this V1NodeSystemInfo. The Operating System reported by the node

Returns The operating_system of this V1NodeSystemInfo.

Return type str

os_image

Gets the os_image of this V1NodeSystemInfo. OS Image reported by the node from /etc/os-release (e.g. Debian GNU/Linux 7 (wheezy)).

Returns The os_image of this V1NodeSystemInfo.

Return type str

swagger_types = {'kernel_version': 'str', 'operating_system': 'str', 'os_image': 'str', 'architecture': 'str', 'boot_id': 'str'}

system_uuid

Gets the system_uuid of this V1NodeSystemInfo. SystemUUID reported by the node. For unique machine identification MachineID is preferred. This field is specific to Red Hat hosts https://access.redhat.com/documentation/en-US/Red_Hat_Subscription_Management/1/html/RHSM/getting-system-uuid.html

Returns The system_uuid of this V1NodeSystemInfo.

Return type str

```

to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_object_field_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector (api_version=None,
                                                                                   field_path=None)

```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ObjectFieldSelector. Version of the schema the FieldPath is written in terms of, defaults to "v1".

Returns The api_version of this V1ObjectFieldSelector.

Return type str

attribute_map = {'field_path': 'fieldPath', 'api_version': 'apiVersion'}

field_path

Gets the field_path of this V1ObjectFieldSelector. Path of the field to select in the specified API version.

Returns The field_path of this V1ObjectFieldSelector.

Return type str

swagger_types = {'field_path': 'str', 'api_version': 'str'}

```

to_dict()
    Returns the model properties as a dict

```

```

to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_object_meta module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_object_meta.V1ObjectMeta (annotations=None,
                                                            cluster_name=None, cre-
                                                            ation_timestamp=None,
                                                            dele-
                                                            tion_grace_period_seconds=None,
                                                            dele-
                                                            tion_timestamp=None,
                                                            finalizers=None,   gen-
                                                            erate_name=None,
                                                            generation=None,
                                                            initializers=None,   la-
                                                            bels=None, name=None,
                                                            namespace=None,
                                                            owner_references=None,
                                                            resource_version=None,
                                                            self_link=None,
                                                            uid=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

annotations

Gets the annotations of this V1ObjectMeta. Annotations is an unstructured key value map stored with a resource that may be set by external tools to store and retrieve arbitrary metadata. They are not queryable and should be preserved when modifying objects. More info: <http://kubernetes.io/docs/user-guide/annotations>

Returns The annotations of this V1ObjectMeta.

Return type dict(str, str)

attribute_map = {'name': 'name', 'owner_references': 'ownerReferences', 'generation': 'generation', 'namespace': 'namespace'}

cluster_name

Gets the cluster_name of this V1ObjectMeta. The name of the cluster which the object belongs to. This is used to distinguish resources with same name and namespace in different clusters. This field is not set anywhere right now and apiserver is going to ignore it if set in create or update request.

Returns The cluster_name of this V1ObjectMeta.

Return type str

creation_timestamp

Gets the creation_timestamp of this V1ObjectMeta. CreationTimestamp is a timestamp representing the server time when this object was created. It is not guaranteed to be set in happens-before order across separate operations. Clients may not set this value. It is represented in RFC3339 form and is in UTC. Populated by the system. Read-only. Null for lists. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The creation_timestamp of this V1ObjectMeta.

Return type datetime

deletion_grace_period_seconds

Gets the deletion_grace_period_seconds of this V1ObjectMeta. Number of seconds allowed for this object to gracefully terminate before it will be removed from the system. Only set when deletionTimestamp is also set. May only be shortened. Read-only.

Returns The deletion_grace_period_seconds of this V1ObjectMeta.

Return type int

deletion_timestamp

Gets the `deletion_timestamp` of this `V1ObjectMeta`. `DeletionTimestamp` is RFC 3339 date and time at which this resource will be deleted. This field is set by the server when a graceful deletion is requested by the user, and is not directly settable by a client. The resource is expected to be deleted (no longer visible from resource lists, and not reachable by name) after the time in this field. Once set, this value may not be unset or be set further into the future, although it may be shortened or the resource may be deleted prior to this time. For example, a user may request that a pod is deleted in 30 seconds. The Kubelet will react by sending a graceful termination signal to the containers in the pod. After that 30 seconds, the Kubelet will send a hard termination signal (SIGKILL) to the container and after cleanup, remove the pod from the API. In the presence of network partitions, this object may still exist after this timestamp, until an administrator or automated process can determine the resource is fully terminated. If not set, graceful deletion of the object has not been requested. Populated by the system when a graceful deletion is requested. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The `deletion_timestamp` of this `V1ObjectMeta`.

Return type `datetime`

finalizers

Gets the finalizers of this `V1ObjectMeta`. Must be empty before the object is deleted from the registry. Each entry is an identifier for the responsible component that will remove the entry from the list. If the `deletionTimestamp` of the object is non-nil, entries in this list can only be removed.

Returns The finalizers of this `V1ObjectMeta`.

Return type `list[str]`

generate_name

Gets the `generate_name` of this `V1ObjectMeta`. `GenerateName` is an optional prefix, used by the server, to generate a unique name ONLY IF the `Name` field has not been provided. If this field is used, the name returned to the client will be different than the name passed. This value will also be combined with a unique suffix. The provided value has the same validation rules as the `Name` field, and may be truncated by the length of the suffix required to make the value unique on the server. If this field is specified and the generated name exists, the server will NOT return a 409 - instead, it will either return 201 Created or 500 with Reason `ServerTimeout` indicating a unique name could not be found in the time allotted, and the client should retry (optionally after the time indicated in the `Retry-After` header). Applied only if `Name` is not specified. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#idempotency>

Returns The `generate_name` of this `V1ObjectMeta`.

Return type `str`

generation

Gets the generation of this `V1ObjectMeta`. A sequence number representing a specific generation of the desired state. Populated by the system. Read-only.

Returns The generation of this `V1ObjectMeta`.

Return type `int`

initializers

Gets the initializers of this `V1ObjectMeta`. An initializer is a controller which enforces some system invariant at object creation time. This field is a list of initializers that have not yet acted on this object. If nil or empty, this object has been completely initialized. Otherwise, the object is considered uninitialized and is hidden (in list/watch and get calls) from clients that haven't explicitly asked to observe uninitialized objects. When an object is created, the system will populate this list with the current set of initializers. Only privileged users may set or modify this list. Once it is empty, it may not be modified further by any user.

Returns The initializers of this `V1ObjectMeta`.

Return type V1Initializers

labels

Gets the labels of this V1ObjectMeta. Map of string keys and values that can be used to organize and categorize (scope and select) objects. May match selectors of replication controllers and services. More info: <http://kubernetes.io/docs/user-guide/labels>

Returns The labels of this V1ObjectMeta.

Return type dict(str, str)

name

Gets the name of this V1ObjectMeta. Name must be unique within a namespace. Is required when creating resources, although some resources may allow a client to request the generation of an appropriate name automatically. Name is primarily intended for creation idempotence and configuration definition. Cannot be updated. More info: <http://kubernetes.io/docs/user-guide/identifiers#names>

Returns The name of this V1ObjectMeta.

Return type str

namespace

Gets the namespace of this V1ObjectMeta. Namespace defines the space within each name must be unique. An empty namespace is equivalent to the “default” namespace, but “default” is the canonical representation. Not all objects are required to be scoped to a namespace - the value of this field for those objects will be empty. Must be a DNS_LABEL. Cannot be updated. More info: <http://kubernetes.io/docs/user-guide/namespaces>

Returns The namespace of this V1ObjectMeta.

Return type str

owner_references

Gets the owner_references of this V1ObjectMeta. List of objects depended by this object. If ALL objects in the list have been deleted, this object will be garbage collected. If this object is managed by a controller, then an entry in this list will point to this controller, with the controller field set to true. There cannot be more than one managing controller.

Returns The owner_references of this V1ObjectMeta.

Return type list[V1OwnerReference]

resource_version

Gets the resource_version of this V1ObjectMeta. An opaque value that represents the internal version of this object that can be used by clients to determine when objects have changed. May be used for optimistic concurrency, change detection, and the watch operation on a resource or set of resources. Clients must treat these values as opaque and passed unmodified back to the server. They may only be valid for a particular resource or set of resources. Populated by the system. Read-only. Value must be treated as opaque by clients and . More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#concurrency-control-and-consistency>

Returns The resource_version of this V1ObjectMeta.

Return type str

self_link

Gets the self_link of this V1ObjectMeta. SelfLink is a URL representing this object. Populated by the system. Read-only.

Returns The self_link of this V1ObjectMeta.

Return type str

```
swagger_types = {'name': 'str', 'owner_references': 'list[V1OwnerReference]', 'generation': 'int', 'namespace': 'str'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

```
uid
```

Gets the uid of this V1ObjectMeta. UID is the unique in time and space value for this object. It is typically generated by the server on successful creation of a resource and is not allowed to change on PUT operations. Populated by the system. Read-only. More info: <http://kubernetes.io/docs/user-guide/identifiers#uids>

Returns The uid of this V1ObjectMeta.

Return type str

kubernetes.client.models.v1_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_object_reference.V1ObjectReference (api_version=None,
                                                                    field_path=None,
                                                                    kind=None,
                                                                    name=None,
                                                                    namespace=None,
                                                                    resource_version=None,
                                                                    uid=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
api_version
```

Gets the api_version of this V1ObjectReference. API version of the referent.

Returns The api_version of this V1ObjectReference.

Return type str

```
attribute_map = {'kind': 'kind', 'name': 'name', 'namespace': 'namespace', 'resource_version': 'resourceVersion', 'f
```

```
field_path
```

Gets the field_path of this V1ObjectReference. If referring to a piece of an object instead of an entire object, this string should contain a valid JSON/Go field access statement, such as desired-State.manifest.containers[2]. For example, if the object reference is to a container within a pod, this would take on a value like: "spec.containers{name}" (where "name" refers to the name of the container that triggered the event) or if no container name is specified "spec.containers[2]" (container with index 2 in this pod). This syntax is chosen only to have some well-defined way of referencing a part of an object.

Returns The field_path of this V1ObjectReference.

Return type str

kind

Gets the kind of this V1ObjectReference. Kind of the referent. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ObjectReference.

Return type str

name

Gets the name of this V1ObjectReference. Name of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1ObjectReference.

Return type str

namespace

Gets the namespace of this V1ObjectReference. Namespace of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/namespaces/>

Returns The namespace of this V1ObjectReference.

Return type str

resource_version

Gets the resource_version of this V1ObjectReference. Specific resourceVersion to which this reference is made, if any. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#concurrency-control-and-consistency>

Returns The resource_version of this V1ObjectReference.

Return type str

swagger_types = {'kind': 'str', 'name': 'str', 'namespace': 'str', 'resource_version': 'str', 'field_path': 'str', 'api_version': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

uid

Gets the uid of this V1ObjectReference. UID of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#uids>

Returns The uid of this V1ObjectReference.

Return type str

kubernetes.client.models.v1_owner_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_owner_reference.V1OwnerReference (api_version=None,
                                                                    block_owner_deletion=None,
                                                                    con-
                                                                    troller=None,
                                                                    kind=None,
                                                                    name=None,
                                                                    uid=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1OwnerReference. API version of the referent.

Returns The api_version of this V1OwnerReference.

Return type str

attribute_map = {'kind': 'kind', 'uid': 'uid', 'controller': 'controller', 'block_owner_deletion': 'blockOwnerDeletion'}

block_owner_deletion

Gets the block_owner_deletion of this V1OwnerReference. If true, AND if the owner has the “foregroundDeletion” finalizer, then the owner cannot be deleted from the key-value store until this reference is removed. Defaults to false. To set this field, a user needs “delete” permission of the owner, otherwise 422 (Unprocessable Entity) will be returned.

Returns The block_owner_deletion of this V1OwnerReference.

Return type bool

controller

Gets the controller of this V1OwnerReference. If true, this reference points to the managing controller.

Returns The controller of this V1OwnerReference.

Return type bool

kind

Gets the kind of this V1OwnerReference. Kind of the referent. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1OwnerReference.

Return type str

name

Gets the name of this V1OwnerReference. Name of the referent. More info: <http://kubernetes.io/docs/user-guide/identifiers#names>

Returns The name of this V1OwnerReference.

Return type str

swagger_types = {'kind': 'str', 'uid': 'str', 'controller': 'bool', 'block_owner_deletion': 'bool', 'api_version': 'str', 'name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

uid

Gets the uid of this V1OwnerReference. UID of the referent. More info: <http://kubernetes.io/docs/user-guide/identifiers#uids>

Returns The uid of this V1OwnerReference.

Return type str

kubernetes.client.models.v1_persistent_volume module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_persistent_volume.V1PersistentVolume (api_version=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None,
                                                                    spec=None,
                                                                    sta-
                                                                    tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PersistentVolume. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PersistentVolume.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1PersistentVolume. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PersistentVolume.

Return type str

metadata

Gets the metadata of this V1PersistentVolume. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1PersistentVolume.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1PersistentVolume. Spec defines a specification of a persistent volume owned by the cluster. Provisioned by an administrator. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistent-volumes>

Returns The spec of this V1PersistentVolume.

Return type *V1PersistentVolumeSpec*

status

Gets the status of this V1PersistentVolume. Status represents the current information/status for the persistent volume. Populated by the system. Read-only. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistent-volumes>

Returns The status of this V1PersistentVolume.

Return type *V1PersistentVolumeStatus*

swagger_types = {'status': 'V1PersistentVolumeStatus', 'kind': 'str', 'spec': 'V1PersistentVolumeSpec', 'api_version':

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_claim module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim (api_version=None,
                                          kind=None,
                                          meta-
                                          data=None,
                                          spec=None,
                                          sta-
                                          tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PersistentVolumeClaim. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PersistentVolumeClaim.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1PersistentVolumeClaim. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PersistentVolumeClaim.

Return type str

metadata

Gets the metadata of this V1PersistentVolumeClaim. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1PersistentVolumeClaim.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1PersistentVolumeClaim. Spec defines the desired characteristics of a volume requested by a pod author. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistentvolumeclaims>

Returns The spec of this V1PersistentVolumeClaim.

Return type *V1PersistentVolumeClaimSpec*

status

Gets the status of this V1PersistentVolumeClaim. Status represents the current information/status of a persistent volume claim. Read-only. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistentvolumeclaims>

Returns The status of this V1PersistentVolumeClaim.

Return type *V1PersistentVolumeClaimStatus*

swagger_types = {'status': 'V1PersistentVolumeClaimStatus', 'kind': 'str', 'spec': 'V1PersistentVolumeClaimSpec', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_claim_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList` (*ap*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PersistentVolumeClaimList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PersistentVolumeClaimList.

Return type `str`


```
attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

items

Gets the items of this V1PersistentVolumeClaimList. A list of persistent volume claims. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistentvolumeclaims>

Returns The items of this V1PersistentVolumeClaimList.

Return type list[V1PersistentVolumeClaim]

kind

Gets the kind of this V1PersistentVolumeClaimList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PersistentVolumeClaimList.

Return type str

metadata

Gets the metadata of this V1PersistentVolumeClaimList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1PersistentVolumeClaimList.

Return type V1ListMeta

```
swagger_types = {'items': 'list[V1PersistentVolumeClaim]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_claim_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec (ac
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

access_modes

Gets the access_modes of this V1PersistentVolumeClaimSpec. AccessModes contains the desired

access modes the volume should have. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#access-modes-1>

Returns The access_modes of this V1PersistentVolumeClaimSpec.

Return type list[str]

attribute_map = {'volume_name': 'volumeName', 'selector': 'selector', 'storage_class_name': 'storageClassName', 'r

resources

Gets the resources of this V1PersistentVolumeClaimSpec. Resources represents the minimum resources the volume should have. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#resources>

Returns The resources of this V1PersistentVolumeClaimSpec.

Return type *V1ResourceRequirements*

selector

Gets the selector of this V1PersistentVolumeClaimSpec. A label query over volumes to consider for binding.

Returns The selector of this V1PersistentVolumeClaimSpec.

Return type V1LabelSelector

storage_class_name

Gets the storage_class_name of this V1PersistentVolumeClaimSpec. Name of the StorageClass required by the claim. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#class-1>

Returns The storage_class_name of this V1PersistentVolumeClaimSpec.

Return type str

swagger_types = {'volume_name': 'str', 'selector': 'V1LabelSelector', 'storage_class_name': 'str', 'resources': 'V1Re

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

volume_name

Gets the volume_name of this V1PersistentVolumeClaimSpec. VolumeName is the binding reference to the PersistentVolume backing this claim.

Returns The volume_name of this V1PersistentVolumeClaimSpec.

Return type str

kubernetes.client.models.v1_persistent_volume_claim_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

access_modes

Gets the access_modes of this V1PersistentVolumeClaimStatus. AccessModes contains the actual access modes the volume backing the PVC has. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#access-modes-1>

Returns The access_modes of this V1PersistentVolumeClaimStatus.

Return type list[str]

attribute_map = {'phase': 'phase', 'conditions': 'conditions', 'capacity': 'capacity', 'access_modes': 'accessModes'}

capacity

Gets the capacity of this V1PersistentVolumeClaimStatus. Represents the actual resources of the underlying volume.

Returns The capacity of this V1PersistentVolumeClaimStatus.

Return type dict(str, str)

conditions

Gets the conditions of this V1PersistentVolumeClaimStatus. Current Condition of persistent volume claim. If underlying persistent volume is being resized then the Condition will be set to 'ResizeStarted'.

Returns The conditions of this V1PersistentVolumeClaimStatus.

Return type list[V1PersistentVolumeClaimCondition]

phase

Gets the phase of this V1PersistentVolumeClaimStatus. Phase represents the current phase of PersistentVolumeClaim.

Returns The phase of this V1PersistentVolumeClaimStatus.

Return type str

swagger_types = {'phase': 'str', 'conditions': 'list[V1PersistentVolumeClaimCondition]', 'capacity': 'dict(str, str)', 'a

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_claim_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_persistent_volume_claim_volume_source.V1PersistentVolumeClaimVolumeSource`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'claim_name': 'claimName'}

claim_name

Gets the claim_name of this V1PersistentVolumeClaimVolumeSource. ClaimName is the name of a PersistentVolumeClaim in the same namespace as the pod using this volume. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistentvolumeclaims>

Returns The claim_name of this V1PersistentVolumeClaimVolumeSource.

Return type str

read_only

Gets the read_only of this V1PersistentVolumeClaimVolumeSource. Will force the ReadOnly setting in VolumeMounts. Default false.

Returns The read_only of this V1PersistentVolumeClaimVolumeSource.

Return type bool

swagger_types = {'read_only': 'bool', 'claim_name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList` (*api_version=None,*

items=None,
kind=None,
meta-
data=None)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PersistentVolumeList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PersistentVolumeList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1PersistentVolumeList. List of persistent volumes. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes>

Returns The items of this V1PersistentVolumeList.

Return type list[V1PersistentVolume]

kind

Gets the kind of this V1PersistentVolumeList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PersistentVolumeList.

Return type str

metadata

Gets the metadata of this V1PersistentVolumeList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1PersistentVolumeList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1PersistentVolume]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_persistent_volume_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec (access_modes=None,
aws_elastic_block_store=None,
azure_disk=None,
azure_file=None,
ca_certificate=None,
cephfs=None,
cinder=None,
claim_ref=None,
fc=None,
flex_volume=None,
flocker=None,
gce_persistent_disk=None,
glusterfs=None,
host_path=None,
iscsi=None,
local=None,
local_volume=None,
mount_options=None,
nfs=None,
persistent_volume_claim_ref=None,
photon_persistent_disk=None,
portworx_volume=None,
quobyte=None,
rbd=None,
scale_io=None,
storage_class_name=None,
storage_a-
geos=None,
vsphere_volume=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

access_modes

Gets the access_modes of this V1PersistentVolumeSpec. AccessModes contains all ways the volume can be mounted. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#access-modes>

Returns The access_modes of this V1PersistentVolumeSpec.

Return type list[str]

attribute_map = {'gce_persistent_disk': 'gcePersistentDisk', 'portworx_volume': 'portworxVolume', 'azure_disk': 'a

aws_elastic_block_store

Gets the aws_elastic_block_store of this V1PersistentVolumeSpec. AWSElasticBlockStore represents an AWS Disk resource that is attached to a kubelet's host machine and then exposed to the pod. More info: <https://kubernetes.io/docs/concepts/storage/volumes#awselasticblockstore>

Returns The `aws_elastic_block_store` of this `V1PersistentVolumeSpec`.

Return type *V1AWSElasticBlockStoreVolumeSource*

azure_disk

Gets the `azure_disk` of this `V1PersistentVolumeSpec`. `AzureDisk` represents an Azure Data Disk mount on the host and bind mount to the pod.

Returns The `azure_disk` of this `V1PersistentVolumeSpec`.

Return type *V1AzureDiskVolumeSource*

azure_file

Gets the `azure_file` of this `V1PersistentVolumeSpec`. `AzureFile` represents an Azure File Service mount on the host and bind mount to the pod.

Returns The `azure_file` of this `V1PersistentVolumeSpec`.

Return type *V1AzureFilePersistentVolumeSource*

capacity

Gets the capacity of this `V1PersistentVolumeSpec`. A description of the persistent volume's resources and capacity. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#capacity>

Returns The capacity of this `V1PersistentVolumeSpec`.

Return type `dict(str, str)`

cephfs

Gets the `cephfs` of this `V1PersistentVolumeSpec`. `CephFS` represents a Ceph FS mount on the host that shares a pod's lifetime

Returns The `cephfs` of this `V1PersistentVolumeSpec`.

Return type *V1CephFSPersistentVolumeSource*

cinder

Gets the `cinder` of this `V1PersistentVolumeSpec`. `Cinder` represents a cinder volume attached and mounted on kubelets host machine More info: <https://releases.k8s.io/HEAD/examples/mysql-cinder-pd/README.md>

Returns The `cinder` of this `V1PersistentVolumeSpec`.

Return type *V1CinderVolumeSource*

claim_ref

Gets the `claim_ref` of this `V1PersistentVolumeSpec`. `ClaimRef` is part of a bi-directional binding between `PersistentVolume` and `PersistentVolumeClaim`. Expected to be non-nil when bound. `claim.VolumeName` is the authoritative bind between PV and PVC. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#binding>

Returns The `claim_ref` of this `V1PersistentVolumeSpec`.

Return type *V1ObjectReference*

fc

Gets the `fc` of this `V1PersistentVolumeSpec`. `FC` represents a Fibre Channel resource that is attached to a kubelet's host machine and then exposed to the pod.

Returns The `fc` of this `V1PersistentVolumeSpec`.

Return type *V1FCVolumeSource*

flex_volume

Gets the flex_volume of this V1PersistentVolumeSpec. FlexVolume represents a generic volume resource that is provisioned/attached using an exec based plugin. This is an alpha feature and may change in future.

Returns The flex_volume of this V1PersistentVolumeSpec.

Return type *V1FlexVolumeSource*

flocker

Gets the flocker of this V1PersistentVolumeSpec. Flocker represents a Flocker volume attached to a kubelet's host machine and exposed to the pod for its usage. This depends on the Flocker control service being running

Returns The flocker of this V1PersistentVolumeSpec.

Return type *V1FlockerVolumeSource*

gce_persistent_disk

Gets the gce_persistent_disk of this V1PersistentVolumeSpec. GCEPersistentDisk represents a GCE Disk resource that is attached to a kubelet's host machine and then exposed to the pod. Provisioned by an admin. More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The gce_persistent_disk of this V1PersistentVolumeSpec.

Return type *V1GCEPersistentDiskVolumeSource*

glusterfs

Gets the glusterfs of this V1PersistentVolumeSpec. Glusterfs represents a Glusterfs volume that is attached to a host and exposed to the pod. Provisioned by an admin. More info: <https://releases.k8s.io/HEAD/examples/volumes/glusterfs/README.md>

Returns The glusterfs of this V1PersistentVolumeSpec.

Return type *V1GlusterfsVolumeSource*

host_path

Gets the host_path of this V1PersistentVolumeSpec. HostPath represents a directory on the host. Provisioned by a developer or tester. This is useful for single-node development and testing only! On-host storage is not supported in any way and WILL NOT WORK in a multi-node cluster. More info: <https://kubernetes.io/docs/concepts/storage/volumes#hostpath>

Returns The host_path of this V1PersistentVolumeSpec.

Return type *V1HostPathVolumeSource*

iscsi

Gets the iscsi of this V1PersistentVolumeSpec. ISCSI represents an ISCSI Disk resource that is attached to a kubelet's host machine and then exposed to the pod. Provisioned by an admin.

Returns The iscsi of this V1PersistentVolumeSpec.

Return type *V1ISCSIVolumeSource*

local

Gets the local of this V1PersistentVolumeSpec. Local represents directly-attached storage with node affinity

Returns The local of this V1PersistentVolumeSpec.

Return type *V1LocalVolumeSource*

mount_options

Gets the mount_options of this V1PersistentVolumeSpec. A list of mount options, e.g. ["ro", "soft"]. Not

validated - mount will simply fail if one is invalid. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes/#mount-options>

Returns The mount_options of this V1PersistentVolumeSpec.

Return type list[str]

nfs

Gets the nfs of this V1PersistentVolumeSpec. NFS represents an NFS mount on the host. Provisioned by an admin. More info: <https://kubernetes.io/docs/concepts/storage/volumes#nfs>

Returns The nfs of this V1PersistentVolumeSpec.

Return type *V1NFSVolumeSource*

persistent_volume_reclaim_policy

Gets the persistent_volume_reclaim_policy of this V1PersistentVolumeSpec. What happens to a persistent volume when released from its claim. Valid options are Retain (default) and Recycle. Recycling must be supported by the volume plugin underlying this persistent volume. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#reclaiming>

Returns The persistent_volume_reclaim_policy of this V1PersistentVolumeSpec.

Return type str

photon_persistent_disk

Gets the photon_persistent_disk of this V1PersistentVolumeSpec. PhotonPersistentDisk represents a PhotonController persistent disk attached and mounted on kubelets host machine

Returns The photon_persistent_disk of this V1PersistentVolumeSpec.

Return type *V1PhotonPersistentDiskVolumeSource*

portworx_volume

Gets the portworx_volume of this V1PersistentVolumeSpec. PortworxVolume represents a portworx volume attached and mounted on kubelets host machine

Returns The portworx_volume of this V1PersistentVolumeSpec.

Return type *V1PortworxVolumeSource*

quobyte

Gets the quobyte of this V1PersistentVolumeSpec. Quobyte represents a Quobyte mount on the host that shares a pod's lifetime

Returns The quobyte of this V1PersistentVolumeSpec.

Return type *V1QuobyteVolumeSource*

rbd

Gets the rbd of this V1PersistentVolumeSpec. RBD represents a Rados Block Device mount on the host that shares a pod's lifetime. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md>

Returns The rbd of this V1PersistentVolumeSpec.

Return type *V1RBDVolumeSource*

scale_io

Gets the scale_io of this V1PersistentVolumeSpec. ScaleIO represents a ScaleIO persistent volume attached and mounted on Kubernetes nodes.

Returns The scale_io of this V1PersistentVolumeSpec.

Return type *V1ScaleIOVolumeSource*

storage_class_name

Gets the storage_class_name of this V1PersistentVolumeSpec. Name of StorageClass to which this persistent volume belongs. Empty value means that this volume does not belong to any StorageClass.

Returns The storage_class_name of this V1PersistentVolumeSpec.

Return type str

storageos

Gets the storageos of this V1PersistentVolumeSpec. StorageOS represents a StorageOS volume that is attached to the kubelet's host machine and mounted into the pod More info: <https://releases.k8s.io/HEAD/examples/volumes/storageos/README.md>

Returns The storageos of this V1PersistentVolumeSpec.

Return type V1StorageOSPersistentVolumeSource

swagger_types = {'gce_persistent_disk': 'V1GCEPersistentDiskVolumeSource', 'portworx_volume': 'V1PortworxVol

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

vsphere_volume

Gets the vsphere_volume of this V1PersistentVolumeSpec. VsphereVolume represents a vSphere volume attached and mounted on kubelets host machine

Returns The vsphere_volume of this V1PersistentVolumeSpec.

Return type *V1VsphereVirtualDiskVolumeSource*

kubernetes.client.models.v1_persistent_volume_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_persistent_volume_status.**V1PersistentVolumeStatus** (*message=None*
phase=None
reason=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'phase': 'phase', 'message': 'message', 'reason': 'reason'}

message

Gets the message of this V1PersistentVolumeStatus. A human-readable message indicating details about why the volume is in this state.

Returns The message of this V1PersistentVolumeStatus.

Return type str

phase

Gets the phase of this V1PersistentVolumeStatus. Phase indicates if a volume is available, bound to a claim, or released by a claim. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#phase>

Returns The phase of this V1PersistentVolumeStatus.

Return type str

reason

Gets the reason of this V1PersistentVolumeStatus. Reason is a brief CamelCase string that describes any failure and is meant for machine parsing and tidy display in the CLI.

Returns The reason of this V1PersistentVolumeStatus.

Return type str

swagger_types = {'phase': 'str', 'message': 'str', 'reason': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_photon_persistent_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDisk

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'pd_id': 'pdID', 'fs_type': 'fsType'}

fs_type

Gets the fs_type of this V1PhotonPersistentDiskVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Ex. "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified.

Returns The fs_type of this V1PhotonPersistentDiskVolumeSource.

Return type str

pd_id

Gets the pd_id of this V1PhotonPersistentDiskVolumeSource. ID that identifies Photon Controller persistent disk

Returns The pd_id of this V1PhotonPersistentDiskVolumeSource.

Return type str

swagger_types = {'pd_id': 'str', 'fs_type': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_pod module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod.V1Pod(api_version=None, kind=None, meta-
data=None, spec=None, status=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Pod. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Pod.

Return type str

```
attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1Pod. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Pod.

Return type str

metadata

Gets the metadata of this V1Pod. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Pod.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Pod. Specification of the desired behavior of the pod. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1Pod.

Return type *V1PodSpec*

status

Gets the status of this V1Pod. Most recently observed status of the pod. This data may not be up to date. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1Pod.

Return type *V1PodStatus*

```
swagger_types = {'status': 'V1PodStatus', 'kind': 'str', 'spec': 'V1PodSpec', 'api_version': 'str', 'metadata': 'V1Obj
```

to_dict()

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.v1_pod_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_condition.V1PodCondition (last_probe_time=None,
                                                                last_transition_time=None,
                                                                message=None,
                                                                reason=None,
                                                                status=None,
                                                                type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'status': 'status', 'last_transition_time': 'lastTransitionTime', 'reason': 'reason', 'message': 'message'}
```

last_probe_time

Gets the last_probe_time of this V1PodCondition. Last time we probed the condition.

Returns The last_probe_time of this V1PodCondition.

Return type datetime

last_transition_time

Gets the last_transition_time of this V1PodCondition. Last time the condition transitioned from one status to another.

Returns The last_transition_time of this V1PodCondition.

Return type datetime

message

Gets the message of this V1PodCondition. Human-readable message indicating details about last transition.

Returns The message of this V1PodCondition.

Return type str

reason

Gets the reason of this V1PodCondition. Unique, one-word, CamelCase reason for the condition's last transition.

Returns The reason of this V1PodCondition.

Return type str

status

Gets the status of this V1PodCondition. Status is the status of the condition. Can be True, False, Unknown. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-conditions>

Returns The status of this V1PodCondition.

Return type str

```
swagger_types = {'status': 'str', 'last_transition_time': 'datetime', 'reason': 'str', 'message': 'str', 'type': 'str', 'last_
```

to_dict()
Returns the model properties as a dict

to_str()
Returns the string representation of the model

type
Gets the type of this V1PodCondition. Type is the type of the condition. Currently only Ready. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-conditions>

Returns The type of this V1PodCondition.

Return type str

kubernetes.client.models.v1_pod_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_list.V1PodList (api_version=None, items=None,
                                                    kind=None, metadata=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PodList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PodList.

Return type str

```
attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

items

Gets the items of this V1PodList. List of pods. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md>

Returns The items of this V1PodList.

Return type list[V1Pod]

kind

Gets the kind of this V1PodList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PodList.

Return type str

metadata

Gets the metadata of this V1PodList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1PodList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Pod]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_pod_security_context module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext (fs_group=None,
                                     run_as_non_root=None,
                                     run_as_user=None,
                                     se_linux_options=None,
                                     supplemental_groups=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'run_as_non_root': 'runAsNonRoot', 'supplemental_groups': 'supplementalGroups', 'se_linux_options': 'seLinuxOptions'}

fs_group

Gets the fs_group of this V1PodSecurityContext. A special supplemental group that applies to all containers in a pod. Some volume types allow the Kubelet to change the ownership of that volume to be owned by the pod: 1. The owning GID will be the FSGroup 2. The setgid bit is set (new files created in the volume will be owned by FSGroup) 3. The permission bits are OR'd with rw-rw— If unset, the Kubelet will not modify the ownership and permissions of any volume.

Returns The fs_group of this V1PodSecurityContext.

Return type int

run_as_non_root

Gets the run_as_non_root of this V1PodSecurityContext. Indicates that the container must run as a non-root user. If true, the Kubelet will validate the image at runtime to ensure that it does not run as UID 0 (root) and fail to start the container if it does. If unset or false, no such validation will be performed. May also be set in SecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.

Returns The run_as_non_root of this V1PodSecurityContext.

Return type bool

run_as_user

Gets the run_as_user of this V1PodSecurityContext. The UID to run the entrypoint of the container process. Defaults to user specified in image metadata if unspecified. May also be set in SecurityContext. If set

in both `SecurityContext` and `PodSecurityContext`, the value specified in `SecurityContext` takes precedence for that container.

Returns The `run_as_user` of this `V1PodSecurityContext`.

Return type `int`

`se_linux_options`

Gets the `se_linux_options` of this `V1PodSecurityContext`. The SELinux context to be applied to all containers. If unspecified, the container runtime will allocate a random SELinux context for each container. May also be set in `SecurityContext`. If set in both `SecurityContext` and `PodSecurityContext`, the value specified in `SecurityContext` takes precedence for that container.

Returns The `se_linux_options` of this `V1PodSecurityContext`.

Return type *`V1SELinuxOptions`*

`supplemental_groups`

Gets the `supplemental_groups` of this `V1PodSecurityContext`. A list of groups applied to the first process run in each container, in addition to the container's primary GID. If unspecified, no groups will be added to any container.

Returns The `supplemental_groups` of this `V1PodSecurityContext`.

Return type `list[int]`

swagger_types = {'run_as_non_root': 'bool', 'supplemental_groups': 'list[int]', 'se_linux_options': 'V1SELinuxOptions'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

`kubernetes.client.models.v1_pod_spec` module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>


```
class kubernetes.client.models.v1_pod_spec.V1PodSpec (active_deadline_seconds=None,
                                                       affinity=None,          auto-
                                                       mount_service_account_token=None,
                                                       containers=None,
                                                       dns_policy=None,
                                                       host_aliases=None,
                                                       host_ipc=None,
                                                       host_network=None,
                                                       host_pid=None, hostname=None,
                                                       image_pull_secrets=None,
                                                       init_containers=None,
                                                       node_name=None,
                                                       node_selector=None,
                                                       priority=None,          prior-
                                                       ity_class_name=None,
                                                       restart_policy=None,      sched-
                                                       uler_name=None,          secu-
                                                       rity_context=None,       ser-
                                                       vice_account=None,      ser-
                                                       vice_account_name=None,
                                                       subdomain=None,         termina-
                                                       tion_grace_period_seconds=None,
                                                       tolerations=None,       vol-
                                                       umes=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

active_deadline_seconds

Gets the active_deadline_seconds of this V1PodSpec. Optional duration in seconds the pod may be active on the node relative to StartTime before the system will actively try to mark it failed and kill associated containers. Value must be a positive integer.

Returns The active_deadline_seconds of this V1PodSpec.

Return type int

affinity

Gets the affinity of this V1PodSpec. If specified, the pod's scheduling constraints

Returns The affinity of this V1PodSpec.

Return type V1Affinity

attribute_map = {'termination_grace_period_seconds': 'terminationGracePeriodSeconds', 'init_containers': 'initContainers'}

automount_service_account_token

Gets the automount_service_account_token of this V1PodSpec. AutomountServiceAccountToken indicates whether a service account token should be automatically mounted.

Returns The automount_service_account_token of this V1PodSpec.

Return type bool

containers

Gets the containers of this V1PodSpec. List of containers belonging to the pod. Containers cannot currently be added or removed. There must be at least one container in a Pod. Cannot be updated.

Returns The containers of this V1PodSpec.

Return type list[V1Container]

dns_policy

Gets the `dns_policy` of this `V1PodSpec`. Set DNS policy for containers within the pod. One of 'ClusterFirstWithHostNet', 'ClusterFirst' or 'Default'. Defaults to "ClusterFirst". To have DNS options set along with `hostNetwork`, you have to specify DNS policy explicitly to 'ClusterFirstWithHostNet'.

Returns The `dns_policy` of this `V1PodSpec`.

Return type str

host_aliases

Gets the `host_aliases` of this `V1PodSpec`. `HostAliases` is an optional list of hosts and IPs that will be injected into the pod's hosts file if specified. This is only valid for non-`hostNetwork` pods.

Returns The `host_aliases` of this `V1PodSpec`.

Return type list[V1HostAlias]

host_ipc

Gets the `host_ipc` of this `V1PodSpec`. Use the host's ipc namespace. Optional: Default to false.

Returns The `host_ipc` of this `V1PodSpec`.

Return type bool

host_network

Gets the `host_network` of this `V1PodSpec`. Host networking requested for this pod. Use the host's network namespace. If this option is set, the ports that will be used must be specified. Default to false.

Returns The `host_network` of this `V1PodSpec`.

Return type bool

host_pid

Gets the `host_pid` of this `V1PodSpec`. Use the host's pid namespace. Optional: Default to false.

Returns The `host_pid` of this `V1PodSpec`.

Return type bool

hostname

Gets the `hostname` of this `V1PodSpec`. Specifies the hostname of the Pod If not specified, the pod's hostname will be set to a system-defined value.

Returns The `hostname` of this `V1PodSpec`.

Return type str

image_pull_secrets

Gets the `image_pull_secrets` of this `V1PodSpec`. `ImagePullSecrets` is an optional list of references to secrets in the same namespace to use for pulling any of the images used by this `PodSpec`. If specified, these secrets will be passed to individual puller implementations for them to use. For example, in the case of docker, only `DockerConfig` type secrets are honored. More info: <https://kubernetes.io/docs/concepts/containers/images#specifying-imagepullsecrets-on-a-pod>

Returns The `image_pull_secrets` of this `V1PodSpec`.

Return type list[V1LocalObjectReference]

init_containers

Gets the `init_containers` of this `V1PodSpec`. List of initialization containers belonging to the pod. Init containers are executed in order prior to containers being started. If any init container fails, the pod is considered to have failed and is handled according to its `restartPolicy`. The name for an init container or normal container must be unique among all containers. Init containers may not have `Lifecycle` actions, `Readiness` probes, or `Liveness` probes. The `resourceRequirements` of an init container are taken

into account during scheduling by finding the highest request/limit for each resource type, and then using the max of of that value or the sum of the normal containers. Limits are applied to init containers in a similar fashion. Init containers cannot currently be added or removed. Cannot be updated. More info: <https://kubernetes.io/docs/concepts/workloads/pods/init-containers/>

Returns The init_containers of this V1PodSpec.

Return type list[V1Container]

node_name

Gets the node_name of this V1PodSpec. NodeName is a request to schedule this pod onto a specific node. If it is non-empty, the scheduler simply schedules this pod onto that node, assuming that it fits resource requirements.

Returns The node_name of this V1PodSpec.

Return type str

node_selector

Gets the node_selector of this V1PodSpec. NodeSelector is a selector which must be true for the pod to fit on a node. Selector which must match a node's labels for the pod to be scheduled on that node. More info: <https://kubernetes.io/docs/concepts/configuration/assign-pod-node/>

Returns The node_selector of this V1PodSpec.

Return type dict(str, str)

priority

Gets the priority of this V1PodSpec. The priority value. Various system components use this field to find the priority of the pod. When Priority Admission Controller is enabled, it prevents users from setting this field. The admission controller populates this field from PriorityClassName. The higher the value, the higher the priority.

Returns The priority of this V1PodSpec.

Return type int

priority_class_name

Gets the priority_class_name of this V1PodSpec. If specified, indicates the pod's priority. "SYSTEM" is a special keyword which indicates the highest priority. Any other name must be defined by creating a PriorityClass object with that name. If not specified, the pod priority will be default or zero if there is no default.

Returns The priority_class_name of this V1PodSpec.

Return type str

restart_policy

Gets the restart_policy of this V1PodSpec. Restart policy for all containers within the pod. One of Always, OnFailure, Never. Default to Always. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/#restart-policy>

Returns The restart_policy of this V1PodSpec.

Return type str

scheduler_name

Gets the scheduler_name of this V1PodSpec. If specified, the pod will be dispatched by specified scheduler. If not specified, the pod will be dispatched by default scheduler.

Returns The scheduler_name of this V1PodSpec.

Return type str

security_context

Gets the security_context of this V1PodSpec. SecurityContext holds pod-level security attributes and common container settings. Optional: Defaults to empty. See type description for default values of each field.

Returns The security_context of this V1PodSpec.

Return type *V1PodSecurityContext*

service_account

Gets the service_account of this V1PodSpec. DeprecatedServiceAccount is a deprecated alias for ServiceAccountName. Deprecated: Use serviceAccountName instead.

Returns The service_account of this V1PodSpec.

Return type str

service_account_name

Gets the service_account_name of this V1PodSpec. ServiceAccountName is the name of the ServiceAccount to use to run this pod. More info: <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

Returns The service_account_name of this V1PodSpec.

Return type str

subdomain

Gets the subdomain of this V1PodSpec. If specified, the fully qualified Pod hostname will be “<host-name>.<subdomain>.<pod namespace>.svc.<cluster domain>”. If not specified, the pod will not have a domainname at all.

Returns The subdomain of this V1PodSpec.

Return type str

swagger_types = {'termination_grace_period_seconds': 'int', 'init_containers': 'list[V1Container]', 'priority_class_name': 'str'}

termination_grace_period_seconds

Gets the termination_grace_period_seconds of this V1PodSpec. Optional duration in seconds the pod needs to terminate gracefully. May be decreased in delete request. Value must be non-negative integer. The value zero indicates delete immediately. If this value is nil, the default grace period will be used instead. The grace period is the duration in seconds after the processes running in the pod are sent a termination signal and the time when the processes are forcibly halted with a kill signal. Set this value longer than the expected cleanup time for your process. Defaults to 30 seconds.

Returns The termination_grace_period_seconds of this V1PodSpec.

Return type int

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

tolerations

Gets the tolerations of this V1PodSpec. If specified, the pod's tolerations.

Returns The tolerations of this V1PodSpec.

Return type list[V1Toleration]

volumes

Gets the volumes of this V1PodSpec. List of volumes that can be mounted by containers belonging to the pod. More info: <https://kubernetes.io/docs/concepts/storage/volumes>

Returns The volumes of this V1PodSpec.

Return type list[*V1Volume*]

kubernetes.client.models.v1_pod_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_status.V1PodStatus (conditions=None, container_statuses=None,
                                                         host_ip=None,
                                                         init_container_statuses=None,
                                                         message=None,
                                                         phase=None, pod_ip=None,
                                                         qos_class=None,
                                                         reason=None,
                                                         start_time=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'phase': 'phase', 'reason': 'reason', 'qos_class': 'qosClass', 'start_time': 'startTime', 'pod_ip': 'podIP'}

conditions

Gets the conditions of this V1PodStatus. Current service state of pod. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-conditions>

Returns The conditions of this V1PodStatus.

Return type list[*V1PodCondition*]

container_statuses

Gets the container_statuses of this V1PodStatus. The list has one entry per container in the manifest. Each entry is currently the output of *docker inspect*. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-and-container-status>

Returns The container_statuses of this V1PodStatus.

Return type list[*V1ContainerStatus*]

host_ip

Gets the host_ip of this V1PodStatus. IP address of the host to which the pod is assigned. Empty if not yet scheduled.

Returns The host_ip of this V1PodStatus.

Return type str

init_container_statuses

Gets the init_container_statuses of this V1PodStatus. The list has one entry per init container in the manifest. The most recent successful init container will have ready = true, the most recently started container

will have `startTime` set. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-and-container-status>

Returns The `init_container_statuses` of this `V1PodStatus`.

Return type `list[V1ContainerStatus]`

message

Gets the message of this `V1PodStatus`. A human readable message indicating details about why the pod is in this condition.

Returns The message of this `V1PodStatus`.

Return type `str`

phase

Gets the phase of this `V1PodStatus`. Current condition of the pod. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#pod-phase>

Returns The phase of this `V1PodStatus`.

Return type `str`

pod_ip

Gets the `pod_ip` of this `V1PodStatus`. IP address allocated to the pod. Routable at least within the cluster. Empty if not yet allocated.

Returns The `pod_ip` of this `V1PodStatus`.

Return type `str`

qos_class

Gets the `qos_class` of this `V1PodStatus`. The Quality of Service (QOS) classification assigned to the pod based on resource requirements See `PodQOSClass` type for available QOS classes More info: <https://github.com/kubernetes/kubernetes/blob/master/docs/design/resource-qos.md>

Returns The `qos_class` of this `V1PodStatus`.

Return type `str`

reason

Gets the reason of this `V1PodStatus`. A brief CamelCase message indicating details about why the pod is in this state. e.g. 'Evicted'

Returns The reason of this `V1PodStatus`.

Return type `str`

start_time

Gets the `start_time` of this `V1PodStatus`. RFC 3339 date and time at which the object was acknowledged by the Kubelet. This is before the Kubelet pulled the container image(s) for the pod.

Returns The `start_time` of this `V1PodStatus`.

Return type `datetime`

swagger_types = {'phase': 'str', 'reason': 'str', 'qos_class': 'str', 'start_time': 'datetime', 'pod_ip': 'str', 'init_contain

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_pod_template module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_template.V1PodTemplate (api_version=None,
                                                             kind=None,      meta-
                                                             data=None,      tem-
                                                             plate=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PodTemplate. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PodTemplate.

Return type str

attribute_map = {'kind': 'kind', 'template': 'template', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1PodTemplate. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PodTemplate.

Return type str

metadata

Gets the metadata of this V1PodTemplate. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1PodTemplate.

Return type *V1ObjectMeta*

swagger_types = {'kind': 'str', 'template': 'V1PodTemplateSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

template

Gets the template of this V1PodTemplate. Template defines the pods that will be created from this pod template. <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The template of this V1PodTemplate.

Return type *V1PodTemplateSpec*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_pod_template_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_template_list.V1PodTemplateList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1PodTemplateList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1PodTemplateList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1PodTemplateList. List of pod templates

Returns The items of this V1PodTemplateList.

Return type list[V1PodTemplate]

kind

Gets the kind of this V1PodTemplateList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1PodTemplateList.

Return type str

metadata

Gets the metadata of this V1PodTemplateList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1PodTemplateList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1PodTemplate]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_pod_template_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec (metadata=None,
                                                                    spec=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'spec': 'spec', 'metadata': 'metadata'}

metadata

Gets the metadata of this V1PodTemplateSpec. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1PodTemplateSpec.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1PodTemplateSpec. Specification of the desired behavior of the pod. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1PodTemplateSpec.

Return type *V1PodSpec*

swagger_types = {'spec': 'V1PodSpec', 'metadata': 'V1ObjectMeta'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_preconditions module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_preconditions.V1Preconditions (uid=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'uid': 'uid'}

swagger_types = {'uid': 'str'}

to_dict ()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

uid

Gets the uid of this V1Preconditions. Specifies the target UID.

Returns The uid of this V1Preconditions.

Return type str

kubernetes.client.models.v1_probe module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_probe.V1Probe(_exec=None, failure_threshold=None,
                                                http_get=None, initial_delay_seconds=None,
                                                period_seconds=None, success_threshold=None,
                                                tcp_socket=None, timeout_seconds=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'tcp_socket': 'tcpSocket', 'timeout_seconds': 'timeoutSeconds', 'failure_threshold': 'failureThreshold'}

failure_threshold

Gets the failure_threshold of this V1Probe. Minimum consecutive failures for the probe to be considered failed after having succeeded. Defaults to 3. Minimum value is 1.

Returns The failure_threshold of this V1Probe.

Return type int

http_get

Gets the http_get of this V1Probe. HTTPGet specifies the http request to perform.

Returns The http_get of this V1Probe.

Return type *V1HTTPGetAction*

initial_delay_seconds

Gets the initial_delay_seconds of this V1Probe. Number of seconds after the container has started before liveness probes are initiated. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#container-probes>

Returns The initial_delay_seconds of this V1Probe.

Return type int

period_seconds

Gets the period_seconds of this V1Probe. How often (in seconds) to perform the probe. Default to 10 seconds. Minimum value is 1.

Returns The period_seconds of this V1Probe.

Return type int

success_threshold

Gets the success_threshold of this V1Probe. Minimum consecutive successes for the probe to be considered successful after having failed. Defaults to 1. Must be 1 for liveness. Minimum value is 1.

Returns The success_threshold of this V1Probe.

Return type int

swagger_types = {'tcp_socket': 'V1TCPSocketAction', 'timeout_seconds': 'int', 'failure_threshold': 'int', '_exec': 'V1Probe'

tcp_socket

Gets the tcp_socket of this V1Probe. TCPSocket specifies an action involving a TCP port. TCP hooks not yet supported

Returns The tcp_socket of this V1Probe.

Return type *V1TCPSocketAction*

timeout_seconds

Gets the timeout_seconds of this V1Probe. Number of seconds after which the probe times out. Defaults to 1 second. Minimum value is 1. More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle#container-probes>

Returns The timeout_seconds of this V1Probe.

Return type int

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_quobyte_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource (group=None,
                                                                              read_only=None,
                                                                              registry=None,
                                                                              user=None,
                                                                              volume=None)

Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'volume': 'volume', 'group': 'group', 'registry': 'registry', 'user': 'user'}

group

Gets the group of this V1QuobyteVolumeSource. Group to map volume access to Default is no group

Returns The group of this V1QuobyteVolumeSource.

Return type str

read_only

Gets the read_only of this V1QuobyteVolumeSource. ReadOnly here will force the Quobyte volume to be mounted with read-only permissions. Defaults to false.

Returns The read_only of this V1QuobyteVolumeSource.

Return type bool

registry

Gets the registry of this V1QuobyteVolumeSource. Registry represents a single or multiple Quobyte Registry services specified as a string as host:port pair (multiple entries are separated with commas) which acts as the central registry for volumes

Returns The registry of this V1QuobyteVolumeSource.

Return type str

swagger_types = {'read_only': 'bool', 'volume': 'str', 'group': 'str', 'registry': 'str', 'user': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

user

Gets the user of this V1QuobyteVolumeSource. User to map volume access to Defaults to serviceaccount user

Returns The user of this V1QuobyteVolumeSource.

Return type str

volume

Gets the volume of this V1QuobyteVolumeSource. Volume is a string that references an already created Quobyte volume by name.

Returns The volume of this V1QuobyteVolumeSource.

Return type str

kubernetes.client.models.v1_rbd_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource (fs_type=None,
im-
age=None,
keyring=None,
monitors=None,
pool=None,
read_only=None,
secret_ref=None,
user=None)
```

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'secret_ref': 'secretRef', 'fs_type': 'fsType', 'user': 'user', 'image': 'image'}

fs_type

Gets the fs_type of this V1RBDVolumeSource. Filesystem type of the volume that you want to mount. Tip: Ensure that the filesystem type is supported by the host operating system. Examples: "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified. More info: <https://kubernetes.io/docs/concepts/storage/volumes#rbd>

Returns The fs_type of this V1RBDVolumeSource.

Return type str

image

Gets the image of this V1RBDVolumeSource. The rados image name. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The image of this V1RBDVolumeSource.

Return type str

keyring

Gets the keyring of this V1RBDVolumeSource. Keyring is the path to key ring for RBDUser. Default is /etc/ceph/keyring. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The keyring of this V1RBDVolumeSource.

Return type str

monitors

Gets the monitors of this V1RBDVolumeSource. A collection of Ceph monitors. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The monitors of this V1RBDVolumeSource.

Return type list[str]

pool

Gets the pool of this V1RBDVolumeSource. The rados pool name. Default is rbd. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The pool of this V1RBDVolumeSource.

Return type str

read_only

Gets the read_only of this V1RBDVolumeSource. ReadOnly here will force the ReadOnly setting in VolumeMounts. Defaults to false. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The read_only of this V1RBDVolumeSource.

Return type bool

secret_ref

Gets the secret_ref of this V1RBDVolumeSource. SecretRef is name of the authentication secret for RBDUser. If provided overrides keyring. Default is nil. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The secret_ref of this V1RBDVolumeSource.

Return type *V1LocalObjectReference*

swagger_types = {'read_only': 'bool', 'secret_ref': 'V1LocalObjectReference', 'fs_type': 'str', 'user': 'str', 'image': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

user

Gets the user of this V1RBDVolumeSource. The rados user name. Default is admin. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md#how-to-use-it>

Returns The user of this V1RBDVolumeSource.

Return type str

kubernetes.client.models.v1_replication_controller module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_replication_controller.V1ReplicationController` (*api_version=None, kind=None, metadata=None, spec=None, status=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ReplicationController. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ReplicationController.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1ReplicationController. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ReplicationController.

Return type str

metadata

Gets the metadata of this V1ReplicationController. If the Labels of a ReplicationController are empty, they are defaulted to be the same as the Pod(s) that the replication controller manages. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ReplicationController.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1ReplicationController. Spec defines the specification of the desired behavior of the replication controller. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1ReplicationController.

Return type *V1ReplicationControllerSpec*

status

Gets the status of this V1ReplicationController. Status is the most recently observed status of the replication controller. This data may be out of date by some window of time. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1ReplicationController.

Return type *V1ReplicationControllerStatus*

swagger_types = {'status': 'V1ReplicationControllerStatus', 'kind': 'str', 'spec': 'V1ReplicationControllerSpec', 'api

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_replication_controller_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCon

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'status': 'status', 'message': 'message', 'type': 'type', 'reason': 'reason', 'last_transition_time': 'las

last_transition_time

Gets the last_transition_time of this V1ReplicationControllerCondition. The last time the condition transitioned from one status to another.

Returns The last_transition_time of this V1ReplicationControllerCondition.

Return type datetime

message

Gets the message of this V1ReplicationControllerCondition. A human readable message indicating details about the transition.

Returns The message of this V1ReplicationControllerCondition.

Return type str

reason

Gets the reason of this V1ReplicationControllerCondition. The reason for the condition's last transition.

Returns The reason of this V1ReplicationControllerCondition.

Return type str

status

Gets the status of this V1ReplicationControllerCondition. Status of the condition, one of True, False, Unknown.

Returns The status of this V1ReplicationControllerCondition.

Return type str

swagger_types = {'status': 'str', 'message': 'str', 'type': 'str', 'reason': 'str', 'last_transition_time': 'datetime'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1ReplicationControllerCondition. Type of replication controller condition.

Returns The type of this V1ReplicationControllerCondition.

Return type str

kubernetes.client.models.v1_replication_controller_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList (api_...
item...
kind...
meta...
data...

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ReplicationControllerList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may

reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ReplicationControllerList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1ReplicationControllerList. List of replication controllers. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller>

Returns The items of this V1ReplicationControllerList.

Return type list[V1ReplicationController]

kind

Gets the kind of this V1ReplicationControllerList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ReplicationControllerList.

Return type str

metadata

Gets the metadata of this V1ReplicationControllerList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1ReplicationControllerList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1ReplicationController]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_replication_controller_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec (min

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'selector': 'selector', 'replicas': 'replicas', 'template': 'template', 'min_ready_seconds': 'minReadySeconds'}

min_ready_seconds

Gets the min_ready_seconds of this V1ReplicationControllerSpec. Minimum number of seconds for which a newly created pod should be ready without any of its container crashing, for it to be considered available. Defaults to 0 (pod will be considered available as soon as it is ready)

Returns The min_ready_seconds of this V1ReplicationControllerSpec.

Return type int

replicas

Gets the replicas of this V1ReplicationControllerSpec. Replicas is the number of desired replicas. This is a pointer to distinguish between explicit zero and unspecified. Defaults to 1. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller#what-is-a-replicationcontroller>

Returns The replicas of this V1ReplicationControllerSpec.

Return type int

selector

Gets the selector of this V1ReplicationControllerSpec. Selector is a label query over pods that should match the Replicas count. If Selector is empty, it is defaulted to the labels present on the Pod template. Label keys and values that must match in order to be controlled by this replication controller, if empty defaulted to labels on Pod template. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#label-selectors>

Returns The selector of this V1ReplicationControllerSpec.

Return type dict(str, str)

swagger_types = {'selector': 'dict(str, str)', 'replicas': 'int', 'template': 'V1PodTemplateSpec', 'min_ready_seconds': 'int'}

template

Gets the template of this V1ReplicationControllerSpec. Template is the object that describes the pod that will be created if insufficient replicas are detected. This takes precedence over a TemplateRef. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller#pod-template>

Returns The template of this V1ReplicationControllerSpec.

Return type *V1PodTemplateSpec*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_replication_controller_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'replicas': 'replicas', 'observed_generation': 'observedGeneration', 'available_replicas': 'availableReplicas'}

available_replicas

Gets the available_replicas of this V1ReplicationControllerStatus. The number of available replicas (ready for at least minReadySeconds) for this replication controller.

Returns The available_replicas of this V1ReplicationControllerStatus.

Return type int

conditions

Gets the conditions of this V1ReplicationControllerStatus. Represents the latest available observations of a replication controller's current state.

Returns The conditions of this V1ReplicationControllerStatus.

Return type list[V1ReplicationControllerCondition]

fully_labeled_replicas

Gets the fully_labeled_replicas of this V1ReplicationControllerStatus. The number of pods that have labels matching the labels of the pod template of the replication controller.

Returns The fully_labeled_replicas of this V1ReplicationControllerStatus.

Return type int

observed_generation

Gets the observed_generation of this V1ReplicationControllerStatus. ObservedGeneration reflects the generation of the most recently observed replication controller.

Returns The observed_generation of this V1ReplicationControllerStatus.

Return type int

ready_replicas

Gets the ready_replicas of this V1ReplicationControllerStatus. The number of ready replicas for this replication controller.

Returns The ready_replicas of this V1ReplicationControllerStatus.

Return type int

replicas

Gets the replicas of this V1ReplicationControllerStatus. Replicas is the most recently observed number of replicas. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller#what-is-a-replicationcontroller>

Returns The replicas of this V1ReplicationControllerStatus.

Return type int

```

swagger_types = {'replicas': 'int', 'observed_generation': 'int', 'available_replicas': 'int', 'ready_replicas': 'int', 'full
to_dict()
    Returns the model properties as a dict
to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_resource_field_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector(container_name,
                                         di-
                                         vi-
                                         sor=None,
                                         re-
                                         source=None)

```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'resource': 'resource', 'divisor': 'divisor', 'container_name': 'containerName'}
```

container_name

Gets the container_name of this V1ResourceFieldSelector. Container name: required for volumes, optional for env vars

Returns The container_name of this V1ResourceFieldSelector.

Return type str

divisor

Gets the divisor of this V1ResourceFieldSelector. Specifies the output format of the exposed resources, defaults to "1"

Returns The divisor of this V1ResourceFieldSelector.

Return type str

resource

Gets the resource of this V1ResourceFieldSelector. Required: resource to select

Returns The resource of this V1ResourceFieldSelector.

Return type str

```
swagger_types = {'resource': 'str', 'divisor': 'str', 'container_name': 'str'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.v1_resource_quota module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_resource_quota.V1ResourceQuota (api_version=None,
                                                                    kind=None,
                                                                    metadata=None,
                                                                    spec=None,
                                                                    status=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ResourceQuota. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ResourceQuota.

Return type str

```
attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1ResourceQuota. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ResourceQuota.

Return type str

metadata

Gets the metadata of this V1ResourceQuota. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ResourceQuota.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1ResourceQuota. Spec defines the desired quota. <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1ResourceQuota.

Return type *V1ResourceQuotaSpec*

status

Gets the status of this V1ResourceQuota. Status defines the actual enforced quota and its current usage. <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1ResourceQuota.

Return type *V1ResourceQuotaStatus*

```
swagger_types = {'status': 'V1ResourceQuotaStatus', 'kind': 'str', 'spec': 'V1ResourceQuotaSpec', 'api_version': 'str'}
```

```

to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model

```

kubernetes.client.models.v1_resource_quota_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList (api_version=None,
                                                                              items=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None)

```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ResourceQuotaList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ResourceQuotaList.

Return type str

```
attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

items

Gets the items of this V1ResourceQuotaList. Items is a list of ResourceQuota objects. More info: https://git.k8s.io/community/contributors/design-proposals/admission_control_resource_quota.md

Returns The items of this V1ResourceQuotaList.

Return type list[V1ResourceQuota]

kind

Gets the kind of this V1ResourceQuotaList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ResourceQuotaList.

Return type str

metadata

Gets the metadata of this V1ResourceQuotaList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1ResourceQuotaList.

Return type V1ListMeta

```
swagger_types = {'items': 'list[V1ResourceQuota]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}
to_dict()
    Returns the model properties as a dict
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1_resource_quota_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec (hard=None,
                                                                    scopes=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'scopes': 'scopes', 'hard': 'hard'}
```

hard

Gets the hard of this V1ResourceQuotaSpec. Hard is the set of desired hard limits for each named resource. More info: https://git.k8s.io/community/contributors/design-proposals/admission_control_resource_quota.md

Returns The hard of this V1ResourceQuotaSpec.

Return type dict(str, str)

scopes

Gets the scopes of this V1ResourceQuotaSpec. A collection of filters that must match each object tracked by a quota. If not specified, the quota matches all objects.

Returns The scopes of this V1ResourceQuotaSpec.

Return type list[str]

```
swagger_types = {'scopes': 'list[str]', 'hard': 'dict(str, str)'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.v1_resource_quota_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_resource_quota_status.V1ResourceQuotaStatus` (*hard=None, used=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'hard': 'hard', 'used': 'used'}

hard

Gets the hard of this V1ResourceQuotaStatus. Hard is the set of enforced hard limits for each named resource. More info: https://git.k8s.io/community/contributors/design-proposals/admission_control_resource_quota.md

Returns The hard of this V1ResourceQuotaStatus.

Return type `dict(str, str)`

swagger_types = {'hard': 'dict(str, str)', 'used': 'dict(str, str)'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

used

Gets the used of this V1ResourceQuotaStatus. Used is the current observed total usage of the resource in the namespace.

Returns The used of this V1ResourceQuotaStatus.

Return type `dict(str, str)`

kubernetes.client.models.v1_resource_requirements module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_resource_requirements.V1ResourceRequirements` (*limits=None, re-quests=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'requests': 'requests', 'limits': 'limits'}

limits

Gets the limits of this V1ResourceRequirements. Limits describes the maximum amount of compute resources allowed. More info: <https://kubernetes.io/docs/concepts/configuration/manage-compute-resources-container/>

Returns The limits of this V1ResourceRequirements.

Return type `dict(str, str)`

requests

Gets the requests of this V1ResourceRequirements. Requests describes the minimum amount of compute resources required. If Requests is omitted for a container, it defaults to Limits if that is explicitly

specified, otherwise to an implementation-defined value. More info: <https://kubernetes.io/docs/concepts/configuration/manage-compute-resources-container/>

Returns The requests of this V1ResourceRequirements.

Return type dict(str, str)

swagger_types = {'requests': 'dict(str, str)', 'limits': 'dict(str, str)'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_scale module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_scale.V1Scale` (*api_version=None, kind=None, metadata=None, spec=None, status=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Scale. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Scale.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1Scale. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Scale.

Return type str

metadata

Gets the metadata of this V1Scale. Standard object metadata; More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>.

Returns The metadata of this V1Scale.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Scale. defines the behavior of the scale. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>.

Returns The spec of this V1Scale.

Return type *V1ScaleSpec*

status

Gets the status of this V1Scale. current status of the scale. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>. Read-only.

Returns The status of this V1Scale.

Return type *V1ScaleStatus*

swagger_types = {'status': 'V1ScaleStatus', 'kind': 'str', 'spec': 'V1ScaleSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_scale_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_scale_spec.V1ScaleSpec` (*replicas=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'replicas': 'replicas'}

replicas

Gets the replicas of this V1ScaleSpec. desired number of instances for the scaled object.

Returns The replicas of this V1ScaleSpec.

Return type `int`

swagger_types = {'replicas': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_scale_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_scale_status.V1ScaleStatus (replicas=None, selector=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'selector': 'selector', 'replicas': 'replicas'}

replicas

Gets the replicas of this V1ScaleStatus. actual number of observed instances of the scaled object.

Returns The replicas of this V1ScaleStatus.

Return type int

selector

Gets the selector of this V1ScaleStatus. label query over pods that should match the replicas count. This is same as the label selector but in the string format to avoid introspection by clients. The string will be in the same format as the query-param syntax. More info about label selectors: <http://kubernetes.io/docs/user-guide/labels#label-selectors>

Returns The selector of this V1ScaleStatus.

Return type str

swagger_types = {'selector': 'str', 'replicas': 'int'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1_se_linux_options module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_se_linux_options.V1SELinuxOptions (level=None,
                                                                    role=None,
                                                                    type=None,
                                                                    user=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'type': 'type', 'role': 'role', 'user': 'user', 'level': 'level'}

level

Gets the level of this V1SELinuxOptions. Level is SELinux level label that applies to the container.

Returns The level of this V1SELinuxOptions.

Return type str

role

Gets the role of this V1SELinuxOptions. Role is a SELinux role label that applies to the container.

Returns The role of this V1SELinuxOptions.

Return type str

swagger_types = {'type': 'str', 'role': 'str', 'user': 'str', 'level': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1SELinuxOptions. Type is a SELinux type label that applies to the container.

Returns The type of this V1SELinuxOptions.

Return type str

user

Gets the user of this V1SELinuxOptions. User is a SELinux user label that applies to the container.

Returns The user of this V1SELinuxOptions.

Return type str

kubernetes.client.models.v1_secret module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_secret.V1Secret (api_version=None, data=None,
                                                    kind=None, metadata=None,
                                                    string_data=None, type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Secret. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Secret.

Return type str

attribute_map = {'kind': 'kind', 'data': 'data', 'string_data': 'stringData', 'type': 'type', 'api_version': 'apiVersion',

data

Gets the data of this V1Secret. Data contains the secret data. Each key must consist of alphanumeric characters, '-', '_' or '.'. The serialized form of the secret data is a base64 encoded string, representing the arbitrary (possibly non-string) data value here. Described in <https://tools.ietf.org/html/rfc4648#section-4>

Returns The data of this V1Secret.

Return type dict(str, str)

kind

Gets the kind of this V1Secret. Kind is a string value representing the REST resource this object represents.

Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Secret.

Return type str

metadata

Gets the metadata of this V1Secret. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Secret.

Return type *V1ObjectMeta*

string_data

Gets the string_data of this V1Secret. stringData allows specifying non-binary secret data in string form. It is provided as a write-only convenience method. All keys and values are merged into the data field on write, overwriting any existing values. It is never output when reading from the API.

Returns The string_data of this V1Secret.

Return type dict(str, str)

swagger_types = {'kind': 'str', 'data': 'dict(str, str)', 'string_data': 'dict(str, str)', 'type': 'str', 'api_version': 'str', 'm

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1Secret. Used to facilitate programmatic handling of secret data.

Returns The type of this V1Secret.

Return type str

kubernetes.client.models.v1_secret_key_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector(key=None,
                                                                           name=None,
                                                                           optional=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'optional': 'optional', 'name': 'name', 'key': 'key'}

key

Gets the key of this V1SecretKeySelector. The key of the secret to select from. Must be a valid secret key.

Returns The key of this V1SecretKeySelector.

Return type str

name

Gets the name of this V1SecretKeySelector. Name of the referent. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1SecretKeySelector.

Return type str

optional

Gets the optional of this V1SecretKeySelector. Specify whether the Secret or it's key must be defined

Returns The optional of this V1SecretKeySelector.

Return type bool

swagger_types = {'optional': 'bool', 'name': 'str', 'key': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_secret_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_secret_list.V1SecretList (*api_version=None, items=None, kind=None, metadata=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1SecretList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1SecretList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1SecretList. Items is a list of secret objects. More info: <https://kubernetes.io/docs/concepts/configuration/secret>

Returns The items of this V1SecretList.

Return type list[V1Secret]

kind

Gets the kind of this V1SecretList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1SecretList.

Return type str

metadata

Gets the metadata of this V1SecretList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1SecretList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Secret]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_secret_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource(default_mode=None,
                                                                              items=None,
                                                                              op-
                                                                              tional=None,
                                                                              se-
                                                                              cret_name=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'items': 'items', 'default_mode': 'defaultMode', 'optional': 'optional', 'secret_name': 'secretName'}

default_mode

Gets the default_mode of this V1SecretVolumeSource. Optional: mode bits to use on created files by default. Must be a value between 0 and 0777. Defaults to 0644. Directories within the path are not affected by this setting. This might be in conflict with other options that affect the file mode, like fsGroup, and the result can be other mode bits set.

Returns The default_mode of this V1SecretVolumeSource.

Return type int

items

Gets the items of this V1SecretVolumeSource. If unspecified, each key-value pair in the Data field of the referenced Secret will be projected into the volume as a file whose name is the key and content is the value. If specified, the listed keys will be projected into the specified paths, and unlisted keys will not be present. If a key is specified which is not present in the Secret, the volume setup will error unless it is marked optional. Paths must be relative and may not contain the '..' path or start with '..'.

Returns The items of this V1SecretVolumeSource.

Return type list[V1KeyToPath]

optional

Gets the optional of this V1SecretVolumeSource. Specify whether the Secret or it's keys must be defined

Returns The optional of this V1SecretVolumeSource.

Return type bool

secret_name

Gets the secret_name of this V1SecretVolumeSource. Name of the secret in the pod's namespace to use.

More info: <https://kubernetes.io/docs/concepts/storage/volumes#secret>

Returns The secret_name of this V1SecretVolumeSource.

Return type str

swagger_types = {'items': 'list[V1KeyToPath]', 'default_mode': 'int', 'optional': 'bool', 'secret_name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_security_context module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_security_context.V1SecurityContext (allow_privilege_escalation=None,
                                                                    capabilities=None,
                                                                    privileged=None,
                                                                    read_only_root_filesystem=None,
                                                                    run_as_non_root=None,
                                                                    run_as_user=None,
                                                                    se_linux_options=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

allow_privilege_escalation

Gets the allow_privilege_escalation of this V1SecurityContext. AllowPrivilegeEscalation controls whether a process can gain more privileges than its parent process. This bool directly controls if the no_new_privs flag will be set on the container process. AllowPrivilegeEscalation is true always when the container is: 1) run as Privileged 2) has CAP_SYS_ADMIN

Returns The allow_privilege_escalation of this V1SecurityContext.

Return type bool

attribute_map = {'allow_privilege_escalation': 'allowPrivilegeEscalation', 'capabilities': 'capabilities', 'read_only_root_filesystem': 'readOnlyRootFilesystem'}

capabilities

Gets the capabilities of this V1SecurityContext. The capabilities to add/drop when running containers. Defaults to the default set of capabilities granted by the container runtime.

Returns The capabilities of this V1SecurityContext.

Return type *V1Capabilities*

privileged

Gets the privileged of this V1SecurityContext. Run container in privileged mode. Processes in privileged containers are essentially equivalent to root on the host. Defaults to false.

Returns The privileged of this V1SecurityContext.

Return type bool

read_only_root_filesystem

Gets the read_only_root_filesystem of this V1SecurityContext. Whether this container has a read-only root filesystem. Default is false.

Returns The read_only_root_filesystem of this V1SecurityContext.

Return type bool

run_as_non_root

Gets the run_as_non_root of this V1SecurityContext. Indicates that the container must run as a non-root user. If true, the Kubelet will validate the image at runtime to ensure that it does not run as UID 0 (root) and fail to start the container if it does. If unset or false, no such validation will be performed. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.

Returns The run_as_non_root of this V1SecurityContext.

Return type bool

run_as_user

Gets the run_as_user of this V1SecurityContext. The UID to run the entrypoint of the container process. Defaults to user specified in image metadata if unspecified. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.

Returns The run_as_user of this V1SecurityContext.

Return type int

se_linux_options

Gets the se_linux_options of this V1SecurityContext. The SELinux context to be applied to the container. If unspecified, the container runtime will allocate a random SELinux context for each container. May also be set in PodSecurityContext. If set in both SecurityContext and PodSecurityContext, the value specified in SecurityContext takes precedence.

Returns The se_linux_options of this V1SecurityContext.

Return type *V1SELinuxOptions*

swagger_types = {'allow_privilege_escalation': 'bool', 'capabilities': 'V1Capabilities', 'read_only_root_filesystem': 'b

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service.V1Service (api_version=None,      kind=None,
                                                    metadata=None,      spec=None,
                                                    status=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1Service. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1Service.

Return type str

```
attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1Service. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1Service.

Return type str

metadata

Gets the metadata of this V1Service. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1Service.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1Service. Spec defines the behavior of a service. <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1Service.

Return type *V1ServiceSpec*

status

Gets the status of this V1Service. Most recently observed status of the service. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1Service.

Return type *V1ServiceStatus*

```
swagger_types = {'status': 'V1ServiceStatus', 'kind': 'str', 'spec': 'V1ServiceSpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service_account module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service_account.V1ServiceAccount (api_version=None,
                                                                    auto-
                                                                    mount_service_account_token=None,
                                                                    im-
                                                                    age_pull_secrets=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None,
                                                                    se-
                                                                    crets=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ServiceAccount. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ServiceAccount.

Return type str

attribute_map = {'kind': 'kind', 'image_pull_secrets': 'imagePullSecrets', 'secrets': 'secrets', 'automount_service_ac

automount_service_account_token

Gets the automount_service_account_token of this V1ServiceAccount. AutomountServiceAccountToken indicates whether pods running as this service account should have an API token automatically mounted. Can be overridden at the pod level.

Returns The automount_service_account_token of this V1ServiceAccount.

Return type bool

image_pull_secrets

Gets the image_pull_secrets of this V1ServiceAccount. ImagePullSecrets is a list of references to secrets in the same namespace to use for pulling any images in pods that reference this ServiceAccount. ImagePullSecrets are distinct from Secrets because Secrets can be mounted in the pod, but ImagePullSecrets are only accessed by the kubelet. More info: <https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod>

Returns The image_pull_secrets of this V1ServiceAccount.

Return type list[V1LocalObjectReference]

kind

Gets the kind of this V1ServiceAccount. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ServiceAccount.

Return type str

metadata

Gets the metadata of this V1ServiceAccount. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1ServiceAccount.

Return type *V1ObjectMeta*

secrets

Gets the secrets of this V1ServiceAccount. Secrets is the list of secrets allowed to be used by pods running using this ServiceAccount. More info: <https://kubernetes.io/docs/concepts/configuration/secret>

Returns The secrets of this V1ServiceAccount.

Return type list[*V1ObjectReference*]

swagger_types = {'kind': 'str', 'image_pull_secrets': 'list[V1LocalObjectReference]', 'secrets': 'list[V1ObjectReferen

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service_account_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service_account_list.V1ServiceAccountList (api_version=None,
                                                                              items=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ServiceAccountList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ServiceAccountList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1ServiceAccountList. List of ServiceAccounts. More info: <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

Returns The items of this V1ServiceAccountList.

Return type list[*V1ServiceAccount*]

kind

Gets the kind of this V1ServiceAccountList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ServiceAccountList.

Return type str

metadata

Gets the metadata of this V1ServiceAccountList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1ServiceAccountList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1ServiceAccount]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service_list.V1ServiceList (api_version=None,
                                                             items=None,
                                                             kind=None,      meta-
                                                             data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1ServiceList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1ServiceList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1ServiceList. List of services

Returns The items of this V1ServiceList.

Return type list[V1Service]

kind

Gets the kind of this V1ServiceList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1ServiceList.

Return type str

metadata

Gets the metadata of this V1ServiceList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1ServiceList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1Service]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service_port.V1ServicePort (name=None,
                                                             node_port=None,
                                                             port=None,          pro-
                                                             tocol=None,          tar-
                                                             get_port=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'port': 'port', 'target_port': 'targetPort', 'protocol': 'protocol', 'name': 'name', 'node_port': 'nodePort'}

name

Gets the name of this V1ServicePort. The name of this port within the service. This must be a DNS_LABEL. All ports within a ServiceSpec must have unique names. This maps to the 'Name' field in EndpointPort objects. Optional if only one ServicePort is defined on this service.

Returns The name of this V1ServicePort.

Return type str

node_port

Gets the node_port of this V1ServicePort. The port on each node on which this service is exposed when type=NodePort or LoadBalancer. Usually assigned by the system. If specified, it will be allocated to the service if unused or else creation of the service will fail. Default is to auto-allocate a port if the ServiceType of this Service requires one. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#type-nodeport>

Returns The node_port of this V1ServicePort.

Return type int

port

Gets the port of this V1ServicePort. The port that will be exposed by this service.

Returns The port of this V1ServicePort.

Return type int

protocol

Gets the protocol of this V1ServicePort. The IP protocol for this port. Supports “TCP” and “UDP”. Default is TCP.

Returns The protocol of this V1ServicePort.

Return type str

swagger_types = {'port': 'int', 'target_port': 'object', 'protocol': 'str', 'name': 'str', 'node_port': 'int'}

target_port

Gets the target_port of this V1ServicePort. Number or name of the port to access on the pods targeted by the service. Number must be in the range 1 to 65535. Name must be an IANA_SVC_NAME. If this is a string, it will be looked up as a named port in the target Pod’s container ports. If this is not specified, the value of the ‘port’ field is used (an identity map). This field is ignored for services with clusterIP=None, and should be omitted or set equal to the ‘port’ field. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#defining-a-service>

Returns The target_port of this V1ServicePort.

Return type object

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_service_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_service_spec.V1ServiceSpec (cluster_ip=None,
                                                             external_i_ps=None,
                                                             external_name=None,
                                                             external_traffic_policy=None,
                                                             health_check_node_port=None,
                                                             load_balancer_ip=None,
                                                             load_balancer_source_ranges=None,
                                                             ports=None,
                                                             publish_not_ready_addresses=None,
                                                             selector=None,
                                                             session_affinity=None,
                                                             session_affinity_config=None,
                                                             type=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'cluster_ip': 'clusterIP', 'external_i_ps': 'externalIPs', 'external_name': 'externalName', 'session_a

cluster_ip

Gets the cluster_ip of this V1ServiceSpec. clusterIP is the IP address of the service and is usually assigned randomly by the master. If an address is specified manually and is not in use by others, it will be allocated to the service; otherwise, creation of the service will fail. This field can not be changed through updates. Valid values are “None”, empty string (“”), or a valid IP address. “None” can be specified for headless services when proxying is not required. Only applies to types ClusterIP, NodePort, and LoadBalancer. Ignored if type is ExternalName. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#virtual-ips-and-service-proxies>

Returns The cluster_ip of this V1ServiceSpec.

Return type str

external_i_ps

Gets the external_i_ps of this V1ServiceSpec. externalIPs is a list of IP addresses for which nodes in the cluster will also accept traffic for this service. These IPs are not managed by Kubernetes. The user is responsible for ensuring that traffic arrives at a node with this IP. A common example is external load-balancers that are not part of the Kubernetes system.

Returns The external_i_ps of this V1ServiceSpec.

Return type list[str]

external_name

Gets the external_name of this V1ServiceSpec. externalName is the external reference that kubedns or equivalent will return as a CNAME record for this service. No proxying will be involved. Must be a valid DNS name and requires Type to be ExternalName.

Returns The external_name of this V1ServiceSpec.

Return type str

external_traffic_policy

Gets the external_traffic_policy of this V1ServiceSpec. externalTrafficPolicy denotes if this Service desires to route external traffic to node-local or cluster-wide endpoints. “Local” preserves the client source IP and avoids a second hop for LoadBalancer and Nodeport type services, but risks potentially imbalanced traffic spreading. “Cluster” obscures the client source IP and may cause a second hop to another node, but should have good overall load-spreading.

Returns The external_traffic_policy of this V1ServiceSpec.

Return type str

health_check_node_port

Gets the health_check_node_port of this V1ServiceSpec. healthCheckNodePort specifies the healthcheck nodePort for the service. If not specified, HealthCheckNodePort is created by the service api backend with the allocated nodePort. Will use user-specified nodePort value if specified by the client. Only effects when Type is set to LoadBalancer and ExternalTrafficPolicy is set to Local.

Returns The health_check_node_port of this V1ServiceSpec.

Return type int

load_balancer_ip

Gets the load_balancer_ip of this V1ServiceSpec. Only applies to Service Type: LoadBalancer LoadBalancer will get created with the IP specified in this field. This feature depends on whether the underlying cloud-provider supports specifying the loadBalancerIP when a load balancer is created. This field will be ignored if the cloud-provider does not support the feature.

Returns The load_balancer_ip of this V1ServiceSpec.

Return type str

load_balancer_source_ranges

Gets the load_balancer_source_ranges of this V1ServiceSpec. If specified and supported by the platform, this will restrict traffic through the cloud-provider load-balancer will be restricted to the specified client IPs. This field will be ignored if the cloud-provider does not support the feature.” More info: <https://kubernetes.io/docs/tasks/access-application-cluster/configure-cloud-provider-firewall/>

Returns The load_balancer_source_ranges of this V1ServiceSpec.

Return type list[str]

ports

Gets the ports of this V1ServiceSpec. The list of ports that are exposed by this service. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#virtual-ips-and-service-proxies>

Returns The ports of this V1ServiceSpec.

Return type list[V1ServicePort]

publish_not_ready_addresses

Gets the publish_not_ready_addresses of this V1ServiceSpec. publishNotReadyAddresses, when set to true, indicates that DNS implementations must publish the notReadyAddresses of subsets for the Endpoints associated with the Service. The default value is false. The primary use case for setting this field is to use a StatefulSet’s Headless Service to propagate SRV records for its Pods without respect to their readiness for purpose of peer discovery. This field will replace the service.alpha.kubernetes.io/tolerate-unready-endpoints when that annotation is deprecated and all clients have been converted to use this field.

Returns The publish_not_ready_addresses of this V1ServiceSpec.

Return type bool

selector

Gets the selector of this V1ServiceSpec. Route service traffic to pods with label keys and values matching this selector. If empty or not present, the service is assumed to have an external process managing its endpoints, which Kubernetes will not modify. Only applies to types ClusterIP, NodePort, and LoadBalancer. Ignored if type is ExternalName. More info: <https://kubernetes.io/docs/concepts/services-networking/service/>

Returns The selector of this V1ServiceSpec.

Return type dict(str, str)

session_affinity

Gets the session_affinity of this V1ServiceSpec. Supports “ClientIP” and “None”. Used to maintain session affinity. Enable client IP based session affinity. Must be ClientIP or None. Defaults to None. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#virtual-ips-and-service-proxies>

Returns The session_affinity of this V1ServiceSpec.

Return type str

session_affinity_config

Gets the session_affinity_config of this V1ServiceSpec. sessionAffinityConfig contains the configurations of session affinity.

Returns The session_affinity_config of this V1ServiceSpec.

Return type V1SessionAffinityConfig

swagger_types = {'cluster_ip': 'str', 'external_ips': 'list[str]', 'external_name': 'str', 'session_affinity_config': 'V1S

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1ServiceSpec. type determines how the Service is exposed. Defaults to ClusterIP. Valid options are ExternalName, ClusterIP, NodePort, and LoadBalancer. “ExternalName” maps to the specified externalName. “ClusterIP” allocates a cluster-internal IP address for load-balancing to endpoints. Endpoints are determined by the selector or if that is not specified, by manual construction of an Endpoints object. If clusterIP is “None”, no virtual IP is allocated and the endpoints are published as a set of endpoints rather than a stable IP. “NodePort” builds on ClusterIP and allocates a port on every node which routes to the clusterIP. “LoadBalancer” builds on NodePort and creates an external load-balancer (if supported in the current cloud) which routes to the clusterIP. More info: <https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services---service-types>

Returns The type of this V1ServiceSpec.

Return type str

kubernetes.client.models.v1_service_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_service_status.V1ServiceStatus(*load_balancer=None*)
Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'load_balancer': 'loadBalancer'}

load_balancer

Gets the load_balancer of this V1ServiceStatus. LoadBalancer contains the current status of the load-balancer, if one is present.

Returns The load_balancer of this V1ServiceStatus.

Return type *V1LoadBalancerStatus*

swagger_types = {'load_balancer': 'V1LoadBalancerStatus'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_tcp_socket_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction` (*host=None*,
port=None)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'host': 'host', 'port': 'port'}

host

Gets the host of this V1TCPSocketAction. Optional: Host name to connect to, defaults to the pod IP.

Returns The host of this V1TCPSocketAction.

Return type str

port

Gets the port of this V1TCPSocketAction. Number or name of the port to access on the container. Number must be in the range 1 to 65535. Name must be an IANA_SVC_NAME.

Returns The port of this V1TCPSocketAction.

Return type object

swagger_types = {'host': 'str', 'port': 'object'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1_volume module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_volume.V1Volume (aws_elastic_block_store=None,
                                                    azure_disk=None, azure_file=None,
                                                    cephfs=None, cinder=None, con-
                                                    fig_map=None, downward_api=None,
                                                    empty_dir=None, fc=None,
                                                    flex_volume=None, flocker=None,
                                                    gce_persistent_disk=None,
                                                    git_repo=None, glusterfs=None,
                                                    host_path=None, iscsi=None,
                                                    name=None, nfs=None, persis-
                                                    tent_volume_claim=None, pho-
                                                    ton_persistent_disk=None, port-
                                                    worx_volume=None, projected=None,
                                                    quobyte=None, rbd=None, stor-
                                                    ageos=None, vsphere_volume=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'gce_persistent_disk': 'gcePersistentDisk', 'persistent_volume_claim': 'persistentVolumeClaim', 'p

aws_elastic_block_store

Gets the aws_elastic_block_store of this V1Volume. AWSElasticBlockStore represents an AWS Disk resource that is attached to a kubelet's host machine and then exposed to the pod. More info: <https://kubernetes.io/docs/concepts/storage/volumes#awselasticblockstore>

Returns The aws_elastic_block_store of this V1Volume.

Return type *VIAWSElasticBlockStoreVolumeSource*

azure_disk

Gets the azure_disk of this V1Volume. AzureDisk represents an Azure Data Disk mount on the host and bind mount to the pod.

Returns The azure_disk of this V1Volume.

Return type *VIAzureDiskVolumeSource*

azure_file

Gets the azure_file of this V1Volume. AzureFile represents an Azure File Service mount on the host and bind mount to the pod.

Returns The azure_file of this V1Volume.

Return type *VIAzureFileVolumeSource*

cephfs

Gets the cephfs of this V1Volume. CephFS represents a Ceph FS mount on the host that shares a pod's lifetime

Returns The cephfs of this V1Volume.

Return type *VICephFSVolumeSource*

cinder

Gets the cinder of this V1Volume. Cinder represents a cinder volume attached and mounted on kubelets host machine More info: <https://releases.k8s.io/HEAD/examples/mysql-cinder-pd/README.md>

Returns The cinder of this V1Volume.

Return type *VICinderVolumeSource*

config_map

Gets the config_map of this V1Volume. ConfigMap represents a configMap that should populate this volume

Returns The config_map of this V1Volume.

Return type *V1ConfigMapVolumeSource*

downward_api

Gets the downward_api of this V1Volume. DownwardAPI represents downward API about the pod that should populate this volume

Returns The downward_api of this V1Volume.

Return type *V1DownwardAPIVolumeSource*

empty_dir

Gets the empty_dir of this V1Volume. EmptyDir represents a temporary directory that shares a pod's lifetime. More info: <https://kubernetes.io/docs/concepts/storage/volumes#emptydir>

Returns The empty_dir of this V1Volume.

Return type *V1EmptyDirVolumeSource*

fc

Gets the fc of this V1Volume. FC represents a Fibre Channel resource that is attached to a kubelet's host machine and then exposed to the pod.

Returns The fc of this V1Volume.

Return type *V1FCVolumeSource*

flex_volume

Gets the flex_volume of this V1Volume. FlexVolume represents a generic volume resource that is provisioned/attached using an exec based plugin. This is an alpha feature and may change in future.

Returns The flex_volume of this V1Volume.

Return type *V1FlexVolumeSource*

flocker

Gets the flocker of this V1Volume. Flocker represents a Flocker volume attached to a kubelet's host machine. This depends on the Flocker control service being running

Returns The flocker of this V1Volume.

Return type *V1FlockerVolumeSource*

gce_persistent_disk

Gets the gce_persistent_disk of this V1Volume. GCEPersistentDisk represents a GCE Disk resource that is attached to a kubelet's host machine and then exposed to the pod. More info: <https://kubernetes.io/docs/concepts/storage/volumes#gcepersistentdisk>

Returns The gce_persistent_disk of this V1Volume.

Return type *V1GCEPersistentDiskVolumeSource*

git_repo

Gets the git_repo of this V1Volume. GitRepo represents a git repository at a particular revision.

Returns The git_repo of this V1Volume.

Return type *V1GitRepoVolumeSource*

glusterfs

Gets the glusterfs of this V1Volume. Glusterfs represents a Glusterfs mount on the host that shares a pod's lifetime. More info: <https://releases.k8s.io/HEAD/examples/volumes/glusterfs/README.md>

Returns The glusterfs of this V1Volume.

Return type *V1GlusterfsVolumeSource*

host_path

Gets the host_path of this V1Volume. HostPath represents a pre-existing file or directory on the host machine that is directly exposed to the container. This is generally used for system agents or other privileged things that are allowed to see the host machine. Most containers will NOT need this. More info: <https://kubernetes.io/docs/concepts/storage/volumes#hostpath>

Returns The host_path of this V1Volume.

Return type *V1HostPathVolumeSource*

iscsi

Gets the iscsi of this V1Volume. ISCSI represents an ISCSI Disk resource that is attached to a kubelet's host machine and then exposed to the pod. More info: <https://releases.k8s.io/HEAD/examples/volumes/iscsi/README.md>

Returns The iscsi of this V1Volume.

Return type *V1ISCSIVolumeSource*

name

Gets the name of this V1Volume. Volume's name. Must be a DNS_LABEL and unique within the pod. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/names/#names>

Returns The name of this V1Volume.

Return type str

nfs

Gets the nfs of this V1Volume. NFS represents an NFS mount on the host that shares a pod's lifetime. More info: <https://kubernetes.io/docs/concepts/storage/volumes#nfs>

Returns The nfs of this V1Volume.

Return type *V1NFSSource*

persistent_volume_claim

Gets the persistent_volume_claim of this V1Volume. PersistentVolumeClaimVolumeSource represents a reference to a PersistentVolumeClaim in the same namespace. More info: <https://kubernetes.io/docs/concepts/storage/persistent-volumes#persistentvolumeclaims>

Returns The persistent_volume_claim of this V1Volume.

Return type *V1PersistentVolumeClaimVolumeSource*

photon_persistent_disk

Gets the photon_persistent_disk of this V1Volume. PhotonPersistentDisk represents a PhotonController persistent disk attached and mounted on kubelets host machine

Returns The photon_persistent_disk of this V1Volume.

Return type *V1PhotonPersistentDiskVolumeSource*

portworx_volume

Gets the portworx_volume of this V1Volume. PortworxVolume represents a portworx volume attached and mounted on kubelets host machine

Returns The portworx_volume of this V1Volume.

Return type `V1PortworxVolumeSource`

projected

Gets the projected of this `V1Volume`. Items for all in one resources secrets, configmaps, and downward API

Returns The projected of this `V1Volume`.

Return type `V1ProjectedVolumeSource`

quobyte

Gets the quobyte of this `V1Volume`. Quobyte represents a Quobyte mount on the host that shares a pod's lifetime

Returns The quobyte of this `V1Volume`.

Return type `V1QuobyteVolumeSource`

rbd

Gets the rbd of this `V1Volume`. RBD represents a Rados Block Device mount on the host that shares a pod's lifetime. More info: <https://releases.k8s.io/HEAD/examples/volumes/rbd/README.md>

Returns The rbd of this `V1Volume`.

Return type `V1RBDVolumeSource`

scale_io

Gets the scale_io of this `V1Volume`. ScaleIO represents a ScaleIO persistent volume attached and mounted on Kubernetes nodes.

Returns The scale_io of this `V1Volume`.

Return type `V1ScaleIOVolumeSource`

secret

Gets the secret of this `V1Volume`. Secret represents a secret that should populate this volume. More info: <https://kubernetes.io/docs/concepts/storage/volumes#secret>

Returns The secret of this `V1Volume`.

Return type `V1SecretVolumeSource`

storageos

Gets the storageos of this `V1Volume`. StorageOS represents a StorageOS volume attached and mounted on Kubernetes nodes.

Returns The storageos of this `V1Volume`.

Return type `V1StorageOSVolumeSource`

swagger_types = {'gce_persistent_disk': 'V1GCEPersistentDiskVolumeSource', 'persistent_volume_claim': 'V1PersistentVolumeClaim', 'portworx': 'V1PortworxVolumeSource', 'projected': 'V1ProjectedVolumeSource', 'quobyte': 'V1QuobyteVolumeSource', 'rbd': 'V1RBDVolumeSource', 'scale_io': 'V1ScaleIOVolumeSource', 'secret': 'V1SecretVolumeSource', 'storageos': 'V1StorageOSVolumeSource', 'vsphere_virtual_disk': 'V1VsphereVirtualDiskVolumeSource'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

vsphere_volume

Gets the vsphere_volume of this `V1Volume`. VsphereVolume represents a vSphere volume attached and mounted on kubelets host machine

Returns The vsphere_volume of this `V1Volume`.

Return type `V1VsphereVirtualDiskVolumeSource`

kubernetes.client.models.v1_volume_mount module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1_volume_mount.V1VolumeMount (mount_path=None,  
                                                             mount_propagation=None,  
                                                             name=None,  
                                                             read_only=None,  
                                                             sub_path=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'read_only': 'readOnly', 'sub_path': 'subPath', 'mount_path': 'mountPath', 'name': 'name', 'mount_propagation': 'mountPropagation'}

mount_path

Gets the mount_path of this V1VolumeMount. Path within the container at which the volume should be mounted. Must not contain '.'.

Returns The mount_path of this V1VolumeMount.

Return type str

mount_propagation

Gets the mount_propagation of this V1VolumeMount. mountPropagation determines how mounts are propagated from the host to container and the other way around. When not set, MountPropagationHostToContainer is used. This field is alpha in 1.8 and can be reworked or removed in a future release.

Returns The mount_propagation of this V1VolumeMount.

Return type str

name

Gets the name of this V1VolumeMount. This must match the Name of a Volume.

Returns The name of this V1VolumeMount.

Return type str

read_only

Gets the read_only of this V1VolumeMount. Mounted read-only if true, read-write otherwise (false or unspecified). Defaults to false.

Returns The read_only of this V1VolumeMount.

Return type bool

sub_path

Gets the sub_path of this V1VolumeMount. Path within the volume from which the container's volume should be mounted. Defaults to "" (volume's root).

Returns The sub_path of this V1VolumeMount.

Return type str

swagger_types = {'read_only': 'bool', 'sub_path': 'str', 'mount_path': 'str', 'name': 'str', 'mount_propagation': 'str'}

to_dict()

Returns the model properties as a dict

to_str()
Returns the string representation of the model

kubernetes.client.models.v1_vsphere_virtual_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'volume_path': 'volumePath', 'fs_type': 'fsType', 'storage_policy_id': 'storagePolicyID', 'storage_policy_name': 'storagePolicyName'}

fs_type

Gets the fs_type of this V1VsphereVirtualDiskVolumeSource. Filesystem type to mount. Must be a filesystem type supported by the host operating system. Ex. "ext4", "xfs", "ntfs". Implicitly inferred to be "ext4" if unspecified.

Returns The fs_type of this V1VsphereVirtualDiskVolumeSource.

Return type str

storage_policy_id

Gets the storage_policy_id of this V1VsphereVirtualDiskVolumeSource. Storage Policy Based Management (SPBM) profile ID associated with the StoragePolicyName.

Returns The storage_policy_id of this V1VsphereVirtualDiskVolumeSource.

Return type str

storage_policy_name

Gets the storage_policy_name of this V1VsphereVirtualDiskVolumeSource. Storage Policy Based Management (SPBM) profile name.

Returns The storage_policy_name of this V1VsphereVirtualDiskVolumeSource.

Return type str

swagger_types = {'volume_path': 'str', 'fs_type': 'str', 'storage_policy_id': 'str', 'storage_policy_name': 'str'}

to_dict()
Returns the model properties as a dict

to_str()
Returns the string representation of the model

volume_path

Gets the volume_path of this V1VsphereVirtualDiskVolumeSource. Path that identifies vSphere volume vmdk

Returns The volume_path of this V1VsphereVirtualDiskVolumeSource.

Return type str

kubernetes.client.models.v1alpha1_certificate_signing_request module

kubernetes.client.models.v1alpha1_certificate_signing_request_condition module

kubernetes.client.models.v1alpha1_certificate_signing_request_list module

kubernetes.client.models.v1alpha1_certificate_signing_request_spec module

kubernetes.client.models.v1alpha1_certificate_signing_request_status module

kubernetes.client.models.v1alpha1_cluster_role module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole (api_version=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None,
                                                                              rules=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1ClusterRole. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1ClusterRole.

Return type str

attribute_map = {'rules': 'rules', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1alpha1ClusterRole. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1ClusterRole.

Return type str

metadata

Gets the metadata of this V1alpha1ClusterRole. Standard object's metadata.

Returns The metadata of this V1alpha1ClusterRole.

Return type *V1ObjectMeta*

rules

Gets the rules of this V1alpha1ClusterRole. Rules holds all the PolicyRules for this ClusterRole

Returns The rules of this V1alpha1ClusterRole.

Return type *list[V1alpha1PolicyRule]*

swagger_types = {'rules': 'list[V1alpha1PolicyRule]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_cluster_role_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding` (*api_version=*
kind=
meta-
data=
role_ref
sub-
jects=

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1ClusterRoleBinding. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1ClusterRoleBinding.

Return type `str`

attribute_map = {'kind': 'kind', 'subjects': 'subjects', 'role_ref': 'roleRef', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1alpha1ClusterRoleBinding. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1ClusterRoleBinding.

Return type `str`

metadata

Gets the metadata of this V1alpha1ClusterRoleBinding. Standard object's metadata.

Returns The metadata of this V1alpha1ClusterRoleBinding.

Return type *V1ObjectMeta*

role_ref

Gets the role_ref of this V1alpha1ClusterRoleBinding. RoleRef can only reference a ClusterRole in the global namespace. If the RoleRef cannot be resolved, the Authorizer must return an error.

Returns The role_ref of this V1alpha1ClusterRoleBinding.

Return type *V1alpha1RoleRef*

subjects

Gets the subjects of this V1alpha1ClusterRoleBinding. Subjects holds references to the objects the role applies to.

Returns The subjects of this V1alpha1ClusterRoleBinding.

Return type list[*V1alpha1Subject*]

swagger_types = {'kind': 'str', 'subjects': 'list[V1alpha1Subject]', 'role_ref': 'V1alpha1RoleRef', 'api_version': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_cluster_role_binding_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1ClusterRoleBindingList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1ClusterRoleBindingList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1alpha1ClusterRoleBindingList. Items is a list of ClusterRoleBindings

Returns The items of this V1alpha1ClusterRoleBindingList.

Return type list[V1alpha1ClusterRoleBinding]

kind

Gets the kind of this V1alpha1ClusterRoleBindingList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1ClusterRoleBindingList.

Return type str

metadata

Gets the metadata of this V1alpha1ClusterRoleBindingList. Standard object's metadata.

Returns The metadata of this V1alpha1ClusterRoleBindingList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1alpha1ClusterRoleBinding]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_cluster_role_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList (api_version=None,
items=None,
kind=None,
meta-
data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1ClusterRoleList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1ClusterRoleList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1alpha1ClusterRoleList. Items is a list of ClusterRoles

Returns The items of this V1alpha1ClusterRoleList.

Return type list[V1alpha1ClusterRole]

kind

Gets the kind of this V1alpha1ClusterRoleList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1ClusterRoleList.

Return type str

metadata

Gets the metadata of this V1alpha1ClusterRoleList. Standard object's metadata.

Returns The metadata of this V1alpha1ClusterRoleList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1alpha1ClusterRole]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_policy_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule (api_groups=None,
                                                                    non_resource_urls=None,
                                                                    re-
                                                                    source_names=None,
                                                                    re-
                                                                    sources=None,
                                                                    verbs=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_groups

Gets the api_groups of this V1alpha1PolicyRule. APIGroups is the name of the APIGroup that contains the resources. If multiple API groups are specified, any action requested against one of the enumerated resources in any API group will be allowed.

Returns The api_groups of this V1alpha1PolicyRule.

Return type list[str]

attribute_map = {'verbs': 'verbs', 'non_resource_urls': 'nonResourceURLs', 'api_groups': 'apiGroups', 'resources': 'resources'}

non_resource_urls

Gets the non_resource_urls of this V1alpha1PolicyRule. NonResourceURLs is a set of partial urls that a user should have access to. *s are allowed, but only as the full, final step in the path This name is intentionally different than the internal type so that the DefaultConvert works nicely and because the ordering may be different. Since non-resource URLs are not namespaced, this field is only applicable

for ClusterRoles referenced from a ClusterRoleBinding. Rules can either apply to API resources (such as “pods” or “secrets”) or non-resource URL paths (such as “/api”), but not both.

Returns The non_resource_ur_ls of this V1alpha1PolicyRule.

Return type list[str]

resource_names

Gets the resource_names of this V1alpha1PolicyRule. ResourceNames is an optional white list of names that the rule applies to. An empty set means that everything is allowed.

Returns The resource_names of this V1alpha1PolicyRule.

Return type list[str]

resources

Gets the resources of this V1alpha1PolicyRule. Resources is a list of resources this rule applies to. ResourceAll represents all resources.

Returns The resources of this V1alpha1PolicyRule.

Return type list[str]

swagger_types = {'verbs': 'list[str]', 'non_resource_ur_ls': 'list[str]', 'api_groups': 'list[str]', 'resources': 'list[str]', 'r

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

verbs

Gets the verbs of this V1alpha1PolicyRule. Verbs is a list of Verbs that apply to ALL the ResourceKinds and AttributeRestrictions contained in this rule. VerbAll represents all kinds.

Returns The verbs of this V1alpha1PolicyRule.

Return type list[str]

kubernetes.client.models.v1alpha1_role module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_role.V1alpha1Role (api_version=None,
                                                            kind=None,          meta-
                                                            data=None, rules=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1Role. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1Role.

Return type str

```
attribute_map = {'rules': 'rules', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1alpha1Role. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1Role.

Return type str

metadata

Gets the metadata of this V1alpha1Role. Standard object's metadata.

Returns The metadata of this V1alpha1Role.

Return type *V1ObjectMeta*

rules

Gets the rules of this V1alpha1Role. Rules holds all the PolicyRules for this Role

Returns The rules of this V1alpha1Role.

Return type list[V1alpha1PolicyRule]

```
swagger_types = {'rules': 'list[V1alpha1PolicyRule]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_role_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding (api_version=None,
                                                                            kind=None,
                                                                            meta-
                                                                            data=None,
                                                                            role_ref=None,
                                                                            sub-
                                                                            jects=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1RoleBinding. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1RoleBinding.

Return type str


```
attribute_map = {'kind': 'kind', 'subjects': 'subjects', 'role_ref': 'roleRef', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1alpha1RoleBinding. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1RoleBinding.

Return type str

metadata

Gets the metadata of this V1alpha1RoleBinding. Standard object's metadata.

Returns The metadata of this V1alpha1RoleBinding.

Return type *V1ObjectMeta*

role_ref

Gets the role_ref of this V1alpha1RoleBinding. RoleRef can reference a Role in the current namespace or a ClusterRole in the global namespace. If the RoleRef cannot be resolved, the Authorizer must return an error.

Returns The role_ref of this V1alpha1RoleBinding.

Return type *V1alpha1RoleRef*

subjects

Gets the subjects of this V1alpha1RoleBinding. Subjects holds references to the objects the role applies to.

Returns The subjects of this V1alpha1RoleBinding.

Return type list[*V1alpha1Subject*]

```
swagger_types = {'kind': 'str', 'subjects': 'list[V1alpha1Subject]', 'role_ref': 'V1alpha1RoleRef', 'api_version': 'str'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_role_binding_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList (api_version=None,
items=None,
kind=None,
meta-
data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1RoleBindingList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1RoleBindingList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1alpha1RoleBindingList. Items is a list of RoleBindings

Returns The items of this V1alpha1RoleBindingList.

Return type list[V1alpha1RoleBinding]

kind

Gets the kind of this V1alpha1RoleBindingList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1RoleBindingList.

Return type str

metadata

Gets the metadata of this V1alpha1RoleBindingList. Standard object's metadata.

Returns The metadata of this V1alpha1RoleBindingList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1alpha1RoleBinding]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_role_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1RoleList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1alpha1RoleList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1alpha1RoleList. Items is a list of Roles

Returns The items of this V1alpha1RoleList.

Return type list[V1alpha1Role]

kind

Gets the kind of this V1alpha1RoleList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1alpha1RoleList.

Return type str

metadata

Gets the metadata of this V1alpha1RoleList. Standard object's metadata.

Returns The metadata of this V1alpha1RoleList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1alpha1Role]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_role_ref module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef(api_group=None,
                                                                kind=None,
                                                                name=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_group

Gets the api_group of this V1alpha1RoleRef. APIGroup is the group for the resource being referenced

Returns The api_group of this V1alpha1RoleRef.

Return type str

attribute_map = {'kind': 'kind', 'api_group': 'apiGroup', 'name': 'name'}

kind

Gets the kind of this V1alpha1RoleRef. Kind is the type of resource being referenced

Returns The kind of this V1alpha1RoleRef.

Return type str

name

Gets the name of this V1alpha1RoleRef. Name is the name of resource being referenced

Returns The name of this V1alpha1RoleRef.

Return type str

swagger_types = {'kind': 'str', 'api_group': 'str', 'name': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1alpha1_subject module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1alpha1_subject.V1alpha1Subject (api_version=None,
                                                                kind=None,
                                                                name=None,
                                                                namespace=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1alpha1Subject. APIVersion holds the API group and version of the referenced subject. Defaults to “v1” for ServiceAccount subjects. Defaults to “rbac.authorization.k8s.io/v1alpha1” for User and Group subjects.

Returns The api_version of this V1alpha1Subject.

Return type str

attribute_map = {'kind': 'kind', 'namespace': 'namespace', 'name': 'name', 'api_version': 'apiVersion'}

kind

Gets the kind of this V1alpha1Subject. Kind of object being referenced. Values defined by this API group are “User”, “Group”, and “ServiceAccount”. If the Authorizer does not recognized the kind value, the Authorizer should report an error.

Returns The kind of this V1alpha1Subject.

Return type str

name

Gets the name of this V1alpha1Subject. Name of the object being referenced.

Returns The name of this V1alpha1Subject.

Return type str

namespace

Gets the namespace of this V1alpha1Subject. Namespace of the referenced object. If the object kind is non-namespace, such as “User” or “Group”, and this value is not empty the Authorizer should report an error.

Returns The namespace of this V1alpha1Subject.

Return type str

swagger_types = {'kind': 'str', 'namespace': 'str', 'name': 'str', 'api_version': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_api_version module

kubernetes.client.models.v1beta1_cpu_target_utilization module

kubernetes.client.models.v1beta1_daemon_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet (api_version=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None,
                                                                    spec=None,
                                                                    sta-
                                                                    tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1DaemonSet. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1DaemonSet.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1DaemonSet. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1DaemonSet.

Return type str

metadata

Gets the metadata of this V1beta1DaemonSet. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1DaemonSet.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1DaemonSet. The desired behavior of this daemon set. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1beta1DaemonSet.

Return type *V1beta1DaemonSetSpec*

status

Gets the status of this V1beta1DaemonSet. The current status of this daemon set. This data may be out of date by some window of time. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1beta1DaemonSet.

Return type *V1beta1DaemonSetStatus*

swagger_types = {'status': 'V1beta1DaemonSetStatus', 'kind': 'str', 'spec': 'V1beta1DaemonSetSpec', 'api_version':

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1beta1_daemon_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList (api_version=None,
                                                                              items=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1DaemonSetList. APIVersion defines the versioned schema of this

representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1DaemonSetList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1DaemonSetList. A list of daemon sets.

Returns The items of this V1beta1DaemonSetList.

Return type list[V1beta1DaemonSet]

kind

Gets the kind of this V1beta1DaemonSetList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1DaemonSetList.

Return type str

metadata

Gets the metadata of this V1beta1DaemonSetList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1DaemonSetList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1DaemonSet]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_daemon_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec (min_ready_seconds=None,
re-
vi-
sion_history_limit=None,
se-
lec-
tor=None,
tem-
plate=None,
tem-
plate_generation=None,
up-
date_strategy=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'min_ready_seconds': 'minReadySeconds', 'update_strategy': 'updateStrategy', 'selector': 'selector'}

min_ready_seconds

Gets the min_ready_seconds of this V1beta1DaemonSetSpec. The minimum number of seconds for which a newly created DaemonSet pod should be ready without any of its container crashing, for it to be considered available. Defaults to 0 (pod will be considered available as soon as it is ready).

Returns The min_ready_seconds of this V1beta1DaemonSetSpec.

Return type int

revision_history_limit

Gets the revision_history_limit of this V1beta1DaemonSetSpec. The number of old history to retain to allow rollback. This is a pointer to distinguish between explicit zero and not specified. Defaults to 10.

Returns The revision_history_limit of this V1beta1DaemonSetSpec.

Return type int

selector

Gets the selector of this V1beta1DaemonSetSpec. A label query over pods that are managed by the daemon set. Must match in order to be controlled. If empty, defaulted to labels on Pod template. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#label-selectors>

Returns The selector of this V1beta1DaemonSetSpec.

Return type V1LabelSelector

swagger_types = {'min_ready_seconds': 'int', 'update_strategy': 'V1beta1DaemonSetUpdateStrategy', 'selector': 'V1LabelSelector'}

template

Gets the template of this V1beta1DaemonSetSpec. An object that describes the pod that will be created. The DaemonSet will create exactly one copy of this pod on every node that matches the template's node selector (or on every node if no node selector is specified). More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller#pod-template>

Returns The template of this V1beta1DaemonSetSpec.

Return type V1PodTemplateSpec

template_generation

Gets the template_generation of this V1beta1DaemonSetSpec. DEPRECATED. A sequence number representing a specific generation of the template. Populated by the system. It can be set only during the creation.

Returns The template_generation of this V1beta1DaemonSetSpec.

Return type int

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

update_strategy

Gets the update_strategy of this V1beta1DaemonSetSpec. An update strategy to replace existing DaemonSet pods with new pods.

Returns The update_strategy of this V1beta1DaemonSetSpec.

Return type V1beta1DaemonSetUpdateStrategy

kubernetes.client.models.v1beta1_daemon_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus (collision_count=None,
                                          current_number_scheduled=None,
                                          desired_number_scheduled=None,
                                          number_available=None,
                                          number_misscheduled=None,
                                          number_ready=None,
                                          number_unavailable=None,
                                          observed_generation=None,
                                          updated_number_scheduled=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'updated_number_scheduled': 'updatedNumberScheduled', 'number_available': 'numberAvailable'}

collision_count

Gets the collision_count of this V1beta1DaemonSetStatus. Count of hash collisions for the DaemonSet. The DaemonSet controller uses this field as a collision avoidance mechanism when it needs to create the name for the newest ControllerRevision.

Returns The collision_count of this V1beta1DaemonSetStatus.

Return type int

current_number_scheduled

Gets the current_number_scheduled of this V1beta1DaemonSetStatus. The number of nodes that are running at least 1 daemon pod and are supposed to run the daemon pod. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

Returns The current_number_scheduled of this V1beta1DaemonSetStatus.

Return type int

desired_number_scheduled

Gets the desired_number_scheduled of this V1beta1DaemonSetStatus. The total number of nodes that should be running the daemon pod (including nodes correctly running the daemon pod). More info: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

Returns The desired_number_scheduled of this V1beta1DaemonSetStatus.

Return type int

number_available

Gets the number_available of this V1beta1DaemonSetStatus. The number of nodes that should be running the daemon pod and have one or more of the daemon pod running and available (ready for at least spec.minReadySeconds)

Returns The number_available of this V1beta1DaemonSetStatus.

Return type int

number_misscheduled

Gets the number_misscheduled of this V1beta1DaemonSetStatus. The number of nodes that are running the daemon pod, but are not supposed to run the daemon pod. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>

Returns The number_misscheduled of this V1beta1DaemonSetStatus.

Return type int

number_ready

Gets the number_ready of this V1beta1DaemonSetStatus. The number of nodes that should be running the daemon pod and have one or more of the daemon pod running and ready.

Returns The number_ready of this V1beta1DaemonSetStatus.

Return type int

number_unavailable

Gets the number_unavailable of this V1beta1DaemonSetStatus. The number of nodes that should be running the daemon pod and have none of the daemon pod running and available (ready for at least spec.minReadySeconds)

Returns The number_unavailable of this V1beta1DaemonSetStatus.

Return type int

observed_generation

Gets the observed_generation of this V1beta1DaemonSetStatus. The most recent generation observed by the daemon set controller.

Returns The observed_generation of this V1beta1DaemonSetStatus.

Return type int

swagger_types = {'updated_number_scheduled': 'int', 'number_available': 'int', 'collision_count': 'int', 'number_un

to_dict ()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

updated_number_scheduled

Gets the updated_number_scheduled of this V1beta1DaemonSetStatus. The total number of nodes that are running updated daemon pod

Returns The updated_number_scheduled of this V1beta1DaemonSetStatus.

Return type int

kubernetes.client.models.v1beta1_deployment module

kubernetes.client.models.v1beta1_deployment_condition module

kubernetes.client.models.v1beta1_deployment_list module

kubernetes.client.models.v1beta1_deployment_rollback module

kubernetes.client.models.v1beta1_deployment_spec module

kubernetes.client.models.v1beta1_deployment_status module

kubernetes.client.models.v1beta1_deployment_strategy module

kubernetes.client.models.v1beta1_eviction module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1beta1_eviction.V1beta1Eviction` (*api_version=None, delete_options=None, kind=None, metadata=None*)

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1Eviction. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1Eviction.

Return type str

attribute_map = {'delete_options': 'deleteOptions', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

delete_options

Gets the delete_options of this V1beta1Eviction. DeleteOptions may be provided

Returns The delete_options of this V1beta1Eviction.

Return type *V1DeleteOptions*

kind

Gets the kind of this V1beta1Eviction. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1Eviction.

Return type str

metadata

Gets the metadata of this V1beta1Eviction. ObjectMeta describes the pod that is being evicted.

Returns The metadata of this V1beta1Eviction.

Return type *V1ObjectMeta*

swagger_types = {'delete_options': 'V1DeleteOptions', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_horizontal_pod_autoscaler module

kubernetes.client.models.v1beta1_horizontal_pod_autoscaler_list module

kubernetes.client.models.v1beta1_horizontal_pod_autoscaler_spec module

kubernetes.client.models.v1beta1_horizontal_pod_autoscaler_status module

kubernetes.client.models.v1beta1_http_ingress_path module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1beta1_http_ingress_path.V1beta1HTTPIngressPath` (*backend=None, path=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'backend': 'backend'}

backend

Gets the backend of this V1beta1HTTPIngressPath. Backend defines the referenced service endpoint to which the traffic will be forwarded to.

Returns The backend of this V1beta1HTTPIngressPath.

Return type *V1beta1IngressBackend*

path

Gets the path of this V1beta1HTTPIngressPath. Path is an extended POSIX regex as defined by IEEE Std 1003.1, (i.e this follows the egrep/unix syntax, not the perl syntax) matched against the path of an incoming request. Currently it can contain characters disallowed from the conventional “path” part of a URL as defined by RFC 3986. Paths must begin with a ‘/’. If unspecified, the path defaults to a catch all sending traffic to the backend.

Returns The path of this V1beta1HTTPIngressPath.

Return type str

swagger_types = {'path': 'str', 'backend': 'V1beta1IngressBackend'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_http_ingress_rule_value module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_http_ingress_rule_value.V1beta1HTTPIngressRuleValue (pa

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'paths': 'paths'}

paths

Gets the paths of this V1beta1HTTPIngressRuleValue. A collection of paths that map requests to backends.

Returns The paths of this V1beta1HTTPIngressRuleValue.

Return type list[V1beta1HTTPIngressPath]

swagger_types = {'paths': 'list[V1beta1HTTPIngressPath]'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_ingress.V1beta1Ingress (api_version=None,
                                                                kind=None,    meta-
                                                                data=None,
                                                                spec=None,    sta-
                                                                tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1Ingress. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1Ingress.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1Ingress. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1Ingress.

Return type str

metadata

Gets the metadata of this V1beta1Ingress. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1Ingress.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1Ingress. Spec is the desired state of the Ingress. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1beta1Ingress.

Return type *V1beta1IngressSpec*

status

Gets the status of this V1beta1Ingress. Status is the current state of the Ingress. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1beta1Ingress.

Return type *V1beta1IngressStatus*

swagger_types = {'status': 'V1beta1IngressStatus', 'kind': 'str', 'spec': 'V1beta1IngressSpec', 'api_version': 'str', 'm

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_backend module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_ingress_backend.V1beta1IngressBackend (service_name=None,
                                                                    ser-
                                                                    vice_port=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'service_name': 'serviceName', 'service_port': 'servicePort'}

service_name

Gets the service_name of this V1beta1IngressBackend. Specifies the name of the referenced service.

Returns The service_name of this V1beta1IngressBackend.

Return type str

service_port

Gets the service_port of this V1beta1IngressBackend. Specifies the port of the referenced service.

Returns The service_port of this V1beta1IngressBackend.

Return type object

swagger_types = {'service_name': 'str', 'service_port': 'object'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList (api_version=None,
                                                                    items=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1IngressList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may

reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1IngressList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1IngressList. Items is the list of Ingress.

Returns The items of this V1beta1IngressList.

Return type list[V1beta1Ingress]

kind

Gets the kind of this V1beta1IngressList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1IngressList.

Return type str

metadata

Gets the metadata of this V1beta1IngressList. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1IngressList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1Ingress]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_ingress_rule.V1beta1IngressRule (*host=None, http=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'host': 'host', 'http': 'http'}

host

Gets the host of this V1beta1IngressRule. Host is the fully qualified domain name of a network host, as defined by RFC 3986. Note the following deviations from the "host" part of the URI as defined in the RFC: 1. IPs are not allowed. Currently an IngressRuleValue can only apply to the IP in the Spec of the parent Ingress. 2. The : delimiter is not respected because ports are not allowed. Currently the port of

an Ingress is implicitly :80 for http and :443 for https. Both these may change in the future. Incoming requests are matched against the host before the IngressRuleValue. If the host is unspecified, the Ingress routes all traffic based on the specified IngressRuleValue.

Returns The host of this V1beta1IngressRule.

Return type str

http

Gets the http of this V1beta1IngressRule.

Returns The http of this V1beta1IngressRule.

Return type *V1beta1HTTPIngressRuleValue*

swagger_types = {'host': 'str', 'http': 'V1beta1HTTPIngressRuleValue'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec (*backend=None, rules=None, tls=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'rules': 'rules', 'tls': 'tls', 'backend': 'backend'}

backend

Gets the backend of this V1beta1IngressSpec. A default backend capable of servicing requests that don't match any rule. At least one of 'backend' or 'rules' must be specified. This field is optional to allow the loadbalancer controller or defaulting logic to specify a global default.

Returns The backend of this V1beta1IngressSpec.

Return type *V1beta1IngressBackend*

rules

Gets the rules of this V1beta1IngressSpec. A list of host rules used to configure the Ingress. If unspecified, or no rule matches, all traffic is sent to the default backend.

Returns The rules of this V1beta1IngressSpec.

Return type list[*V1beta1IngressRule*]

swagger_types = {'rules': 'list[V1beta1IngressRule]', 'tls': 'list[V1beta1IngressTLS]', 'backend': 'V1beta1IngressBackend'}

tls

Gets the tls of this V1beta1IngressSpec. TLS configuration. Currently the Ingress only supports a single TLS port, 443. If multiple members of this list specify different hosts, they will be multiplexed on the

same port according to the hostname specified through the SNI TLS extension, if the ingress controller fulfilling the ingress supports SNI.

Returns The tls of this `V1beta1IngressSpec`.

Return type list[*V1beta1IngressTLS*]

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1beta1_ingress_status.V1beta1IngressStatus` (*load_balancer=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'load_balancer': 'loadBalancer'}

load_balancer

Gets the load_balancer of this `V1beta1IngressStatus`. LoadBalancer contains the current status of the load-balancer.

Returns The load_balancer of this `V1beta1IngressStatus`.

Return type *V1LoadBalancerStatus*

swagger_types = {'load_balancer': 'V1LoadBalancerStatus'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_ingress_tls module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v1beta1_ingress_tls.V1beta1IngressTLS` (*hosts=None, se-cret_name=None*)
Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'hosts': 'hosts', 'secret_name': 'secretName'}
```

hosts

Gets the hosts of this V1beta1IngressTLS. Hosts are a list of hosts included in the TLS certificate. The values in this list must match the name/s used in the tlsSecret. Defaults to the wildcard host setting for the loadbalancer controller fulfilling this Ingress, if left unspecified.

Returns The hosts of this V1beta1IngressTLS.

Return type list[str]

secret_name

Gets the secret_name of this V1beta1IngressTLS. SecretName is the name of the secret used to terminate SSL traffic on 443. Field is left optional to allow SSL routing based on SNI hostname alone. If the SNI host in a listener conflicts with the “Host” header field used by an IngressRule, the SNI host is used for termination and value of the Host header is used for routing.

Returns The secret_name of this V1beta1IngressTLS.

Return type str

```
swagger_types = {'hosts': 'list[str]', 'secret_name': 'str'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_job module

kubernetes.client.models.v1beta1_job_condition module

kubernetes.client.models.v1beta1_job_list module

kubernetes.client.models.v1beta1_job_spec module

kubernetes.client.models.v1beta1_job_status module

kubernetes.client.models.v1beta1_local_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the `api_version` of this `V1beta1LocalSubjectAccessReview`. `APIVersion` defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The `api_version` of this `V1beta1LocalSubjectAccessReview`.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the `kind` of this `V1beta1LocalSubjectAccessReview`. `Kind` is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The `kind` of this `V1beta1LocalSubjectAccessReview`.

Return type str

metadata

Gets the `metadata` of this `V1beta1LocalSubjectAccessReview`.

Returns The `metadata` of this `V1beta1LocalSubjectAccessReview`.

Return type *V1ObjectMeta*

spec

Gets the `spec` of this `V1beta1LocalSubjectAccessReview`. `Spec` holds information about the request being evaluated. `spec.namespace` must be equal to the namespace you made the request against. If empty, it is defaulted.

Returns The `spec` of this `V1beta1LocalSubjectAccessReview`.

Return type *V1beta1SubjectAccessReviewSpec*

status

Gets the `status` of this `V1beta1LocalSubjectAccessReview`. `Status` is filled in by the server and indicates whether the request is allowed or not

Returns The `status` of this `V1beta1LocalSubjectAccessReview`.

Return type *V1beta1SubjectAccessReviewStatus*

swagger_types = {'status': 'V1beta1SubjectAccessReviewStatus', 'kind': 'str', 'spec': 'V1beta1SubjectAccessReviewS

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy (api_version=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None,
                                                                              spec=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1NetworkPolicy. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1NetworkPolicy.

Return type str

```
attribute_map = {'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1beta1NetworkPolicy. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1NetworkPolicy.

Return type str

metadata

Gets the metadata of this V1beta1NetworkPolicy. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1NetworkPolicy.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1NetworkPolicy. Specification of the desired behavior for this NetworkPolicy.

Returns The spec of this V1beta1NetworkPolicy.

Return type *V1beta1NetworkPolicySpec*

```
swagger_types = {'kind': 'str', 'spec': 'V1beta1NetworkPolicySpec', 'api_version': 'str', 'metadata': 'V1ObjectMeta'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy_ingress_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_network_policy_ingress_rule.V1beta1NetworkPolicyIngressRule
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'_from': 'from', 'ports': 'ports'}

ports

Gets the ports of this V1beta1NetworkPolicyIngressRule. List of ports which should be made accessible on the pods selected for this rule. Each item in this list is combined using a logical OR. If this field is empty or missing, this rule matches all ports (traffic not restricted by port). If this field is present and contains at least one item, then this rule allows traffic only if the traffic matches at least one port in the list.

Returns The ports of this V1beta1NetworkPolicyIngressRule.

Return type list[V1beta1NetworkPolicyPort]

swagger_types = {'_from': 'list[V1beta1NetworkPolicyPeer]', 'ports': 'list[V1beta1NetworkPolicyPort]'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList (api_version=api_version,
                                                                                       items=items,
                                                                                       kind=kind,
                                                                                       meta-data=meta-data)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1NetworkPolicyList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1NetworkPolicyList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1NetworkPolicyList. Items is a list of schema objects.

Returns The items of this V1beta1NetworkPolicyList.

Return type list[V1beta1NetworkPolicy]

kind

Gets the kind of this V1beta1NetworkPolicyList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1NetworkPolicyList.

Return type str

metadata

Gets the metadata of this V1beta1NetworkPolicyList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1NetworkPolicyList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1NetworkPolicy]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy_peer module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_network_policy_peer.V1beta1NetworkPolicyPeer (*ip_block=None, namespace_selector=None, pod_selector=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'pod_selector': 'podSelector', 'ip_block': 'ipBlock', 'namespace_selector': 'namespaceSelector'}

ip_block

Gets the ip_block of this V1beta1NetworkPolicyPeer. IPBlock defines policy on a particular IPBlock

Returns The ip_block of this V1beta1NetworkPolicyPeer.

Return type V1beta1IPBlock

namespace_selector

Gets the namespace_selector of this V1beta1NetworkPolicyPeer. Selects Namespaces using cluster scoped-labels. This matches all pods in all namespaces selected by this label selector. This field follows standard label selector semantics. If present but empty, this selector selects all namespaces.

Returns The namespace_selector of this V1beta1NetworkPolicyPeer.

Return type V1LabelSelector

pod_selector

Gets the pod_selector of this V1beta1NetworkPolicyPeer. This is a label selector which selects Pods in this namespace. This field follows standard label selector semantics. If present but empty, this selector selects all pods in this namespace.

Returns The pod_selector of this V1beta1NetworkPolicyPeer.

Return type V1LabelSelector

swagger_types = {'pod_selector': 'V1LabelSelector', 'ip_block': 'V1beta1IPBlock', 'namespace_selector': 'V1LabelS

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_network_policy_port.V1beta1NetworkPolicyPort (*port=None, pro-to-col=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'protocol': 'protocol', 'port': 'port'}

port

Gets the port of this V1beta1NetworkPolicyPort. If specified, the port on the given protocol. This can either be a numerical or named port on a pod. If this field is not provided, this matches all port names and numbers. If present, only traffic on the specified protocol AND port will be matched.

Returns The port of this V1beta1NetworkPolicyPort.

Return type object

protocol

Gets the protocol of this V1beta1NetworkPolicyPort. Optional. The protocol (TCP or UDP) which traffic must match. If not specified, this field defaults to TCP.

Returns The protocol of this V1beta1NetworkPolicyPort.

Return type str

swagger_types = {'protocol': 'str', 'port': 'object'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_network_policy_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_network_policy_spec.V1beta1NetworkPolicySpec (egress=None,
                                             ingress=None,
                                             pod_selector=None,
                                             policy_types=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'policy_types': 'policyTypes', 'ingress': 'ingress', 'egress': 'egress', 'pod_selector': 'podSelector'}

egress

Gets the egress of this V1beta1NetworkPolicySpec. List of egress rules to be applied to the selected pods. Outgoing traffic is allowed if there are no NetworkPolicies selecting the pod (and cluster policy otherwise allows the traffic), OR if the traffic matches at least one egress rule across all of the NetworkPolicy objects whose podSelector matches the pod. If this field is empty then this NetworkPolicy limits all outgoing traffic (and serves solely to ensure that the pods it selects are isolated by default). This field is beta-level in 1.8

Returns The egress of this V1beta1NetworkPolicySpec.

Return type list[V1beta1NetworkPolicyEgressRule]

ingress

Gets the ingress of this V1beta1NetworkPolicySpec. List of ingress rules to be applied to the selected pods. Traffic is allowed to a pod if there are no NetworkPolicies selecting the pod OR if the traffic source is the pod's local node, OR if the traffic matches at least one ingress rule across all of the NetworkPolicy objects whose podSelector matches the pod. If this field is empty then this NetworkPolicy does not allow any traffic (and serves solely to ensure that the pods it selects are isolated by default).

Returns The ingress of this V1beta1NetworkPolicySpec.

Return type list[V1beta1NetworkPolicyIngressRule]

pod_selector

Gets the pod_selector of this V1beta1NetworkPolicySpec. Selects the pods to which this NetworkPolicy object applies. The array of ingress rules is applied to any pods selected by this field. Multiple network policies can select the same set of pods. In this case, the ingress rules for each are combined additively. This field is NOT optional and follows standard label selector semantics. An empty podSelector matches all pods in this namespace.

Returns The pod_selector of this V1beta1NetworkPolicySpec.

Return type V1LabelSelector

policy_types

Gets the policy_types of this V1beta1NetworkPolicySpec. List of rule types that the NetworkPolicy relates to. Valid options are Ingress, Egress, or Ingress,Egress. If this field is not specified, it will default based on the existence of Ingress or Egress rules; policies that contain an Egress section are assumed to affect Egress, and all policies (whether or not they contain an Ingress section) are assumed to affect Ingress. If you want to write an egress-only policy, you must explicitly specify policyTypes ["Egress"]. Likewise, if you want to write a policy that specifies that no egress is allowed, you must specify a policyTypes value

that include “Egress” (since such a policy would not include an Egress section and would otherwise default to just [“Ingress”]). This field is beta-level in 1.8

Returns The policy_types of this V1beta1NetworkPolicySpec.

Return type list[str]

swagger_types = {'policy_types': 'list[str]', 'ingress': 'list[V1beta1NetworkPolicyIngressRule]', 'egress': 'list[V1beta1NetworkPolicyEgressRule]'

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1beta1_non_resource_attributes module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes (p

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'path': 'path', 'verb': 'verb'}

path

Gets the path of this V1beta1NonResourceAttributes. Path is the URL path of the request

Returns The path of this V1beta1NonResourceAttributes.

Return type str

swagger_types = {'path': 'str', 'verb': 'str'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

verb

Gets the verb of this V1beta1NonResourceAttributes. Verb is the standard HTTP verb

Returns The verb of this V1beta1NonResourceAttributes.

Return type str

kubernetes.client.models.v1beta1_pod_disruption_budget module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget (api_version=api_version,
kind=kind,
meta=meta,
data=data,
spec=spec,
status=status)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1PodDisruptionBudget. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1PodDisruptionBudget.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1PodDisruptionBudget. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1PodDisruptionBudget.

Return type str

metadata

Gets the metadata of this V1beta1PodDisruptionBudget.

Returns The metadata of this V1beta1PodDisruptionBudget.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1PodDisruptionBudget. Specification of the desired behavior of the PodDisruptionBudget.

Returns The spec of this V1beta1PodDisruptionBudget.

Return type *V1beta1PodDisruptionBudgetSpec*

status

Gets the status of this V1beta1PodDisruptionBudget. Most recently observed status of the PodDisruptionBudget.

Returns The status of this V1beta1PodDisruptionBudget.

Return type *V1beta1PodDisruptionBudgetStatus*

swagger_types = {'status': 'V1beta1PodDisruptionBudgetStatus', 'kind': 'str', 'spec': 'V1beta1PodDisruptionBudgetSpec'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_pod_disruption_budget_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1PodDisruptionBudgetList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1PodDisruptionBudgetList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1PodDisruptionBudgetList.

Returns The items of this V1beta1PodDisruptionBudgetList.

Return type list[V1beta1PodDisruptionBudget]

kind

Gets the kind of this V1beta1PodDisruptionBudgetList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1PodDisruptionBudgetList.

Return type str

metadata

Gets the metadata of this V1beta1PodDisruptionBudgetList.

Returns The metadata of this V1beta1PodDisruptionBudgetList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1PodDisruptionBudget]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_pod_disruption_budget_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'min_available': 'minAvailable', 'max_unavailable': 'maxUnavailable', 'selector': 'selector'}

max_unavailable

Gets the max_unavailable of this V1beta1PodDisruptionBudgetSpec. An eviction is allowed if at most “maxUnavailable” pods selected by “selector” are unavailable after the eviction, i.e. even in absence of the evicted pod. For example, one can prevent all voluntary evictions by specifying 0. This is a mutually exclusive setting with “minAvailable”.

Returns The max_unavailable of this V1beta1PodDisruptionBudgetSpec.

Return type object

min_available

Gets the min_available of this V1beta1PodDisruptionBudgetSpec. An eviction is allowed if at least “minAvailable” pods selected by “selector” will still be available after the eviction, i.e. even in the absence of the evicted pod. So for example you can prevent all voluntary evictions by specifying “100%”.

Returns The min_available of this V1beta1PodDisruptionBudgetSpec.

Return type object

selector

Gets the selector of this V1beta1PodDisruptionBudgetSpec. Label query over pods whose evictions are managed by the disruption budget.

Returns The selector of this V1beta1PodDisruptionBudgetSpec.

Return type V1LabelSelector

swagger_types = {'min_available': 'object', 'max_unavailable': 'object', 'selector': 'V1LabelSelector'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_pod_disruption_budget_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

`class kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'disrupted_pods': 'disruptedPods', 'desired_healthy': 'desiredHealthy', 'current_healthy': 'currentHealthy'}

current_healthy

Gets the current_healthy of this V1beta1PodDisruptionBudgetStatus. current number of healthy pods

Returns The current_healthy of this V1beta1PodDisruptionBudgetStatus.

Return type `int`

desired_healthy

Gets the desired_healthy of this V1beta1PodDisruptionBudgetStatus. minimum desired number of healthy pods

Returns The desired_healthy of this V1beta1PodDisruptionBudgetStatus.

Return type `int`

disrupted_pods

Gets the disrupted_pods of this V1beta1PodDisruptionBudgetStatus. DisruptedPods contains information about pods whose eviction was processed by the API server eviction subresource handler but has not yet been observed by the PodDisruptionBudget controller. A pod will be in this map from the time when the API server processed the eviction request to the time when the pod is seen by PDB controller as having been marked for deletion (or after a timeout). The key in the map is the name of the pod and the value is the time when the API server processed the eviction request. If the deletion didn't occur and a pod is still there it will be removed from the list automatically by PodDisruptionBudget controller after some time. If everything goes smooth this map should be empty for the most of the time. Large number of entries in the map may indicate problems with pod deletions.

Returns The disrupted_pods of this V1beta1PodDisruptionBudgetStatus.

Return type `dict(str, datetime)`

disruptions_allowed

Gets the disruptions_allowed of this V1beta1PodDisruptionBudgetStatus. Number of pod disruptions that are currently allowed.

Returns The disruptions_allowed of this V1beta1PodDisruptionBudgetStatus.

Return type `int`

expected_pods

Gets the expected_pods of this V1beta1PodDisruptionBudgetStatus. total number of pods counted by this disruption budget

Returns The expected_pods of this V1beta1PodDisruptionBudgetStatus.

Return type int

observed_generation

Gets the observed_generation of this V1beta1PodDisruptionBudgetStatus. Most recent generation observed when updating this PDB status. PodDisruptionsAllowed and other status informatio is valid only if observedGeneration equals to PDB's object generation.

Returns The observed_generation of this V1beta1PodDisruptionBudgetStatus.

Return type int

swagger_types = {'disrupted_pods': 'dict(str, datetime)', 'desired_healthy': 'int', 'current_healthy': 'int', 'observed_g

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_replica_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet (api_version=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None,
                                                                    spec=None,
                                                                    sta-
                                                                    tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1ReplicaSet. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1ReplicaSet.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1ReplicaSet. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1ReplicaSet.

Return type str

metadata

Gets the metadata of this V1beta1ReplicaSet. If the Labels of a ReplicaSet are empty, they are defaulted to be the same as the Pod(s) that the ReplicaSet manages. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1ReplicaSet.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1ReplicaSet. Spec defines the specification of the desired behavior of the ReplicaSet. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V1beta1ReplicaSet.

Return type *V1beta1ReplicaSetSpec*

status

Gets the status of this V1beta1ReplicaSet. Status is the most recently observed status of the ReplicaSet. This data may be out of date by some window of time. Populated by the system. Read-only. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V1beta1ReplicaSet.

Return type *V1beta1ReplicaSetStatus*

swagger_types = {'status': 'V1beta1ReplicaSetStatus', 'kind': 'str', 'spec': 'V1beta1ReplicaSetSpec', 'api_version': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_replica_set_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition (*last_transition_time*, *message*, *reason*, *status*, *type*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'status': 'status', 'message': 'message', 'type': 'type', 'reason': 'reason', 'last_transition_time': 'last_transition_time'}

last_transition_time

Gets the last_transition_time of this V1beta1ReplicaSetCondition. The last time the condition transitioned from one status to another.

Returns The last_transition_time of this V1beta1ReplicaSetCondition.

Return type datetime

message

Gets the message of this V1beta1ReplicaSetCondition. A human readable message indicating details about the transition.

Returns The message of this V1beta1ReplicaSetCondition.

Return type str

reason

Gets the reason of this V1beta1ReplicaSetCondition. The reason for the condition's last transition.

Returns The reason of this V1beta1ReplicaSetCondition.

Return type str

status

Gets the status of this V1beta1ReplicaSetCondition. Status of the condition, one of True, False, Unknown.

Returns The status of this V1beta1ReplicaSetCondition.

Return type str

swagger_types = {'status': 'str', 'message': 'str', 'type': 'str', 'reason': 'str', 'last_transition_time': 'datetime'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

type

Gets the type of this V1beta1ReplicaSetCondition. Type of replica set condition.

Returns The type of this V1beta1ReplicaSetCondition.

Return type str

kubernetes.client.models.v1beta1_replica_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList (api_version=None,
                                         items=None,
                                         kind=None,
                                         meta-
                                         data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1ReplicaSetList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1ReplicaSetList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1ReplicaSetList. List of ReplicaSets. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller>

Returns The items of this V1beta1ReplicaSetList.

Return type list[V1beta1ReplicaSet]

kind

Gets the kind of this V1beta1ReplicaSetList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1ReplicaSetList.

Return type str

metadata

Gets the metadata of this V1beta1ReplicaSetList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The metadata of this V1beta1ReplicaSetList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1ReplicaSet]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_replica_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec (min_ready_seconds=repl-
                                                                                   cas=None,
                                                                                   se-
                                                                                   lec-
                                                                                   tor=None,
                                                                                   tem-
                                                                                   plate=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'selector': 'selector', 'replicas': 'replicas', 'template': 'template', 'min_ready_seconds': 'minReadySeconds'}

min_ready_seconds

Gets the min_ready_seconds of this V1beta1ReplicaSetSpec. Minimum number of seconds for which a newly created pod should be ready without any of its container crashing, for it to be considered available. Defaults to 0 (pod will be considered available as soon as it is ready)

Returns The min_ready_seconds of this V1beta1ReplicaSetSpec.

Return type int

replicas

Gets the replicas of this V1beta1ReplicaSetSpec. Replicas is the number of desired replicas. This is a pointer to distinguish between explicit zero and unspecified. Defaults to 1. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller/#what-is-a-replicationcontroller>

Returns The replicas of this V1beta1ReplicaSetSpec.

Return type int

selector

Gets the selector of this V1beta1ReplicaSetSpec. Selector is a label query over pods that should match the replica count. If the selector is empty, it is defaulted to the labels present on the pod template. Label keys and values that must match in order to be controlled by this replica set. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#label-selectors>

Returns The selector of this V1beta1ReplicaSetSpec.

Return type V1LabelSelector

swagger_types = {'selector': 'V1LabelSelector', 'replicas': 'int', 'template': 'V1PodTemplateSpec', 'min_ready_seconds': 'int'}

template

Gets the template of this V1beta1ReplicaSetSpec. Template is the object that describes the pod that will be created if insufficient replicas are detected. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller#pod-template>

Returns The template of this V1beta1ReplicaSetSpec.

Return type *V1PodTemplateSpec*

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_replica_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus (available_replicas=
con-
di-
tions=None,
fully_labeled_re
ob-
served_generati
ready_replicas=
repli-
cas=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'replicas': 'replicas', 'observed_generation': 'observedGeneration', 'available_replicas': 'availableReplicas'}

available_replicas

Gets the available_replicas of this V1beta1ReplicaSetStatus. The number of available replicas (ready for at least minReadySeconds) for this replica set.

Returns The available_replicas of this V1beta1ReplicaSetStatus.

Return type int

conditions

Gets the conditions of this V1beta1ReplicaSetStatus. Represents the latest available observations of a replica set's current state.

Returns The conditions of this V1beta1ReplicaSetStatus.

Return type list[V1beta1ReplicaSetCondition]

fully_labeled_replicas

Gets the fully_labeled_replicas of this V1beta1ReplicaSetStatus. The number of pods that have labels matching the labels of the pod template of the replicaset.

Returns The fully_labeled_replicas of this V1beta1ReplicaSetStatus.

Return type int

observed_generation

Gets the observed_generation of this V1beta1ReplicaSetStatus. ObservedGeneration reflects the generation of the most recently observed ReplicaSet.

Returns The observed_generation of this V1beta1ReplicaSetStatus.

Return type int

ready_replicas

Gets the ready_replicas of this V1beta1ReplicaSetStatus. The number of ready replicas for this replica set.

Returns The ready_replicas of this V1beta1ReplicaSetStatus.

Return type int

replicas

Gets the replicas of this V1beta1ReplicaSetStatus. Replicas is the most recently observed number of replicas. More info: <https://kubernetes.io/docs/concepts/workloads/controllers/replicationcontroller/#what-is-a-replicationcontroller>

Returns The replicas of this V1beta1ReplicaSetStatus.

Return type int

swagger_types = {'replicas': 'int', 'observed_generation': 'int', 'available_replicas': 'int', 'ready_replicas': 'int', 'fully_labeled_replicas': 'int'}

```
to_dict()
    Returns the model properties as a dict

to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v1beta1_resource_attributes module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes (group=None,
name=None, namespace=None, verb=None, version=None)
    name=None
    namespace=None
    verb=None
    version=None
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'resource': 'resource', 'name': 'name', 'namespace': 'namespace', 'verb': 'verb', 'version': 'version'}
```

group

Gets the group of this V1beta1ResourceAttributes. Group is the API Group of the Resource. “*” means all.

Returns The group of this V1beta1ResourceAttributes.

Return type str

name

Gets the name of this V1beta1ResourceAttributes. Name is the name of the resource being requested for a “get” or deleted for a “delete”. “” (empty) means all.

Returns The name of this V1beta1ResourceAttributes.

Return type str

namespace

Gets the namespace of this V1beta1ResourceAttributes. Namespace is the namespace of the action being requested. Currently, there is no distinction between no namespace and all namespaces “” (empty) is defaulted for LocalSubjectAccessReviews “” (empty) is empty for cluster-scoped resources “” (empty) means “all” for namespace scoped resources from a SubjectAccessReview or SelfSubjectAccessReview

Returns The namespace of this V1beta1ResourceAttributes.

Return type str

resource

Gets the resource of this V1beta1ResourceAttributes. Resource is one of the existing resource types. “*” means all.

Returns The resource of this V1beta1ResourceAttributes.

Return type str

subresource

Gets the subresource of this V1beta1ResourceAttributes. Subresource is one of the existing resource types. “” means none.

Returns The subresource of this V1beta1ResourceAttributes.

Return type str

swagger_types = {'resource': 'str', 'name': 'str', 'namespace': 'str', 'verb': 'str', 'version': 'str', 'group': 'str', 'subresource': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

verb

Gets the verb of this V1beta1ResourceAttributes. Verb is a kubernetes resource API verb, like: get, list, watch, create, update, delete, proxy. “*” means all.

Returns The verb of this V1beta1ResourceAttributes.

Return type str

version

Gets the version of this V1beta1ResourceAttributes. Version is the API Version of the Resource. “*” means all.

Returns The version of this V1beta1ResourceAttributes.

Return type str

kubernetes.client.models.v1beta1_rollback_config module

kubernetes.client.models.v1beta1_rolling_update_deployment module

kubernetes.client.models.v1beta1_scale module

kubernetes.client.models.v1beta1_scale_spec module

kubernetes.client.models.v1beta1_scale_status module

kubernetes.client.models.v1beta1_self_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1SelfSubjectAccessReview. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1SelfSubjectAccessReview.

Return type str

```
attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V1beta1SelfSubjectAccessReview. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1SelfSubjectAccessReview.

Return type str

metadata

Gets the metadata of this V1beta1SelfSubjectAccessReview.

Returns The metadata of this V1beta1SelfSubjectAccessReview.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1SelfSubjectAccessReview. Spec holds information about the request being evaluated. user and groups must be empty

Returns The spec of this V1beta1SelfSubjectAccessReview.

Return type *V1beta1SelfSubjectAccessReviewSpec*

status

Gets the status of this V1beta1SelfSubjectAccessReview. Status is filled in by the server and indicates whether the request is allowed or not

Returns The status of this V1beta1SelfSubjectAccessReview.

Return type *V1beta1SubjectAccessReviewStatus*

```
swagger_types = {'status': 'V1beta1SubjectAccessReviewStatus', 'kind': 'str', 'spec': 'V1beta1SelfSubjectAccessReviewSpec'}
```

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_self_subject_access_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_self_subject_access_review_spec.V1beta1SelfSubjectAccessReviewSpec

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'resource_attributes': 'resourceAttributes', 'non_resource_attributes': 'nonResourceAttributes'}

non_resource_attributes

Gets the non_resource_attributes of this V1beta1SelfSubjectAccessReviewSpec. NonResourceAttributes describes information for a non-resource access request

Returns The non_resource_attributes of this V1beta1SelfSubjectAccessReviewSpec.

Return type *V1beta1NonResourceAttributes*

resource_attributes

Gets the resource_attributes of this V1beta1SelfSubjectAccessReviewSpec. ResourceAuthorizationAttributes describes information for a resource access request

Returns The resource_attributes of this V1beta1SelfSubjectAccessReviewSpec.

Return type *V1beta1ResourceAttributes*

swagger_types = {'resource_attributes': 'V1beta1ResourceAttributes', 'non_resource_attributes': 'V1beta1NonResourceAttributes'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_stateful_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet (*api_version=None, kind=None, meta_data=None, spec=None, status=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1StatefulSet. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1StatefulSet.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1StatefulSet. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1StatefulSet.

Return type str

metadata

Gets the metadata of this V1beta1StatefulSet.

Returns The metadata of this V1beta1StatefulSet.

Return type *V1ObjectMeta*

spec

Gets the spec of this V1beta1StatefulSet. Spec defines the desired identities of pods in this set.

Returns The spec of this V1beta1StatefulSet.

Return type *V1beta1StatefulSetSpec*

status

Gets the status of this V1beta1StatefulSet. Status is the current status of Pods in this StatefulSet. This data may be out of date by some window of time.

Returns The status of this V1beta1StatefulSet.

Return type *V1beta1StatefulSetStatus*

swagger_types = {'status': 'V1beta1StatefulSetStatus', 'kind': 'str', 'spec': 'V1beta1StatefulSetSpec', 'api_version': 'V1beta1StatefulSetSpec'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_stateful_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList (api_version=None,
                                                                                   items=None,
                                                                                   kind=None,
                                                                                   meta-
                                                                                   data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1StatefulSetList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1StatefulSetList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1StatefulSetList.

Returns The items of this V1beta1StatefulSetList.

Return type list[V1beta1StatefulSet]

kind

Gets the kind of this V1beta1StatefulSetList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1StatefulSetList.

Return type str

metadata

Gets the metadata of this V1beta1StatefulSetList.

Returns The metadata of this V1beta1StatefulSetList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1StatefulSet]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_stateful_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec (pod_management_
    repli-
    cas=None,
    re-
    vi-
    sion_history_limit=
    se-
    lec-
    tor=None,
    ser-
    vice_name=None,
    tem-
    plate=None,
    up-
    date_strategy=None,
    vol-
    ume_claim_templa
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'volume_claim_templates': 'volumeClaimTemplates', 'pod_management_policy': 'podManagementPolicy'}

pod_management_policy

Gets the pod_management_policy of this V1beta1StatefulSetSpec. podManagementPolicy controls how pods are created during initial scale up, when replacing pods on nodes, or when scaling down. The default policy is *OrderedReady*, where pods are created in increasing order (pod-0, then pod-1, etc) and the controller will wait until each pod is ready before continuing. When scaling down, the pods are removed in the opposite order. The alternative policy is *Parallel* which will create pods in parallel to match the desired scale without waiting, and on scale down will delete all pods at once.

Returns The pod_management_policy of this V1beta1StatefulSetSpec.

Return type str

replicas

Gets the replicas of this V1beta1StatefulSetSpec. replicas is the desired number of replicas of the given Template. These are replicas in the sense that they are instantiations of the same Template, but individual replicas also have a consistent identity. If unspecified, defaults to 1.

Returns The replicas of this V1beta1StatefulSetSpec.

Return type int

revision_history_limit

Gets the revision_history_limit of this V1beta1StatefulSetSpec. revisionHistoryLimit is the maximum number of revisions that will be maintained in the StatefulSet's revision history. The revision history consists of all revisions not represented by a currently applied StatefulSetSpec version. The default value is 10.

Returns The revision_history_limit of this V1beta1StatefulSetSpec.

Return type int

selector

Gets the selector of this V1beta1StatefulSetSpec. selector is a label query over pods that should match the replica count. If empty, defaulted to labels on the pod template. More info: <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/#label-selectors>

Returns The selector of this V1beta1StatefulSetSpec.

Return type `V1LabelSelector`

service_name

Gets the `service_name` of this `V1beta1StatefulSetSpec`. `serviceName` is the name of the service that governs this `StatefulSet`. This service must exist before the `StatefulSet`, and is responsible for the network identity of the set. Pods get DNS/hostnames that follow the pattern: `pod-specific-string.serviceName.default.svc.cluster.local` where “`pod-specific-string`” is managed by the `StatefulSet` controller.

Returns The `service_name` of this `V1beta1StatefulSetSpec`.

Return type `str`

swagger_types = {'volume_claim_templates': 'list[V1PersistentVolumeClaim]', 'pod_management_policy': 'str', 'revision_history_limit': 'int', 'update_strategy': 'StatefulSetUpdateStrategy'}

template

Gets the `template` of this `V1beta1StatefulSetSpec`. `template` is the object that describes the pod that will be created if insufficient replicas are detected. Each pod stamped out by the `StatefulSet` will fulfill this `Template`, but have a unique identity from the rest of the `StatefulSet`.

Returns The `template` of this `V1beta1StatefulSetSpec`.

Return type `V1PodTemplateSpec`

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

update_strategy

Gets the `update_strategy` of this `V1beta1StatefulSetSpec`. `updateStrategy` indicates the `StatefulSetUpdateStrategy` that will be employed to update Pods in the `StatefulSet` when a revision is made to `Template`.

Returns The `update_strategy` of this `V1beta1StatefulSetSpec`.

Return type `V1beta1StatefulSetUpdateStrategy`

volume_claim_templates

Gets the `volume_claim_templates` of this `V1beta1StatefulSetSpec`. `volumeClaimTemplates` is a list of claims that pods are allowed to reference. The `StatefulSet` controller is responsible for mapping network identities to claims in a way that maintains the identity of a pod. Every claim in this list must have at least one matching (by name) `volumeMount` in one container in the `template`. A claim in this list takes precedence over any volumes in the `template`, with the same name.

Returns The `volume_claim_templates` of this `V1beta1StatefulSetSpec`.

Return type `list[V1PersistentVolumeClaim]`

kubernetes.client.models.v1beta1_stateful_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus (collision_count=collision_count,
current_replicas=current_replicas,
current_revision=current_revision,
observed_generation=observed_generation,
ready_replicas=ready_replicas,
replicas=replicas,
cas=None,
update_revision=update_revision,
update_replicas=update_replicas)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'collision_count': 'collisionCount', 'update_revision': 'updateRevision', 'current_replicas': 'currentReplicas'}
```

collision_count

Gets the collision_count of this V1beta1StatefulSetStatus. collisionCount is the count of hash collisions for the StatefulSet. The StatefulSet controller uses this field as a collision avoidance mechanism when it needs to create the name for the newest ControllerRevision.

Returns The collision_count of this V1beta1StatefulSetStatus.

Return type int

current_replicas

Gets the current_replicas of this V1beta1StatefulSetStatus. currentReplicas is the number of Pods created by the StatefulSet controller from the StatefulSet version indicated by currentRevision.

Returns The current_replicas of this V1beta1StatefulSetStatus.

Return type int

current_revision

Gets the current_revision of this V1beta1StatefulSetStatus. currentRevision, if not empty, indicates the version of the StatefulSet used to generate Pods in the sequence [0,currentReplicas).

Returns The current_revision of this V1beta1StatefulSetStatus.

Return type str

observed_generation

Gets the observed_generation of this V1beta1StatefulSetStatus. observedGeneration is the most recent generation observed for this StatefulSet. It corresponds to the StatefulSet's generation, which is updated on mutation by the API Server.

Returns The observed_generation of this V1beta1StatefulSetStatus.

Return type int

ready_replicas

Gets the ready_replicas of this V1beta1StatefulSetStatus. readyReplicas is the number of Pods created by the StatefulSet controller that have a Ready Condition.

Returns The ready_replicas of this V1beta1StatefulSetStatus.

Return type int

replicas

Gets the replicas of this V1beta1StatefulSetStatus. replicas is the number of Pods created by the StatefulSet controller.

Returns The replicas of this V1beta1StatefulSetStatus.

Return type int

swagger_types = {'collision_count': 'int', 'update_revision': 'str', 'current_replicas': 'int', 'replicas': 'int', 'ready_replicas': 'int'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

update_revision

Gets the update_revision of this V1beta1StatefulSetStatus. updateRevision, if not empty, indicates the version of the StatefulSet used to generate Pods in the sequence [replicas-updatedReplicas,replicas)

Returns The update_revision of this V1beta1StatefulSetStatus.

Return type str

updated_replicas

Gets the updated_replicas of this V1beta1StatefulSetStatus. updatedReplicas is the number of Pods created by the StatefulSet controller from the StatefulSet version indicated by updateRevision.

Returns The updated_replicas of this V1beta1StatefulSetStatus.

Return type int

kubernetes.client.models.v1beta1_storage_class module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass (allow_volume_expansion=None,
    api_version=None,
    kind=None,
    meta-
    data=None,
    mount_options=None,
    pa-
    ram-
    e-
    ters=None,
    pro-
    vi-
    sioner=None,
    re-
    claim_policy=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

allow_volume_expansion

Gets the allow_volume_expansion of this V1beta1StorageClass. AllowVolumeExpansion shows whether the storage class allow volume expand

Returns The allow_volume_expansion of this V1beta1StorageClass.

Return type bool

api_version

Gets the api_version of this V1beta1StorageClass. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1StorageClass.

Return type str

attribute_map = {'kind': 'kind', 'allow_volume_expansion': 'allowVolumeExpansion', 'mount_options': 'mountOptions'}

kind

Gets the kind of this V1beta1StorageClass. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1StorageClass.

Return type str

metadata

Gets the metadata of this V1beta1StorageClass. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1StorageClass.

Return type *V1ObjectMeta*

mount_options

Gets the mount_options of this V1beta1StorageClass. Dynamically provisioned PersistentVolumes of this storage class are created with these mountOptions, e.g. ["ro", "soft"]. Not validated - mount of the PVs will simply fail if one is invalid.

Returns The mount_options of this V1beta1StorageClass.

Return type list[str]

parameters

Gets the parameters of this V1beta1StorageClass. Parameters holds the parameters for the provisioner that should create volumes of this storage class.

Returns The parameters of this V1beta1StorageClass.

Return type dict(str, str)

provisioner

Gets the provisioner of this V1beta1StorageClass. Provisioner indicates the type of the provisioner.

Returns The provisioner of this V1beta1StorageClass.

Return type str

reclaim_policy

Gets the reclaim_policy of this V1beta1StorageClass. Dynamically provisioned PersistentVolumes of this storage class are created with this reclaimPolicy. Defaults to Delete.

Returns The reclaim_policy of this V1beta1StorageClass.

Return type str

swagger_types = {'kind': 'str', 'allow_volume_expansion': 'bool', 'mount_options': 'list[str]', 'parameters': 'dict(str,

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_storage_class_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList (*api_version=None, items=None, kind=None, meta-data=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1StorageClassList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1StorageClassList.

Return type str

attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}

items

Gets the items of this V1beta1StorageClassList. Items is the list of StorageClasses

Returns The items of this V1beta1StorageClassList.

Return type list[V1beta1StorageClass]

kind

Gets the kind of this V1beta1StorageClassList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1StorageClassList.

Return type str

metadata

Gets the metadata of this V1beta1StorageClassList. Standard list metadata More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V1beta1StorageClassList.

Return type V1ListMeta

swagger_types = {'items': 'list[V1beta1StorageClass]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview (api_version=
kind=N
meta-
data=
spec=
sta-
tus=No
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V1beta1SubjectAccessReview. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V1beta1SubjectAccessReview.

Return type str

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the kind of this V1beta1SubjectAccessReview. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V1beta1SubjectAccessReview.

Return type str

metadata

Gets the metadata of this V1beta1SubjectAccessReview.

Returns The metadata of this V1beta1SubjectAccessReview.

Return type V1ObjectMeta

spec

Gets the spec of this V1beta1SubjectAccessReview. Spec holds information about the request being evaluated

Returns The spec of this V1beta1SubjectAccessReview.

Return type *V1beta1SubjectAccessReviewSpec*

status

Gets the status of this V1beta1SubjectAccessReview. Status is filled in by the server and indicates whether the request is allowed or not

Returns The status of this V1beta1SubjectAccessReview.

Return type *V1beta1SubjectAccessReviewStatus*

swagger_types = {'status': 'V1beta1SubjectAccessReviewStatus', 'kind': 'str', 'spec': 'V1beta1SubjectAccessReviewSpec'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_subject_access_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'group': 'group', 'uid': 'uid', 'extra': 'extra', 'non_resource_attributes': 'nonResourceAttributes', 'spec': 'spec', 'status': 'status'}

extra

Gets the extra of this V1beta1SubjectAccessReviewSpec. Extra corresponds to the user.Info.GetExtra() method from the authenticator. Since that is input to the authorizer it needs a reflection here.

Returns The extra of this V1beta1SubjectAccessReviewSpec.

Return type dict(str, list[str])

group

Gets the group of this V1beta1SubjectAccessReviewSpec. Groups is the groups you're testing for.

Returns The group of this V1beta1SubjectAccessReviewSpec.

Return type list[str]

non_resource_attributes

Gets the non_resource_attributes of this V1beta1SubjectAccessReviewSpec. NonResourceAttributes describes information for a non-resource access request

Returns The non_resource_attributes of this V1beta1SubjectAccessReviewSpec.

Return type *V1beta1NonResourceAttributes*

resource_attributes

Gets the resource_attributes of this V1beta1SubjectAccessReviewSpec. ResourceAuthorizationAttributes describes information for a resource access request

Returns The resource_attributes of this V1beta1SubjectAccessReviewSpec.

Return type *V1beta1ResourceAttributes*

swagger_types = {'group': 'list[str]', 'uid': 'str', 'extra': 'dict(str, list[str])', 'non_resource_attributes': 'V1beta1Non

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

uid

Gets the uid of this V1beta1SubjectAccessReviewSpec. UID information about the requesting user.

Returns The uid of this V1beta1SubjectAccessReviewSpec.

Return type str

user

Gets the user of this V1beta1SubjectAccessReviewSpec. User is the user you're testing for. If you specify "User" but not "Group", then is it interpreted as "What if User were not a member of any groups"

Returns The user of this V1beta1SubjectAccessReviewSpec.

Return type str

kubernetes.client.models.v1beta1_subject_access_review_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v1beta1_subject_access_review_status.V1beta1SubjectAccessReviewStatus

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

allowed

Gets the allowed of this V1beta1SubjectAccessReviewStatus. Allowed is required. True if the action would be allowed, false otherwise.

Returns The allowed of this V1beta1SubjectAccessReviewStatus.

Return type bool

attribute_map = {'reason': 'reason', 'evaluation_error': 'evaluationError', 'allowed': 'allowed'}

evaluation_error

Gets the evaluation_error of this V1beta1SubjectAccessReviewStatus. EvaluationError is an indication that some error occurred during the authorization check. It is entirely possible to get an error and be able to continue determine authorization status in spite of it. For instance, RBAC can be missing a role, but enough roles are still present and bound to reason about the request.

Returns The evaluation_error of this V1beta1SubjectAccessReviewStatus.

Return type str

reason

Gets the reason of this V1beta1SubjectAccessReviewStatus. Reason is optional. It indicates why a request was allowed or denied.

Returns The reason of this V1beta1SubjectAccessReviewStatus.

Return type str

swagger_types = {'reason': 'str', 'evaluation_error': 'str', 'allowed': 'bool'}

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v1beta1_subresource_reference module

kubernetes.client.models.v1beta1_third_party_resource module

kubernetes.client.models.v1beta1_third_party_resource_list module

kubernetes.client.models.v1beta1_token_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview (api_version=None,
                                                                    kind=None,
                                                                    meta-
                                                                    data=None,
                                                                    spec=None,
                                                                    sta-
                                                                    tus=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the `api_version` of this `V1beta1TokenReview`. `APIVersion` defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The `api_version` of this `V1beta1TokenReview`.

Return type `str`

attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}

kind

Gets the `kind` of this `V1beta1TokenReview`. `Kind` is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The `kind` of this `V1beta1TokenReview`.

Return type `str`

metadata

Gets the `metadata` of this `V1beta1TokenReview`.

Returns The `metadata` of this `V1beta1TokenReview`.

Return type *`V1ObjectMeta`*

spec

Gets the `spec` of this `V1beta1TokenReview`. `Spec` holds information about the request being evaluated

Returns The `spec` of this `V1beta1TokenReview`.

Return type *`V1beta1TokenReviewSpec`*

status

Gets the `status` of this `V1beta1TokenReview`. `Status` is filled in by the server and indicates whether the request can be authenticated.

Returns The `status` of this `V1beta1TokenReview`.

Return type *`V1beta1TokenReviewStatus`*

swagger_types = {'status': 'V1beta1TokenReviewStatus', 'kind': 'str', 'spec': 'V1beta1TokenReviewSpec', 'api_version': 'str'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v1beta1_token_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_token_review_spec.V1beta1TokenReviewSpec (token=None)
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'token': 'token'}
```

```
swagger_types = {'token': 'str'}
```

```
to_dict ()
```

Returns the model properties as a dict

```
to_str ()
```

Returns the string representation of the model

```
token
```

Gets the token of this V1beta1TokenReviewSpec. Token is the opaque bearer token.

Returns The token of this V1beta1TokenReviewSpec.

Return type str

kubernetes.client.models.v1beta1_token_review_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus (authenticated=er-
                                                                                      ror=None,
                                                                                      user=None)
```

```
    Bases: object
```

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'authenticated': 'authenticated', 'user': 'user', 'error': 'error'}
```

```
authenticated
```

Gets the authenticated of this V1beta1TokenReviewStatus. Authenticated indicates that the token was associated with a known user.

Returns The authenticated of this V1beta1TokenReviewStatus.

Return type bool

```
error
```

Gets the error of this V1beta1TokenReviewStatus. Error indicates that the token couldn't be checked

Returns The error of this V1beta1TokenReviewStatus.

Return type str

```
swagger_types = {'authenticated': 'bool', 'user': 'V1beta1UserInfo', 'error': 'str'}
```

```
to_dict ()
```

Returns the model properties as a dict

```
to_str ()
```

Returns the string representation of the model

user

Gets the user of this V1beta1TokenReviewStatus. User is the UserInfo associated with the provided token.

Returns The user of this V1beta1TokenReviewStatus.

Return type *V1beta1UserInfo*

kubernetes.client.models.v1beta1_user_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v1beta1_user_info.V1beta1UserInfo (extra=None,
                                                                    groups=None,
                                                                    uid=None, user-
                                                                    name=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'username': 'username', 'uid': 'uid', 'groups': 'groups', 'extra': 'extra'}

extra

Gets the extra of this V1beta1UserInfo. Any additional information provided by the authenticator.

Returns The extra of this V1beta1UserInfo.

Return type dict(str, list[str])

groups

Gets the groups of this V1beta1UserInfo. The names of groups this user is a part of.

Returns The groups of this V1beta1UserInfo.

Return type list[str]

swagger_types = {'username': 'str', 'uid': 'str', 'groups': 'list[str]', 'extra': 'dict(str, list[str])'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

uid

Gets the uid of this V1beta1UserInfo. A unique value that identifies this user across time. If this user is deleted and another user by the same name is added, they will have different UIDs.

Returns The uid of this V1beta1UserInfo.

Return type str

username

Gets the username of this V1beta1UserInfo. The name that uniquely identifies this user among all active users.

Returns The username of this V1beta1UserInfo.

Return type str

kubernetes.client.models.v2alpha1_cron_job module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob(api_version=None,  
                                                                kind=None,  
                                                                metadata=None,  
                                                                spec=None,  
                                                                status=None)
```

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V2alpha1CronJob. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V2alpha1CronJob.

Return type `str`

```
attribute_map = {'status': 'status', 'kind': 'kind', 'spec': 'spec', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

kind

Gets the kind of this V2alpha1CronJob. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V2alpha1CronJob.

Return type `str`

metadata

Gets the metadata of this V2alpha1CronJob. Standard object's metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V2alpha1CronJob.

Return type `V1ObjectMeta`

spec

Gets the spec of this V2alpha1CronJob. Specification of the desired behavior of a cron job, including the schedule. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V2alpha1CronJob.

Return type `V2alpha1CronJobSpec`

status

Gets the status of this V2alpha1CronJob. Current status of a cron job. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The status of this V2alpha1CronJob.

Return type `V2alpha1CronJobStatus`


```
swagger_types = {'status': 'V2alpha1CronJobStatus', 'kind': 'str', 'spec': 'V2alpha1CronJobSpec', 'api_version': 'str'}
to_dict()
    Returns the model properties as a dict
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v2alpha1_cron_job_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList (api_version=None,
                                                                              items=None,
                                                                              kind=None,
                                                                              meta-
                                                                              data=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

api_version

Gets the api_version of this V2alpha1CronJobList. APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#resources>

Returns The api_version of this V2alpha1CronJobList.

Return type str

```
attribute_map = {'items': 'items', 'kind': 'kind', 'api_version': 'apiVersion', 'metadata': 'metadata'}
```

items

Gets the items of this V2alpha1CronJobList. items is the list of CronJobs.

Returns The items of this V2alpha1CronJobList.

Return type list[V2alpha1CronJob]

kind

Gets the kind of this V2alpha1CronJobList. Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#types-kinds>

Returns The kind of this V2alpha1CronJobList.

Return type str

metadata

Gets the metadata of this V2alpha1CronJobList. Standard list metadata. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V2alpha1CronJobList.

Return type V1ListMeta

```
swagger_types = {'items': 'list[V2alpha1CronJob]', 'kind': 'str', 'api_version': 'str', 'metadata': 'V1ListMeta'}
to_dict()
    Returns the model properties as a dict
to_str()
    Returns the string representation of the model
```

kubernetes.client.models.v2alpha1_cron_job_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec (concurrency_policy=None,
                                                                              failed_jobs_history_limit=None,
                                                                              job_template=None,
                                                                              schedule=None,
                                                                              starting_deadline_seconds=None,
                                                                              successful_jobs_history_limit=None,
                                                                              suspend=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'suspend': 'suspend', 'job_template': 'jobTemplate', 'schedule': 'schedule', 'successful_jobs_history_limit': 'successfulJobsHistoryLimit'}
```

concurrency_policy

Gets the concurrency_policy of this V2alpha1CronJobSpec. Specifies how to treat concurrent executions of a Job. Defaults to Allow.

Returns The concurrency_policy of this V2alpha1CronJobSpec.

Return type str

failed_jobs_history_limit

Gets the failed_jobs_history_limit of this V2alpha1CronJobSpec. The number of failed finished jobs to retain. This is a pointer to distinguish between explicit zero and not specified.

Returns The failed_jobs_history_limit of this V2alpha1CronJobSpec.

Return type int

job_template

Gets the job_template of this V2alpha1CronJobSpec. Specifies the job that will be created when executing a CronJob.

Returns The job_template of this V2alpha1CronJobSpec.

Return type *V2alpha1JobTemplateSpec*

schedule

Gets the schedule of this V2alpha1CronJobSpec. The schedule in Cron format, see <https://en.wikipedia.org/wiki/Cron>.

Returns The schedule of this V2alpha1CronJobSpec.

Return type str

starting_deadline_seconds

Gets the starting_deadline_seconds of this V2alpha1CronJobSpec. Optional deadline in seconds for starting the job if it misses scheduled time for any reason. Missed jobs executions will be counted as failed ones.

Returns The starting_deadline_seconds of this V2alpha1CronJobSpec.

Return type int

successful_jobs_history_limit

Gets the successful_jobs_history_limit of this V2alpha1CronJobSpec. The number of successful finished jobs to retain. This is a pointer to distinguish between explicit zero and not specified.

Returns The successful_jobs_history_limit of this V2alpha1CronJobSpec.

Return type int

suspend

Gets the suspend of this V2alpha1CronJobSpec. This flag tells the controller to suspend subsequent executions, it does not apply to already started executions. Defaults to false.

Returns The suspend of this V2alpha1CronJobSpec.

Return type bool

swagger_types = {'suspend': 'bool', 'job_template': 'V2alpha1JobTemplateSpec', 'schedule': 'str', 'successful_jobs_h

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.v2alpha1_cron_job_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class kubernetes.client.models.v2alpha1_cron_job_status.V2alpha1CronJobStatus (*active=None, last_schedule_time=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

active

Gets the active of this V2alpha1CronJobStatus. A list of pointers to currently running jobs.

Returns The active of this V2alpha1CronJobStatus.

Return type list[V1ObjectReference]

attribute_map = {'active': 'active', 'last_schedule_time': 'lastScheduleTime'}

last_schedule_time

Gets the last_schedule_time of this V2alpha1CronJobStatus. Information when was the last time the job was successfully scheduled.

Returns The last_schedule_time of this V2alpha1CronJobStatus.

Return type datetime

swagger_types = {'active': 'list[V1ObjectReference]', 'last_schedule_time': 'datetime'}

to_dict()

Returns the model properties as a dict

to_str()

Returns the string representation of the model

kubernetes.client.models.v2alpha1_job module
kubernetes.client.models.v2alpha1_job_condition module
kubernetes.client.models.v2alpha1_job_list module
kubernetes.client.models.v2alpha1_job_spec module
kubernetes.client.models.v2alpha1_job_status module
kubernetes.client.models.v2alpha1_job_template_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.models.v2alpha1_job_template_spec.V2alpha1JobTemplateSpec` (*metadata=None, spec=None*)

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

attribute_map = {'spec': 'spec', 'metadata': 'metadata'}

metadata

Gets the metadata of this V2alpha1JobTemplateSpec. Standard object's metadata of the jobs created from this template. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#metadata>

Returns The metadata of this V2alpha1JobTemplateSpec.

Return type *V1ObjectMeta*

spec

Gets the spec of this V2alpha1JobTemplateSpec. Specification of the desired behavior of the job. More info: <https://git.k8s.io/community/contributors/devel/api-conventions.md#spec-and-status>

Returns The spec of this V2alpha1JobTemplateSpec.

Return type *V1JobSpec*

```
swagger_types = {'spec': 'V1JobSpec', 'metadata': 'V1ObjectMeta'}
```

```
to_dict()
```

Returns the model properties as a dict

```
to_str()
```

Returns the string representation of the model

kubernetes.client.models.version_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.models.version_info.VersionInfo (build_date=None,
                                                         compiler=None,
                                                         git_commit=None,
                                                         git_tree_state=None,
                                                         git_version=None,
                                                         go_version=None,      ma-
                                                         jor=None,      minor=None,
                                                         platform=None)
```

Bases: object

NOTE: This class is auto generated by the swagger code generator program. Do not edit the class manually.

```
attribute_map = {'build_date': 'buildDate', 'major': 'major', 'minor': 'minor', 'platform': 'platform', 'go_version':
```

```
build_date
```

Gets the build_date of this VersionInfo.

Returns The build_date of this VersionInfo.

Return type str

```
compiler
```

Gets the compiler of this VersionInfo.

Returns The compiler of this VersionInfo.

Return type str

```
git_commit
```

Gets the git_commit of this VersionInfo.

Returns The git_commit of this VersionInfo.

Return type str

```
git_tree_state
```

Gets the git_tree_state of this VersionInfo.

Returns The git_tree_state of this VersionInfo.

Return type str

```
git_version
```

Gets the git_version of this VersionInfo.

Returns The git_version of this VersionInfo.

Return type str

go_version

Gets the go_version of this VersionInfo.

Returns The go_version of this VersionInfo.

Return type str

major

Gets the major of this VersionInfo.

Returns The major of this VersionInfo.

Return type str

minor

Gets the minor of this VersionInfo.

Returns The minor of this VersionInfo.

Return type str

platform

Gets the platform of this VersionInfo.

Returns The platform of this VersionInfo.

Return type str

swagger_types = {'build_date': 'str', 'major': 'str', 'minor': 'str', 'platform': 'str', 'go_version': 'str', 'git_commit':

to_dict ()

Returns the model properties as a dict

to_str ()

Returns the string representation of the model

kubernetes.client.models.versioned_event module

Module contents

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

Submodules

kubernetes.client.api_client module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.client.api_client.ApiClient(configuration=None, header_name=None,
                                             header_value=None, cookie=None)
```

Bases: object

Generic API client for Swagger client library builds.

Swagger generic API client. This client handles the client- server communication, and is invariant across implementations. Specifics of the methods and models for each application are generated from the Swagger templates.

NOTE: This class is auto generated by the swagger code generator program. Ref: <https://github.com/swagger-api/swagger-codegen> Do not edit the class manually.

Parameters

- **host** – The base path for the server to call.
- **header_name** – a header to pass when making calls to the API.
- **header_value** – a header value to pass when making calls to the API.

```
NATIVE_TYPES_MAPPING = {'date': <type 'datetime.date'>, 'object': <type 'object'>, 'bool': <type 'bool'>, 'str': <type 'str'>
```

```
PRIMITIVE_TYPES = (<type 'float'>, <type 'bool'>, <type 'str'>, <type 'unicode'>, <type 'int'>, <type 'long'>)
```

```
call_api(resource_path, method, path_params=None, query_params=None, header_params=None,
         body=None, post_params=None, files=None, response_type=None, auth_settings=None,
         async=None, _return_http_data_only=None, collection_formats=None,
         _preload_content=True, _request_timeout=None)
```

Makes the HTTP request (synchronous) and return the deserialized data. To make an async request, set the async parameter.

Parameters

- **resource_path** – Path to method endpoint.
- **method** – Method to call.
- **path_params** – Path parameters in the url.
- **query_params** – Query parameters in the url.
- **header_params** – Header parameters to be placed in the request header.
- **body** – Request body.
- **dict** (*files*) – Request post form parameters, for *application/x-www-form-urlencoded*, *multipart/form-data*.
- **list** (*auth_settings*) – Auth Settings names for the request.
- **response** – Response data type.
- **dict** – key -> filename, value -> filepath, for *multipart/form-data*.
- **bool** (*async*) – execute request asynchronously
- **_return_http_data_only** – response data without head status code and headers
- **collection_formats** – dict of collection formats for path, query, header, and post parameters.
- **_preload_content** – if False, the urllib3.HTTPResponse object will be returned without reading/decoding response data. Default is True.
- **_request_timeout** – timeout setting for this request. If one number provided, it will be total request timeout. It can also be a pair (tuple) of (connection, read) timeouts.

Returns If `async` parameter is `True`, the request will be called asynchronously. The method will return the request thread. If parameter `async` is `False` or missing, then the method will return the response directly.

deserialize (*response, response_type*)

Deserializes response into an object.

Parameters

- **response** – `RESTResponse` object to be deserialized.
- **response_type** – class literal for deserialized object, or string of class name.

Returns deserialized object.

parameters_to_tuples (*params, collection_formats*)

Get parameters as list of tuples, formatting collections.

Parameters

- **params** – Parameters as dict or list of two-tuples
- **collection_formats** (*dict*) – Parameter collection formats

Returns Parameters as list of tuples, collections formatted

prepare_post_parameters (*post_params=None, files=None*)

Builds form parameters.

Parameters

- **post_params** – Normal form parameters.
- **files** – File parameters.

Returns Form parameters with files.

request (*method, url, query_params=None, headers=None, post_params=None, body=None, _preload_content=True, _request_timeout=None*)

Makes the HTTP request using `RESTClient`.

sanitize_for_serialization (*obj*)

Builds a JSON POST object.

If `obj` is `None`, return `None`. If `obj` is `str`, `int`, `long`, `float`, `bool`, return directly. If `obj` is `datetime.datetime`, `datetime.date`

convert to string in iso8601 format.

If `obj` is list, sanitize each element in the list. If `obj` is dict, return the dict. If `obj` is swagger model, return the properties dict.

Parameters `obj` – The data to serialize.

Returns The serialized form of data.

select_header_accept (*accepts*)

Returns `Accept` based on an array of accepts provided.

Parameters `accepts` – List of headers.

Returns `Accept` (e.g. `application/json`).

select_header_content_type (*content_types*)

Returns `Content-Type` based on an array of `content_types` provided.

Parameters `content_types` – List of content-types.

Returns Content-Type (e.g. application/json).

set_default_header (*header_name, header_value*)

update_params_for_auth (*headers, queries, auth_settings*)

Updates header and query params based on authentication setting.

Parameters

- **headers** – Header parameters dict to be updated.
- **queries** – Query parameters tuple list to be updated.
- **auth_settings** – Authentication setting identifiers list.

user_agent

Gets user agent.

kubernetes.client.configuration module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

class `kubernetes.client.configuration.Configuration`

Bases: `object`

NOTE: This class is auto generated by the swagger code generator program. Ref: <https://github.com/swagger-api/swagger-codegen> Do not edit the class manually.

auth_settings ()

Gets Auth Settings dict for api client.

Returns The Auth Settings information dict.

debug

Gets the debug status.

get_api_key_with_prefix (*identifier*)

Gets API key (with prefix if set).

Parameters **identifier** – The identifier of apiKey.

Returns The token for api key authentication.

get_basic_auth_token ()

Gets HTTP basic authentication header (string).

Returns The token for basic HTTP authentication.

logger_file

Gets the logger_file.

logger_format

Gets the logger_format.

to_debug_report ()

Gets the essential information for debugging.

Returns The report for debugging.

```
class kubernetes.client.configuration.TypeWithDefault (name, bases, dct)
    Bases: type
    set_default (default)
```

kubernetes.client.rest module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
exception kubernetes.client.rest.ApiException (status=None, reason=None,
                                              http_resp=None)
    Bases: exceptions.Exception
```

```
class kubernetes.client.rest.RESTClientObject (configuration, pools_size=4, maxsize=None)
    Bases: object
```

```
DELETE (url, headers=None, query_params=None, body=None, _preload_content=True, _request_timeout=None)
```

```
GET (url, headers=None, query_params=None, _preload_content=True, _request_timeout=None)
```

```
HEAD (url, headers=None, query_params=None, _preload_content=True, _request_timeout=None)
```

```
OPTIONS (url, headers=None, query_params=None, post_params=None, body=None,
         _preload_content=True, _request_timeout=None)
```

```
PATCH (url, headers=None, query_params=None, post_params=None, body=None,
        _preload_content=True, _request_timeout=None)
```

```
POST (url, headers=None, query_params=None, post_params=None, body=None,
       _preload_content=True, _request_timeout=None)
```

```
PUT (url, headers=None, query_params=None, post_params=None, body=None,
     _preload_content=True, _request_timeout=None)
```

```
request (method, url, query_params=None, headers=None, body=None, post_params=None,
         _preload_content=True, _request_timeout=None)
```

Parameters

- **method** – http request method
- **url** – http request url
- **query_params** – query parameters in the url
- **headers** – http request headers
- **body** – request json body, for *application/json*
- **post_params** – request post parameters, *application/x-www-form-urlencoded* and *multipart/form-data*
- **_preload_content** – if False, the urllib3.HTTPResponse object will be returned without reading/decoding response data. Default is True.
- **_request_timeout** – timeout setting for this request. If one number provided, it will be total request timeout. It can also be a pair (tuple) of (connection, read) timeouts.

```
class kubernetes.client.rest.RESTResponse(resp)
    Bases: io.IOBase

    getheader (name, default=None)
        Returns a given response header.

    getheaders ()
        Returns a dictionary of the response headers.
```

Module contents

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

kubernetes.config package

Submodules

kubernetes.config.config_exception module

```
exception kubernetes.config.config_exception.ConfigException
    Bases: exceptions.Exception
```

kubernetes.config.incluster_config module

```
class kubernetes.config.incluster_config.InClusterConfigLoader (token_filename,
                                                                cert_filename, env-
                                                                ron={ 'LANG':
                                                                'C.UTF-8',
                                                                'READTHE-
                                                                DOCS_PROJECT':
                                                                'kubernetes',
                                                                'READTHE-
                                                                DOCS':      'True',
                                                                'APPDIR':
                                                                '/app',      'DE-
                                                                BIAN_FRONTEND':
                                                                'noninteractive',
                                                                'OLDPWD':
                                                                '/home/docs',
                                                                'HOSTNAME':
                                                                'build-6207035-
                                                                project-77016-
                                                                kubernetes',
                                                                'PWD':
                                                                '/home/docs/checkouts/readthedocs.org/user
                                                                'BIN_PATH':
                                                                '/home/docs/checkouts/readthedocs.org/user
                                                                'READTHE-
                                                                DOCS_VERSION':
                                                                'latest',    'PATH':
                                                                '/home/docs/checkouts/readthedocs.org/user
                                                                'HOME':
                                                                '/home/docs'})
```

Bases: object

load_and_set()

kubernetes.config.incluster_config.load_incluster_config()

Use the service account kubernetes gives to pods to connect to kubernetes cluster. It's intended for clients that expect to be running inside a pod running on kubernetes. It will raise an exception if called from a process not running in a kubernetes environment.

kubernetes.config.incluster_config_test module

```
class kubernetes.config.incluster_config_test.InClusterConfigTest (methodName='runTest')
```

Bases: unittest.case.TestCase

```
get_test_loader (token_filename=None, cert_filename=None, env-
ron={ 'KUBERNETES_SERVICE_HOST':      '127.0.0.1',    'KUBER-
NETES_SERVICE_PORT': '80'})
```

setUp()

tearDown()

test_empty_cert_file()

test_empty_host()

```

test_empty_port()
test_empty_token_file()
test_join_host_port()
test_load_config()
test_no_cert_file()
test_no_host()
test_no_port()
test_no_token_file()

```

kubernetes.config.kube_config module

class `kubernetes.config.kube_config.ConfigNode` (*name, value*)

Bases: `object`

Remembers each config key's path and construct a relevant exception message in case of missing keys. The assumption is all access keys are present in a well-formed kube-config.

get_with_name (*name, safe=False*)

safe_get (*key*)

class `kubernetes.config.kube_config.FileOrData` (*obj, file_key_name, data_key_name=None, file_base_path='', base64_file_content=True*)

Bases: `object`

Utility class to read content of `obj[%data_key_name]` or file's content of `obj[%file_key_name]` and represent it as file or data. Note that the data is preferred. The `obj[%file_key_name]` will be used iff `obj['%data_key_name']` is not set or empty. Assumption is file content is raw data and data field is base64 string. The assumption can be changed with `base64_file_content` flag. If set to `False`, the content of the file will assumed to be base64 and read as is. The default `True` value will result in base64 encode of the file content after read.

as_data ()

If `obj[%data_key_name]` exists, Return `obj[%data_key_name]` otherwise base64 encoded string of `obj[%file_key_name]` file content.

as_file ()

If `obj[%data_key_name]` exists, return name of a file with base64 decoded `obj[%data_key_name]` content otherwise `obj[%file_key_name]`.

class `kubernetes.config.kube_config.KubeConfigLoader` (*config_dict, active_context=None, get_google_credentials=None, config_base_path='', config_persister=None*)

Bases: `object`

current_context

list_contexts ()

load_and_set (*client_configuration*)

set_active_context (*context_name=None*)

`kubernetes.config.kube_config.list_kube_config_contexts` (*config_file=None*)

```
kubernetes.config.kube_config.load_kube_config (config_file=None, context=None,
                                              client_configuration=None, per-
                                              sist_config=True)
```

Loads authentication and cluster information from kube-config file and stores them in kubernetes.client.configuration.

Parameters

- **config_file** – Name of the kube-config file.
- **context** – set the active context. If is set to None, current_context from config file will be used.
- **client_configuration** – The kubernetes.client.Configuration to set configs to.
- **persist_config** – If True, config file will be updated when changed (e.g GCP token refresh).

```
kubernetes.config.kube_config.new_client_from_config (config_file=None,
                                                    context=None, per-
                                                    sist_config=True)
```

Loads configuration the same as load_kube_config but returns an ApiClient to be used with any API object. This will allow the caller to concurrently talk with multiple clusters.

kubernetes.config.kube_config_test module

```
class kubernetes.config.kube_config_test.BaseTestCase (methodName='runTest')
```

Bases: unittest.case.TestCase

```
expect_exception (func, message_part, *args, **kwargs)
```

```
setUp ()
```

```
tearDown ()
```

```
class kubernetes.config.kube_config_test.FakeConfig (token=None, **kwargs)
```

```
FILE_KEYS = ['ssl_ca_cert', 'key_file', 'cert_file']
```

```
class kubernetes.config.kube_config_test.TestConfigNode (methodName='runTest')
```

Bases: *kubernetes.config.kube_config_test.BaseTestCase*

```
setUp ()
```

```
test_get_with_name ()
```

```
test_get_with_name_on_invalid_object ()
```

```
test_get_with_name_on_name_does_not_exists ()
```

```
test_get_with_name_on_non_list_object ()
```

```
test_key_does_not_exists ()
```

```
test_normal_map_array_operations ()
```

```
test_obj = {'key3': {'inner_key': 'inner_value'}, 'key2': ['a', 'b', 'c'], 'key1': 'test', 'with_names': [{'name': 'test_nam
```

```
class kubernetes.config.kube_config_test.TestFileOrData (methodName='runTest')
```

Bases: *kubernetes.config.kube_config_test.BaseTestCase*

```
static get_file_content (filename)
```

```
test_create_temp_file_with_content ()
```

```

test_data_given_data()
test_data_given_file()
test_data_given_file_and_data()
test_data_given_file_no_base64()
test_file_given_data()
test_file_given_data_no_base64()
test_file_given_file()
test_file_given_file_and_data()
test_file_given_non_existing_file()
test_file_with_custom_dirname()

class kubernetes.config.kube_config_test.TestKubeConfigLoader (methodName='runTest')
    Bases: kubernetes.config.kube_config_test.BaseTestCase
    TEST_KUBE_CONFIG = {'contexts': [{'name': 'no_user', 'context': {'cluster': 'default'}}, {'name': 'simple_token', 'context': {'cluster': 'default'}}]}

    test_current_context()
    test_gcp_no_refresh()
    test_list_contexts()
    test_list_kube_config_contexts()
    test_load_gcp_token_no_refresh()
    test_load_gcp_token_with_refresh()
    test_load_kube_config()
    test_load_user_pass_token()
    test_load_user_token()
    test_new_client_from_config()
    test_no_user_context()
    test_no_users_section()
    test_non_existing_user()
    test_set_active_context()
    test_simple_token()
    test_ssl()
    test_ssl_no_cert_files()
    test_ssl_no_verification()
    test_ssl_with_relative_ssl_files()
    test_user_pass()

```

Module contents

kubernetes.test package

Submodules

kubernetes.test.test_apis_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_apis_api.TestApisApi (methodName='runTest')
    Bases: unittest.case.TestCase
    ApisApi unit test stubs
    setUp ()
    tearDown ()
    test_get_api_versions ()
        Test case for get_api_versions
```

kubernetes.test.test_apps_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_apps_api.TestAppsApi (methodName='runTest')
    Bases: unittest.case.TestCase
    AppsApi unit test stubs
    setUp ()
    tearDown ()
    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_apps_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>


```

class kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    AppsV1beta1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_namespaced_controller_revision ()
        Test case for create_namespaced_controller_revision

    test_create_namespaced_deployment ()
        Test case for create_namespaced_deployment

    test_create_namespaced_deployment_rollback ()
        Test case for create_namespaced_deployment_rollback

    test_create_namespaced_stateful_set ()
        Test case for create_namespaced_stateful_set

    test_delete_collection_namespaced_controller_revision ()
        Test case for delete_collection_namespaced_controller_revision

    test_delete_collection_namespaced_deployment ()
        Test case for delete_collection_namespaced_deployment

    test_delete_collection_namespaced_stateful_set ()
        Test case for delete_collection_namespaced_stateful_set

    test_delete_namespaced_controller_revision ()
        Test case for delete_namespaced_controller_revision

    test_delete_namespaced_deployment ()
        Test case for delete_namespaced_deployment

    test_delete_namespaced_stateful_set ()
        Test case for delete_namespaced_stateful_set

    test_get_api_resources ()
        Test case for get_api_resources

    test_list_controller_revision_for_all_namespaces ()
        Test case for list_controller_revision_for_all_namespaces

    test_list_deployment_for_all_namespaces ()
        Test case for list_deployment_for_all_namespaces

    test_list_namespaced_controller_revision ()
        Test case for list_namespaced_controller_revision

    test_list_namespaced_deployment ()
        Test case for list_namespaced_deployment

    test_list_namespaced_stateful_set ()
        Test case for list_namespaced_stateful_set

    test_list_stateful_set_for_all_namespaces ()
        Test case for list_stateful_set_for_all_namespaces

    test_patch_namespaced_controller_revision ()
        Test case for patch_namespaced_controller_revision

    test_patch_namespaced_deployment ()
        Test case for patch_namespaced_deployment

```

```

test_patch_namespaced_deployment_scale ()
    Test case for patch_namespaced_deployment_scale

test_patch_namespaced_deployment_status ()
    Test case for patch_namespaced_deployment_status

test_patch_namespaced_stateful_set ()
    Test case for patch_namespaced_stateful_set

test_patch_namespaced_stateful_set_scale ()
    Test case for patch_namespaced_stateful_set_scale

test_patch_namespaced_stateful_set_status ()
    Test case for patch_namespaced_stateful_set_status

test_read_namespaced_controller_revision ()
    Test case for read_namespaced_controller_revision

test_read_namespaced_deployment ()
    Test case for read_namespaced_deployment

test_read_namespaced_deployment_scale ()
    Test case for read_namespaced_deployment_scale

test_read_namespaced_deployment_status ()
    Test case for read_namespaced_deployment_status

test_read_namespaced_stateful_set ()
    Test case for read_namespaced_stateful_set

test_read_namespaced_stateful_set_scale ()
    Test case for read_namespaced_stateful_set_scale

test_read_namespaced_stateful_set_status ()
    Test case for read_namespaced_stateful_set_status

test_replace_namespaced_controller_revision ()
    Test case for replace_namespaced_controller_revision

test_replace_namespaced_deployment ()
    Test case for replace_namespaced_deployment

test_replace_namespaced_deployment_scale ()
    Test case for replace_namespaced_deployment_scale

test_replace_namespaced_deployment_status ()
    Test case for replace_namespaced_deployment_status

test_replace_namespaced_stateful_set ()
    Test case for replace_namespaced_stateful_set

test_replace_namespaced_stateful_set_scale ()
    Test case for replace_namespaced_stateful_set_scale

test_replace_namespaced_stateful_set_status ()
    Test case for replace_namespaced_stateful_set_status

```

kubernetes.test.test_authentication_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_authentication_api.TestAuthenticationApi (methodName='runTest')
    Bases: unittest.case.TestCase

    AuthenticationApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_authentication_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_authentication_v1beta1_api.TestAuthenticationV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    AuthenticationV1beta1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_token_review ()
        Test case for create_token_review

    test_get_api_resources ()
        Test case for get_api_resources
```

kubernetes.test.test_authorization_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_authorization_api.TestAuthorizationApi (methodName='runTest')
    Bases: unittest.case.TestCase

    AuthorizationApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_authorization_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase
        AuthorizationV1beta1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_namespaced_local_subject_access_review ()
        Test case for create_namespaced_local_subject_access_review

    test_create_self_subject_access_review ()
        Test case for create_self_subject_access_review

    test_create_self_subject_rules_review ()
        Test case for create_self_subject_rules_review

    test_create_subject_access_review ()
        Test case for create_subject_access_review

    test_get_api_resources ()
        Test case for get_api_resources
```

kubernetes.test.test_autoscaling_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_autoscaling_api.TestAutoscalingApi (methodName='runTest')
    Bases: unittest.case.TestCase
        AutoscalingApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_autoscaling_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    AutoscalingV1Api unit test stubs

    setUp()

    tearDown()

    test_create_namespaced_horizontal_pod_autoscaler()
        Test case for create_namespaced_horizontal_pod_autoscaler

    test_delete_collection_namespaced_horizontal_pod_autoscaler()
        Test case for delete_collection_namespaced_horizontal_pod_autoscaler

    test_delete_namespaced_horizontal_pod_autoscaler()
        Test case for delete_namespaced_horizontal_pod_autoscaler

    test_get_api_resources()
        Test case for get_api_resources

    test_list_horizontal_pod_autoscaler_for_all_namespaces()
        Test case for list_horizontal_pod_autoscaler_for_all_namespaces

    test_list_namespaced_horizontal_pod_autoscaler()
        Test case for list_namespaced_horizontal_pod_autoscaler

    test_patch_namespaced_horizontal_pod_autoscaler()
        Test case for patch_namespaced_horizontal_pod_autoscaler

    test_patch_namespaced_horizontal_pod_autoscaler_status()
        Test case for patch_namespaced_horizontal_pod_autoscaler_status

    test_read_namespaced_horizontal_pod_autoscaler()
        Test case for read_namespaced_horizontal_pod_autoscaler

    test_read_namespaced_horizontal_pod_autoscaler_status()
        Test case for read_namespaced_horizontal_pod_autoscaler_status

    test_replace_namespaced_horizontal_pod_autoscaler()
        Test case for replace_namespaced_horizontal_pod_autoscaler

    test_replace_namespaced_horizontal_pod_autoscaler_status()
        Test case for replace_namespaced_horizontal_pod_autoscaler_status
```

kubernetes.test.test_batch_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_batch_api.TestBatchApi (methodName='runTest')
    Bases: unittest.case.TestCase

    BatchApi unit test stubs

    setUp()
```

```
tearDown()  
  
test_get_api_group()  
    Test case for get_api_group
```

kubernetes.test.test_batch_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_batch_v1_api.TestBatchV1Api (methodName='runTest')  
    Bases: unittest.case.TestCase  
  
    BatchV1Api unit test stubs  
  
    setUp()  
  
    tearDown()  
  
    test_create_namespaced_job()  
        Test case for create_namespaced_job  
  
    test_delete_collection_namespaced_job()  
        Test case for delete_collection_namespaced_job  
  
    test_delete_namespaced_job()  
        Test case for delete_namespaced_job  
  
    test_get_api_resources()  
        Test case for get_api_resources  
  
    test_list_job_for_all_namespaces()  
        Test case for list_job_for_all_namespaces  
  
    test_list_namespaced_job()  
        Test case for list_namespaced_job  
  
    test_patch_namespaced_job()  
        Test case for patch_namespaced_job  
  
    test_patch_namespaced_job_status()  
        Test case for patch_namespaced_job_status  
  
    test_read_namespaced_job()  
        Test case for read_namespaced_job  
  
    test_read_namespaced_job_status()  
        Test case for read_namespaced_job_status  
  
    test_replace_namespaced_job()  
        Test case for replace_namespaced_job  
  
    test_replace_namespaced_job_status()  
        Test case for replace_namespaced_job_status
```

kubernetes.test.test_batch_v2alpha1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api (methodName='runTest')
    Bases: unittest.case.TestCase
    BatchV2alpha1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_namespaced_cron_job ()
        Test case for create_namespaced_cron_job

    test_delete_collection_namespaced_cron_job ()
        Test case for delete_collection_namespaced_cron_job

    test_delete_namespaced_cron_job ()
        Test case for delete_namespaced_cron_job

    test_get_api_resources ()
        Test case for get_api_resources

    test_list_cron_job_for_all_namespaces ()
        Test case for list_cron_job_for_all_namespaces

    test_list_namespaced_cron_job ()
        Test case for list_namespaced_cron_job

    test_patch_namespaced_cron_job ()
        Test case for patch_namespaced_cron_job

    test_patch_namespaced_cron_job_status ()
        Test case for patch_namespaced_cron_job_status

    test_read_namespaced_cron_job ()
        Test case for read_namespaced_cron_job

    test_read_namespaced_cron_job_status ()
        Test case for read_namespaced_cron_job_status

    test_replace_namespaced_cron_job ()
        Test case for replace_namespaced_cron_job

    test_replace_namespaced_cron_job_status ()
        Test case for replace_namespaced_cron_job_status

```

kubernetes.test.test_certificates_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_certificates_api.TestCertificatesApi (methodName='runTest')
    Bases: unittest.case.TestCase
    CertificatesApi unit test stubs
    setUp ()
    tearDown ()
    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_certificates_v1alpha1_api module

kubernetes.test.test_core_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_core_api.TestCoreApi (methodName='runTest')
    Bases: unittest.case.TestCase
    CoreApi unit test stubs
    setUp ()
    tearDown ()
    test_get_api_versions ()
        Test case for get_api_versions
```

kubernetes.test.test_core_v1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_core_v1_api.TestCoreV1Api (methodName='runTest')
    Bases: unittest.case.TestCase
    CoreV1Api unit test stubs
    setUp ()
    tearDown ()
    test_connect_delete_namespaced_pod_proxy ()
        Test case for connect_delete_namespaced_pod_proxy
    test_connect_delete_namespaced_pod_proxy_with_path ()
        Test case for connect_delete_namespaced_pod_proxy_with_path
    test_connect_delete_namespaced_service_proxy ()
        Test case for connect_delete_namespaced_service_proxy
```



```

test_connect_delete_namespaced_service_proxy_with_path()
    Test case for connect_delete_namespaced_service_proxy_with_path

test_connect_delete_node_proxy()
    Test case for connect_delete_node_proxy

test_connect_delete_node_proxy_with_path()
    Test case for connect_delete_node_proxy_with_path

test_connect_get_namespaced_pod_attach()
    Test case for connect_get_namespaced_pod_attach

test_connect_get_namespaced_pod_exec()
    Test case for connect_get_namespaced_pod_exec

test_connect_get_namespaced_pod_portforward()
    Test case for connect_get_namespaced_pod_portforward

test_connect_get_namespaced_pod_proxy()
    Test case for connect_get_namespaced_pod_proxy

test_connect_get_namespaced_pod_proxy_with_path()
    Test case for connect_get_namespaced_pod_proxy_with_path

test_connect_get_namespaced_service_proxy()
    Test case for connect_get_namespaced_service_proxy

test_connect_get_namespaced_service_proxy_with_path()
    Test case for connect_get_namespaced_service_proxy_with_path

test_connect_get_node_proxy()
    Test case for connect_get_node_proxy

test_connect_get_node_proxy_with_path()
    Test case for connect_get_node_proxy_with_path

test_connect_head_namespaced_pod_proxy()
    Test case for connect_head_namespaced_pod_proxy

test_connect_head_namespaced_pod_proxy_with_path()
    Test case for connect_head_namespaced_pod_proxy_with_path

test_connect_head_namespaced_service_proxy()
    Test case for connect_head_namespaced_service_proxy

test_connect_head_namespaced_service_proxy_with_path()
    Test case for connect_head_namespaced_service_proxy_with_path

test_connect_head_node_proxy()
    Test case for connect_head_node_proxy

test_connect_head_node_proxy_with_path()
    Test case for connect_head_node_proxy_with_path

test_connect_options_namespaced_pod_proxy()
    Test case for connect_options_namespaced_pod_proxy

test_connect_options_namespaced_pod_proxy_with_path()
    Test case for connect_options_namespaced_pod_proxy_with_path

test_connect_options_namespaced_service_proxy()
    Test case for connect_options_namespaced_service_proxy

```

```
test_connect_options_namespaced_service_proxy_with_path()
    Test case for connect_options_namespaced_service_proxy_with_path

test_connect_options_node_proxy()
    Test case for connect_options_node_proxy

test_connect_options_node_proxy_with_path()
    Test case for connect_options_node_proxy_with_path

test_connect_patch_namespaced_pod_proxy()
    Test case for connect_patch_namespaced_pod_proxy

test_connect_patch_namespaced_pod_proxy_with_path()
    Test case for connect_patch_namespaced_pod_proxy_with_path

test_connect_patch_namespaced_service_proxy()
    Test case for connect_patch_namespaced_service_proxy

test_connect_patch_namespaced_service_proxy_with_path()
    Test case for connect_patch_namespaced_service_proxy_with_path

test_connect_patch_node_proxy()
    Test case for connect_patch_node_proxy

test_connect_patch_node_proxy_with_path()
    Test case for connect_patch_node_proxy_with_path

test_connect_post_namespaced_pod_attach()
    Test case for connect_post_namespaced_pod_attach

test_connect_post_namespaced_pod_exec()
    Test case for connect_post_namespaced_pod_exec

test_connect_post_namespaced_pod_portforward()
    Test case for connect_post_namespaced_pod_portforward

test_connect_post_namespaced_pod_proxy()
    Test case for connect_post_namespaced_pod_proxy

test_connect_post_namespaced_pod_proxy_with_path()
    Test case for connect_post_namespaced_pod_proxy_with_path

test_connect_post_namespaced_service_proxy()
    Test case for connect_post_namespaced_service_proxy

test_connect_post_namespaced_service_proxy_with_path()
    Test case for connect_post_namespaced_service_proxy_with_path

test_connect_post_node_proxy()
    Test case for connect_post_node_proxy

test_connect_post_node_proxy_with_path()
    Test case for connect_post_node_proxy_with_path

test_connect_put_namespaced_pod_proxy()
    Test case for connect_put_namespaced_pod_proxy

test_connect_put_namespaced_pod_proxy_with_path()
    Test case for connect_put_namespaced_pod_proxy_with_path

test_connect_put_namespaced_service_proxy()
    Test case for connect_put_namespaced_service_proxy
```

```
test_connect_put_namespaced_service_proxy_with_path()
    Test case for connect_put_namespaced_service_proxy_with_path

test_connect_put_node_proxy()
    Test case for connect_put_node_proxy

test_connect_put_node_proxy_with_path()
    Test case for connect_put_node_proxy_with_path

test_create_namespace()
    Test case for create_namespace

test_create_namespaced_binding()
    Test case for create_namespaced_binding

test_create_namespaced_config_map()
    Test case for create_namespaced_config_map

test_create_namespaced_endpoints()
    Test case for create_namespaced_endpoints

test_create_namespaced_event()
    Test case for create_namespaced_event

test_create_namespaced_limit_range()
    Test case for create_namespaced_limit_range

test_create_namespaced_persistent_volume_claim()
    Test case for create_namespaced_persistent_volume_claim

test_create_namespaced_pod()
    Test case for create_namespaced_pod

test_create_namespaced_pod_binding()
    Test case for create_namespaced_pod_binding

test_create_namespaced_pod_eviction()
    Test case for create_namespaced_pod_eviction

test_create_namespaced_pod_template()
    Test case for create_namespaced_pod_template

test_create_namespaced_replication_controller()
    Test case for create_namespaced_replication_controller

test_create_namespaced_resource_quota()
    Test case for create_namespaced_resource_quota

test_create_namespaced_secret()
    Test case for create_namespaced_secret

test_create_namespaced_service()
    Test case for create_namespaced_service

test_create_namespaced_service_account()
    Test case for create_namespaced_service_account

test_create_node()
    Test case for create_node

test_create_persistent_volume()
    Test case for create_persistent_volume
```

```
test_delete_collection_namespaced_config_map ()
    Test case for delete_collection_namespaced_config_map

test_delete_collection_namespaced_endpoints ()
    Test case for delete_collection_namespaced_endpoints

test_delete_collection_namespaced_event ()
    Test case for delete_collection_namespaced_event

test_delete_collection_namespaced_limit_range ()
    Test case for delete_collection_namespaced_limit_range

test_delete_collection_namespaced_persistent_volume_claim ()
    Test case for delete_collection_namespaced_persistent_volume_claim

test_delete_collection_namespaced_pod ()
    Test case for delete_collection_namespaced_pod

test_delete_collection_namespaced_pod_template ()
    Test case for delete_collection_namespaced_pod_template

test_delete_collection_namespaced_replication_controller ()
    Test case for delete_collection_namespaced_replication_controller

test_delete_collection_namespaced_resource_quota ()
    Test case for delete_collection_namespaced_resource_quota

test_delete_collection_namespaced_secret ()
    Test case for delete_collection_namespaced_secret

test_delete_collection_namespaced_service_account ()
    Test case for delete_collection_namespaced_service_account

test_delete_collection_node ()
    Test case for delete_collection_node

test_delete_collection_persistent_volume ()
    Test case for delete_collection_persistent_volume

test_delete_namespace ()
    Test case for delete_namespace

test_delete_namespaced_config_map ()
    Test case for delete_namespaced_config_map

test_delete_namespaced_endpoints ()
    Test case for delete_namespaced_endpoints

test_delete_namespaced_event ()
    Test case for delete_namespaced_event

test_delete_namespaced_limit_range ()
    Test case for delete_namespaced_limit_range

test_delete_namespaced_persistent_volume_claim ()
    Test case for delete_namespaced_persistent_volume_claim

test_delete_namespaced_pod ()
    Test case for delete_namespaced_pod

test_delete_namespaced_pod_template ()
    Test case for delete_namespaced_pod_template
```

```
test_delete_namespaced_replication_controller ()
    Test case for delete_namespaced_replication_controller

test_delete_namespaced_resource_quota ()
    Test case for delete_namespaced_resource_quota

test_delete_namespaced_secret ()
    Test case for delete_namespaced_secret

test_delete_namespaced_service ()
    Test case for delete_namespaced_service

test_delete_namespaced_service_account ()
    Test case for delete_namespaced_service_account

test_delete_node ()
    Test case for delete_node

test_delete_persistent_volume ()
    Test case for delete_persistent_volume

test_get_api_resources ()
    Test case for get_api_resources

test_list_component_status ()
    Test case for list_component_status

test_list_config_map_for_all_namespaces ()
    Test case for list_config_map_for_all_namespaces

test_list_endpoints_for_all_namespaces ()
    Test case for list_endpoints_for_all_namespaces

test_list_event_for_all_namespaces ()
    Test case for list_event_for_all_namespaces

test_list_limit_range_for_all_namespaces ()
    Test case for list_limit_range_for_all_namespaces

test_list_namespace ()
    Test case for list_namespace

test_list_namespaced_config_map ()
    Test case for list_namespaced_config_map

test_list_namespaced_endpoints ()
    Test case for list_namespaced_endpoints

test_list_namespaced_event ()
    Test case for list_namespaced_event

test_list_namespaced_limit_range ()
    Test case for list_namespaced_limit_range

test_list_namespaced_persistent_volume_claim ()
    Test case for list_namespaced_persistent_volume_claim

test_list_namespaced_pod ()
    Test case for list_namespaced_pod

test_list_namespaced_pod_template ()
    Test case for list_namespaced_pod_template
```

```
test_list_namespaced_replication_controller ()
    Test case for list_namespaced_replication_controller

test_list_namespaced_resource_quota ()
    Test case for list_namespaced_resource_quota

test_list_namespaced_secret ()
    Test case for list_namespaced_secret

test_list_namespaced_service ()
    Test case for list_namespaced_service

test_list_namespaced_service_account ()
    Test case for list_namespaced_service_account

test_list_node ()
    Test case for list_node

test_list_persistent_volume ()
    Test case for list_persistent_volume

test_list_persistent_volume_claim_for_all_namespaces ()
    Test case for list_persistent_volume_claim_for_all_namespaces

test_list_pod_for_all_namespaces ()
    Test case for list_pod_for_all_namespaces

test_list_pod_template_for_all_namespaces ()
    Test case for list_pod_template_for_all_namespaces

test_list_replication_controller_for_all_namespaces ()
    Test case for list_replication_controller_for_all_namespaces

test_list_resource_quota_for_all_namespaces ()
    Test case for list_resource_quota_for_all_namespaces

test_list_secret_for_all_namespaces ()
    Test case for list_secret_for_all_namespaces

test_list_service_account_for_all_namespaces ()
    Test case for list_service_account_for_all_namespaces

test_list_service_for_all_namespaces ()
    Test case for list_service_for_all_namespaces

test_patch_namespace ()
    Test case for patch_namespace

test_patch_namespace_status ()
    Test case for patch_namespace_status

test_patch_namespaced_config_map ()
    Test case for patch_namespaced_config_map

test_patch_namespaced_endpoints ()
    Test case for patch_namespaced_endpoints

test_patch_namespaced_event ()
    Test case for patch_namespaced_event

test_patch_namespaced_limit_range ()
    Test case for patch_namespaced_limit_range
```

```
test_patch_namespaced_persistent_volume_claim()  
    Test case for patch_namespaced_persistent_volume_claim  
  
test_patch_namespaced_persistent_volume_claim_status()  
    Test case for patch_namespaced_persistent_volume_claim_status  
  
test_patch_namespaced_pod()  
    Test case for patch_namespaced_pod  
  
test_patch_namespaced_pod_status()  
    Test case for patch_namespaced_pod_status  
  
test_patch_namespaced_pod_template()  
    Test case for patch_namespaced_pod_template  
  
test_patch_namespaced_replication_controller()  
    Test case for patch_namespaced_replication_controller  
  
test_patch_namespaced_replication_controller_scale()  
    Test case for patch_namespaced_replication_controller_scale  
  
test_patch_namespaced_replication_controller_status()  
    Test case for patch_namespaced_replication_controller_status  
  
test_patch_namespaced_resource_quota()  
    Test case for patch_namespaced_resource_quota  
  
test_patch_namespaced_resource_quota_status()  
    Test case for patch_namespaced_resource_quota_status  
  
test_patch_namespaced_secret()  
    Test case for patch_namespaced_secret  
  
test_patch_namespaced_service()  
    Test case for patch_namespaced_service  
  
test_patch_namespaced_service_account()  
    Test case for patch_namespaced_service_account  
  
test_patch_namespaced_service_status()  
    Test case for patch_namespaced_service_status  
  
test_patch_node()  
    Test case for patch_node  
  
test_patch_node_status()  
    Test case for patch_node_status  
  
test_patch_persistent_volume()  
    Test case for patch_persistent_volume  
  
test_patch_persistent_volume_status()  
    Test case for patch_persistent_volume_status  
  
test_proxy_delete_namespaced_pod()  
    Test case for proxy_delete_namespaced_pod  
  
test_proxy_delete_namespaced_pod_with_path()  
    Test case for proxy_delete_namespaced_pod_with_path  
  
test_proxy_delete_namespaced_service()  
    Test case for proxy_delete_namespaced_service
```

```
test_proxy_delete_namespaced_service_with_path()
    Test case for proxy_delete_namespaced_service_with_path

test_proxy_delete_node()
    Test case for proxy_delete_node

test_proxy_delete_node_with_path()
    Test case for proxy_delete_node_with_path

test_proxy_get_namespaced_pod()
    Test case for proxy_get_namespaced_pod

test_proxy_get_namespaced_pod_with_path()
    Test case for proxy_get_namespaced_pod_with_path

test_proxy_get_namespaced_service()
    Test case for proxy_get_namespaced_service

test_proxy_get_namespaced_service_with_path()
    Test case for proxy_get_namespaced_service_with_path

test_proxy_get_node()
    Test case for proxy_get_node

test_proxy_get_node_with_path()
    Test case for proxy_get_node_with_path

test_proxy_head_namespaced_pod()
    Test case for proxy_head_namespaced_pod

test_proxy_head_namespaced_pod_with_path()
    Test case for proxy_head_namespaced_pod_with_path

test_proxy_head_namespaced_service()
    Test case for proxy_head_namespaced_service

test_proxy_head_namespaced_service_with_path()
    Test case for proxy_head_namespaced_service_with_path

test_proxy_head_node()
    Test case for proxy_head_node

test_proxy_head_node_with_path()
    Test case for proxy_head_node_with_path

test_proxy_options_namespaced_pod()
    Test case for proxy_options_namespaced_pod

test_proxy_options_namespaced_pod_with_path()
    Test case for proxy_options_namespaced_pod_with_path

test_proxy_options_namespaced_service()
    Test case for proxy_options_namespaced_service

test_proxy_options_namespaced_service_with_path()
    Test case for proxy_options_namespaced_service_with_path

test_proxy_options_node()
    Test case for proxy_options_node

test_proxy_options_node_with_path()
    Test case for proxy_options_node_with_path
```



```
test_proxy_patch_namespaced_pod()
    Test case for proxy_patch_namespaced_pod

test_proxy_patch_namespaced_pod_with_path()
    Test case for proxy_patch_namespaced_pod_with_path

test_proxy_patch_namespaced_service()
    Test case for proxy_patch_namespaced_service

test_proxy_patch_namespaced_service_with_path()
    Test case for proxy_patch_namespaced_service_with_path

test_proxy_patch_node()
    Test case for proxy_patch_node

test_proxy_patch_node_with_path()
    Test case for proxy_patch_node_with_path

test_proxy_post_namespaced_pod()
    Test case for proxy_post_namespaced_pod

test_proxy_post_namespaced_pod_with_path()
    Test case for proxy_post_namespaced_pod_with_path

test_proxy_post_namespaced_service()
    Test case for proxy_post_namespaced_service

test_proxy_post_namespaced_service_with_path()
    Test case for proxy_post_namespaced_service_with_path

test_proxy_post_node()
    Test case for proxy_post_node

test_proxy_post_node_with_path()
    Test case for proxy_post_node_with_path

test_proxy_put_namespaced_pod()
    Test case for proxy_put_namespaced_pod

test_proxy_put_namespaced_pod_with_path()
    Test case for proxy_put_namespaced_pod_with_path

test_proxy_put_namespaced_service()
    Test case for proxy_put_namespaced_service

test_proxy_put_namespaced_service_with_path()
    Test case for proxy_put_namespaced_service_with_path

test_proxy_put_node()
    Test case for proxy_put_node

test_proxy_put_node_with_path()
    Test case for proxy_put_node_with_path

test_read_component_status()
    Test case for read_component_status

test_read_namespace()
    Test case for read_namespace

test_read_namespace_status()
    Test case for read_namespace_status
```

```
test_read_namespaced_config_map ()
    Test case for read_namespaced_config_map

test_read_namespaced_endpoints ()
    Test case for read_namespaced_endpoints

test_read_namespaced_event ()
    Test case for read_namespaced_event

test_read_namespaced_limit_range ()
    Test case for read_namespaced_limit_range

test_read_namespaced_persistent_volume_claim ()
    Test case for read_namespaced_persistent_volume_claim

test_read_namespaced_persistent_volume_claim_status ()
    Test case for read_namespaced_persistent_volume_claim_status

test_read_namespaced_pod ()
    Test case for read_namespaced_pod

test_read_namespaced_pod_log ()
    Test case for read_namespaced_pod_log

test_read_namespaced_pod_status ()
    Test case for read_namespaced_pod_status

test_read_namespaced_pod_template ()
    Test case for read_namespaced_pod_template

test_read_namespaced_replication_controller ()
    Test case for read_namespaced_replication_controller

test_read_namespaced_replication_controller_scale ()
    Test case for read_namespaced_replication_controller_scale

test_read_namespaced_replication_controller_status ()
    Test case for read_namespaced_replication_controller_status

test_read_namespaced_resource_quota ()
    Test case for read_namespaced_resource_quota

test_read_namespaced_resource_quota_status ()
    Test case for read_namespaced_resource_quota_status

test_read_namespaced_secret ()
    Test case for read_namespaced_secret

test_read_namespaced_service ()
    Test case for read_namespaced_service

test_read_namespaced_service_account ()
    Test case for read_namespaced_service_account

test_read_namespaced_service_status ()
    Test case for read_namespaced_service_status

test_read_node ()
    Test case for read_node

test_read_node_status ()
    Test case for read_node_status
```

```
test_read_persistent_volume()
    Test case for read_persistent_volume

test_read_persistent_volume_status()
    Test case for read_persistent_volume_status

test_replace_namespace()
    Test case for replace_namespace

test_replace_namespace_finalize()
    Test case for replace_namespace_finalize

test_replace_namespace_status()
    Test case for replace_namespace_status

test_replace_namespaced_config_map()
    Test case for replace_namespaced_config_map

test_replace_namespaced_endpoints()
    Test case for replace_namespaced_endpoints

test_replace_namespaced_event()
    Test case for replace_namespaced_event

test_replace_namespaced_limit_range()
    Test case for replace_namespaced_limit_range

test_replace_namespaced_persistent_volume_claim()
    Test case for replace_namespaced_persistent_volume_claim

test_replace_namespaced_persistent_volume_claim_status()
    Test case for replace_namespaced_persistent_volume_claim_status

test_replace_namespaced_pod()
    Test case for replace_namespaced_pod

test_replace_namespaced_pod_status()
    Test case for replace_namespaced_pod_status

test_replace_namespaced_pod_template()
    Test case for replace_namespaced_pod_template

test_replace_namespaced_replication_controller()
    Test case for replace_namespaced_replication_controller

test_replace_namespaced_replication_controller_scale()
    Test case for replace_namespaced_replication_controller_scale

test_replace_namespaced_replication_controller_status()
    Test case for replace_namespaced_replication_controller_status

test_replace_namespaced_resource_quota()
    Test case for replace_namespaced_resource_quota

test_replace_namespaced_resource_quota_status()
    Test case for replace_namespaced_resource_quota_status

test_replace_namespaced_secret()
    Test case for replace_namespaced_secret

test_replace_namespaced_service()
    Test case for replace_namespaced_service
```

```

test_replace_namespaced_service_account ()
    Test case for replace_namespaced_service_account

test_replace_namespaced_service_status ()
    Test case for replace_namespaced_service_status

test_replace_node ()
    Test case for replace_node

test_replace_node_status ()
    Test case for replace_node_status

test_replace_persistent_volume ()
    Test case for replace_persistent_volume

test_replace_persistent_volume_status ()
    Test case for replace_persistent_volume_status

```

kubernetes.test.test_extensions_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_extensions_api.TestExtensionsApi (methodName='runTest')
    Bases: unittest.case.TestCase

    ExtensionsApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group

```

kubernetes.test.test_extensions_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    ExtensionsV1beta1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_namespaced_daemon_set ()
        Test case for create_namespaced_daemon_set

```

```
test_create_namespaced_deployment ()
    Test case for create_namespaced_deployment

test_create_namespaced_deployment_rollback ()
    Test case for create_namespaced_deployment_rollback

test_create_namespaced_ingress ()
    Test case for create_namespaced_ingress

test_create_namespaced_network_policy ()
    Test case for create_namespaced_network_policy

test_create_namespaced_replica_set ()
    Test case for create_namespaced_replica_set

test_create_pod_security_policy ()
    Test case for create_pod_security_policy

test_delete_collection_namespaced_daemon_set ()
    Test case for delete_collection_namespaced_daemon_set

test_delete_collection_namespaced_deployment ()
    Test case for delete_collection_namespaced_deployment

test_delete_collection_namespaced_ingress ()
    Test case for delete_collection_namespaced_ingress

test_delete_collection_namespaced_network_policy ()
    Test case for delete_collection_namespaced_network_policy

test_delete_collection_namespaced_replica_set ()
    Test case for delete_collection_namespaced_replica_set

test_delete_collection_pod_security_policy ()
    Test case for delete_collection_pod_security_policy

test_delete_namespaced_daemon_set ()
    Test case for delete_namespaced_daemon_set

test_delete_namespaced_deployment ()
    Test case for delete_namespaced_deployment

test_delete_namespaced_ingress ()
    Test case for delete_namespaced_ingress

test_delete_namespaced_network_policy ()
    Test case for delete_namespaced_network_policy

test_delete_namespaced_replica_set ()
    Test case for delete_namespaced_replica_set

test_delete_pod_security_policy ()
    Test case for delete_pod_security_policy

test_get_api_resources ()
    Test case for get_api_resources

test_list_daemon_set_for_all_namespaces ()
    Test case for list_daemon_set_for_all_namespaces

test_list_deployment_for_all_namespaces ()
    Test case for list_deployment_for_all_namespaces
```

```
test_list_ingress_for_all_namespaces ()  
    Test case for list_ingress_for_all_namespaces  
  
test_list_namespaced_daemon_set ()  
    Test case for list_namespaced_daemon_set  
  
test_list_namespaced_deployment ()  
    Test case for list_namespaced_deployment  
  
test_list_namespaced_ingress ()  
    Test case for list_namespaced_ingress  
  
test_list_namespaced_network_policy ()  
    Test case for list_namespaced_network_policy  
  
test_list_namespaced_replica_set ()  
    Test case for list_namespaced_replica_set  
  
test_list_network_policy_for_all_namespaces ()  
    Test case for list_network_policy_for_all_namespaces  
  
test_list_pod_security_policy ()  
    Test case for list_pod_security_policy  
  
test_list_replica_set_for_all_namespaces ()  
    Test case for list_replica_set_for_all_namespaces  
  
test_patch_namespaced_daemon_set ()  
    Test case for patch_namespaced_daemon_set  
  
test_patch_namespaced_daemon_set_status ()  
    Test case for patch_namespaced_daemon_set_status  
  
test_patch_namespaced_deployment ()  
    Test case for patch_namespaced_deployment  
  
test_patch_namespaced_deployment_scale ()  
    Test case for patch_namespaced_deployment_scale  
  
test_patch_namespaced_deployment_status ()  
    Test case for patch_namespaced_deployment_status  
  
test_patch_namespaced_ingress ()  
    Test case for patch_namespaced_ingress  
  
test_patch_namespaced_ingress_status ()  
    Test case for patch_namespaced_ingress_status  
  
test_patch_namespaced_network_policy ()  
    Test case for patch_namespaced_network_policy  
  
test_patch_namespaced_replica_set ()  
    Test case for patch_namespaced_replica_set  
  
test_patch_namespaced_replica_set_scale ()  
    Test case for patch_namespaced_replica_set_scale  
  
test_patch_namespaced_replica_set_status ()  
    Test case for patch_namespaced_replica_set_status  
  
test_patch_namespaced_replication_controller_dummy_scale ()  
    Test case for patch_namespaced_replication_controller_dummy_scale
```

```
test_patch_pod_security_policy()  
    Test case for patch_pod_security_policy  
  
test_read_namespaced_daemon_set()  
    Test case for read_namespaced_daemon_set  
  
test_read_namespaced_daemon_set_status()  
    Test case for read_namespaced_daemon_set_status  
  
test_read_namespaced_deployment()  
    Test case for read_namespaced_deployment  
  
test_read_namespaced_deployment_scale()  
    Test case for read_namespaced_deployment_scale  
  
test_read_namespaced_deployment_status()  
    Test case for read_namespaced_deployment_status  
  
test_read_namespaced_ingress()  
    Test case for read_namespaced_ingress  
  
test_read_namespaced_ingress_status()  
    Test case for read_namespaced_ingress_status  
  
test_read_namespaced_network_policy()  
    Test case for read_namespaced_network_policy  
  
test_read_namespaced_replica_set()  
    Test case for read_namespaced_replica_set  
  
test_read_namespaced_replica_set_scale()  
    Test case for read_namespaced_replica_set_scale  
  
test_read_namespaced_replica_set_status()  
    Test case for read_namespaced_replica_set_status  
  
test_read_namespaced_replication_controller_dummy_scale()  
    Test case for read_namespaced_replication_controller_dummy_scale  
  
test_read_pod_security_policy()  
    Test case for read_pod_security_policy  
  
test_replace_namespaced_daemon_set()  
    Test case for replace_namespaced_daemon_set  
  
test_replace_namespaced_daemon_set_status()  
    Test case for replace_namespaced_daemon_set_status  
  
test_replace_namespaced_deployment()  
    Test case for replace_namespaced_deployment  
  
test_replace_namespaced_deployment_scale()  
    Test case for replace_namespaced_deployment_scale  
  
test_replace_namespaced_deployment_status()  
    Test case for replace_namespaced_deployment_status  
  
test_replace_namespaced_ingress()  
    Test case for replace_namespaced_ingress  
  
test_replace_namespaced_ingress_status()  
    Test case for replace_namespaced_ingress_status
```

```

test_replace_namespaced_network_policy()
    Test case for replace_namespaced_network_policy

test_replace_namespaced_replica_set()
    Test case for replace_namespaced_replica_set

test_replace_namespaced_replica_set_scale()
    Test case for replace_namespaced_replica_set_scale

test_replace_namespaced_replica_set_status()
    Test case for replace_namespaced_replica_set_status

test_replace_namespaced_replication_controller_dummy_scale()
    Test case for replace_namespaced_replication_controller_dummy_scale

test_replace_pod_security_policy()
    Test case for replace_pod_security_policy

```

kubernetes.test.test_intstr_int_or_string module

kubernetes.test.test_logs_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_logs_api.TestLogsApi (methodName='runTest')
    Bases: unittest.case.TestCase

    LogsApi unit test stubs

    setUp()

    tearDown()

    test_log_file_handler()
        Test case for log_file_handler

    test_log_file_list_handler()
        Test case for log_file_list_handler

```

kubernetes.test.test_policy_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_policy_api.TestPolicyApi (methodName='runTest')
    Bases: unittest.case.TestCase

    PolicyApi unit test stubs

    setUp()

```



```
tearDown ()

test_get_api_group ()
    Test case for get_api_group
```

kubernetes.test.test_policy_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    PolicyV1beta1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_namespaced_pod_disruption_budget ()
        Test case for create_namespaced_pod_disruption_budget

    test_delete_collection_namespaced_pod_disruption_budget ()
        Test case for delete_collection_namespaced_pod_disruption_budget

    test_delete_namespaced_pod_disruption_budget ()
        Test case for delete_namespaced_pod_disruption_budget

    test_get_api_resources ()
        Test case for get_api_resources

    test_list_namespaced_pod_disruption_budget ()
        Test case for list_namespaced_pod_disruption_budget

    test_list_pod_disruption_budget_for_all_namespaces ()
        Test case for list_pod_disruption_budget_for_all_namespaces

    test_patch_namespaced_pod_disruption_budget ()
        Test case for patch_namespaced_pod_disruption_budget

    test_patch_namespaced_pod_disruption_budget_status ()
        Test case for patch_namespaced_pod_disruption_budget_status

    test_read_namespaced_pod_disruption_budget ()
        Test case for read_namespaced_pod_disruption_budget

    test_read_namespaced_pod_disruption_budget_status ()
        Test case for read_namespaced_pod_disruption_budget_status

    test_replace_namespaced_pod_disruption_budget ()
        Test case for replace_namespaced_pod_disruption_budget

    test_replace_namespaced_pod_disruption_budget_status ()
        Test case for replace_namespaced_pod_disruption_budget_status
```

kubernetes.test.test_rbac_authorization_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_rbac_authorization_api.TestRbacAuthorizationApi (methodName='runTest')
    Bases: unittest.case.TestCase

    RbacAuthorizationApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_rbac_authorization_v1alpha1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    RbacAuthorizationV1alpha1Api unit test stubs

    setUp ()

    tearDown ()

    test_create_cluster_role ()
        Test case for create_cluster_role

    test_create_cluster_role_binding ()
        Test case for create_cluster_role_binding

    test_create_namespaced_role ()
        Test case for create_namespaced_role

    test_create_namespaced_role_binding ()
        Test case for create_namespaced_role_binding

    test_delete_cluster_role ()
        Test case for delete_cluster_role

    test_delete_cluster_role_binding ()
        Test case for delete_cluster_role_binding

    test_delete_collection_cluster_role ()
        Test case for delete_collection_cluster_role

    test_delete_collection_cluster_role_binding ()
        Test case for delete_collection_cluster_role_binding
```

```
test_delete_collection_namespaced_role ()
    Test case for delete_collection_namespaced_role

test_delete_collection_namespaced_role_binding ()
    Test case for delete_collection_namespaced_role_binding

test_delete_namespaced_role ()
    Test case for delete_namespaced_role

test_delete_namespaced_role_binding ()
    Test case for delete_namespaced_role_binding

test_get_api_resources ()
    Test case for get_api_resources

test_list_cluster_role ()
    Test case for list_cluster_role

test_list_cluster_role_binding ()
    Test case for list_cluster_role_binding

test_list_namespaced_role ()
    Test case for list_namespaced_role

test_list_namespaced_role_binding ()
    Test case for list_namespaced_role_binding

test_list_role_binding_for_all_namespaces ()
    Test case for list_role_binding_for_all_namespaces

test_list_role_for_all_namespaces ()
    Test case for list_role_for_all_namespaces

test_patch_cluster_role ()
    Test case for patch_cluster_role

test_patch_cluster_role_binding ()
    Test case for patch_cluster_role_binding

test_patch_namespaced_role ()
    Test case for patch_namespaced_role

test_patch_namespaced_role_binding ()
    Test case for patch_namespaced_role_binding

test_read_cluster_role ()
    Test case for read_cluster_role

test_read_cluster_role_binding ()
    Test case for read_cluster_role_binding

test_read_namespaced_role ()
    Test case for read_namespaced_role

test_read_namespaced_role_binding ()
    Test case for read_namespaced_role_binding

test_replace_cluster_role ()
    Test case for replace_cluster_role

test_replace_cluster_role_binding ()
    Test case for replace_cluster_role_binding
```

```
test_replace_namespaced_role ()
    Test case for replace_namespaced_role

test_replace_namespaced_role_binding ()
    Test case for replace_namespaced_role_binding
```

kubernetes.test.test_resource_quantity module

kubernetes.test.test_runtime_raw_extension module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_runtime_raw_extension.TestRuntimeRawExtension (methodName='runTest')
    Bases: unittest.case.TestCase

    RuntimeRawExtension unit test stubs

    setUp ()

    tearDown ()

    testRuntimeRawExtension ()
        Test RuntimeRawExtension
```

kubernetes.test.test_storage_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_storage_api.TestStorageApi (methodName='runTest')
    Bases: unittest.case.TestCase

    StorageApi unit test stubs

    setUp ()

    tearDown ()

    test_get_api_group ()
        Test case for get_api_group
```

kubernetes.test.test_storage_v1beta1_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api (methodName='runTest')
    Bases: unittest.case.TestCase

    StorageV1beta1Api unit test stubs

    setUp()

    tearDown()

    test_create_storage_class()
        Test case for create_storage_class

    test_delete_collection_storage_class()
        Test case for delete_collection_storage_class

    test_delete_storage_class()
        Test case for delete_storage_class

    test_get_api_resources()
        Test case for get_api_resources

    test_list_storage_class()
        Test case for list_storage_class

    test_patch_storage_class()
        Test case for patch_storage_class

    test_read_storage_class()
        Test case for read_storage_class

    test_replace_storage_class()
        Test case for replace_storage_class
```

`kubernetes.test.test_unversioned_api_group` module

`kubernetes.test.test_unversioned_api_group_list` module

`kubernetes.test.test_unversioned_api_resource` module

`kubernetes.test.test_unversioned_api_resource_list` module

`kubernetes.test.test_unversioned_api_versions` module

`kubernetes.test.test_unversioned_group_version_for_discovery` module

`kubernetes.test.test_unversioned_label_selector` module

`kubernetes.test.test_unversioned_label_selector_requirement` module

`kubernetes.test.test_unversioned_list_meta` module

`kubernetes.test.test_unversioned_server_address_by_client_cidr` module

`kubernetes.test.test_unversioned_status` module

`kubernetes.test.test_unversioned_status_cause` module

`kubernetes.test.test_unversioned_status_details` module

`kubernetes.test.test_unversioned_time` module

`kubernetes.test.test_v1_attached_volume` module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_attached_volume.TestV1AttachedVolume (methodName='runTest')
    Bases: unittest.case.TestCase
        V1AttachedVolume unit test stubs

    setUp ()

    tearDown ()

    testV1AttachedVolume ()
        Test V1AttachedVolume
```

kubernetes.test.test_v1_aws_elastic_block_store_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_aws_elastic_block_store_volume_source.TestV1AWSElasticBlockStoreVolumeSource
    Bases: unittest.case.TestCase
    V1AWSElasticBlockStoreVolumeSource unit test stubs

    setUp ()

    tearDown ()

    testV1AWSElasticBlockStoreVolumeSource ()
        Test V1AWSElasticBlockStoreVolumeSource

```

kubernetes.test.test_v1_azure_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_azure_disk_volume_source.TestV1AzureDiskVolumeSource (methodName=
    Bases: unittest.case.TestCase
    V1AzureDiskVolumeSource unit test stubs

    setUp ()

    tearDown ()

    testV1AzureDiskVolumeSource ()
        Test V1AzureDiskVolumeSource

```

kubernetes.test.test_v1_azure_file_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_azure_file_volume_source.TestV1AzureFileVolumeSource (methodName=
    Bases: unittest.case.TestCase
    V1AzureFileVolumeSource unit test stubs

    setUp ()

    tearDown ()

```

```
testV1AzureFileVolumeSource ()  
    Test V1AzureFileVolumeSource
```

kubernetes.test.test_v1_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_binding.TestV1Binding (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1Binding unit test stubs  
    setUp ()  
    tearDown ()  
    testV1Binding ()  
        Test V1Binding
```

kubernetes.test.test_v1_capabilities module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_capabilities.TestV1Capabilities (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1Capabilities unit test stubs  
    setUp ()  
    tearDown ()  
    testV1Capabilities ()  
        Test V1Capabilities
```

kubernetes.test.test_v1_ceph_fs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_ceph_fs_volume_source.TestV1CephFSVolumeSource (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1CephFSVolumeSource unit test stubs
```



```

setUp ()
tearDown ()
testV1CephFSVolumeSource ()
    Test V1CephFSVolumeSource

```

kubernetes.test.test_v1_cinder_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_cinder_volume_source.TestV1CinderVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1CinderVolumeSource unit test stubs

    setUp ()
    tearDown ()
    testV1CinderVolumeSource ()
        Test V1CinderVolumeSource

```

kubernetes.test.test_v1_component_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_component_condition.TestV1ComponentCondition (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ComponentCondition unit test stubs

    setUp ()
    tearDown ()
    testV1ComponentCondition ()
        Test V1ComponentCondition

```

kubernetes.test.test_v1_component_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_component_status.TestV1ComponentStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ComponentStatus unit test stubs
    setUp()
    tearDown()
    testV1ComponentStatus()
        Test V1ComponentStatus
```

kubernetes.test.test_v1_component_status_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_component_status_list.TestV1ComponentStatusList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ComponentStatusList unit test stubs
    setUp()
    tearDown()
    testV1ComponentStatusList()
        Test V1ComponentStatusList
```

kubernetes.test.test_v1_config_map module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_config_map.TestV1ConfigMap (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ConfigMap unit test stubs
    setUp()
    tearDown()
    testV1ConfigMap()
        Test V1ConfigMap
```

kubernetes.test.test_v1_config_map_key_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_config_map_key_selector.TestV1ConfigMapKeySelector (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ConfigMapKeySelector unit test stubs

    setUp ()

    tearDown ()

    testV1ConfigMapKeySelector ()
        Test V1ConfigMapKeySelector
```

kubernetes.test.test_v1_config_map_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_config_map_list.TestV1ConfigMapList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ConfigMapList unit test stubs

    setUp ()

    tearDown ()

    testV1ConfigMapList ()
        Test V1ConfigMapList
```

kubernetes.test.test_v1_config_map_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_config_map_volume_source.TestV1ConfigMapVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ConfigMapVolumeSource unit test stubs

    setUp ()

    tearDown ()

    testV1ConfigMapVolumeSource ()
        Test V1ConfigMapVolumeSource
```

kubernetes.test.test_v1_container module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container.TestV1Container (methodName='runTest')
    Bases: unittest.case.TestCase
        V1Container unit test stubs

    setUp ()

    tearDown ()

    testV1Container ()
        Test V1Container
```

kubernetes.test.test_v1_container_image module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container_image.TestV1ContainerImage (methodName='runTest')
    Bases: unittest.case.TestCase
        V1ContainerImage unit test stubs

    setUp ()

    tearDown ()

    testV1ContainerImage ()
        Test V1ContainerImage
```

kubernetes.test.test_v1_container_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container_port.TestV1ContainerPort (methodName='runTest')
    Bases: unittest.case.TestCase
        V1ContainerPort unit test stubs

    setUp ()

    tearDown ()
```

```
testV1ContainerPort ()
    Test V1ContainerPort
```

kubernetes.test.test_v1_container_state module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container_state.TestV1ContainerState (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ContainerState unit test stubs
    setUp ()
    tearDown ()
    testV1ContainerState ()
        Test V1ContainerState
```

kubernetes.test.test_v1_container_state_running module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container_state_running.TestV1ContainerStateRunning (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ContainerStateRunning unit test stubs
    setUp ()
    tearDown ()
    testV1ContainerStateRunning ()
        Test V1ContainerStateRunning
```

kubernetes.test.test_v1_container_state_terminated module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_container_state_terminated.TestV1ContainerStateTerminated (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ContainerStateTerminated unit test stubs
```

```

setUp ()
tearDown ()
testV1ContainerStateTerminated ()
    Test V1ContainerStateTerminated

```

kubernetes.test.test_v1_container_state_waiting module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_container_state_waiting.TestV1ContainerStateWaiting (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ContainerStateWaiting unit test stubs

    setUp ()
    tearDown ()
    testV1ContainerStateWaiting ()
        Test V1ContainerStateWaiting

```

kubernetes.test.test_v1_container_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_container_status.TestV1ContainerStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ContainerStatus unit test stubs

    setUp ()
    tearDown ()
    testV1ContainerStatus ()
        Test V1ContainerStatus

```

kubernetes.test.test_v1_cross_version_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_cross_version_object_reference.TestV1CrossVersionObjectReference
    Bases: unittest.case.TestCase
    V1CrossVersionObjectReference unit test stubs
    setUp()
    tearDown()
    testV1CrossVersionObjectReference()
        Test V1CrossVersionObjectReference

```

kubernetes.test.test_v1_daemon_endpoint module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_daemon_endpoint.TestV1DaemonEndpoint (methodName='runTest')
    Bases: unittest.case.TestCase
    V1DaemonEndpoint unit test stubs
    setUp()
    tearDown()
    testV1DaemonEndpoint()
        Test V1DaemonEndpoint

```

kubernetes.test.test_v1_delete_options module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_delete_options.TestV1DeleteOptions (methodName='runTest')
    Bases: unittest.case.TestCase
    V1DeleteOptions unit test stubs
    setUp()
    tearDown()
    testV1DeleteOptions()
        Test V1DeleteOptions

```

kubernetes.test.test_v1_downward_api_volume_file module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_downward_api_volume_file.TestV1DownwardAPIVolumeFile (methodName=  
    Bases: unittest.case.TestCase  
    V1DownwardAPIVolumeFile unit test stubs  
    setUp ()  
    tearDown ()  
    testV1DownwardAPIVolumeFile ()  
        Test V1DownwardAPIVolumeFile
```

kubernetes.test.test_v1_downward_api_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_downward_api_volume_source.TestV1DownwardAPIVolumeSource (methodName=  
    Bases: unittest.case.TestCase  
    V1DownwardAPIVolumeSource unit test stubs  
    setUp ()  
    tearDown ()  
    testV1DownwardAPIVolumeSource ()  
        Test V1DownwardAPIVolumeSource
```

kubernetes.test.test_v1_empty_dir_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_empty_dir_volume_source.TestV1EmptyDirVolumeSource (methodName='ru  
    Bases: unittest.case.TestCase  
    V1EmptyDirVolumeSource unit test stubs  
    setUp ()  
    tearDown ()  
    testV1EmptyDirVolumeSource ()  
        Test V1EmptyDirVolumeSource
```


kubernetes.test.test_v1_endpoint_address module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_endpoint_address.TestV1EndpointAddress (methodName='runTest')
    Bases: unittest.case.TestCase
        V1EndpointAddress unit test stubs
        setUp ()
        tearDown ()
        testV1EndpointAddress ()
            Test V1EndpointAddress
```

kubernetes.test.test_v1_endpoint_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_endpoint_port.TestV1EndpointPort (methodName='runTest')
    Bases: unittest.case.TestCase
        V1EndpointPort unit test stubs
        setUp ()
        tearDown ()
        testV1EndpointPort ()
            Test V1EndpointPort
```

kubernetes.test.test_v1_endpoint_subset module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_endpoint_subset.TestV1EndpointSubset (methodName='runTest')
    Bases: unittest.case.TestCase
        V1EndpointSubset unit test stubs
        setUp ()
        tearDown ()
```

```
testV1EndpointSubset ()  
    Test V1EndpointSubset
```

kubernetes.test.test_v1_endpoints module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_endpoints.TestV1Endpoints (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1Endpoints unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1Endpoints ()  
        Test V1Endpoints
```

kubernetes.test.test_v1_endpoints_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_endpoints_list.TestV1EndpointsList (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1EndpointsList unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1EndpointsList ()  
        Test V1EndpointsList
```

kubernetes.test.test_v1_env_var module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_env_var.TestV1EnvVar (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1EnvVar unit test stubs
```

```

setUp ()
tearDown ()
testV1EnvVar ()
    Test V1EnvVar

```

kubernetes.test.test_v1_env_var_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_env_var_source.TestV1EnvVarSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1EnvVarSource unit test stubs

    setUp ()
    tearDown ()
    testV1EnvVarSource ()
        Test V1EnvVarSource

```

kubernetes.test.test_v1_event module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_event.TestV1Event (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Event unit test stubs

    setUp ()
    tearDown ()
    testV1Event ()
        Test V1Event

```

kubernetes.test.test_v1_event_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_event_list.TestV1EventList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1EventList unit test stubs
    setUp()
    tearDown()
    testV1EventList()
        Test V1EventList
```

kubernetes.test.test_v1_event_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_event_source.TestV1EventSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1EventSource unit test stubs
    setUp()
    tearDown()
    testV1EventSource()
        Test V1EventSource
```

kubernetes.test.test_v1_exec_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_exec_action.TestV1ExecAction (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ExecAction unit test stubs
    setUp()
    tearDown()
    testV1ExecAction()
        Test V1ExecAction
```

kubernetes.test.test_v1_fc_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_fc_volume_source.TestV1FCVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1FCVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1FCVolumeSource()
        Test V1FCVolumeSource
```

kubernetes.test.test_v1_flex_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_flex_volume_source.TestV1FlexVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1FlexVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1FlexVolumeSource()
        Test V1FlexVolumeSource
```

kubernetes.test.test_v1_flocker_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_flocker_volume_source.TestV1FlockerVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1FlockerVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1FlockerVolumeSource()
        Test V1FlockerVolumeSource
```

kubernetes.test.test_v1_gce_persistent_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_gce_persistent_disk_volume_source.TestV1GCEPersistentDiskVolumeSource
    Bases: unittest.case.TestCase
    V1GCEPersistentDiskVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1GCEPersistentDiskVolumeSource()
        Test V1GCEPersistentDiskVolumeSource
```

kubernetes.test.test_v1_git_repo_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_git_repo_volume_source.TestV1GitRepoVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1GitRepoVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1GitRepoVolumeSource()
        Test V1GitRepoVolumeSource
```

kubernetes.test.test_v1_glusterfs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_glusterfs_volume_source.TestV1GlusterfsVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1GlusterfsVolumeSource unit test stubs
    setUp()
    tearDown()
```

```
testV1GlusterfsVolumeSource ()
    Test V1GlusterfsVolumeSource
```

kubernetes.test.test_v1_handler module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_handler.TestV1Handler (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Handler unit test stubs
    setUp ()
    tearDown ()
    testV1Handler ()
        Test V1Handler
```

kubernetes.test.test_v1_horizontal_pod_autoscaler module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_horizontal_pod_autoscaler.TestV1HorizontalPodAutoscaler (methodName='runTest')
    Bases: unittest.case.TestCase
    V1HorizontalPodAutoscaler unit test stubs
    setUp ()
    tearDown ()
    testV1HorizontalPodAutoscaler ()
        Test V1HorizontalPodAutoscaler
```

kubernetes.test.test_v1_horizontal_pod_autoscaler_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_horizontal_pod_autoscaler_list.TestV1HorizontalPodAutoscalerList
    Bases: unittest.case.TestCase
    V1HorizontalPodAutoscalerList unit test stubs
```

```

setUp ()
tearDown ()
testV1HorizontalPodAutoscalerList ()
    Test V1HorizontalPodAutoscalerList

```

kubernetes.test.test_v1_horizontal_pod_autoscaler_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_horizontal_pod_autoscaler_spec.TestV1HorizontalPodAutoscalerSpec
    Bases: unittest.case.TestCase
    V1HorizontalPodAutoscalerSpec unit test stubs
    setUp ()
    tearDown ()
    testV1HorizontalPodAutoscalerSpec ()
        Test V1HorizontalPodAutoscalerSpec

```

kubernetes.test.test_v1_horizontal_pod_autoscaler_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_horizontal_pod_autoscaler_status.TestV1HorizontalPodAutoscalerStatus
    Bases: unittest.case.TestCase
    V1HorizontalPodAutoscalerStatus unit test stubs
    setUp ()
    tearDown ()
    testV1HorizontalPodAutoscalerStatus ()
        Test V1HorizontalPodAutoscalerStatus

```

kubernetes.test.test_v1_host_path_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>


```
class kubernetes.test.test_v1_host_path_volume_source.TestV1HostPathVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1HostPathVolumeSource unit test stubs

    setUp ()

    tearDown ()

    testV1HostPathVolumeSource ()
        Test V1HostPathVolumeSource
```

kubernetes.test.test_v1_http_get_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_http_get_action.TestV1HTTPGetAction (methodName='runTest')
    Bases: unittest.case.TestCase
    V1HTTPGetAction unit test stubs

    setUp ()

    tearDown ()

    testV1HTTPGetAction ()
        Test V1HTTPGetAction
```

kubernetes.test.test_v1_http_header module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_http_header.TestV1HTTPHeader (methodName='runTest')
    Bases: unittest.case.TestCase
    V1HTTPHeader unit test stubs

    setUp ()

    tearDown ()

    testV1HTTPHeader ()
        Test V1HTTPHeader
```

kubernetes.test.test_v1_iscsi_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_iscsi_volume_source.TestV1ISCSIVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ISCSIVolumeSource unit test stubs

    setUp()

    tearDown()

    testV1ISCSIVolumeSource()
        Test V1ISCSIVolumeSource
```

kubernetes.test.test_v1_job module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_job.TestV1Job (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Job unit test stubs

    setUp()

    tearDown()

    testV1Job()
        Test V1Job
```

kubernetes.test.test_v1_job_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_job_condition.TestV1JobCondition (methodName='runTest')
    Bases: unittest.case.TestCase
    V1JobCondition unit test stubs

    setUp()

    tearDown()

    testV1JobCondition()
        Test V1JobCondition
```

kubernetes.test.test_v1_job_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_job_list.TestV1JobList (methodName='runTest')
    Bases: unittest.case.TestCase
        V1JobList unit test stubs

    setUp ()

    tearDown ()

    testV1JobList ()
        Test V1JobList

```

kubernetes.test.test_v1_job_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_job_spec.TestV1JobSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1JobSpec unit test stubs

    setUp ()

    tearDown ()

    testV1JobSpec ()
        Test V1JobSpec

```

kubernetes.test.test_v1_job_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_job_status.TestV1JobStatus (methodName='runTest')
    Bases: unittest.case.TestCase
        V1JobStatus unit test stubs

    setUp ()

    tearDown ()

```

```
testV1JobStatus ()  
    Test V1JobStatus
```

kubernetes.test.test_v1_key_to_path module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_key_to_path.TestV1KeyToPath (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1KeyToPath unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1KeyToPath ()  
        Test V1KeyToPath
```

kubernetes.test.test_v1_lifecycle module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_lifecycle.TestV1Lifecycle (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1Lifecycle unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1Lifecycle ()  
        Test V1Lifecycle
```

kubernetes.test.test_v1_limit_range module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_limit_range.TestV1LimitRange (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1LimitRange unit test stubs
```

```

setUp ()
tearDown ()
testV1LimitRange ()
    Test V1LimitRange

```

kubernetes.test.test_v1_limit_range_item module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_limit_range_item.TestV1LimitRangeItem (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LimitRangeItem unit test stubs
    setUp ()
    tearDown ()
    testV1LimitRangeItem ()
        Test V1LimitRangeItem

```

kubernetes.test.test_v1_limit_range_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_limit_range_list.TestV1LimitRangeList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LimitRangeList unit test stubs
    setUp ()
    tearDown ()
    testV1LimitRangeList ()
        Test V1LimitRangeList

```

kubernetes.test.test_v1_limit_range_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_limit_range_spec.TestV1LimitRangeSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LimitRangeSpec unit test stubs
    setUp ()
    tearDown ()
    testV1LimitRangeSpec ()
        Test V1LimitRangeSpec

```

kubernetes.test.test_v1_load_balancer_ingress module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_load_balancer_ingress.TestV1LoadBalancerIngress (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LoadBalancerIngress unit test stubs
    setUp ()
    tearDown ()
    testV1LoadBalancerIngress ()
        Test V1LoadBalancerIngress

```

kubernetes.test.test_v1_load_balancer_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_load_balancer_status.TestV1LoadBalancerStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LoadBalancerStatus unit test stubs
    setUp ()
    tearDown ()
    testV1LoadBalancerStatus ()
        Test V1LoadBalancerStatus

```

kubernetes.test.test_v1_local_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_local_object_reference.TestV1LocalObjectReference (methodName='runTest')
    Bases: unittest.case.TestCase
    V1LocalObjectReference unit test stubs

    setUp ()

    tearDown ()

    testV1LocalObjectReference ()
        Test V1LocalObjectReference
```

kubernetes.test.test_v1_namespace module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_namespace.TestV1Namespace (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Namespace unit test stubs

    setUp ()

    tearDown ()

    testV1Namespace ()
        Test V1Namespace
```

kubernetes.test.test_v1_namespace_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_namespace_list.TestV1NamespaceList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NamespaceList unit test stubs

    setUp ()

    tearDown ()

    testV1NamespaceList ()
        Test V1NamespaceList
```

kubernetes.test.test_v1_namespace_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_namespace_spec.TestV1NamespaceSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1NamespaceSpec unit test stubs

    setUp ()

    tearDown ()

    testV1NamespaceSpec ()
        Test V1NamespaceSpec
```

kubernetes.test.test_v1_namespace_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_namespace_status.TestV1NamespaceStatus (methodName='runTest')
    Bases: unittest.case.TestCase
        V1NamespaceStatus unit test stubs

    setUp ()

    tearDown ()

    testV1NamespaceStatus ()
        Test V1NamespaceStatus
```

kubernetes.test.test_v1_nfs_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_nfs_volume_source.TestV1NFSVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
        V1NFSVolumeSource unit test stubs

    setUp ()

    tearDown ()
```



```
testV1NFSSource ()
    Test V1NFSSource
```

kubernetes.test.test_v1_node module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_node.TestV1Node (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Node unit test stubs
    setUp ()
    tearDown ()
    testV1Node ()
        Test V1Node
```

kubernetes.test.test_v1_node_address module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_node_address.TestV1NodeAddress (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NodeAddress unit test stubs
    setUp ()
    tearDown ()
    testV1NodeAddress ()
        Test V1NodeAddress
```

kubernetes.test.test_v1_node_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_node_condition.TestV1NodeCondition (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NodeCondition unit test stubs
```

```
setUp()  
tearDown()  
testV1NodeCondition()  
    Test V1NodeCondition
```

kubernetes.test.test_v1_node_daemon_endpoints module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_node_daemon_endpoints.TestV1NodeDaemonEndpoints (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1NodeDaemonEndpoints unit test stubs  
  
    setUp()  
    tearDown()  
  
    testV1NodeDaemonEndpoints()  
        Test V1NodeDaemonEndpoints
```

kubernetes.test.test_v1_node_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_node_list.TestV1NodeList (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1NodeList unit test stubs  
  
    setUp()  
    tearDown()  
  
    testV1NodeList()  
        Test V1NodeList
```

kubernetes.test.test_v1_node_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_node_spec.TestV1NodeSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NodeSpec unit test stubs
    setUp ()
    tearDown ()
    testV1NodeSpec ()
        Test V1NodeSpec

```

kubernetes.test.test_v1_node_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_node_status.TestV1NodeStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NodeStatus unit test stubs
    setUp ()
    tearDown ()
    testV1NodeStatus ()
        Test V1NodeStatus

```

kubernetes.test.test_v1_node_system_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_node_system_info.TestV1NodeSystemInfo (methodName='runTest')
    Bases: unittest.case.TestCase
    V1NodeSystemInfo unit test stubs
    setUp ()
    tearDown ()
    testV1NodeSystemInfo ()
        Test V1NodeSystemInfo

```

kubernetes.test.test_v1_object_field_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_object_field_selector.TestV1ObjectFieldSelector (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ObjectFieldSelector unit test stubs

    setUp ()

    tearDown ()

    testV1ObjectFieldSelector ()
        Test V1ObjectFieldSelector
```

kubernetes.test.test_v1_object_meta module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_object_meta.TestV1ObjectMeta (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ObjectMeta unit test stubs

    setUp ()

    tearDown ()

    testV1ObjectMeta ()
        Test V1ObjectMeta
```

kubernetes.test.test_v1_object_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_object_reference.TestV1ObjectReference (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ObjectReference unit test stubs

    setUp ()

    tearDown ()

    testV1ObjectReference ()
        Test V1ObjectReference
```

kubernetes.test.test_v1_owner_reference module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_owner_reference.TestV1OwnerReference (methodName='runTest')
    Bases: unittest.case.TestCase
        V1OwnerReference unit test stubs

    setUp ()

    tearDown ()

    testV1OwnerReference ()
        Test V1OwnerReference

```

kubernetes.test.test_v1_persistent_volume module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_persistent_volume.TestV1PersistentVolume (methodName='runTest')
    Bases: unittest.case.TestCase
        V1PersistentVolume unit test stubs

    setUp ()

    tearDown ()

    testV1PersistentVolume ()
        Test V1PersistentVolume

```

kubernetes.test.test_v1_persistent_volume_claim module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_persistent_volume_claim.TestV1PersistentVolumeClaim (methodName='runTest')
    Bases: unittest.case.TestCase
        V1PersistentVolumeClaim unit test stubs

    setUp ()

    tearDown ()

```

```
testV1PersistentVolumeClaim()  
    Test V1PersistentVolumeClaim
```

kubernetes.test.test_v1_persistent_volume_claim_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_persistent_volume_claim_list.TestV1PersistentVolumeClaimList(m  
    Bases: unittest.case.TestCase  
    V1PersistentVolumeClaimList unit test stubs  
    setUp()  
    tearDown()  
    testV1PersistentVolumeClaimList()  
        Test V1PersistentVolumeClaimList
```

kubernetes.test.test_v1_persistent_volume_claim_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_persistent_volume_claim_spec.TestV1PersistentVolumeClaimSpec(m  
    Bases: unittest.case.TestCase  
    V1PersistentVolumeClaimSpec unit test stubs  
    setUp()  
    tearDown()  
    testV1PersistentVolumeClaimSpec()  
        Test V1PersistentVolumeClaimSpec
```

kubernetes.test.test_v1_persistent_volume_claim_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_persistent_volume_claim_status.TestV1PersistentVolumeClaimStatu  
    Bases: unittest.case.TestCase  
    V1PersistentVolumeClaimStatus unit test stubs
```

```

setUp()
tearDown()
testV1PersistentVolumeClaimStatus()
    Test V1PersistentVolumeClaimStatus

```

kubernetes.test.test_v1_persistent_volume_claim_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_persistent_volume_claim_volume_source.TestV1PersistentVolumeClaimVolumeSource
    Bases: unittest.case.TestCase
    V1PersistentVolumeClaimVolumeSource unit test stubs

    setUp()
    tearDown()
    testV1PersistentVolumeClaimVolumeSource()
        Test V1PersistentVolumeClaimVolumeSource

```

kubernetes.test.test_v1_persistent_volume_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_persistent_volume_list.TestV1PersistentVolumeList (methodName='run')
    Bases: unittest.case.TestCase
    V1PersistentVolumeList unit test stubs

    setUp()
    tearDown()
    testV1PersistentVolumeList()
        Test V1PersistentVolumeList

```

kubernetes.test.test_v1_persistent_volume_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_persistent_volume_spec.TestV1PersistentVolumeSpec (methodName='run')
    Bases: unittest.case.TestCase
    V1PersistentVolumeSpec unit test stubs
    setUp()
    tearDown()
    testV1PersistentVolumeSpec()
        Test V1PersistentVolumeSpec
```

kubernetes.test.test_v1_persistent_volume_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_persistent_volume_status.TestV1PersistentVolumeStatus (methodName='run')
    Bases: unittest.case.TestCase
    V1PersistentVolumeStatus unit test stubs
    setUp()
    tearDown()
    testV1PersistentVolumeStatus()
        Test V1PersistentVolumeStatus
```

kubernetes.test.test_v1_photon_persistent_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_photon_persistent_disk_volume_source.TestV1PhotonPersistentDiskVolumeSource (methodName='run')
    Bases: unittest.case.TestCase
    V1PhotonPersistentDiskVolumeSource unit test stubs
    setUp()
    tearDown()
    testV1PhotonPersistentDiskVolumeSource()
        Test V1PhotonPersistentDiskVolumeSource
```

kubernetes.test.test_v1_pod module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod.TestV1Pod (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Pod unit test stubs
    setUp ()
    tearDown ()
    testV1Pod ()
        Test V1Pod
```

kubernetes.test.test_v1_pod_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_condition.TestV1PodCondition (methodName='runTest')
    Bases: unittest.case.TestCase
    V1PodCondition unit test stubs
    setUp ()
    tearDown ()
    testV1PodCondition ()
        Test V1PodCondition
```

kubernetes.test.test_v1_pod_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_list.TestV1PodList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1PodList unit test stubs
    setUp ()
    tearDown ()
    testV1PodList ()
        Test V1PodList
```

kubernetes.test.test_v1_pod_security_context module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_security_context.TestV1PodSecurityContext (methodName='runTest')
    Bases: unittest.case.TestCase
        V1PodSecurityContext unit test stubs

    setUp ()

    tearDown ()

    testV1PodSecurityContext ()
        Test V1PodSecurityContext
```

kubernetes.test.test_v1_pod_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_spec.TestV1PodSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1PodSpec unit test stubs

    setUp ()

    tearDown ()

    testV1PodSpec ()
        Test V1PodSpec
```

kubernetes.test.test_v1_pod_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_status.TestV1PodStatus (methodName='runTest')
    Bases: unittest.case.TestCase
        V1PodStatus unit test stubs

    setUp ()

    tearDown ()
```

```
testV1PodStatus ()  
    Test V1PodStatus
```

kubernetes.test.test_v1_pod_template module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_template.TestV1PodTemplate (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1PodTemplate unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1PodTemplate ()  
        Test V1PodTemplate
```

kubernetes.test.test_v1_pod_template_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_template_list.TestV1PodTemplateList (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1PodTemplateList unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1PodTemplateList ()  
        Test V1PodTemplateList
```

kubernetes.test.test_v1_pod_template_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_pod_template_spec.TestV1PodTemplateSpec (methodName='runTest')  
    Bases: unittest.case.TestCase  
        V1PodTemplateSpec unit test stubs
```

```
setUp()  
tearDown()  
testV1PodTemplateSpec()  
    Test V1PodTemplateSpec
```

kubernetes.test.test_v1_preconditions module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_preconditions.TestV1Preconditions (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1Preconditions unit test stubs  
    setUp()  
    tearDown()  
    testV1Preconditions()  
        Test V1Preconditions
```

kubernetes.test.test_v1_probe module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_probe.TestV1Probe (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1Probe unit test stubs  
    setUp()  
    tearDown()  
    testV1Probe()  
        Test V1Probe
```

kubernetes.test.test_v1_quobyte_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_quobyte_volume_source.TestV1QuobyteVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1QuobyteVolumeSource unit test stubs

    setUp()

    tearDown()

    testV1QuobyteVolumeSource()
        Test V1QuobyteVolumeSource
```

kubernetes.test.test_v1_rbd_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_rbd_volume_source.TestV1RBDVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1RBDVolumeSource unit test stubs

    setUp()

    tearDown()

    testV1RBDVolumeSource()
        Test V1RBDVolumeSource
```

kubernetes.test.test_v1_replication_controller module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_replication_controller.TestV1ReplicationController (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ReplicationController unit test stubs

    setUp()

    tearDown()

    testV1ReplicationController()
        Test V1ReplicationController
```

kubernetes.test.test_v1_replication_controller_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_replication_controller_condition.TestV1ReplicationControllerCon
    Bases: unittest.case.TestCase
    V1ReplicationControllerCondition unit test stubs
    setUp()
    tearDown()
    testV1ReplicationControllerCondition()
        Test V1ReplicationControllerCondition
```

kubernetes.test.test_v1_replication_controller_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_replication_controller_list.TestV1ReplicationControllerList (met
    Bases: unittest.case.TestCase
    V1ReplicationControllerList unit test stubs
    setUp()
    tearDown()
    testV1ReplicationControllerList()
        Test V1ReplicationControllerList
```

kubernetes.test.test_v1_replication_controller_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_replication_controller_spec.TestV1ReplicationControllerSpec (met
    Bases: unittest.case.TestCase
    V1ReplicationControllerSpec unit test stubs
    setUp()
    tearDown()
    testV1ReplicationControllerSpec()
        Test V1ReplicationControllerSpec
```

kubernetes.test.test_v1_replication_controller_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_replication_controller_status.TestV1ReplicationControllerStatus
    Bases: unittest.case.TestCase
    V1ReplicationControllerStatus unit test stubs
    setUp()
    tearDown()
    testV1ReplicationControllerStatus()
        Test V1ReplicationControllerStatus
```

kubernetes.test.test_v1_resource_field_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_resource_field_selector.TestV1ResourceFieldSelector (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ResourceFieldSelector unit test stubs
    setUp()
    tearDown()
    testV1ResourceFieldSelector()
        Test V1ResourceFieldSelector
```

kubernetes.test.test_v1_resource_quota module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_resource_quota.TestV1ResourceQuota (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ResourceQuota unit test stubs
    setUp()
    tearDown()
```

```
testV1ResourceQuota ()  
    Test V1ResourceQuota
```

kubernetes.test.test_v1_resource_quota_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_resource_quota_list.TestV1ResourceQuotaList (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1ResourceQuotaList unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1ResourceQuotaList ()  
        Test V1ResourceQuotaList
```

kubernetes.test.test_v1_resource_quota_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_resource_quota_spec.TestV1ResourceQuotaSpec (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1ResourceQuotaSpec unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1ResourceQuotaSpec ()  
        Test V1ResourceQuotaSpec
```

kubernetes.test.test_v1_resource_quota_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_resource_quota_status.TestV1ResourceQuotaStatus (methodName='runTest')  
    Bases: unittest.case.TestCase  
    V1ResourceQuotaStatus unit test stubs
```



```

setUp ()
tearDown ()
testV1ResourceQuotaStatus ()
    Test V1ResourceQuotaStatus

```

kubernetes.test.test_v1_resource_requirements module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_resource_requirements.TestV1ResourceRequirements (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ResourceRequirements unit test stubs

    setUp ()
    tearDown ()
    testV1ResourceRequirements ()
        Test V1ResourceRequirements

```

kubernetes.test.test_v1_scale module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1_scale.TestV1Scale (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Scale unit test stubs

    setUp ()
    tearDown ()
    testV1Scale ()
        Test V1Scale

```

kubernetes.test.test_v1_scale_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_scale_spec.TestV1ScaleSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ScaleSpec unit test stubs
    setUp ()
    tearDown ()
    testV1ScaleSpec ()
        Test V1ScaleSpec
```

kubernetes.test.test_v1_scale_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_scale_status.TestV1ScaleStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ScaleStatus unit test stubs
    setUp ()
    tearDown ()
    testV1ScaleStatus ()
        Test V1ScaleStatus
```

kubernetes.test.test_v1_se_linux_options module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_se_linux_options.TestV1SELinuxOptions (methodName='runTest')
    Bases: unittest.case.TestCase
    V1SELinuxOptions unit test stubs
    setUp ()
    tearDown ()
    testV1SELinuxOptions ()
        Test V1SELinuxOptions
```

kubernetes.test.test_v1_secret module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_secret.TestV1Secret (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Secret unit test stubs
    setUp()
    tearDown()
    testV1Secret()
        Test V1Secret
```

kubernetes.test.test_v1_secret_key_selector module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_secret_key_selector.TestV1SecretKeySelector (methodName='runTest')
    Bases: unittest.case.TestCase
    V1SecretKeySelector unit test stubs
    setUp()
    tearDown()
    testV1SecretKeySelector()
        Test V1SecretKeySelector
```

kubernetes.test.test_v1_secret_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_secret_list.TestV1SecretList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1SecretList unit test stubs
    setUp()
    tearDown()
    testV1SecretList()
        Test V1SecretList
```

kubernetes.test.test_v1_secret_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_secret_volume_source.TestV1SecretVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
        V1SecretVolumeSource unit test stubs

    setUp ()

    tearDown ()

    testV1SecretVolumeSource ()
        Test V1SecretVolumeSource
```

kubernetes.test.test_v1_security_context module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_security_context.TestV1SecurityContext (methodName='runTest')
    Bases: unittest.case.TestCase
        V1SecurityContext unit test stubs

    setUp ()

    tearDown ()

    testV1SecurityContext ()
        Test V1SecurityContext
```

kubernetes.test.test_v1_service module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service.TestV1Service (methodName='runTest')
    Bases: unittest.case.TestCase
        V1Service unit test stubs

    setUp ()

    tearDown ()
```

```
testV1Service()
    Test V1Service
```

kubernetes.test.test_v1_service_account module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_account.TestV1ServiceAccount (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServiceAccount unit test stubs

    setUp()

    tearDown()

    testV1ServiceAccount()
        Test V1ServiceAccount
```

kubernetes.test.test_v1_service_account_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_account_list.TestV1ServiceAccountList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServiceAccountList unit test stubs

    setUp()

    tearDown()

    testV1ServiceAccountList()
        Test V1ServiceAccountList
```

kubernetes.test.test_v1_service_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_list.TestV1ServiceList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServiceList unit test stubs
```

```
setUp ()
tearDown ()
testV1ServiceList ()
    Test V1ServiceList
```

kubernetes.test.test_v1_service_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_port.TestV1ServicePort (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServicePort unit test stubs

    setUp ()
    tearDown ()
    testV1ServicePort ()
        Test V1ServicePort
```

kubernetes.test.test_v1_service_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_spec.TestV1ServiceSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServiceSpec unit test stubs

    setUp ()
    tearDown ()
    testV1ServiceSpec ()
        Test V1ServiceSpec
```

kubernetes.test.test_v1_service_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_service_status.TestV1ServiceStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V1ServiceStatus unit test stubs
    setUp ()
    tearDown ()
    testV1ServiceStatus ()
        Test V1ServiceStatus
```

kubernetes.test.test_v1_tcp_socket_action module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_tcp_socket_action.TestV1TCPSocketAction (methodName='runTest')
    Bases: unittest.case.TestCase
    V1TCPSocketAction unit test stubs
    setUp ()
    tearDown ()
    testV1TCPSocketAction ()
        Test V1TCPSocketAction
```

kubernetes.test.test_v1_volume module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_volume.TestV1Volume (methodName='runTest')
    Bases: unittest.case.TestCase
    V1Volume unit test stubs
    setUp ()
    tearDown ()
    testV1Volume ()
        Test V1Volume
```

kubernetes.test.test_v1_volume_mount module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_volume_mount.TestV1VolumeMount (methodName='runTest')
    Bases: unittest.case.TestCase
    V1VolumeMount unit test stubs

    setUp ()
    tearDown ()
    testV1VolumeMount ()
        Test V1VolumeMount
```

kubernetes.test.test_v1_vsphere_virtual_disk_volume_source module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1_vsphere_virtual_disk_volume_source.TestV1VsphereVirtualDiskVolumeSource (methodName='runTest')
    Bases: unittest.case.TestCase
    V1VsphereVirtualDiskVolumeSource unit test stubs

    setUp ()
    tearDown ()
    testV1VsphereVirtualDiskVolumeSource ()
        Test V1VsphereVirtualDiskVolumeSource
```

kubernetes.test.test_v1alpha1_certificate_signing_request module

kubernetes.test.test_v1alpha1_certificate_signing_request_condition module

kubernetes.test.test_v1alpha1_certificate_signing_request_list module

kubernetes.test.test_v1alpha1_certificate_signing_request_spec module

kubernetes.test.test_v1alpha1_certificate_signing_request_status module

kubernetes.test.test_v1alpha1_cluster_role module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>


```
class kubernetes.test.test_v1alpha1_cluster_role.TestV1alpha1ClusterRole (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1ClusterRole unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1ClusterRole ()
        Test V1alpha1ClusterRole
```

kubernetes.test.test_v1alpha1_cluster_role_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_cluster_role_binding.TestV1alpha1ClusterRoleBinding (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1ClusterRoleBinding unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1ClusterRoleBinding ()
        Test V1alpha1ClusterRoleBinding
```

kubernetes.test.test_v1alpha1_cluster_role_binding_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_cluster_role_binding_list.TestV1alpha1ClusterRoleBindingList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1ClusterRoleBindingList unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1ClusterRoleBindingList ()
        Test V1alpha1ClusterRoleBindingList
```

kubernetes.test.test_v1alpha1_cluster_role_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_cluster_role_list.TestV1alpha1ClusterRoleList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1ClusterRoleList unit test stubs

    setUp()

    tearDown()

    testV1alpha1ClusterRoleList()
        Test V1alpha1ClusterRoleList
```

kubernetes.test.test_v1alpha1_policy_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_policy_rule.TestV1alpha1PolicyRule (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1PolicyRule unit test stubs

    setUp()

    tearDown()

    testV1alpha1PolicyRule()
        Test V1alpha1PolicyRule
```

kubernetes.test.test_v1alpha1_role module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_role.TestV1alpha1Role (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1Role unit test stubs

    setUp()

    tearDown()

    testV1alpha1Role()
        Test V1alpha1Role
```

kubernetes.test.test_v1alpha1_role_binding module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_role_binding.TestV1alpha1RoleBinding (methodName='runTest')
    Bases: unittest.case.TestCase
        V1alpha1RoleBinding unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1RoleBinding ()
        Test V1alpha1RoleBinding
```

kubernetes.test.test_v1alpha1_role_binding_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_role_binding_list.TestV1alpha1RoleBindingList (methodName='runTest')
    Bases: unittest.case.TestCase
        V1alpha1RoleBindingList unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1RoleBindingList ()
        Test V1alpha1RoleBindingList
```

kubernetes.test.test_v1alpha1_role_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_role_list.TestV1alpha1RoleList (methodName='runTest')
    Bases: unittest.case.TestCase
        V1alpha1RoleList unit test stubs
    setUp ()
    tearDown ()
```

```
testV1alpha1RoleList ()
    Test V1alpha1RoleList
```

kubernetes.test.test_v1alpha1_role_ref module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_role_ref.TestV1alpha1RoleRef (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1RoleRef unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1RoleRef ()
        Test V1alpha1RoleRef
```

kubernetes.test.test_v1alpha1_subject module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1alpha1_subject.TestV1alpha1Subject (methodName='runTest')
    Bases: unittest.case.TestCase
    V1alpha1Subject unit test stubs
    setUp ()
    tearDown ()
    testV1alpha1Subject ()
        Test V1alpha1Subject
```

kubernetes.test.test_v1beta1_api_version module

kubernetes.test.test_v1beta1_cpu_target_utilization module

kubernetes.test.test_v1beta1_daemon_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_daemon_set.TestV1beta1DaemonSet (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1DaemonSet unit test stubs
    setUp ()
    tearDown ()
    testV1beta1DaemonSet ()
        Test V1beta1DaemonSet

```

kubernetes.test.test_v1beta1_daemon_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_daemon_set_list.TestV1beta1DaemonSetList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1DaemonSetList unit test stubs
    setUp ()
    tearDown ()
    testV1beta1DaemonSetList ()
        Test V1beta1DaemonSetList

```

kubernetes.test.test_v1beta1_daemon_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_daemon_set_spec.TestV1beta1DaemonSetSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1DaemonSetSpec unit test stubs
    setUp ()
    tearDown ()
    testV1beta1DaemonSetSpec ()
        Test V1beta1DaemonSetSpec

```

kubernetes.test.test_v1beta1_daemon_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_daemon_set_status.TestV1beta1DaemonSetStatus (methodName='run')
    Bases: unittest.case.TestCase
    V1beta1DaemonSetStatus unit test stubs

    setUp()

    tearDown()

    testV1beta1DaemonSetStatus()
        Test V1beta1DaemonSetStatus
```

kubernetes.test.test_v1beta1_deployment module

kubernetes.test.test_v1beta1_deployment_condition module

kubernetes.test.test_v1beta1_deployment_list module

kubernetes.test.test_v1beta1_deployment_rollback module

kubernetes.test.test_v1beta1_deployment_spec module

kubernetes.test.test_v1beta1_deployment_status module

kubernetes.test.test_v1beta1_deployment_strategy module

kubernetes.test.test_v1beta1_eviction module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_eviction.TestV1beta1Eviction (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1Eviction unit test stubs

    setUp()

    tearDown()

    testV1beta1Eviction()
        Test V1beta1Eviction
```

kubernetes.test.test_v1beta1_horizontal_pod_autoscaler module

kubernetes.test.test_v1beta1_horizontal_pod_autoscaler_list module

kubernetes.test.test_v1beta1_horizontal_pod_autoscaler_spec module

kubernetes.test.test_v1beta1_horizontal_pod_autoscaler_status module

kubernetes.test.test_v1beta1_http_ingress_path module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_http_ingress_path.TestV1beta1HTTPIngressPath (methodName='run')
    Bases: unittest.case.TestCase
    V1beta1HTTPIngressPath unit test stubs

    setUp ()

    tearDown ()

    testV1beta1HTTPIngressPath ()
        Test V1beta1HTTPIngressPath
```

kubernetes.test.test_v1beta1_http_ingress_rule_value module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_http_ingress_rule_value.TestV1beta1HTTPIngressRuleValue (methodName='run')
    Bases: unittest.case.TestCase
    V1beta1HTTPIngressRuleValue unit test stubs

    setUp ()

    tearDown ()

    testV1beta1HTTPIngressRuleValue ()
        Test V1beta1HTTPIngressRuleValue
```

kubernetes.test.test_v1beta1_ingress module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_ingress.TestV1beta1Ingress (methodName='runTest')
    Bases: unittest.case.TestCase

    V1beta1Ingress unit test stubs

    setUp ()

    tearDown ()

    testV1beta1Ingress ()
        Test V1beta1Ingress
```

kubernetes.test.test_v1beta1_ingress_backend module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_ingress_backend.TestV1beta1IngressBackend (methodName='runTest')
    Bases: unittest.case.TestCase

    V1beta1IngressBackend unit test stubs

    setUp ()

    tearDown ()

    testV1beta1IngressBackend ()
        Test V1beta1IngressBackend
```

kubernetes.test.test_v1beta1_ingress_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_ingress_list.TestV1beta1IngressList (methodName='runTest')
    Bases: unittest.case.TestCase

    V1beta1IngressList unit test stubs

    setUp ()

    tearDown ()

    testV1beta1IngressList ()
        Test V1beta1IngressList
```


kubernetes.test.test_v1beta1_ingress_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_ingress_rule.TestV1beta1IngressRule (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1IngressRule unit test stubs

    setUp ()

    tearDown ()

    testV1beta1IngressRule ()
        Test V1beta1IngressRule

```

kubernetes.test.test_v1beta1_ingress_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_ingress_spec.TestV1beta1IngressSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1IngressSpec unit test stubs

    setUp ()

    tearDown ()

    testV1beta1IngressSpec ()
        Test V1beta1IngressSpec

```

kubernetes.test.test_v1beta1_ingress_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_ingress_status.TestV1beta1IngressStatus (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1IngressStatus unit test stubs

    setUp ()

    tearDown ()

```

```
testV1beta1IngressStatus ()
    Test V1beta1IngressStatus
```

kubernetes.test.test_v1beta1_ingress_tls module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_ingress_tls.TestV1beta1IngressTLS (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1IngressTLS unit test stubs

    setUp ()
    tearDown ()
    testV1beta1IngressTLS ()
        Test V1beta1IngressTLS
```

kubernetes.test.test_v1beta1_job module

kubernetes.test.test_v1beta1_job_condition module

kubernetes.test.test_v1beta1_job_list module

kubernetes.test.test_v1beta1_job_spec module

kubernetes.test.test_v1beta1_job_status module

kubernetes.test.test_v1beta1_local_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_local_subject_access_review.TestV1beta1LocalSubjectAccessReview (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1LocalSubjectAccessReview unit test stubs

    setUp ()
    tearDown ()
    testV1beta1LocalSubjectAccessReview ()
        Test V1beta1LocalSubjectAccessReview
```

kubernetes.test.test_v1beta1_network_policy module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_network_policy.TestV1beta1NetworkPolicy (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1NetworkPolicy unit test stubs

    setUp()

    tearDown()

    testV1beta1NetworkPolicy()
        Test V1beta1NetworkPolicy

```

kubernetes.test.test_v1beta1_network_policy_ingress_rule module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_network_policy_ingress_rule.TestV1beta1NetworkPolicyIngressRule
    Bases: unittest.case.TestCase
        V1beta1NetworkPolicyIngressRule unit test stubs

    setUp()

    tearDown()

    testV1beta1NetworkPolicyIngressRule()
        Test V1beta1NetworkPolicyIngressRule

```

kubernetes.test.test_v1beta1_network_policy_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_network_policy_list.TestV1beta1NetworkPolicyList (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1NetworkPolicyList unit test stubs

    setUp()

    tearDown()

```

```
testV1beta1NetworkPolicyList ()  
    Test V1beta1NetworkPolicyList
```

kubernetes.test.test_v1beta1_network_policy_peer module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_network_policy_peer.TestV1beta1NetworkPolicyPeer (methodName)  
    Bases: unittest.case.TestCase  
    V1beta1NetworkPolicyPeer unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1beta1NetworkPolicyPeer ()  
        Test V1beta1NetworkPolicyPeer
```

kubernetes.test.test_v1beta1_network_policy_port module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_network_policy_port.TestV1beta1NetworkPolicyPort (methodName)  
    Bases: unittest.case.TestCase  
    V1beta1NetworkPolicyPort unit test stubs  
  
    setUp ()  
  
    tearDown ()  
  
    testV1beta1NetworkPolicyPort ()  
        Test V1beta1NetworkPolicyPort
```

kubernetes.test.test_v1beta1_network_policy_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_network_policy_spec.TestV1beta1NetworkPolicySpec (methodName)  
    Bases: unittest.case.TestCase  
    V1beta1NetworkPolicySpec unit test stubs
```

```

setUp ()
tearDown ()
testV1beta1NetworkPolicySpec ()
    Test V1beta1NetworkPolicySpec

```

kubernetes.test.test_v1beta1_non_resource_attributes module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_non_resource_attributes.TestV1beta1NonResourceAttributes (unittest.case.TestCase)
    V1beta1NonResourceAttributes unit test stubs
    setUp ()
    tearDown ()
    testV1beta1NonResourceAttributes ()
        Test V1beta1NonResourceAttributes

```

kubernetes.test.test_v1beta1_pod_disruption_budget module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_pod_disruption_budget.TestV1beta1PodDisruptionBudget (unittest.case.TestCase)
    V1beta1PodDisruptionBudget unit test stubs
    setUp ()
    tearDown ()
    testV1beta1PodDisruptionBudget ()
        Test V1beta1PodDisruptionBudget

```

kubernetes.test.test_v1beta1_pod_disruption_budget_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_pod_disruption_budget_list.TestV1beta1PodDisruptionBudgetList:
    Bases: unittest.case.TestCase

    V1beta1PodDisruptionBudgetList unit test stubs

    setUp()

    tearDown()

    testV1beta1PodDisruptionBudgetList()
        Test V1beta1PodDisruptionBudgetList
```

kubernetes.test.test_v1beta1_pod_disruption_budget_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_pod_disruption_budget_spec.TestV1beta1PodDisruptionBudgetSpec:
    Bases: unittest.case.TestCase

    V1beta1PodDisruptionBudgetSpec unit test stubs

    setUp()

    tearDown()

    testV1beta1PodDisruptionBudgetSpec()
        Test V1beta1PodDisruptionBudgetSpec
```

kubernetes.test.test_v1beta1_pod_disruption_budget_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_pod_disruption_budget_status.TestV1beta1PodDisruptionBudgetStatus:
    Bases: unittest.case.TestCase

    V1beta1PodDisruptionBudgetStatus unit test stubs

    setUp()

    tearDown()

    testV1beta1PodDisruptionBudgetStatus()
        Test V1beta1PodDisruptionBudgetStatus
```

kubernetes.test.test_v1beta1_replica_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_replica_set.TestV1beta1ReplicaSet (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1ReplicaSet unit test stubs
    setUp()
    tearDown()
    testV1beta1ReplicaSet()
        Test V1beta1ReplicaSet
```

kubernetes.test.test_v1beta1_replica_set_condition module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_replica_set_condition.TestV1beta1ReplicaSetCondition (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1ReplicaSetCondition unit test stubs
    setUp()
    tearDown()
    testV1beta1ReplicaSetCondition()
        Test V1beta1ReplicaSetCondition
```

kubernetes.test.test_v1beta1_replica_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_replica_set_list.TestV1beta1ReplicaSetList (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1ReplicaSetList unit test stubs
    setUp()
    tearDown()
    testV1beta1ReplicaSetList()
        Test V1beta1ReplicaSetList
```

kubernetes.test.test_v1beta1_replica_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_replica_set_spec.TestV1beta1ReplicaSetSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1ReplicaSetSpec unit test stubs

    setUp ()

    tearDown ()

    testV1beta1ReplicaSetSpec ()
        Test V1beta1ReplicaSetSpec
```

kubernetes.test.test_v1beta1_replica_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_replica_set_status.TestV1beta1ReplicaSetStatus (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1ReplicaSetStatus unit test stubs

    setUp ()

    tearDown ()

    testV1beta1ReplicaSetStatus ()
        Test V1beta1ReplicaSetStatus
```

kubernetes.test.test_v1beta1_resource_attributes module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_resource_attributes.TestV1beta1ResourceAttributes (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1ResourceAttributes unit test stubs

    setUp ()

    tearDown ()
```



```
testV1beta1ResourceAttributes ()
    Test V1beta1ResourceAttributes
```

kubernetes.test.test_v1beta1_rollback_config module

kubernetes.test.test_v1beta1_rolling_update_deployment module

kubernetes.test.test_v1beta1_scale module

kubernetes.test.test_v1beta1_scale_spec module

kubernetes.test.test_v1beta1_scale_status module

kubernetes.test.test_v1beta1_self_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_self_subject_access_review.TestV1beta1SelfSubjectAccessReview
    Bases: unittest.case.TestCase
    V1beta1SelfSubjectAccessReview unit test stubs

    setUp ()
    tearDown ()
    testV1beta1SelfSubjectAccessReview ()
        Test V1beta1SelfSubjectAccessReview
```

kubernetes.test.test_v1beta1_self_subject_access_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_self_subject_access_review_spec.TestV1beta1SelfSubjectAccessReviewSpec
    Bases: unittest.case.TestCase
    V1beta1SelfSubjectAccessReviewSpec unit test stubs

    setUp ()
    tearDown ()
    testV1beta1SelfSubjectAccessReviewSpec ()
        Test V1beta1SelfSubjectAccessReviewSpec
```

kubernetes.test.test_v1beta1_stateful_set module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_stateful_set.TestV1beta1StatefulSet (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1StatefulSet unit test stubs

    setUp ()

    tearDown ()

    testV1beta1StatefulSet ()
        Test V1beta1StatefulSet
```

kubernetes.test.test_v1beta1_stateful_set_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_stateful_set_list.TestV1beta1StatefulSetList (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1StatefulSetList unit test stubs

    setUp ()

    tearDown ()

    testV1beta1StatefulSetList ()
        Test V1beta1StatefulSetList
```

kubernetes.test.test_v1beta1_stateful_set_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_stateful_set_spec.TestV1beta1StatefulSetSpec (methodName='runTest')
    Bases: unittest.case.TestCase
        V1beta1StatefulSetSpec unit test stubs

    setUp ()

    tearDown ()
```

```
testV1beta1StatefulSetSpec ()
    Test V1beta1StatefulSetSpec
```

kubernetes.test.test_v1beta1_stateful_set_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_stateful_set_status.TestV1beta1StatefulSetStatus (methodName=
    Bases: unittest.case.TestCase
    V1beta1StatefulSetStatus unit test stubs
    setUp ()
    tearDown ()
    testV1beta1StatefulSetStatus ()
        Test V1beta1StatefulSetStatus
```

kubernetes.test.test_v1beta1_storage_class module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_storage_class.TestV1beta1StorageClass (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1StorageClass unit test stubs
    setUp ()
    tearDown ()
    testV1beta1StorageClass ()
        Test V1beta1StorageClass
```

kubernetes.test.test_v1beta1_storage_class_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_storage_class_list.TestV1beta1StorageClassList (methodName=
    Bases: unittest.case.TestCase
    V1beta1StorageClassList unit test stubs
```

```

setUp ()
tearDown ()
testV1beta1StorageClassList ()
    Test V1beta1StorageClassList

```

kubernetes.test.test_v1beta1_subject_access_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_subject_access_review.TestV1beta1SubjectAccessReview (method
    Bases: unittest.case.TestCase
    V1beta1SubjectAccessReview unit test stubs
    setUp ()
    tearDown ()
    testV1beta1SubjectAccessReview ()
        Test V1beta1SubjectAccessReview

```

kubernetes.test.test_v1beta1_subject_access_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v1beta1_subject_access_review_spec.TestV1beta1SubjectAccessReviewSpec
    Bases: unittest.case.TestCase
    V1beta1SubjectAccessReviewSpec unit test stubs
    setUp ()
    tearDown ()
    testV1beta1SubjectAccessReviewSpec ()
        Test V1beta1SubjectAccessReviewSpec

```

kubernetes.test.test_v1beta1_subject_access_review_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_subject_access_review_status.TestV1beta1SubjectAccessReviewStatus:
    Bases: unittest.case.TestCase

    V1beta1SubjectAccessReviewStatus unit test stubs

    setUp()

    tearDown()

    testV1beta1SubjectAccessReviewStatus()
        Test V1beta1SubjectAccessReviewStatus
```

kubernetes.test.test_v1beta1_subresource_reference module

kubernetes.test.test_v1beta1_third_party_resource module

kubernetes.test.test_v1beta1_third_party_resource_list module

kubernetes.test.test_v1beta1_token_review module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_token_review.TestV1beta1TokenReview(methodName='runTest'):
    Bases: unittest.case.TestCase

    V1beta1TokenReview unit test stubs

    setUp()

    tearDown()

    testV1beta1TokenReview()
        Test V1beta1TokenReview
```

kubernetes.test.test_v1beta1_token_review_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_token_review_spec.TestV1beta1TokenReviewSpec(methodName='runTest'):
    Bases: unittest.case.TestCase

    V1beta1TokenReviewSpec unit test stubs

    setUp()

    tearDown()

    testV1beta1TokenReviewSpec()
        Test V1beta1TokenReviewSpec
```

kubernetes.test.test_v1beta1_token_review_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_token_review_status.TestV1beta1TokenReviewStatus (methodName
    Bases: unittest.case.TestCase
    V1beta1TokenReviewStatus unit test stubs
    setUp ()
    tearDown ()
    testV1beta1TokenReviewStatus ()
        Test V1beta1TokenReviewStatus
```

kubernetes.test.test_v1beta1_user_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v1beta1_user_info.TestV1beta1UserInfo (methodName='runTest')
    Bases: unittest.case.TestCase
    V1beta1UserInfo unit test stubs
    setUp ()
    tearDown ()
    testV1beta1UserInfo ()
        Test V1beta1UserInfo
```

kubernetes.test.test_v2alpha1_cron_job module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v2alpha1_cron_job.TestV2alpha1CronJob (methodName='runTest')
    Bases: unittest.case.TestCase
    V2alpha1CronJob unit test stubs
    setUp ()
    tearDown ()
```

```
testV2alpha1CronJob ()
    Test V2alpha1CronJob
```

kubernetes.test.test_v2alpha1_cron_job_list module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v2alpha1_cron_job_list.TestV2alpha1CronJobList (methodName='runTest')
    Bases: unittest.case.TestCase
    V2alpha1CronJobList unit test stubs
    setUp ()
    tearDown ()
    testV2alpha1CronJobList ()
        Test V2alpha1CronJobList
```

kubernetes.test.test_v2alpha1_cron_job_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v2alpha1_cron_job_spec.TestV2alpha1CronJobSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V2alpha1CronJobSpec unit test stubs
    setUp ()
    tearDown ()
    testV2alpha1CronJobSpec ()
        Test V2alpha1CronJobSpec
```

kubernetes.test.test_v2alpha1_cron_job_status module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_v2alpha1_cron_job_status.TestV2alpha1CronJobStatus (methodName='runTest')
    Bases: unittest.case.TestCase
    V2alpha1CronJobStatus unit test stubs
```

```

setUp ()
tearDown ()
testV2alpha1CronJobStatus ()
    Test V2alpha1CronJobStatus

```

kubernetes.test.test_v2alpha1_job module

kubernetes.test.test_v2alpha1_job_condition module

kubernetes.test.test_v2alpha1_job_list module

kubernetes.test.test_v2alpha1_job_spec module

kubernetes.test.test_v2alpha1_job_status module

kubernetes.test.test_v2alpha1_job_template_spec module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_v2alpha1_job_template_spec.TestV2alpha1JobTemplateSpec (methodName='runTest')
    Bases: unittest.case.TestCase
    V2alpha1JobTemplateSpec unit test stubs

    setUp ()
    tearDown ()
    testV2alpha1JobTemplateSpec ()
        Test V2alpha1JobTemplateSpec

```

kubernetes.test.test_version_api module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```

class kubernetes.test.test_version_api.TestVersionApi (methodName='runTest')
    Bases: unittest.case.TestCase
    VersionApi unit test stubs

    setUp ()
    tearDown ()

```



```
test_get_code ()
    Test case for get_code
```

kubernetes.test.test_version_info module

Kubernetes

No description provided (generated by Swagger Codegen <https://github.com/swagger-api/swagger-codegen>)

OpenAPI spec version: v1.8.2

Generated by: <https://github.com/swagger-api/swagger-codegen.git>

```
class kubernetes.test.test_version_info.TestVersionInfo (methodName='runTest')
    Bases: unittest.case.TestCase
    VersionInfo unit test stubs
    setUp ()
    tearDown ()
    testVersionInfo ()
        Test VersionInfo
```

kubernetes.test.test_versioned_event module

Module contents

kubernetes.watch package

Submodules

kubernetes.watch.watch module

```
class kubernetes.watch.watch.SimpleNamespace (**kwargs)
```

```
class kubernetes.watch.watch.Watch (return_type=None)
    Bases: object
```

```
    get_return_type (func)
```

```
    stop ()
```

```
    stream (func, *args, **kwargs)
```

Watch an API resource and stream the result back via a generator.

Parameters **func** – The API function pointer. Any parameter to the function can be passed after this parameter.

Returns

Event object with these keys: ‘type’: The type of event such as “ADDED”, “DELETED”, etc. ‘raw_object’: a dict representing the watched object. ‘object’: A model representation of raw_object. The name of

model will be determined based on the func’s doc string. If it cannot be determined, ‘object’ value will be the same as ‘raw_object’.

Example: `v1 = kubernetes.client.CoreV1Api() watch = kubernetes.watch.Watch() for e in watch.stream(v1.list_namespace, resource_version=1127):`

```
    type = e['type'] object = e['object'] # object is one of type return_type raw_object =
    e['raw_object'] # raw_object is a dict ... if should_stop:
```

```
        watch.stop()
```

```
    unmarshal_event (data, return_type)
```

```
kubernetes.watch.Watch.iter_resp_lines (resp)
```

kubernetes.watch.watch_test module

```
class kubernetes.watch.watch_test.WatchTests (methodName='runTest')
```

```
    Bases: unittest.case.TestCase
```

```
    test_unmarshal_with_float_object ()
```

```
    test_unmarshal_with_no_return_type ()
```

```
    test_watch_stream_twice ()
```

```
    test_watch_with_decode ()
```

```
    test_watch_with_exception ()
```

Module contents

4.1.2 Module contents

Contributing guidelines

How to become a contributor and submit your own code

Contributor License Agreements

We'd love to accept your patches! Before we can take them, we have to jump a couple of legal hurdles.

Please fill out either the individual or corporate Contributor License Agreement (CLA).

- If you are an individual writing original source code and you're sure you own the intellectual property, then you'll need to sign an [individual CLA](<https://identity.linuxfoundation.org/node/285/node/285/individual-signup>).
- If you work for a company that wants to allow you to contribute your work, then you'll need to sign a [corporate CLA](<https://identity.linuxfoundation.org/node/285/organization-signup>).

Follow either of the two links above to access the appropriate CLA and instructions for how to sign and return it. Once we receive it, we'll be able to accept your pull requests.

Contributing A Patch

1. Submit an issue describing your proposed change to the repo in question. 1. The [repo owners](OWNERS) will respond to your issue promptly. 1. If your proposed change is accepted, and you haven't already done so, sign a Contributor License Agreement (see details above). 1. Fork the desired repo, develop and test your code changes. 1. Submit a pull request.

Adding dependencies

If your patch depends on new packages, add those packages to requirements.txt and setup.py.

CHAPTER 6

Indices and tables

- `genindex`
- `modindex`
- `search`

Python Module Index

k

kubernetes, 638
kubernetes.client, 527
kubernetes.client.api_client, 522
kubernetes.client.apis, 297
kubernetes.client.apis.apis_api, 11
kubernetes.client.apis.apps_api, 12
kubernetes.client.apis.apps_v1beta1_api, 12
kubernetes.client.apis.authentication_api, 38
kubernetes.client.apis.authentication_v1beta1_api, 38
kubernetes.client.apis.authorization_api, 39
kubernetes.client.apis.authorization_v1beta1_api, 40
kubernetes.client.apis.autoscaling_api, 42
kubernetes.client.apis.autoscaling_v1_api, 42
kubernetes.client.apis.batch_api, 51
kubernetes.client.apis.batch_v1_api, 52
kubernetes.client.apis.batch_v2alpha1_api, 60
kubernetes.client.apis.certificates_api, 68
kubernetes.client.apis.core_api, 69
kubernetes.client.apis.core_v1_api, 69
kubernetes.client.apis.extensions_api, 209
kubernetes.client.apis.extensions_v1beta1_api, 210
kubernetes.client.apis.logs_api, 257
kubernetes.client.apis.policy_api, 258
kubernetes.client.apis.policy_v1beta1_api, 258
kubernetes.client.apis.rbac_authorization_api, 267
kubernetes.client.apis.rbac_authorization_v1alpha1_api, 268
kubernetes.client.apis.storage_api, 291
kubernetes.client.apis.storage_v1beta1_api, 291
kubernetes.client.apis.version_api, 296
kubernetes.client.configuration, 525
kubernetes.client.models, 522
kubernetes.client.models.runtime_raw_extension, 297
kubernetes.client.models.v1_attached_volume, 298
kubernetes.client.models.v1_aws_elastic_block_store_volume_source, 299
kubernetes.client.models.v1_azure_disk_volume_source, 300
kubernetes.client.models.v1_azure_file_volume_source, 301
kubernetes.client.models.v1_binding, 302
kubernetes.client.models.v1_capabilities, 303
kubernetes.client.models.v1_ceph_fs_volume_source, 304
kubernetes.client.models.v1_cinder_volume_source, 305
kubernetes.client.models.v1_component_condition, 306
kubernetes.client.models.v1_component_status, 307
kubernetes.client.models.v1_component_status_list, 308
kubernetes.client.models.v1_config_map, 309
kubernetes.client.models.v1_config_map_key_selector, 310
kubernetes.client.models.v1_config_map_list, 310
kubernetes.client.models.v1_config_map_volume_source, 311
kubernetes.client.models.v1_container,

| | | | |
|-----|--|-----|--|
| 312 | kubernetes.client.models.v1_container_image | 341 | kubernetes.client.models.v1_git_repo_volume_source, |
| 316 | kubernetes.client.models.v1_container_port | 342 | kubernetes.client.models.v1_glusterfs_volume_source, |
| 317 | kubernetes.client.models.v1_container_status | 343 | kubernetes.client.models.v1_handler, |
| 318 | kubernetes.client.models.v1_container_state_running, | 344 | kubernetes.client.models.v1_horizontal_pod_autoscaler, |
| 319 | kubernetes.client.models.v1_container_state_terminated, | 345 | kubernetes.client.models.v1_horizontal_pod_autoscaler, |
| 319 | kubernetes.client.models.v1_container_state_waiting, | 347 | kubernetes.client.models.v1_horizontal_pod_autoscaler, |
| 321 | kubernetes.client.models.v1_container_status, | 348 | kubernetes.client.models.v1_host_path_volume_source, |
| 321 | kubernetes.client.models.v1_cross_version_object_reference, | 349 | kubernetes.client.models.v1_http_get_action, |
| 323 | kubernetes.client.models.v1_daemon_endpoint, | 349 | kubernetes.client.models.v1_http_header, |
| 323 | kubernetes.client.models.v1_delete_options, | 350 | kubernetes.client.models.v1_iscsi_volume_source, |
| 324 | kubernetes.client.models.v1_downward_api_volume_file, | 351 | kubernetes.client.models.v1_job, |
| 325 | kubernetes.client.models.v1_downward_api_volume_source, | 353 | kubernetes.client.models.v1_job_condition, |
| 326 | kubernetes.client.models.v1_empty_dir_volume_source, | 354 | kubernetes.client.models.v1_job_list, |
| 327 | kubernetes.client.models.v1_endpoint_address, | 355 | kubernetes.client.models.v1_job_spec, |
| 328 | kubernetes.client.models.v1_endpoint_port, | 356 | kubernetes.client.models.v1_job_status, |
| 329 | kubernetes.client.models.v1_endpoint_subresource, | 358 | kubernetes.client.models.v1_key_to_path, |
| 330 | kubernetes.client.models.v1_endpoints, | 359 | kubernetes.client.models.v1_lifecycle, |
| 330 | kubernetes.client.models.v1_endpoints_list, | 360 | kubernetes.client.models.v1_limit_range, |
| 331 | kubernetes.client.models.v1_env_var, | 360 | kubernetes.client.models.v1_limit_range_item, |
| 332 | kubernetes.client.models.v1_env_var_source, | 361 | kubernetes.client.models.v1_limit_range_list, |
| 333 | kubernetes.client.models.v1_event, | 363 | kubernetes.client.models.v1_limit_range_spec, |
| 334 | kubernetes.client.models.v1_event_list, | 364 | kubernetes.client.models.v1_load_balancer_ingress, |
| 336 | kubernetes.client.models.v1_event_source, | 364 | kubernetes.client.models.v1_load_balancer_status, |
| 337 | kubernetes.client.models.v1_exec_action, | 365 | kubernetes.client.models.v1_local_object_reference, |
| 338 | kubernetes.client.models.v1_fc_volume_source, | 365 | kubernetes.client.models.v1_namespace, |
| 338 | kubernetes.client.models.v1_flex_volume_source, | 366 | kubernetes.client.models.v1_namespace_list, |
| 339 | kubernetes.client.models.v1_flocker_volume_source, | 367 | kubernetes.client.models.v1_namespace_spec, |
| 340 | kubernetes.client.models.v1_gce_persistent_disk_volume_source, | | |

| | |
|--|--|
| 368 | 403 |
| kubernetes.client.models.v1_namespace_status, | kubernetes.client.models.v1_pod_spec, |
| 368 | 404 |
| kubernetes.client.models.v1_nfs_volume_source, | kubernetes.client.models.v1_pod_status, |
| 369 | 409 |
| kubernetes.client.models.v1_node, | 370 |
| kubernetes.client.models.v1_node_address, | 411 |
| 371 | kubernetes.client.models.v1_pod_template, |
| kubernetes.client.models.v1_node_condition, | 412 |
| 371 | kubernetes.client.models.v1_pod_template_list, |
| kubernetes.client.models.v1_node_daemon_endpoint, | 413 |
| 373 | kubernetes.client.models.v1_pod_template_spec, |
| kubernetes.client.models.v1_node_list, | 413 |
| 373 | kubernetes.client.models.v1_preconditions, |
| kubernetes.client.models.v1_node_spec, | kubernetes.client.models.v1_probe, |
| 374 | 414 |
| kubernetes.client.models.v1_node_status, | kubernetes.client.models.v1_quobyte_volume_source, |
| 375 | 415 |
| kubernetes.client.models.v1_node_system_info, | kubernetes.client.models.v1_rbd_volume_source, |
| 377 | 416 |
| kubernetes.client.models.v1_object_field_selector, | kubernetes.client.models.v1_replication_controller, |
| 379 | 418 |
| kubernetes.client.models.v1_object_meta, | kubernetes.client.models.v1_replication_controller, |
| 379 | 419 |
| kubernetes.client.models.v1_object_reference, | kubernetes.client.models.v1_replication_controller, |
| 383 | 420 |
| kubernetes.client.models.v1_owner_reference, | kubernetes.client.models.v1_replication_controller, |
| 384 | 421 |
| kubernetes.client.models.v1_persistent_volume, | kubernetes.client.models.v1_replication_controller, |
| 386 | 422 |
| kubernetes.client.models.v1_persistent_volume_claim, | kubernetes.client.models.v1_resource_field_selector, |
| 387 | 424 |
| kubernetes.client.models.v1_persistent_volume_claim_spec, | kubernetes.client.models.v1_resource_quota, |
| 388 | 425 |
| kubernetes.client.models.v1_persistent_volume_claim_status, | kubernetes.client.models.v1_resource_quota_list, |
| 389 | 426 |
| kubernetes.client.models.v1_persistent_volume_claim_status, | kubernetes.client.models.v1_resource_quota_spec, |
| 390 | 427 |
| kubernetes.client.models.v1_persistent_volume_claim_status, | kubernetes.client.models.v1_resource_quota_status, |
| 391 | 427 |
| kubernetes.client.models.v1_persistent_volume_claim_status, | kubernetes.client.models.v1_resource_quota_status, |
| 392 | 428 |
| kubernetes.client.models.v1_persistent_volume_spec, | kubernetes.client.models.v1_scale, |
| 393 | 429 |
| kubernetes.client.models.v1_persistent_volume_status, | kubernetes.client.models.v1_scale_spec, |
| 398 | kubernetes.client.models.v1_scale_status, |
| kubernetes.client.models.v1_photon_persistent_volume_source, | kubernetes.client.models.v1_se_linux_options, |
| 399 | 431 |
| kubernetes.client.models.v1_pod, | kubernetes.client.models.v1_secret, |
| 400 | 432 |
| kubernetes.client.models.v1_pod_condition, | kubernetes.client.models.v1_secret_key_selector, |
| 401 | 433 |
| kubernetes.client.models.v1_pod_list, | kubernetes.client.models.v1_secret_list, |
| 402 | 434 |
| kubernetes.client.models.v1_pod_security_context, | kubernetes.client.models.v1_secret_volume_source, |
| | 435 |

| | |
|--|--|
| kubernetes.client.models.v1_security_context, | kubernetes.client.models.v1beta1_http_ingress_path, |
| 436 | 472 |
| kubernetes.client.models.v1_service, | kubernetes.client.models.v1beta1_http_ingress_rule, |
| 437 | |
| kubernetes.client.models.v1_service_account, | 473 |
| 439 | kubernetes.client.models.v1beta1_ingress, |
| kubernetes.client.models.v1_service_account_list, | 473 |
| 440 | kubernetes.client.models.v1beta1_ingress_backend, |
| kubernetes.client.models.v1_service_list, | 475 |
| 441 | kubernetes.client.models.v1beta1_ingress_list, |
| kubernetes.client.models.v1_service_port, | 475 |
| 442 | kubernetes.client.models.v1beta1_ingress_rule, |
| kubernetes.client.models.v1_service_spec, | 476 |
| 443 | kubernetes.client.models.v1beta1_ingress_spec, |
| kubernetes.client.models.v1_service_status, | 477 |
| 446 | kubernetes.client.models.v1beta1_ingress_status, |
| kubernetes.client.models.v1_tcp_socket_action, | 478 |
| 447 | kubernetes.client.models.v1beta1_ingress_tls, |
| kubernetes.client.models.v1_volume, | 447 478 |
| kubernetes.client.models.v1_volume_mount, | kubernetes.client.models.v1beta1_local_subject_access, |
| 452 | 479 |
| kubernetes.client.models.v1_vsphere_virtual_disk_volume, | kubernetes.client.models.v1beta1_network_policy, |
| 453 | 480 |
| kubernetes.client.models.v1alpha1_cluster_info, | kubernetes.client.models.v1beta1_network_policy_ingress, |
| 454 | 481 |
| kubernetes.client.models.v1alpha1_cluster_info_list, | kubernetes.client.models.v1beta1_network_policy_list, |
| 455 | 482 |
| kubernetes.client.models.v1alpha1_cluster_info_spec, | kubernetes.client.models.v1beta1_network_policy_peer, |
| 456 | 483 |
| kubernetes.client.models.v1alpha1_cluster_info_status, | kubernetes.client.models.v1beta1_network_policy_port, |
| 457 | 484 |
| kubernetes.client.models.v1alpha1_policy, | kubernetes.client.models.v1beta1_network_policy_spec, |
| 458 | 485 |
| kubernetes.client.models.v1alpha1_role, | kubernetes.client.models.v1beta1_non_resource_attribute, |
| 459 | 486 |
| kubernetes.client.models.v1alpha1_role_binding, | kubernetes.client.models.v1beta1_pod_disruption_budget, |
| 460 | 486 |
| kubernetes.client.models.v1alpha1_role_binding_list, | kubernetes.client.models.v1beta1_pod_disruption_budget, |
| 461 | 488 |
| kubernetes.client.models.v1alpha1_role_list, | kubernetes.client.models.v1beta1_pod_disruption_budget, |
| 462 | 489 |
| kubernetes.client.models.v1alpha1_role_ref, | kubernetes.client.models.v1beta1_pod_disruption_budget, |
| 463 | 489 |
| kubernetes.client.models.v1alpha1_subject, | kubernetes.client.models.v1beta1_replica_set, |
| 464 | 491 |
| kubernetes.client.models.v1beta1_daemon_set, | kubernetes.client.models.v1beta1_replica_set_condition, |
| 465 | 492 |
| kubernetes.client.models.v1beta1_daemon_set_list, | kubernetes.client.models.v1beta1_replica_set_list, |
| 466 | 493 |
| kubernetes.client.models.v1beta1_daemon_set_spec, | kubernetes.client.models.v1beta1_replica_set_spec, |
| 467 | 494 |
| kubernetes.client.models.v1beta1_daemon_set_status, | kubernetes.client.models.v1beta1_replica_set_status, |
| 469 | 495 |
| kubernetes.client.models.v1beta1_eviction, | kubernetes.client.models.v1beta1_resource_attribute, |
| 471 | 497 |

[kubernetes.client.models.v1beta1_self_subject_access_review](#),
[498](#) [kubernetes.test.test_authentication_api](#),
[kubernetes.client.models.v1beta1_self_subject_access_review_spec](#),
[500](#) [kubernetes.test.test_authentication_v1beta1_api](#),
[kubernetes.client.models.v1beta1_stateful_set](#), [535](#)
[500](#) [kubernetes.test.test_authorization_api](#),
[kubernetes.client.models.v1beta1_stateful_set_spec](#),
[501](#) [kubernetes.test.test_authorization_v1beta1_api](#),
[kubernetes.client.models.v1beta1_stateful_set_spec](#),
[502](#) [kubernetes.test.test_autoscaling_api](#),
[kubernetes.client.models.v1beta1_stateful_set_status](#),
[504](#) [kubernetes.test.test_autoscaling_v1_api](#),
[kubernetes.client.models.v1beta1_storage_class](#),[536](#)
[506](#) [kubernetes.test.test_batch_api](#),[537](#)
[kubernetes.client.models.v1beta1_storage_class_spec](#),[508](#) [kubernetes.test.test_batch_v1_api](#),[538](#)
[508](#) [kubernetes.test.test_batch_v2alpha1_api](#),
[kubernetes.client.models.v1beta1_subject_access_review](#),
[509](#) [kubernetes.test.test_certificates_api](#),
[kubernetes.client.models.v1beta1_subject_access_review_spec](#),
[510](#) [kubernetes.test.test_core_api](#),[540](#)
[kubernetes.client.models.v1beta1_subject_access_review_status](#),
[511](#) [kubernetes.test.test_core_v1_api](#),[540](#)
[511](#) [kubernetes.test.test_extensions_api](#),[552](#)
[kubernetes.client.models.v1beta1_token_request](#),[512](#) [kubernetes.test.test_extensions_v1beta1_api](#),
[512](#) [552](#)
[kubernetes.client.models.v1beta1_token_request_spec](#),
[513](#) [kubernetes.test.test_logs_api](#),[556](#)
[513](#) [kubernetes.test.test_policy_api](#),[556](#)
[kubernetes.client.models.v1beta1_token_request_status](#),
[514](#) [kubernetes.test.test_policy_v1beta1_api](#),
[514](#) [557](#)
[kubernetes.client.models.v1beta1_user_info](#),[515](#) [kubernetes.test.test_rbac_authorization_api](#),
[515](#) [558](#)
[kubernetes.client.models.v2alpha1_cron_job](#),[516](#) [kubernetes.test.test_rbac_authorization_v1alpha1_api](#),
[516](#) [558](#)
[kubernetes.client.models.v2alpha1_cron_job_spec](#),
[517](#) [kubernetes.test.test_runtime_raw_extension](#),
[517](#) [560](#)
[kubernetes.client.models.v2alpha1_cron_job_spec](#),
[518](#) [kubernetes.test.test_storage_api](#),[560](#)
[518](#) [kubernetes.test.test_storage_v1beta1_api](#),
[kubernetes.client.models.v2alpha1_cron_job_status](#),
[519](#) [560](#)
[519](#) [kubernetes.test.test_v1_attached_volume](#),
[kubernetes.client.models.v2alpha1_job_template_spec](#),
[520](#) [562](#)
[520](#) [kubernetes.test.test_v1_aws_elastic_block_store_volume_source](#),
[kubernetes.client.models.version_info](#),
[521](#) [563](#)
[521](#) [kubernetes.test.test_v1_azure_disk_volume_source](#),
[kubernetes.client.rest](#),[526](#) [563](#)
[kubernetes.config](#),[532](#) [kubernetes.test.test_v1_azure_file_volume_source](#),
[532](#) [563](#)
[kubernetes.config.config_exception](#),[527](#) [563](#)
[kubernetes.config.incluster_config](#),[528](#) [kubernetes.test.test_v1_binding](#),[564](#)
[528](#) [kubernetes.test.test_v1_capabilities](#),
[528](#) [564](#)
[kubernetes.config.kube_config](#),[529](#) [kubernetes.test.test_v1_ceph_fs_volume_source](#),
[529](#) [564](#)
[kubernetes.config.kube_config_test](#),[530](#) [564](#)
[kubernetes.test](#),[637](#) [kubernetes.test.test_v1_cinder_volume_source](#),
[637](#) [565](#)
[kubernetes.test.test_apis_api](#),[532](#) [565](#)
[kubernetes.test.test_apps_api](#),[532](#) [kubernetes.test.test_v1_component_condition](#),
[532](#) [565](#)
[kubernetes.test.test_apps_v1beta1_api](#),
[532](#) [565](#)

[kubernetes.test.test_v1_component_status](#)[kubernetes.test.test_v1_exec_action](#), 576
[565](#) [kubernetes.test.test_v1_fc_volume_source](#),
[kubernetes.test.test_v1_component_status_list](#), 576
[566](#) [kubernetes.test.test_v1_flex_volume_source](#),
[kubernetes.test.test_v1_config_map](#), 566 [577](#)
[kubernetes.test.test_v1_config_map_key_selector](#), 577
[566](#) [kubernetes.test.test_v1_flocker_volume_source](#),
[kubernetes.test.test_v1_config_map_list](#), [kubernetes.test.test_v1_gce_persistent_disk_volume](#),
[567](#) [578](#)
[kubernetes.test.test_v1_config_map_volume_source](#), [kubernetes.test.test_v1_git_repo_volume_source](#),
[567](#) [578](#)
[kubernetes.test.test_v1_container](#), 568 [kubernetes.test.test_v1_glusterfs_volume_source](#),
[kubernetes.test.test_v1_container_image](#), [578](#)
[568](#) [kubernetes.test.test_v1_handler](#), 579
[kubernetes.test.test_v1_container_port](#), [kubernetes.test.test_v1_horizontal_pod_autoscaler](#),
[568](#) [579](#)
[kubernetes.test.test_v1_container_state](#), [kubernetes.test.test_v1_horizontal_pod_autoscaler_list](#),
[569](#) [579](#)
[kubernetes.test.test_v1_container_state_waiting](#), [kubernetes.test.test_v1_horizontal_pod_autoscaler_spec](#),
[569](#) [580](#)
[kubernetes.test.test_v1_container_state_terminated](#), [kubernetes.test.test_v1_horizontal_pod_autoscaler_status](#),
[569](#) [580](#)
[kubernetes.test.test_v1_container_state_unhealthy](#), [kubernetes.test.test_v1_host_path_volume_source](#),
[570](#) [580](#)
[kubernetes.test.test_v1_container_status](#), [kubernetes.test.test_v1_http_get_action](#),
[570](#) [581](#)
[kubernetes.test.test_v1_cross_version_object_reference](#), [kubernetes.test.test_v1_http_header](#), 581
[570](#) [kubernetes.test.test_v1_iscsi_volume_source](#),
[kubernetes.test.test_v1_daemon_endpoint](#), [581](#)
[571](#) [kubernetes.test.test_v1_job](#), 582
[kubernetes.test.test_v1_delete_options](#), [kubernetes.test.test_v1_job_condition](#),
[571](#) [582](#)
[kubernetes.test.test_v1_downward_api_volume_file](#), [kubernetes.test.test_v1_job_list](#), 583
[571](#) [kubernetes.test.test_v1_job_spec](#), 583
[kubernetes.test.test_v1_downward_api_volume_source](#), [kubernetes.test.test_v1_job_status](#), 583
[572](#) [kubernetes.test.test_v1_key_to_path](#), 584
[kubernetes.test.test_v1_empty_dir_volume_source](#), [kubernetes.test.test_v1_lifecycle](#), 584
[572](#) [kubernetes.test.test_v1_limit_range](#), 584
[kubernetes.test.test_v1_endpoint_address](#), [kubernetes.test.test_v1_limit_range_item](#),
[573](#) [585](#)
[kubernetes.test.test_v1_endpoint_port](#), [kubernetes.test.test_v1_limit_range_list](#),
[573](#) [585](#)
[kubernetes.test.test_v1_endpoint_subset](#), [kubernetes.test.test_v1_limit_range_spec](#),
[573](#) [585](#)
[kubernetes.test.test_v1_endpoints](#), 574 [kubernetes.test.test_v1_load_balancer_ingress](#),
[kubernetes.test.test_v1_endpoints_list](#), [586](#)
[574](#) [kubernetes.test.test_v1_load_balancer_status](#),
[kubernetes.test.test_v1_env_var](#), 574 [586](#)
[kubernetes.test.test_v1_env_var_source](#), [kubernetes.test.test_v1_local_object_reference](#),
[575](#) [586](#)
[kubernetes.test.test_v1_event](#), 575 [kubernetes.test.test_v1_namespace](#), 587
[kubernetes.test.test_v1_event_list](#), 575 [kubernetes.test.test_v1_namespace_list](#),
[kubernetes.test.test_v1_event_source](#), [587](#)
[576](#) [kubernetes.test.test_v1_namespace_spec](#),

| | |
|--|---|
| 588 | kubernetes.test.test_v1_pod_template_list, |
| kubernetes.test.test_v1_namespace_status, | 599 |
| 588 | kubernetes.test.test_v1_pod_template_spec, |
| kubernetes.test.test_v1_nfs_volume_source, | 599 |
| 588 | kubernetes.test.test_v1_preconditions, |
| kubernetes.test.test_v1_node, | 589 600 |
| kubernetes.test.test_v1_node_address, | kubernetes.test.test_v1_probe, 600 |
| 589 | kubernetes.test.test_v1_quobyte_volume_source, |
| kubernetes.test.test_v1_node_condition, | 600 |
| 589 | kubernetes.test.test_v1_rbd_volume_source, |
| kubernetes.test.test_v1_node_daemon_endpoints, | 601 |
| 590 | kubernetes.test.test_v1_replication_controller, |
| kubernetes.test.test_v1_node_list, | 590 601 |
| kubernetes.test.test_v1_node_spec, | 590 kubernetes.test.test_v1_replication_controller_con |
| kubernetes.test.test_v1_node_status, | 591 601 |
| kubernetes.test.test_v1_node_system_info, | kubernetes.test.test_v1_replication_controller_list |
| 591 | 602 |
| kubernetes.test.test_v1_object_field_selector, | kubernetes.test.test_v1_replication_controller_spec |
| 591 | 602 |
| kubernetes.test.test_v1_object_meta, | 592 kubernetes.test.test_v1_replication_controller_stat |
| kubernetes.test.test_v1_object_reference, | 603 |
| 592 | kubernetes.test.test_v1_resource_field_selector, |
| kubernetes.test.test_v1_owner_reference, | 603 |
| 593 | kubernetes.test.test_v1_resource_quota, |
| kubernetes.test.test_v1_persistent_volume, | 603 |
| 593 | kubernetes.test.test_v1_resource_quota_list, |
| kubernetes.test.test_v1_persistent_volume_claim, | 604 |
| 593 | kubernetes.test.test_v1_resource_quota_spec, |
| kubernetes.test.test_v1_persistent_volume_claim_list, | 604 |
| 594 | kubernetes.test.test_v1_resource_quota_status, |
| kubernetes.test.test_v1_persistent_volume_claim_spec, | 604 |
| 594 | kubernetes.test.test_v1_resource_requirements, |
| kubernetes.test.test_v1_persistent_volume_claim_status, | 605 |
| 594 | kubernetes.test.test_v1_scale, 605 |
| kubernetes.test.test_v1_persistent_volume_claim_status, | kubernetes.test.test_v1_scale_spec, 605 |
| 595 | kubernetes.test.test_v1_scale_status, |
| kubernetes.test.test_v1_persistent_volume_list, | 606 |
| 595 | kubernetes.test.test_v1_se_linux_options, |
| kubernetes.test.test_v1_persistent_volume_spec, | 606 |
| 595 | kubernetes.test.test_v1_secret, 606 |
| kubernetes.test.test_v1_persistent_volume_spec, | kubernetes.test.test_v1_secret_key_selector, |
| 596 | 607 |
| kubernetes.test.test_v1_photon_persistent_volume_source, | kubernetes.test.test_v1_secret_list, 607 |
| 596 | kubernetes.test.test_v1_secret_volume_source, |
| kubernetes.test.test_v1_pod, | 596 608 |
| kubernetes.test.test_v1_pod_condition, | kubernetes.test.test_v1_security_context, |
| 597 | 608 |
| kubernetes.test.test_v1_pod_list, | 597 kubernetes.test.test_v1_service, 608 |
| kubernetes.test.test_v1_pod_security_context, | kubernetes.test.test_v1_service_account, |
| 598 | 609 |
| kubernetes.test.test_v1_pod_spec, | 598 kubernetes.test.test_v1_service_account_list, |
| kubernetes.test.test_v1_pod_status, | 598 609 |
| kubernetes.test.test_v1_pod_template, | kubernetes.test.test_v1_service_list, |
| 599 | 609 |

| | |
|--|---|
| kubernetes.test.test_v1_service_port, 610 | kubernetes.test.test_v1beta1_ingress_rule, 621 |
| kubernetes.test.test_v1_service_spec, 610 | kubernetes.test.test_v1beta1_ingress_spec, 621 |
| kubernetes.test.test_v1_service_status, 610 | kubernetes.test.test_v1beta1_ingress_status, 621 |
| kubernetes.test.test_v1_tcp_socket_action, 611 | kubernetes.test.test_v1beta1_ingress_tls, 622 |
| kubernetes.test.test_v1_volume, 611 | kubernetes.test.test_v1beta1_local_subject_access_review, 622 |
| kubernetes.test.test_v1_volume_mount, 611 | kubernetes.test.test_v1beta1_network_policy, 623 |
| kubernetes.test.test_v1_vsphere_virtual_disk_volume_source, 612 | kubernetes.test.test_v1beta1_network_policy_ingress_rule, 623 |
| kubernetes.test.test_v1alpha1_cluster_role, 612 | kubernetes.test.test_v1beta1_network_policy_list, 623 |
| kubernetes.test.test_v1alpha1_cluster_role_binding, 613 | kubernetes.test.test_v1beta1_network_policy_peer, 624 |
| kubernetes.test.test_v1alpha1_cluster_role_binding_list, 613 | kubernetes.test.test_v1beta1_network_policy_port, 624 |
| kubernetes.test.test_v1alpha1_cluster_role_list, 613 | kubernetes.test.test_v1beta1_network_policy_spec, 624 |
| kubernetes.test.test_v1alpha1_policy_rule, 614 | kubernetes.test.test_v1beta1_non_resource_attributes, 625 |
| kubernetes.test.test_v1alpha1_role, 614 | kubernetes.test.test_v1beta1_pod_disruption_budget, 625 |
| kubernetes.test.test_v1alpha1_role_binding, 615 | kubernetes.test.test_v1beta1_pod_disruption_budget_list, 625 |
| kubernetes.test.test_v1alpha1_role_binding_list, 615 | kubernetes.test.test_v1beta1_pod_disruption_budget_status, 626 |
| kubernetes.test.test_v1alpha1_role_list, 615 | kubernetes.test.test_v1beta1_pod_disruption_budget_status, 626 |
| kubernetes.test.test_v1alpha1_role_ref, 616 | kubernetes.test.test_v1beta1_replica_set, 626 |
| kubernetes.test.test_v1alpha1_subject, 616 | kubernetes.test.test_v1beta1_replica_set_condition, 627 |
| kubernetes.test.test_v1beta1_daemon_set, 616 | kubernetes.test.test_v1beta1_replica_set_list, 627 |
| kubernetes.test.test_v1beta1_daemon_set_kube_pod_template_ref, 617 | kubernetes.test.test_v1beta1_replica_set_spec, 628 |
| kubernetes.test.test_v1beta1_daemon_set_kube_pod_template_status_ref, 617 | kubernetes.test.test_v1beta1_replica_set_status, 628 |
| kubernetes.test.test_v1beta1_eviction, 618 | kubernetes.test.test_v1beta1_resource_attributes, 628 |
| kubernetes.test.test_v1beta1_http_ingress_backend, 619 | kubernetes.test.test_v1beta1_self_subject_access_review, 629 |
| kubernetes.test.test_v1beta1_http_ingress_backend_value, 619 | kubernetes.test.test_v1beta1_self_subject_access_review_value, 629 |
| kubernetes.test.test_v1beta1_ingress, 619 | kubernetes.test.test_v1beta1_stateful_set, 630 |
| kubernetes.test.test_v1beta1_ingress_backend, 620 | kubernetes.test.test_v1beta1_stateful_set_list, 630 |
| kubernetes.test.test_v1beta1_ingress_list, 620 | kubernetes.test.test_v1beta1_stateful_set_spec, 630 |

```

kubernetes.test.test_v1beta1_stateful_set_status,
    631
kubernetes.test.test_v1beta1_storage_class,
    631
kubernetes.test.test_v1beta1_storage_class_list,
    631
kubernetes.test.test_v1beta1_subject_access_review,
    632
kubernetes.test.test_v1beta1_subject_access_review_spec,
    632
kubernetes.test.test_v1beta1_subject_access_review_status,
    632
kubernetes.test.test_v1beta1_token_review,
    633
kubernetes.test.test_v1beta1_token_review_spec,
    633
kubernetes.test.test_v1beta1_token_review_status,
    634
kubernetes.test.test_v1beta1_user_info,
    634
kubernetes.test.test_v2alpha1_cron_job,
    634
kubernetes.test.test_v2alpha1_cron_job_list,
    635
kubernetes.test.test_v2alpha1_cron_job_spec,
    635
kubernetes.test.test_v2alpha1_cron_job_status,
    635
kubernetes.test.test_v2alpha1_job_template_spec,
    636
kubernetes.test.test_version_api, 636
kubernetes.test.test_version_info, 637
kubernetes.watch, 638
kubernetes.watch.watch, 637
kubernetes.watch.watch_test, 638

```


A

- access_modes (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 389
- access_modes (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute), 391
- access_modes (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 394
- active (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 358
- active (kubernetes.client.models.v2alpha1_cron_job_status.V2alpha1CronJobStatus attribute), 519
- active_deadline_seconds (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 356
- active_deadline_seconds (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 405
- add (kubernetes.client.models.v1_capabilities.V1Capabilities attribute), 303
- address (kubernetes.client.models.v1_node_address.V1NodeAddress attribute), 371
- addresses (kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset attribute), 330
- addresses (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 375
- affinity (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 405
- allocatable (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376
- allow_privilege_escalation (kubernetes.client.models.v1_security_context.V1SecurityContext attribute), 436
- allow_volume_expansion (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass attribute), 506
- allowed (kubernetes.client.models.v1beta1_subject_access_review_status.V1beta1SubjectAccessReviewStatus attribute), 511
- annotations (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 380
- api_group (kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef attribute), 463
- api_groups (kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule attribute), 458
- api_version (kubernetes.client.models.v1_binding.V1Binding attribute), 392
- api_version (kubernetes.client.models.v1_component_status.V1ComponentStatus attribute), 307
- api_version (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList attribute), 308
- api_version (kubernetes.client.models.v1_config_map.V1ConfigMap attribute), 309
- api_version (kubernetes.client.models.v1_config_map_list.V1ConfigMapList attribute), 311
- api_version (kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference attribute), 323
- api_version (kubernetes.client.models.v1_delete_options.V1DeleteOptions attribute), 324
- api_version (kubernetes.client.models.v1_endpoints.V1Endpoints attribute), 331
- api_version (kubernetes.client.models.v1_endpoints_list.V1EndpointsList attribute), 332
- api_version (kubernetes.client.models.v1_event.V1Event attribute), 334
- api_version (kubernetes.client.models.v1_event_list.V1EventList attribute), 336
- api_version (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute), 345
- api_version (kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList attribute), 346
- api_version (kubernetes.client.models.v1_job.V1Job attribute), 353
- api_version (kubernetes.client.models.v1_job_list.V1JobList attribute), 355
- api_version (kubernetes.client.models.v1_limit_range.V1LimitRange attribute), 361
- api_version (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList attribute), 363
- api_version (kubernetes.client.models.v1_namespace.V1Namespace attribute), 366

[api_version \(kubernetes.client.models.v1_namespace_list.V1NamespaceList attribute\), 367](#)
[api_version \(kubernetes.client.models.v1_node.V1Node attribute\), 370](#)
[api_version \(kubernetes.client.models.v1_node_list.V1NodeList attribute\), 373](#)
[api_version \(kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector attribute\), 379](#)
[api_version \(kubernetes.client.models.v1_object_reference.V1ObjectReference attribute\), 383](#)
[api_version \(kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute\), 385](#)
[api_version \(kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute\), 386](#)
[api_version \(kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute\), 387](#)
[api_version \(kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList attribute\), 388](#)
[api_version \(kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute\), 392](#)
[api_version \(kubernetes.client.models.v1_pod.V1Pod attribute\), 400](#)
[api_version \(kubernetes.client.models.v1_pod_list.V1PodList attribute\), 402](#)
[api_version \(kubernetes.client.models.v1_pod_template.V1PodTemplate attribute\), 411](#)
[api_version \(kubernetes.client.models.v1_pod_template_list.V1PodTemplateList attribute\), 412](#)
[api_version \(kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute\), 418](#)
[api_version \(kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList attribute\), 420](#)
[api_version \(kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute\), 425](#)
[api_version \(kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList attribute\), 426](#)
[api_version \(kubernetes.client.models.v1_scale.V1Scale attribute\), 429](#)
[api_version \(kubernetes.client.models.v1_secret.V1Secret attribute\), 432](#)
[api_version \(kubernetes.client.models.v1_secret_list.V1SecretList attribute\), 434](#)
[api_version \(kubernetes.client.models.v1_service.V1Service attribute\), 438](#)
[api_version \(kubernetes.client.models.v1_service_account.V1ServiceAccount attribute\), 439](#)
[api_version \(kubernetes.client.models.v1_service_account_list.V1ServiceAccountList attribute\), 440](#)
[api_version \(kubernetes.client.models.v1_service_list.V1ServiceList attribute\), 441](#)
[api_version \(kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole attribute\), 454](#)
[api_version \(kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute\), 455](#)
[api_version \(kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList attribute\), 456](#)
[api_version \(kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList attribute\), 457](#)
[api_version \(kubernetes.client.models.v1alpha1_role.V1alpha1Role attribute\), 459](#)
[api_version \(kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding attribute\), 460](#)
[api_version \(kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList attribute\), 461](#)
[api_version \(kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList attribute\), 462](#)
[api_version \(kubernetes.client.models.v1alpha1_subject.V1alpha1Subject attribute\), 464](#)
[api_version \(kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet attribute\), 465](#)
[api_version \(kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList attribute\), 466](#)
[api_version \(kubernetes.client.models.v1beta1_eviction.V1beta1Eviction attribute\), 471](#)
[api_version \(kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute\), 474](#)
[api_version \(kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList attribute\), 475](#)
[api_version \(kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute\), 479](#)
[api_version \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute\), 481](#)
[api_version \(kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList attribute\), 482](#)
[api_version \(kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute\), 487](#)
[api_version \(kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList attribute\), 488](#)
[api_version \(kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute\), 491](#)
[api_version \(kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList attribute\), 493](#)
[api_version \(kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute\), 499](#)
[api_version \(kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute\), 500](#)
[api_version \(kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList attribute\), 502](#)
[api_version \(kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass attribute\), 507](#)
[api_version \(kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList attribute\), 508](#)
[api_version \(kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview attribute\), 509](#)
[api_version \(kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview attribute\), 512](#)
[api_version \(kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob attribute\), 516](#)

[api_version \(kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList attribute\), 517](#)
[ApiClient \(class in kubernetes.client.api_client\), 522](#)
[ApiException, 526](#)
[ApisApi \(class in kubernetes.client.apis.apis_api\), 11](#)
[AppsApi \(class in kubernetes.client.apis.apps_api\), 12](#)
[AppsV1beta1Api \(class in kubernetes.client.apis.apps_v1beta1_api\), 12](#)
[architecture \(kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute\), 377](#)
[args \(kubernetes.client.models.v1_container.V1Container attribute\), 313](#)
[as_data\(\) \(kubernetes.config.kube_config.FileOrData method\), 529](#)
[as_file\(\) \(kubernetes.config.kube_config.FileOrData method\), 529](#)
[attribute_map \(kubernetes.client.models.runtime_raw_extension.V1RuntimeRawExtension attribute\), 297](#)
[attribute_map \(kubernetes.client.models.v1_attached_volume.V1AttachedVolume attribute\), 298](#)
[attribute_map \(kubernetes.client.models.v1_aws_elastic_block_store.V1AwsElasticBlockStoreVolumeSource attribute\), 299](#)
[attribute_map \(kubernetes.client.models.v1_azure_disk_volume.V1AzureDiskVolumeSource attribute\), 300](#)
[attribute_map \(kubernetes.client.models.v1_azure_file_volume.V1AzureFileVolumeSource attribute\), 301](#)
[attribute_map \(kubernetes.client.models.v1_binding.V1Binding attribute\), 302](#)
[attribute_map \(kubernetes.client.models.v1_capabilities.V1Capabilities attribute\), 303](#)
[attribute_map \(kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute\), 304](#)
[attribute_map \(kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource attribute\), 305](#)
[attribute_map \(kubernetes.client.models.v1_component_condition.V1ComponentCondition attribute\), 306](#)
[attribute_map \(kubernetes.client.models.v1_component_status.V1ComponentStatus attribute\), 307](#)
[attribute_map \(kubernetes.client.models.v1_component_status_list.V1ComponentStatusList attribute\), 308](#)
[attribute_map \(kubernetes.client.models.v1_config_map.V1ConfigMap attribute\), 309](#)
[attribute_map \(kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector attribute\), 310](#)
[attribute_map \(kubernetes.client.models.v1_config_map_list.V1ConfigMapList attribute\), 311](#)
[attribute_map \(kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource attribute\), 312](#)
[attribute_map \(kubernetes.client.models.v1_container.V1Container attribute\), 313](#)
[attribute_map \(kubernetes.client.models.v1_container_image.V1ContainerImage attribute\), 316](#)
[attribute_map \(kubernetes.client.models.v1_container_port.V1ContainerPort attribute\), 317](#)
[attribute_map \(kubernetes.client.models.v1_container_state.V1ContainerState attribute\), 318](#)
[attribute_map \(kubernetes.client.models.v1_container_state_running.V1ContainerStateRunning attribute\), 319](#)
[attribute_map \(kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated attribute\), 319](#)
[attribute_map \(kubernetes.client.models.v1_container_state_waiting.V1ContainerStateWaiting attribute\), 321](#)
[attribute_map \(kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute\), 322](#)
[attribute_map \(kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference attribute\), 323](#)
[attribute_map \(kubernetes.client.models.v1_daemon_endpoint.V1DaemonEndpoint attribute\), 324](#)
[attribute_map \(kubernetes.client.models.v1_delete_options.V1DeleteOptions attribute\), 324](#)
[attribute_map \(kubernetes.client.models.runtime_raw_extension.V1RuntimeRawExtension attribute\), 326](#)
[attribute_map \(kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile attribute\), 326](#)
[attribute_map \(kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource attribute\), 327](#)
[attribute_map \(kubernetes.client.models.v1_aws_elastic_block_store.V1AwsElasticBlockStoreVolumeSource attribute\), 327](#)
[attribute_map \(kubernetes.client.models.v1_azure_disk_volume.V1AzureDiskVolumeSource attribute\), 328](#)
[attribute_map \(kubernetes.client.models.v1_azure_file_volume.V1AzureFileVolumeSource attribute\), 329](#)
[attribute_map \(kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset attribute\), 330](#)
[attribute_map \(kubernetes.client.models.v1_endpoints.V1Endpoints attribute\), 331](#)
[attribute_map \(kubernetes.client.models.v1_endpoints_list.V1EndpointsList attribute\), 332](#)
[attribute_map \(kubernetes.client.models.v1_env_var.V1EnvVar attribute\), 333](#)
[attribute_map \(kubernetes.client.models.v1_env_var_source.V1EnvVarSource attribute\), 333](#)
[attribute_map \(kubernetes.client.models.v1_event.V1Event attribute\), 335](#)
[attribute_map \(kubernetes.client.models.v1_event_list.V1EventList attribute\), 336](#)
[attribute_map \(kubernetes.client.models.v1_event_source.V1EventSource attribute\), 337](#)
[attribute_map \(kubernetes.client.models.v1_exec_action.V1ExecAction attribute\), 338](#)
[attribute_map \(kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource attribute\), 338](#)
[attribute_map \(kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource attribute\), 339](#)
[attribute_map \(kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource attribute\), 340](#)
[attribute_map \(kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource attribute\), 341](#)
[attribute_map \(kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource attribute\), 342](#)

[attribute_map \(kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterfsVolumeSource attribute\), 343](#)
[attribute_map \(kubernetes.client.models.v1_handler.V1Handler attribute\), 344](#)
[attribute_map \(kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute\), 345](#)
[attribute_map \(kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute\), 346](#)
[attribute_map \(kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute\), 347](#)
[attribute_map \(kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute\), 348](#)
[attribute_map \(kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource attribute\), 349](#)
[attribute_map \(kubernetes.client.models.v1_http_get_action.V1HTTPGetAction attribute\), 350](#)
[attribute_map \(kubernetes.client.models.v1_http_header.V1HTTPHeader attribute\), 351](#)
[attribute_map \(kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute\), 351](#)
[attribute_map \(kubernetes.client.models.v1_job.V1Job attribute\), 353](#)
[attribute_map \(kubernetes.client.models.v1_job_condition.V1JobCondition attribute\), 354](#)
[attribute_map \(kubernetes.client.models.v1_job_list.V1JobList attribute\), 356](#)
[attribute_map \(kubernetes.client.models.v1_job_spec.V1JobSpec attribute\), 357](#)
[attribute_map \(kubernetes.client.models.v1_job_status.V1JobStatus attribute\), 358](#)
[attribute_map \(kubernetes.client.models.v1_key_to_path.V1KeyToPath attribute\), 359](#)
[attribute_map \(kubernetes.client.models.v1_lifecycle.V1Lifecycle attribute\), 360](#)
[attribute_map \(kubernetes.client.models.v1_limit_range.V1LimitRange attribute\), 361](#)
[attribute_map \(kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem attribute\), 362](#)
[attribute_map \(kubernetes.client.models.v1_limit_range_list.V1LimitRangeList attribute\), 363](#)
[attribute_map \(kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec attribute\), 364](#)
[attribute_map \(kubernetes.client.models.v1_load_balancer.V1LoadBalancer attribute\), 364](#)
[attribute_map \(kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus attribute\), 365](#)
[attribute_map \(kubernetes.client.models.v1_local_object_reference.V1LocalObjectReference attribute\), 365](#)
[attribute_map \(kubernetes.client.models.v1_namespace.V1Namespace attribute\), 366](#)
[attribute_map \(kubernetes.client.models.v1_namespace_list.V1NamespaceList attribute\), 367](#)
[attribute_map \(kubernetes.client.models.v1_namespace_spec.V1NamespaceSpec attribute\), 368](#)
[attribute_map \(kubernetes.client.models.v1_namespace_status.V1NamespaceStatus attribute\), 369](#)
[attribute_map \(kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource attribute\), 369](#)
[attribute_map \(kubernetes.client.models.v1_node.V1Node attribute\), 370](#)
[attribute_map \(kubernetes.client.models.v1_node_address.V1NodeAddress attribute\), 371](#)
[attribute_map \(kubernetes.client.models.v1_node_condition.V1NodeCondition attribute\), 372](#)
[attribute_map \(kubernetes.client.models.v1_node_status.V1NodeStatus attribute\), 373](#)
[attribute_map \(kubernetes.client.models.v1_node_list.V1NodeList attribute\), 373](#)
[attribute_map \(kubernetes.client.models.v1_node_spec.V1NodeSpec attribute\), 374](#)
[attribute_map \(kubernetes.client.models.v1_node_status.V1NodeStatus attribute\), 376](#)
[attribute_map \(kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute\), 377](#)
[attribute_map \(kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector attribute\), 379](#)
[attribute_map \(kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute\), 380](#)
[attribute_map \(kubernetes.client.models.v1_object_reference.V1ObjectReference attribute\), 383](#)
[attribute_map \(kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute\), 385](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute\), 386](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute\), 387](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList attribute\), 388](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute\), 390](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute\), 391](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_claim_volume_source.V1PersistentVolumeClaimVolumeSource attribute\), 392](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList attribute\), 392](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute\), 394](#)
[attribute_map \(kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus attribute\), 398](#)
[attribute_map \(kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource attribute\), 399](#)
[attribute_map \(kubernetes.client.models.v1_pod.V1Pod attribute\), 400](#)
[attribute_map \(kubernetes.client.models.v1_pod_condition.V1PodCondition attribute\), 401](#)
[attribute_map \(kubernetes.client.models.v1_pod_list.V1PodList attribute\), 402](#)

[attribute_map \(kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute\), 403](#)
[attribute_map \(kubernetes.client.models.v1_pod_spec.V1PodSpec attribute\), 405](#)
[attribute_map \(kubernetes.client.models.v1_pod_status.V1PodStatus attribute\), 409](#)
[attribute_map \(kubernetes.client.models.v1_pod_template.V1PodTemplate attribute\), 411](#)
[attribute_map \(kubernetes.client.models.v1_pod_template_list.V1PodTemplateList attribute\), 412](#)
[attribute_map \(kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec attribute\), 413](#)
[attribute_map \(kubernetes.client.models.v1_preconditions.V1Preconditions attribute\), 413](#)
[attribute_map \(kubernetes.client.models.v1_probe.V1Probe attribute\), 414](#)
[attribute_map \(kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource attribute\), 415](#)
[attribute_map \(kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource attribute\), 417](#)
[attribute_map \(kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute\), 418](#)
[attribute_map \(kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition attribute\), 419](#)
[attribute_map \(kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList attribute\), 421](#)
[attribute_map \(kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec attribute\), 422](#)
[attribute_map \(kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus attribute\), 423](#)
[attribute_map \(kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector attribute\), 424](#)
[attribute_map \(kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute\), 425](#)
[attribute_map \(kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList attribute\), 426](#)
[attribute_map \(kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec attribute\), 427](#)
[attribute_map \(kubernetes.client.models.v1_resource_quota_status.V1ResourceQuotaStatus attribute\), 428](#)
[attribute_map \(kubernetes.client.models.v1_resource_requirement.V1ResourceRequirement attribute\), 428](#)
[attribute_map \(kubernetes.client.models.v1_scale.V1Scale attribute\), 429](#)
[attribute_map \(kubernetes.client.models.v1_scale_spec.V1ScaleSpec attribute\), 430](#)
[attribute_map \(kubernetes.client.models.v1_scale_status.V1ScaleStatus attribute\), 431](#)
[attribute_map \(kubernetes.client.models.v1_selinux_options.V1SELinuxOptions attribute\), 431](#)
[attribute_map \(kubernetes.client.models.v1_secret.V1Secret attribute\), 432](#)
[attribute_map \(kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector attribute\), 433](#)
[attribute_map \(kubernetes.client.models.v1_secret_list.V1SecretList attribute\), 434](#)
[attribute_map \(kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource attribute\), 435](#)
[attribute_map \(kubernetes.client.models.v1_security_context.V1SecurityContext attribute\), 436](#)
[attribute_map \(kubernetes.client.models.v1_service.V1Service attribute\), 438](#)
[attribute_map \(kubernetes.client.models.v1_service_account.V1ServiceAccount attribute\), 439](#)
[attribute_map \(kubernetes.client.models.v1_service_account_list.V1ServiceAccountList attribute\), 440](#)
[attribute_map \(kubernetes.client.models.v1_service_list.V1ServiceList attribute\), 441](#)
[attribute_map \(kubernetes.client.models.v1_service_port.V1ServicePort attribute\), 442](#)
[attribute_map \(kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute\), 444](#)
[attribute_map \(kubernetes.client.models.v1_service_status.V1ServiceStatus attribute\), 446](#)
[attribute_map \(kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction attribute\), 447](#)
[attribute_map \(kubernetes.client.models.v1_volume.V1Volume attribute\), 448](#)
[attribute_map \(kubernetes.client.models.v1_volume_mount.V1VolumeMount attribute\), 452](#)
[attribute_map \(kubernetes.client.models.v1_volume_spec.V1VolumeSpec attribute\), 453](#)
[attribute_map \(kubernetes.client.models.v1_volume_status.V1VolumeStatus attribute\), 454](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_cluster_role.V1alpha1ClusterRole attribute\), 455](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute\), 456](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_cluster_role_list.V1alpha1ClusterRoleList attribute\), 457](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_policy_rule.V1alpha1PolicyRule attribute\), 458](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_role.V1alpha1Role attribute\), 459](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_role_binding.V1alpha1RoleBinding attribute\), 460](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_role_binding_list.V1alpha1RoleBindingList attribute\), 462](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_role_list.V1alpha1RoleList attribute\), 463](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_role_ref.V1alpha1RoleRef attribute\), 464](#)
[attribute_map \(kubernetes.client.models.v1_alpha1_subject.V1alpha1Subject attribute\), 464](#)
[attribute_map \(kubernetes.client.models.v1_beta1_daemon_set.V1beta1DaemonSet attribute\), 465](#)
[attribute_map \(kubernetes.client.models.v1_beta1_daemon_set_list.V1beta1DaemonSetList attribute\), 467](#)

[attribute_map \(kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus attribute\), 468](#)
[attribute_map \(kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus attribute\), 469](#)
[attribute_map \(kubernetes.client.models.v1beta1_eviction.V1beta1Eviction attribute\), 471](#)
[attribute_map \(kubernetes.client.models.v1beta1_http_ingress_rule.V1beta1HTTPIngressRule attribute\), 472](#)
[attribute_map \(kubernetes.client.models.v1beta1_http_ingress_rule.V1beta1HTTPIngressRule attribute\), 473](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute\), 474](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_backend.V1beta1IngressBackend attribute\), 475](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList attribute\), 476](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_rule.V1beta1HTTPIngressRule attribute\), 476](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec attribute\), 477](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_status.V1beta1IngressStatus attribute\), 478](#)
[attribute_map \(kubernetes.client.models.v1beta1_ingress_tls.V1beta1IngressTLS attribute\), 478](#)
[attribute_map \(kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute\), 480](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute\), 481](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute\), 482](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicyList attribute\), 482](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicyList attribute\), 483](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicyList attribute\), 484](#)
[attribute_map \(kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicyList attribute\), 485](#)
[attribute_map \(kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes attribute\), 486](#)
[attribute_map \(kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute\), 487](#)
[attribute_map \(kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudgetList attribute\), 488](#)
[attribute_map \(kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudgetSpec attribute\), 489](#)
[attribute_map \(kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudgetSpec attribute\), 490](#)
[attribute_map \(kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute\), 491](#)
[attribute_map \(kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute\), 492](#)
[attribute_map \(kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSetList attribute\), 494](#)

AuthenticationV1beta1Api (class in kubernetes.client.apis.authentication_v1beta1_api), 38

AuthorizationApi (class in kubernetes.client.apis.authorization_api), 39

AuthorizationV1beta1Api (class in kubernetes.client.apis.authorization_v1beta1_api), 40

automount_service_account_token (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 405

automount_service_account_token (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 439

AutoscalingApi (class in kubernetes.client.apis.autoscaling_api), 42

AutoscalingV1Api (class in kubernetes.client.apis.autoscaling_v1_api), 42

available_replicas (kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus attribute), 423

available_replicas (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus attribute), 496

aws_elastic_block_store (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 394

aws_elastic_block_store (kubernetes.client.models.v1_volume.V1Volume attribute), 448

azure_disk (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 395

azure_disk (kubernetes.client.models.v1_volume.V1Volume attribute), 448

azure_file (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 395

azure_file (kubernetes.client.models.v1_volume.V1Volume attribute), 448

B

backend (kubernetes.client.models.v1beta1_http_ingress_path.V1beta1HTTPIngressPath attribute), 472

backend (kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec attribute), 477

backoff_limit (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 357

BaseTestCase (class in kubernetes.config.kube_config_test), 530

BatchApi (class in kubernetes.client.apis.batch_api), 51

BatchV1Api (class in kubernetes.client.apis.batch_v1_api), 52

BatchV2alpha1Api (class in kubernetes.client.apis.batch_v2alpha1_api), 60

block_owner_deletion (kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute), 385

boot_id (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 377

build_date (kubernetes.client.models.version_info.VersionInfo attribute), 521

C

caching_mode (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 300

call_api() (kubernetes.client.api_client.ApiClient method), 523

capabilities (kubernetes.client.models.v1_security_context.V1SecurityContext attribute), 436

capacity (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376

capacity (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute), 391

capacity (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 395

cephfs (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 448

cephfs (kubernetes.client.models.v1_volume.V1Volume attribute), 448

CertificateApi (class in kubernetes.client.apis.certificates_api), 68

chap_auth_discovery (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352

chap_auth_session (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352

cinder (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 448

cinder (kubernetes.client.models.v1_volume.V1Volume attribute), 448

claim_name (kubernetes.client.models.v1_persistent_volume_claim_volume_source.V1PersistentVolumeClaimVolumeSource attribute), 392

claim_ref (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 395

cluster_ip (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 444

cluster_name (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 380

collision_count (kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus attribute), 469

collision_count (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus attribute), 505

command (kubernetes.client.models.v1_container.V1Container attribute), 313

`command` (kubernetes.client.models.v1_exec_action.V1ExecAction attribute), 338
`connect_delete_namespaced_service_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`compiler` (kubernetes.client.models.version_info.VersionInfo attribute), 521
`completion_time` (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 358
`completions` (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 357
`component` (kubernetes.client.models.v1_event_source.V1EventSource attribute), 337
`concurrency_policy` (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 518
`conditions` (kubernetes.client.models.v1_component_status.V1ComponentStatus attribute), 307
`conditions` (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 358
`conditions` (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376
`conditions` (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute), 391
`conditions` (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 409
`conditions` (kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus attribute), 423
`conditions` (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus attribute), 496
`config_map` (kubernetes.client.models.v1_volume.V1Volume attribute), 448
`config_map_key_ref` (kubernetes.client.models.v1_env_var_source.V1EnvVarSource attribute), 333
`config_source` (kubernetes.client.models.v1_node_spec.V1NodeSpec attribute), 374
`ConfigException`, 527
`ConfigNode` (class in kubernetes.config.kube_config), 529
`Configuration` (class in kubernetes.client.configuration), 525
`connect_delete_namespaced_pod_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 69
`connect_delete_namespaced_pod_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 70
`connect_delete_namespaced_pod_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 70
`connect_delete_namespaced_pod_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 70
`connect_delete_namespaced_service_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`connect_delete_namespaced_service_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`connect_delete_namespaced_service_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`connect_delete_namespaced_service_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`connect_delete_node_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 71
`connect_delete_node_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 72
`connect_delete_node_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 72
`connect_delete_node_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 72
`connect_get_namespaced_pod_attach()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 72
`connect_get_namespaced_pod_attach_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 72
`connect_get_namespaced_pod_exec()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 73
`connect_get_namespaced_pod_exec_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 73
`connect_get_namespaced_pod_portforward()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 73
`connect_get_namespaced_pod_portforward_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 73
`connect_get_namespaced_pod_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 74
`connect_get_namespaced_pod_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 74
`connect_get_namespaced_pod_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 74
`connect_get_namespaced_pod_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 74
`connect_get_namespaced_service_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 74

method), 74

`connect_get_namespaced_service_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 75

`connect_get_namespaced_service_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 75

`connect_get_namespaced_service_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 75

`connect_get_node_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 75

`connect_get_node_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 75

`connect_get_node_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 76

`connect_get_node_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 76

`connect_head_namespaced_pod_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 76

`connect_head_namespaced_pod_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 76

`connect_head_namespaced_pod_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 76

`connect_head_namespaced_pod_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 77

`connect_head_namespaced_service_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 77

`connect_head_namespaced_service_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 77

`connect_head_namespaced_service_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 77

`connect_head_namespaced_service_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 77

`connect_head_node_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 78

`connect_head_node_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 78

`connect_head_node_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 78

`connect_head_node_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 78

`connect_options_namespaced_pod_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 78

`connect_options_namespaced_pod_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 79

`connect_options_namespaced_pod_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 79

`connect_options_namespaced_pod_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 79

`connect_options_namespaced_service_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 79

`connect_options_namespaced_service_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 80

`connect_options_namespaced_service_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 80

`connect_options_namespaced_service_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 80

`connect_options_node_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 80

`connect_options_node_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_options_node_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_options_node_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_patch_namespaced_pod_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_patch_namespaced_pod_proxy_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_patch_namespaced_pod_proxy_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 81

`connect_patch_namespaced_pod_proxy_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 82

`connect_patch_namespaced_service_proxy()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 82

| | |
|---|--|
| method), 82 | method), 86 |
| connect_patch_namespaced_service_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 82 | connect_post_namespaced_service_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 86 |
| connect_patch_namespaced_service_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 82 | connect_post_namespaced_service_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 86 |
| connect_patch_namespaced_service_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 83 | connect_post_namespaced_service_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 87 |
| connect_patch_node_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 83 | connect_post_node_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 87 |
| connect_patch_node_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 83 | connect_post_node_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 87 |
| connect_patch_node_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 83 | connect_post_node_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 87 |
| connect_patch_node_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 83 | connect_post_node_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 87 |
| connect_post_namespaced_pod_attach() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 84 | connect_put_namespaced_pod_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 88 |
| connect_post_namespaced_pod_attach_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 84 | connect_put_namespaced_pod_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 88 |
| connect_post_namespaced_pod_exec() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 84 | connect_put_namespaced_pod_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 88 |
| connect_post_namespaced_pod_exec_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 84 | connect_put_namespaced_pod_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 88 |
| connect_post_namespaced_pod_portforward() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 85 | connect_put_namespaced_service_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 88 |
| connect_post_namespaced_pod_portforward_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 85 | connect_put_namespaced_service_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 89 |
| connect_post_namespaced_pod_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 85 | connect_put_namespaced_service_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 89 |
| connect_post_namespaced_pod_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 85 | connect_put_namespaced_service_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 89 |
| connect_post_namespaced_pod_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 85 | connect_put_node_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 89 |
| connect_post_namespaced_pod_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 86 | connect_put_node_proxy_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 89 |
| connect_post_namespaced_service_proxy() (kubernetes.client.apis.core_v1_api.CoreV1Api | connect_put_node_proxy_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api |

method), 90

connect_put_node_proxy_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 90

container_id (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated attribute), 319

container_id (kubernetes.client.models.v1_container_status.V1ContainerStatus attribute), 322

container_name (kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector attribute), 424

container_port (kubernetes.client.models.v1_container_port.V1ContainerPort attribute), 317

container_runtime_version (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378

container_statuses (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 409

containers (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 405

controller (kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute), 385

CoreApi (class in kubernetes.client.apis.core_api), 69

CoreV1Api (class in kubernetes.client.apis.core_v1_api), 69

count (kubernetes.client.models.v1_event.V1Event attribute), 335

create_cluster_role() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 268

create_cluster_role_binding() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 268

create_cluster_role_binding_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 268

create_cluster_role_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 268

create_namespace() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 90

create_namespace_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 90

create_namespaced_binding() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 90

create_namespaced_binding_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 90

create_namespaced_config_map() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 91

create_namespaced_config_map_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 91

create_namespaced_controller_revision_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 12

create_namespaced_cron_job() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 60

create_namespaced_cron_job_with_http_info() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 60

create_namespaced_daemon_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 210

create_namespaced_daemon_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 210

create_namespaced_deployment() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 13

create_namespaced_deployment() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 210

create_namespaced_deployment_rollback() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 13

create_namespaced_deployment_rollback() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 211

create_namespaced_deployment_rollback_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 13

create_namespaced_deployment_rollback_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 211

create_namespaced_deployment_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 13

create_namespaced_deployment_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 211

create_namespaced_endpoints() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 91

create_namespaced_endpoints_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 91

create_namespaced_event() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 91

| | |
|--|---|
| method), 91 | method), 93 |
| create_namespaced_event_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 92 | create_namespaced_pod_disruption_budget() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 258 |
| create_namespaced_horizontal_pod_autoscaler() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 42 | create_namespaced_pod_disruption_budget_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 259 |
| create_namespaced_horizontal_pod_autoscaler_with_http_info() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 43 | create_namespaced_pod_eviction() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 |
| create_namespaced_ingress() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 211 | create_namespaced_pod_eviction_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 |
| create_namespaced_ingress_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 211 | create_namespaced_pod_template() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 |
| create_namespaced_job() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 52 | create_namespaced_pod_template_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 94 |
| create_namespaced_job_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 52 | create_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 94 |
| create_namespaced_limit_range() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 92 | create_namespaced_replica_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 212 |
| create_namespaced_limit_range_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 92 | create_namespaced_replica_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 212 |
| create_namespaced_local_subject_access_review() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 40 | create_namespaced_replication_controller() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 94 |
| create_namespaced_local_subject_access_review_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 40 | create_namespaced_replication_controller_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 94 |
| create_namespaced_network_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 212 | create_namespaced_resource_quota() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 94 |
| create_namespaced_network_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 212 | create_namespaced_resource_quota_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 95 |
| create_namespaced_persistent_volume_claim() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 92 | create_namespaced_role() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 268 |
| create_namespaced_persistent_volume_claim_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 92 | create_namespaced_role_binding() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 269 |
| create_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 | create_namespaced_role_binding_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 269 |
| create_namespaced_pod_binding() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 | create_namespaced_role_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 269 |
| create_namespaced_pod_binding_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 | create_namespaced_secret() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 93 |

method), 95

create_namespaced_secret_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 95

create_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 95

create_namespaced_service_account() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 95

create_namespaced_service_account_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 95

create_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 96

create_namespaced_stateful_set() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 14

create_namespaced_stateful_set_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 14

create_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 96

create_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 96

create_persistent_volume() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 96

create_persistent_volume_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 96

create_pod_security_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 212

create_pod_security_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 213

create_self_subject_access_review() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 40

create_self_subject_access_review_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41

create_self_subject_rules_review() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41

create_self_subject_rules_review_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41

create_storage_class() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 291

create_storage_class_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 291

create_subject_access_review() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41

create_subject_access_review_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41

create_token_review() (kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api method), 39

create_token_review_with_http_info() (kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api method), 39

creation_timestamp (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 380

current_context (kubernetes.config.kube_config.KubeConfigLoader attribute), 529

current_cpu_utilization_percentage (kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus attribute), 348

current_healthy (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus attribute), 490

current_number_scheduled (kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus attribute), 469

current_replicas (kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus attribute), 348

current_replicas (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus attribute), 505

current_revision (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus attribute), 505

D

daemon_endpoints (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 316

data (kubernetes.client.models.v1_config_map.V1ConfigMap attribute), 309

data (kubernetes.client.models.v1_secret.V1Secret attribute), 432

dataset_name (kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource attribute), 340

dataset_uuid (kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource attribute), 341

debug (kubernetes.client.configuration.Configuration attribute), 525

[default \(kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem attribute\), 362](#)
[default_mode \(kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource attribute\), 312](#)
[default_mode \(kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource attribute\), 327](#)
[default_mode \(kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource attribute\), 435](#)
[default_request \(kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem attribute\), 362](#)
[DELETE\(\) \(kubernetes.client.rest.RESTClientObject method\), 526](#)
[delete_cluster_role\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 269](#)
[delete_cluster_role_binding\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 269](#)
[delete_cluster_role_binding_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 270](#)
[delete_cluster_role_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 270](#)
[delete_collection_cluster_role\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 270](#)
[delete_collection_cluster_role_binding\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 271](#)
[delete_collection_cluster_role_binding_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 272](#)
[delete_collection_cluster_role_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method\), 273](#)
[delete_collection_namespaced_config_map\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 96](#)
[delete_collection_namespaced_config_map_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 97](#)
[delete_collection_namespaced_controller_revision\(\) \(kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method\), 14](#)
[delete_collection_namespaced_controller_revision_with_http_info\(\) \(kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method\), 15](#)
[delete_collection_namespaced_cron_job\(\) \(kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method\), 60](#)
[delete_collection_namespaced_cron_job_with_http_info\(\) \(kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method\), 61](#)
[delete_collection_namespaced_daemon_set\(\) \(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method\), 213](#)
[delete_collection_namespaced_deployment\(\) \(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method\), 214](#)
[delete_collection_namespaced_deployment_with_http_info\(\) \(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method\), 215](#)
[delete_collection_namespaced_endpoints\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 98](#)
[delete_collection_namespaced_endpoints_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 99](#)
[delete_collection_namespaced_event\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 99](#)
[delete_collection_namespaced_event_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 100](#)
[delete_collection_namespaced_horizontal_pod_autoscaler\(\) \(kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method\), 43](#)
[delete_collection_namespaced_horizontal_pod_autoscaler_with_http_info\(\) \(kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method\), 44](#)
[delete_collection_namespaced_ingress\(\) \(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method\), 216](#)
[delete_collection_namespaced_ingress_with_http_info\(\) \(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method\), 216](#)
[delete_collection_namespaced_job\(\) \(kubernetes.client.apis.batch_v1_api.BatchV1Api method\), 52](#)
[delete_collection_namespaced_job_with_http_info\(\) \(kubernetes.client.apis.batch_v1_api.BatchV1Api method\), 53](#)
[delete_collection_namespaced_limit_range\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 101](#)
[delete_collection_namespaced_limit_range_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 101](#)

method), 102

delete_collection_namespaced_network_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 217

delete_collection_namespaced_network_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 218

delete_collection_namespaced_persistent_volume_claim() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 102

delete_collection_namespaced_persistent_volume_claim_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 103

delete_collection_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 104

delete_collection_namespaced_pod_disruption_budget() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 259

delete_collection_namespaced_pod_disruption_budget_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 260

delete_collection_namespaced_pod_template() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 105

delete_collection_namespaced_pod_template_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 105

delete_collection_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 106

delete_collection_namespaced_replica_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 219

delete_collection_namespaced_replica_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 219

delete_collection_namespaced_replication_controller() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 107

delete_collection_namespaced_replication_controller_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 107

delete_collection_namespaced_resource_quota() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 108

delete_collection_namespaced_resource_quota_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 109

delete_collection_namespaced_role() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 273

delete_collection_namespaced_role_binding() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 274

delete_collection_namespaced_role_binding_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 275

delete_collection_namespaced_role_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 275

delete_collection_namespaced_secret() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 110

delete_collection_namespaced_secret_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 110

delete_collection_namespaced_service_account() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 111

delete_collection_namespaced_service_account_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 112

delete_collection_namespaced_stateful_set() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 17

delete_collection_namespaced_stateful_set_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 18

delete_collection_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 113

delete_collection_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 113

delete_collection_persistent_volume() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 114

delete_collection_persistent_volume_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 115

delete_collection_pod_security_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 220

delete_collection_pod_security_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 221

delete_collection_storage_class() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 292

delete_collection_storage_class_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 292

delete_namespace() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 115

delete_namespace_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 115

method), 116

delete_namespaced_config_map() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 116

delete_namespaced_config_map_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 116

delete_namespaced_controller_revision() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 18

delete_namespaced_controller_revision_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 19

delete_namespaced_cron_job() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 62

delete_namespaced_cron_job_with_http_info() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 62

delete_namespaced_daemon_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 221

delete_namespaced_daemon_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 222

delete_namespaced_deployment() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 19

delete_namespaced_deployment() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 222

delete_namespaced_deployment_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 20

delete_namespaced_deployment_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 222

delete_namespaced_endpoints() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 117

delete_namespaced_endpoints_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 117

delete_namespaced_event() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 117

delete_namespaced_event_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 118

delete_namespaced_horizontal_pod_autoscaler() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 44

delete_namespaced_horizontal_pod_autoscaler_with_http_info() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 45

delete_namespaced_ingress() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 223

delete_namespaced_ingress_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 223

delete_namespaced_job() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 54

delete_namespaced_job_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 54

delete_namespaced_limit_range() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 118

delete_namespaced_limit_range_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 118

delete_namespaced_network_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 224

delete_namespaced_network_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 224

delete_namespaced_persistent_volume_claim() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 119

delete_namespaced_persistent_volume_claim_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 119

delete_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 120

delete_namespaced_pod_disruption_budget() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 260

delete_namespaced_pod_disruption_budget_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 261

delete_namespaced_pod_template() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 120

delete_namespaced_pod_template_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 120

delete_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 121

delete_namespaced_replica_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 224

delete_namespaced_replica_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 224

- method), 225
- delete_namespaced_replication_controller() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 121
- delete_namespaced_replication_controller_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 121
- delete_namespaced_resource_quota() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 122
- delete_namespaced_resource_quota_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 122
- delete_namespaced_role() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 276
- delete_namespaced_role_binding() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 276
- delete_namespaced_role_binding_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 277
- delete_namespaced_role_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 277
- delete_namespaced_secret() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 122
- delete_namespaced_secret_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 123
- delete_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 123
- delete_namespaced_service_account() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 123
- delete_namespaced_service_account_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 124
- delete_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 124
- delete_namespaced_stateful_set() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 20
- delete_namespaced_stateful_set_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 20
- delete_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 124
- delete_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 124
- delete_options (kubernetes.client.models.v1beta1_eviction.V1beta1Eviction attribute), 471
- delete_persistent_volume() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 125
- delete_persistent_volume_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 125
- delete_pod_security_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 225
- delete_pod_security_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 225
- delete_storage_class() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 293
- delete_storage_class_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 293
- delete_token_review() (kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api method), 380
- deserialize() (kubernetes.client.api_client.ApiClient method), 524
- desired_healthy (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus attribute), 490
- desired_number_scheduled (kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus attribute), 470
- desired_replicas (kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus attribute), 348
- device_path (kubernetes.client.models.v1_attached_volume.V1AttachedVolume attribute), 299
- directory (kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource attribute), 342
- disk_name (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 300
- disk_uri (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 301
- disrupted_pods (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus attribute), 490
- disruptions_allowed (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus attribute), 490
- divisor (kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector attribute), 424
- dns_policy (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 405

downward_api (kubernetes.client.models.v1_volume.V1VolumeSource attribute), 449

driver (kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource attribute), 339

drop (kubernetes.client.models.v1_capabilities.V1Capabilities attribute), 303

E

egress (kubernetes.client.models.v1beta1_network_policy_spec.V1Beta1NetworkPolicySpec attribute), 485

empty_dir (kubernetes.client.models.v1_volume.V1VolumeSource attribute), 449

endpoints (kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterfsVolumeSource attribute), 343

env (kubernetes.client.models.v1_container.V1Container attribute), 313

env_from (kubernetes.client.models.v1_container.V1Container attribute), 313

error (kubernetes.client.models.v1_component_condition.V1ComponentCondition attribute), 306

error (kubernetes.client.models.v1beta1_token_review_status.V1Beta1TokenReviewStatus attribute), 514

evaluation_error (kubernetes.client.models.v1beta1_subject_access_review_status.V1Beta1SubjectAccessReviewStatus attribute), 512

exit_code (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated attribute), 320

expect_exception() (kubernetes.config.kube_config_test.BaseTestCase method), 530

expected_pods (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1Beta1PodDisruptionBudgetStatus attribute), 490

ExtensionsApi (class in kubernetes.client.apis.extensions_api), 209

ExtensionsV1beta1Api (class in kubernetes.client.apis.extensions_v1beta1_api), 210

external_ips (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 444

external_id (kubernetes.client.models.v1_node_spec.V1NodeSpec attribute), 374

external_name (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 444

external_traffic_policy (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 444

extra (kubernetes.client.models.v1beta1_subject_access_review_spec.V1Beta1SubjectAccessReviewSpec attribute), 510

extra (kubernetes.client.models.v1beta1_user_info.V1Beta1UserInfo attribute), 515

F

failed (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 358

failed_jobs_history_limit (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 518

failure_threshold (kubernetes.client.models.v1_probe.V1Probe attribute), 414

FakeConfig (class in kubernetes.config.kube_config_test), 530

flex_volume (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 395

flex_volume (kubernetes.client.models.v1_volume.V1VolumeSource attribute), 449

flocker (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 396

flocker (kubernetes.client.models.v1_volume.V1VolumeSource attribute), 449

fs_group (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 403

fs_type (kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AwsElasticBlockStoreVolumeSource attribute), 299

fs_type (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 301

fs_type (kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource attribute), 305

fs_type (kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource attribute), 338

fs_type (kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource attribute), 340

fs_type (kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GcePersistentDiskVolumeSource attribute), 341

fs_type (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352
 fs_type (kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource attribute), 399
 fs_type (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource attribute), 417
 fs_type (kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource attribute), 453
 fully_labeled_replicas (kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus attribute), 423
 fully_labeled_replicas (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus attribute), 496
G
 gce_persistent_disk (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 396
 gce_persistent_disk (kubernetes.client.models.v1_volume.V1Volume attribute), 449
 generate_name (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 381
 generation (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 381
 GET() (kubernetes.client.rest.RESTClientObject method), 526
 get_api_group() (kubernetes.client.apis.apps_api.AppsApi method), 12
 get_api_group() (kubernetes.client.apis.authentication_api.AuthenticationApi method), 38
 get_api_group() (kubernetes.client.apis.authorization_api.AuthorizationApi method), 39
 get_api_group() (kubernetes.client.apis.autoscaling_api.AutoscalingApi method), 42
 get_api_group() (kubernetes.client.apis.batch_api.BatchApi method), 51
 get_api_group() (kubernetes.client.apis.certificates_api.CertificatesApi method), 68
 get_api_group() (kubernetes.client.apis.extensions_api.ExtensionsApi method), 210
 get_api_group() (kubernetes.client.apis.policy_api.PolicyApi method), 258
 get_api_group() (kubernetes.client.apis.rbac_authorization_api.RbacAuthorizationApi method), 267
 get_api_group() (kubernetes.client.apis.storage_api.StorageApi method), 291
 get_api_key_with_prefix() (kubernetes.client.configuration.Configuration method), 525
 get_api_resources() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 21
 get_api_resources() (kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api method), 39
 get_api_resources() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41
 get_api_resources() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 45
 get_api_resources() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 54
 get_api_resources() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 55

| | | |
|--|--|--|
| method), 63 | 69 | |
| get_api_resources() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 125 | get_api_versions_with_http_info() (kubernetes.client.apis.apis_api.ApisApi method), 12 | |
| get_api_resources() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 226 | get_api_versions_with_http_info() (kubernetes.client.apis.core_api.CoreApi method), 69 | |
| get_api_resources() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 261 | get_basic_auth_token() (kubernetes.client.configuration.Configuration method), 525 | |
| get_api_resources() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 278 | get_code() (kubernetes.client.apis.version_api.VersionApi method), 297 | |
| get_api_resources() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 294 | get_code_with_http_info() (kubernetes.client.apis.version_api.VersionApi method), 297 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 21 | get_file_content() (kubernetes.config.kube_config_test.TestFileOrData static method), 530 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.authentication_v1beta1_api.AuthenticationV1beta1Api method), 39 | get_return_type() (kubernetes.watch.watch.Watch method), 637 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.authorization_v1beta1_api.AuthorizationV1beta1Api method), 41 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 45 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 55 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 63 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 126 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 226 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 261 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 278 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_resources_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 294 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_versions() (kubernetes.client.apis.apis_api.ApisApi method), 11 | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| get_api_versions() (kubernetes.client.apis.core_api.CoreApi method), | get_status_value() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |

- attribute), 515
- ## H
- hard (kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec attribute), 427
- hard (kubernetes.client.models.v1_resource_quota_status.V1ResourceQuotaStatus attribute), 428
- HEAD() (kubernetes.client.rest.RESTClientObject method), 526
- health_check_node_port (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445
- host (kubernetes.client.models.v1_event_source.V1EventSource attribute), 337
- host (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction attribute), 350
- host (kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction attribute), 447
- host (kubernetes.client.models.v1beta1_ingress_rule.V1beta1IngressRule attribute), 476
- host_aliases (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- host_ip (kubernetes.client.models.v1_container_port.V1ContainerPort attribute), 317
- host_ip (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 409
- host_ipc (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- host_network (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- host_path (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 396
- host_path (kubernetes.client.models.v1_volume.V1Volume attribute), 450
- host_pid (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- host_port (kubernetes.client.models.v1_container_port.V1ContainerPort attribute), 317
- hostname (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress attribute), 328
- hostname (kubernetes.client.models.v1_load_balancer_ingress.V1LoadBalancerIngress attribute), 364
- hostname (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- hosts (kubernetes.client.models.v1beta1_ingress_tls.V1beta1IngressTLS attribute), 479
- http (kubernetes.client.models.v1beta1_ingress_rule.V1beta1IngressRule attribute), 477
- http_get (kubernetes.client.models.v1_handler.V1Handler attribute), 344
- http_get (kubernetes.client.models.v1_probe.V1Probe attribute), 414
- http_headers (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction attribute), 350
- image (kubernetes.client.models.v1_container.V1Container attribute), 314
- image (kubernetes.client.models.v1_container_status.V1ContainerStatus attribute), 322
- image (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource attribute), 417
- image_id (kubernetes.client.models.v1_container_status.V1ContainerStatus attribute), 322
- image_pull_policy (kubernetes.client.models.v1_container.V1Container attribute), 314
- image_pull_secrets (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- image_pull_secrets (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 439
- images (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376
- in_cluster_loader (class in kubernetes.config.incluster_config), 528
- in_cluster_loader (class in kubernetes.config.incluster_config_test), 528
- ingress (kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus attribute), 365
- ingress (kubernetes.client.models.v1beta1_network_policy_spec.V1beta1NetworkPolicySpec attribute), 485
- init_containers (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 406
- initial_delay_seconds (kubernetes.client.models.v1_probe.V1Probe attribute), 414
- initiators (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 381
- initiators (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352
- ip (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress attribute), 328
- ip (kubernetes.client.models.v1_load_balancer_ingress.V1LoadBalancerIngress attribute), 364
- ip_block (kubernetes.client.models.v1beta1_network_policy_peer.V1beta1NetworkPolicyPeer attribute), 483
- iqn (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352
- iscsi (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 396
- iscsi (kubernetes.client.models.v1_volume.V1Volume attribute), 450

attribute), 450
 iscsi_interface (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource, attribute), 352
 items (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList, attribute), 308
 items (kubernetes.client.models.v1_config_map_list.V1ConfigMapList, attribute), 311
 items (kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource, attribute), 312
 items (kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource, attribute), 327
 items (kubernetes.client.models.v1_endpoints_list.V1EndpointsList, attribute), 332
 items (kubernetes.client.models.v1_event_list.V1EventList, attribute), 336
 items (kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList, attribute), 346
 items (kubernetes.client.models.v1_job_list.V1JobList, attribute), 356
 items (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList, attribute), 363
 items (kubernetes.client.models.v1_namespace_list.V1NamespaceList, attribute), 367
 items (kubernetes.client.models.v1_node_list.V1NodeList, attribute), 373
 items (kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList, attribute), 389
 items (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList, attribute), 393
 items (kubernetes.client.models.v1_pod_list.V1PodList, attribute), 402
 items (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList, attribute), 412
 items (kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList, attribute), 421
 items (kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList, attribute), 426
 items (kubernetes.client.models.v1_secret_list.V1SecretList, attribute), 434
 items (kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource, attribute), 435
 items (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList, attribute), 440
 items (kubernetes.client.models.v1_service_list.V1ServiceList, attribute), 441
 items (kubernetes.client.models.v1alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList, attribute), 456
 items (kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList, attribute), 457
 items (kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList, attribute), 462
 items (kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList, attribute), 463
 items (kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList, attribute), 467
 items (kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList, attribute), 476
 items (kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList, attribute), 482
 items (kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList, attribute), 488
 items (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList, attribute), 494
 items (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList, attribute), 502
 items (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList, attribute), 508
 items (kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList, attribute), 517
 items (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler, attribute), 638
 items (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec, attribute), 518

J

items (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec, attribute), 518

K

kernel_version (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo, attribute), 378
 key (kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector, attribute), 310
 key (kubernetes.client.models.v1_key_to_path.V1KeyToPath, attribute), 359
 key (kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector, attribute), 433
 keyring (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource, attribute), 417
 kind (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource, attribute), 301
 kind (kubernetes.client.models.v1_binding.V1Binding, attribute), 302
 kind (kubernetes.client.models.v1_component_status.V1ComponentStatus, attribute), 307
 kind (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList, attribute), 308
 kind (kubernetes.client.models.v1_config_map.V1ConfigMap, attribute), 309
 kind (kubernetes.client.models.v1_config_map_list.V1ConfigMapList, attribute), 311
 kind (kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference, attribute), 323
 kind (kubernetes.client.models.v1_delete_options.V1DeleteOptions, attribute), 325
 kind (kubernetes.client.models.v1_endpoints.V1Endpoints, attribute), 331
 kind (kubernetes.client.models.v1_endpoints_list.V1EndpointsList, attribute), 332

| | |
|--|--|
| kind (kubernetes.client.models.v1_event.V1Event attribute), 335 | kind (kubernetes.client.models.v1_secret.V1Secret attribute), 432 |
| kind (kubernetes.client.models.v1_event_list.V1EventList attribute), 336 | kind (kubernetes.client.models.v1_secret_list.V1SecretList attribute), 434 |
| kind (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute), 345 | kind (kubernetes.client.models.v1_service.V1Service attribute), 438 |
| kind (kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList attribute), 346 | kind (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 439 |
| kind (kubernetes.client.models.v1_job.V1Job attribute), 353 | kind (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList attribute), 441 |
| kind (kubernetes.client.models.v1_job_list.V1JobList attribute), 356 | kind (kubernetes.client.models.v1_service_list.V1ServiceList attribute), 441 |
| kind (kubernetes.client.models.v1_limit_range.V1LimitRange attribute), 361 | kind (kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole attribute), 454 |
| kind (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList attribute), 363 | kind (kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute), 455 |
| kind (kubernetes.client.models.v1_namespace.V1Namespace attribute), 366 | kind (kubernetes.client.models.v1alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList attribute), 457 |
| kind (kubernetes.client.models.v1_namespace_list.V1NamespaceList attribute), 367 | kind (kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList attribute), 458 |
| kind (kubernetes.client.models.v1_node.V1Node attribute), 370 | kind (kubernetes.client.models.v1alpha1_role.V1alpha1Role attribute), 460 |
| kind (kubernetes.client.models.v1_node_list.V1NodeList attribute), 374 | kind (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding attribute), 461 |
| kind (kubernetes.client.models.v1_object_reference.V1ObjectReference attribute), 383 | kind (kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList attribute), 462 |
| kind (kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute), 385 | kind (kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList attribute), 463 |
| kind (kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute), 386 | kind (kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef attribute), 464 |
| kind (kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute), 387 | kind (kubernetes.client.models.v1alpha1_subject.V1alpha1Subject attribute), 464 |
| kind (kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList attribute), 389 | kind (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet attribute), 465 |
| kind (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList attribute), 393 | kind (kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList attribute), 467 |
| kind (kubernetes.client.models.v1_pod.V1Pod attribute), 400 | kind (kubernetes.client.models.v1beta1_eviction.V1beta1Eviction attribute), 472 |
| kind (kubernetes.client.models.v1_pod_list.V1PodList attribute), 402 | kind (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute), 474 |
| kind (kubernetes.client.models.v1_pod_template.V1PodTemplate attribute), 411 | kind (kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList attribute), 476 |
| kind (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList attribute), 412 | kind (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute), 480 |
| kind (kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute), 418 | kind (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute), 481 |
| kind (kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList attribute), 421 | kind (kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList attribute), 482 |
| kind (kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute), 425 | kind (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 487 |
| kind (kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList attribute), 426 | kind (kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList attribute), 488 |
| kind (kubernetes.client.models.v1_scale.V1Scale attribute), 429 | kind (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 491 |

kind (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList attribute), 494

kind (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 499

kind (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute), 501

kind (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList attribute), 502

kind (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass attribute), 507

kind (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList attribute), 508

kind (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 509

kind (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview attribute), 513

kind (kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob attribute), 516

kind (kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList attribute), 517

kube_proxy_version (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378

KubeConfigLoader (class in kubernetes.config.kube_config), 529

kubelet_endpoint (kubernetes.client.models.v1_node_daemon_endpoints.V1NodeDaemonEndpoints attribute), 373

kubelet_version (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378

kubernetes (module), 638

kubernetes.client (module), 527

kubernetes.client.api_client (module), 522

kubernetes.client.apis (module), 297

kubernetes.client.apis.apis_api (module), 11

kubernetes.client.apis.apps_api (module), 12

kubernetes.client.apis.apps_v1beta1_api (module), 12

kubernetes.client.apis.authentication_api (module), 38

kubernetes.client.apis.authentication_v1beta1_api (module), 38

kubernetes.client.apis.authorization_api (module), 39

kubernetes.client.apis.authorization_v1beta1_api (module), 40

kubernetes.client.apis.autoscaling_api (module), 42

kubernetes.client.apis.autoscaling_v1_api (module), 42

kubernetes.client.apis.batch_api (module), 51

kubernetes.client.apis.batch_v1_api (module), 52

kubernetes.client.apis.batch_v2alpha1_api (module), 60

kubernetes.client.apis.certificates_api (module), 68

kubernetes.client.apis.core_api (module), 69

kubernetes.client.apis.core_v1_api (module), 69

kubernetes.client.apis.extensions_api (module), 209

kubernetes.client.apis.extensions_v1beta1_api (module), 209

kubernetes.client.apis.logs_api (module), 257

kubernetes.client.apis.metrics_api (module), 258

kubernetes.client.apis.policy_v1beta1_api (module), 258

kubernetes.client.apis.rbac_authorization_api (module), 267

kubernetes.client.apis.rbac_authorization_v1alpha1_api (module), 268

kubernetes.client.apis.storage_api (module), 291

kubernetes.client.apis.storage_v1beta1_api (module), 291

kubernetes.client.apis.version_api (module), 296

kubernetes.client.models (module), 522

kubernetes.client.models.v1_attached_volume (module), 298

kubernetes.client.models.v1_aws_elastic_block_store_volume_source (module), 299

kubernetes.client.models.v1_azure_disk_volume_source (module), 300

kubernetes.client.models.v1_azure_file_volume_source (module), 301

kubernetes.client.models.v1_binding (module), 302

kubernetes.client.models.v1_capabilities (module), 303

kubernetes.client.models.v1_ceph_fs_volume_source (module), 304

kubernetes.client.models.v1_cinder_volume_source (module), 305

kubernetes.client.models.v1_component_condition (module), 306

kubernetes.client.models.v1_component_status (module), 307

kubernetes.client.models.v1_component_status_list (module), 308

kubernetes.client.models.v1_config_map (module), 309

kubernetes.client.models.v1_config_map_key_selector (module), 310

kubernetes.client.models.v1_config_map_list (module), 310

kubernetes.client.models.v1_config_map_volume_source (module), 311

kubernetes.client.models.v1_container (module), 312

kubernetes.client.models.v1_container_image (module), 316

kubernetes.client.models.v1_container_port (module), 317

kubernetes.client.models.v1_container_state (module), 318

kubernetes.client.models.v1_container_state_running (module), 319

kubernetes.client.models.v1_container_state_terminated (module), 319

[kubernetes.client.models.v1_container_state_waiting \(module\)](#), 321
[kubernetes.client.models.v1_container_status \(module\)](#), 321
[kubernetes.client.models.v1_cross_version_object_reference \(module\)](#), 323
[kubernetes.client.models.v1_daemon_endpoint \(module\)](#), 323
[kubernetes.client.models.v1_delete_options \(module\)](#), 324
[kubernetes.client.models.v1_downward_api_volume_file \(module\)](#), 325
[kubernetes.client.models.v1_downward_api_volume_source \(module\)](#), 326
[kubernetes.client.models.v1_empty_dir_volume_source \(module\)](#), 327
[kubernetes.client.models.v1_endpoint_address \(module\)](#), 328
[kubernetes.client.models.v1_endpoint_port \(module\)](#), 329
[kubernetes.client.models.v1_endpoint_subset \(module\)](#), 330
[kubernetes.client.models.v1_endpoints \(module\)](#), 330
[kubernetes.client.models.v1_endpoints_list \(module\)](#), 331
[kubernetes.client.models.v1_env_var \(module\)](#), 332
[kubernetes.client.models.v1_env_var_source \(module\)](#), 333
[kubernetes.client.models.v1_event \(module\)](#), 334
[kubernetes.client.models.v1_event_list \(module\)](#), 336
[kubernetes.client.models.v1_event_source \(module\)](#), 337
[kubernetes.client.models.v1_exec_action \(module\)](#), 338
[kubernetes.client.models.v1_fc_volume_source \(module\)](#), 338
[kubernetes.client.models.v1_flex_volume_source \(module\)](#), 339
[kubernetes.client.models.v1_flocker_volume_source \(module\)](#), 340
[kubernetes.client.models.v1_gce_persistent_disk_volume_source \(module\)](#), 341
[kubernetes.client.models.v1_git_repo_volume_source \(module\)](#), 342
[kubernetes.client.models.v1_glusterfs_volume_source \(module\)](#), 343
[kubernetes.client.models.v1_handler \(module\)](#), 344
[kubernetes.client.models.v1_horizontal_pod_autoscaler \(module\)](#), 344
[kubernetes.client.models.v1_horizontal_pod_autoscaler_list \(module\)](#), 346
[kubernetes.client.models.v1_horizontal_pod_autoscaler_spec \(module\)](#), 347
[kubernetes.client.models.v1_horizontal_pod_autoscaler_status \(module\)](#), 348
[kubernetes.client.models.v1_host_path_volume_source \(module\)](#), 349
[kubernetes.client.models.v1_http_get_action \(module\)](#), 349
[kubernetes.client.models.v1_http_header \(module\)](#), 350
[kubernetes.client.models.v1_iscsi_volume_source \(module\)](#), 351
[kubernetes.client.models.v1_job \(module\)](#), 353
[kubernetes.client.models.v1_job_condition \(module\)](#), 354
[kubernetes.client.models.v1_job_list \(module\)](#), 355
[kubernetes.client.models.v1_job_spec \(module\)](#), 356
[kubernetes.client.models.v1_job_status \(module\)](#), 358
[kubernetes.client.models.v1_key_to_path \(module\)](#), 359
[kubernetes.client.models.v1_lifecycle \(module\)](#), 360
[kubernetes.client.models.v1_limit_range \(module\)](#), 360
[kubernetes.client.models.v1_limit_range_item \(module\)](#), 361
[kubernetes.client.models.v1_limit_range_list \(module\)](#), 363
[kubernetes.client.models.v1_limit_range_spec \(module\)](#), 364
[kubernetes.client.models.v1_load_balancer_ingress \(module\)](#), 364
[kubernetes.client.models.v1_load_balancer_status \(module\)](#), 365
[kubernetes.client.models.v1_local_object_reference \(module\)](#), 365
[kubernetes.client.models.v1_namespace \(module\)](#), 366
[kubernetes.client.models.v1_namespace_list \(module\)](#), 367
[kubernetes.client.models.v1_namespace_spec \(module\)](#), 368
[kubernetes.client.models.v1_namespace_status \(module\)](#), 368
[kubernetes.client.models.v1_nfs_volume_source \(module\)](#), 369
[kubernetes.client.models.v1_node \(module\)](#), 370
[kubernetes.client.models.v1_node_address \(module\)](#), 371
[kubernetes.client.models.v1_node_condition \(module\)](#), 371
[kubernetes.client.models.v1_node_daemon_endpoints \(module\)](#), 373
[kubernetes.client.models.v1_node_list \(module\)](#), 373
[kubernetes.client.models.v1_node_spec \(module\)](#), 374
[kubernetes.client.models.v1_node_status \(module\)](#), 375
[kubernetes.client.models.v1_node_system_info \(module\)](#), 377
[kubernetes.client.models.v1_object_field_selector \(module\)](#), 379
[kubernetes.client.models.v1_object_meta \(module\)](#), 379
[kubernetes.client.models.v1_object_reference \(module\)](#), 383
[kubernetes.client.models.v1_owner_reference \(module\)](#), 384

[kubernetes.client.models.v1_persistent_volume \(module\), 386](#)
[kubernetes.client.models.v1_persistent_volume_claim \(module\), 387](#)
[kubernetes.client.models.v1_persistent_volume_claim_list \(module\), 388](#)
[kubernetes.client.models.v1_persistent_volume_claim_spec \(module\), 389](#)
[kubernetes.client.models.v1_persistent_volume_claim_status \(module\), 390](#)
[kubernetes.client.models.v1_persistent_volume_claim_volume_source \(module\), 391](#)
[kubernetes.client.models.v1_persistent_volume_list \(module\), 392](#)
[kubernetes.client.models.v1_persistent_volume_spec \(module\), 393](#)
[kubernetes.client.models.v1_persistent_volume_status \(module\), 398](#)
[kubernetes.client.models.v1_photon_persistent_disk_volume_source \(module\), 399](#)
[kubernetes.client.models.v1_pod \(module\), 400](#)
[kubernetes.client.models.v1_pod_condition \(module\), 401](#)
[kubernetes.client.models.v1_pod_list \(module\), 402](#)
[kubernetes.client.models.v1_pod_security_context \(module\), 403](#)
[kubernetes.client.models.v1_pod_spec \(module\), 404](#)
[kubernetes.client.models.v1_pod_status \(module\), 409](#)
[kubernetes.client.models.v1_pod_template \(module\), 411](#)
[kubernetes.client.models.v1_pod_template_list \(module\), 412](#)
[kubernetes.client.models.v1_pod_template_spec \(module\), 413](#)
[kubernetes.client.models.v1_preconditions \(module\), 413](#)
[kubernetes.client.models.v1_probe \(module\), 414](#)
[kubernetes.client.models.v1_quobyte_volume_source \(module\), 415](#)
[kubernetes.client.models.v1_rbd_volume_source \(module\), 416](#)
[kubernetes.client.models.v1_replication_controller \(module\), 418](#)
[kubernetes.client.models.v1_replication_controller_condition \(module\), 419](#)
[kubernetes.client.models.v1_replication_controller_list \(module\), 420](#)
[kubernetes.client.models.v1_replication_controller_spec \(module\), 421](#)
[kubernetes.client.models.v1_replication_controller_status \(module\), 422](#)
[kubernetes.client.models.v1_resource_field_selector \(module\), 424](#)
[kubernetes.client.models.v1_resource_quota \(module\), 425](#)
[kubernetes.client.models.v1_resource_quota_list \(module\), 426](#)
[kubernetes.client.models.v1_resource_quota_spec \(module\), 427](#)
[kubernetes.client.models.v1_resource_quota_status \(module\), 427](#)
[kubernetes.client.models.v1_resource_requirements \(module\), 428](#)
[kubernetes.client.models.v1_scale \(module\), 429](#)
[kubernetes.client.models.v1_scale_spec \(module\), 430](#)
[kubernetes.client.models.v1_scale_status \(module\), 430](#)
[kubernetes.client.models.v1_se_linux_options \(module\), 431](#)
[kubernetes.client.models.v1_secret \(module\), 432](#)
[kubernetes.client.models.v1_secret_key_selector \(module\), 433](#)
[kubernetes.client.models.v1_secret_list \(module\), 434](#)
[kubernetes.client.models.v1_secret_volume_source \(module\), 435](#)
[kubernetes.client.models.v1_security_context \(module\), 436](#)
[kubernetes.client.models.v1_service \(module\), 437](#)
[kubernetes.client.models.v1_service_account \(module\), 439](#)
[kubernetes.client.models.v1_service_account_list \(module\), 440](#)
[kubernetes.client.models.v1_service_list \(module\), 441](#)
[kubernetes.client.models.v1_service_port \(module\), 442](#)
[kubernetes.client.models.v1_service_spec \(module\), 443](#)
[kubernetes.client.models.v1_service_status \(module\), 446](#)
[kubernetes.client.models.v1_tcp_socket_action \(module\), 447](#)
[kubernetes.client.models.v1_volume \(module\), 447](#)
[kubernetes.client.models.v1_volume_mount \(module\), 452](#)
[kubernetes.client.models.v1_vsphere_virtual_disk_volume_source \(module\), 453](#)
[kubernetes.client.models.v1alpha1_cluster_role \(module\), 454](#)
[kubernetes.client.models.v1alpha1_cluster_role_binding \(module\), 455](#)
[kubernetes.client.models.v1alpha1_cluster_role_binding_list \(module\), 456](#)
[kubernetes.client.models.v1alpha1_cluster_role_list \(module\), 457](#)
[kubernetes.client.models.v1alpha1_policy_rule \(module\), 458](#)
[kubernetes.client.models.v1alpha1_role \(module\), 459](#)
[kubernetes.client.models.v1alpha1_role_binding \(module\), 460](#)
[kubernetes.client.models.v1alpha1_role_binding_list \(module\), 461](#)
[kubernetes.client.models.v1alpha1_role_list \(module\), 462](#)

| | |
|---|---|
| <code>kubernetes.client.models.v1alpha1_role_ref</code> (module), 463 | <code>kubernetes.client.models.v1beta1_replica_set</code> (module), 491 |
| <code>kubernetes.client.models.v1alpha1_subject</code> (module), 464 | <code>kubernetes.client.models.v1beta1_replica_set_condition</code> (module), 492 |
| <code>kubernetes.client.models.v1beta1_daemon_set</code> (module), 465 | <code>kubernetes.client.models.v1beta1_replica_set_list</code> (module), 493 |
| <code>kubernetes.client.models.v1beta1_daemon_set_list</code> (module), 466 | <code>kubernetes.client.models.v1beta1_replica_set_spec</code> (module), 494 |
| <code>kubernetes.client.models.v1beta1_daemon_set_spec</code> (module), 467 | <code>kubernetes.client.models.v1beta1_replica_set_status</code> (module), 495 |
| <code>kubernetes.client.models.v1beta1_daemon_set_status</code> (module), 469 | <code>kubernetes.client.models.v1beta1_resource_attributes</code> (module), 497 |
| <code>kubernetes.client.models.v1beta1_eviction</code> (module), 471 | <code>kubernetes.client.models.v1beta1_self_subject_access_review</code> (module), 498 |
| <code>kubernetes.client.models.v1beta1_http_ingress_path</code> (module), 472 | <code>kubernetes.client.models.v1beta1_self_subject_access_review_spec</code> (module), 500 |
| <code>kubernetes.client.models.v1beta1_http_ingress_rule_value</code> (module), 473 | <code>kubernetes.client.models.v1beta1_stateful_set</code> (module), 500 |
| <code>kubernetes.client.models.v1beta1_ingress</code> (module), 473 | <code>kubernetes.client.models.v1beta1_stateful_set_list</code> (module), 501 |
| <code>kubernetes.client.models.v1beta1_ingress_backend</code> (module), 475 | <code>kubernetes.client.models.v1beta1_stateful_set_spec</code> (module), 502 |
| <code>kubernetes.client.models.v1beta1_ingress_list</code> (module), 475 | <code>kubernetes.client.models.v1beta1_stateful_set_status</code> (module), 504 |
| <code>kubernetes.client.models.v1beta1_ingress_rule</code> (module), 476 | <code>kubernetes.client.models.v1beta1_storage_class</code> (module), 506 |
| <code>kubernetes.client.models.v1beta1_ingress_spec</code> (module), 477 | <code>kubernetes.client.models.v1beta1_storage_class_list</code> (module), 508 |
| <code>kubernetes.client.models.v1beta1_ingress_status</code> (module), 478 | <code>kubernetes.client.models.v1beta1_subject_access_review</code> (module), 509 |
| <code>kubernetes.client.models.v1beta1_ingress_tls</code> (module), 478 | <code>kubernetes.client.models.v1beta1_subject_access_review_spec</code> (module), 510 |
| <code>kubernetes.client.models.v1beta1_local_subject_access_review</code> (module), 479 | <code>kubernetes.client.models.v1beta1_subject_access_review_status</code> (module), 511 |
| <code>kubernetes.client.models.v1beta1_network_policy</code> (module), 480 | <code>kubernetes.client.models.v1beta1_token_review</code> (module), 512 |
| <code>kubernetes.client.models.v1beta1_network_policy_ingress_rule_value</code> (module), 481 | <code>kubernetes.client.models.v1beta1_token_review_spec</code> (module), 513 |
| <code>kubernetes.client.models.v1beta1_network_policy_list</code> (module), 482 | <code>kubernetes.client.models.v1beta1_token_review_status</code> (module), 514 |
| <code>kubernetes.client.models.v1beta1_network_policy_peer_selector</code> (module), 483 | <code>kubernetes.client.models.v1beta1_user_info</code> (module), 515 |
| <code>kubernetes.client.models.v1beta1_network_policy_port</code> (module), 484 | <code>kubernetes.client.models.v2alpha1_cron_job</code> (module), 516 |
| <code>kubernetes.client.models.v1beta1_network_policy_spec</code> (module), 485 | <code>kubernetes.client.models.v2alpha1_cron_job_list</code> (module), 517 |
| <code>kubernetes.client.models.v1beta1_non_resource_attributes</code> (module), 486 | <code>kubernetes.client.models.v2alpha1_cron_job_spec</code> (module), 518 |
| <code>kubernetes.client.models.v1beta1_pod_disruption_budget</code> (module), 486 | <code>kubernetes.client.models.v2alpha1_cron_job_status</code> (module), 519 |
| <code>kubernetes.client.models.v1beta1_pod_disruption_budget_list</code> (module), 488 | <code>kubernetes.client.models.v2alpha1_job_template_spec</code> (module), 520 |
| <code>kubernetes.client.models.v1beta1_pod_disruption_budget_spec</code> (module), 489 | <code>kubernetes.client.models.version_info</code> (module), 521 |
| <code>kubernetes.client.models.v1beta1_pod_disruption_budget_status</code> (module), 489 | <code>kubernetes.client.rest</code> (module), 526 |

- kubernetes.config (module), 532
- kubernetes.config.config_exception (module), 527
- kubernetes.config.incluster_config (module), 528
- kubernetes.config.incluster_config_test (module), 528
- kubernetes.config.kube_config (module), 529
- kubernetes.config.kube_config_test (module), 530
- kubernetes.test (module), 637
- kubernetes.test.test_apis_api (module), 532
- kubernetes.test.test_apps_api (module), 532
- kubernetes.test.test_apps_v1beta1_api (module), 532
- kubernetes.test.test_authentication_api (module), 534
- kubernetes.test.test_authentication_v1beta1_api (module), 535
- kubernetes.test.test_authorization_api (module), 535
- kubernetes.test.test_authorization_v1beta1_api (module), 536
- kubernetes.test.test_autoscaling_api (module), 536
- kubernetes.test.test_autoscaling_v1_api (module), 536
- kubernetes.test.test_batch_api (module), 537
- kubernetes.test.test_batch_v1_api (module), 538
- kubernetes.test.test_batch_v2alpha1_api (module), 539
- kubernetes.test.test_certificates_api (module), 539
- kubernetes.test.test_core_api (module), 540
- kubernetes.test.test_core_v1_api (module), 540
- kubernetes.test.test_extensions_api (module), 552
- kubernetes.test.test_extensions_v1beta1_api (module), 552
- kubernetes.test.test_logs_api (module), 556
- kubernetes.test.test_policy_api (module), 556
- kubernetes.test.test_policy_v1beta1_api (module), 557
- kubernetes.test.test_rbac_authorization_api (module), 558
- kubernetes.test.test_rbac_authorization_v1alpha1_api (module), 558
- kubernetes.test.test_runtime_raw_extension (module), 560
- kubernetes.test.test_storage_api (module), 560
- kubernetes.test.test_storage_v1beta1_api (module), 560
- kubernetes.test.test_v1_attached_volume (module), 562
- kubernetes.test.test_v1_aws_elastic_block_store_volume_source (module), 563
- kubernetes.test.test_v1_azure_disk_volume_source (module), 563
- kubernetes.test.test_v1_azure_file_volume_source (module), 563
- kubernetes.test.test_v1_binding (module), 564
- kubernetes.test.test_v1_capabilities (module), 564
- kubernetes.test.test_v1_ceph_fs_volume_source (module), 564
- kubernetes.test.test_v1_cinder_volume_source (module), 565
- kubernetes.test.test_v1_component_condition (module), 565
- kubernetes.test.test_v1_component_status (module), 565
- kubernetes.test.test_v1_component_status_list (module), 566
- kubernetes.test.test_v1_config_map (module), 566
- kubernetes.test.test_v1_config_map_key_selector (module), 566
- kubernetes.test.test_v1_config_map_list (module), 567
- kubernetes.test.test_v1_config_map_volume_source (module), 567
- kubernetes.test.test_v1_container (module), 568
- kubernetes.test.test_v1_container_image (module), 568
- kubernetes.test.test_v1_container_port (module), 568
- kubernetes.test.test_v1_container_state (module), 569
- kubernetes.test.test_v1_container_state_running (module), 569
- kubernetes.test.test_v1_container_state_terminated (module), 569
- kubernetes.test.test_v1_container_state_waiting (module), 570
- kubernetes.test.test_v1_container_status (module), 570
- kubernetes.test.test_v1_cross_version_object_reference (module), 570
- kubernetes.test.test_v1_daemon_endpoint (module), 571
- kubernetes.test.test_v1_delete_options (module), 571
- kubernetes.test.test_v1_downward_api_volume_file (module), 571
- kubernetes.test.test_v1_downward_api_volume_source (module), 572
- kubernetes.test.test_v1_empty_dir_volume_source (module), 572
- kubernetes.test.test_v1_endpoint_address (module), 573
- kubernetes.test.test_v1_endpoint_port (module), 573
- kubernetes.test.test_v1_endpoint_subset (module), 573
- kubernetes.test.test_v1_endpoints (module), 574
- kubernetes.test.test_v1_endpoints_list (module), 574
- kubernetes.test.test_v1_env_var (module), 574
- kubernetes.test.test_v1_env_var_source (module), 575
- kubernetes.test.test_v1_event (module), 575
- kubernetes.test.test_v1_event_list (module), 575
- kubernetes.test.test_v1_event_source (module), 576
- kubernetes.test.test_v1_exec_action (module), 576
- kubernetes.test.test_v1_fc_volume_source (module), 576
- kubernetes.test.test_v1_flex_volume_source (module), 577
- kubernetes.test.test_v1_flocker_volume_source (module), 577
- kubernetes.test.test_v1_gce_persistent_disk_volume_source (module), 578
- kubernetes.test.test_v1_git_repo_volume_source (module), 578
- kubernetes.test.test_v1_glusterfs_volume_source (module), 578
- kubernetes.test.test_v1_handler (module), 579
- kubernetes.test.test_v1_horizontal_pod_autoscaler (module), 579

- kubernetes.test.test_v1_horizontal_pod_autoscaler_list (module), 579
- kubernetes.test.test_v1_horizontal_pod_autoscaler_spec (module), 580
- kubernetes.test.test_v1_horizontal_pod_autoscaler_status (module), 580
- kubernetes.test.test_v1_host_path_volume_source (module), 580
- kubernetes.test.test_v1_http_get_action (module), 581
- kubernetes.test.test_v1_http_header (module), 581
- kubernetes.test.test_v1_iscsi_volume_source (module), 581
- kubernetes.test.test_v1_job (module), 582
- kubernetes.test.test_v1_job_condition (module), 582
- kubernetes.test.test_v1_job_list (module), 583
- kubernetes.test.test_v1_job_spec (module), 583
- kubernetes.test.test_v1_job_status (module), 583
- kubernetes.test.test_v1_key_to_path (module), 584
- kubernetes.test.test_v1_lifecycle (module), 584
- kubernetes.test.test_v1_limit_range (module), 584
- kubernetes.test.test_v1_limit_range_item (module), 585
- kubernetes.test.test_v1_limit_range_list (module), 585
- kubernetes.test.test_v1_limit_range_spec (module), 585
- kubernetes.test.test_v1_load_balancer_ingress (module), 586
- kubernetes.test.test_v1_load_balancer_status (module), 586
- kubernetes.test.test_v1_local_object_reference (module), 586
- kubernetes.test.test_v1_namespace (module), 587
- kubernetes.test.test_v1_namespace_list (module), 587
- kubernetes.test.test_v1_namespace_spec (module), 588
- kubernetes.test.test_v1_namespace_status (module), 588
- kubernetes.test.test_v1_nfs_volume_source (module), 588
- kubernetes.test.test_v1_node (module), 589
- kubernetes.test.test_v1_node_address (module), 589
- kubernetes.test.test_v1_node_condition (module), 589
- kubernetes.test.test_v1_node_daemon_endpoints (module), 590
- kubernetes.test.test_v1_node_list (module), 590
- kubernetes.test.test_v1_node_spec (module), 590
- kubernetes.test.test_v1_node_status (module), 591
- kubernetes.test.test_v1_node_system_info (module), 591
- kubernetes.test.test_v1_object_field_selector (module), 591
- kubernetes.test.test_v1_object_meta (module), 592
- kubernetes.test.test_v1_object_reference (module), 592
- kubernetes.test.test_v1_owner_reference (module), 593
- kubernetes.test.test_v1_persistent_volume (module), 593
- kubernetes.test.test_v1_persistent_volume_claim (module), 593
- kubernetes.test.test_v1_persistent_volume_claim_list (module), 594
- kubernetes.test.test_v1_persistent_volume_claim_spec (module), 594
- kubernetes.test.test_v1_persistent_volume_claim_status (module), 594
- kubernetes.test.test_v1_persistent_volume_claim_volume_source (module), 595
- kubernetes.test.test_v1_persistent_volume_list (module), 595
- kubernetes.test.test_v1_persistent_volume_spec (module), 595
- kubernetes.test.test_v1_persistent_volume_status (module), 596
- kubernetes.test.test_v1_photon_persistent_disk_volume_source (module), 596
- kubernetes.test.test_v1_pod (module), 596
- kubernetes.test.test_v1_pod_condition (module), 597
- kubernetes.test.test_v1_pod_list (module), 597
- kubernetes.test.test_v1_pod_security_context (module), 598
- kubernetes.test.test_v1_pod_spec (module), 598
- kubernetes.test.test_v1_pod_status (module), 598
- kubernetes.test.test_v1_pod_template (module), 599
- kubernetes.test.test_v1_pod_template_list (module), 599
- kubernetes.test.test_v1_pod_template_spec (module), 599
- kubernetes.test.test_v1_preconditions (module), 600
- kubernetes.test.test_v1_probe (module), 600
- kubernetes.test.test_v1_quobyte_volume_source (module), 600
- kubernetes.test.test_v1_rbd_volume_source (module), 601
- kubernetes.test.test_v1_replication_controller (module), 601
- kubernetes.test.test_v1_replication_controller_condition (module), 601
- kubernetes.test.test_v1_replication_controller_list (module), 602
- kubernetes.test.test_v1_replication_controller_spec (module), 602
- kubernetes.test.test_v1_replication_controller_status (module), 603
- kubernetes.test.test_v1_resource_field_selector (module), 603
- kubernetes.test.test_v1_resource_quota (module), 603
- kubernetes.test.test_v1_resource_quota_list (module), 604
- kubernetes.test.test_v1_resource_quota_spec (module), 604
- kubernetes.test.test_v1_resource_quota_status (module), 604
- kubernetes.test.test_v1_resource_requirements (module), 605
- kubernetes.test.test_v1_scale (module), 605
- kubernetes.test.test_v1_scale_spec (module), 605

- kubernetes.test.test_v1_scale_status (module), 606
- kubernetes.test.test_v1_se_linux_options (module), 606
- kubernetes.test.test_v1_secret (module), 606
- kubernetes.test.test_v1_secret_key_selector (module), 607
- kubernetes.test.test_v1_secret_list (module), 607
- kubernetes.test.test_v1_secret_volume_source (module), 608
- kubernetes.test.test_v1_security_context (module), 608
- kubernetes.test.test_v1_service (module), 608
- kubernetes.test.test_v1_service_account (module), 609
- kubernetes.test.test_v1_service_account_list (module), 609
- kubernetes.test.test_v1_service_list (module), 609
- kubernetes.test.test_v1_service_port (module), 610
- kubernetes.test.test_v1_service_spec (module), 610
- kubernetes.test.test_v1_service_status (module), 610
- kubernetes.test.test_v1_tcp_socket_action (module), 611
- kubernetes.test.test_v1_volume (module), 611
- kubernetes.test.test_v1_volume_mount (module), 611
- kubernetes.test.test_v1_vsphere_virtual_disk_volume_source (module), 612
- kubernetes.test.test_v1alpha1_cluster_role (module), 612
- kubernetes.test.test_v1alpha1_cluster_role_binding (module), 613
- kubernetes.test.test_v1alpha1_cluster_role_binding_list (module), 613
- kubernetes.test.test_v1alpha1_cluster_role_list (module), 613
- kubernetes.test.test_v1alpha1_policy_rule (module), 614
- kubernetes.test.test_v1alpha1_role (module), 614
- kubernetes.test.test_v1alpha1_role_binding (module), 615
- kubernetes.test.test_v1alpha1_role_binding_list (module), 615
- kubernetes.test.test_v1alpha1_role_list (module), 615
- kubernetes.test.test_v1alpha1_role_ref (module), 616
- kubernetes.test.test_v1alpha1_subject (module), 616
- kubernetes.test.test_v1beta1_daemon_set (module), 616
- kubernetes.test.test_v1beta1_daemon_set_list (module), 617
- kubernetes.test.test_v1beta1_daemon_set_spec (module), 617
- kubernetes.test.test_v1beta1_daemon_set_status (module), 617
- kubernetes.test.test_v1beta1_eviction (module), 618
- kubernetes.test.test_v1beta1_http_ingress_path (module), 619
- kubernetes.test.test_v1beta1_http_ingress_rule_value (module), 619
- kubernetes.test.test_v1beta1_ingress (module), 619
- kubernetes.test.test_v1beta1_ingress_backend (module), 620
- kubernetes.test.test_v1beta1_ingress_list (module), 620
- kubernetes.test.test_v1beta1_ingress_rule (module), 621
- kubernetes.test.test_v1beta1_ingress_spec (module), 621
- kubernetes.test.test_v1beta1_ingress_status (module), 621
- kubernetes.test.test_v1beta1_ingress_tls (module), 622
- kubernetes.test.test_v1beta1_local_subject_access_review (module), 622
- kubernetes.test.test_v1beta1_network_policy (module), 623
- kubernetes.test.test_v1beta1_network_policy_ingress_rule (module), 623
- kubernetes.test.test_v1beta1_network_policy_list (module), 623
- kubernetes.test.test_v1beta1_network_policy_peer (module), 624
- kubernetes.test.test_v1beta1_network_policy_port (module), 624
- kubernetes.test.test_v1beta1_network_policy_spec (module), 624
- kubernetes.test.test_v1beta1_non_resource_attributes (module), 625
- kubernetes.test.test_v1beta1_pod_disruption_budget (module), 625
- kubernetes.test.test_v1beta1_pod_disruption_budget_list (module), 625
- kubernetes.test.test_v1beta1_pod_disruption_budget_spec (module), 626
- kubernetes.test.test_v1beta1_pod_disruption_budget_status (module), 626
- kubernetes.test.test_v1beta1_replica_set (module), 626
- kubernetes.test.test_v1beta1_replica_set_condition (module), 627
- kubernetes.test.test_v1beta1_replica_set_list (module), 627
- kubernetes.test.test_v1beta1_replica_set_spec (module), 628
- kubernetes.test.test_v1beta1_replica_set_status (module), 628
- kubernetes.test.test_v1beta1_resource_attributes (module), 628
- kubernetes.test.test_v1beta1_self_subject_access_review (module), 629
- kubernetes.test.test_v1beta1_self_subject_access_review_spec (module), 629
- kubernetes.test.test_v1beta1_stateful_set (module), 630
- kubernetes.test.test_v1beta1_stateful_set_list (module), 630
- kubernetes.test.test_v1beta1_stateful_set_spec (module), 630
- kubernetes.test.test_v1beta1_stateful_set_status (module), 631
- kubernetes.test.test_v1beta1_storage_class (module), 631
- kubernetes.test.test_v1beta1_storage_class_list (module), 631

[kubernetes.test.test_v1beta1_subject_access_review \(module\), 632](#)
[kubernetes.test.test_v1beta1_subject_access_review_spec \(module\), 632](#)
[kubernetes.test.test_v1beta1_subject_access_review_status \(module\), 632](#)
[kubernetes.test.test_v1beta1_token_review \(module\), 633](#)
[kubernetes.test.test_v1beta1_token_review_spec \(module\), 633](#)
[kubernetes.test.test_v1beta1_token_review_status \(module\), 634](#)
[kubernetes.test.test_v1beta1_user_info \(module\), 634](#)
[kubernetes.test.test_v2alpha1_cron_job \(module\), 634](#)
[kubernetes.test.test_v2alpha1_cron_job_list \(module\), 635](#)
[kubernetes.test.test_v2alpha1_cron_job_spec \(module\), 635](#)
[kubernetes.test.test_v2alpha1_cron_job_status \(module\), 635](#)
[kubernetes.test.test_v2alpha1_job_template_spec \(module\), 636](#)
[kubernetes.test.test_version_api \(module\), 636](#)
[kubernetes.test.test_version_info \(module\), 637](#)
[kubernetes.watch \(module\), 638](#)
[kubernetes.watch.watch \(module\), 637](#)
[kubernetes.watch.watch_test \(module\), 638](#)

L

[labels \(kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute\), 382](#)
[last_heartbeat_time \(kubernetes.client.models.v1_node_condition.V1NodeCondition attribute\), 372](#)
[last_probe_time \(kubernetes.client.models.v1_job_condition.V1JobCondition attribute\), 354](#)
[last_probe_time \(kubernetes.client.models.v1_pod_condition.V1PodCondition attribute\), 401](#)
[last_scale_time \(kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus attribute\), 348](#)
[last_schedule_time \(kubernetes.client.models.v2alpha1_cron_job_status.V2alpha1CronJobStatus attribute\), 519](#)
[last_state \(kubernetes.client.models.v1_container_status.V1ContainerStatus attribute\), 322](#)
[last_timestamp \(kubernetes.client.models.v1_event.V1Event attribute\), 335](#)
[last_transition_time \(kubernetes.client.models.v1_job_condition.V1JobCondition attribute\), 354](#)
[last_transition_time \(kubernetes.client.models.v1_node_condition.V1NodeCondition attribute\), 372](#)
[last_transition_time \(kubernetes.client.models.v1_pod_condition.V1PodCondition attribute\), 401](#)
[last_transition_time \(kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition attribute\), 419](#)
[last_transition_time \(kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition attribute\), 492](#)
[level \(kubernetes.client.models.v1_selinux_options.V1SELinuxOptions attribute\), 431](#)
[lifecycle \(kubernetes.client.models.v1_container.V1Container attribute\), 314](#)
[limits \(kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec attribute\), 364](#)
[limits \(kubernetes.client.models.v1_resource_requirements.V1ResourceRequirements attribute\), 428](#)
[list_cluster_role\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1Api method\), 278](#)
[list_cluster_role_binding\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1Api method\), 278](#)
[list_cluster_role_binding_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1Api method\), 279](#)
[list_cluster_role_with_http_info\(\) \(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1Api method\), 280](#)
[list_component_status\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 126](#)
[list_component_status_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 126](#)
[list_config_map_for_all_namespaces\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 127](#)
[list_config_map_for_all_namespaces_with_http_info\(\) \(kubernetes.client.apis.core_v1_api.CoreV1Api method\), 128](#)
[list_contexts\(\) \(kubernetes.config.kube_config.KubeConfigLoader method\), 529](#)
[list_controller_revision_for_all_namespaces\(\) \(kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method\), 21](#)
[list_controller_revision_for_all_namespaces_with_http_info\(\) \(kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method\), 22](#)
[list_cron_job_for_all_namespaces\(\) \(kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method\), 63](#)
[list_cron_job_for_all_namespaces_with_http_info\(\) \(kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method\), 63](#)

```

    kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 131
    method), 64
list_daemon_set_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api
    kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 132
    method), 226
list_daemon_set_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
    kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 133
    method), 227
list_deployment_for_all_namespaces() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api
    method), 133
    method), 22
list_deployment_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api
    kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 134
    method), 227
list_deployment_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
    kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 135
    method), 23
list_deployment_for_all_namespaces_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
    method), 228
list_endpoints_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api
    method), 128
list_endpoints_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
    method), 129
list_event_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api
    method), 130
list_event_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
    method), 131
list_horizontal_pod_autoscaler_for_all_namespaces() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api
    method), 45
list_horizontal_pod_autoscaler_for_all_namespaces_with_http_info() (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api
    method), 46
list_ingress_for_all_namespaces() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
    method), 229
list_ingress_for_all_namespaces_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
    method), 229
list_job_for_all_namespaces() (kubernetes.client.apis.batch_v1_api.BatchV1Api
    method), 55
list_job_for_all_namespaces_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api
    method), 55
list_kube_config_contexts() (in module kubernetes.config.kube_config), 529
list_limit_range_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api
    kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 64
    list_limit_range_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
        kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 64
    list_namespace() (kubernetes.client.apis.core_v1_api.CoreV1Api
        kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 133
    list_namespace_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
        method), 133
    list_namespaced_config_map() (kubernetes.client.apis.core_v1_api.CoreV1Api
        kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 134
    list_namespaced_config_map_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
        method), 135
    list_namespaced_controller_revision() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api
        kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 134
    list_namespaced_controller_revision_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api
        method), 24
    list_namespaced_cron_job() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api
        method), 64
    list_namespaced_cron_job_with_http_info() (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api
        method), 65
    list_namespaced_daemon_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
        method), 230
    list_namespaced_daemon_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
        method), 231
    list_namespaced_deployment() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api
        kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 231
    list_namespaced_deployment_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api
        method), 25
    list_namespaced_deployment() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
        kubernetes.client.apis.batch_v1_api.BatchV1Api method), 231
    list_namespaced_deployment_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
        method), 232
    list_namespaced_endpoints() (kubernetes.client.apis.core_v1_api.CoreV1Api
        method), 135
    list_namespaced_endpoints_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api
        method), 136
    list_namespaced_event() (kubernetes.client.apis.core_v1_api.CoreV1Api

```


- method), 137
- list_namespaced_event_with_http_info() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 138
- list_namespaced_horizontal_pod_autoscaler() (kuber-
netes.client.apis.autoscaling_v1_api.AutoscalingV1Api
method), 47
- list_namespaced_horizontal_pod_autoscaler_with_http_info() (kuber-
(kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api
method), 48
- list_namespaced_ingress() (kuber-
netes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 233
- list_namespaced_ingress_with_http_info() (kuber-
netes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 234
- list_namespaced_job() (kuber-
netes.client.apis.batch_v1_api.BatchV1Api
method), 56
- list_namespaced_job_with_http_info() (kuber-
netes.client.apis.batch_v1_api.BatchV1Api
method), 57
- list_namespaced_limit_range() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 138
- list_namespaced_limit_range_with_http_info() (ku-
bernetes.client.apis.core_v1_api.CoreV1Api
method), 139
- list_namespaced_network_policy() (kuber-
netes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 234
- list_namespaced_network_policy_with_http_info() (ku-
bernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 235
- list_namespaced_persistent_volume_claim() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 140
- list_namespaced_persistent_volume_claim_with_http_info() (kuber-
(kubernetes.client.apis.core_v1_api.CoreV1Api
method), 140
- list_namespaced_pod() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 141
- list_namespaced_pod_disruption_budget() (kuber-
netes.client.apis.policy_v1beta1_api.PolicyV1beta1Api
method), 261
- list_namespaced_pod_disruption_budget_with_http_info() (kuber-
(kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api
method), 262
- list_namespaced_pod_template() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 142
- list_namespaced_pod_template_with_http_info() (ku-
bernetes.client.apis.core_v1_api.CoreV1Api
method), 143
- list_namespaced_pod_with_http_info() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 143
- list_namespaced_replica_set() (kuber-
netes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 236
- list_namespaced_replica_set_with_http_info() (kuber-
(kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api
method), 236
- list_namespaced_replication_controller() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 144
- list_namespaced_replication_controller_with_http_info() (kuber-
(kubernetes.client.apis.core_v1_api.CoreV1Api
method), 145
- list_namespaced_resource_quota() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 146
- list_namespaced_resource_quota_with_http_info() (ku-
bernetes.client.apis.core_v1_api.CoreV1Api
method), 146
- list_namespaced_role() (kuber-
netes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthoriza
method), 281
- list_namespaced_role_binding() (kuber-
netes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthoriza
method), 281
- list_namespaced_role_binding_with_http_info() (kuber-
(kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthoriza
method), 282
- list_namespaced_role_with_http_info() (kuber-
netes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthoriza
method), 283
- list_namespaced_secret() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 147
- list_namespaced_secret_with_http_info() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 148
- list_namespaced_service() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 148
- list_namespaced_service_account() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 149
- list_namespaced_service_account_with_http_info() (ku-
(kubernetes.client.apis.core_v1_api.CoreV1Api
method), 150
- list_namespaced_service_with_http_info() (kuber-
netes.client.apis.core_v1_api.CoreV1Api
method), 151
- list_namespaced_stateful_set() (kuber-
netes.client.apis.apps_v1beta1_api.AppsV1beta1Api
method), 151

method), 27

list_namespaced_stateful_set_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 27

list_network_policy_for_all_namespaces() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 237

list_network_policy_for_all_namespaces_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 238

list_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 151

list_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 152

list_persistent_volume() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 153

list_persistent_volume_claim_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 153

list_persistent_volume_claim_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 154

list_persistent_volume_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 155

list_pod_disruption_budget_for_all_namespaces() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 263

list_pod_disruption_budget_for_all_namespaces_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 264

list_pod_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 155

list_pod_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 156

list_pod_security_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 239

list_pod_security_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 239

list_pod_template_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 157

list_pod_template_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 158

list_replica_set_for_all_namespaces() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 240

list_replica_set_for_all_namespaces_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 241

list_replication_controller_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 158

list_replication_controller_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 159

list_resource_quota_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 160

list_resource_quota_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 160

list_role_binding_for_all_namespaces() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 283

list_role_binding_for_all_namespaces_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 284

list_role_for_all_namespaces() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 285

list_role_for_all_namespaces_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 286

list_secret_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 161

list_secret_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 162

list_service_account_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 162

list_service_account_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 163

list_service_for_all_namespaces() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 164

list_service_for_all_namespaces_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 165

list_stateful_set_for_all_namespaces() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 28

list_stateful_set_for_all_namespaces_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 29

list_storage_class() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 294

list_storage_class_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 295

liveness_probe (kubernetes.client.models.v1_container.V1Container attribute), 314

load_and_set() (kubernetes.config.incluster_config.InClusterConfigLoader attribute), 362

load_and_set() (kubernetes.config.kube_config.KubeConfigLoader attribute), 347

load_balancer (kubernetes.client.models.v1_service_status.V1ServiceStatus attribute), 446

load_balancer (kubernetes.client.models.v1beta1_ingress_status.V1beta1IngressStatus attribute), 478

load_balancer_ip (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445

load_balancer_source_ranges (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445

load_incluster_config() (in module kubernetes.config.incluster_config), 528

load_kube_config() (in module kubernetes.config.kube_config), 529

local (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 396

log_file_handler() (kubernetes.client.apis.logs_api.LogsApi method), 257

log_file_handler_with_http_info() (kubernetes.client.apis.logs_api.LogsApi method), 257

log_file_list_handler() (kubernetes.client.apis.logs_api.LogsApi method), 257

log_file_list_handler_with_http_info() (kubernetes.client.apis.logs_api.LogsApi method), 258

logger_file (kubernetes.client.configuration.Configuration attribute), 525

logger_format (kubernetes.client.configuration.Configuration attribute), 525

LogsApi (class in kubernetes.client.apis.logs_api), 257

lun (kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource attribute), 339

lun (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352

M

machine_id (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378

major (kubernetes.client.models.version_info.VersionInfo attribute), 522

manual_selector (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 357

max_replicas (kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec attribute), 347

max_unavailable (kubernetes.client.models.v1beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec attribute), 489

message (kubernetes.client.models.v1_component_condition.V1ComponentCondition attribute), 306

message (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated attribute), 320

message (kubernetes.client.models.v1_container_state_waiting.V1ContainerStateWaiting attribute), 321

message (kubernetes.client.models.v1_event.V1Event attribute), 335

message (kubernetes.client.models.v1_job_condition.V1JobCondition attribute), 355

message (kubernetes.client.models.v1_node_condition.V1NodeCondition attribute), 372

message (kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus attribute), 398

message (kubernetes.client.models.v1_pod_condition.V1PodCondition attribute), 401

message (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 410

message (kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition attribute), 420

message (kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition attribute), 493

metadata (kubernetes.client.models.v1_binding.V1Binding attribute), 302

metadata (kubernetes.client.models.v1_component_status.V1ComponentStatus attribute), 307

metadata (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList attribute), 308

metadata (kubernetes.client.models.v1_config_map.V1ConfigMap attribute), 309

metadata (kubernetes.client.models.v1_config_map_list.V1ConfigMapList attribute), 311

metadata (kubernetes.client.models.v1_endpoints.V1Endpoints attribute), 331

metadata (kubernetes.client.models.v1_endpoints_list.V1EndpointsList attribute), 332

metadata (kubernetes.client.models.v1_event.V1Event attribute), 335

metadata (kubernetes.client.models.v1_event_list.V1EventList attribute), 337

metadata (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute), 347

attribute), 345
 metadata (kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList attribute), 346
 metadata (kubernetes.client.models.v1_job.V1Job attribute), 354
 metadata (kubernetes.client.models.v1_job_list.V1JobList attribute), 356
 metadata (kubernetes.client.models.v1_limit_range.V1LimitRange attribute), 361
 metadata (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList attribute), 363
 metadata (kubernetes.client.models.v1_namespace.V1Namespace attribute), 366
 metadata (kubernetes.client.models.v1_namespace_list.V1NamespaceList attribute), 367
 metadata (kubernetes.client.models.v1_node.V1Node attribute), 370
 metadata (kubernetes.client.models.v1_node_list.V1NodeList attribute), 374
 metadata (kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute), 386
 metadata (kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute), 387
 metadata (kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList attribute), 389
 metadata (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList attribute), 393
 metadata (kubernetes.client.models.v1_pod.V1Pod attribute), 400
 metadata (kubernetes.client.models.v1_pod_list.V1PodList attribute), 402
 metadata (kubernetes.client.models.v1_pod_template.V1PodTemplate attribute), 411
 metadata (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList attribute), 412
 metadata (kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec attribute), 413
 metadata (kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute), 418
 metadata (kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList attribute), 421
 metadata (kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute), 425
 metadata (kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList attribute), 426
 metadata (kubernetes.client.models.v1_scale.V1Scale attribute), 429
 metadata (kubernetes.client.models.v1_secret.V1Secret attribute), 433
 metadata (kubernetes.client.models.v1_secret_list.V1SecretList attribute), 435
 metadata (kubernetes.client.models.v1_service.V1Service attribute), 438
 metadata (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 440
 metadata (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList attribute), 441
 metadata (kubernetes.client.models.v1_service_list.V1ServiceList attribute), 442
 metadata (kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole attribute), 454
 metadata (kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute), 455
 metadata (kubernetes.client.models.v1alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList attribute), 457
 metadata (kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList attribute), 458
 metadata (kubernetes.client.models.v1alpha1_role.V1alpha1Role attribute), 460
 metadata (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding attribute), 461
 metadata (kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList attribute), 462
 metadata (kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList attribute), 463
 metadata (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet attribute), 466
 metadata (kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList attribute), 467
 metadata (kubernetes.client.models.v1beta1_eviction.V1beta1Eviction attribute), 472
 metadata (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute), 474
 metadata (kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList attribute), 476
 metadata (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute), 480
 metadata (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute), 481
 metadata (kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList attribute), 483
 metadata (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 487
 metadata (kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList attribute), 488
 metadata (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 491
 metadata (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList attribute), 494
 metadata (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 499
 metadata (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute), 501
 metadata (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList attribute), 502
 metadata (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass attribute), 507
 metadata (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList attribute), 507

attribute), 508
 metadata (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview
 attribute), 509
 metadata (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview attribute), 513
 metadata (kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob attribute), 322
 metadata (kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList attribute), 323
 metadata (kubernetes.client.models.v2alpha1_job_template_spec.V2alpha1JobTemplateSpec
 attribute), 520
 min (kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem attribute), 333
 min_available (kubernetes.client.models.v1beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec
 attribute), 489
 min_ready_seconds (kuber- attribute), 365
 netes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec.v1_object_meta.V1ObjectMeta
 attribute), 422
 min_ready_seconds (kuber- name (kubernetes.client.models.v1_object_reference.V1ObjectReference
 netes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec attribute), 384
 attribute), 468
 min_ready_seconds (kuber- attribute), 385
 netes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec.v1_secret_key_selector.V1SecretKeySelect
 attribute), 494
 min_replicas (kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec.V1ServicePort
 attribute), 347
 minor (kubernetes.client.models.version_info.VersionInfo name (kubernetes.client.models.v1_volume.V1Volume
 attribute), 522
 mode (kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile.v1_volume_mount.V1VolumeMount
 attribute), 326
 mode (kubernetes.client.models.v1_key_to_path.V1KeyToPath name (kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef
 attribute), 359
 monitors (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource.v1alpha1_subject.V1alpha1Subject
 attribute), 304
 monitors (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource.v1beta1_resource_attributes.V1beta1Resour
 attribute), 417
 mount_options (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec.v1_container_image.V1ContainerImage
 attribute), 396
 mount_options (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass.v1_object_meta.V1ObjectMeta
 attribute), 507
 mount_path (kubernetes.client.models.v1_volume_mount.V1VolumeMount name (kubernetes.client.models.v1_object_reference.V1ObjectReferen
 attribute), 452
 mount_propagation (kuber- namespace (kubernetes.client.models.v1alpha1_subject.V1alpha1Subject
 netes.client.models.v1_volume_mount.V1VolumeMount attribute), 465
 attribute), 452
 namespace (kubernetes.client.models.v1beta1_resource_attributes.V1beta1Resour
 attribute), 497
 namespace_selector (kuber-
 netes.client.models.v1beta1_network_policy_peer.V1beta1Netwo
 attribute), 483
 name (kubernetes.client.models.v1_attached_volume.V1AttachedVolume
 attribute), 299
 name (kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector
 attribute), 310
 name (kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource
 attribute), 312
 new_client_from_config() (in module kuber-
 netes.config.kube_config), 530

| | |
|---|---|
| <code>patch_namespace()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 165 | <code>patch_namespaced_deployment_scale()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 |
| <code>patch_namespace_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 165 | <code>patch_namespaced_deployment_scale()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 |
| <code>patch_namespace_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 | <code>patch_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 |
| <code>patch_namespace_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 | <code>patch_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 |
| <code>patch_namespaced_config_map()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 | <code>patch_namespaced_deployment_status()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 |
| <code>patch_namespaced_config_map_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 | <code>patch_namespaced_deployment_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 243 |
| <code>patch_namespaced_controller_revision()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 29 | <code>patch_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 |
| <code>patch_namespaced_controller_revision_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 | <code>patch_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 243 |
| <code>patch_namespaced_cron_job()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 66 | <code>patch_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 31 |
| <code>patch_namespaced_cron_job_status()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 66 | <code>patch_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 243 |
| <code>patch_namespaced_cron_job_status_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 66 | <code>patch_namespaced_endpoints()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 |
| <code>patch_namespaced_cron_job_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 66 | <code>patch_namespaced_endpoints_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 166 |
| <code>patch_namespaced_daemon_set()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 241 | <code>patch_namespaced_event()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 167 |
| <code>patch_namespaced_daemon_set_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 | <code>patch_namespaced_event_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 167 |
| <code>patch_namespaced_daemon_set_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 | <code>patch_namespaced_horizontal_pod_autoscaler()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 48 |
| <code>patch_namespaced_daemon_set_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 | <code>patch_namespaced_horizontal_pod_autoscaler_status()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 48 |
| <code>patch_namespaced_deployment()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 30 | <code>patch_namespaced_horizontal_pod_autoscaler_status_with_http_info()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 49 |
| <code>patch_namespaced_deployment()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 242 | <code>patch_namespaced_horizontal_pod_autoscaler_with_http_info()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 49 |

| | |
|--|--|
| patch_namespaced_ingress() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 243 | patch_namespaced_pod_disruption_budget_status() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 264 |
| patch_namespaced_ingress_status() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 243 | patch_namespaced_pod_disruption_budget_status_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 265 |
| patch_namespaced_ingress_status_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 244 | patch_namespaced_pod_disruption_budget_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 265 |
| patch_namespaced_ingress_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 244 | patch_namespaced_pod_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 168 |
| patch_namespaced_job() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 57 | patch_namespaced_pod_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 169 |
| patch_namespaced_job_status() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 58 | patch_namespaced_pod_template() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 169 |
| patch_namespaced_job_status_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 58 | patch_namespaced_pod_template_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 169 |
| patch_namespaced_job_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 58 | patch_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 169 |
| patch_namespaced_limit_range() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 167 | patch_namespaced_replica_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 244 |
| patch_namespaced_limit_range_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 167 | patch_namespaced_replica_set_scale() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 245 |
| patch_namespaced_network_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 244 | patch_namespaced_replica_set_scale_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 245 |
| patch_namespaced_network_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 244 | patch_namespaced_replica_set_status() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 245 |
| patch_namespaced_persistent_volume_claim() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 167 | patch_namespaced_replica_set_status_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 245 |
| patch_namespaced_persistent_volume_claim_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 168 | patch_namespaced_replica_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 245 |
| patch_namespaced_persistent_volume_claim_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 168 | patch_namespaced_replication_controller() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 169 |
| patch_namespaced_persistent_volume_claim_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 168 | patch_namespaced_replication_controller_dummy_scale() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 246 |
| patch_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 168 | patch_namespaced_replication_controller_dummy_scale_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 246 |
| patch_namespaced_pod_disruption_budget() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 264 | patch_namespaced_replication_controller_scale() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 170 |

| | |
|--|---|
| patch_namespaced_replication_controller_scale_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 170 | patch_namespaced_service_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_replication_controller_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 170 | patch_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_replication_controller_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 170 | patch_namespaced_stateful_set() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 31 |
| patch_namespaced_replication_controller_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 170 | patch_namespaced_stateful_set_scale() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 31 |
| patch_namespaced_resource_quota() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 171 | patch_namespaced_stateful_set_scale_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 31 |
| patch_namespaced_resource_quota_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 171 | patch_namespaced_stateful_set_status() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 31 |
| patch_namespaced_resource_quota_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 171 | patch_namespaced_stateful_set_status_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 32 |
| patch_namespaced_resource_quota_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 171 | patch_namespaced_stateful_set_with_http_info() (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 32 |
| patch_namespaced_role() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 287 | patch_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_role_binding() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 287 | patch_node_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_role_binding_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 287 | patch_node_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_role_with_http_info() (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 287 | patch_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 173 |
| patch_namespaced_secret() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 171 | patch_persistent_volume() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 174 |
| patch_namespaced_secret_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 172 | patch_persistent_volume_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 174 |
| patch_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 172 | patch_persistent_volume_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 174 |
| patch_namespaced_service_account() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 172 | patch_persistent_volume_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 174 |
| patch_namespaced_service_account_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 172 | patch_pod_security_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 246 |
| patch_namespaced_service_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 172 | patch_pod_security_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 246 |
| | patch_storage_class() (kubernetes.client.apis.storage_v1_api.StorageV1Api method), 175 |

netes.client.apis.storage_v1beta1_api.StorageV1beta1Api (kubernetes.client.models.version_info.VersionInfo
method), 295
platform (kubernetes.client.models.version_info.VersionInfo attribute), 522
patch_storage_class_with_http_info() (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api attribute), 374
method), 296
pod_cidr (kubernetes.client.models.v1_node_spec.V1NodeSpec attribute), 374
pod_ip (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 374
path (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 304
pod_management_policy (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 304
path (kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile attribute), 503
v1_beta1_stateful_set_spec.V1beta1StatefulSetSpec attribute), 503
path (kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterFSVolumeSource attribute), 483
v1_beta1_network_policy_peer.V1beta1NetworkPolicyPeer attribute), 483
path (kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource attribute), 485
v1_beta1_network_policy_spec.V1beta1NetworkPolicySpec attribute), 485
path (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction attribute), 485
v1_beta1_network_policy_spec.V1beta1NetworkPolicySpec attribute), 485
path (kubernetes.client.models.v1_key_to_path.V1KeyToPath attribute), 258
PolicyApi (class in kubernetes.client.apis.policy_api), 258
path (kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource attribute), 369
(class in kubernetes.client.apis.policy_v1beta1_api), 258
path (kubernetes.client.models.v1_beta1_http_ingress_path.V1beta1HTTPIngressPath attribute), 417
v1_rbd_volume_source.V1RBDVolumeSource attribute), 417
path (kubernetes.client.models.v1_beta1_non_resource_attribute.V1beta1NonResourceAttribute attribute), 324
daemon_endpoint.V1DaemonEndpoint attribute), 324
paths (kubernetes.client.models.v1_beta1_http_ingress_rule_port.V1beta1HTTPIngressRuleValue attribute), 329
v1_endpoint_port.V1EndpointPort attribute), 329
pd_id (kubernetes.client.models.v1_photon_persistent_disk_port.V1PhotonPersistentDiskVolumeSource attribute), 350
v1_http_get_action.V1HTTPGetAction attribute), 350
pd_name (kubernetes.client.models.v1_gce_persistent_disk_port.V1GCEPersistentDiskVolumeSource attribute), 443
v1_tcp_socket_action.V1TCPSocketAction attribute), 447
period_seconds (kubernetes.client.models.v1_probe.V1Probe attribute), 414
port (kubernetes.client.models.v1_beta1_network_policy_port.V1beta1NetworkPolicyPort attribute), 484
persistent_volume_claim (kubernetes.client.models.v1_volume.V1Volume attribute), 450
ports (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 352
persistent_volume_reclaim_policy (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
ports (kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset attribute), 330
phase (kubernetes.client.models.v1_namespace_status.V1NamespaceStatus attribute), 369
ports (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445
phase (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376
ports (kubernetes.client.models.v1_beta1_network_policy_ingress_rule.V1beta1NetworkPolicyIngressRule attribute), 445
phase (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute), 391
portworx_volume (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
phase (kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus attribute), 398
ports (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
phase (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 410
portworx_volume (kubernetes.client.models.v1_volume.V1Volume attribute), 450
photon_persistent_disk (kubernetes.client.models.v1_persistent_volume_spec.V1PhotonPersistentDiskVolumeSource attribute), 397
PPS() (kubernetes.client.rest.RESTClientObject attribute), 526
method), 526
photon_persistent_disk (kubernetes.client.models.v1_volume.V1Volume attribute), 450
post_start (kubernetes.client.models.v1_lifecycle.V1Lifecycle attribute), 360
pre_stop (kubernetes.client.models.v1_lifecycle.V1Lifecycle attribute), 360

attribute), 360

proxy_delete_namespaced_service_with_path_with_http_info()
 (kubernetes.client.apis.core_v1_api.CoreV1Api
 method), 176

preconditions (kubernetes.client.models.v1_delete_options.V1DeleteOptions
 attribute), 325

prepare_post_parameters() (kuber- proxy_delete_node() (kuber-
 netes.client.api_client.ApiClient method), netes.client.apis.core_v1_api.CoreV1Api
 524 method), 176

PRIMITIVE_TYPES (kuber- proxy_delete_node_with_http_info() (kuber-
 netes.client.api_client.ApiClient attribute), netes.client.apis.core_v1_api.CoreV1Api
 523 method), 176

priority (kubernetes.client.models.v1_pod_spec.V1PodSpecproxy_delete_node_with_path() (kuber-
 attribute), 407 netes.client.apis.core_v1_api.CoreV1Api
 method), 176

priority_class_name (kuber- proxy_delete_node_with_path_with_http_info() (ku-
 netes.client.models.v1_pod_spec.V1PodSpec bernetes.client.apis.core_v1_api.CoreV1Api
 attribute), 407 method), 176

privileged (kubernetes.client.models.v1_security_context.V1SecurityContextmethod), 176
 attribute), 437

projected (kubernetes.client.models.v1_volume.V1Volume netes.client.apis.core_v1_api.CoreV1Api
 attribute), 451 method), 176

propagation_policy (kuber- proxy_get_namespaced_pod_with_http_info() (ku-
 netes.client.models.v1_delete_options.V1DeleteOptions bernetes.client.apis.core_v1_api.CoreV1Api
 attribute), 325 method), 177

protocol (kubernetes.client.models.v1_container_port.V1ContainerPortnamespaced_pod_with_path() (kuber-
 attribute), 318 netes.client.apis.core_v1_api.CoreV1Api
 method), 177

protocol (kubernetes.client.models.v1_endpoint_port.V1EndpointPortmethod), 177
 attribute), 329

protocol (kubernetes.client.models.v1_service_port.V1ServicePort proxy_get_namespaced_pod_with_path_with_http_info() (kuber-
 attribute), 443 netes.client.apis.core_v1_api.CoreV1Api
 method), 177

protocol (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicyPortmethod), 177
 attribute), 484

provider_id (kubernetes.client.models.v1_node_spec.V1NodeSpec netes.client.apis.core_v1_api.CoreV1Api
 attribute), 375 method), 177

provisioner (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClassproxy_get_namespaced_service_with_http_info() (ku-
 attribute), 507 netes.client.apis.core_v1_api.CoreV1Api
 method), 177

proxy_delete_namespaced_pod() (kuber- proxy_get_namespaced_service_with_path() (kuber-
 netes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 174 method), 177

proxy_delete_namespaced_pod_with_http_info() (ku- proxy_get_namespaced_service_with_path_with_http_info() (kuber-
 bernetes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 174 method), 178

proxy_delete_namespaced_pod_with_path() (kuber- proxy_get_node() (kuber-
 netes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 175 method), 178

proxy_delete_namespaced_pod_with_path_with_http_info(proxy_get_node_with_http_info() (kuber-
 (kubernetes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 175 method), 178

proxy_delete_namespaced_service() (kuber- proxy_get_node_with_path() (kuber-
 netes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 175 method), 178

proxy_delete_namespaced_service_with_http_info() (ku- proxy_get_node_with_path_with_http_info() (kuber-
 bernetes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 175 method), 178

proxy_delete_namespaced_service_with_path() (ku- proxy_head_namespaced_pod() (kuber-
 bernetes.client.apis.core_v1_api.CoreV1Api netes.client.apis.core_v1_api.CoreV1Api
 method), 175 method), 178

| | |
|---|--|
| proxy_head_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_options_namespaced_service_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 |
| proxy_head_namespaced_pod_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_options_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 |
| proxy_head_namespaced_pod_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_options_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 |
| proxy_head_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_options_node_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 |
| proxy_head_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_options_node_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 |
| proxy_head_namespaced_service_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 179 | proxy_patch_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_head_namespaced_service_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_head_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_pod_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_head_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_pod_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_head_node_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_head_node_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 183 |
| proxy_options_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 180 | proxy_patch_namespaced_service_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 184 |
| proxy_options_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 181 | proxy_patch_namespaced_service_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 184 |
| proxy_options_namespaced_pod_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 181 | proxy_patch_node() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 184 |
| proxy_options_namespaced_pod_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 181 | proxy_patch_node_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 184 |
| proxy_options_namespaced_service() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 181 | proxy_patch_node_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 184 |
| proxy_options_namespaced_service_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 181 | proxy_patch_node_with_path_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185 |
| proxy_options_namespaced_service_with_path() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 182 | proxy_post_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185 |

- `proxy_post_namespaced_pod_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185
- `proxy_post_namespaced_pod_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185
- `proxy_post_namespaced_pod_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185
- `proxy_post_namespaced_service()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 185
- `proxy_post_namespaced_service_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_namespaced_service_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_namespaced_service_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_node()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_node_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_node_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 186
- `proxy_post_node_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_pod()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_pod_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_pod_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_pod_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_service()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 187
- `proxy_put_namespaced_service_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_namespaced_service_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_namespaced_service_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_node()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_node_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_node_with_path()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 188
- `proxy_put_node_with_path_with_http_info()` (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189
- `publish_not_ready_addresses` (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445
- `PUT()` (kubernetes.client.rest.RESTClientObject method), 526
- ## Q
- `qos_class` (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 410
- `quobyte` (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
- `quobyte` (kubernetes.client.models.v1_volume.V1Volume attribute), 451
- ## R
- `raw` (kubernetes.client.models.runtime_raw_extension.RuntimeRawExtension attribute), 297
- `RbacAuthorizationApi` (class in kubernetes.client.apis.rbac_authorization_api), 267
- `RbacAuthorizationV1alpha1Api` (class in kubernetes.client.apis.rbac_authorization_v1alpha1_api), 268
- `rbd` (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
- `rbd` (kubernetes.client.models.v1_volume.V1Volume attribute), 451
- `read_cluster_role()` (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationApi method), 288
- `read_cluster_role_binding()` (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationApi method), 288
- `read_cluster_role_binding_with_http_info()` (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationApi method), 288
- `read_cluster_role_with_http_info()` (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationApi method), 288

| | |
|--|--|
| <code>read_component_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189 | <code>read_namespaced_deployment()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_component_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189 | <code>read_namespaced_deployment()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 247 |
| <code>read_namespace()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189 | <code>read_namespaced_deployment_scale()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_namespace_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189 | <code>read_namespaced_deployment_scale()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 247 |
| <code>read_namespace_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 189 | <code>read_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_namespace_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 190 | <code>read_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 248 |
| <code>read_namespaced_config_map()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 190 | <code>read_namespaced_deployment_status()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_namespaced_config_map_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 190 | <code>read_namespaced_deployment_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 248 |
| <code>read_namespaced_controller_revision()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 32 | <code>read_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_namespaced_controller_revision_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 32 | <code>read_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 248 |
| <code>read_namespaced_cron_job()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 66 | <code>read_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 33 |
| <code>read_namespaced_cron_job_status()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 67 | <code>read_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 248 |
| <code>read_namespaced_cron_job_status_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 67 | <code>read_namespaced_endpoints()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 190 |
| <code>read_namespaced_cron_job_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 67 | <code>read_namespaced_endpoints_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 190 |
| <code>read_namespaced_daemon_set()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 246 | <code>read_namespaced_event()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 191 |
| <code>read_namespaced_daemon_set_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 247 | <code>read_namespaced_event_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 191 |
| <code>read_namespaced_daemon_set_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 247 | <code>read_namespaced_horizontal_pod_autoscaler()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 49 |
| <code>read_namespaced_daemon_set_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 247 | <code>read_namespaced_horizontal_pod_autoscaler_status()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 49 |

[read_namespaced_horizontal_pod_autoscaler_status_with_http_info\(\)](#) (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 50
[read_namespaced_horizontal_pod_autoscaler_with_http_info\(\)](#) (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 50
[read_namespaced_ingress\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 248
[read_namespaced_ingress_status\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 249
[read_namespaced_ingress_status_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 249
[read_namespaced_ingress_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 249
[read_namespaced_job\(\)](#) (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 58
[read_namespaced_job_status\(\)](#) (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 58
[read_namespaced_job_status_with_http_info\(\)](#) (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 59
[read_namespaced_job_with_http_info\(\)](#) (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 59
[read_namespaced_limit_range\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 191
[read_namespaced_limit_range_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 191
[read_namespaced_network_policy\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 249
[read_namespaced_network_policy_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 249
[read_namespaced_persistent_volume_claim\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 192
[read_namespaced_persistent_volume_claim_status\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 192
[read_namespaced_persistent_volume_claim_status_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 192
[read_namespaced_persistent_volume_claim_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 192
[read_namespaced_pod\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 192
[read_namespaced_pod_disruption_budget\(\)](#) (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 265
[read_namespaced_pod_disruption_budget_status\(\)](#) (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 265
[read_namespaced_pod_disruption_budget_status_with_http_info\(\)](#) (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 265
[read_namespaced_pod_disruption_budget_with_http_info\(\)](#) (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 266
[read_namespaced_pod_log\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 193
[read_namespaced_pod_log_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 193
[read_namespaced_pod_status\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 193
[read_namespaced_pod_status_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 194
[read_namespaced_pod_template\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 194
[read_namespaced_pod_template_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 194
[read_namespaced_pod_with_http_info\(\)](#) (kubernetes.client.apis.core_v1_api.CoreV1Api method), 194
[read_namespaced_replica_set\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 250
[read_namespaced_replica_set_scale\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 250
[read_namespaced_replica_set_scale_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 250
[read_namespaced_replica_set_status\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 250
[read_namespaced_replica_set_status_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 250
[read_namespaced_replica_set_with_http_info\(\)](#) (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 251

| | |
|---|--|
| <code>read_namespaced_replication_controller()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 194 | <code>read_namespaced_service()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_replication_controller_dummy_scale()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 251 | <code>read_namespaced_service_account()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_replication_controller_dummy_scale_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 251 | <code>read_namespaced_service_account_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_replication_controller_scale()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 195 | <code>read_namespaced_service_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 198 |
| <code>read_namespaced_replication_controller_scale_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 195 | <code>read_namespaced_service_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 198 |
| <code>read_namespaced_replication_controller_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 195 | <code>read_namespaced_service_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 198 |
| <code>read_namespaced_replication_controller_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 195 | <code>read_namespaced_stateful_set()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 34 |
| <code>read_namespaced_replication_controller_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 196 | <code>read_namespaced_stateful_set_scale()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 34 |
| <code>read_namespaced_resource_quota()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 196 | <code>read_namespaced_stateful_set_scale_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 34 |
| <code>read_namespaced_resource_quota_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 196 | <code>read_namespaced_stateful_set_status()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 34 |
| <code>read_namespaced_resource_quota_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 196 | <code>read_namespaced_stateful_set_status_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 34 |
| <code>read_namespaced_resource_quota_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 196 | <code>read_namespaced_stateful_set_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 35 |
| <code>read_namespaced_role()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 288 | <code>read_node()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_role_binding()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 288 | <code>read_node_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_role_binding_with_http_info()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 289 | <code>read_node_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_role_with_http_info()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 289 | <code>read_node_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 |
| <code>read_namespaced_secret()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 | <code>read_only</code> (kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AwsElasticBlockStoreVolumeSource attribute), 300 |
| <code>read_namespaced_secret_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 197 | <code>read_only</code> (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 301 |
| | <code>read_only</code> (kubernetes.client.models.v1_azure_file_volume_source.V1AzureFileVolumeSource attribute), 301 |
| | <code>read_only</code> (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 301 |

| | |
|--------------|------------|
| Index | 701 |
|--------------|------------|

| | |
|---|---|
| <code>replace_namespace_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 200 | <code>replace_namespaced_deployment_scale()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 253 |
| <code>replace_namespace_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 200 | <code>replace_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 36 |
| <code>replace_namespace_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 200 | <code>replace_namespaced_deployment_scale_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 253 |
| <code>replace_namespaced_config_map()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 201 | <code>replace_namespaced_deployment_status()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 36 |
| <code>replace_namespaced_config_map_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 201 | <code>replace_namespaced_deployment_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 253 |
| <code>replace_namespaced_controller_revision()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 35 | <code>replace_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 36 |
| <code>replace_namespaced_controller_revision_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 35 | <code>replace_namespaced_deployment_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 253 |
| <code>replace_namespaced_cron_job()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 67 | <code>replace_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 36 |
| <code>replace_namespaced_cron_job_status()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 67 | <code>replace_namespaced_deployment_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 |
| <code>replace_namespaced_cron_job_status_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 68 | <code>replace_namespaced_endpoints()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 201 |
| <code>replace_namespaced_cron_job_with_http_info()</code> (kubernetes.client.apis.batch_v2alpha1_api.BatchV2alpha1Api method), 68 | <code>replace_namespaced_endpoints_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 201 |
| <code>replace_namespaced_daemon_set()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 252 | <code>replace_namespaced_event()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 201 |
| <code>replace_namespaced_daemon_set_status()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 252 | <code>replace_namespaced_event_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 202 |
| <code>replace_namespaced_daemon_set_status_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 252 | <code>replace_namespaced_horizontal_pod_autoscaler()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 50 |
| <code>replace_namespaced_daemon_set_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 252 | <code>replace_namespaced_horizontal_pod_autoscaler_status()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 50 |
| <code>replace_namespaced_deployment()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 35 | <code>replace_namespaced_horizontal_pod_autoscaler_status_with_http_info()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 51 |
| <code>replace_namespaced_deployment()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 252 | <code>replace_namespaced_horizontal_pod_autoscaler_with_http_info()</code> (kubernetes.client.apis.autoscaling_v1_api.AutoscalingV1Api method), 51 |
| <code>replace_namespaced_deployment_scale()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 36 | <code>replace_namespaced_ingress()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 |

| | | | |
|---|---|---|---|
| replace_namespaced_ingress_status() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 | replace_namespaced_ingress_status_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 | replace_namespaced_pod_disruption_budget_status_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 266 | replace_namespaced_pod_disruption_budget_with_http_info() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 267 |
| replace_namespaced_ingress_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 | replace_namespaced_job() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 59 | replace_namespaced_pod_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 203 | replace_namespaced_pod_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 203 |
| replace_namespaced_job_status() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 59 | replace_namespaced_job_status_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 59 | replace_namespaced_pod_template() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 204 | replace_namespaced_pod_template_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 204 |
| replace_namespaced_job_with_http_info() (kubernetes.client.apis.batch_v1_api.BatchV1Api method), 60 | replace_namespaced_limit_range() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 202 | replace_namespaced_pod_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 204 | replace_namespaced_replica_set() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 255 |
| replace_namespaced_limit_range_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 202 | replace_namespaced_network_policy() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 254 | replace_namespaced_replica_set_scale() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 255 | replace_namespaced_replica_set_scale_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 255 |
| replace_namespaced_network_policy_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 255 | replace_namespaced_persistent_volume_claim() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 202 | replace_namespaced_replica_set_status() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 255 | replace_namespaced_replica_set_status_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 256 |
| replace_namespaced_persistent_volume_claim_status() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 202 | replace_namespaced_persistent_volume_claim_status_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 203 | replace_namespaced_replica_set_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 256 | replace_namespaced_replication_controller() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 204 |
| replace_namespaced_persistent_volume_claim_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 203 | replace_namespaced_pod() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 203 | replace_namespaced_replication_controller_dummy_scale() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 256 | replace_namespaced_replication_controller_dummy_scale_with_http_info() (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 256 |
| replace_namespaced_pod_disruption_budget() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 266 | replace_namespaced_pod_disruption_budget_status() (kubernetes.client.apis.policy_v1beta1_api.PolicyV1beta1Api method), 266 | replace_namespaced_replication_controller_scale() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 204 | replace_namespaced_replication_controller_scale_with_http_info() (kubernetes.client.apis.core_v1_api.CoreV1Api method), 205 |

| | |
|---|--|
| <code>replace_namespaced_replication_controller_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 205 | <code>replace_namespaced_service_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 |
| <code>replace_namespaced_replication_controller_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 205 | <code>replace_namespaced_stateful_set()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 37 |
| <code>replace_namespaced_replication_controller_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 205 | <code>replace_namespaced_stateful_set_scale()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 37 |
| <code>replace_namespaced_resource_quota()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 206 | <code>replace_namespaced_stateful_set_scale_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 37 |
| <code>replace_namespaced_resource_quota_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 206 | <code>replace_namespaced_stateful_set_status()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 37 |
| <code>replace_namespaced_resource_quota_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 206 | <code>replace_namespaced_stateful_set_status_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 37 |
| <code>replace_namespaced_resource_quota_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 206 | <code>replace_namespaced_stateful_set_with_http_info()</code> (kubernetes.client.apis.apps_v1beta1_api.AppsV1beta1Api method), 38 |
| <code>replace_namespaced_role()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 290 | <code>replace_node()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 |
| <code>replace_namespaced_role_binding()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 290 | <code>replace_node_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 |
| <code>replace_namespaced_role_binding_with_http_info()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 290 | <code>replace_node_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 |
| <code>replace_namespaced_role_with_http_info()</code> (kubernetes.client.apis.rbac_authorization_v1alpha1_api.RbacAuthorizationV1alpha1Api method), 290 | <code>replace_node_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 |
| <code>replace_namespaced_secret()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_persistent_volume()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 209 |
| <code>replace_namespaced_secret_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_persistent_volume_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 209 |
| <code>replace_namespaced_service()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_persistent_volume_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 209 |
| <code>replace_namespaced_service_account()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_persistent_volume_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 209 |
| <code>replace_namespaced_service_account_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_pod_security_policy()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 257 |
| <code>replace_namespaced_service_status()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 207 | <code>replace_pod_security_policy_with_http_info()</code> (kubernetes.client.apis.extensions_v1beta1_api.ExtensionsV1beta1Api method), 257 |
| <code>replace_namespaced_service_status_with_http_info()</code> (kubernetes.client.apis.core_v1_api.CoreV1Api method), 208 | <code>replace_storage_class()</code> (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 296 |
| | <code>replace_storage_class_with_http_info()</code> (kubernetes.client.apis.storage_v1beta1_api.StorageV1beta1Api method), 296 |

netes.client.apis.storage_v1beta1_api.StorageV1beta1Api attribute), 390
 method), 296
 resources (kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule attribute), 390
 restart_count (kubernetes.client.models.v1_container_status.V1ContainerStatus attribute), 372
 restart_policy (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 407
 RESTClientObject (class in kubernetes.client.rest), 526
 RESTResponse (class in kubernetes.client.rest), 526
 revision (kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource attribute), 342
 revision_history_limit (kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec attribute), 468
 revision_history_limit (kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec attribute), 503
 role (kubernetes.client.models.v1_selinux_options.V1SELinuxOptions attribute), 505
 role_ref (kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute), 456
 role_ref (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding attribute), 461
 rules (kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole attribute), 456
 rules (kubernetes.client.models.v1alpha1_role.V1alpha1Role attribute), 461
 rules (kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec attribute), 471
 run_as_non_root (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 500
 run_as_non_root (kubernetes.client.models.v1_security_context.V1SecurityContext attribute), 511
 run_as_user (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 403
 running (kubernetes.client.models.v1_container_state.V1ContainerState attribute), 318
 RuntimeRawExtension (class in kubernetes.client.models.runtime_raw_extension), 527
 sanitize_for_serialization() (kubernetes.config.kube_config.ConfigNode method), 529
 scale_io (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
 scale_io (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 397

S

safe_get() (kubernetes.config.kube_config.ConfigNode method), 529
 sanitize_for_serialization() (kubernetes.config.kube_config.ConfigNode method), 529
 scale_io (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397
 scale_io (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 397

| | |
|--|---|
| scale_io (kubernetes.client.models.v1_volume.V1Volume attribute), 451 | selector (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 357 |
| scale_target_ref (kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec attribute), 347 | selector (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 360 |
| schedule (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 518 | selector (kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec attribute), 322 |
| scheduler_name (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 407 | selector (kubernetes.client.models.v1_scale_status.V1ScaleStatus attribute), 431 |
| scheme (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction attribute), 350 | selector (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 445 |
| scopes (kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec attribute), 427 | selector (kubernetes.client.models.v1_beta1_daemon_set_spec.V1beta1DaemonSetSpec attribute), 468 |
| se_linux_options (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 404 | selector (kubernetes.client.models.v1_beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec attribute), 489 |
| se_linux_options (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 404 | selector (kubernetes.client.models.v1_beta1_replica_set_spec.V1beta1ReplicaSetSpec attribute), 495 |
| se_linux_options (kubernetes.client.models.v1_security_context.V1SecurityContext attribute), 437 | selector (kubernetes.client.models.v1_beta1_stateful_set_spec.V1beta1StatefulSetSpec attribute), 503 |
| secret (kubernetes.client.models.v1_volume.V1Volume attribute), 451 | server (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 382 |
| secret_file (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 304 | server (kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource attribute), 369 |
| secret_key_ref (kubernetes.client.models.v1_env_var_source.V1EnvVarSource attribute), 334 | service (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 304 |
| secret_name (kubernetes.client.models.v1_azure_file_volume_source.V1AzureFileVolumeSource attribute), 302 | service_account_name (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 408 |
| secret_name (kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource attribute), 436 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_name (kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource attribute), 436 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_name (kubernetes.client.models.v1_beta1_ingress_tls_virtual_host.V1beta1IngressVirtualHost attribute), 479 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_ref (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 304 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_ref (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource attribute), 304 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_ref (kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource attribute), 340 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_ref (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 353 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secret_ref (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource attribute), 417 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| secrets (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 440 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| security_context (kubernetes.client.models.v1_container.V1Container attribute), 315 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| security_context (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 407 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| select_header_accept() (kubernetes.client.api_client.ApiClient method), 524 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |
| select_header_content_type() (kubernetes.client.api_client.ApiClient method), 524 | service_account_name (kubernetes.client.models.v1_beta1_ingress_spec.V1beta1IngressSpec attribute), 408 |

setUp() (kubernetes.test.test_apis_api.TestApisApi method), 532

setUp() (kubernetes.test.test_apps_api.TestAppsApi method), 532

setUp() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533

setUp() (kubernetes.test.test_authentication_api.TestAuthenticationApi method), 535

setUp() (kubernetes.test.test_authentication_v1beta1_api.TestAuthenticationV1beta1Api method), 535

setUp() (kubernetes.test.test_authorization_api.TestAuthorizationApi method), 535

setUp() (kubernetes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api method), 536

setUp() (kubernetes.test.test_autoscaling_api.TestAutoscalingApi method), 536

setUp() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537

setUp() (kubernetes.test.test_batch_api.TestBatchApi method), 537

setUp() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538

setUp() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539

setUp() (kubernetes.test.test_certificates_api.TestCertificatesApi method), 540

setUp() (kubernetes.test.test_core_api.TestCoreApi method), 540

setUp() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 540

setUp() (kubernetes.test.test_extensions_api.TestExtensionsApi method), 552

setUp() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 552

setUp() (kubernetes.test.test_logs_api.TestLogsApi method), 556

setUp() (kubernetes.test.test_policy_api.TestPolicyApi method), 556

setUp() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557

setUp() (kubernetes.test.test_rbac_authorization_api.TestRbacAuthorizationApi method), 558

setUp() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

setUp() (kubernetes.test.test_runtime_raw_extension.TestRuntimeRawExtension method), 560

setUp() (kubernetes.test.test_storage_api.TestStorageApi method), 560

setUp() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561

setUp() (kubernetes.test.test_v1_attached_volume.TestV1AttachedVolume method), 562

setUp() (kubernetes.test.test_v1_aws_elastic_block_store_volume_source.TestV1AwsElasticBlockStoreVolumeSource method), 563

setUp() (kubernetes.test.test_v1_azure_disk_volume_source.TestV1AzureDiskVolumeSource method), 563

setUp() (kubernetes.test.test_v1_azure_file_volume_source.TestV1AzureFileVolumeSource method), 563

setUp() (kubernetes.test.test_v1_binding.TestV1Binding method), 564

setUp() (kubernetes.test.test_v1_capabilities.TestV1Capabilities method), 564

setUp() (kubernetes.test.test_v1_ceph_fs_volume_source.TestV1CephFSVolumeSource method), 564

setUp() (kubernetes.test.test_v1_cinder_volume_source.TestV1CinderVolumeSource method), 565

setUp() (kubernetes.test.test_v1_component_condition.TestV1ComponentCondition method), 565

setUp() (kubernetes.test.test_v1_component_status.TestV1ComponentStatus method), 566

setUp() (kubernetes.test.test_v1_component_status_list.TestV1ComponentStatusList method), 566

setUp() (kubernetes.test.test_v1_config_map.TestV1ConfigMap method), 566

setUp() (kubernetes.test.test_v1_config_map_key_selector.TestV1ConfigMapKeySelector method), 567

setUp() (kubernetes.test.test_v1_config_map_list.TestV1ConfigMapList method), 567

setUp() (kubernetes.test.test_v1_config_map_volume_source.TestV1ConfigMapVolumeSource method), 567

setUp() (kubernetes.test.test_v1_container.TestV1Container method), 568

setUp() (kubernetes.test.test_v1_container_image.TestV1ContainerImage method), 568

setUp() (kubernetes.test.test_v1_container_port.TestV1ContainerPort method), 568

setUp() (kubernetes.test.test_v1_container_state.TestV1ContainerState method), 569

setUp() (kubernetes.test.test_v1_container_state_running.TestV1ContainerStateRunning method), 569

setUp() (kubernetes.test.test_v1_container_state_terminated.TestV1ContainerStateTerminated method), 569

setUp() (kubernetes.test.test_v1_container_state_waiting.TestV1ContainerStateWaiting method), 570

setUp() (kubernetes.test.test_v1_container_status.TestV1ContainerStatus method), 570

setUp() (kubernetes.test.test_v1_cron_job.TestV1CronJob method), 570

setUp() (kubernetes.test.test_v1_cron_job_list.TestV1CronJobList method), 571

setUp() (kubernetes.test.test_v1_daemon_endpoint.TestV1DaemonEndpoint method), 571

setUp() (kubernetes.test.test_v1_delete_options.TestV1DeleteOptions method), 571

setUp() (kubernetes.test.test_v1_downward_api_volume_file.TestV1DownwardApiVolumeFile method), 572

setUp() (kubernetes.test.test_v1_downward_api_volume_source.TestV1DownwardApiVolumeSource method), 572

setUp() (kubernetes.test.test_v1_empty_dir_volume_source.TestV1EmptyDirVolumeSource method), 572

setUp() (kubernetes.test.test_v1_aws_elastic_block_store_volume_source.TestV1AwsElasticBlockStoreVolumeSource method), 572

setUp() (kubernetes.test.test_v1_endpoint_address.TestV1EndpointAddress method), 573
 setUp() (kubernetes.test.test_v1_endpoint_port.TestV1EndpointPort method), 573
 setUp() (kubernetes.test.test_v1_endpoint_subset.TestV1EndpointSubset method), 573
 setUp() (kubernetes.test.test_v1_endpoints.TestV1Endpoints method), 574
 setUp() (kubernetes.test.test_v1_endpoints_list.TestV1EndpointsList method), 574
 setUp() (kubernetes.test.test_v1_env_var.TestV1EnvVar method), 574
 setUp() (kubernetes.test.test_v1_env_var_source.TestV1EnvVarSource method), 575
 setUp() (kubernetes.test.test_v1_event.TestV1Event method), 575
 setUp() (kubernetes.test.test_v1_event_list.TestV1EventList method), 576
 setUp() (kubernetes.test.test_v1_event_source.TestV1EventSource method), 576
 setUp() (kubernetes.test.test_v1_exec_action.TestV1ExecAction method), 576
 setUp() (kubernetes.test.test_v1_fc_volume_source.TestV1FCVolumeSource method), 577
 setUp() (kubernetes.test.test_v1_flex_volume_source.TestV1FlexVolumeSource method), 577
 setUp() (kubernetes.test.test_v1_flocker_volume_source.TestV1FlockerVolumeSource method), 577
 setUp() (kubernetes.test.test_v1_gce_persistent_disk_volume_source.TestV1GCEPersistentDiskVolumeSource method), 578
 setUp() (kubernetes.test.test_v1_git_repo_volume_source.TestV1GitRepoVolumeSource method), 578
 setUp() (kubernetes.test.test_v1_glusterfs_volume_source.TestV1GlusterfsVolumeSource method), 578
 setUp() (kubernetes.test.test_v1_handler.TestV1Handler method), 579
 setUp() (kubernetes.test.test_v1_horizontal_pod_autoscaler.TestV1HorizontalPodAutoscaler method), 579
 setUp() (kubernetes.test.test_v1_horizontal_pod_autoscaler_list.TestV1HorizontalPodAutoscalerList method), 579
 setUp() (kubernetes.test.test_v1_horizontal_pod_autoscaler_spec.TestV1HorizontalPodAutoscalerSpec method), 580
 setUp() (kubernetes.test.test_v1_horizontal_pod_autoscaler_status.TestV1HorizontalPodAutoscalerStatus method), 580
 setUp() (kubernetes.test.test_v1_host_path_volume_source.TestV1HostPathVolumeSource method), 581
 setUp() (kubernetes.test.test_v1_http_get_action.TestV1HTTPGetAction method), 581
 setUp() (kubernetes.test.test_v1_http_header.TestV1HTTPHeader method), 581
 setUp() (kubernetes.test.test_v1_iscsi_volume_source.TestV1ISCSIVolumeSource method), 582
 setUp() (kubernetes.test.test_v1_job.TestV1Job method), 582
 setUp() (kubernetes.test.test_v1_job_condition.TestV1JobCondition method), 582
 setUp() (kubernetes.test.test_v1_job_list.TestV1JobList method), 583
 setUp() (kubernetes.test.test_v1_job_spec.TestV1JobSpec method), 583
 setUp() (kubernetes.test.test_v1_job_status.TestV1JobStatus method), 583
 setUp() (kubernetes.test.test_v1_key_to_path.TestV1KeyToPath method), 584
 setUp() (kubernetes.test.test_v1_lifecycle.TestV1Lifecycle method), 584
 setUp() (kubernetes.test.test_v1_limit_range.TestV1LimitRange method), 584
 setUp() (kubernetes.test.test_v1_limit_range_item.TestV1LimitRangeItem method), 585
 setUp() (kubernetes.test.test_v1_limit_range_list.TestV1LimitRangeList method), 585
 setUp() (kubernetes.test.test_v1_limit_range_spec.TestV1LimitRangeSpec method), 586
 setUp() (kubernetes.test.test_v1_load_balancer_ingress.TestV1LoadBalancerIngress method), 586
 setUp() (kubernetes.test.test_v1_load_balancer_status.TestV1LoadBalancerStatus method), 586
 setUp() (kubernetes.test.test_v1_local_object_reference.TestV1LocalObjectReference method), 587
 setUp() (kubernetes.test.test_v1_namespace.TestV1Namespace method), 587
 setUp() (kubernetes.test.test_v1_namespace_list.TestV1NamespaceList method), 587
 setUp() (kubernetes.test.test_v1_namespace_spec.TestV1NamespaceSpec method), 588
 setUp() (kubernetes.test.test_v1_namespace_status.TestV1NamespaceStatus method), 588
 setUp() (kubernetes.test.test_v1_nfs_volume_source.TestV1NFSVolumeSource method), 588
 setUp() (kubernetes.test.test_v1_node.TestV1Node method), 589
 setUp() (kubernetes.test.test_v1_node_address.TestV1NodeAddress method), 589
 setUp() (kubernetes.test.test_v1_node_condition.TestV1NodeCondition method), 589
 setUp() (kubernetes.test.test_v1_node_daemon_endpoints.TestV1NodeDaemonEndpoints method), 590
 setUp() (kubernetes.test.test_v1_node_list.TestV1NodeList method), 590
 setUp() (kubernetes.test.test_v1_node_spec.TestV1NodeSpec method), 591
 setUp() (kubernetes.test.test_v1_node_status.TestV1NodeStatus method), 591
 setUp() (kubernetes.test.test_v1_node_system_info.TestV1NodeSystemInfo method), 591
 setUp() (kubernetes.test.test_v1_object_field_selector.TestV1ObjectFieldSelector method), 592

[setUp\(\) \(kubernetes.test.test_v1_object_meta.TestV1ObjectMeta method\), 592](#)
[setUp\(\) \(kubernetes.test.test_v1_object_reference.TestV1ObjectReference method\), 592](#)
[setUp\(\) \(kubernetes.test.test_v1_owner_reference.TestV1OwnerReference method\), 593](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume.TestV1PersistentVolume method\), 593](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_claim.TestV1PersistentVolumeClaim method\), 593](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_claim_list.TestV1PersistentVolumeClaimList method\), 594](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_claim_spec.TestV1PersistentVolumeClaimSpec method\), 594](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_claim_status.TestV1PersistentVolumeClaimStatus method\), 594](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_claim_volume_source.TestV1PersistentVolumeClaimVolumeSource method\), 595](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_list.TestV1PersistentVolumeList method\), 595](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_spec.TestV1PersistentVolumeSpec method\), 596](#)
[setUp\(\) \(kubernetes.test.test_v1_persistent_volume_status.TestV1PersistentVolumeStatus method\), 596](#)
[setUp\(\) \(kubernetes.test.test_v1_photon_persistent_disk_volume_source.TestV1PhotonPersistentDiskVolumeSource method\), 596](#)
[setUp\(\) \(kubernetes.test.test_v1_pod.TestV1Pod method\), 597](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_condition.TestV1PodCondition method\), 597](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_list.TestV1PodList method\), 597](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_security_context.TestV1PodSecurityContext method\), 598](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_spec.TestV1PodSpec method\), 598](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_status.TestV1PodStatus method\), 598](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_template.TestV1PodTemplate method\), 599](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_template_list.TestV1PodTemplateList method\), 599](#)
[setUp\(\) \(kubernetes.test.test_v1_pod_template_spec.TestV1PodTemplateSpec method\), 599](#)
[setUp\(\) \(kubernetes.test.test_v1_preconditions.TestV1Preconditions method\), 600](#)
[setUp\(\) \(kubernetes.test.test_v1_probe.TestV1Probe method\), 600](#)
[setUp\(\) \(kubernetes.test.test_v1_quobyte_volume_source.TestV1QuobyteVolumeSource method\), 601](#)
[setUp\(\) \(kubernetes.test.test_v1_rbd_volume_source.TestV1RBDVolumeSource method\), 601](#)
[setUp\(\) \(kubernetes.test.test_v1_replication_controller.TestV1ReplicationController method\), 601](#)
[setUp\(\) \(kubernetes.test.test_v1_replication_controller_condition.TestV1ReplicationControllerCondition method\), 602](#)
[setUp\(\) \(kubernetes.test.test_v1_replication_controller_list.TestV1ReplicationControllerList method\), 602](#)
[setUp\(\) \(kubernetes.test.test_v1_replication_controller_spec.TestV1ReplicationControllerSpec method\), 602](#)
[setUp\(\) \(kubernetes.test.test_v1_replication_controller_status.TestV1ReplicationControllerStatus method\), 603](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_field_selector.TestV1ResourceFieldSelector method\), 603](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_requirements.TestV1ResourceRequirements method\), 603](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_quota.TestV1ResourceQuota method\), 603](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_quota_list.TestV1ResourceQuotaList method\), 604](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_quota_spec.TestV1ResourceQuotaSpec method\), 604](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_quota_status.TestV1ResourceQuotaStatus method\), 604](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_quota_status_spec.TestV1ResourceQuotaStatusSpec method\), 604](#)
[setUp\(\) \(kubernetes.test.test_v1_resource_requirements.TestV1ResourceRequirements method\), 605](#)
[setUp\(\) \(kubernetes.test.test_v1_scale.TestV1Scale method\), 605](#)
[setUp\(\) \(kubernetes.test.test_v1_scale_spec.TestV1ScaleSpec method\), 606](#)
[setUp\(\) \(kubernetes.test.test_v1_scaled_disk_volume_source.TestV1ScaledDiskVolumeSource method\), 606](#)
[setUp\(\) \(kubernetes.test.test_v1_selinux_options.TestV1SELinuxOptions method\), 606](#)
[setUp\(\) \(kubernetes.test.test_v1_secret.TestV1Secret method\), 607](#)
[setUp\(\) \(kubernetes.test.test_v1_secret_key_selector.TestV1SecretKeySelector method\), 607](#)
[setUp\(\) \(kubernetes.test.test_v1_secret_list.TestV1SecretList method\), 607](#)
[setUp\(\) \(kubernetes.test.test_v1_secret_volume_source.TestV1SecretVolumeSource method\), 608](#)
[setUp\(\) \(kubernetes.test.test_v1_security_context.TestV1SecurityContext method\), 608](#)
[setUp\(\) \(kubernetes.test.test_v1_service.TestV1Service method\), 608](#)
[setUp\(\) \(kubernetes.test.test_v1_service_account.TestV1ServiceAccount method\), 609](#)
[setUp\(\) \(kubernetes.test.test_v1_service_account_list.TestV1ServiceAccountList method\), 609](#)
[setUp\(\) \(kubernetes.test.test_v1_service_list.TestV1ServiceList method\), 609](#)
[setUp\(\) \(kubernetes.test.test_v1_service_port.TestV1ServicePort method\), 610](#)
[setUp\(\) \(kubernetes.test.test_v1_service_spec.TestV1ServiceSpec method\), 610](#)
[setUp\(\) \(kubernetes.test.test_v1_service_status.TestV1ServiceStatus method\), 611](#)
[setUp\(\) \(kubernetes.test.test_v1_tcp_socket_action.TestV1TCPSocketAction method\), 611](#)

```

setUp() (kubernetes.test.test_v1_volume.TestV1Volume
method), 611
setUp() (kubernetes.test.test_v1_volume_mount.TestV1VolumeMount
method), 612
setUp() (kubernetes.test.test_v1_vsphere_virtual_disk_volume.TestV1VsphereVirtualDiskVolume
method), 612
setUp() (kubernetes.test.test_v1alpha1_cluster_role.TestV1alpha1ClusterRole
method), 613
setUp() (kubernetes.test.test_v1alpha1_cluster_role_bindings.TestV1alpha1ClusterRoleBinding
method), 613
setUp() (kubernetes.test.test_v1alpha1_cluster_role_bindings_list.TestV1alpha1ClusterRoleBindingList
method), 613
setUp() (kubernetes.test.test_v1alpha1_cluster_role_list.TestV1alpha1ClusterRoleList
method), 614
setUp() (kubernetes.test.test_v1alpha1_policy_rule.TestV1alpha1PolicyRule
method), 614
setUp() (kubernetes.test.test_v1alpha1_role.TestV1alpha1Role
method), 614
setUp() (kubernetes.test.test_v1alpha1_role_binding.TestV1alpha1RoleBinding
method), 615
setUp() (kubernetes.test.test_v1alpha1_role_binding_list.TestV1alpha1RoleBindingList
method), 615
setUp() (kubernetes.test.test_v1alpha1_role_list.TestV1alpha1RoleList
method), 615
setUp() (kubernetes.test.test_v1alpha1_role_ref.TestV1alpha1RoleRef
method), 616
setUp() (kubernetes.test.test_v1alpha1_subject.TestV1alpha1Subject
method), 616
setUp() (kubernetes.test.test_v1beta1_daemon_set.TestV1beta1DaemonSet
method), 617
setUp() (kubernetes.test.test_v1beta1_daemon_set_list.TestV1beta1DaemonSetList
method), 617
setUp() (kubernetes.test.test_v1beta1_daemon_set_spec.TestV1beta1DaemonSetSpec
method), 617
setUp() (kubernetes.test.test_v1beta1_daemon_set_status.TestV1beta1DaemonSetStatus
method), 618
setUp() (kubernetes.test.test_v1beta1_eviction.TestV1beta1Eviction
method), 618
setUp() (kubernetes.test.test_v1beta1_http_ingress_path.TestV1beta1HTTPIngressPath
method), 619
setUp() (kubernetes.test.test_v1beta1_http_ingress_rule_value.TestV1beta1HTTPIngressRuleValue
method), 619
setUp() (kubernetes.test.test_v1beta1_ingress.TestV1beta1Ingress
method), 620
setUp() (kubernetes.test.test_v1beta1_ingress_backend.TestV1beta1IngressBackend
method), 620
setUp() (kubernetes.test.test_v1beta1_ingress_list.TestV1beta1IngressList
method), 620
setUp() (kubernetes.test.test_v1beta1_ingress_rule.TestV1beta1IngressRule
method), 621
setUp() (kubernetes.test.test_v1beta1_ingress_spec.TestV1beta1IngressSpec
method), 621
setUp() (kubernetes.test.test_v1beta1_ingress_status.TestV1beta1IngressStatus
method), 621
setUp() (kubernetes.test.test_v1beta1_ingress_tls.TestV1beta1IngressTLS
method), 622
setUp() (kubernetes.test.test_v1beta1_local_subject_access_review.TestV1beta1LocalSubjectAccessReview
method), 622
setUp() (kubernetes.test.test_v1beta1_network_policy.TestV1beta1NetworkPolicy
method), 623
setUp() (kubernetes.test.test_v1beta1_network_policy_ingress_rule.TestV1beta1NetworkPolicyIngressRule
method), 623
setUp() (kubernetes.test.test_v1beta1_network_policy_list.TestV1beta1NetworkPolicyList
method), 623
setUp() (kubernetes.test.test_v1beta1_network_policy_peer.TestV1beta1NetworkPolicyPeer
method), 624
setUp() (kubernetes.test.test_v1beta1_network_policy_port.TestV1beta1NetworkPolicyPort
method), 624
setUp() (kubernetes.test.test_v1beta1_network_policy_spec.TestV1beta1NetworkPolicySpec
method), 624
setUp() (kubernetes.test.test_v1beta1_non_resource_attributes.TestV1beta1NonResourceAttributes
method), 625
setUp() (kubernetes.test.test_v1beta1_pod_disruption_budget.TestV1beta1PodDisruptionBudget
method), 625
setUp() (kubernetes.test.test_v1beta1_pod_disruption_budget_list.TestV1beta1PodDisruptionBudgetList
method), 626
setUp() (kubernetes.test.test_v1beta1_pod_disruption_budget_spec.TestV1beta1PodDisruptionBudgetSpec
method), 626
setUp() (kubernetes.test.test_v1beta1_pod_disruption_budget_status.TestV1beta1PodDisruptionBudgetStatus
method), 626
setUp() (kubernetes.test.test_v1beta1_replica_set.TestV1beta1ReplicaSet
method), 627
setUp() (kubernetes.test.test_v1beta1_replica_set_condition.TestV1beta1ReplicaSetCondition
method), 627
setUp() (kubernetes.test.test_v1beta1_replica_set_list.TestV1beta1ReplicaSetList
method), 627
setUp() (kubernetes.test.test_v1beta1_replica_set_spec.TestV1beta1ReplicaSetSpec
method), 628
setUp() (kubernetes.test.test_v1beta1_replica_set_status.TestV1beta1ReplicaSetStatus
method), 628
setUp() (kubernetes.test.test_v1beta1_resource_attributes.TestV1beta1ResourceAttributes
method), 628
setUp() (kubernetes.test.test_v1beta1_self_subject_access_review.TestV1beta1SelfSubjectAccessReview
method), 629
setUp() (kubernetes.test.test_v1beta1_self_subject_access_review_spec.TestV1beta1SelfSubjectAccessReviewSpec
method), 629
setUp() (kubernetes.test.test_v1beta1_stateful_set.TestV1beta1StatefulSet
method), 630
setUp() (kubernetes.test.test_v1beta1_stateful_set_list.TestV1beta1StatefulSetList
method), 630
setUp() (kubernetes.test.test_v1beta1_stateful_set_spec.TestV1beta1StatefulSetSpec
method), 630
setUp() (kubernetes.test.test_v1beta1_stateful_set_status.TestV1beta1StatefulSetStatus
method), 631
setUp() (kubernetes.test.test_v1beta1_storage_class.TestV1beta1StorageClass
method), 631
setUp() (kubernetes.test.test_v1beta1_storage_class_list.TestV1beta1StorageClassList
method), 631

```

| | |
|---|--|
| setUp() (kubernetes.test.test_v1beta1_subject_access_review. method), 632 | TestV1beta1SubjectAccessReview spec (kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec attribute), 440 |
| setUp() (kubernetes.test.test_v1beta1_subject_access_review. method), 632 | TestV1beta1SubjectAccessReviewSpec spec (kubernetes.client.models.v1_replication_controller.V1ReplicationCon troller attribute), 427 |
| setUp() (kubernetes.test.test_v1beta1_subject_access_review. method), 633 | TestV1beta1SubjectAccessReviewStatus spec (kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute), 425 |
| setUp() (kubernetes.test.test_v1beta1_token_review. method), 633 | TestV1beta1TokenReview spec (kubernetes.client.models.v1_scale.V1Scale at- tribute), 426 |
| setUp() (kubernetes.test.test_v1beta1_token_review_spec. method), 633 | TestV1beta1TokenReviewSpec spec (kubernetes.client.models.v1_service.V1Service at- tribute), 426 |
| setUp() (kubernetes.test.test_v1beta1_token_review_status. method), 634 | TestV1beta1TokenReviewStatus spec (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet attribute), 426 |
| setUp() (kubernetes.test.test_v1beta1_user_info. method), 634 | TestV1beta1UserInfo spec (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute), 466 |
| setUp() (kubernetes.test.test_v2alpha1_cron_job. method), 634 | TestV2alpha1CronJob spec (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute), 474 |
| setUp() (kubernetes.test.test_v2alpha1_cron_job_list. method), 635 | TestV2alpha1CronJobList spec (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute), 480 |
| setUp() (kubernetes.test.test_v2alpha1_cron_job_spec. method), 635 | TestV2alpha1CronJobSpec spec (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 481 |
| setUp() (kubernetes.test.test_v2alpha1_cron_job_status. method), 635 | TestV2alpha1CronJobStatus spec (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 483 |
| setUp() (kubernetes.test.test_v2alpha1_job_template_spec. method), 636 | TestV2alpha1JobTemplateSpec spec (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 493 |
| setUp() (kubernetes.test.test_version_api. method), 636 | TestVersionApi spec (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute), 499 |
| setUp() (kubernetes.test.test_version_info. method), 637 | TestVersionInfo spec (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview attribute), 501 |
| share_name (kubernetes.client.models.v1_azure_file_volume. source attribute), 302 | V1AzureFileVolumeSource spec (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview attribute), 426 |
| signal (kubernetes.client.models.v1_container_state_terminated. attribute), 320 | V1ContainerStateTerminated spec (kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob attribute), 516 |
| SimpleNamespace (class in kubernetes.watch.watch), 637 | |
| size_bytes (kubernetes.client.models.v1_container_image. V1ContainerImage attribute), 316 | V1ContainerImage spec (kubernetes.client.models.v2alpha1_job_template_spec.V2alpha1JobTemplateSpec attribute), 520 |
| size_limit (kubernetes.client.models.v1_empty_dir_volume. source attribute), 327 | V1EmptyDirVolumeSource spec (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 358 |
| source (kubernetes.client.models.v1_event. V1Event attribute), 336 | V1Event start_time (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 410 |
| spec (kubernetes.client.models.v1_horizontal_pod_autoscaler. V1HorizontalPodAutoscaler attribute), 345 | V1HorizontalPodAutoscaler spec (kubernetes.client.models.v1_container_state_running.V1ContainerStateRunning attribute), 319 |
| spec (kubernetes.client.models.v1_job. V1Job attribute), 354 | V1Job started_at (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated attribute), 320 |
| spec (kubernetes.client.models.v1_limit_range. V1LimitRange attribute), 361 | V1LimitRange starting_deadline_seconds (kuber- netes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 519 |
| spec (kubernetes.client.models.v1_namespace. V1Namespace attribute), 366 | V1Namespace state (kubernetes.client.models.v1_container_status.V1ContainerStatus attribute), 322 |
| spec (kubernetes.client.models.v1_node. V1Node attribute), 370 | V1Node status (kubernetes.client.models.v1_component_condition.V1ComponentCondition attribute), 306 |
| spec (kubernetes.client.models.v1_persistent_volume. V1PersistentVolume attribute), 386 | V1PersistentVolume status (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler attribute), 427 |
| spec (kubernetes.client.models.v1_persistent_volume_claim. V1PersistentVolumeClaim attribute), 388 | V1PersistentVolumeClaim status (kubernetes.client.models.v1_job.V1Job attribute), 354 |
| spec (kubernetes.client.models.v1_pod. V1Pod attribute), 354 | V1Pod |

status (kubernetes.client.models.v1_job_condition.V1JobCondition attribute), 355

status (kubernetes.client.models.v1_namespace.V1Namespace attribute), 366

status (kubernetes.client.models.v1_node.V1Node attribute), 370

status (kubernetes.client.models.v1_node_condition.V1NodeCondition attribute), 372

status (kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute), 386

status (kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute), 388

status (kubernetes.client.models.v1_pod.V1Pod attribute), 400

status (kubernetes.client.models.v1_pod_condition.V1PodCondition attribute), 401

status (kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute), 419

status (kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition attribute), 420

status (kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute), 425

status (kubernetes.client.models.v1_scale.V1Scale attribute), 430

status (kubernetes.client.models.v1_service.V1Service attribute), 438

status (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet attribute), 466

status (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress attribute), 474

status (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview attribute), 480

status (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 487

status (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 492

status (kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition attribute), 493

status (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 499

status (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute), 501

status (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview attribute), 510

status (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview attribute), 513

status (kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob attribute), 516

stdin (kubernetes.client.models.v1_container.V1Container attribute), 315

stdin_once (kubernetes.client.models.v1_container.V1Container attribute), 315

stop() (kubernetes.watch.watch.Watch method), 637

storage_class_name (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 390

storage_class_name (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 397

storage_policy_id (kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource attribute), 453

storage_policy_name (kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource attribute), 453

StorageApi (class in kubernetes.client.apis.storage_api), 291

storageos (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 398

storageos (kubernetes.client.models.v1_volume.V1Volume attribute), 398

StorageV1beta1Api (class in kubernetes.client.apis.storage_v1beta1), 291

stream() (kubernetes.watch.watch.Watch method), 637

stream_data (kubernetes.client.models.v1_secret.V1Secret attribute), 433

sub_path (kubernetes.client.models.v1_volume_mount.V1VolumeMount attribute), 452

subdomain (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 408

subPaths (kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding attribute), 456

subjects (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding attribute), 461

subresource_volumes (kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes attribute), 498

subresource_volumes (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 331

subresource_volumes (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 359

subresource_volumes (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 493

subresource_volumes (kubernetes.client.models.v1_probe.V1Probe attribute), 404

successful_jobs_history_limit (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 519

supply_minimal_subjects (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 404

suspend (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec attribute), 519

swagger_types (kubernetes.client.models.runtime_raw_extension.RuntimeRawExtension attribute), 297

swagger_types (kubernetes.client.models.v1_attached_volume.V1AttachedVolume attribute), 299

swagger_types (kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AwsElasticBlockStoreVolumeSource attribute), 300

swagger_types (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource attribute), 300

attribute), 301
 swagger_types (kubernetes.client.models.v1_azure_file_volume_source.V1AzureFileVolumeSource, attribute), 302
 swagger_types (kubernetes.client.models.v1_binding.V1Binding, attribute), 303
 swagger_types (kubernetes.client.models.v1_capabilities.V1Capabilities, attribute), 303
 swagger_types (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource, attribute), 305
 swagger_types (kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource, attribute), 305
 swagger_types (kubernetes.client.models.v1_component_condition.V1ComponentCondition, attribute), 306
 swagger_types (kubernetes.client.models.v1_component_status.V1ComponentStatus, attribute), 307
 swagger_types (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList, attribute), 308
 swagger_types (kubernetes.client.models.v1_config_map.V1ConfigMap, attribute), 309
 swagger_types (kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector, attribute), 310
 swagger_types (kubernetes.client.models.v1_config_map_list.V1ConfigMapList, attribute), 311
 swagger_types (kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource, attribute), 312
 swagger_types (kubernetes.client.models.v1_container.V1Container, attribute), 315
 swagger_types (kubernetes.client.models.v1_container_image_source.V1ContainerImageSource, attribute), 317
 swagger_types (kubernetes.client.models.v1_container_ports.V1ContainerPort, attribute), 318
 swagger_types (kubernetes.client.models.v1_container_state.V1ContainerState, attribute), 318
 swagger_types (kubernetes.client.models.v1_container_state_waiting.V1ContainerStateWaiting, attribute), 319
 swagger_types (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated, attribute), 320
 swagger_types (kubernetes.client.models.v1_container_state_running.V1ContainerStateRunning, attribute), 321
 swagger_types (kubernetes.client.models.v1_container_status.V1ContainerStatus, attribute), 322
 swagger_types (kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference, attribute), 323
 swagger_types (kubernetes.client.models.v1_daemon_endpoint.V1DaemonEndpoint, attribute), 324
 swagger_types (kubernetes.client.models.v1_delete_options.V1DeleteOptions, attribute), 325
 swagger_types (kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource, attribute), 326
 swagger_types (kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile, attribute), 327
 swagger_types (kubernetes.client.models.v1_empty_dir_volume_source.V1EmptyDirVolumeSource, attribute), 328
 swagger_types (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress, attribute), 328
 swagger_types (kubernetes.client.models.v1_endpoint_port.V1EndpointPort, attribute), 329
 swagger_types (kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset, attribute), 330
 swagger_types (kubernetes.client.models.v1_endpoints.V1Endpoints, attribute), 331
 swagger_types (kubernetes.client.models.v1_endpoints_list.V1EndpointsList, attribute), 332
 swagger_types (kubernetes.client.models.v1_env_var.V1EnvVar, attribute), 333
 swagger_types (kubernetes.client.models.v1_env_var_source.V1EnvVarSource, attribute), 334
 swagger_types (kubernetes.client.models.v1_event.V1Event, attribute), 336
 swagger_types (kubernetes.client.models.v1_event_list.V1EventList, attribute), 337
 swagger_types (kubernetes.client.models.v1_event_source.V1EventSource, attribute), 337
 swagger_types (kubernetes.client.models.v1_exec_action.V1ExecAction, attribute), 338
 swagger_types (kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource, attribute), 339
 swagger_types (kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource, attribute), 340
 swagger_types (kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource, attribute), 341
 swagger_types (kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource, attribute), 342
 swagger_types (kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource, attribute), 343
 swagger_types (kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterFSVolumeSource, attribute), 343
 swagger_types (kubernetes.client.models.v1_handler.V1Handler, attribute), 344
 swagger_types (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler, attribute), 345
 swagger_types (kubernetes.client.models.v1_horizontal_pod_autoscaler_list.V1HorizontalPodAutoscalerList, attribute), 346
 swagger_types (kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec, attribute), 347
 swagger_types (kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource, attribute), 349
 swagger_types (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction, attribute), 350
 swagger_types (kubernetes.client.models.v1_http_header.V1HTTPHeader, attribute), 351
 swagger_types (kubernetes.client.models.v1_iiscsi_volume_source.V1ISCSIVolumeSource, attribute), 353
 swagger_types (kubernetes.client.models.v1_job.V1Job, attribute), 354
 swagger_types (kubernetes.client.models.v1_job_condition.V1JobCondition, attribute), 354

attribute), 355
 swagger_types (kubernetes.client.models.v1_job_list.V1JobList attribute), 356
 swagger_types (kubernetes.client.models.v1_job_spec.V1JobSpec attribute), 357
 swagger_types (kubernetes.client.models.v1_job_status.V1JobStatus attribute), 359
 swagger_types (kubernetes.client.models.v1_key_to_path.V1KeyToPath attribute), 359
 swagger_types (kubernetes.client.models.v1_lifecycle.V1Lifecycle attribute), 360
 swagger_types (kubernetes.client.models.v1_limit_range.V1LimitRange attribute), 361
 swagger_types (kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem attribute), 362
 swagger_types (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList attribute), 363
 swagger_types (kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec attribute), 364
 swagger_types (kubernetes.client.models.v1_load_balancer.V1LoadBalancer attribute), 364
 swagger_types (kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus attribute), 365
 swagger_types (kubernetes.client.models.v1_local_object_reference.V1LocalObjectReference attribute), 366
 swagger_types (kubernetes.client.models.v1_namespace.V1Namespace attribute), 367
 swagger_types (kubernetes.client.models.v1_namespace_list.V1NamespaceList attribute), 368
 swagger_types (kubernetes.client.models.v1_namespace_spec.V1NamespaceSpec attribute), 368
 swagger_types (kubernetes.client.models.v1_namespace_status.V1NamespaceStatus attribute), 369
 swagger_types (kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource attribute), 370
 swagger_types (kubernetes.client.models.v1_node.V1Node attribute), 371
 swagger_types (kubernetes.client.models.v1_node_address.V1NodeAddress attribute), 371
 swagger_types (kubernetes.client.models.v1_node_condition.V1NodeCondition attribute), 372
 swagger_types (kubernetes.client.models.v1_node_daemon_set.V1NodeDaemonSet attribute), 373
 swagger_types (kubernetes.client.models.v1_node_list.V1NodeList attribute), 374
 swagger_types (kubernetes.client.models.v1_node_spec.V1NodeSpec attribute), 375
 swagger_types (kubernetes.client.models.v1_node_status.V1NodeStatus attribute), 376
 swagger_types (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378
 swagger_types (kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector attribute), 379
 swagger_types (kubernetes.client.models.v1_object_meta.V1ObjectMeta attribute), 382
 swagger_types (kubernetes.client.models.v1_object_reference.V1ObjectReference attribute), 384
 swagger_types (kubernetes.client.models.v1_owner_reference.V1OwnerReference attribute), 385
 swagger_types (kubernetes.client.models.v1_persistent_volume.V1PersistentVolume attribute), 387
 swagger_types (kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim attribute), 388
 swagger_types (kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList attribute), 389
 swagger_types (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec attribute), 390
 swagger_types (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus attribute), 391
 swagger_types (kubernetes.client.models.v1_persistent_volume_claim_volume_source.V1PersistentVolumeClaimVolumeSource attribute), 392
 swagger_types (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList attribute), 393
 swagger_types (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec attribute), 398
 swagger_types (kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus attribute), 399
 swagger_types (kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource attribute), 399
 swagger_types (kubernetes.client.models.v1_pod.V1Pod attribute), 400
 swagger_types (kubernetes.client.models.v1_pod_condition.V1PodCondition attribute), 401
 swagger_types (kubernetes.client.models.v1_pod_list.V1PodList attribute), 403
 swagger_types (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext attribute), 404
 swagger_types (kubernetes.client.models.v1_pod_spec.V1PodSpec attribute), 408
 swagger_types (kubernetes.client.models.v1_pod_status.V1PodStatus attribute), 410
 swagger_types (kubernetes.client.models.v1_pod_template.V1PodTemplate attribute), 411
 swagger_types (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList attribute), 412
 swagger_types (kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec attribute), 413
 swagger_types (kubernetes.client.models.v1_preconditions.V1Preconditions attribute), 413
 swagger_types (kubernetes.client.models.v1_probe.V1Probe attribute), 415
 swagger_types (kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource attribute), 416
 swagger_types (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource attribute), 418
 swagger_types (kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute), 419
 swagger_types (kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition attribute), 419

attribute), 420
 swagger_types (kubernetes.client.models.v1_replication_controller.V1ReplicationController attribute), 421
 swagger_types (kubernetes.client.models.v1_replication_controller.V1ReplicationControllerSpec attribute), 422
 swagger_types (kubernetes.client.models.v1_replication_controller.V1ReplicationControllerStatus attribute), 423
 swagger_types (kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector attribute), 424
 swagger_types (kubernetes.client.models.v1_resource_quota.V1ResourceQuota attribute), 425
 swagger_types (kubernetes.client.models.v1_resource_quota.V1ResourceQuotaList attribute), 426
 swagger_types (kubernetes.client.models.v1_resource_quota.V1ResourceQuotaSpec attribute), 427
 swagger_types (kubernetes.client.models.v1_resource_quota.V1ResourceQuotaStatus attribute), 428
 swagger_types (kubernetes.client.models.v1_resource_requirement.V1ResourceRequirement attribute), 429
 swagger_types (kubernetes.client.models.v1_scale.V1Scale attribute), 430
 swagger_types (kubernetes.client.models.v1_scale_spec.V1ScaleSpec attribute), 430
 swagger_types (kubernetes.client.models.v1_scale_status.V1ScaleStatus attribute), 431
 swagger_types (kubernetes.client.models.v1_se_linux_options.V1SELinuxOptions attribute), 432
 swagger_types (kubernetes.client.models.v1_secret.V1Secret attribute), 433
 swagger_types (kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector attribute), 434
 swagger_types (kubernetes.client.models.v1_secret_list.V1SecretList attribute), 435
 swagger_types (kubernetes.client.models.v1_secret_volumes.V1SecretVolumes attribute), 436
 swagger_types (kubernetes.client.models.v1_security_context.V1SecurityContext attribute), 437
 swagger_types (kubernetes.client.models.v1_service.V1Service attribute), 438
 swagger_types (kubernetes.client.models.v1_service_account.V1ServiceAccount attribute), 440
 swagger_types (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList attribute), 441
 swagger_types (kubernetes.client.models.v1_service_list.V1ServiceList attribute), 442
 swagger_types (kubernetes.client.models.v1_service_port.V1ServicePort attribute), 443
 swagger_types (kubernetes.client.models.v1_service_spec.V1ServiceSpec attribute), 446
 swagger_types (kubernetes.client.models.v1_service_status.V1ServiceStatus attribute), 447
 swagger_types (kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction attribute), 447
 swagger_types (kubernetes.client.models.v1_volume.V1Volume attribute), 451
 swagger_types (kubernetes.client.models.v1_volume_mount.V1VolumeMount attribute), 452
 swagger_types (kubernetes.client.models.v1_volume_resource.V1VolumeResource attribute), 453
 swagger_types (kubernetes.client.models.v1_volume_resource_list.V1VolumeResourceList attribute), 455
 swagger_types (kubernetes.client.models.v1_volume_resource_selector.V1VolumeResourceSelector attribute), 456
 swagger_types (kubernetes.client.models.v1_volume_resource_spec.V1VolumeResourceSpec attribute), 457
 swagger_types (kubernetes.client.models.v1_volume_resource_status.V1VolumeResourceStatus attribute), 458
 swagger_types (kubernetes.client.models.v1_volume_resource_status_list.V1VolumeResourceStatusList attribute), 459
 swagger_types (kubernetes.client.models.v1_volume_resource_status_spec.V1VolumeResourceStatusSpec attribute), 460
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status.V1VolumeResourceStatusStatus attribute), 461
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_list.V1VolumeResourceStatusStatusList attribute), 462
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_spec.V1VolumeResourceStatusStatusSpec attribute), 463
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status.V1VolumeResourceStatusStatusStatus attribute), 464
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_list.V1VolumeResourceStatusStatusStatusList attribute), 465
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_spec.V1VolumeResourceStatusStatusStatusSpec attribute), 466
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status.V1VolumeResourceStatusStatusStatusStatus attribute), 467
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_list.V1VolumeResourceStatusStatusStatusStatusList attribute), 468
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_spec.V1VolumeResourceStatusStatusStatusStatusSpec attribute), 470
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status.V1VolumeResourceStatusStatusStatusStatusStatus attribute), 472
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_list.V1VolumeResourceStatusStatusStatusStatusStatusList attribute), 473
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_spec.V1VolumeResourceStatusStatusStatusStatusStatusSpec attribute), 473
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status.V1VolumeResourceStatusStatusStatusStatusStatusStatus attribute), 474
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_list.V1VolumeResourceStatusStatusStatusStatusStatusStatusList attribute), 475
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_spec.V1VolumeResourceStatusStatusStatusStatusStatusStatusSpec attribute), 476
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_status.V1VolumeResourceStatusStatusStatusStatusStatusStatusStatus attribute), 477
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_status_list.V1VolumeResourceStatusStatusStatusStatusStatusStatusStatusList attribute), 477
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_status_spec.V1VolumeResourceStatusStatusStatusStatusStatusStatusStatusSpec attribute), 478
 swagger_types (kubernetes.client.models.v1_volume_resource_status_status_status_status_status_status_status_status.V1VolumeResourceStatusStatusStatusStatusStatusStatusStatusStatus attribute), 478

attribute), 479
 swagger_types (kubernetes.client.models.v1beta1_local_subject_access_review_spec.V1beta1LocalSubjectAccessReviewSpec attribute), 480
 swagger_types (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy attribute), 481
 swagger_types (kubernetes.client.models.v1beta1_network_policy_entry.V1beta1NetworkPolicyEntry attribute), 482
 swagger_types (kubernetes.client.models.v1beta1_network_policy_rule.V1beta1NetworkPolicyRule attribute), 483
 swagger_types (kubernetes.client.models.v1beta1_network_policy_rule.V1beta1NetworkPolicyRule attribute), 484
 swagger_types (kubernetes.client.models.v1beta1_network_policy_rule.V1beta1NetworkPolicyRule attribute), 484
 swagger_types (kubernetes.client.models.v1beta1_network_policy_rule.V1beta1NetworkPolicyRule attribute), 486
 swagger_types (kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes attribute), 486
 swagger_types (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 487
 swagger_types (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 488
 swagger_types (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 489
 swagger_types (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget attribute), 491
 swagger_types (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet attribute), 492
 swagger_types (kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition attribute), 493
 swagger_types (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList attribute), 494
 swagger_types (kubernetes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec attribute), 495
 swagger_types (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus attribute), 496
 swagger_types (kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes attribute), 498
 swagger_types (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview attribute), 499
 swagger_types (kubernetes.client.models.v1beta1_self_subject_access_review_spec.V1beta1SelfSubjectAccessReviewSpec attribute), 500
 swagger_types (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet attribute), 501
 swagger_types (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList attribute), 502
 swagger_types (kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec attribute), 504
 swagger_types (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus attribute), 506
 swagger_types (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass attribute), 508
 swagger_types (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList attribute), 509
 swagger_types (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview attribute), 510
 swagger_types (kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec attribute), 511
 swagger_types (kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec attribute), 512
 swagger_types (kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec attribute), 513
 swagger_types (kubernetes.client.models.v1beta1_token_review_spec.V1beta1TokenReviewSpec attribute), 514
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 514
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 515
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 516
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 517
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 519
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 520
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 520
 swagger_types (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus attribute), 522
 swagger_types (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo attribute), 378
 target (kubernetes.client.models.v1_node_spec.V1NodeSpec attribute), 375
 target (kubernetes.client.models.v1_binding.V1Binding attribute), 303
 target_cpu_utilization_percentage (kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec attribute), 340
 target_port (kubernetes.client.models.v1_service_port.V1ServicePort attribute), 495
 target_portal (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource attribute), 328
 target_ref (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress attribute), 328
 target_ww_ns (kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource attribute), 328
 tcp_socket (kubernetes.client.models.v1_handler.V1Handler attribute), 341
 tcp_socket (kubernetes.client.models.v1_probe.V1Probe attribute), 495
 tearDown() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 935
 tearDown() (kubernetes.config.kube_config_test.BaseTestCase method), 935
 tearDown() (kubernetes.test.test_apis_api.TestApisApi method), 935

tearDown() (kubernetes.test.test_apps_api.TestAppsApi method), 532

tearDown() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1Beta1Api method), 533

tearDown() (kubernetes.test.test_authentication_api.TestAuthenticationApi method), 535

tearDown() (kubernetes.test.test_authentication_v1beta1_api.TestAuthenticationV1Beta1Api method), 535

tearDown() (kubernetes.test.test_authorization_api.TestAuthorizationApi method), 535

tearDown() (kubernetes.test.test_authorization_v1beta1_api.TestAuthorizationV1Beta1Api method), 536

tearDown() (kubernetes.test.test_autoscaling_api.TestAutoscalingApi method), 536

tearDown() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537

tearDown() (kubernetes.test.test_batch_api.TestBatchApi method), 537

tearDown() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538

tearDown() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2Alpha1Api method), 539

tearDown() (kubernetes.test.test_certificates_api.TestCertificatesApi method), 540

tearDown() (kubernetes.test.test_core_api.TestCoreApi method), 540

tearDown() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 540

tearDown() (kubernetes.test.test_extensions_api.TestExtensionsApi method), 552

tearDown() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1Beta1Api method), 552

tearDown() (kubernetes.test.test_logs_api.TestLogsApi method), 556

tearDown() (kubernetes.test.test_policy_api.TestPolicyApi method), 556

tearDown() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1Beta1Api method), 557

tearDown() (kubernetes.test.test_rbac_authorization_api.TestRbacAuthorizationApi method), 558

tearDown() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1Alpha1Api method), 558

tearDown() (kubernetes.test.test_runtime_raw_extension.TestRuntimeRawExtension method), 560

tearDown() (kubernetes.test.test_storage_api.TestStorageApi method), 560

tearDown() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1Beta1Api method), 561

tearDown() (kubernetes.test.test_v1_attached_volume.TestV1AttachedVolume method), 562

tearDown() (kubernetes.test.test_v1_aws_elastic_block_store.TestV1AwsElasticBlockStore method), 563

tearDown() (kubernetes.test.test_v1_azure_disk_volume_source.TestV1AzureDiskVolumeSource method), 563

tearDown() (kubernetes.test.test_v1_azure_file_volume_source.TestV1AzureFileVolumeSource method), 563

tearDown() (kubernetes.test.test_v1_binding.TestV1Binding method), 564

tearDown() (kubernetes.test.test_v1_capabilities.TestV1Capabilities method), 564

tearDown() (kubernetes.test.test_v1_ceph_fs_volume_source.TestV1CephFsVolumeSource method), 565

tearDown() (kubernetes.test.test_v1_cinder_volume_source.TestV1CinderVolumeSource method), 565

tearDown() (kubernetes.test.test_v1_component_condition.TestV1ComponentCondition method), 565

tearDown() (kubernetes.test.test_v1_component_status.TestV1ComponentStatus method), 566

tearDown() (kubernetes.test.test_v1_component_status_list.TestV1ComponentStatusList method), 566

tearDown() (kubernetes.test.test_v1_config_map.TestV1ConfigMap method), 566

tearDown() (kubernetes.test.test_v1_config_map_key_selector.TestV1ConfigMapKeySelector method), 567

tearDown() (kubernetes.test.test_v1_config_map_list.TestV1ConfigMapList method), 567

tearDown() (kubernetes.test.test_v1_config_map_volume_source.TestV1ConfigMapVolumeSource method), 567

tearDown() (kubernetes.test.test_v1_container.TestV1Container method), 568

tearDown() (kubernetes.test.test_v1_container_image.TestV1ContainerImage method), 568

tearDown() (kubernetes.test.test_v1_container_port.TestV1ContainerPort method), 568

tearDown() (kubernetes.test.test_v1_container_state.TestV1ContainerState method), 569

tearDown() (kubernetes.test.test_v1_container_state_running.TestV1ContainerStateRunning method), 569

tearDown() (kubernetes.test.test_v1_container_state_terminated.TestV1ContainerStateTerminated method), 570

tearDown() (kubernetes.test.test_v1_container_state_waiting.TestV1ContainerStateWaiting method), 570

tearDown() (kubernetes.test.test_v1_container_status.TestV1ContainerStatus method), 570

tearDown() (kubernetes.test.test_v1_daemon_endpoint.TestV1DaemonEndpoint method), 571

tearDown() (kubernetes.test.test_v1_delete_options.TestV1DeleteOptions method), 571

tearDown() (kubernetes.test.test_v1_downward_api_volume_file.TestV1DownwardApiVolumeFile method), 572

tearDown() (kubernetes.test.test_v1_downward_api_volume_source.TestV1DownwardApiVolumeSource method), 572

tearDown() (kubernetes.test.test_v1_emptier_block_store_volume_source.TestV1EmptyDirVolumeSource method), 572

tearDown() (kubernetes.test.test_v1_endpoint_address.TestV1EndpointAddress method), 573

tearDown() (kubernetes.test.test_v1_endpoint_port.TestV1EndpointPort method), 573

tearDown() (kubernetes.test.test_v1_endpoint_subset.TestV1EndpointSubset method), 573

tearDown() (kubernetes.test.test_v1_endpoints.TestV1Endpoints method), 574

tearDown() (kubernetes.test.test_v1_endpoints_list.TestV1EndpointsList method), 574

tearDown() (kubernetes.test.test_v1_env_var.TestV1EnvVar method), 575

tearDown() (kubernetes.test.test_v1_env_var_source.TestV1EnvVarSource method), 575

tearDown() (kubernetes.test.test_v1_event.TestV1Event method), 575

tearDown() (kubernetes.test.test_v1_event_list.TestV1EventList method), 576

tearDown() (kubernetes.test.test_v1_event_source.TestV1EventSource method), 576

tearDown() (kubernetes.test.test_v1_exec_action.TestV1ExecAction method), 576

tearDown() (kubernetes.test.test_v1_fc_volume_source.TestV1FCVolumeSource method), 577

tearDown() (kubernetes.test.test_v1_flex_volume_source.TestV1FlexVolumeSource method), 577

tearDown() (kubernetes.test.test_v1_flocker_volume_source.TestV1FlockerVolumeSource method), 577

tearDown() (kubernetes.test.test_v1_gce_persistent_disk_volume_source.TestV1GCEPersistentDiskVolumeSource method), 578

tearDown() (kubernetes.test.test_v1_git_repo_volume_source.TestV1GitRepoVolumeSource method), 578

tearDown() (kubernetes.test.test_v1_glusterfs_volume_source.TestV1GlusterfsVolumeSource method), 578

tearDown() (kubernetes.test.test_v1_handler.TestV1Handler method), 579

tearDown() (kubernetes.test.test_v1_horizontal_pod_autoscaler.TestV1HorizontalPodAutoscaler method), 579

tearDown() (kubernetes.test.test_v1_horizontal_pod_autoscaler_list.TestV1HorizontalPodAutoscalerList method), 580

tearDown() (kubernetes.test.test_v1_horizontal_pod_autoscaler_spec.TestV1HorizontalPodAutoscalerSpec method), 580

tearDown() (kubernetes.test.test_v1_horizontal_pod_autoscaler_status.TestV1HorizontalPodAutoscalerStatus method), 580

tearDown() (kubernetes.test.test_v1_host_path_volume_source.TestV1HostPathVolumeSource method), 581

tearDown() (kubernetes.test.test_v1_http_get_action.TestV1HTTPGetAction method), 581

tearDown() (kubernetes.test.test_v1_http_header.TestV1HTTPHeader method), 581

tearDown() (kubernetes.test.test_v1_iscsi_volume_source.TestV1ISCSIVolumeSource method), 582

tearDown() (kubernetes.test.test_v1_job.TestV1Job method), 582

tearDown() (kubernetes.test.test_v1_job_condition.TestV1JobCondition method), 582

tearDown() (kubernetes.test.test_v1_job_list.TestV1JobList method), 583

tearDown() (kubernetes.test.test_v1_job_spec.TestV1JobSpec method), 583

tearDown() (kubernetes.test.test_v1_job_status.TestV1JobStatus method), 583

tearDown() (kubernetes.test.test_v1_key_to_path.TestV1KeyToPath method), 584

tearDown() (kubernetes.test.test_v1_lifecycle.TestV1Lifecycle method), 584

tearDown() (kubernetes.test.test_v1_limit_range.TestV1LimitRange method), 585

tearDown() (kubernetes.test.test_v1_limit_range_item.TestV1LimitRangeItem method), 585

tearDown() (kubernetes.test.test_v1_limit_range_list.TestV1LimitRangeList method), 585

tearDown() (kubernetes.test.test_v1_limit_range_spec.TestV1LimitRangeSpec method), 586

tearDown() (kubernetes.test.test_v1_load_balancer_ingress.TestV1LoadBalancerIngress method), 586

tearDown() (kubernetes.test.test_v1_load_balancer_status.TestV1LoadBalancerStatus method), 586

tearDown() (kubernetes.test.test_v1_local_object_reference.TestV1LocalObjectReference method), 587

tearDown() (kubernetes.test.test_v1_namespace.TestV1Namespace method), 587

tearDown() (kubernetes.test.test_v1_namespace_list.TestV1NamespaceList method), 587

tearDown() (kubernetes.test.test_v1_namespace_spec.TestV1NamespaceSpec method), 588

tearDown() (kubernetes.test.test_v1_namespace_status.TestV1NamespaceStatus method), 588

tearDown() (kubernetes.test.test_v1_nfs_volume_source.TestV1NFSVolumeSource method), 588

tearDown() (kubernetes.test.test_v1_node.TestV1Node method), 589

tearDown() (kubernetes.test.test_v1_node_address.TestV1NodeAddress method), 589

tearDown() (kubernetes.test.test_v1_node_condition.TestV1NodeCondition method), 590

tearDown() (kubernetes.test.test_v1_node_list.TestV1NodeList method), 590

tearDown() (kubernetes.test.test_v1_node_spec.TestV1NodeSpec method), 591

tearDown() (kubernetes.test.test_v1_node_status.TestV1NodeStatus method), 591

tearDown() (kubernetes.test.test_v1_node_system_info.TestV1NodeSystemInfo method), 591

tearDown() (kubernetes.test.test_v1_object_field_selector.TestV1ObjectFieldSelector method), 592

tearDown() (kubernetes.test.test_v1_object_meta.TestV1ObjectMeta method), 592

[tearDown\(\) \(kubernetes.test.test_v1_object_reference.TestV1ObjectReferenceReference method\), 592](#)
[tearDown\(\) \(kubernetes.test.test_v1_owner_reference.TestV1OwnerReferenceReference method\), 593](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume.TestV1PersistentVolume method\), 593](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_claim.TestV1PersistentVolumeClaim method\), 593](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_claim_resource_quota.TestV1PersistentVolumeClaimResourceQuota method\), 594](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_claim_resource_quota_list.TestV1PersistentVolumeClaimResourceQuotaList method\), 594](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_claim_resource_quota_spec.TestV1PersistentVolumeClaimResourceQuotaSpec method\), 595](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_claim_volume_status.TestV1PersistentVolumeClaimVolumeStatus method\), 595](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_list.TestV1PersistentVolumeList method\), 595](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_spec.TestV1PersistentVolumeSpec method\), 596](#)
[tearDown\(\) \(kubernetes.test.test_v1_persistent_volume_status.TestV1PersistentVolumeStatus method\), 596](#)
[tearDown\(\) \(kubernetes.test.test_v1_photon_persistent_disk_test_v1_scale.TestV1PhotonPersistentDiskTestV1ScaleStatus method\), 596](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod.TestV1Pod method\), 597](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_condition.TestV1PodCondition method\), 597](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_list.TestV1PodList method\), 597](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_security_context.TestV1PodSecurityContext method\), 598](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_spec.TestV1PodSpec method\), 598](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_status.TestV1PodStatus method\), 598](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_template.TestV1PodTemplate method\), 599](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_template_list.TestV1PodTemplateList method\), 599](#)
[tearDown\(\) \(kubernetes.test.test_v1_pod_template_spec.TestV1PodTemplateSpec method\), 600](#)
[tearDown\(\) \(kubernetes.test.test_v1_preconditions.TestV1Preconditions method\), 600](#)
[tearDown\(\) \(kubernetes.test.test_v1_probe.TestV1Probe method\), 600](#)
[tearDown\(\) \(kubernetes.test.test_v1_quobyte_volume_source.TestV1QuobyteVolumeSource method\), 601](#)
[tearDown\(\) \(kubernetes.test.test_v1_rbd_volume_source.TestV1RBDVolumeSource method\), 601](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller.TestV1ReplicationController method\), 601](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller_condition.TestV1ReplicationControllerCondition method\), 602](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller_list.TestV1ReplicationControllerList method\), 602](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller_spec.TestV1ReplicationControllerSpec method\), 602](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller_status.TestV1ReplicationControllerStatus method\), 603](#)
[tearDown\(\) \(kubernetes.test.test_v1_replication_controller_status.TestV1ReplicationControllerStatus method\), 603](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_field_selector.TestV1ResourceFieldSelector method\), 603](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_quota.TestV1ResourceQuota method\), 603](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_quota_list.TestV1ResourceQuotaList method\), 604](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_quota_spec.TestV1ResourceQuotaSpec method\), 604](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_quota_spec.TestV1ResourceQuotaSpec method\), 605](#)
[tearDown\(\) \(kubernetes.test.test_v1_resource_requirements.TestV1ResourceRequirements method\), 605](#)
[tearDown\(\) \(kubernetes.test.test_v1_scale.TestV1Scale method\), 605](#)
[tearDown\(\) \(kubernetes.test.test_v1_scale_spec.TestV1ScaleSpec method\), 606](#)
[tearDown\(\) \(kubernetes.test.test_v1_secret.TestV1Secret method\), 607](#)
[tearDown\(\) \(kubernetes.test.test_v1_secret_key_selector.TestV1SecretKeySelector method\), 607](#)
[tearDown\(\) \(kubernetes.test.test_v1_secret_list.TestV1SecretList method\), 607](#)
[tearDown\(\) \(kubernetes.test.test_v1_secret_volume_source.TestV1SecretVolumeSource method\), 608](#)
[tearDown\(\) \(kubernetes.test.test_v1_security_context.TestV1SecurityContext method\), 608](#)
[tearDown\(\) \(kubernetes.test.test_v1_service.TestV1Service method\), 608](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_account.TestV1ServiceAccount method\), 609](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_account_list.TestV1ServiceAccountList method\), 609](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_list.TestV1ServiceList method\), 610](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_port.TestV1ServicePort method\), 610](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_spec.TestV1ServiceSpec method\), 610](#)
[tearDown\(\) \(kubernetes.test.test_v1_service_status.TestV1ServiceStatus method\), 611](#)
[tearDown\(\) \(kubernetes.test.test_v1_tcp_socket_action.TestV1TCPSocketAction method\), 611](#)
[tearDown\(\) \(kubernetes.test.test_v1_volume.TestV1Volume method\), 611](#)
[tearDown\(\) \(kubernetes.test.test_v1_volume_test_v1_scale.TestV1VolumeTestV1Scale method\), 611](#)

tearDown() (kubernetes.test.test_v1_volume_mount.TestV1VolumeMount method), 612

tearDown() (kubernetes.test.test_v1_vsphere_virtual_disk_volume.TestV1VsphereVirtualDiskVolume method), 612

tearDown() (kubernetes.test.test_v1alpha1_cluster_role.TestV1Alpha1ClusterRole method), 613

tearDown() (kubernetes.test.test_v1alpha1_cluster_role_binding.TestV1Alpha1ClusterRoleBinding method), 613

tearDown() (kubernetes.test.test_v1alpha1_cluster_role_binding_distribution.TestV1Alpha1ClusterRoleBindingDistribution method), 613

tearDown() (kubernetes.test.test_v1alpha1_cluster_role_list.TestV1Alpha1ClusterRoleList method), 614

tearDown() (kubernetes.test.test_v1alpha1_policy_rule.TestV1Alpha1PolicyRule method), 614

tearDown() (kubernetes.test.test_v1alpha1_role.TestV1Alpha1Role method), 614

tearDown() (kubernetes.test.test_v1alpha1_role_binding.TestV1Alpha1RoleBinding method), 615

tearDown() (kubernetes.test.test_v1alpha1_role_binding_list.TestV1Alpha1RoleBindingList method), 615

tearDown() (kubernetes.test.test_v1alpha1_role_list.TestV1Alpha1RoleList method), 615

tearDown() (kubernetes.test.test_v1alpha1_role_ref.TestV1Alpha1RoleRef method), 616

tearDown() (kubernetes.test.test_v1alpha1_subject.TestV1Alpha1Subject method), 616

tearDown() (kubernetes.test.test_v1beta1_daemon_set.TestV1Beta1DaemonSet method), 617

tearDown() (kubernetes.test.test_v1beta1_daemon_set_list.TestV1Beta1DaemonSetList method), 617

tearDown() (kubernetes.test.test_v1beta1_daemon_set_spec.TestV1Beta1DaemonSetSpec method), 617

tearDown() (kubernetes.test.test_v1beta1_daemon_set_status.TestV1Beta1DaemonSetStatus method), 618

tearDown() (kubernetes.test.test_v1beta1_eviction.TestV1Beta1Eviction method), 618

tearDown() (kubernetes.test.test_v1beta1_http_ingress_path.TestV1Beta1HttpIngressPath method), 619

tearDown() (kubernetes.test.test_v1beta1_http_ingress_rule.TestV1Beta1HttpIngressRule method), 619

tearDown() (kubernetes.test.test_v1beta1_ingress.TestV1Beta1Ingress method), 620

tearDown() (kubernetes.test.test_v1beta1_ingress_backend.TestV1Beta1IngressBackend method), 620

tearDown() (kubernetes.test.test_v1beta1_ingress_list.TestV1Beta1IngressList method), 620

tearDown() (kubernetes.test.test_v1beta1_ingress_rule.TestV1Beta1IngressRule method), 621

tearDown() (kubernetes.test.test_v1beta1_ingress_spec.TestV1Beta1IngressSpec method), 621

tearDown() (kubernetes.test.test_v1beta1_ingress_status.TestV1Beta1IngressStatus method), 621

tearDown() (kubernetes.test.test_v1beta1_ingress_tls.TestV1Beta1IngressTls method), 622

tearDown() (kubernetes.test.test_v1beta1_local_subject_access_review.TestV1Beta1LocalSubjectAccessReview method), 622

tearDown() (kubernetes.test.test_v1beta1_network_policy.TestV1Beta1NetworkPolicy method), 623

tearDown() (kubernetes.test.test_v1beta1_network_policy_ingress_rule.TestV1Beta1NetworkPolicyIngressRule method), 623

tearDown() (kubernetes.test.test_v1beta1_network_policy_list.TestV1Beta1NetworkPolicyList method), 623

tearDown() (kubernetes.test.test_v1beta1_network_policy_peer.TestV1Beta1NetworkPolicyPeer method), 624

tearDown() (kubernetes.test.test_v1beta1_network_policy_port.TestV1Beta1NetworkPolicyPort method), 624

tearDown() (kubernetes.test.test_v1beta1_network_policy_spec.TestV1Beta1NetworkPolicySpec method), 625

tearDown() (kubernetes.test.test_v1beta1_non_resource_attributes.TestV1Beta1NonResourceAttributes method), 625

tearDown() (kubernetes.test.test_v1beta1_pod_disruption_budget.TestV1Beta1PodDisruptionBudget method), 625

tearDown() (kubernetes.test.test_v1beta1_pod_disruption_budget_list.TestV1Beta1PodDisruptionBudgetList method), 626

tearDown() (kubernetes.test.test_v1beta1_pod_disruption_budget_spec.TestV1Beta1PodDisruptionBudgetSpec method), 626

tearDown() (kubernetes.test.test_v1beta1_pod_disruption_budget_status.TestV1Beta1PodDisruptionBudgetStatus method), 626

tearDown() (kubernetes.test.test_v1beta1_replica_set.TestV1Beta1ReplicaSet method), 627

tearDown() (kubernetes.test.test_v1beta1_replica_set_condition.TestV1Beta1ReplicaSetCondition method), 627

tearDown() (kubernetes.test.test_v1beta1_replica_set_list.TestV1Beta1ReplicaSetList method), 627

tearDown() (kubernetes.test.test_v1beta1_replica_set_spec.TestV1Beta1ReplicaSetSpec method), 628

tearDown() (kubernetes.test.test_v1beta1_replica_set_status.TestV1Beta1ReplicaSetStatus method), 628

tearDown() (kubernetes.test.test_v1beta1_resource_attributes.TestV1Beta1ResourceAttributes method), 628

tearDown() (kubernetes.test.test_v1beta1_self_subject_access_review.TestV1Beta1SelfSubjectAccessReview method), 629

tearDown() (kubernetes.test.test_v1beta1_self_subject_access_review_spec.TestV1Beta1SelfSubjectAccessReviewSpec method), 629

tearDown() (kubernetes.test.test_v1beta1_stateful_set.TestV1Beta1StatefulSet method), 630

tearDown() (kubernetes.test.test_v1beta1_stateful_set_list.TestV1Beta1StatefulSetList method), 630

tearDown() (kubernetes.test.test_v1beta1_stateful_set_spec.TestV1Beta1StatefulSetSpec method), 630

tearDown() (kubernetes.test.test_v1beta1_stateful_set_status.TestV1Beta1StatefulSetStatus method), 631

tearDown() (kubernetes.test.test_v1beta1_storage_class.TestV1Beta1StorageClass method), 631

tearDown() (kubernetes.test.test_v1beta1_storage_class_list.TestV1Beta1StorageClassList method), 632

tearDown() (kubernetes.test.test_v1beta1_subject_access_review.TestV1Beta1SubjectAccessReview method), 632

[tearDown\(\) \(kubernetes.test.test_v1beta1_subject_access_review_spec.TestV1beta1SubjectAccessReviewSpec method\), 632](#)
[tearDown\(\) \(kubernetes.test.test_v1beta1_subject_access_review_status.TestV1beta1SubjectAccessReviewStatus method\), 633](#)
[tearDown\(\) \(kubernetes.test.test_v1beta1_token_review.TestV1beta1TokenReview method\), 633](#)
[tearDown\(\) \(kubernetes.test.test_v1beta1_token_review_spec.TestV1beta1TokenReviewSpec method\), 633](#)
[tearDown\(\) \(kubernetes.test.test_v1beta1_token_review_status.TestV1beta1TokenReviewStatus method\), 634](#)
[tearDown\(\) \(kubernetes.test.test_v1beta1_user_info.TestV1beta1UserInfo method\), 634](#)
[tearDown\(\) \(kubernetes.test.test_v2alpha1_cron_job.TestV2alpha1CronJob method\), 634](#)
[tearDown\(\) \(kubernetes.test.test_v2alpha1_cron_job_list.TestV2alpha1CronJobList method\), 635](#)
[tearDown\(\) \(kubernetes.test.test_v2alpha1_cron_job_spec.TestV2alpha1CronJobSpec method\), 635](#)
[tearDown\(\) \(kubernetes.test.test_v2alpha1_cron_job_status.TestV2alpha1CronJobStatus method\), 636](#)
[tearDown\(\) \(kubernetes.test.test_v2alpha1_job_template_spec.TestV2alpha1JobTemplateSpec method\), 636](#)
[tearDown\(\) \(kubernetes.test.test_version_api.TestVersionApi method\), 636](#)
[tearDown\(\) \(kubernetes.test.test_version_info.TestVersionInfo method\), 637](#)
[template \(kubernetes.client.models.v1_job_spec.V1JobSpec attribute\), 357](#)
[template \(kubernetes.client.models.v1_pod_template.V1PodTemplate attribute\), 411](#)
[template \(kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec attribute\), 422](#)
[template \(kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec attribute\), 468](#)
[template \(kubernetes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec attribute\), 495](#)
[template \(kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec attribute\), 504](#)
[template_generation \(kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec attribute\), 468](#)
[terminated \(kubernetes.client.models.v1_container_state.V1ContainerState attribute\), 318](#)
[termination_grace_period_seconds \(kubernetes.client.models.v1_pod_spec.V1PodSpec attribute\), 408](#)
[termination_message_path \(kubernetes.client.models.v1_container.V1Container attribute\), 315](#)
[termination_message_policy \(kubernetes.client.models.v1_container.V1Container attribute\), 315](#)
[test_connect_delete_namespaced_pod_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_delete_namespaced_service_proxy_with_path\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_delete_node_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_delete_node_proxy_with_path\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_delete_namespaced_pod_attach\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_namespaced_pod_exec\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_namespaced_pod_portforward\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_namespaced_pod_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_namespaced_service_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_namespaced_service_proxy_with_path\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_get_node_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_head_namespaced_pod_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_head_namespaced_pod_proxy_with_path\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_head_namespaced_service_proxy\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)
[test_connect_head_namespaced_service_proxy_with_path\(\) \(kubernetes.test.test_core_v1_api.TestCoreV1Api method\), 541](#)

| | |
|--|--|
| method), 541 | method), 542 |
| test_connect_head_node_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_post_namespaced_pod_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_head_node_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_post_namespaced_service_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_namespaced_pod_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_post_namespaced_service_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_namespaced_pod_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_post_node_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_namespaced_service_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_post_node_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_namespaced_service_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 541 | test_connect_put_namespaced_pod_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_node_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_connect_put_namespaced_pod_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_options_node_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_connect_put_namespaced_service_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_patch_namespaced_pod_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_connect_put_namespaced_service_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 |
| test_connect_patch_namespaced_pod_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_connect_put_node_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_connect_patch_namespaced_service_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_connect_put_node_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_connect_patch_namespaced_service_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_cluster_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1Alpha1Api method), 558 |
| test_connect_patch_node_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_cluster_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1Alpha1Api method), 558 |
| test_connect_patch_node_proxy_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_namespace() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_connect_post_namespaced_pod_attach() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_namespaced_binding() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_connect_post_namespaced_pod_exec() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_namespaced_config_map() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_connect_post_namespaced_pod_portforward() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_namespaced_controller_revision() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1Beta1Api method), 533 |
| test_connect_post_namespaced_pod_proxy() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 542 | test_create_namespaced_cron_job() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2Alpha1Api method), 542 |

| | |
|---|--|
| method), 539 | method), 543 |
| test_create_namespaced_daemon_set() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 552 | test_create_namespaced_pod_template() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_deployment() (kuber- netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | test_create_namespaced_replica_set() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 |
| test_create_namespaced_deployment() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 552 | test_create_namespaced_replication_controller() (ku- bernetes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_deployment_rollback() (kuber- netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | test_create_namespaced_resource_quota() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_deployment_rollback() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | test_create_namespaced_role() (kuber- netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori- zationV1alpha1Api method), 558 |
| test_create_namespaced_endpoints() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_namespaced_role_binding() (kuber- netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori- zationV1alpha1Api method), 558 |
| test_create_namespaced_event() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_namespaced_secret() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_horizontal_pod_autoscaler() (kuber- netes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537 | test_create_namespaced_service() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_ingress() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | test_create_namespaced_service_account() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_job() (kuber- netes.test.test_batch_v1_api.TestBatchV1Api method), 538 | test_create_namespaced_stateful_set() (kuber- netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 |
| test_create_namespaced_limit_range() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_node() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_local_subject_access_review() (kuber- netes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api method), 536 | test_create_persistent_volume() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 |
| test_create_namespaced_network_policy() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | test_create_pod_security_policy() (kuber- netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 |
| test_create_namespaced_persistent_volume_claim() (ku- bernetes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_self_subject_access_review() (kuber- netes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api method), 536 |
| test_create_namespaced_pod() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_self_subject_rules_review() (kuber- netes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api method), 536 |
| test_create_namespaced_pod_binding() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_storage_class() (kuber- netes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561 |
| test_create_namespaced_pod_disruption_budget() (ku- bernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557 | test_create_subject_access_review() (kuber- netes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api method), 536 |
| test_create_namespaced_pod_eviction() (kuber- netes.test.test_core_v1_api.TestCoreV1Api method), 543 | test_create_temp_file_with_content() (kuber- netes.config.kube_config_test.TestFileOrData method), 543 |

method), 530

test_create_token_review() (kubernetes.test.test_authentication_v1beta1_api.TestAuthenticationV1beta1Api method), 535

test_current_context() (kubernetes.config.kube_config_test.TestKubeConfigLoader method), 531

test_data_given_data() (kubernetes.config.kube_config_test.TestFileOrData method), 530

test_data_given_file() (kubernetes.config.kube_config_test.TestFileOrData method), 531

test_data_given_file_and_data() (kubernetes.config.kube_config_test.TestFileOrData method), 531

test_data_given_file_no_base64() (kubernetes.config.kube_config_test.TestFileOrData method), 531

test_delete_cluster_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

test_delete_cluster_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

test_delete_collection_cluster_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

test_delete_collection_cluster_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

test_delete_collection_namespaced_config_map() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 543

test_delete_collection_namespaced_controller_revision() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533

test_delete_collection_namespaced_cron_job() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539

test_delete_collection_namespaced_daemon_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553

test_delete_collection_namespaced_deployment() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533

test_delete_collection_namespaced_deployment() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553

test_delete_collection_namespaced_endpoints() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_event() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_horizontal_pod_autoscaler() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537

test_delete_collection_namespaced_ingress() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553

test_delete_collection_namespaced_job() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538

test_delete_collection_namespaced_limit_range() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_network_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553

test_delete_collection_namespaced_persistent_volume_claim() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_pod() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_pod_disruption_budget() (kubernetes.test.test_extensions_v1beta1_api.TestPolicyV1beta1Api method), 557

test_delete_collection_namespaced_pod_template() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_replica_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553

test_delete_collection_namespaced_replication_controller() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_resource_quota() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 558

test_delete_collection_namespaced_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_delete_collection_namespaced_secret() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_service_account() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

test_delete_collection_namespaced_stateful_set() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533

test_delete_collection_node() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544

| | | | |
|---|--|---|--|
| method), 544 | | method), 544 | |
| test_delete_collection_persistent_volume() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_namespaced_pod() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | |
| test_delete_collection_pod_security_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | | test_delete_namespaced_pod_disruption_budget() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557 | |
| test_delete_collection_storage_class() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561 | | test_delete_namespaced_pod_template() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | |
| test_delete_namespace() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_namespaced_replica_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | |
| test_delete_namespaced_config_map() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_namespaced_replication_controller() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | |
| test_delete_namespaced_controller_revision() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | | test_delete_namespaced_resource_quota() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_cron_job() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539 | | test_delete_namespaced_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559 | |
| test_delete_namespaced_daemon_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | | test_delete_namespaced_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559 | |
| test_delete_namespaced_deployment() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | | test_delete_namespaced_secret() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_deployment() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | | test_delete_namespaced_service() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_endpoints() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_namespaced_service_account() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_event() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_namespaced_stateful_set() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | |
| test_delete_namespaced_horizontal_pod_autoscaler() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537 | | test_delete_node() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_ingress() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | | test_delete_persistent_volume() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 545 | |
| test_delete_namespaced_job() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538 | | test_delete_pod_security_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | |
| test_delete_namespaced_limit_range() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_delete_storage_class() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561 | |
| test_delete_namespaced_network_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 553 | | test_empty_cert_file() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |
| test_delete_namespaced_persistent_volume_claim() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 544 | | test_empty_host() (kubernetes.config.incluster_config_test.InClusterConfigTest method), 528 | |

| | | | |
|---|----------|--|----------|
| method), 528 | | method), 558 | |
| test_empty_port() | (kuber- | test_get_api_group() | (kuber- |
| netes.config.incluster_config_test.InClusterConfigTest | | netes.test.test_storage_api.TestStorageApi | |
| method), 528 | | method), 560 | |
| test_empty_token_file() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.incluster_config_test.InClusterConfigTest | | netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | |
| method), 529 | | method), 533 | |
| test_file_given_data() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_authentication_v1beta1_api.TestAuthenticationV1beta1Api | |
| method), 531 | | method), 535 | |
| test_file_given_data_no_base64() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_authorization_v1beta1_api.TestAuthorizationV1beta1Api | |
| method), 531 | | method), 536 | |
| test_file_given_file() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_autoscaling_v1_api.TestAutoscalingV1Api | |
| method), 531 | | method), 537 | |
| test_file_given_file_and_data() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_batch_v1_api.TestBatchV1Api | |
| method), 531 | | method), 538 | |
| test_file_given_non_existing_file() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api | |
| method), 531 | | method), 539 | |
| test_file_with_custom_dirname() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestFileOrData | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 531 | | method), 545 | |
| test_gcp_no_refresh() | (kuber- | test_get_api_resources() | (kuber- |
| netes.config.kube_config_test.TestKubeConfigLoader | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 531 | | method), 553 | |
| test_get_api_group() | (kuber- | test_get_api_resources() | (kuber- |
| netes.test.test_apps_api.TestAppsApi | method), | netes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api | |
| 532 | | method), 557 | |
| test_get_api_group() | (kuber- | test_get_api_resources() | (kuber- |
| netes.test.test_authentication_api.TestAuthenticationApi | | netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api | |
| method), 535 | | method), 559 | |
| test_get_api_group() | (kuber- | test_get_api_resources() | (kuber- |
| netes.test.test_authorization_api.TestAuthorizationApi | | netes.test.test_storage_v1beta1_api.TestStorageV1beta1Api | |
| method), 535 | | method), 561 | |
| test_get_api_group() | (kuber- | test_get_api_versions() | (kuber- |
| netes.test.test_autoscaling_api.TestAutoscalingApi | | netes.test.test_apis_api.TestApisApi | method), |
| method), 536 | | 532 | |
| test_get_api_group() | (kuber- | test_get_api_versions() | (kuber- |
| netes.test.test_batch_api.TestBatchApi | | netes.test.test_core_api.TestCoreApi | method), |
| method), 538 | | 540 | |
| test_get_api_group() | (kuber- | test_get_code() | (kuber- |
| netes.test.test_certificates_api.TestCertificatesApi | | netes.test.test_version_api.TestVersionApi | |
| method), 540 | | method), 636 | |
| test_get_api_group() | (kuber- | test_get_with_name() | (kuber- |
| netes.test.test_extensions_api.TestExtensionsApi | | netes.config.kube_config_test.TestConfigNode | |
| method), 552 | | method), 530 | |
| test_get_api_group() | (kuber- | test_get_with_name_on_invalid_object() | (kuber- |
| netes.test.test_policy_api.TestPolicyApi | | netes.config.kube_config_test.TestConfigNode | |
| method), 557 | | method), 530 | |
| test_get_api_group() | (kuber- | test_get_with_name_on_name_does_not_exists() | (kuber- |
| netes.test.test_rbac_authorization_api.TestRbacAuthorizationApi | | netes.config.kube_config_test.TestConfigNode | |

method), 530

test_get_with_name_on_non_list_object() (kuber-
netes.config.kube_config_test.TestConfigNode
method), 530

test_join_host_port() (kuber-
netes.config.incluster_config_test.InClusterConfigTest
method), 529

test_key_does_not_exists() (kuber-
netes.config.kube_config_test.TestConfigNode
method), 530

TEST_KUBE_CONFIG (kuber-
netes.config.kube_config_test.TestKubeConfigLoader
attribute), 531

test_list_cluster_role() (kuber-
netes.test.test_rbac_authorization_v1alpha1_api.TestRbacA
method), 559

test_list_cluster_role_binding() (kuber-
netes.test.test_rbac_authorization_v1alpha1_api.TestRbacA
method), 559

test_list_component_status() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_config_map_for_all_namespaces() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_contexts() (kuber-
netes.config.kube_config_test.TestKubeConfigLoader
method), 531

test_list_controller_revision_for_all_namespaces() (ku-
bernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api
method), 533

test_list_cron_job_for_all_namespaces() (kuber-
netes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api
method), 539

test_list_daemon_set_for_all_namespaces() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 553

test_list_deployment_for_all_namespaces() (kuber-
netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api
method), 533

test_list_deployment_for_all_namespaces() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 553

test_list_endpoints_for_all_namespaces() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_event_for_all_namespaces() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_horizontal_pod_autoscaler_for_all_namespaces() (kuber-
netes.test.test_autoscaling_v1_api.TestAutoscalingV1Api
method), 537

test_list_ingress_for_all_namespaces() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 553

method), 553

test_list_job_for_all_namespaces() (kuber-
netes.test.test_batch_v1_api.TestBatchV1Api
method), 538

test_list_kube_config_contexts() (kuber-
netes.config.kube_config_test.TestKubeConfigLoader
method), 531

test_list_limit_range_for_all_namespaces() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_namespace() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_namespaced_config_map() (kuber-
netes.test.test_rbac_authorization_v1alpha1_api.TestCoreV1Api
method), 545

test_list_namespaced_controller_revision() (kuber-
netes.test.test_rbac_authorization_v1alpha1_api.TestAppsV1beta1Api
method), 533

test_list_namespaced_cron_job() (kuber-
netes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api
method), 539

test_list_namespaced_daemon_set() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 554

test_list_namespaced_deployment() (kuber-
netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api
method), 533

test_list_namespaced_deployment() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 554

test_list_namespaced_endpoints() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_namespaced_event() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_namespaced_horizontal_pod_autoscaler() (ku-
bernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api
method), 537

test_list_namespaced_ingress() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 554

test_list_namespaced_job() (kuber-
netes.test.test_batch_v1_api.TestBatchV1Api
method), 538

test_list_namespaced_limit_range() (kuber-
netes.test.test_core_v1_api.TestCoreV1Api
method), 545

test_list_namespaced_network_policy() (kuber-
netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Ap
method), 554

test_list_namespaced_persistent_volume_claim() (ku-
bernetes.test.test_core_v1_api.TestCoreV1Api
method), 545

| | | | |
|---|---|--------------|--|
| method), 545 | | method), 546 | |
| test_list_namespaced_pod() (kuber- | test_list_pod_security_policy() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | |
| method), 545 | method), 554 | | |
| test_list_namespaced_pod_disruption_budget() (kuber- | test_list_pod_template_for_all_namespaces() (kuber- | | |
| netes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 557 | method), 546 | | |
| test_list_namespaced_pod_template() (kuber- | test_list_replica_set_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | |
| method), 545 | method), 554 | | |
| test_list_namespaced_replica_set() (kuber- | test_list_replication_controller_for_all_namespaces() (kuber- | | |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 554 | method), 546 | | |
| test_list_namespaced_replication_controller() (kuber- | test_list_resource_quota_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 545 | method), 546 | | |
| test_list_namespaced_resource_quota() (kuber- | test_list_role_binding_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori | | |
| method), 546 | method), 559 | | |
| test_list_namespaced_role() (kuber- | test_list_role_for_all_namespaces() (kuber- | | |
| netes.test.test_rbac_authorization_v1alpha1_api.TestRbacA | netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori | | |
| method), 559 | method), 559 | | |
| test_list_namespaced_role_binding() (kuber- | test_list_secret_for_all_namespaces() (kuber- | | |
| netes.test.test_rbac_authorization_v1alpha1_api.TestRbacA | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 559 | method), 546 | | |
| test_list_namespaced_secret() (kuber- | test_list_service_account_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 546 | method), 546 | | |
| test_list_namespaced_service() (kuber- | test_list_service_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_core_v1_api.TestCoreV1Api | | |
| method), 546 | method), 546 | | |
| test_list_namespaced_service_account() (kuber- | test_list_stateful_set_for_all_namespaces() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | |
| method), 546 | method), 533 | | |
| test_list_namespaced_stateful_set() (kuber- | test_list_storage_class() (kuber- | | |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | netes.test.test_storage_v1beta1_api.TestStorageV1beta1Api | | |
| method), 533 | method), 561 | | |
| test_list_network_policy_for_all_namespaces() (kuber- | test_load_config() (kuber- | | |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | netes.config.incluster_config_test.InClusterConfigTest | | |
| method), 554 | method), 529 | | |
| test_list_node() (kuber- | test_load_gcp_token_no_refresh() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.config.kube_config_test.TestKubeConfigLoader | | |
| method), 546 | method), 531 | | |
| test_list_persistent_volume() (kuber- | test_load_gcp_token_with_refresh() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.config.kube_config_test.TestKubeConfigLoader | | |
| method), 546 | method), 531 | | |
| test_list_persistent_volume_claim_for_all_namespaces() (kuber- | test_load_kube_config() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.config.kube_config_test.TestKubeConfigLoader | | |
| method), 546 | method), 531 | | |
| test_list_pod_disruption_budget_for_all_namespaces() (kuber- | test_load_user_pass_token() (kuber- | | |
| netes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api | netes.config.kube_config_test.TestKubeConfigLoader | | |
| method), 557 | method), 531 | | |
| test_list_pod_for_all_namespaces() (kuber- | test_load_user_token() (kuber- | | |
| netes.test.test_core_v1_api.TestCoreV1Api | netes.config.kube_config_test.TestKubeConfigLoader | | |

| | | | |
|---|---|--|---|
| method), 531 | | method), 539 | |
| test_log_file_handler() | (kubernetes.test.test_logs_api.TestLogsApi method), 556 | test_patch_namespaced_cron_job_status() | (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539 |
| test_log_file_list_handler() | (kubernetes.test.test_logs_api.TestLogsApi method), 556 | test_patch_namespaced_daemon_set() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_new_client_from_config() | (kubernetes.config.kube_config_test.TestKubeConfigLoader method), 531 | test_patch_namespaced_daemon_set_status() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_no_cert_file() | (kubernetes.config.incluster_config_test.InClusterConfigTest method), 529 | test_patch_namespaced_deployment() | (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 |
| test_no_host() | (kubernetes.config.incluster_config_test.InClusterConfigTest method), 529 | test_patch_namespaced_deployment_status() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_no_port() | (kubernetes.config.incluster_config_test.InClusterConfigTest method), 529 | test_patch_namespaced_deployment_scale() | (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 |
| test_no_token_file() | (kubernetes.config.incluster_config_test.InClusterConfigTest method), 529 | test_patch_namespaced_deployment_scale() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_no_user_context() | (kubernetes.config.kube_config_test.TestKubeConfigLoader method), 531 | test_patch_namespaced_deployment_status() | (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534 |
| test_no_users_section() | (kubernetes.config.kube_config_test.TestKubeConfigLoader method), 531 | test_patch_namespaced_deployment_status() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_non_existing_user() | (kubernetes.config.kube_config_test.TestKubeConfigLoader method), 531 | test_patch_namespaced_endpoints() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546 |
| test_normal_map_array_operations() | (kubernetes.config.kube_config_test.TestConfigNode method), 530 | test_patch_namespaced_event() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546 |
| test_obj | (kubernetes.config.kube_config_test.TestConfigNode attribute), 530 | test_patch_namespaced_horizontal_pod_autoscaler() | (kubernetes.test.test_extensions_v1beta1_api.TestAutoscalingV1beta1Api method), 537 |
| test_patch_cluster_role() | (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559 | test_patch_namespaced_horizontal_pod_autoscaler_status() | (kubernetes.test.test_extensions_v1beta1_api.TestAutoscalingV1beta1Api method), 537 |
| test_patch_cluster_role_binding() | (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559 | test_patch_namespaced_ingress() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_patch_namespace() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546 | test_patch_namespaced_ingress_status() | (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554 |
| test_patch_namespace_status() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546 | test_patch_namespaced_job() | (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538 |
| test_patch_namespaced_config_map() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546 | test_patch_namespaced_job_status() | (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538 |
| test_patch_namespaced_controller_revision() | (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 533 | test_patch_namespaced_limit_range() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 539 |
| test_patch_namespaced_cron_job() | (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539 | | |

method), 546

test_patch_namespaced_network_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_namespaced_persistent_volume_claim() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 546

test_patch_namespaced_persistent_volume_claim_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_pod() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_pod_disruption_budget() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557

test_patch_namespaced_pod_disruption_budget_status() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557

test_patch_namespaced_pod_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_pod_template() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_replica_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_namespaced_replica_set_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_namespaced_replica_set_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_namespaced_replication_controller() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_replication_controller_dummy_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_namespaced_replication_controller_scale() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_replication_controller_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_resource_quota() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_resource_quota_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_patch_namespaced_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_patch_namespaced_secret() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_service() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_service_account() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_service_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_namespaced_stateful_set() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_patch_namespaced_stateful_set_scale() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_patch_namespaced_stateful_set_status() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_patch_node() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_node_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_persistent_volume() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_persistent_volume_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_patch_pod_security_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 554

test_patch_storage_class() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561

test_proxy_delete_namespaced_pod() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_proxy_delete_namespaced_pod_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_proxy_delete_namespaced_service() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

test_proxy_delete_namespaced_service_with_path() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 547

| | | | |
|---|--|---|--|
| method), 547 | | method), 548 | |
| test_proxy_delete_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_options_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 |
| test_proxy_delete_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_options_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 |
| test_proxy_get_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 |
| test_proxy_get_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_get_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_get_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_get_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_get_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_patch_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_node() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_head_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_post_node_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_options_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_put_namespaced_pod() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_options_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_put_namespaced_pod_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_options_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_put_namespaced_service() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |
| test_proxy_options_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 548 | test_proxy_put_namespaced_service_with_path() | (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 549 |

| | | | |
|--|---------|---|---------|
| method), 549 | | method), 534 | |
| test_proxy_put_node() | (kuber- | test_read_namespaced_deployment_status() | (kuber- |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 549 | | method), 555 | |
| test_proxy_put_node_with_path() | (kuber- | test_read_namespaced_endpoints() | (kuber- |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 549 | | method), 550 | |
| test_read_cluster_role() | (kuber- | test_read_namespaced_event() | (kuber- |
| netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 559 | | method), 550 | |
| test_read_cluster_role_binding() | (kuber- | test_read_namespaced_horizontal_pod_autoscaler() | (kuber- |
| netes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 559 | | method), 537 | |
| test_read_component_status() | (kuber- | test_read_namespaced_horizontal_pod_autoscaler_status() | |
| netes.test.test_core_v1_api.TestCoreV1Api | | (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api | |
| method), 549 | | method), 537 | |
| test_read_namespace() | (kuber- | test_read_namespaced_ingress() | (kuber- |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 549 | | method), 555 | |
| test_read_namespace_status() | (kuber- | test_read_namespaced_ingress_status() | (kuber- |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 549 | | method), 555 | |
| test_read_namespaced_config_map() | (kuber- | test_read_namespaced_job() | (kuber- |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_batch_v1_api.TestBatchV1Api | |
| method), 549 | | method), 538 | |
| test_read_namespaced_controller_revision() | (kuber- | test_read_namespaced_job_status() | (kuber- |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | netes.test.test_batch_v1_api.TestBatchV1Api | |
| method), 534 | | method), 538 | |
| test_read_namespaced_cron_job() | (kuber- | test_read_namespaced_limit_range() | (kuber- |
| netes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 539 | | method), 550 | |
| test_read_namespaced_cron_job_status() | (kuber- | test_read_namespaced_network_policy() | (kuber- |
| netes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api | | netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | |
| method), 539 | | method), 555 | |
| test_read_namespaced_daemon_set() | (kuber- | test_read_namespaced_persistent_volume_claim() | (kuber- |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 555 | | method), 550 | |
| test_read_namespaced_daemon_set_status() | (kuber- | test_read_namespaced_persistent_volume_claim_status() | (kuber- |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 555 | | method), 550 | |
| test_read_namespaced_deployment() | (kuber- | test_read_namespaced_pod() | (kuber- |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 534 | | method), 550 | |
| test_read_namespaced_deployment() | (kuber- | test_read_namespaced_pod_disruption_budget() | (kuber- |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | netes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api | |
| method), 555 | | method), 557 | |
| test_read_namespaced_deployment_scale() | (kuber- | test_read_namespaced_pod_disruption_budget_status() | |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api | |
| method), 534 | | method), 557 | |
| test_read_namespaced_deployment_scale() | (kuber- | test_read_namespaced_pod_log() | (kuber- |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |
| method), 555 | | method), 550 | |
| test_read_namespaced_deployment_status() | (kuber- | test_read_namespaced_pod_status() | (kuber- |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | netes.test.test_core_v1_api.TestCoreV1Api | |

method), 550

test_read_namespaced_pod_template() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_replica_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_read_namespaced_replica_set_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_read_namespaced_replica_set_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_read_namespaced_replication_controller() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_replication_controller_dummy_scaled() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_read_namespaced_replication_controller_scale() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_replication_controller_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_resource_quota() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_resource_quota_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_read_namespaced_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_read_namespaced_secret() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_service() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_service_account() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_service_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_namespaced_stateful_set() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_read_namespaced_stateful_set_scale() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_read_namespaced_stateful_set_status() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_read_node() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_node_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_persistent_volume() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 550

test_read_persistent_volume_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551

test_read_pod_security_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_read_storage_class() (kubernetes.test.test_storage_v1beta1_api.TestStorageV1beta1Api method), 561

test_replace_cluster_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_replace_cluster_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthorizationV1alpha1Api method), 559

test_replace_namespace() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551

test_replace_namespace_finalize() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551

test_replace_namespace_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551

test_replace_namespaced_config_map() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551

test_replace_namespaced_controller_revision() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534

test_replace_namespaced_cron_job() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539

test_replace_namespaced_cron_job_status() (kubernetes.test.test_batch_v2alpha1_api.TestBatchV2alpha1Api method), 539

test_replace_namespaced_daemon_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

test_replace_namespaced_daemon_set_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555

| | |
|---|---|
| method), 555 | method), 551 |
| test_replace_namespaced_deployment() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534 | test_replace_namespaced_pod() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_deployment() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_pod_disruption_budget() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557 |
| test_replace_namespaced_deployment_scale() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534 | test_replace_namespaced_pod_disruption_budget_status() (kubernetes.test.test_policy_v1beta1_api.TestPolicyV1beta1Api method), 557 |
| test_replace_namespaced_deployment_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_pod_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_deployment_status() (kubernetes.test.test_apps_v1beta1_api.TestAppsV1beta1Api method), 534 | test_replace_namespaced_pod_template() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_deployment_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_replica_set() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 556 |
| test_replace_namespaced_endpoints() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 | test_replace_namespaced_replica_set_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 556 |
| test_replace_namespaced_event() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 | test_replace_namespaced_replica_set_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 556 |
| test_replace_namespaced_horizontal_pod_autoscaler() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537 | test_replace_namespaced_replication_controller() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_horizontal_pod_autoscaler_status() (kubernetes.test.test_autoscaling_v1_api.TestAutoscalingV1Api method), 537 | test_replace_namespaced_replication_controller_dummy_scale() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 556 |
| test_replace_namespaced_ingress() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_replication_controller_scale() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_ingress_status() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_replication_controller_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_job() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538 | test_replace_namespaced_resource_quota() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_job_status() (kubernetes.test.test_batch_v1_api.TestBatchV1Api method), 538 | test_replace_namespaced_resource_quota_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_limit_range() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 | test_replace_namespaced_role() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori method), 559 |
| test_replace_namespaced_network_policy() (kubernetes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api method), 555 | test_replace_namespaced_role_binding() (kubernetes.test.test_rbac_authorization_v1alpha1_api.TestRbacAuthori method), 560 |
| test_replace_namespaced_persistent_volume_claim() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 | test_replace_namespaced_secret() (kubernetes.test.test_core_v1_api.TestCoreV1Api method), 551 |
| test_replace_namespaced_persistent_volume_claim_status() (kubernetes.test.test_core_v1_api.TestCoreV1Api | test_replace_namespaced_service() (kubernetes.test.test_core_v1_api.TestCoreV1Api |

| | | | |
|---|--|---|--|
| method), 551 | | test_unmarshal_with_no_return_type() (kuber- | |
| test_replace_namespaced_service_account() (kuber- | | netes.watch.watch_test.WatchTests method), | |
| netes.test.test_core_v1_api.TestCoreV1Api | | 638 | |
| method), 551 | | test_user_pass() (kuber- | |
| test_replace_namespaced_service_status() (kuber- | | netes.config.kube_config_test.TestKubeConfigLoader | |
| netes.test.test_core_v1_api.TestCoreV1Api | | method), 531 | |
| method), 552 | | test_watch_stream_twice() (kuber- | |
| test_replace_namespaced_stateful_set() (kuber- | | netes.watch.watch_test.WatchTests method), | |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | 638 | |
| method), 534 | | test_watch_with_decode() (kuber- | |
| test_replace_namespaced_stateful_set_scale() (kuber- | | netes.watch.watch_test.WatchTests method), | |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | 638 | |
| method), 534 | | test_watch_with_exception() (kuber- | |
| test_replace_namespaced_stateful_set_status() (kuber- | | netes.watch.watch_test.WatchTests method), | |
| netes.test.test_apps_v1beta1_api.TestAppsV1beta1Api | | 638 | |
| method), 534 | | TestApisApi (class in kubernetes.test.test_apis_api), 532 | |
| test_replace_node() (kuber- | | TestAppsApi (class in kubernetes.test.test_apps_api), 532 | |
| netes.test.test_core_v1_api.TestCoreV1Api | | TestAppsV1beta1Api (class in kuber- | |
| method), 552 | | netes.test.test_apps_v1beta1_api), 532 | |
| test_replace_node_status() (kuber- | | TestAuthenticationApi (class in kuber- | |
| netes.test.test_core_v1_api.TestCoreV1Api | | netes.test.test_authentication_api), 535 | |
| method), 552 | | TestAuthenticationV1beta1Api (class in kuber- | |
| test_replace_persistent_volume() (kuber- | | netes.test.test_authentication_v1beta1_api), | |
| netes.test.test_core_v1_api.TestCoreV1Api | | 535 | |
| method), 552 | | TestAuthorizationApi (class in kuber- | |
| test_replace_persistent_volume_status() (kuber- | | netes.test.test_authorization_api), 535 | |
| netes.test.test_core_v1_api.TestCoreV1Api | | TestAuthorizationV1beta1Api (class in kuber- | |
| method), 552 | | netes.test.test_authorization_v1beta1_api), | |
| test_replace_pod_security_policy() (kuber- | | 536 | |
| netes.test.test_extensions_v1beta1_api.TestExtensionsV1beta1Api | | TestAutoscalingApi (class in kuber- | |
| method), 556 | | netes.test.test_autoscaling_api), 536 | |
| test_replace_storage_class() (kuber- | | TestAutoscalingV1Api (class in kuber- | |
| netes.test.test_storage_v1beta1_api.TestStorageV1beta1Api | | netes.test.test_autoscaling_v1_api), 537 | |
| method), 561 | | TestBatchApi (class in kubernetes.test.test_batch_api), | |
| test_set_active_context() (kuber- | | 537 | |
| netes.config.kube_config_test.TestKubeConfigLoader | | TestBatchV1Api (class in kuber- | |
| method), 531 | | netes.test.test_batch_v1_api), 538 | |
| test_simple_token() (kuber- | | TestBatchV2alpha1Api (class in kuber- | |
| netes.config.kube_config_test.TestKubeConfigLoader | | netes.test.test_batch_v2alpha1_api), 539 | |
| method), 531 | | TestCertificatesApi (class in kuber- | |
| test_ssl() (kubernetes.config.kube_config_test.TestKubeConfigLoader | | netes.test.test_certificates_api), 539 | |
| method), 531 | | TestConfigNode (class in kuber- | |
| test_ssl_no_cert_files() (kuber- | | netes.config.kube_config_test), 530 | |
| netes.config.kube_config_test.TestKubeConfigLoader | | TestCoreApi (class in kubernetes.test.test_core_api), 540 | |
| method), 531 | | TestCoreV1Api (class in kuber- | |
| test_ssl_no_verification() (kuber- | | netes.test.test_core_v1_api), 540 | |
| netes.config.kube_config_test.TestKubeConfigLoader | | TestExtensionsApi (class in kuber- | |
| method), 531 | | netes.test.test_extensions_api), 552 | |
| test_ssl_with_relative_ssl_files() (kuber- | | TestExtensionsV1beta1Api (class in kuber- | |
| netes.config.kube_config_test.TestKubeConfigLoader | | netes.test.test_extensions_v1beta1_api), | |
| method), 531 | | 552 | |
| test_unmarshal_with_float_object() (kuber- | | TestFileOrData (class in kuber- | |
| netes.watch.watch_test.WatchTests method), | | netes.config.kube_config_test), 530 | |
| 638 | | TestKubeConfigLoader (class in kuber- | |

netes.config.kube_config_test), 531

TestLogsApi (class in kubernetes.test.test_logs_api), 556

TestPolicyApi (class in kubernetes.test.test_policy_api), 556

TestPolicyV1beta1Api (class in kubernetes.test.test_policy_v1beta1_api), 557

TestRbacAuthorizationApi (class in kubernetes.test.test_rbac_authorization_api), 558

TestRbacAuthorizationV1alpha1Api (class in kubernetes.test.test_rbac_authorization_v1alpha1_api), 558

TestRuntimeRawExtension (class in kubernetes.test.test_runtime_raw_extension), 560

testRuntimeRawExtension() (kubernetes.test.test_runtime_raw_extension.TestRuntimeRawExtension method), 560

TestStorageApi (class in kubernetes.test.test_storage_api), 560

TestStorageV1beta1Api (class in kubernetes.test.test_storage_v1beta1_api), 560

TestV1alpha1ClusterRole (class in kubernetes.test.test_v1alpha1_cluster_role), 612

testV1alpha1ClusterRole() (kubernetes.test.test_v1alpha1_cluster_role.TestV1alpha1ClusterRole method), 613

TestV1alpha1ClusterRoleBinding (class in kubernetes.test.test_v1alpha1_cluster_role_binding), 613

testV1alpha1ClusterRoleBinding() (kubernetes.test.test_v1alpha1_cluster_role_binding.TestV1alpha1ClusterRoleBinding method), 613

TestV1alpha1ClusterRoleBindingList (class in kubernetes.test.test_v1alpha1_cluster_role_binding_list), 613

testV1alpha1ClusterRoleBindingList() (kubernetes.test.test_v1alpha1_cluster_role_binding_list.TestV1alpha1ClusterRoleBindingList method), 613

TestV1alpha1ClusterRoleList (class in kubernetes.test.test_v1alpha1_cluster_role_list), 614

testV1alpha1ClusterRoleList() (kubernetes.test.test_v1alpha1_cluster_role_list.TestV1alpha1ClusterRoleList method), 614

TestV1alpha1PolicyRule (class in kubernetes.test.test_v1alpha1_policy_rule), 614

testV1alpha1PolicyRule() (kubernetes.test.test_v1alpha1_policy_rule.TestV1alpha1PolicyRule method), 614

TestV1alpha1Role (class in kubernetes.test.test_v1alpha1_role), 614

testV1alpha1Role() (kubernetes.test.test_v1alpha1_role.TestV1alpha1Role method), 614

TestV1alpha1RoleBinding (class in kubernetes.test.test_v1alpha1_role_binding), 615

testV1alpha1RoleBinding() (kubernetes.test.test_v1alpha1_role_binding.TestV1alpha1RoleBinding method), 615

TestV1alpha1RoleBindingList (class in kubernetes.test.test_v1alpha1_role_binding_list), 615

testV1alpha1RoleBindingList() (kubernetes.test.test_v1alpha1_role_binding_list.TestV1alpha1RoleBindingList method), 615

TestV1alpha1RoleList (class in kubernetes.test.test_v1alpha1_role_list), 615

testV1alpha1RoleList() (kubernetes.test.test_v1alpha1_role_list.TestV1alpha1RoleList method), 615

TestV1alpha1RoleRef (class in kubernetes.test.test_v1alpha1_role_ref), 616

testV1alpha1RoleRef() (kubernetes.test.test_v1alpha1_role_ref.TestV1alpha1RoleRef method), 616

TestV1alpha1Subject (class in kubernetes.test.test_v1alpha1_subject), 616

testV1alpha1Subject() (kubernetes.test.test_v1alpha1_subject.TestV1alpha1Subject method), 616

TestV1AttachedVolume (class in kubernetes.test.test_v1_attached_volume), 562

testV1AttachedVolume() (kubernetes.test.test_v1_attached_volume.TestV1AttachedVolume method), 562

TestV1AWSElasticBlockStoreVolumeSource (class in kubernetes.test.test_v1_aws_elastic_block_store_volume_source), 563

testV1AWSElasticBlockStoreVolumeSource() (kubernetes.test.test_v1_aws_elastic_block_store_volume_source.TestV1AWSElasticBlockStoreVolumeSource method), 563

TestV1AzureDiskVolumeSource (class in kubernetes.test.test_v1_azure_disk_volume_source), 563

testV1AzureDiskVolumeSource() (kubernetes.test.test_v1_azure_disk_volume_source.TestV1AzureDiskVolumeSource method), 563

TestV1AzureFileVolumeSource (class in kubernetes.test.test_v1_azure_file_volume_source), 563

testV1AzureFileVolumeSource() (kubernetes.test.test_v1_azure_file_volume_source.TestV1AzureFileVolumeSource method), 563

TestV1beta1DaemonSet (class in kubernetes.test.test_v1beta1_daemon_set), 616

testV1beta1DaemonSet() (kubernetes.test.test_v1beta1_daemon_set.TestV1beta1DaemonSet method), 617

TestV1beta1DaemonSetList (class in kubernetes.test.test_v1beta1_daemon_set_list), 617

testV1beta1DaemonSetList() (kubernetes.test.test_v1beta1_daemon_set_list.TestV1beta1DaemonSetList method), 617

TestV1beta1DaemonSetSpec (class in kubernetes.test.test_v1beta1_daemon_set_spec), 617

testV1beta1DaemonSetSpec() (kubernetes.test.test_v1beta1_daemon_set_spec.TestV1beta1DaemonSetSpec method), 617

TestV1beta1DaemonSetStatus (class in kubernetes.test.test_v1beta1_daemon_set_status), 618

testV1beta1DaemonSetStatus() (kubernetes.test.test_v1beta1_daemon_set_status.TestV1beta1DaemonSetStatus method), 618

TestV1beta1Eviction (class in kubernetes.test.test_v1beta1_eviction), 618

testV1beta1Eviction() (kubernetes.test.test_v1beta1_eviction.TestV1beta1Eviction method), 618

TestV1beta1HTTPIngressPath (class in kubernetes.test.test_v1beta1_http_ingress_path), 619

testV1beta1HTTPIngressPath() (kubernetes.test.test_v1beta1_http_ingress_path.TestV1beta1HTTPIngressPath method), 619

TestV1beta1HTTPIngressRuleValue (class in kubernetes.test.test_v1beta1_http_ingress_rule_value), 619

testV1beta1HTTPIngressRuleValue() (kubernetes.test.test_v1beta1_http_ingress_rule_value.TestV1beta1HTTPIngressRuleValue method), 619

TestV1beta1Ingress (class in kubernetes.test.test_v1beta1_ingress), 620

testV1beta1Ingress() (kubernetes.test.test_v1beta1_ingress.TestV1beta1Ingress method), 620

TestV1beta1IngressBackend (class in kubernetes.test.test_v1beta1_ingress_backend), 620

testV1beta1IngressBackend() (kubernetes.test.test_v1beta1_ingress_backend.TestV1beta1IngressBackend method), 620

TestV1beta1IngressList (class in kubernetes.test.test_v1beta1_ingress_list), 620

testV1beta1IngressList() (kubernetes.test.test_v1beta1_ingress_list.TestV1beta1IngressList method), 620

TestV1beta1IngressRule (class in kubernetes.test.test_v1beta1_ingress_rule), 621

testV1beta1IngressRule() (kubernetes.test.test_v1beta1_ingress_rule.TestV1beta1IngressRule method), 621

TestV1beta1IngressSpec (class in kubernetes.test.test_v1beta1_ingress_spec), 621

testV1beta1IngressSpec() (kubernetes.test.test_v1beta1_ingress_spec.TestV1beta1IngressSpec method), 621

TestV1beta1IngressStatus (class in kubernetes.test.test_v1beta1_ingress_status), 621

testV1beta1IngressStatus() (kubernetes.test.test_v1beta1_ingress_status.TestV1beta1IngressStatus method), 621

TestV1beta1IngressTLS (class in kubernetes.test.test_v1beta1_ingress_tls), 622

testV1beta1IngressTLS() (kubernetes.test.test_v1beta1_ingress_tls.TestV1beta1IngressTLS method), 622

TestV1beta1LocalSubjectAccessReview (class in kubernetes.test.test_v1beta1_local_subject_access_review), 622

testV1beta1LocalSubjectAccessReview() (kubernetes.test.test_v1beta1_local_subject_access_review.TestV1beta1LocalSubjectAccessReview method), 622

TestV1beta1NetworkPolicy (class in kubernetes.test.test_v1beta1_network_policy), 623

testV1beta1NetworkPolicy() (kubernetes.test.test_v1beta1_network_policy.TestV1beta1NetworkPolicy method), 623

TestV1beta1NetworkPolicyIngressRule (class in kubernetes.test.test_v1beta1_network_policy_ingress_rule), 623

testV1beta1NetworkPolicyIngressRule() (kubernetes.test.test_v1beta1_network_policy_ingress_rule.TestV1beta1NetworkPolicyIngressRule method), 623

TestV1beta1NetworkPolicyList (class in kubernetes.test.test_v1beta1_network_policy_list), 623

testV1beta1NetworkPolicyList() (kubernetes.test.test_v1beta1_network_policy_list.TestV1beta1NetworkPolicyList method), 623

TestV1beta1NetworkPolicyPeer (class in kubernetes.test.test_v1beta1_network_policy_peer), 624

testV1beta1NetworkPolicyPeer() (kubernetes.test.test_v1beta1_network_policy_peer.TestV1beta1NetworkPolicyPeer method), 624

TestV1beta1NetworkPolicyPort (class in kubernetes.test.test_v1beta1_network_policy_port), 624

testV1beta1NetworkPolicyPort() (kubernetes.test.test_v1beta1_network_policy_port.TestV1beta1NetworkPolicyPort method), 624

TestV1beta1NetworkPolicySpec (class in kubernetes.test.test_v1beta1_network_policy_spec), 624

testV1beta1NetworkPolicySpec() (kubernetes.test.test_v1beta1_network_policy_spec.TestV1beta1NetworkPolicySpec method), 624

netes.test.test_v1beta1_network_policy_spec), 624
testV1beta1NetworkPolicySpec() (kuber- netes.test.test_v1beta1_network_policy_spec.TestV1beta1NetworkPolicySpec method), 625
TestV1beta1NonResourceAttributes (class in kuber- netes.test.test_v1beta1_non_resource_attributes), 625
testV1beta1NonResourceAttributes() (kuber- netes.test.test_v1beta1_non_resource_attributes.TestV1beta1NonResourceAttributes method), 625
TestV1beta1PodDisruptionBudget (class in kuber- netes.test.test_v1beta1_pod_disruption_budget), 625
testV1beta1PodDisruptionBudget() (kuber- netes.test.test_v1beta1_pod_disruption_budget.TestV1beta1PodDisruptionBudget method), 625
TestV1beta1PodDisruptionBudgetList (class in kuber- netes.test.test_v1beta1_pod_disruption_budget_list), 625
testV1beta1PodDisruptionBudgetList() (kuber- netes.test.test_v1beta1_pod_disruption_budget_list.TestV1beta1PodDisruptionBudgetList method), 626
TestV1beta1PodDisruptionBudgetSpec (class in kuber- netes.test.test_v1beta1_pod_disruption_budget_spec), 626
testV1beta1PodDisruptionBudgetSpec() (kuber- netes.test.test_v1beta1_pod_disruption_budget_spec.TestV1beta1PodDisruptionBudgetSpec method), 626
TestV1beta1PodDisruptionBudgetStatus (class in kuber- netes.test.test_v1beta1_pod_disruption_budget_status), 626
testV1beta1PodDisruptionBudgetStatus() (kuber- netes.test.test_v1beta1_pod_disruption_budget_status.TestV1beta1PodDisruptionBudgetStatus method), 626
TestV1beta1ReplicaSet (class in kuber- netes.test.test_v1beta1_replica_set), 627
testV1beta1ReplicaSet() (kuber- netes.test.test_v1beta1_replica_set.TestV1beta1ReplicaSet method), 627
TestV1beta1ReplicaSetCondition (class in kuber- netes.test.test_v1beta1_replica_set_condition), 627
testV1beta1ReplicaSetCondition() (kuber- netes.test.test_v1beta1_replica_set_condition.TestV1beta1ReplicaSetCondition method), 627
TestV1beta1ReplicaSetList (class in kuber- netes.test.test_v1beta1_replica_set_list), 627
testV1beta1ReplicaSetList() (kuber- netes.test.test_v1beta1_replica_set_list.TestV1beta1ReplicaSetList method), 627
TestV1beta1ReplicaSetSpec (class in kuber- netes.test.test_v1beta1_replica_set_spec), 628
testV1beta1ReplicaSetSpec() (kuber- netes.test.test_v1beta1_replica_set_spec.TestV1beta1ReplicaSetSpec method), 628
TestV1beta1ReplicaSetStatus (class in kuber- netes.test.test_v1beta1_replica_set_status), 628
testV1beta1ReplicaSetStatus() (kuber- netes.test.test_v1beta1_replica_set_status.TestV1beta1ReplicaSetStatus method), 628
TestV1beta1ResourceAttributes (class in kuber- netes.test.test_v1beta1_resource_attributes), 628
testV1beta1ResourceAttributes() (kuber- netes.test.test_v1beta1_resource_attributes.TestV1beta1ResourceAttributes method), 628
TestV1beta1SelfSubjectAccessReview (class in kuber- netes.test.test_v1beta1_self_subject_access_review), 629
testV1beta1SelfSubjectAccessReview() (kuber- netes.test.test_v1beta1_self_subject_access_review.TestV1beta1SelfSubjectAccessReview method), 629
TestV1beta1SelfSubjectAccessReviewSpec (class in kuber- netes.test.test_v1beta1_self_subject_access_review_spec), 629
testV1beta1SelfSubjectAccessReviewSpec() (kuber- netes.test.test_v1beta1_self_subject_access_review_spec.TestV1beta1SelfSubjectAccessReviewSpec method), 629
TestV1beta1StatefulSet (class in kuber- netes.test.test_v1beta1_stateful_set), 630
testV1beta1StatefulSet() (kuber- netes.test.test_v1beta1_stateful_set.TestV1beta1StatefulSet method), 630
TestV1beta1StatefulSetList (class in kuber- netes.test.test_v1beta1_stateful_set_list), 630
testV1beta1StatefulSetList() (kuber- netes.test.test_v1beta1_stateful_set_list.TestV1beta1StatefulSetList method), 630
TestV1beta1StatefulSetSpec (class in kuber- netes.test.test_v1beta1_stateful_set_spec), 630
testV1beta1StatefulSetSpec() (kuber- netes.test.test_v1beta1_stateful_set_spec.TestV1beta1StatefulSetSpec method), 630
TestV1beta1StatefulSetStatus (class in kuber- netes.test.test_v1beta1_stateful_set_status), 631
testV1beta1StatefulSetStatus() (kuber- netes.test.test_v1beta1_stateful_set_status.TestV1beta1StatefulSetStatus method), 631
TestV1beta1StorageClass (class in kuber- netes.test.test_v1beta1_storage_class), 631

| | | | |
|---|---|--|------------------|
| testV1beta1StorageClass() | (kuber- | TestV1Capabilities | (class in kuber- |
| netes.test.test_v1beta1_storage_class.TestV1beta1StorageClass | netes.test.test_v1_capabilities), 564 | testV1Capabilities() | (kuber- |
| method), 631 | | netes.test.test_v1_capabilities.TestV1Capabilities | method), 564 |
| TestV1beta1StorageClassList | (class in kuber- | TestV1CephFSVolumeSource | (class in kuber- |
| netes.test.test_v1beta1_storage_class_list), | | netes.test.test_v1_ceph_fs_volume_source), | |
| 631 | | TestV1CephFSVolumeSourceList | 664 |
| testV1beta1StorageClassList() | (kuber- | testV1CephFSVolumeSource() | (kuber- |
| netes.test.test_v1beta1_storage_class_list.TestV1beta1Storage | netes.test.test_v1_ceph_fs_volume_source.TestV1CephFSVolume | method), 565 | |
| method), 632 | | TestV1CinderVolumeSource | (class in kuber- |
| TestV1beta1SubjectAccessReview | (class in kuber- | netes.test.test_v1_cinder_volume_source), | |
| netes.test.test_v1beta1_subject_access_review), | | testV1CinderVolumeSource() | (kuber- |
| 632 | | netes.test.test_v1_cinder_volume_source.TestV1CinderVolumeS | method), 565 |
| testV1beta1SubjectAccessReview() | (kuber- | TestV1ComponentCondition | (class in kuber- |
| netes.test.test_v1beta1_subject_access_review.TestV1beta1SubjectAccessReview | netes.test.test_v1_component_condition), | testV1ComponentCondition() | (kuber- |
| method), 632 | | netes.test.test_v1_component_condition.TestV1ComponentCondi | method), 565 |
| TestV1beta1SubjectAccessReviewSpec | (class in kuber- | TestV1ComponentStatus | (class in kuber- |
| netes.test.test_v1beta1_subject_access_review_spec), | | netes.test.test_v1_component_status), 565 | |
| 632 | | TestV1ComponentStatusList | (class in kuber- |
| testV1beta1SubjectAccessReviewSpec() | (kuber- | netes.test.test_v1_component_status_list), | |
| netes.test.test_v1beta1_subject_access_review_spec.TestV1beta1SubjectAccessReviewSpec | testV1ComponentStatusList() | (kuber- | |
| method), 632 | | netes.test.test_v1_component_status_list.TestV1ComponentStatu | method), 566 |
| TestV1beta1SubjectAccessReviewStatus | (class in kuber- | TestV1ConfigMap | (class in kuber- |
| netes.test.test_v1beta1_subject_access_review_status), | | netes.test.test_v1_config_map), 566 | |
| 632 | | TestV1ConfigMapList | (class in kuber- |
| testV1beta1SubjectAccessReviewStatus() | (kuber- | netes.test.test_v1_config_map_list), 567 | |
| netes.test.test_v1beta1_subject_access_review_status.TestV1beta1SubjectAccessReviewStatus | testV1ConfigMapKeySelector | (class in kuber- | |
| method), 633 | | netes.test.test_v1_config_map_key_selector), | |
| TestV1beta1TokenReview | (class in kuber- | TestV1ConfigMapVolumeSource | (class in kuber- |
| netes.test.test_v1beta1_token_review), 633 | | netes.test.test_v1_config_map_volume_source), | |
| testV1beta1TokenReview() | (kuber- | testV1ConfigMapVolumeSource() | (kuber- |
| netes.test.test_v1beta1_token_review.TestV1beta1TokenReview | method), 566 | | |
| method), 633 | | TestV1ConfigMapVolumeSourceList | (class in kuber- |
| TestV1beta1TokenReviewSpec | (class in kuber- | netes.test.test_v1_config_map_volume_source_list), 567 | |
| netes.test.test_v1beta1_token_review_spec), | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| 633 | | netes.test.test_v1_config_map_volume_source_list.TestV1ConfigMapVolumeSourceList | method), 567 |
| testV1beta1TokenReviewSpec() | (kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1beta1_token_review_spec.TestV1beta1TokenReviewSpec | TestV1ConfigMapVolumeSourceList() | (kuber- | |
| method), 633 | | netes.test.test_v1_config_map_volume_source_list.TestV1ConfigMapVolumeSourceList | method), 567 |
| TestV1beta1TokenReviewStatus | (class in kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1beta1_token_review_status), | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| 634 | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| testV1beta1TokenReviewStatus() | (kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1beta1_token_review_status.TestV1beta1TokenReviewStatus | TestV1ConfigMapVolumeSourceList() | (kuber- | |
| method), 634 | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| TestV1beta1UserInfo | (class in kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1beta1_user_info), 634 | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| testV1beta1UserInfo() | (kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1beta1_user_info.TestV1beta1UserInfo | TestV1ConfigMapVolumeSourceList() | (kuber- | |
| method), 634 | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| TestV1Binding | (class in kubernetes.test.test_v1_binding), | TestV1ConfigMapVolumeSourceList() | (kuber- |
| 564 | | TestV1ConfigMapVolumeSourceList() | (kuber- |
| testV1Binding() | (kuber- | TestV1ConfigMapVolumeSourceList() | (kuber- |
| netes.test.test_v1_binding.TestV1Binding | TestV1ConfigMapVolumeSourceList() | (kuber- | |
| method), 564 | | TestV1ConfigMapVolumeSourceList() | (kuber- |

| | |
|--|--|
| netes.test.test_v1_config_map_volume_source.TestV1ConfigMapVolumeSource (class in kubernetes.test.test_v1_config_map_volume_source), 567 | TestV1DaemonEndpoint (class in kubernetes.test.test_v1_daemon_endpoint), 571 |
| TestV1Container (class in kubernetes.test.test_v1_container), 568 | TestV1DeleteOptions (class in kubernetes.test.test_v1_delete_options), 571 |
| testV1Container() (kubernetes.test.test_v1_container.TestV1Container method), 568 | testV1DeleteOptions() (kubernetes.test.test_v1_delete_options.TestV1DeleteOptions method), 571 |
| TestV1ContainerImage (class in kubernetes.test.test_v1_container_image), 568 | TestV1DownwardAPIVolumeFile (class in kubernetes.test.test_v1_downward_api_volume_file), 572 |
| testV1ContainerImage() (kubernetes.test.test_v1_container_image.TestV1ContainerImage method), 568 | testV1DownwardAPIVolumeFile() (kubernetes.test.test_v1_downward_api_volume_file.TestV1DownwardAPIVolumeFile method), 572 |
| TestV1ContainerPort (class in kubernetes.test.test_v1_container_port), 568 | TestV1DownwardAPIVolumeSource (class in kubernetes.test.test_v1_downward_api_volume_source), 572 |
| testV1ContainerPort() (kubernetes.test.test_v1_container_port.TestV1ContainerPort method), 568 | testV1DownwardAPIVolumeSource() (kubernetes.test.test_v1_downward_api_volume_source.TestV1DownwardAPIVolumeSource method), 572 |
| TestV1ContainerState (class in kubernetes.test.test_v1_container_state), 569 | TestV1EmptyDirVolumeSource (class in kubernetes.test.test_v1_empty_dir_volume_source), 572 |
| testV1ContainerState() (kubernetes.test.test_v1_container_state.TestV1ContainerState method), 569 | testV1EmptyDirVolumeSource() (kubernetes.test.test_v1_empty_dir_volume_source.TestV1EmptyDirVolumeSource method), 572 |
| TestV1ContainerStateRunning (class in kubernetes.test.test_v1_container_state_running), 569 | TestV1EndpointAddress (class in kubernetes.test.test_v1_endpoint_address), 573 |
| testV1ContainerStateRunning() (kubernetes.test.test_v1_container_state_running.TestV1ContainerStateRunning method), 569 | testV1EndpointAddress() (kubernetes.test.test_v1_endpoint_address.TestV1EndpointAddress method), 573 |
| TestV1ContainerStateTerminated (class in kubernetes.test.test_v1_container_state_terminated), 569 | TestV1EndpointPort (class in kubernetes.test.test_v1_endpoint_port), 573 |
| testV1ContainerStateTerminated() (kubernetes.test.test_v1_container_state_terminated.TestV1ContainerStateTerminated method), 570 | testV1EndpointPort() (kubernetes.test.test_v1_endpoint_port.TestV1EndpointPort method), 573 |
| TestV1ContainerStateWaiting (class in kubernetes.test.test_v1_container_state_waiting), 570 | TestV1Endpoints (class in kubernetes.test.test_v1_endpoints), 574 |
| testV1ContainerStateWaiting() (kubernetes.test.test_v1_container_state_waiting.TestV1ContainerStateWaiting method), 570 | testV1Endpoints() (kubernetes.test.test_v1_endpoints.TestV1Endpoints method), 574 |
| TestV1ContainerStatus (class in kubernetes.test.test_v1_container_status), 570 | TestV1EndpointsList (class in kubernetes.test.test_v1_endpoints_list), 574 |
| testV1ContainerStatus() (kubernetes.test.test_v1_container_status.TestV1ContainerStatus method), 570 | testV1EndpointsList() (kubernetes.test.test_v1_endpoints_list.TestV1EndpointsList method), 574 |
| TestV1CrossVersionObjectReference (class in kubernetes.test.test_v1_cross_version_object_reference), 570 | TestV1EndpointSubset (class in kubernetes.test.test_v1_endpoint_subset), 573 |
| testV1CrossVersionObjectReference() (kubernetes.test.test_v1_cross_version_object_reference.TestV1CrossVersionObjectReference method), 571 | testV1EndpointSubset() (kubernetes.test.test_v1_endpoint_subset.TestV1EndpointSubset method), 573 |
| TestV1DaemonEndpoint (class in kubernetes.test.test_v1_daemon_endpoint), 571 | TestV1EnvVar (class in kubernetes.test.test_v1_env_var), 574 |
| testV1DaemonEndpoint() (kubernetes.test.test_v1_daemon_endpoint.TestV1DaemonEndpoint method), 571 | testV1EnvVar() (kubernetes.test.test_v1_env_var.TestV1EnvVar method), 574 |

method), 575

TestV1EnvVarSource (class in kubernetes.test.test_v1_env_var_source), 575

testV1EnvVarSource() (kubernetes.test.test_v1_env_var_source.TestV1EnvVarSource method), 575

TestV1Event (class in kubernetes.test.test_v1_event), 575

testV1Event() (kubernetes.test.test_v1_event.TestV1Event method), 575

TestV1EventList (class in kubernetes.test.test_v1_event_list), 575

testV1EventList() (kubernetes.test.test_v1_event_list.TestV1EventList method), 576

TestV1EventSource (class in kubernetes.test.test_v1_event_source), 576

testV1EventSource() (kubernetes.test.test_v1_event_source.TestV1EventSource method), 576

TestV1ExecAction (class in kubernetes.test.test_v1_exec_action), 576

testV1ExecAction() (kubernetes.test.test_v1_exec_action.TestV1ExecAction method), 576

TestV1FCVolumeSource (class in kubernetes.test.test_v1_fc_volume_source), 577

testV1FCVolumeSource() (kubernetes.test.test_v1_fc_volume_source.TestV1FCVolumeSource method), 577

TestV1FlexVolumeSource (class in kubernetes.test.test_v1_flex_volume_source), 577

testV1FlexVolumeSource() (kubernetes.test.test_v1_flex_volume_source.TestV1FlexVolumeSource method), 577

TestV1FlockerVolumeSource (class in kubernetes.test.test_v1_flocker_volume_source), 577

testV1FlockerVolumeSource() (kubernetes.test.test_v1_flocker_volume_source.TestV1FlockerVolumeSource method), 577

TestV1GCEPersistentDiskVolumeSource (class in kubernetes.test.test_v1_gce_persistent_disk_volume_source), 578

testV1GCEPersistentDiskVolumeSource() (kubernetes.test.test_v1_gce_persistent_disk_volume_source.TestV1GCEPersistentDiskVolumeSource method), 578

TestV1GitRepoVolumeSource (class in kubernetes.test.test_v1_git_repo_volume_source), 578

testV1GitRepoVolumeSource() (kubernetes.test.test_v1_git_repo_volume_source.TestV1GitRepoVolumeSource method), 578

TestV1GlusterfsVolumeSource (class in kubernetes.test.test_v1_glusterfs_volume_source), 578

testV1GlusterfsVolumeSource() (kubernetes.test.test_v1_glusterfs_volume_source.TestV1GlusterfsVolumeSource method), 578

TestV1Handler (class in kubernetes.test.test_v1_handler), 579

testV1Handler() (kubernetes.test.test_v1_handler.TestV1Handler method), 579

TestV1HorizontalPodAutoscaler (class in kubernetes.test.test_v1_horizontal_pod_autoscaler), 579

testV1HorizontalPodAutoscaler() (kubernetes.test.test_v1_horizontal_pod_autoscaler.TestV1HorizontalPodAutoscaler method), 579

TestV1HorizontalPodAutoscalerList (class in kubernetes.test.test_v1_horizontal_pod_autoscaler_list), 579

testV1HorizontalPodAutoscalerList() (kubernetes.test.test_v1_horizontal_pod_autoscaler_list.TestV1HorizontalPodAutoscalerList method), 580

TestV1HorizontalPodAutoscalerSpec (class in kubernetes.test.test_v1_horizontal_pod_autoscaler_spec), 580

testV1HorizontalPodAutoscalerSpec() (kubernetes.test.test_v1_horizontal_pod_autoscaler_spec.TestV1HorizontalPodAutoscalerSpec method), 580

TestV1HorizontalPodAutoscalerStatus (class in kubernetes.test.test_v1_horizontal_pod_autoscaler_status), 580

testV1HorizontalPodAutoscalerStatus() (kubernetes.test.test_v1_horizontal_pod_autoscaler_status.TestV1HorizontalPodAutoscalerStatus method), 580

TestV1HostPathVolumeSource (class in kubernetes.test.test_v1_host_path_volume_source), 580

testV1HostPathVolumeSource() (kubernetes.test.test_v1_host_path_volume_source.TestV1HostPathVolumeSource method), 581

TestV1HTTPGetAction (class in kubernetes.test.test_v1_http_get_action), 581

testV1HTTPGetAction() (kubernetes.test.test_v1_http_get_action.TestV1HTTPGetAction method), 581

TestV1HTTPHeader (class in kubernetes.test.test_v1_http_header), 581

testV1HTTPHeader() (kubernetes.test.test_v1_http_header.TestV1HTTPHeader method), 581

TestV1ISCSIVolumeSource (class in kubernetes.test.test_v1_iscsi_volume_source), 582

testV1ISCSIVolumeSource() (kubernetes.test.test_v1_iscsi_volume_source.TestV1ISCSIVolumeSource method), 582

[netes.test.test_v1_iscsi_volume_source.TestV1ISCSIVolumeSource](#), 586
[method](#)), 582
[TestV1Job](#) (class in [kubernetes.test.test_v1_job](#)), 582
[testV1Job\(\)](#) ([kubernetes.test.test_v1_job.TestV1Job](#)
[method](#)), 582
[TestV1JobCondition](#) (class in [kuber-](#)
[netes.test.test_v1_job_condition](#)), 582
[testV1JobCondition\(\)](#) ([kuber-](#)
[netes.test.test_v1_job_condition.TestV1JobCondition](#)
[method](#)), 582
[TestV1JobList](#) (class in [kubernetes.test.test_v1_job_list](#)),
583
[testV1JobList\(\)](#) ([kuber-](#)
[netes.test.test_v1_job_list.TestV1JobList](#)
[method](#)), 583
[TestV1JobSpec](#) (class in [kuber-](#)
[netes.test.test_v1_job_spec](#)), 583
[testV1JobSpec\(\)](#) ([kuber-](#)
[netes.test.test_v1_job_spec.TestV1JobSpec](#)
[method](#)), 583
[TestV1JobStatus](#) (class in [kuber-](#)
[netes.test.test_v1_job_status](#)), 583
[testV1JobStatus\(\)](#) ([kuber-](#)
[netes.test.test_v1_job_status.TestV1JobStatus](#)
[method](#)), 583
[TestV1KeyToPath](#) (class in [kuber-](#)
[netes.test.test_v1_key_to_path](#)), 584
[testV1KeyToPath\(\)](#) ([kuber-](#)
[netes.test.test_v1_key_to_path.TestV1KeyToPath](#)
[method](#)), 584
[TestV1Lifecycle](#) (class in [kuber-](#)
[netes.test.test_v1_lifecycle](#)), 584
[testV1Lifecycle\(\)](#) ([kuber-](#)
[netes.test.test_v1_lifecycle.TestV1Lifecycle](#)
[method](#)), 584
[TestV1LimitRange](#) (class in [kuber-](#)
[netes.test.test_v1_limit_range](#)), 584
[testV1LimitRange\(\)](#) ([kuber-](#)
[netes.test.test_v1_limit_range.TestV1LimitRange](#)
[method](#)), 585
[TestV1LimitRangeItem](#) (class in [kuber-](#)
[netes.test.test_v1_limit_range_item](#)), 585
[testV1LimitRangeItem\(\)](#) ([kuber-](#)
[netes.test.test_v1_limit_range_item.TestV1LimitRangeItem](#)
[method](#)), 585
[TestV1LimitRangeList](#) (class in [kuber-](#)
[netes.test.test_v1_limit_range_list](#)), 585
[testV1LimitRangeList\(\)](#) ([kuber-](#)
[netes.test.test_v1_limit_range_list.TestV1LimitRangeList](#)
[method](#)), 585
[TestV1LimitRangeSpec](#) (class in [kuber-](#)
[netes.test.test_v1_limit_range_spec](#)), 585
[testV1LimitRangeSpec\(\)](#) ([kuber-](#)
[netes.test.test_v1_limit_range_spec.TestV1LimitRangeSpec](#)
[netes.test.test_v1_node_condition](#)), 589
[TestV1ISCSIVolumeSource](#) (class in [kuber-](#)
[netes.test.test_v1_iscsi_volume_source](#)), 586
[TestV1LoadBalancerIngress](#) (class in [kuber-](#)
[netes.test.test_v1_load_balancer_ingress](#)),
586
[testV1LoadBalancerIngress\(\)](#) ([kuber-](#)
[netes.test.test_v1_load_balancer_ingress.TestV1LoadBalancerIngress](#)
[method](#)), 586
[TestV1LoadBalancerStatus](#) (class in [kuber-](#)
[netes.test.test_v1_load_balancer_status](#)),
586
[testV1LoadBalancerStatus\(\)](#) ([kuber-](#)
[netes.test.test_v1_load_balancer_status.TestV1LoadBalancerStatus](#)
[method](#)), 586
[TestV1LocalObjectReference](#) (class in [kuber-](#)
[netes.test.test_v1_local_object_reference](#)),
587
[testV1LocalObjectReference\(\)](#) ([kuber-](#)
[netes.test.test_v1_local_object_reference.TestV1LocalObjectReference](#)
[method](#)), 587
[TestV1Namespace](#) (class in [kuber-](#)
[netes.test.test_v1_namespace](#)), 587
[testV1Namespace\(\)](#) ([kuber-](#)
[netes.test.test_v1_namespace.TestV1Namespace](#)
[method](#)), 587
[TestV1NamespaceList](#) (class in [kuber-](#)
[netes.test.test_v1_namespace_list](#)), 587
[testV1NamespaceList\(\)](#) ([kuber-](#)
[netes.test.test_v1_namespace_list.TestV1NamespaceList](#)
[method](#)), 587
[TestV1NamespaceSpec](#) (class in [kuber-](#)
[netes.test.test_v1_namespace_spec](#)), 588
[testV1NamespaceSpec\(\)](#) ([kuber-](#)
[netes.test.test_v1_namespace_spec.TestV1NamespaceSpec](#)
[method](#)), 588
[TestV1NamespaceStatus](#) (class in [kuber-](#)
[netes.test.test_v1_namespace_status](#)), 588
[testV1NamespaceStatus\(\)](#) ([kuber-](#)
[netes.test.test_v1_namespace_status.TestV1NamespaceStatus](#)
[method](#)), 588
[TestV1NFSVolumeSource](#) (class in [kuber-](#)
[netes.test.test_v1_nfs_volume_source](#)), 588
[testV1NFSVolumeSource\(\)](#) ([kuber-](#)
[netes.test.test_v1_nfs_volume_source.TestV1NFSVolumeSource](#)
[method](#)), 588
[TestV1Node](#) (class in [kubernetes.test.test_v1_node](#)), 589
[testV1Node\(\)](#) ([kubernetes.test.test_v1_node.TestV1Node](#)
[method](#)), 589
[TestV1NodeAddress](#) (class in [kuber-](#)
[netes.test.test_v1_node_address](#)), 589
[testV1NodeAddress\(\)](#) ([kuber-](#)
[netes.test.test_v1_node_address.TestV1NodeAddress](#)
[method](#)), 589
[TestV1NodeCondition](#) (class in [kuber-](#)
[netes.test.test_v1_node_condition](#)), 589

testV1NodeCondition() (kubernetes.test.test_v1_node_condition.TestV1NodeCondition method), 593

TestV1NodeCondition (class in kubernetes.test.test_v1_node_condition), 593

TestV1NodeDaemonEndpoints (class in kubernetes.test.test_v1_node_daemon_endpoints), 590

testV1NodeDaemonEndpoints() (kubernetes.test.test_v1_node_daemon_endpoints.TestV1NodeDaemonEndpoints method), 590

TestV1NodeList (class in kubernetes.test.test_v1_node_list), 590

testV1NodeList() (kubernetes.test.test_v1_node_list.TestV1NodeList method), 590

TestV1NodeSpec (class in kubernetes.test.test_v1_node_spec), 590

testV1NodeSpec() (kubernetes.test.test_v1_node_spec.TestV1NodeSpec method), 591

TestV1NodeStatus (class in kubernetes.test.test_v1_node_status), 591

testV1NodeStatus() (kubernetes.test.test_v1_node_status.TestV1NodeStatus method), 591

TestV1NodeSystemInfo (class in kubernetes.test.test_v1_node_system_info), 591

testV1NodeSystemInfo() (kubernetes.test.test_v1_node_system_info.TestV1NodeSystemInfo method), 591

TestV1ObjectFieldSelector (class in kubernetes.test.test_v1_object_field_selector), 592

testV1ObjectFieldSelector() (kubernetes.test.test_v1_object_field_selector.TestV1ObjectFieldSelector method), 592

TestV1ObjectMeta (class in kubernetes.test.test_v1_object_meta), 592

testV1ObjectMeta() (kubernetes.test.test_v1_object_meta.TestV1ObjectMeta method), 592

TestV1ObjectReference (class in kubernetes.test.test_v1_object_reference), 592

testV1ObjectReference() (kubernetes.test.test_v1_object_reference.TestV1ObjectReference method), 592

TestV1OwnerReference (class in kubernetes.test.test_v1_owner_reference), 593

testV1OwnerReference() (kubernetes.test.test_v1_owner_reference.TestV1OwnerReference method), 593

TestV1PersistentVolume (class in kubernetes.test.test_v1_persistent_volume), 593

testV1PersistentVolume() (kubernetes.test.test_v1_persistent_volume.TestV1PersistentVolume method), 593

TestV1PersistentVolumeClaim (class in kubernetes.test.test_v1_persistent_volume_claim), 593

testV1PersistentVolumeClaim() (kubernetes.test.test_v1_persistent_volume_claim.TestV1PersistentVolumeClaim method), 593

TestV1PersistentVolumeClaimList (class in kubernetes.test.test_v1_persistent_volume_claim_list), 594

testV1PersistentVolumeClaimList() (kubernetes.test.test_v1_persistent_volume_claim_list.TestV1PersistentVolumeClaimList method), 594

TestV1PersistentVolumeClaimSpec (class in kubernetes.test.test_v1_persistent_volume_claim_spec), 594

testV1PersistentVolumeClaimSpec() (kubernetes.test.test_v1_persistent_volume_claim_spec.TestV1PersistentVolumeClaimSpec method), 594

TestV1PersistentVolumeClaimStatus (class in kubernetes.test.test_v1_persistent_volume_claim_status), 594

testV1PersistentVolumeClaimStatus() (kubernetes.test.test_v1_persistent_volume_claim_status.TestV1PersistentVolumeClaimStatus method), 595

TestV1PersistentVolumeClaimVolumeSource (class in kubernetes.test.test_v1_persistent_volume_claim_volume_source), 595

testV1PersistentVolumeClaimVolumeSource() (kubernetes.test.test_v1_persistent_volume_claim_volume_source.TestV1PersistentVolumeClaimVolumeSource method), 595

TestV1PersistentVolumeList (class in kubernetes.test.test_v1_persistent_volume_list), 595

testV1PersistentVolumeList() (kubernetes.test.test_v1_persistent_volume_list.TestV1PersistentVolumeList method), 595

TestV1PersistentVolumeSpec (class in kubernetes.test.test_v1_persistent_volume_spec), 595

testV1PersistentVolumeSpec() (kubernetes.test.test_v1_persistent_volume_spec.TestV1PersistentVolumeSpec method), 596

TestV1PersistentVolumeStatus (class in kubernetes.test.test_v1_persistent_volume_status), 596

testV1PersistentVolumeStatus() (kubernetes.test.test_v1_persistent_volume_status.TestV1PersistentVolumeStatus method), 596

TestV1PhotonPersistentDiskVolumeSource (class in kubernetes.test.test_v1_photon_persistent_disk_volume_source), 596

| | |
|--|--|
| testV1PhotonPersistentDiskVolumeSource() (kubernetes.test.test_v1_photon_persistent_disk_volume_source.TestV1PhotonPersistentDiskVolumeSource method), 596 | testV1Probe() (kubernetes.test.test_v1_probe.TestV1Probe method), 600 |
| TestV1Pod (class in kubernetes.test.test_v1_pod), 597 | TestV1QuobyteVolumeSource (class in kubernetes.test.test_v1_quobyte_volume_source), 600 |
| testV1Pod() (kubernetes.test.test_v1_pod.TestV1Pod method), 597 | testV1QuobyteVolumeSource() (kubernetes.test.test_v1_quobyte_volume_source.TestV1QuobyteVolumeSource method), 601 |
| TestV1PodCondition (class in kubernetes.test.test_v1_pod_condition), 597 | TestV1RBDVolumeSource (class in kubernetes.test.test_v1_rbd_volume_source), 601 |
| testV1PodCondition() (kubernetes.test.test_v1_pod_condition.TestV1PodCondition method), 597 | testV1RBDVolumeSource() (kubernetes.test.test_v1_rbd_volume_source.TestV1RBDVolumeSource method), 601 |
| TestV1PodList (class in kubernetes.test.test_v1_pod_list), 597 | TestV1ReplicationController (class in kubernetes.test.test_v1_replication_controller), 601 |
| testV1PodList() (kubernetes.test.test_v1_pod_list.TestV1PodList method), 597 | testV1ReplicationController() (kubernetes.test.test_v1_replication_controller.TestV1ReplicationController method), 601 |
| TestV1PodSecurityContext (class in kubernetes.test.test_v1_pod_security_context), 598 | TestV1ReplicationControllerCondition (class in kubernetes.test.test_v1_replication_controller_condition), 602 |
| testV1PodSecurityContext() (kubernetes.test.test_v1_pod_security_context.TestV1PodSecurityContext method), 598 | testV1ReplicationControllerCondition() (kubernetes.test.test_v1_replication_controller_condition.TestV1ReplicationControllerCondition method), 602 |
| TestV1PodSpec (class in kubernetes.test.test_v1_pod_spec), 598 | TestV1ReplicationControllerList (class in kubernetes.test.test_v1_replication_controller_list), 602 |
| testV1PodSpec() (kubernetes.test.test_v1_pod_spec.TestV1PodSpec method), 598 | testV1ReplicationControllerList() (kubernetes.test.test_v1_replication_controller_list.TestV1ReplicationControllerList method), 602 |
| TestV1PodStatus (class in kubernetes.test.test_v1_pod_status), 598 | TestV1ReplicationControllerSpec (class in kubernetes.test.test_v1_replication_controller_spec), 602 |
| testV1PodStatus() (kubernetes.test.test_v1_pod_status.TestV1PodStatus method), 598 | testV1ReplicationControllerSpec() (kubernetes.test.test_v1_replication_controller_spec.TestV1ReplicationControllerSpec method), 602 |
| TestV1PodTemplate (class in kubernetes.test.test_v1_pod_template), 599 | TestV1ReplicationControllerStatus (class in kubernetes.test.test_v1_replication_controller_status), 603 |
| testV1PodTemplate() (kubernetes.test.test_v1_pod_template.TestV1PodTemplate method), 599 | testV1ReplicationControllerStatus() (kubernetes.test.test_v1_replication_controller_status.TestV1ReplicationControllerStatus method), 603 |
| TestV1PodTemplateList (class in kubernetes.test.test_v1_pod_template_list), 599 | TestV1ResourceFieldSelector (class in kubernetes.test.test_v1_resource_field_selector), 603 |
| testV1PodTemplateList() (kubernetes.test.test_v1_pod_template_list.TestV1PodTemplateList method), 599 | testV1ResourceFieldSelector() (kubernetes.test.test_v1_resource_field_selector.TestV1ResourceFieldSelector method), 603 |
| TestV1PodTemplateSpec (class in kubernetes.test.test_v1_pod_template_spec), 599 | TestV1ResourceQuota (class in kubernetes.test.test_v1_resource_quota), 603 |
| testV1PodTemplateSpec() (kubernetes.test.test_v1_pod_template_spec.TestV1PodTemplateSpec method), 600 | testV1ResourceQuota() (kubernetes.test.test_v1_resource_quota.TestV1ResourceQuota method), 603 |
| TestV1Preconditions (class in kubernetes.test.test_v1_preconditions), 600 | |
| testV1Preconditions() (kubernetes.test.test_v1_preconditions.TestV1Preconditions method), 600 | |
| TestV1Probe (class in kubernetes.test.test_v1_probe), 600 | |

TestV1ResourceQuotaList (class in kubernetes.test.test_v1_resource_quota_list), 604
 testV1ResourceQuotaList() (kubernetes.test.test_v1_resource_quota_list.TestV1ResourceQuotaList method), 604
 TestV1ResourceQuotaSpec (class in kubernetes.test.test_v1_resource_quota_spec), 604
 testV1ResourceQuotaSpec() (kubernetes.test.test_v1_resource_quota_spec.TestV1ResourceQuotaSpec method), 604
 TestV1ResourceQuotaStatus (class in kubernetes.test.test_v1_resource_quota_status), 604
 testV1ResourceQuotaStatus() (kubernetes.test.test_v1_resource_quota_status.TestV1ResourceQuotaStatus method), 605
 TestV1ResourceRequirements (class in kubernetes.test.test_v1_resource_requirements), 605
 testV1ResourceRequirements() (kubernetes.test.test_v1_resource_requirements.TestV1ResourceRequirements method), 605
 TestV1Scale (class in kubernetes.test.test_v1_scale), 605
 testV1Scale() (kubernetes.test.test_v1_scale.TestV1Scale method), 605
 TestV1ScaleSpec (class in kubernetes.test.test_v1_scale_spec), 605
 testV1ScaleSpec() (kubernetes.test.test_v1_scale_spec.TestV1ScaleSpec method), 606
 TestV1ScaleStatus (class in kubernetes.test.test_v1_scale_status), 606
 testV1ScaleStatus() (kubernetes.test.test_v1_scale_status.TestV1ScaleStatus method), 606
 TestV1Secret (class in kubernetes.test.test_v1_secret), 607
 testV1Secret() (kubernetes.test.test_v1_secret.TestV1Secret method), 607
 TestV1SecretKeySelector (class in kubernetes.test.test_v1_secret_key_selector), 607
 testV1SecretKeySelector() (kubernetes.test.test_v1_secret_key_selector.TestV1SecretKeySelector method), 607
 TestV1SecretList (class in kubernetes.test.test_v1_secret_list), 607
 testV1SecretList() (kubernetes.test.test_v1_secret_list.TestV1SecretList method), 607
 TestV1SecretVolumeSource (class in kubernetes.test.test_v1_secret_volume_source), 608
 testV1SecretVolumeSource() (kubernetes.test.test_v1_secret_volume_source.TestV1SecretVolumeSource method), 608
 TestV1SecurityContext (class in kubernetes.test.test_v1_security_context), 608
 testV1SecurityContext() (kubernetes.test.test_v1_security_context.TestV1SecurityContext method), 608
 TestV1SELinuxOptions (class in kubernetes.test.test_v1_selinux_options), 606
 testV1SELinuxOptions() (kubernetes.test.test_v1_selinux_options.TestV1SELinuxOptions method), 606
 TestV1Service (class in kubernetes.test.test_v1_service), 608
 testV1Service() (kubernetes.test.test_v1_service.TestV1Service method), 608
 TestV1ServiceAccount (class in kubernetes.test.test_v1_service_account), 609
 testV1ServiceAccount() (kubernetes.test.test_v1_service_account.TestV1ServiceAccount method), 609
 TestV1ServiceAccountList (class in kubernetes.test.test_v1_service_account_list), 609
 testV1ServiceAccountList() (kubernetes.test.test_v1_service_account_list.TestV1ServiceAccountList method), 609
 TestV1ServiceList (class in kubernetes.test.test_v1_service_list), 609
 testV1ServiceList() (kubernetes.test.test_v1_service_list.TestV1ServiceList method), 610
 TestV1ServicePort (class in kubernetes.test.test_v1_service_port), 610
 testV1ServicePort() (kubernetes.test.test_v1_service_port.TestV1ServicePort method), 610
 TestV1ServiceSpec (class in kubernetes.test.test_v1_service_spec), 610
 testV1ServiceSpec() (kubernetes.test.test_v1_service_spec.TestV1ServiceSpec method), 610
 TestV1ServiceStatus (class in kubernetes.test.test_v1_service_status), 610
 testV1ServiceStatus() (kubernetes.test.test_v1_service_status.TestV1ServiceStatus method), 611
 TestV1TCPSocketAction (class in kubernetes.test.test_v1_tcp_socket_action), 611
 testV1TCPSocketAction() (kubernetes.test.test_v1_tcp_socket_action.TestV1TCPSocketAction method), 611
 TestV1Volume (class in kubernetes.test.test_v1_volume),

| | | |
|---|---|---|
| 611 | | attribute), 477 |
| testV1Volume() | (kuber- netes.test.test_v1_volume.TestV1Volume method), 611 | to_debug_report() (kuber- netes.client.configuration.Configuration method), 525 |
| TestV1VolumeMount | (class in kuber- netes.test.test_v1_volume_mount), 612 | to_dict() (kubernetes.client.models.runtime_raw_extension.RuntimeRawEx method), 298 |
| testV1VolumeMount() | (kuber- netes.test.test_v1_volume_mount.TestV1VolumeMount method), 612 | to_dict() (kubernetes.client.models.v1_attached_volume.V1AttachedVolum method), 299 |
| TestV1VsphereVirtualDiskVolumeSource | (class in kuber- netes.test.test_v1_vsphere_virtual_disk_volume_source), 612 | to_dict() (kubernetes.client.models.v1_aws_elastic_block_store_volume_so method), 300 |
| testV1VsphereVirtualDiskVolumeSource() | (kuber- netes.test.test_v1_vsphere_virtual_disk_volume_source.TestV1VsphereVirtualDiskVolumeSource method), 612 | to_dict() (kubernetes.client.models.v1_azure_disk_volume_source.V1Azure method), 301 |
| TestV2alpha1CronJob | (class in kuber- netes.test.test_v2alpha1_cron_job), 634 | to_dict() (kubernetes.client.models.v1_azure_file_volume_source.V1Azure method), 303 |
| testV2alpha1CronJob() | (kuber- netes.test.test_v2alpha1_cron_job.TestV2alpha1CronJob method), 634 | to_dict() (kubernetes.client.models.v1_binding.V1Binding method), 303 |
| TestV2alpha1CronJobList | (class in kuber- netes.test.test_v2alpha1_cron_job_list), 635 | to_dict() (kubernetes.client.models.v1_capabilities.V1Capabilities method), 303 |
| testV2alpha1CronJobList() | (kuber- netes.test.test_v2alpha1_cron_job_list.TestV2alpha1CronJobList method), 635 | to_dict() (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFS method), 305 |
| TestV2alpha1CronJobSpec | (class in kuber- netes.test.test_v2alpha1_cron_job_spec), 635 | to_dict() (kubernetes.client.models.v1_cinder_volume_source.V1CinderVol method), 305 |
| testV2alpha1CronJobSpec() | (kuber- netes.test.test_v2alpha1_cron_job_spec.TestV2alpha1CronJobSpec method), 635 | to_dict() (kubernetes.client.models.v1_cinder_volume_source.V1CinderVol method), 305 |
| TestV2alpha1CronJobStatus | (class in kuber- netes.test.test_v2alpha1_cron_job_status), 635 | to_dict() (kubernetes.client.models.v1_component_condition.V1Component method), 306 |
| testV2alpha1CronJobStatus() | (kuber- netes.test.test_v2alpha1_cron_job_status.TestV2alpha1CronJobStatus method), 636 | to_dict() (kubernetes.client.models.v1_component_status.V1ComponentSta method), 307 |
| TestV2alpha1JobTemplateSpec | (class in kuber- netes.test.test_v2alpha1_job_template_spec), 636 | to_dict() (kubernetes.client.models.v1_component_status_list.V1Component method), 308 |
| testV2alpha1JobTemplateSpec() | (kuber- netes.test.test_v2alpha1_job_template_spec.TestV2alpha1JobTemplateSpec method), 636 | to_dict() (kubernetes.client.models.v1_config_map.V1ConfigMap method), 309 |
| TestVersionApi | (class in kuber- netes.test.test_version_api), 636 | to_dict() (kubernetes.client.models.v1_config_map_key_selector.V1Config method), 310 |
| TestVersionInfo | (class in kuber- netes.test.test_version_info), 637 | to_dict() (kubernetes.client.models.v1_config_map_list.V1ConfigMapList method), 311 |
| testVersionInfo() | (kuber- netes.test.test_version_info.TestVersionInfo method), 637 | to_dict() (kubernetes.client.models.v1_config_map_list.V1ConfigMapList method), 311 |
| timeout_seconds | (kuber- netes.client.models.v1_probe.V1Probe attribute), 415 | to_dict() (kubernetes.client.models.v1_config_map_volume_source.V1Con method), 312 |
| tls (kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec | | to_dict() (kubernetes.client.models.v1_config_map_key_selector.V1Config method), 310 |
| | | to_dict() (kubernetes.client.models.v1_config_map_list.V1ConfigMapList method), 311 |
| | | to_dict() (kubernetes.client.models.v1_config_map_volume_source.V1Con method), 312 |
| | | to_dict() (kubernetes.client.models.v1_container.V1Container method), 316 |
| | | to_dict() (kubernetes.client.models.v1_container_image.V1ContainerImage method), 317 |
| | | to_dict() (kubernetes.client.models.v1_container_port.V1ContainerPort method), 318 |
| | | to_dict() (kubernetes.client.models.v1_container_state.V1ContainerState method), 318 |
| | | to_dict() (kubernetes.client.models.v1_container_state_running.V1Contain method), 319 |
| | | to_dict() (kubernetes.client.models.v1_container_state_terminated.V1Conta method), 320 |
| | | to_dict() (kubernetes.client.models.v1_container_state_waiting.V1Contain method), 321 |
| | | to_dict() (kubernetes.client.models.v1_container_status.V1ContainerStatus method), 322 |
| | | to_dict() (kubernetes.client.models.v1_cross_version_object_reference.V1C method), 323 |

to_dict() (kubernetes.client.models.v1_daemon_endpoint.V1DaemonEndpoint method), 324

to_dict() (kubernetes.client.models.v1_delete_options.V1DeleteOptions method), 325

to_dict() (kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource method), 326

to_dict() (kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource method), 327

to_dict() (kubernetes.client.models.v1_empty_dir_volume_source.V1EmptyDirVolumeSource method), 328

to_dict() (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress method), 329

to_dict() (kubernetes.client.models.v1_endpoint_port.V1EndpointPort method), 329

to_dict() (kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset method), 330

to_dict() (kubernetes.client.models.v1_endpoints.V1Endpoints method), 331

to_dict() (kubernetes.client.models.v1_endpoints_list.V1EndpointsList method), 332

to_dict() (kubernetes.client.models.v1_env_var.V1EnvVar method), 333

to_dict() (kubernetes.client.models.v1_env_var_source.V1EnvVarSource method), 334

to_dict() (kubernetes.client.models.v1_event.V1Event method), 336

to_dict() (kubernetes.client.models.v1_event_list.V1EventList method), 337

to_dict() (kubernetes.client.models.v1_event_source.V1EventSource method), 337

to_dict() (kubernetes.client.models.v1_exec_action.V1ExecAction method), 338

to_dict() (kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource method), 339

to_dict() (kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource method), 340

to_dict() (kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource method), 341

to_dict() (kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource method), 342

to_dict() (kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource method), 343

to_dict() (kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterfsVolumeSource method), 344

to_dict() (kubernetes.client.models.v1_handler.V1Handler method), 344

to_dict() (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler method), 345

to_dict() (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler method), 346

to_dict() (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler method), 347

to_dict() (kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler method), 349

to_dict() (kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource method), 349

to_dict() (kubernetes.client.models.v1_http_get_action.V1HTTPGetAction method), 350

to_dict() (kubernetes.client.models.v1_http_header.V1HTTPHeader method), 351

to_dict() (kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource method), 353

to_dict() (kubernetes.client.models.v1_job.V1Job method), 354

to_dict() (kubernetes.client.models.v1_job_condition.V1JobCondition method), 355

to_dict() (kubernetes.client.models.v1_job_list.V1JobList method), 356

to_dict() (kubernetes.client.models.v1_job_spec.V1JobSpec method), 358

to_dict() (kubernetes.client.models.v1_job_status.V1JobStatus method), 359

to_dict() (kubernetes.client.models.v1_key_to_path.V1KeyToPath method), 359

to_dict() (kubernetes.client.models.v1_lifecycle.V1Lifecycle method), 360

to_dict() (kubernetes.client.models.v1_limit_range.V1LimitRange method), 361

to_dict() (kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem method), 362

to_dict() (kubernetes.client.models.v1_limit_range_list.V1LimitRangeList method), 363

to_dict() (kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec method), 364

to_dict() (kubernetes.client.models.v1_load_balancer_ingress.V1LoadBalancerIngress method), 365

to_dict() (kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus method), 365

to_dict() (kubernetes.client.models.v1_local_object_reference.V1LocalObjectReference method), 366

to_dict() (kubernetes.client.models.v1_namespace.V1Namespace method), 367

to_dict() (kubernetes.client.models.v1_namespace_list.V1NamespaceList method), 368

to_dict() (kubernetes.client.models.v1_namespace_spec.V1NamespaceSpec method), 368

to_dict() (kubernetes.client.models.v1_namespace_status.V1NamespaceStatus method), 369

to_dict() (kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource method), 370

to_dict() (kubernetes.client.models.v1_node.V1Node method), 371

to_dict() (kubernetes.client.models.v1_node_address.V1NodeAddress method), 371

to_dict() (kubernetes.client.models.v1_node_condition.V1NodeCondition method), 372

to_dict() (kubernetes.client.models.v1_node_daemon_endpoints.V1NodeDaemonEndpoints method), 373

to_dict() (kubernetes.client.models.v1_node_list.V1NodeList method), 374

to_dict() (kubernetes.client.models.v1_node_spec.V1NodeSpec method), 375

to_dict() (kubernetes.client.models.v1_node_status.V1NodeStatus method), 376

to_dict() (kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo method), 378

to_dict() (kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector method), 379

to_dict() (kubernetes.client.models.v1_object_meta.V1ObjectMeta method), 383

to_dict() (kubernetes.client.models.v1_object_reference.V1ObjectReference method), 384

to_dict() (kubernetes.client.models.v1_owner_reference.V1OwnerReference method), 385

to_dict() (kubernetes.client.models.v1_persistent_volume.V1PersistentVolume method), 387

to_dict() (kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim method), 388

to_dict() (kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList method), 389

to_dict() (kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec method), 390

to_dict() (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus method), 391

to_dict() (kubernetes.client.models.v1_persistent_volume_claim_status_list.V1PersistentVolumeClaimStatusList method), 392

to_dict() (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList method), 393

to_dict() (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec method), 398

to_dict() (kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus method), 399

to_dict() (kubernetes.client.models.v1_photon_persistent_disk.V1PhotonPersistentDisk method), 399

to_dict() (kubernetes.client.models.v1_pod.V1Pod method), 400

to_dict() (kubernetes.client.models.v1_pod_condition.V1PodCondition method), 402

to_dict() (kubernetes.client.models.v1_pod_list.V1PodList method), 403

to_dict() (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext method), 404

to_dict() (kubernetes.client.models.v1_pod_spec.V1PodSpec method), 408

to_dict() (kubernetes.client.models.v1_pod_status.V1PodStatus method), 410

to_dict() (kubernetes.client.models.v1_pod_template.V1PodTemplate method), 411

to_dict() (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList method), 412

to_dict() (kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec method), 413

to_dict() (kubernetes.client.models.v1_preconditions.V1Preconditions method), 413

to_dict() (kubernetes.client.models.v1_probe.V1Probe method), 415

to_dict() (kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource method), 416

to_dict() (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource method), 418

to_dict() (kubernetes.client.models.v1_replication_controller.V1ReplicationController method), 419

to_dict() (kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition method), 420

to_dict() (kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList method), 421

to_dict() (kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec method), 422

to_dict() (kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus method), 424

to_dict() (kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector method), 424

to_dict() (kubernetes.client.models.v1_resource_quota.V1ResourceQuota method), 425

to_dict() (kubernetes.client.models.v1_resource_quota_list.V1ResourceQuotaList method), 427

to_dict() (kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec method), 427

to_dict() (kubernetes.client.models.v1_resource_quota_status.V1ResourceQuotaStatus method), 428

to_dict() (kubernetes.client.models.v1_resource_requirements.V1ResourceRequirements method), 429

to_dict() (kubernetes.client.models.v1_scale.V1Scale method), 430

to_dict() (kubernetes.client.models.v1_scale_spec.V1ScaleSpec method), 430

to_dict() (kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource method), 431

to_dict() (kubernetes.client.models.v1_selinux_options.V1SELinuxOptions method), 432

to_dict() (kubernetes.client.models.v1_secret.V1Secret method), 433

to_dict() (kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector method), 434

to_dict() (kubernetes.client.models.v1_secret_list.V1SecretList method), 435

to_dict() (kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource method), 436

to_dict() (kubernetes.client.models.v1_security_context.V1SecurityContext method), 437

to_dict() (kubernetes.client.models.v1_service.V1Service method), 438

to_dict() (kubernetes.client.models.v1_service_account.V1ServiceAccount method), 440

to_dict() (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList method), 441

to_dict() (kubernetes.client.models.v1_service_list.V1ServiceList method), 442

to_dict() (kubernetes.client.models.v1_service_port.V1ServicePort method), 443

to_dict() (kubernetes.client.models.v1_service_spec.V1ServiceSpec method), 446

to_dict() (kubernetes.client.models.v1_service_status.V1ServiceStatus method), 447

to_dict() (kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction method), 447

to_dict() (kubernetes.client.models.v1_volume.V1Volume method), 451

to_dict() (kubernetes.client.models.v1_volume_mount.V1VolumeMount method), 452

to_dict() (kubernetes.client.models.v1_vsphere_virtual_disk_to_volume.V1VsphereVirtualDiskToVolume method), 453

to_dict() (kubernetes.client.models.v1alpha1_cluster_role.V1alpha1ClusterRole method), 455

to_dict() (kubernetes.client.models.v1alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding method), 456

to_dict() (kubernetes.client.models.v1alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList method), 457

to_dict() (kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList method), 458

to_dict() (kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule method), 459

to_dict() (kubernetes.client.models.v1alpha1_role.V1alpha1Role method), 460

to_dict() (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding method), 461

to_dict() (kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList method), 462

to_dict() (kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList method), 463

to_dict() (kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef method), 464

to_dict() (kubernetes.client.models.v1alpha1_subject.V1alpha1Subject method), 465

to_dict() (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet method), 466

to_dict() (kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList method), 467

to_dict() (kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec method), 469

to_dict() (kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus method), 470

to_dict() (kubernetes.client.models.v1beta1_eviction.V1beta1Eviction method), 472

to_dict() (kubernetes.client.models.v1beta1_http_ingress_path.V1beta1HTTPIngressPath method), 473

to_dict() (kubernetes.client.models.v1beta1_http_ingress_rule.V1beta1HTTPIngressRule method), 473

to_dict() (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress method), 474

to_dict() (kubernetes.client.models.v1beta1_ingress_backend.V1beta1IngressBackend method), 475

to_dict() (kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList method), 476

to_dict() (kubernetes.client.models.v1beta1_ingress_rule.V1beta1IngressRule method), 477

to_dict() (kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec method), 478

to_dict() (kubernetes.client.models.v1beta1_ingress_status.V1beta1IngressStatus method), 478

to_dict() (kubernetes.client.models.v1beta1_ingress_tls.V1beta1IngressTLS method), 479

to_dict() (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview method), 480

to_dict() (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy method), 481

to_dict() (kubernetes.client.models.v1beta1_network_policy_ingress_rule.V1beta1NetworkPolicyIngressRule method), 482

to_dict() (kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList method), 483

to_dict() (kubernetes.client.models.v1beta1_network_policy_peer.V1beta1NetworkPolicyPeer method), 484

to_dict() (kubernetes.client.models.v1beta1_network_policy_port.V1beta1NetworkPolicyPort method), 484

to_dict() (kubernetes.client.models.v1beta1_network_policy_spec.V1beta1NetworkPolicySpec method), 486

to_dict() (kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes method), 486

to_dict() (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget method), 487

to_dict() (kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList method), 488

to_dict() (kubernetes.client.models.v1beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec method), 489

to_dict() (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus method), 491

to_dict() (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet method), 492

to_dict() (kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition method), 493

to_dict() (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList method), 494

to_dict() (kubernetes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec method), 495

to_dict() (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus method), 497

to_dict() (kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes method), 498

to_dict() (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview method), 499

to_dict() (kubernetes.client.models.v1beta1_self_subject_access_review_spec.V1beta1SelfSubjectAccessReviewSpec method), 500

to_dict() (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet method), 501

to_dict() (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList method), 502

to_dict() (kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec method), 504

to_dict() (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus method), 506

to_dict() (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass method), 508

to_dict() (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList method), 509

to_dict() (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview method), 510

to_dict() (kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec method), 511

to_dict() (kubernetes.client.models.v1beta1_subject_access_review_status.V1beta1SubjectAccessReviewStatus method), 512

to_dict() (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview method), 513

to_dict() (kubernetes.client.models.v1beta1_token_review_spec.V1beta1TokenReviewSpec method), 514

to_dict() (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus method), 514

to_dict() (kubernetes.client.models.v1beta1_user_info.V1beta1UserInfo method), 515

to_dict() (kubernetes.client.models.v2alpha1_cron_job.V2alpha1CronJob method), 517

to_dict() (kubernetes.client.models.v2alpha1_cron_job_list.V2alpha1CronJobList method), 518

to_dict() (kubernetes.client.models.v2alpha1_cron_job_spec.V2alpha1CronJobSpec method), 519

to_dict() (kubernetes.client.models.v2alpha1_cron_job_status.V2alpha1CronJobStatus method), 520

to_dict() (kubernetes.client.models.v2alpha1_job_template_spec.V2alpha1JobTemplateSpec method), 521

to_dict() (kubernetes.client.models.version_info.VersionInfo method), 522

to_str() (kubernetes.client.models.runtime_raw_extension.RuntimeRawExtension method), 298

to_str() (kubernetes.client.models.v1_attached_volume.V1AttachedVolume method), 299

to_str() (kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AWSElasticBlockStoreVolumeSource method), 300

to_str() (kubernetes.client.models.v1_azure_disk_volume_source.V1AzureDiskVolumeSource method), 301

to_str() (kubernetes.client.models.v1_azure_file_volume_source.V1AzureFileVolumeSource method), 302

to_str() (kubernetes.client.models.v1_binding.V1Binding method), 303

to_str() (kubernetes.client.models.v1_capabilities.V1Capabilities method), 303

to_str() (kubernetes.client.models.v1_ceph_fs_volume_source.V1CephFSVolumeSource method), 305

to_str() (kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource method), 305

to_str() (kubernetes.client.models.v1_component_condition.V1ComponentCondition method), 306

to_str() (kubernetes.client.models.v1_component_status.V1ComponentStatus method), 308

to_str() (kubernetes.client.models.v1_component_status_list.V1ComponentStatusList method), 308

to_str() (kubernetes.client.models.v1_config_map.V1ConfigMap method), 309

to_str() (kubernetes.client.models.v1_config_map_key_selector.V1ConfigMapKeySelector method), 310

to_str() (kubernetes.client.models.v1_config_map_list.V1ConfigMapList method), 311

to_str() (kubernetes.client.models.v1_config_map_volume_source.V1ConfigMapVolumeSource method), 312

to_str() (kubernetes.client.models.v1_container.V1Container method), 316

to_str() (kubernetes.client.models.v1_container_image.V1ContainerImage method), 317

to_str() (kubernetes.client.models.v1_container_port.V1ContainerPort method), 318

to_str() (kubernetes.client.models.v1_container_state.V1ContainerState method), 318

to_str() (kubernetes.client.models.v1_container_state_running.V1ContainerStateRunning method), 319

to_str() (kubernetes.client.models.v1_container_state_terminated.V1ContainerStateTerminated method), 320

to_str() (kubernetes.client.models.v1_container_state_waiting.V1ContainerStateWaiting method), 321

to_str() (kubernetes.client.models.v1_container_status.V1ContainerStatus method), 323

to_str() (kubernetes.client.models.v1_cross_version_object_reference.V1CrossVersionObjectReference method), 323

to_str() (kubernetes.client.models.v1_daemon_endpoint.V1DaemonEndpoint method), 324

to_str() (kubernetes.client.models.v1_delete_options.V1DeleteOptions method), 325

to_str() (kubernetes.client.models.v1_downward_api_volume_file.V1DownwardAPIVolumeFile method), 326

to_str() (kubernetes.client.models.v1_downward_api_volume_source.V1DownwardAPIVolumeSource method), 327

to_str() (kubernetes.client.models.v1_empty_dir_volume_source.V1EmptyDirVolumeSource method), 328

to_str() (kubernetes.client.models.v1_endpoint_address.V1EndpointAddress method), 329

to_str() (kubernetes.client.models.v1_endpoint_port.V1EndpointPort method), 329

to_str() (kubernetes.client.models.v1_endpoint_subset.V1EndpointSubset method), 330

to_str() (kubernetes.client.models.v1_endpoints.V1Endpoints method), 331

to_str() (kubernetes.client.models.v1_endpoints_list.V1EndpointsList method), 332

to_str() (kubernetes.client.models.v1_env_var.V1EnvVar method), 333

[to_str\(\) \(kubernetes.client.models.v1_env_var_source.V1EnvVarSource method\), 334](#)
[to_str\(\) \(kubernetes.client.models.v1_event.V1Event method\), 336](#)
[to_str\(\) \(kubernetes.client.models.v1_event_list.V1EventList method\), 337](#)
[to_str\(\) \(kubernetes.client.models.v1_event_source.V1EventSource method\), 337](#)
[to_str\(\) \(kubernetes.client.models.v1_exec_action.V1ExecAction method\), 338](#)
[to_str\(\) \(kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource method\), 339](#)
[to_str\(\) \(kubernetes.client.models.v1_flex_volume_source.V1FlexVolumeSource method\), 340](#)
[to_str\(\) \(kubernetes.client.models.v1_flocker_volume_source.V1FlockerVolumeSource method\), 341](#)
[to_str\(\) \(kubernetes.client.models.v1_gce_persistent_disk_volume_source.V1GCEPersistentDiskVolumeSource method\), 342](#)
[to_str\(\) \(kubernetes.client.models.v1_git_repo_volume_source.V1GitRepoVolumeSource method\), 343](#)
[to_str\(\) \(kubernetes.client.models.v1_glusterfs_volume_source.V1GlusterfsVolumeSource method\), 344](#)
[to_str\(\) \(kubernetes.client.models.v1_handler.V1Handler method\), 344](#)
[to_str\(\) \(kubernetes.client.models.v1_horizontal_pod_autoscaler.V1HorizontalPodAutoscaler method\), 345](#)
[to_str\(\) \(kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus method\), 346](#)
[to_str\(\) \(kubernetes.client.models.v1_horizontal_pod_autoscaler_spec.V1HorizontalPodAutoscalerSpec method\), 347](#)
[to_str\(\) \(kubernetes.client.models.v1_horizontal_pod_autoscaler_status.V1HorizontalPodAutoscalerStatus method\), 349](#)
[to_str\(\) \(kubernetes.client.models.v1_host_path_volume_source.V1HostPathVolumeSource method\), 349](#)
[to_str\(\) \(kubernetes.client.models.v1_http_get_action.V1HTTPGetAction method\), 350](#)
[to_str\(\) \(kubernetes.client.models.v1_http_header.V1HTTPHeader method\), 351](#)
[to_str\(\) \(kubernetes.client.models.v1_iscsi_volume_source.V1ISCSIVolumeSource method\), 353](#)
[to_str\(\) \(kubernetes.client.models.v1_job.V1Job method\), 354](#)
[to_str\(\) \(kubernetes.client.models.v1_job_condition.V1JobCondition method\), 355](#)
[to_str\(\) \(kubernetes.client.models.v1_job_list.V1JobList method\), 356](#)
[to_str\(\) \(kubernetes.client.models.v1_job_spec.V1JobSpec method\), 358](#)
[to_str\(\) \(kubernetes.client.models.v1_job_status.V1JobStatus method\), 359](#)
[to_str\(\) \(kubernetes.client.models.v1_key_to_path.V1KeyToPath method\), 359](#)
[to_str\(\) \(kubernetes.client.models.v1_lifecycle.V1Lifecycle method\), 360](#)
[to_str\(\) \(kubernetes.client.models.v1_limit_range.V1LimitRange method\), 361](#)
[to_str\(\) \(kubernetes.client.models.v1_limit_range_item.V1LimitRangeItem method\), 362](#)
[to_str\(\) \(kubernetes.client.models.v1_limit_range_list.V1LimitRangeList method\), 363](#)
[to_str\(\) \(kubernetes.client.models.v1_limit_range_spec.V1LimitRangeSpec method\), 364](#)
[to_str\(\) \(kubernetes.client.models.v1_load_balancer_ingress.V1LoadBalancerIngress method\), 365](#)
[to_str\(\) \(kubernetes.client.models.v1_load_balancer_status.V1LoadBalancerStatus method\), 365](#)
[to_str\(\) \(kubernetes.client.models.v1_local_object_reference.V1LocalObjectReference method\), 366](#)
[to_str\(\) \(kubernetes.client.models.v1_namespace.V1Namespace method\), 367](#)
[to_str\(\) \(kubernetes.client.models.v1_namespace_list.V1NamespaceList method\), 368](#)
[to_str\(\) \(kubernetes.client.models.v1_namespace_spec.V1NamespaceSpec method\), 368](#)
[to_str\(\) \(kubernetes.client.models.v1_namespace_status.V1NamespaceStatus method\), 369](#)
[to_str\(\) \(kubernetes.client.models.v1_nfs_volume_source.V1NFSVolumeSource method\), 370](#)
[to_str\(\) \(kubernetes.client.models.v1_node.V1Node method\), 371](#)
[to_str\(\) \(kubernetes.client.models.v1_node_address.V1NodeAddress method\), 371](#)
[to_str\(\) \(kubernetes.client.models.v1_node_condition.V1NodeCondition method\), 372](#)
[to_str\(\) \(kubernetes.client.models.v1_node_autoscaler_status.V1NodeAutoscalerStatus method\), 373](#)
[to_str\(\) \(kubernetes.client.models.v1_node_list.V1NodeList method\), 374](#)
[to_str\(\) \(kubernetes.client.models.v1_node_spec.V1NodeSpec method\), 375](#)
[to_str\(\) \(kubernetes.client.models.v1_node_status.V1NodeStatus method\), 377](#)
[to_str\(\) \(kubernetes.client.models.v1_node_system_info.V1NodeSystemInfo method\), 379](#)
[to_str\(\) \(kubernetes.client.models.v1_object_field_selector.V1ObjectFieldSelector method\), 379](#)
[to_str\(\) \(kubernetes.client.models.v1_object_meta.V1ObjectMeta method\), 383](#)
[to_str\(\) \(kubernetes.client.models.v1_object_reference.V1ObjectReference method\), 384](#)
[to_str\(\) \(kubernetes.client.models.v1_owner_reference.V1OwnerReference method\), 385](#)
[to_str\(\) \(kubernetes.client.models.v1_persistent_volume.V1PersistentVolume method\), 387](#)
[to_str\(\) \(kubernetes.client.models.v1_persistent_volume_claim.V1PersistentVolumeClaim method\), 388](#)
[to_str\(\) \(kubernetes.client.models.v1_persistent_volume_claim_list.V1PersistentVolumeClaimList method\), 389](#)

to_str() (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus method), 390

to_str() (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus method), 391

to_str() (kubernetes.client.models.v1_persistent_volume_claim_status.V1PersistentVolumeClaimStatus method), 392

to_str() (kubernetes.client.models.v1_persistent_volume_list.V1PersistentVolumeList method), 393

to_str() (kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec method), 398

to_str() (kubernetes.client.models.v1_persistent_volume_status.V1PersistentVolumeStatus method), 399

to_str() (kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource method), 399

to_str() (kubernetes.client.models.v1_pod.V1Pod method), 400

to_str() (kubernetes.client.models.v1_pod_condition.V1PodCondition method), 402

to_str() (kubernetes.client.models.v1_pod_list.V1PodList method), 403

to_str() (kubernetes.client.models.v1_pod_security_context.V1PodSecurityContext method), 404

to_str() (kubernetes.client.models.v1_pod_spec.V1PodSpec method), 408

to_str() (kubernetes.client.models.v1_pod_status.V1PodStatus method), 410

to_str() (kubernetes.client.models.v1_pod_template.V1PodTemplate method), 411

to_str() (kubernetes.client.models.v1_pod_template_list.V1PodTemplateList method), 412

to_str() (kubernetes.client.models.v1_pod_template_spec.V1PodTemplateSpec method), 413

to_str() (kubernetes.client.models.v1_preconditions.V1Preconditions method), 413

to_str() (kubernetes.client.models.v1_probe.V1Probe method), 415

to_str() (kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource method), 416

to_str() (kubernetes.client.models.v1_rbd_volume_source.V1RBDVolumeSource method), 418

to_str() (kubernetes.client.models.v1_replication_controller.V1ReplicationController method), 419

to_str() (kubernetes.client.models.v1_replication_controller_condition.V1ReplicationControllerCondition method), 420

to_str() (kubernetes.client.models.v1_replication_controller_list.V1ReplicationControllerList method), 421

to_str() (kubernetes.client.models.v1_replication_controller_spec.V1ReplicationControllerSpec method), 422

to_str() (kubernetes.client.models.v1_replication_controller_status.V1ReplicationControllerStatus method), 424

to_str() (kubernetes.client.models.v1_resource_field_selector.V1ResourceFieldSelector method), 424

to_str() (kubernetes.client.models.v1_resource_quota.V1ResourceQuota method), 426

to_str() (kubernetes.client.models.v1_resource_quota_spec.V1ResourceQuotaSpec method), 427

to_str() (kubernetes.client.models.v1_resource_requirements.V1ResourceRequirements method), 428

to_str() (kubernetes.client.models.v1_resource_requirements.V1ResourceRequirements method), 429

to_str() (kubernetes.client.models.v1_scale.V1Scale method), 430

to_str() (kubernetes.client.models.v1_scale_spec.V1ScaleSpec method), 430

to_str() (kubernetes.client.models.v1_photon_persistent_disk_volume_source.V1PhotonPersistentDiskVolumeSource method), 431

to_str() (kubernetes.client.models.v1_selinux_options.V1SELinuxOptions method), 432

to_str() (kubernetes.client.models.v1_secret.V1Secret method), 433

to_str() (kubernetes.client.models.v1_secret_key_selector.V1SecretKeySelector method), 434

to_str() (kubernetes.client.models.v1_secret_list.V1SecretList method), 435

to_str() (kubernetes.client.models.v1_secret_volume_source.V1SecretVolumeSource method), 436

to_str() (kubernetes.client.models.v1_security_context.V1SecurityContext method), 437

to_str() (kubernetes.client.models.v1_service.V1Service method), 438

to_str() (kubernetes.client.models.v1_service_account.V1ServiceAccount method), 440

to_str() (kubernetes.client.models.v1_service_account_list.V1ServiceAccountList method), 441

to_str() (kubernetes.client.models.v1_service_list.V1ServiceList method), 442

to_str() (kubernetes.client.models.v1_service_port.V1ServicePort method), 443

to_str() (kubernetes.client.models.v1_service_spec.V1ServiceSpec method), 446

to_str() (kubernetes.client.models.v1_service_status.V1ServiceStatus method), 447

to_str() (kubernetes.client.models.v1_tcp_socket_action.V1TCPSocketAction method), 447

to_str() (kubernetes.client.models.v1_replication_controller_condition.V1VolumeCondition method), 451

to_str() (kubernetes.client.models.v1_volume_mount.V1VolumeMount method), 452

to_str() (kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource method), 453

to_str() (kubernetes.client.models.v1_alpha1_cluster_role.V1alpha1ClusterRole method), 455

to_str() (kubernetes.client.models.v1_alpha1_cluster_role_binding.V1alpha1ClusterRoleBinding method), 456

to_str() (kubernetes.client.models.v1_alpha1_cluster_role_binding_list.V1alpha1ClusterRoleBindingList method), 457

to_str() (kubernetes.client.models.v1alpha1_cluster_role_list.V1alpha1ClusterRoleList method), 458

to_str() (kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule method), 459

to_str() (kubernetes.client.models.v1alpha1_role.V1alpha1Role method), 460

to_str() (kubernetes.client.models.v1alpha1_role_binding.V1alpha1RoleBinding method), 461

to_str() (kubernetes.client.models.v1alpha1_role_binding_list.V1alpha1RoleBindingList method), 462

to_str() (kubernetes.client.models.v1alpha1_role_list.V1alpha1RoleList method), 463

to_str() (kubernetes.client.models.v1alpha1_role_ref.V1alpha1RoleRef method), 464

to_str() (kubernetes.client.models.v1alpha1_subject.V1alpha1Subject method), 465

to_str() (kubernetes.client.models.v1beta1_daemon_set.V1beta1DaemonSet method), 466

to_str() (kubernetes.client.models.v1beta1_daemon_set_list.V1beta1DaemonSetList method), 467

to_str() (kubernetes.client.models.v1beta1_daemon_set_spec.V1beta1DaemonSetSpec method), 469

to_str() (kubernetes.client.models.v1beta1_daemon_set_status.V1beta1DaemonSetStatus method), 470

to_str() (kubernetes.client.models.v1beta1_eviction.V1beta1Eviction method), 472

to_str() (kubernetes.client.models.v1beta1_http_ingress_path.V1beta1HTTPIngressPath method), 473

to_str() (kubernetes.client.models.v1beta1_http_ingress_rule.V1beta1HTTPIngressRule method), 473

to_str() (kubernetes.client.models.v1beta1_ingress.V1beta1Ingress method), 474

to_str() (kubernetes.client.models.v1beta1_ingress_backend.V1beta1IngressBackend method), 475

to_str() (kubernetes.client.models.v1beta1_ingress_list.V1beta1IngressList method), 476

to_str() (kubernetes.client.models.v1beta1_ingress_rule.V1beta1IngressRule method), 477

to_str() (kubernetes.client.models.v1beta1_ingress_spec.V1beta1IngressSpec method), 478

to_str() (kubernetes.client.models.v1beta1_ingress_status.V1beta1IngressStatus method), 478

to_str() (kubernetes.client.models.v1beta1_ingress_tls.V1beta1IngressTLS method), 479

to_str() (kubernetes.client.models.v1beta1_local_subject_access_review.V1beta1LocalSubjectAccessReview method), 480

to_str() (kubernetes.client.models.v1beta1_network_policy.V1beta1NetworkPolicy method), 481

to_str() (kubernetes.client.models.v1beta1_network_policy_ingress.V1beta1NetworkPolicyIngressRule method), 482

to_str() (kubernetes.client.models.v1beta1_network_policy_list.V1beta1NetworkPolicyList method), 483

to_str() (kubernetes.client.models.v1beta1_network_policy_peer.V1beta1NetworkPolicyPeer method), 484

to_str() (kubernetes.client.models.v1beta1_network_policy_port.V1beta1NetworkPolicyPort method), 484

to_str() (kubernetes.client.models.v1beta1_network_policy_rule.V1beta1NetworkPolicyRule method), 486

to_str() (kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes method), 486

to_str() (kubernetes.client.models.v1beta1_pod_disruption_budget.V1beta1PodDisruptionBudget method), 487

to_str() (kubernetes.client.models.v1beta1_pod_disruption_budget_list.V1beta1PodDisruptionBudgetList method), 488

to_str() (kubernetes.client.models.v1beta1_pod_disruption_budget_spec.V1beta1PodDisruptionBudgetSpec method), 489

to_str() (kubernetes.client.models.v1beta1_pod_disruption_budget_status.V1beta1PodDisruptionBudgetStatus method), 491

to_str() (kubernetes.client.models.v1beta1_replica_set.V1beta1ReplicaSet method), 492

to_str() (kubernetes.client.models.v1beta1_replica_set_condition.V1beta1ReplicaSetCondition method), 493

to_str() (kubernetes.client.models.v1beta1_replica_set_list.V1beta1ReplicaSetList method), 494

to_str() (kubernetes.client.models.v1beta1_replica_set_spec.V1beta1ReplicaSetSpec method), 495

to_str() (kubernetes.client.models.v1beta1_replica_set_status.V1beta1ReplicaSetStatus method), 497

to_str() (kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes method), 498

to_str() (kubernetes.client.models.v1beta1_self_subject_access_review.V1beta1SelfSubjectAccessReview method), 499

to_str() (kubernetes.client.models.v1beta1_self_subject_access_review_spec.V1beta1SelfSubjectAccessReviewSpec method), 500

to_str() (kubernetes.client.models.v1beta1_stateful_set.V1beta1StatefulSet method), 501

to_str() (kubernetes.client.models.v1beta1_stateful_set_list.V1beta1StatefulSetList method), 502

to_str() (kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec method), 504

to_str() (kubernetes.client.models.v1beta1_stateful_set_status.V1beta1StatefulSetStatus method), 506

to_str() (kubernetes.client.models.v1beta1_storage_class.V1beta1StorageClass method), 508

to_str() (kubernetes.client.models.v1beta1_storage_class_list.V1beta1StorageClassList method), 509

to_str() (kubernetes.client.models.v1beta1_subject_access_review.V1beta1SubjectAccessReview method), 510

to_str() (kubernetes.client.models.v1beta1_subject_access_review_spec.V1beta1SubjectAccessReviewSpec method), 511

to_str() (kubernetes.client.models.v1beta1_subject_access_review_status.V1beta1SubjectAccessReviewStatus method), 512

to_str() (kubernetes.client.models.v1beta1_token_review.V1beta1TokenReview method), 513

to_str() (kubernetes.client.models.v1beta1_token_review_spec.V1beta1TokenReviewSpec method), 514

to_str() (kubernetes.client.models.v1beta1_token_review_status.V1beta1TokenReviewStatus method), 514

V

Index

| | |
|--|--|
| V1alpha1ClusterRoleBinding (class in kubernetes.client.models.v1alpha1_cluster_role_binding), 455 | V1beta1HTTPIngressRuleValue (class in kubernetes.client.models.v1beta1_http_ingress_rule_value), 473 |
| V1alpha1ClusterRoleBindingList (class in kubernetes.client.models.v1alpha1_cluster_role_binding_list), 456 | V1beta1Ingress (class in kubernetes.client.models.v1beta1_ingress), 473 |
| V1alpha1ClusterRoleList (class in kubernetes.client.models.v1alpha1_cluster_role_list), 457 | V1beta1IngressBackend (class in kubernetes.client.models.v1beta1_ingress_backend), 475 |
| V1alpha1PolicyRule (class in kubernetes.client.models.v1alpha1_policy_rule), 458 | V1beta1IngressList (class in kubernetes.client.models.v1beta1_ingress_list), 475 |
| V1alpha1Role (class in kubernetes.client.models.v1alpha1_role), 459 | V1beta1IngressRule (class in kubernetes.client.models.v1beta1_ingress_rule), 476 |
| V1alpha1RoleBinding (class in kubernetes.client.models.v1alpha1_role_binding), 460 | V1beta1IngressSpec (class in kubernetes.client.models.v1beta1_ingress_spec), 477 |
| V1alpha1RoleBindingList (class in kubernetes.client.models.v1alpha1_role_binding_list), 461 | V1beta1IngressStatus (class in kubernetes.client.models.v1beta1_ingress_status), 478 |
| V1alpha1RoleList (class in kubernetes.client.models.v1alpha1_role_list), 462 | V1beta1IngressTLS (class in kubernetes.client.models.v1beta1_ingress_tls), 478 |
| V1alpha1RoleRef (class in kubernetes.client.models.v1alpha1_role_ref), 463 | V1beta1LocalSubjectAccessReview (class in kubernetes.client.models.v1beta1_local_subject_access_review), 479 |
| V1alpha1Subject (class in kubernetes.client.models.v1alpha1_subject), 464 | V1beta1NetworkPolicy (class in kubernetes.client.models.v1beta1_network_policy), 480 |
| V1AttachedVolume (class in kubernetes.client.models.v1_attached_volume), 298 | V1beta1NetworkPolicyIngressRule (class in kubernetes.client.models.v1beta1_network_policy_ingress_rule), 481 |
| V1AWSElasticBlockStoreVolumeSource (class in kubernetes.client.models.v1_aws_elastic_block_store_volume_source), 299 | V1beta1NetworkPolicyList (class in kubernetes.client.models.v1beta1_network_policy_list), 482 |
| V1AzureDiskVolumeSource (class in kubernetes.client.models.v1_azure_disk_volume_source), 300 | V1beta1NetworkPolicyPeer (class in kubernetes.client.models.v1beta1_network_policy_peer), 483 |
| V1AzureFileVolumeSource (class in kubernetes.client.models.v1_azure_file_volume_source), 301 | V1beta1NetworkPolicyPort (class in kubernetes.client.models.v1beta1_network_policy_port), 484 |
| V1beta1DaemonSet (class in kubernetes.client.models.v1beta1_daemon_set), 465 | V1beta1NetworkPolicySpec (class in kubernetes.client.models.v1beta1_network_policy_spec), 485 |
| V1beta1DaemonSetList (class in kubernetes.client.models.v1beta1_daemon_set_list), 466 | V1beta1NonResourceAttributes (class in kubernetes.client.models.v1beta1_non_resource_attributes), 486 |
| V1beta1DaemonSetSpec (class in kubernetes.client.models.v1beta1_daemon_set_spec), 467 | V1beta1PodDisruptionBudget (class in kubernetes.client.models.v1beta1_pod_disruption_budget), 486 |
| V1beta1DaemonSetStatus (class in kubernetes.client.models.v1beta1_daemon_set_status), 469 | V1beta1PodDisruptionBudgetList (class in kubernetes.client.models.v1beta1_pod_disruption_budget_list), 488 |
| V1beta1Eviction (class in kubernetes.client.models.v1beta1_eviction), 471 | |
| V1beta1HTTPIngressPath (class in kubernetes.client.models.v1beta1_http_ingress_path), 472 | |

| | |
|--|---|
| V1beta1PodDisruptionBudgetSpec (class in kubernetes.client.models.v1beta1_pod_disruption_budget_spec), 489 | V1beta1SubjectAccessReviewStatus (class in kubernetes.client.models.v1beta1_subject_access_review_status), 511 |
| V1beta1PodDisruptionBudgetStatus (class in kubernetes.client.models.v1beta1_pod_disruption_budget_status), 490 | V1beta1TokenReview (class in kubernetes.client.models.v1beta1_token_review), 512 |
| V1beta1ReplicaSet (class in kubernetes.client.models.v1beta1_replica_set), 491 | V1beta1TokenReviewSpec (class in kubernetes.client.models.v1beta1_token_review_spec), 513 |
| V1beta1ReplicaSetCondition (class in kubernetes.client.models.v1beta1_replica_set_condition), 492 | V1beta1TokenReviewStatus (class in kubernetes.client.models.v1beta1_token_review_status), 514 |
| V1beta1ReplicaSetList (class in kubernetes.client.models.v1beta1_replica_set_list), 493 | V1beta1UserInfo (class in kubernetes.client.models.v1beta1_user_info), 515 |
| V1beta1ReplicaSetSpec (class in kubernetes.client.models.v1beta1_replica_set_spec), 494 | V1Binding (class in kubernetes.client.models.v1_binding), 302 |
| V1beta1ReplicaSetStatus (class in kubernetes.client.models.v1beta1_replica_set_status), 495 | V1Capabilities (class in kubernetes.client.models.v1_capabilities), 303 |
| V1beta1ResourceAttributes (class in kubernetes.client.models.v1beta1_resource_attributes), 497 | V1CephFSVolumeSource (class in kubernetes.client.models.v1_ceph_fs_volume_source), 304 |
| V1beta1SelfSubjectAccessReview (class in kubernetes.client.models.v1beta1_self_subject_access_review), 498 | V1CinderVolumeSource (class in kubernetes.client.models.v1_cinder_volume_source), 305 |
| V1beta1SelfSubjectAccessReviewSpec (class in kubernetes.client.models.v1beta1_self_subject_access_review_spec), 500 | V1ComponentCondition (class in kubernetes.client.models.v1_component_condition), 306 |
| V1beta1StatefulSet (class in kubernetes.client.models.v1beta1_stateful_set), 500 | V1ComponentStatus (class in kubernetes.client.models.v1_component_status), 307 |
| V1beta1StatefulSetList (class in kubernetes.client.models.v1beta1_stateful_set_list), 501 | V1ComponentStatusList (class in kubernetes.client.models.v1_component_status_list), 308 |
| V1beta1StatefulSetSpec (class in kubernetes.client.models.v1beta1_stateful_set_spec), 502 | V1ConfigMap (class in kubernetes.client.models.v1_config_map), 309 |
| V1beta1StatefulSetStatus (class in kubernetes.client.models.v1beta1_stateful_set_status), 504 | V1ConfigMapKeySelector (class in kubernetes.client.models.v1_config_map_key_selector), 310 |
| V1beta1StorageClass (class in kubernetes.client.models.v1beta1_storage_class), 506 | V1ConfigMapList (class in kubernetes.client.models.v1_config_map_list), 310 |
| V1beta1StorageClassList (class in kubernetes.client.models.v1beta1_storage_class_list), 508 | V1ConfigMapVolumeSource (class in kubernetes.client.models.v1_config_map_volume_source), 311 |
| V1beta1SubjectAccessReview (class in kubernetes.client.models.v1beta1_subject_access_review), 509 | V1Container (class in kubernetes.client.models.v1_container), 313 |
| V1beta1SubjectAccessReviewSpec (class in kubernetes.client.models.v1beta1_subject_access_review_spec), 510 | V1ContainerImage (class in kubernetes.client.models.v1_container_image), 316 |
| | V1ContainerPort (class in kubernetes.client.models.v1_container_port), 317 |
| | V1ContainerState (class in kubernetes.client.models.v1_container_state), 318 |

| | |
|---|--|
| V1ContainerStateRunning (class in kubernetes.client.models.v1_container_state_running), 319 | V1FCVolumeSource (class in kubernetes.client.models.v1_fc_volume_source), 338 |
| V1ContainerStateTerminated (class in kubernetes.client.models.v1_container_state_terminated), 319 | V1FlexVolumeSource (class in kubernetes.client.models.v1_flex_volume_source), 339 |
| V1ContainerStateWaiting (class in kubernetes.client.models.v1_container_state_waiting), 321 | V1FlockerVolumeSource (class in kubernetes.client.models.v1_flocker_volume_source), 340 |
| V1ContainerStatus (class in kubernetes.client.models.v1_container_status), 321 | V1GCEPersistentDiskVolumeSource (class in kubernetes.client.models.v1_gce_persistent_disk_volume_source), 341 |
| V1CrossVersionObjectReference (class in kubernetes.client.models.v1_cross_version_object_reference), 323 | V1GitRepoVolumeSource (class in kubernetes.client.models.v1_git_repo_volume_source), 342 |
| V1DaemonEndpoint (class in kubernetes.client.models.v1_daemon_endpoint), 324 | V1GlusterfsVolumeSource (class in kubernetes.client.models.v1_glusterfs_volume_source), 343 |
| V1DeleteOptions (class in kubernetes.client.models.v1_delete_options), 324 | V1Handler (class in kubernetes.client.models.v1_handler), 344 |
| V1DownwardAPIVolumeFile (class in kubernetes.client.models.v1_downward_api_volume_file), 325 | V1HorizontalPodAutoscaler (class in kubernetes.client.models.v1_horizontal_pod_autoscaler), 344 |
| V1DownwardAPIVolumeSource (class in kubernetes.client.models.v1_downward_api_volume_source), 326 | V1HorizontalPodAutoscalerList (class in kubernetes.client.models.v1_horizontal_pod_autoscaler_list), 346 |
| V1EmptyDirVolumeSource (class in kubernetes.client.models.v1_empty_dir_volume_source), 327 | V1HorizontalPodAutoscalerSpec (class in kubernetes.client.models.v1_horizontal_pod_autoscaler_spec), 347 |
| V1EndpointAddress (class in kubernetes.client.models.v1_endpoint_address), 328 | V1HorizontalPodAutoscalerStatus (class in kubernetes.client.models.v1_horizontal_pod_autoscaler_status), 348 |
| V1EndpointPort (class in kubernetes.client.models.v1_endpoint_port), 329 | V1HostPathVolumeSource (class in kubernetes.client.models.v1_host_path_volume_source), 349 |
| V1Endpoints (class in kubernetes.client.models.v1_endpoints), 330 | V1HTTPGetAction (class in kubernetes.client.models.v1_http_get_action), 349 |
| V1EndpointsList (class in kubernetes.client.models.v1_endpoints_list), 331 | V1HTTPHeader (class in kubernetes.client.models.v1_http_header), 351 |
| V1EndpointSubset (class in kubernetes.client.models.v1_endpoint_subset), 330 | V1ISCSIVolumeSource (class in kubernetes.client.models.v1_iscsi_volume_source), 351 |
| V1EnvVar (class in kubernetes.client.models.v1_env_var), 332 | V1Job (class in kubernetes.client.models.v1_job), 353 |
| V1EnvVarSource (class in kubernetes.client.models.v1_env_var_source), 333 | V1JobCondition (class in kubernetes.client.models.v1_job_condition), 354 |
| V1Event (class in kubernetes.client.models.v1_event), 334 | V1JobList (class in kubernetes.client.models.v1_job_list), 355 |
| V1EventList (class in kubernetes.client.models.v1_event_list), 336 | V1JobSpec (class in kubernetes.client.models.v1_job_spec), 356 |
| V1EventSource (class in kubernetes.client.models.v1_event_source), 337 | V1JobStatus (class in kubernetes.client.models.v1_job_status), 358 |
| V1ExecAction (class in kubernetes.client.models.v1_exec_action), 338 | V1KeyToPath (class in kubernetes.client.models.v1_key_to_path), 359 |

| | | | | | |
|------------------------|--|-----|-------------------------------------|---|-----|
| V1Lifecycle | (class in kubernetes.client.models.v1_lifecycle), | 360 | V1ObjectFieldSelector | (class in kubernetes.client.models.v1_object_field_selector), | 379 |
| V1LimitRange | (class in kubernetes.client.models.v1_limit_range), | 360 | V1ObjectMeta | (class in kubernetes.client.models.v1_object_meta), | 379 |
| V1LimitRangeItem | (class in kubernetes.client.models.v1_limit_range_item), | 361 | V1ObjectReference | (class in kubernetes.client.models.v1_object_reference), | 383 |
| V1LimitRangeList | (class in kubernetes.client.models.v1_limit_range_list), | 363 | V1OwnerReference | (class in kubernetes.client.models.v1_owner_reference), | 384 |
| V1LimitRangeSpec | (class in kubernetes.client.models.v1_limit_range_spec), | 364 | V1PersistentVolume | (class in kubernetes.client.models.v1_persistent_volume), | 386 |
| V1LoadBalancerIngress | (class in kubernetes.client.models.v1_load_balancer_ingress), | 364 | V1PersistentVolumeClaim | (class in kubernetes.client.models.v1_persistent_volume_claim), | 387 |
| V1LoadBalancerStatus | (class in kubernetes.client.models.v1_load_balancer_status), | 365 | V1PersistentVolumeClaimList | (class in kubernetes.client.models.v1_persistent_volume_claim_list), | 388 |
| V1LocalObjectReference | (class in kubernetes.client.models.v1_local_object_reference), | 365 | V1PersistentVolumeClaimSpec | (class in kubernetes.client.models.v1_persistent_volume_claim_spec), | 389 |
| V1Namespace | (class in kubernetes.client.models.v1_namespace), | 366 | V1PersistentVolumeClaimStatus | (class in kubernetes.client.models.v1_persistent_volume_claim_status), | 390 |
| V1NamespaceList | (class in kubernetes.client.models.v1_namespace_list), | 367 | V1PersistentVolumeClaimVolumeSource | (class in kubernetes.client.models.v1_persistent_volume_claim_volume_source), | 392 |
| V1NamespaceSpec | (class in kubernetes.client.models.v1_namespace_spec), | 368 | V1PersistentVolumeList | (class in kubernetes.client.models.v1_persistent_volume_list), | 392 |
| V1NamespaceStatus | (class in kubernetes.client.models.v1_namespace_status), | 368 | V1PersistentVolumeSpec | (class in kubernetes.client.models.v1_persistent_volume_spec), | 393 |
| V1NFSVolumeSource | (class in kubernetes.client.models.v1_nfs_volume_source), | 369 | V1PersistentVolumeStatus | (class in kubernetes.client.models.v1_persistent_volume_status), | 398 |
| V1Node | (class in kubernetes.client.models.v1_node), | 370 | V1PhotonPersistentDiskVolumeSource | (class in kubernetes.client.models.v1_photon_persistent_disk_volume_source), | 399 |
| V1NodeAddress | (class in kubernetes.client.models.v1_node_address), | 371 | V1Pod | (class in kubernetes.client.models.v1_pod), | 400 |
| V1NodeCondition | (class in kubernetes.client.models.v1_node_condition), | 372 | V1PodCondition | (class in kubernetes.client.models.v1_pod_condition), | 401 |
| V1NodeDaemonEndpoints | (class in kubernetes.client.models.v1_node_daemon_endpoints), | 373 | V1PodList | (class in kubernetes.client.models.v1_pod_list), | 402 |
| V1NodeList | (class in kubernetes.client.models.v1_node_list), | 373 | V1PodSecurityContext | (class in kubernetes.client.models.v1_pod_security_context), | 403 |
| V1NodeSpec | (class in kubernetes.client.models.v1_node_spec), | 374 | V1PodSpec | (class in kubernetes.client.models.v1_pod_spec), | 404 |
| V1NodeStatus | (class in kubernetes.client.models.v1_node_status), | 375 | V1PodStatus | (class in kubernetes.client.models.v1_pod_status), | 409 |
| V1NodeSystemInfo | (class in kubernetes.client.models.v1_node_system_info), | 377 | V1PodTemplate | (class in kuber- | |

| | |
|---|---|
| netes.client.models.v1_pod_template), 411 | V1ScaleStatus (class in kuber- |
| V1PodTemplateList (class in kuber- | netes.client.models.v1_scale_status), 430 |
| netes.client.models.v1_pod_template_list), 412 | V1Secret (class in kubernetes.client.models.v1_secret), 432 |
| V1PodTemplateSpec (class in kuber- | V1SecretKeySelector (class in kuber- |
| netes.client.models.v1_pod_template_spec), 413 | netes.client.models.v1_secret_key_selector), 433 |
| V1Preconditions (class in kuber- | V1SecretList (class in kuber- |
| netes.client.models.v1_preconditions), 413 | netes.client.models.v1_secret_list), 434 |
| V1Probe (class in kubernetes.client.models.v1_probe), 414 | V1SecretVolumeSource (class in kuber- |
| V1QuobyteVolumeSource (class in kuber- | netes.client.models.v1_secret_volume_source), 435 |
| netes.client.models.v1_quobyte_volume_source), 415 | V1SecurityContext (class in kuber- |
| V1RBDVolumeSource (class in kuber- | netes.client.models.v1_security_context), 436 |
| netes.client.models.v1_rbd_volume_source), 416 | V1SELinuxOptions (class in kuber- |
| V1ReplicationController (class in kuber- | netes.client.models.v1_se_linux_options), 431 |
| netes.client.models.v1_replication_controller), 418 | V1Service (class in kubernetes.client.models.v1_service), 438 |
| V1ReplicationControllerCondition (class in kuber- | V1ServiceAccount (class in kuber- |
| netes.client.models.v1_replication_controller_condition), 419 | netes.client.models.v1_service_account), 439 |
| V1ReplicationControllerList (class in kuber- | V1ServiceAccountList (class in kuber- |
| netes.client.models.v1_replication_controller_list), 420 | netes.client.models.v1_service_account_list), 440 |
| V1ReplicationControllerSpec (class in kuber- | V1ServiceList (class in kuber- |
| netes.client.models.v1_replication_controller_spec), 421 | netes.client.models.v1_service_list), 441 |
| V1ReplicationControllerStatus (class in kuber- | V1ServicePort (class in kuber- |
| netes.client.models.v1_replication_controller_status), 422 | netes.client.models.v1_service_port), 442 |
| V1ResourceFieldSelector (class in kuber- | V1ServiceSpec (class in kuber- |
| netes.client.models.v1_resource_field_selector), 424 | netes.client.models.v1_service_spec), 443 |
| V1ResourceQuota (class in kuber- | V1ServiceStatus (class in kuber- |
| netes.client.models.v1_resource_quota), 425 | netes.client.models.v1_service_status), 446 |
| V1ResourceQuotaList (class in kuber- | V1TCPSocketAction (class in kuber- |
| netes.client.models.v1_resource_quota_list), 426 | netes.client.models.v1_tcp_socket_action), 447 |
| V1ResourceQuotaSpec (class in kuber- | V1Volume (class in kuber- |
| netes.client.models.v1_resource_quota_spec), 427 | netes.client.models.v1_volume), 447 |
| V1ResourceQuotaStatus (class in kuber- | V1VolumeMount (class in kuber- |
| netes.client.models.v1_resource_quota_status), 427 | netes.client.models.v1_volume_mount), 452 |
| V1ResourceRequirements (class in kuber- | V1VsphereVirtualDiskVolumeSource (class in kuber- |
| netes.client.models.v1_resource_requirements), 428 | netes.client.models.v1_vsphere_virtual_disk_volume_source), 453 |
| V1Scale (class in kubernetes.client.models.v1_scale), 429 | V2alpha1CronJob (class in kuber- |
| V1ScaleSpec (class in kuber- | netes.client.models.v2alpha1_cron_job), 516 |
| netes.client.models.v1_scale_spec), 430 | V2alpha1CronJobList (class in kuber- |
| | netes.client.models.v2alpha1_cron_job_list), 517 |
| | V2alpha1CronJobSpec (class in kuber- |
| | netes.client.models.v2alpha1_cron_job_spec), 518 |

[V2alpha1CronJobStatus](#) (class in [kubernetes.client.models.v2alpha1_cron_job_status](#)), [519](#)
[V2alpha1JobTemplateSpec](#) (class in [kubernetes.client.models.v2alpha1_job_template_spec](#)), [520](#)
[value](#) ([kubernetes.client.models.v1_env_var.V1EnvVar](#) attribute), [333](#)
[value](#) ([kubernetes.client.models.v1_http_header.V1HTTPHeader](#) attribute), [351](#)
[value_from](#) ([kubernetes.client.models.v1_env_var.V1EnvVar](#) attribute), [333](#)
[verb](#) ([kubernetes.client.models.v1beta1_non_resource_attributes.V1beta1NonResourceAttributes](#) attribute), [486](#)
[verb](#) ([kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes](#) attribute), [498](#)
[verbs](#) ([kubernetes.client.models.v1alpha1_policy_rule.V1alpha1PolicyRule](#) attribute), [459](#)
[version](#) ([kubernetes.client.models.v1beta1_resource_attributes.V1beta1ResourceAttributes](#) attribute), [498](#)
[VersionApi](#) (class in [kubernetes.client.apis.version_api](#)), [297](#)
[VersionInfo](#) (class in [kubernetes.client.models.version_info](#)), [521](#)
[volume](#) ([kubernetes.client.models.v1_quobyte_volume_source.V1QuobyteVolumeSource](#) attribute), [416](#)
[volume_claim_templates](#) ([kubernetes.client.models.v1beta1_stateful_set_spec.V1beta1StatefulSetSpec](#) attribute), [504](#)
[volume_id](#) ([kubernetes.client.models.v1_aws_elastic_block_store_volume_source.V1AWSElasticBlockStoreVolumeSource](#) attribute), [300](#)
[volume_id](#) ([kubernetes.client.models.v1_cinder_volume_source.V1CinderVolumeSource](#) attribute), [305](#)
[volume_mounts](#) ([kubernetes.client.models.v1_container.V1Container](#) attribute), [316](#)
[volume_name](#) ([kubernetes.client.models.v1_persistent_volume_claim_spec.V1PersistentVolumeClaimSpec](#) attribute), [390](#)
[volume_path](#) ([kubernetes.client.models.v1_vsphere_virtual_disk_volume_source.V1VsphereVirtualDiskVolumeSource](#) attribute), [453](#)
[volumes](#) ([kubernetes.client.models.v1_pod_spec.V1PodSpec](#) attribute), [408](#)
[volumes_attached](#) ([kubernetes.client.models.v1_node_status.V1NodeStatus](#) attribute), [377](#)
[volumes_in_use](#) ([kubernetes.client.models.v1_node_status.V1NodeStatus](#) attribute), [377](#)
[vsphere_volume](#) ([kubernetes.client.models.v1_persistent_volume_spec.V1PersistentVolumeSpec](#) attribute), [398](#)
[vsphere_volume](#) ([kubernetes.client.models.v1_volume.V1Volume](#) attribute), [451](#)

W

[waiting](#) ([kubernetes.client.models.v1_container_state.V1ContainerState](#) attribute), [318](#)
[Watch](#) (class in [kubernetes.watch.watch](#)), [637](#)
[WatchTests](#) (class in [kubernetes.watch.watch_test](#)), [638](#)
[working_dir](#) ([kubernetes.client.models.v1_container.V1Container](#) attribute), [316](#)
[wwids](#) ([kubernetes.client.models.v1_fc_volume_source.V1FCVolumeSource](#) attribute), [339](#)