

INTRODUCTION TO CONVOLUTIONAL NETWORKS USING TENSORFLOW

Jesús Fernández Bes, jfbes@ing.uc3m.es

8 de febrero de 2016

Install

What is Tensorflow?

Implementing Softmax Regression

Deep Convolutional Networks in Tensorflow

What else?

INSTALL INSTRUCTIONS



- ◆ Create a virtual environment with anaconda (it takes some time)

```
$ conda update conda
```

```
$ conda create -n tensorflow python=2.7 anaconda
```

(tensorflow is the name of the environment, it can be whatever we want)

- ◆ Activate our new environment, prompt changes to (tensorflow)\$

```
$ source activate tensorflow
```

- ◆ To deactivate the environment you have to write (do it at the end of the session)

```
$ source deactivate
```

INSTALL INSTRUCTIONS

- ◆ Install tensorflow but only in the new environment (also takes time)

Ubuntu/Linux 64-bit, CPU only:

```
$ pip install --upgrade  
https://storage.googleapis.com/tensorflow/linux/  
cpu/tensorflow-0.5.0-cp27-none-linux\_x86\_64.whl
```

Ubuntu/Linux 64-bit, GPU enabled:

```
$ pip install --upgrade  
https://storage.googleapis.com/tensorflow/linux/  
gpu/tensorflow-0.5.0-cp27-none-linux\_x86\_64.whl
```

Mac OS X, CPU only:

```
$ pip install --upgrade  
https://storage.googleapis.com/tensorflow/mac/  
tensorflow-0.5.0-py2-none-any.whl
```

Test the new instalation

```
(tensorflow)$ python
```

```
>>> import tensorflow as tf
>>> hello = tf.constant('Hello, TensorFlow!')
>>> sess = tf.Session()
>>> print sess.run(hello)
Hello, TensorFlow!
>>> a = tf.constant(10)
>>> b = tf.constant(32)
>>> print sess.run(a + b)
42
>>>
```

Install

What is Tensorflow?

Implementing Softmax Regression

Deep Convolutional Networks in Tensorflow

What else?

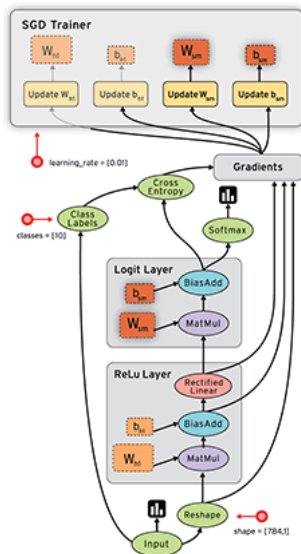
Developed by **Google Brain Team** and **Google's Machine Intelligence** research organization.



- ◆ Interface for expressing machine learning algorithms, and an implementation for executing them.
- ◆ Deep Flexibility: If you can write a computation graph.
- ◆ True Portability.
- ◆ Connect Research and Production.
- ◆ Auto-Differentiation.
- ◆ Language Options: Python, C++.
- ◆ Maximize Performance.

- ◆ In *tensorflow* computation represented using **Graphs**.
Each node is an **operation (op)**.
- ◆ Data is represented a **Tensors**.
Op takes **Tensors** and returns **Tensors**.
- ◆ Variables maintain state across executions of the graph.
- ◆ Two phases in the program:
Construct the computation **graph**.
Executes a **graph** in the context of a **Session**.
- ◆ Feed/fetch data to/from the graph.

EXAMPLE OF COMPUTATION GRAPH



Always the same 3-steps pattern:

1. `inference()` - Builds the graph as far as is required for running the network forward to make predictions.
2. `loss()` - Adds to the inference graph the ops required to generate loss.
3. `training()` - Adds to the loss graph the ops required to compute and apply gradients.

Install

What is Tensorflow?

Implementing Softmax Regression

Deep Convolutional Networks in Tensorflow

What else?

MNIST DATASET: THE HELLO WORLD OF ML



Each image is 28 pixels by 28 pixels.

55,000 data points of training data (`mnist.train`)


10,000 points of test data (`mnist.test`)

5,000 points of validation data (`mnist.validation`).

Use `tensorflow.googlesource.com/tensorflow/+/master/tensorflow/examples/tutorials/mnist/input_data.py` to download the data.

$$y = \text{softmax}(Wx + b)$$

$$\text{softmax}(x)_i = \frac{\exp x_i}{\sum_j \exp x_j}$$



mnist_softmax.py

```
import tensorflow as tf
import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

Fichero descarga de datos:

```
tensorflow.googlesource.com/tensorflow/+/master/  
tensorflow/examples/tutorials/mnist/input_data.py
```

`mnist_softmax.py`: Variable declaration

```
x = tf.placeholder(tf.float32, [None, 784])  
y_ = tf.placeholder(tf.float32, [None, 10])  
  
W = tf.Variable(tf.zeros([784, 10]))  
b = tf.Variable(tf.zeros([10]))
```


◆ Inference

mnist_softmax.py: Inference

```
y = tf.nn.softmax(tf.matmul(x, W) + b)
```

◆ Loss

mnist_softmax.py: Loss computation

```
cross_entropy = -tf.reduce_sum(y*tf.log(y))
```

◆ Training

mnist_softmax.py : Training

```
train_step = tf.train.GradientDescentOptimizer(0.01)
               .minimize(cross_entropy)
```

TensorFlow also provides many other **optimization algorithms**: using one is as simple as tweaking one line:

```
class tf.train.AdagradOptimizer
```

```
class tf.train.MomentumOptimizer
```

```
class tf.train.AdamOptimizer
```

How well we are predicting the correct label?

`mnist_softmax.py`: Evaluation

```
correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_,1))  
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

EXECUTE THE COMPUTATION GRAPH

mnist_softmax.py: Initialize all the variables

```
init = tf.initialize_all_variables()
```

mnist_softmax.py: Start a new session

```
sess = tf.Session()
```

```
sess.run(init)
```

```
# Let's train -- we'll run the training step 1000 times!
```

```
for i in range(1000):
```

```
    batch_xs, batch_ys = mnist.train.next_batch(100)
```

```
    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
```

```
print sess.run(accuracy, feed_dict={x: mnist.test.images,  
                                     y_: mnist.test.labels})
```

Result around 91 %: **VERY BAD** for MNIST

Install

What is Tensorflow?

Implementing Softmax Regression

Deep Convolutional Networks in Tensorflow

What else?

- ◆ State-of-the-art of Image Recognition.
- ◆ Traditional Approach: **Handmade** features.




Figura 1: <http://www.cs.umd.edu/~djacobs>

- ◆ ¿Can we learn the features+classifier?
Complex non-linear map from pixels to labels.
Do it in several *simple* layers: Function composition.

TWO SLIDES COURSE IN CONVOLUTIONAL NETWORKS

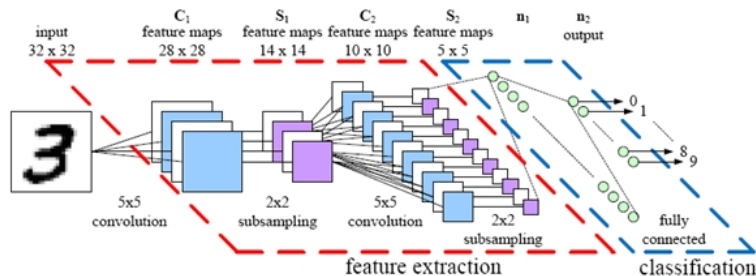


Figura 2: <http://parse.ele.tue.nl/>

Explanation of each layer.

Train using **Backpropagation**.

This works very well. Why?

Rick Baraniuk “opinion”: **A Probabilistic Theory of Deep Learning.**

Reuse the data loading part of `mnist_softmax.py`.

`mnist_cnn.py`

```
import tensorflow as tf
import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)

x = tf.placeholder("float", shape=[None, 784])
y_ = tf.placeholder("float", shape=[None, 10])
```


We will have to initialize a lot of weights.

`mnist_cnn.py`: weight initialization

```
def weight_variable(shape):  
    initial = tf.truncated_normal(shape, stddev=0.1)  
    return tf.Variable(initial)  
  
def bias_variable(shape):  
    initial = tf.constant(0.1, shape=shape)  
    return tf.Variable(initial)
```

Convolution and pooling operations. We will use them in different layers.

`mnist_cnn.py`: Convolution and pooling

```
def conv2d(x, W):  
    return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')  
  
def max_pool_2x2(x):  
    return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],  
                           strides=[1, 2, 2, 1], padding='SAME')
```

First layer: From input data to second layer

mnist_cnn.py: First Convolutional Layer

```
W_conv1 = weight_variable([5, 5, 1, 32])
b_conv1 = bias_variable([32])

x_image = tf.reshape(x, [-1, 28, 28, 1])

h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
h_pool1 = max_pool_2x2(h_conv1)
```

Second layer: From output of first layer to FC layer

mnist_cnn.py: Second Convolutional Layer

```
W_conv2 = weight_variable([5, 5, 32, 64])
b_conv2 = bias_variable([64])

h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
h_pool2 = max_pool_2x2(h_conv2)
```

mnist_cnn.py: Fully Connected layer

```
# Densely connected layer
W_fc1 = weight_variable([7 * 7 * 64, 1024])
b_fc1 = bias_variable([1024])

h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
```

Dropout trick and output layer.

mnist_cnn.py: Dropout

```
# Add dropout
keep_prob = tf.placeholder("float")
h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
```

mnist_cnn.py: Readout layer

```
# Readout layer
W_fc2 = weight_variable([1024, 10])
b_fc2 = bias_variable([10])

y_conv = tf.nn.softmax(tf.matmul(h_fc1_drop, W_fc2) + b_fc2)
```

mnist_cnn.py: Train and Evaluate

```
cross_entropy = -tf.reduce_sum(y_*tf.log(y_conv))  
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)  
correct_prediction = tf.equal(tf.argmax(y_conv,1), tf.argmax(y_,1))  
accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

mnist_cnn.py: Execute

```
init = tf.initialize_all_variables()
sess = tf.InteractiveSession()
sess.run(init)
for i in range(20000):
    batch = mnist.train.next_batch(50)
    if i%100 == 0:
        train_accuracy = accuracy.eval(feed_dict={
            x:batch[0], y_: batch[1], keep_prob: 1.0})
        print "step %d, training accuracy %g"%(i, train_accuracy)
        train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5})

print "test accuracy %g"%accuracy.eval(feed_dict={
    x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0})
```

Go back to work while it finishes: Accuracy $\approx 99,2\%$.

Install

What is Tensorflow?

Implementing Softmax Regression

Deep Convolutional Networks in Tensorflow

What else?

- ◆ Tensorboard.

`https://www.tensorflow.org/versions/0.6.0/how_tos/summaries_and_tensorboard/index.html`

- ◆ Vector Representation of Words (word2vec).
- ◆ Recurrent Neural Networks (Long short-term memory Networks, seq2seq models).
- ◆ General Mathematics (Mandelbrot Set, Partial Differential Equations)
- ◆ **Udacity free online course**